
DICE Embeddings

Release 0.1.3.2

Caglar Demir

Apr 30, 2026

Contents:

1	Dicee Manual	2
2	Installation	3
2.1	Installation from Source	3
3	Download Knowledge Graphs	3
4	Knowledge Graph Embedding Models	3
5	How to Train	3
6	Creating an Embedding Vector Database	5
6.1	Learning Embeddings	5
6.2	Loading Embeddings into Qdrant Vector Database	6
6.3	Launching Webservice	6
7	Answering Complex Queries	6
8	Predicting Missing Links	8
9	Downloading Pretrained Models	8
10	How to Deploy	8
11	Docker	8
12	Coverage Report	8
13	How to cite	10
14	dicee	12
14.1	Submodules	12
14.2	Attributes	225
14.3	Classes	225
14.4	Package Contents	225
	Python Module Index	239

DICE Embeddings¹: Hardware-agnostic Framework for Large-scale Knowledge Graph Embeddings:

1 Dicee Manual

Version: dicee 0.3.2

GitHub repository: <https://github.com/dice-group/dice-embeddings>

Publisher and maintainer: Caglar Demir²

Contact: caglar.demir@upb.de

License: OSI Approved :: MIT License

Dicee is a hardware-agnostic framework for large-scale knowledge graph embeddings.

Knowledge graph embedding research has mainly focused on learning continuous representations of knowledge graphs towards the link prediction problem. Recently developed frameworks can be effectively applied in a wide range of research-related applications. Yet, using these frameworks in real-world applications becomes more challenging as the size of the knowledge graph grows

We developed the DICE Embeddings framework (dicee) to compute embeddings for large-scale knowledge graphs in a hardware-agnostic manner. To achieve this goal, we rely on

1. **Pandas**³ & Co. to use parallelism at preprocessing a large knowledge graph,
2. **PyTorch**⁴ & Co. to learn knowledge graph embeddings via multi-CPU, GPUs, TPUs or computing cluster, and
3. **Huggingface**⁵ to ease the deployment of pre-trained models.

Why Pandas⁶ & Co. ? A large knowledge graph can be read and preprocessed (e.g. removing literals) by pandas, modin, or polars in parallel. Through polars, a knowledge graph having more than 1 billion triples can be read in parallel fashion. Importantly, using these frameworks allow us to perform all necessary computations on a single CPU as well as a cluster of computers.

Why PyTorch⁷ & Co. ? PyTorch is one of the most popular machine learning frameworks available at the time of writing. PytorchLightning facilitates scaling the training procedure of PyTorch without boilerplate. In our framework, we combine **PyTorch**⁸ & **PytorchLightning**⁹. Users can choose the trainer class (e.g., DDP by Pytorch) to train large knowledge graph embedding models with billions of parameters. PytorchLightning allows us to use state-of-the-art model parallelism techniques (e.g. Fully Sharded Training, FairScale, or DeepSpeed) without extra effort. With our framework, practitioners can directly use PytorchLightning for model parallelism to train gigantic embedding models.

Why Huggingface¹⁰? Seamlessly deploy and share pre-trained embedding models through the Huggingface ecosystem.

¹ <https://github.com/dice-group/dice-embeddings>

² <https://github.com/Demirrr>

³ <https://pandas.pydata.org/>

⁴ <https://pytorch.org/>

⁵ <https://huggingface.co/>

⁶ <https://pandas.pydata.org/>

⁷ <https://pytorch.org/>

⁸ <https://pytorch.org/>

⁹ <https://www.pytorchlightning.ai/>

¹⁰ <https://huggingface.co/>

2 Installation

2.1 Installation from Source

```
git clone https://github.com/dice-group/dice-embeddings.git
conda create -n dice python=3.10.13 --no-default-packages && conda activate dice &&
↪ cd dice-embeddings &&
pip3 install -e .
```

or

```
pip install dicee
```

3 Download Knowledge Graphs

```
wget https://files.dice-research.org/datasets/dice-embeddings/KGs.zip --no-check-
↪ certificate && unzip KGs.zip
```

To test the Installation

```
python -m pytest -p no:warnings -x # Runs >114 tests leading to > 15 mins
python -m pytest -p no:warnings --lf # run only the last failed test
python -m pytest -p no:warnings --ff # to run the failures first and then the rest of
↪ the tests.
```

4 Knowledge Graph Embedding Models

1. TransE, DistMult, ComplEx, ConEx, QMult, OMult, ConvO, ConvQ, Keci
2. All 44 models available in <https://github.com/pykeen/pykeen#models>

For more, please refer to examples.

5 How to Train

To Train a KGE model (KECI) and evaluate it on the train, validation, and test sets of the UMLS benchmark dataset.

```
from dicee.executer import Execute
from dicee.config import Namespace
args = Namespace()
args.model = 'Keci'
args.scoring_technique = "KvsAll" # 1vsAll, or AllvsAll, or NegSample
args.dataset_dir = "KGs/UMLS"
args.path_to_store_single_run = "Keci_UMLS"
args.num_epochs = 100
args.embedding_dim = 32
args.batch_size = 1024
reports = Execute(args).start()
print(reports["Train"]["MRR"]) # => 0.9912
print(reports["Test"]["MRR"]) # => 0.8155
# See the Keci_UMLS folder embeddings and all other files
```

where the data is in the following form

```
$ head -3 KGs/UMLS/train.txt
acquired_abnormality    location_of             experimental_model_of_disease
anatomical_abnormality  manifestation_of        physiologic_function
alga    isa      entity
```

A KGE model can also be trained from the command line

```
dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
```

dicee automatically detects available GPUs and trains a model with distributed data parallels technique. Under the hood, dicee uses lightning as a default trainer.

```
# Train a model by only using the GPU-0
CUDA_VISIBLE_DEVICES=0 dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model
↪ "train_val_test"
# Train a model by only using GPU-1
CUDA_VISIBLE_DEVICES=1 dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model
↪ "train_val_test"
NCCL_P2P_DISABLE=1 CUDA_VISIBLE_DEVICES=0,1 python dicee/scripts/run.py --trainer PL -
↪ --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
```

Under the hood, dicee executes run.py script and uses lightning as a default trainer

```
# Two equivalent executions
# (1)
dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
↪ 9753123402351737}
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
↪ 'MRR': 0.8072362996241839}
# Evaluate Keci on Test set: Evaluate Keci on Test set
# {'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
↪ 'MRR': 0.8064032293278861}

# (2)
CUDA_VISIBLE_DEVICES=0,1 python dicee/scripts/run.py --trainer PL --dataset_dir "KGs/
↪ UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
↪ 9753123402351737}
# Evaluate Keci on Train set: Evaluate Keci on Train set
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
↪ 'MRR': 0.8072362996241839}
# Evaluate Keci on Test set: Evaluate Keci on Test set
# {'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
↪ 'MRR': 0.8064032293278861}
```

Similarly, models can be easily trained with torchrun

```
torchrun --standalone --nnodes=1 --nproc_per_node=gpu dicee/scripts/run.py --trainer_
→torchDDP --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
→9753123402351737}
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
→'MRR': 0.8072499937521418}
# Evaluate Keci on Test set: Evaluate Keci on Test set
{'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
→'MRR': 0.8064032293278861}
```

You can also train a model in multi-node multi-gpu setting.

```
torchrun --nnodes 2 --nproc_per_node=gpu --node_rank 0 --rdzv_id 455 --rdzv_backend_
→c10d --rdzv_endpoint=nebula dicee/scripts/run.py --trainer torchDDP --dataset_dir_
→KGs/UMLS
torchrun --nnodes 2 --nproc_per_node=gpu --node_rank 1 --rdzv_id 455 --rdzv_backend_
→c10d --rdzv_endpoint=nebula dicee/scripts/run.py --trainer torchDDP --dataset_dir_
→KGs/UMLS
```

Train a KGE model by providing the path of a single file and store all parameters under newly created directory called KeciFamilyRun.

```
dicee --path_single_kg "KGs/Family/family-benchmark_rich_background.owl" --model Keci_
→--path_to_store_single_run KeciFamilyRun --backend rdflib
```

where the data is in the following form

```
$ head -3 KGs/Family/train.txt
_:1 <http://www.w3.org/1999/02/22-rdf-syntax-ns#type> <http://www.w3.org/2002/07/owl
→#Ontology> .
<http://www.benchmark.org/family#hasChild> <http://www.w3.org/1999/02/22-rdf-syntax-ns
→#type> <http://www.w3.org/2002/07/owl#ObjectProperty> .
<http://www.benchmark.org/family#hasParent> <http://www.w3.org/1999/02/22-rdf-syntax-
→ns#type> <http://www.w3.org/2002/07/owl#ObjectProperty> .
```

Apart from n-triples or standard link prediction dataset formats, we support ["owl", "nt", "turtle", "rdf/xml", "n3"]*. Moreover, a KGE model can be also trained by providing an endpoint of a triple store.

```
dicee --sparql_endpoint "http://localhost:3030/mutagenesis/" --model Keci
```

For more, please refer to examples.

6 Creating an Embedding Vector Database

6.1 Learning Embeddings

```
# Train an embedding model
dicee --dataset_dir KGs/Countries-S1 --path_to_store_single_run CountryEmbeddings --
→model Keci --p 0 --q 1 --embedding_dim 32 --adaptive_swa
```

6.2 Loading Embeddings into Qdrant Vector Database

```
# Ensure that Qdrant available
# docker pull qdrant/qdrant && docker run -p 6333:6333 -p 6334:6334 -v $(pwd)/
↪qdrant_storage:/qdrant/storage:z qdrant/qdrant
diceeindex --path_model "CountryEmbeddings" --collection_name "dummy" --location
↪"localhost"
```

6.3 Launching Webservice

```
diceeserve --path_model "CountryEmbeddings" --collection_name "dummy" --collection_
↪location "localhost"
```

Retrieve and Search

Get embedding of germany

```
curl -X 'GET' 'http://0.0.0.0:8000/api/get?q=germany' -H 'accept: application/json'
```

Get most similar things to europe

```
curl -X 'GET' 'http://0.0.0.0:8000/api/search?q=europe' -H 'accept: application/json'
{"result":[{"hit":"europe","score":1.0},
{"hit":"northern_europe","score":0.67126536},
{"hit":"western_europe","score":0.6010134},
{"hit":"puerto_rico","score":0.5051694},
{"hit":"southern_europe","score":0.4829831}]}
```

7 Answering Complex Queries

```
# pip install dicee
# wget https://files.dice-research.org/datasets/dice-embeddings/KGs.zip --no-check-
↪certificate & unzip KGs.zip
from dicee.executer import Execute
from dicee.config import Namespace
from dicee.knowledge_graph_embeddings import KGE
# (1) Train a KGE model
args = Namespace()
args.model = 'Keci'
args.p=0
args.q=1
args.optim = 'Adam'
args.scoring_technique = "AllvsAll"
args.path_single_kg = "KGs/Family/family-benchmark_rich_background.owl"
args.backend = "rdflib"
args.num_epochs = 200
args.batch_size = 1024
args.lr = 0.1
args.embedding_dim = 512
result = Execute(args).start()
# (2) Load the pre-trained model
```

(continues on next page)

(continued from previous page)

```
pre_trained_kge = KGE(path=result['path_experiment_folder'])
# (3) Single-hop query answering
# Query: ?E : \exist E.hasSibling(E, F9M167)
# Question: Who are the siblings of F9M167?
# Answer: [F9M157, F9F141], as (F9M167, hasSibling, F9M157) and (F9M167, hasSibling, ↵
↵F9F141)
predictions = pre_trained_kge.answer_multi_hop_query(query_type="1p",
                                                       query=('http://www.benchmark.org/
↵family#F9M167',
                                                       ('http://www.benchmark.
↵org/family#hasSibling',)),
                                                       tnorm="min", k=3)
top_entities = [topk_entity for topk_entity, query_score in predictions]
assert "http://www.benchmark.org/family#F9F141" in top_entities
assert "http://www.benchmark.org/family#F9M157" in top_entities
# (2) Two-hop query answering
# Query: ?D : \exist E.Married(D, E) \land hasSibling(E, F9M167)
# Question: To whom a sibling of F9M167 is married to?
# Answer: [F9F158, F9M142] as (F9M157 #married F9F158) and (F9F141 #married F9M142)
predictions = pre_trained_kge.answer_multi_hop_query(query_type="2p",
                                                       query=("http://www.benchmark.org/
↵family#F9M167",
                                                       ("http://www.benchmark.
↵org/family#hasSibling",
                                                       "http://www.benchmark.
↵org/family#married")),
                                                       tnorm="min", k=3)
top_entities = [topk_entity for topk_entity, query_score in predictions]
assert "http://www.benchmark.org/family#F9M142" in top_entities
assert "http://www.benchmark.org/family#F9F158" in top_entities
# (3) Three-hop query answering
# Query: ?T : \exist D.type(D,T) \land Married(D,E) \land hasSibling(E, F9M167)
# Question: What are the type of people who are married to a sibling of F9M167?
# (3) Answer: [Person, Male, Father] since F9M157 is [Brother Father Grandfather ↵
↵Male] and F9M142 is [Male Grandfather Father]
predictions = pre_trained_kge.answer_multi_hop_query(query_type="3p", query=("http://
↵www.benchmark.org/family#F9M167",
                                                       ("http://
↵www.benchmark.org/family#hasSibling",
                                                       "http://
↵www.benchmark.org/family#married",
                                                       "http://
↵www.w3.org/1999/02/22-rdf-syntax-ns#type")),
                                                       tnorm="min", k=5)
top_entities = [topk_entity for topk_entity, query_score in predictions]
print(top_entities)
assert "http://www.benchmark.org/family#Person" in top_entities
assert "http://www.benchmark.org/family#Father" in top_entities
assert "http://www.benchmark.org/family#Male" in top_entities
```

For more, please refer to examples/multi_hop_query_answering.

8 Predicting Missing Links

```
from dicee import KGE
# (1) Train a knowledge graph embedding model..
# (2) Load a pretrained model
pre_trained_kge = KGE(path='../')
# (3) Predict missing links through head entity rankings
pre_trained_kge.predict_topk(h=[".."],r=[".."],topk=10)
# (4) Predict missing links through relation rankings
pre_trained_kge.predict_topk(h=[".."],t=[".."],topk=10)
# (5) Predict missing links through tail entity rankings
pre_trained_kge.predict_topk(r=[".."],t=[".."],topk=10)
```

9 Downloading Pretrained Models

```
from dicee import KGE
# (1) Load a pretrained ConEx on DBpedia
model = KGE(url="https://files.dice-research.org/projects/DiceEmbeddings/KINSHIP-Keci-
↳dim128-epoch256-KvsAll")
```

- For more please look at dice-research.org/projects/DiceEmbeddings/¹¹

10 How to Deploy

```
from dicee import KGE
KGE(path='../').deploy(share=True, top_k=10)
```

11 Docker

To build the Docker image:

```
docker build -t dice-embeddings .
```

To test the Docker image:

```
docker run --rm -v ~/.local/share/dicee/KGs:/dicee/KGs dice-embeddings ./main.py --
↳model AConEx --embedding_dim 16
```

12 Coverage Report

The coverage report is generated using `coverage.py`¹²:

Name	Stmts	Miss	Cover	Missing
-----	-----	-----	-----	-----
dicee/___init___ .py	7	0	100%	
dicee/abstracts.py	338	115	66%	112-113, 115

(continues on next page)

¹¹ <https://files.dice-research.org/projects/DiceEmbeddings/>

¹² <https://coverage.readthedocs.io/en/7.6.0/>

(continued from previous page)

↪131, 154-155, 160, 173, 197, 240-254, 290, 303-306, 309-313, 353-364, 379-387, 402, ↪				
↪413-417, 427-428, 434-436, 442-445, 448-453, 576-596, 602-606, 610-612, 631, 658-696				
dicee/callbacks.py	248	103	58%	50-55, ↪
↪67-73, 76, 88-93, 98-103, 106-109, 116-133, 138-142, 146-147, 247, 281-285, 291-292, ↪				
↪310-316, 319, 324-325, 337-343, 349-358, 363-365, 410, 421-434, 438-473, 485-491				
dicee/config.py	97	2	98%	146-147
dicee/dataset_classes.py	430	146	66%	16, 44, ↪
↪57, 89-98, 104, 111-118, 121, 124, 127-151, 207-213, 216, 219-221, 324, 335-338, ↪				
↪354, 420-421, 439, 562-581, 583, 587-599, 606-615, 618, 622-636, 780-787, 790-794, ↪				
↪845, 866-878, 902-915, 937, 941-954, 964-967, 973, 985, 987, 989, 1012-1022				
dicee/eval_static_funcs.py	256	100	61%	104, 109, ↪
↪114, 261-356, 363-414, 442, 465-468				
dicee/evaluator.py	267	48	82%	48, 53, ↪
↪58, 77, 82-83, 86, 102, 119, 130, 134, 139, 173-184, 191-202, 310, 340-358, 452, ↪				
↪462, 480-485				
dicee/executer.py	134	16	88%	53-57, ↪
↪166-176, 235-236, 283				
dicee/knowledge_graph.py	82	10	88%	84, 94- ↪
↪95, 124, 128, 132-134, 137-138, 140				
dicee/knowledge_graph_embeddings.py	654	415	37%	25, 28- ↪
↪29, 37-50, 55-88, 91-125, 129-137, 171, 173-229, 261, 265, 276-277, 301-303, 311, ↪				
↪339-362, 493, 497-519, 523-547, 580, 656, 665, 710-716, 748, 806-1171, 1202-1263, ↪				
↪1267-1295, 1326, 1332				
dicee/models/___init___py	9	0	100%	
dicee/models/adopt.py	187	172	8%	50-86, ↪
↪99-110, 129-185, 195-242, 266-322, 346-448, 484-517				
dicee/models/base_model.py	240	35	85%	30-35, ↪
↪64, 66, 92, 99-116, 171, 204, 244, 250, 259, 262, 266, 273, 277, 279, 294, 307-308, ↪				
↪362, 365, 438, 450				
dicee/models/clifford.py	470	278	41%	10, 12, ↪
↪16, 24-25, 52-56, 79-87, 101-103, 108-109, 140-160, 184, 191, 195-256, 273-277, 289, ↪				
↪292, 297, 302, 346-361, 377-444, 464-470, 483, 486, 491, 496, 525-531, 544, 547, ↪				
↪552, 557, 567-576, 592-593, 613-685, 696-699, 724-749, 773-806, 842-846, 859, 869, ↪				
↪872, 877, 882, 887, 891, 895, 904-905, 935, 942, 947, 975-979, 1007-1016, 1026-1034, ↪				
↪1052-1054, 1072-1074, 1090-1092				
dicee/models/complex.py	162	25	85%	86-109, ↪
↪273-287				
dicee/models/dualE.py	59	10	83%	93-102, ↪
↪142-156				
dicee/models/ensemble.py	89	67	25%	7-29, 31, ↪
↪34, 37, 40, 43, 46, 49, 52-54, 56-58, 64-68, 71-90, 93-94, 97-112, 131				
dicee/models/function_space.py	262	221	16%	10-23, ↪
↪27-36, 39-48, 52-69, 76-87, 90-99, 102-111, 115-127, 135-157, 160-166, 169-186, 189- ↪				
↪195, 198-206, 209, 214-235, 244-247, 251-255, 259-268, 272-293, 302-308, 312-329, ↪				
↪333-336, 345-353, 356, 367-373, 393-407, 425-439, 444-454, 462-466, 475-479				
dicee/models/literal.py	33	1	97%	82
dicee/models/octonion.py	227	83	63%	21-44, ↪
↪320-329, 334-345, 348-370, 374-416, 426-474				
dicee/models/pykeen_models.py	55	5	91%	77-80, ↪
↪135				
dicee/models/quaternion.py	192	69	64%	7-21, 30- ↪
↪55, 68-72, 107, 185, 328-342, 345-364, 368-389, 399-426				

(continues on next page)

(continued from previous page)

dicee/models/real.py	61	12	80%	37-42, ↵
↪70-73, 91, 107-110				
dicee/models/static_funcs.py	10	0	100%	
dicee/models/transformers.py	234	189	19%	20-39, ↵
↪42, 56-71, 80-98, 101-112, 119-121, 124, 130-147, 151-176, 182-186, 189-193, 199-				
↪203, 206-208, 225-252, 261-264, 267-272, 275-300, 306-311, 315-368, 372-394, 400-410				
dicee/query_generator.py	374	346	7%	17-51, ↵
↪55, 61-64, 68-69, 77-91, 99-146, 154-187, 191-205, 211-268, 273-302, 306-442, 452-				
↪471, 479-502, 509-513, 518, 523-529				
dicee/read_preprocess_save_load_kg/___init___py	3	0	100%	
dicee/read_preprocess_save_load_kg/preprocess.py	243	40	84%	33, 39, ↵
↪76, 100-125, 131, 136-149, 175, 205, 380-381				
dicee/read_preprocess_save_load_kg/read_from_disk.py	36	11	69%	34, 38-
↪40, 47, 55, 58-72				
dicee/read_preprocess_save_load_kg/save_load_disk.py	53	21	60%	29-30, ↵
↪38, 47-68				
dicee/read_preprocess_save_load_kg/util.py	236	125	47%	159, 173-
↪175, 179-180, 198-204, 207-209, 214-216, 230, 244-247, 252-260, 265-271, 276-281, ↵				
↪286-291, 303-324, 330-386, 390-394, 398-399, 403, 407-408, 436, 441, 448-449				
dicee/sanity_checkers.py	47	19	60%	8-12, 21-
↪31, 46, 51, 58, 69-79				
dicee/static_funcs.py	483	194	60%	42, 52, ↵
↪58-63, 85, 92-96, 109-119, 129-131, 136, 143, 167, 172, 184, 190, 198, 202, 229-233,				
↪295, 303-309, 320-330, 341-361, 389, 413-414, 419-420, 437-438, 440-441, 443-444, ↵				
↪452, 470-474, 491-494, 498-503, 507-511, 515-516, 522-524, 539-553, 558-561, 566-				
↪569, 578-629, 634-646, 663-680, 683-691, 695-713, 724				
dicee/static_funcs_training.py	155	66	57%	7-10, ↵
↪222-319, 327-328				
dicee/static_preprocess_funcs.py	98	43	56%	17-25, ↵
↪50, 57, 59, 70, 83-107, 112-115, 120-123, 128-131				
dicee/trainer/___init___py	1	0	100%	
dicee/trainer/dice_trainer.py	151	18	88%	22, 30-
↪31, 33-35, 97, 104, 109-114, 152, 237, 280-283				
dicee/trainer/model_parallelism.py	99	87	12%	10-25, ↵
↪30-116, 121-132, 136, 141-197				
dicee/trainer/torch_trainer.py	77	6	92%	31, 102, ↵
↪168, 179-181				
dicee/trainer/torch_trainer_ddp.py	89	71	20%	11-14, ↵
↪43, 47-67, 78-94, 113-122, 126-136, 151-158, 168-191				

TOTAL	6948	3169	54%	

13 How to cite

Currently, we are working on our manuscript describing our framework. If you really like our work and want to cite it now, feel free to chose one :)

```
# Keci
@inproceedings{demir2023clifford,
  title={Clifford Embeddings--A Generalized Approach for Embedding in Normed Algebras}
  ↪,
```

(continues on next page)

```

    author={Demir, Caglar and Ngonga Ngomo, Axel-Cyrille},
    booktitle={Joint European Conference on Machine Learning and Knowledge Discovery in
↪Databases},
    pages={567--582},
    year={2023},
    organization={Springer}
}
# LitCQD
@inproceedings{demir2023litcqd,
    title={LitCQD: Multi-Hop Reasoning in Incomplete Knowledge Graphs with Numeric
↪Literals},
    author={Demir, Caglar and Wiebesiek, Michel and Lu, Renzhong and Ngonga Ngomo, Axel-
↪Cyrille and Heindorf, Stefan},
    booktitle={Joint European Conference on Machine Learning and Knowledge Discovery in
↪Databases},
    pages={617--633},
    year={2023},
    organization={Springer}
}
# DICE Embedding Framework
@article{demir2022hardware,
    title={Hardware-agnostic computation for large-scale knowledge graph embeddings},
    author={Demir, Caglar and Ngomo, Axel-Cyrille Ngonga},
    journal={Software Impacts},
    year={2022},
    publisher={Elsevier}
}
# KronE
@inproceedings{demir2022kronecker,
    title={Kronecker decomposition for knowledge graph embeddings},
    author={Demir, Caglar and Lienen, Julian and Ngonga Ngomo, Axel-Cyrille},
    booktitle={Proceedings of the 33rd ACM Conference on Hypertext and Social Media},
    pages={1--10},
    year={2022}
}
# QMult, OMult, ConvQ, ConvO
@InProceedings{pmlr-v157-demir21a,
    title = {Convolutional Hypercomplex Embeddings for Link Prediction},
    author = {Demir, Caglar and Moussallem, Diego and Heindorf, Stefan and Ngonga
↪Ngomo, Axel-Cyrille},
    booktitle = {Proceedings of The 13th Asian Conference on Machine Learning},
    pages = {656--671},
    year = {2021},
    editor = {Balasubramanian, Vineeth N. and Tsang, Ivor},
    volume = {157},
    series = {Proceedings of Machine Learning Research},
    month = {17--19 Nov},
    publisher = {PMLR},
    pdf = {https://proceedings.mlr.press/v157/demir21a/demir21a.pdf},
    url = {https://proceedings.mlr.press/v157/demir21a.html},
}
# ConEx

```

```
@inproceedings{demir2021convolutional,
  title={Convolutional Complex Knowledge Graph Embeddings},
  author={Caglar Demir and Axel-Cyrille Ngonga Ngomo},
  booktitle={Eighteenth Extended Semantic Web Conference - Research Track},
  year={2021},
  url={https://openreview.net/forum?id=6T45-4TFqaX}}
# Shallom
@inproceedings{demir2021shallow,
  title={A shallow neural model for relation prediction},
  author={Demir, Caglar and Moussallem, Diego and Ngomo, Axel-Cyrille Ngonga},
  booktitle={2021 IEEE 15th International Conference on Semantic Computing (ICSC)},
  pages={179--182},
  year={2021},
  organization={IEEE}
```

For any questions or wishes, please contact: caglar.demir@upb.de

14 dicee

DICE Embeddings - Knowledge Graph Embedding Library.

A library for training and using knowledge graph embedding models with support for various scoring techniques and training strategies.

Submodules:

evaluation: Model evaluation functions and Evaluator class models: KGE model implementations trainer: Training orchestration scripts: Utility scripts

14.1 Submodules

`dicee.__main__`

`dicee.abstracts`

Classes

<code>AbstractTrainer</code>	Abstract base class for KGE model trainers.
<code>BaseInteractiveKGE</code>	Base class for interactive, post-training use of KGE models.
<code>InteractiveQueryDecomposition</code>	Mixin that provides fuzzy-logic operators for multi-hop EPFO query answering.
<code>AbstractCallback</code>	Abstract base class for KGE training lifecycle callbacks.
<code>AbstractPPECallback</code>	Abstract base class for Post-training Parameter Ensembling (PPE) callbacks.
<code>BaseInteractiveTrainKGE</code>	Abstract/base class for training knowledge graph embedding models interactively.

Module Contents

class `dicee.abstracts.AbstractTrainer` (*args*, *callbacks*)

Abstract base class for KGE model trainers.

Provides the callback dispatch mechanism shared by all concrete trainer implementations (TorchTrainer, TorchD-
DPTrainer, etc.). Sub-classes call the `on_*` hooks at the appropriate points in the training loop so that any registered `AbstractCallback` can react.

Parameters

- **args** (*argparse.Namespace or similar*) – Processed configuration object. Must expose at least `random_seed` (int).
- **callbacks** (*list of AbstractCallback*) – Ordered list of callback instances to invoke at each lifecycle hook.

attributes

callbacks

is_global_zero = True

global_rank = 0

local_rank = 0

strategy = None

on_fit_start (*args, **kwargs)

Dispatch `on_fit_start` to all registered callbacks.

Called once before the first training epoch begins.

on_fit_end (*args, **kwargs)

Dispatch `on_fit_end` to all registered callbacks.

Called once after the last training epoch completes.

on_train_epoch_start (*args, **kwargs)

Dispatch `on_train_epoch_start` to all registered callbacks.

Called at the beginning of every epoch.

on_train_epoch_end (*args, **kwargs)

Dispatch `on_train_epoch_end` to all registered callbacks.

Called at the end of every epoch after the loss has been accumulated.

on_train_batch_end (*args, **kwargs)

Dispatch `on_train_batch_end` to all registered callbacks.

Called after each mini-batch gradient update.

static save_checkpoint (*full_path: str, model*) → None

Persist model weights to disk.

Parameters

- **full_path** (*str*) – Absolute or relative file path (including filename) where the `state_dict` will be written, e.g. `'Experiments/run1/model.pt'`.
- **model** (*torch.nn.Module*) – The model whose `state_dict` is to be saved.

```
class dicee.abstracts.BaseInteractiveKGE (path: str = None, url: str = None,  
    construct_ensemble: bool = False, model_name: str = None,  
    apply_semantic_constraint: bool = False)
```

Base class for interactive, post-training use of KGE models.

Loads a pre-trained model from disk (or a remote URL) together with its entity/relation index mappings and exposes the prediction API used by [KGE](#).

Parameters

- **path**(*str*, *optional*) – Path to the experiment directory produced by `Execute`. Must contain `model.pt`, `configuration.json`, `entity_to_idx.csv` and `relation_to_idx.csv`.
- **url**(*str*, *optional*) – Remote URL of a pre-trained model to download. Mutually exclusive with *path*.
- **construct_ensemble**(*bool*, *optional*) – When `True`, load all checkpoint files in *path* and average their weights to form an ensemble model. Defaults to `False`.
- **model_name**(*str*, *optional*) – Filename (without extension) of the checkpoint to load when multiple `.pt` files exist in *path*.
- **apply_semantic_constraint**(*bool*, *optional*) – Reserved for future use. Defaults to `False`.

`construct_ensemble = False`

`apply_semantic_constraint = False`

`configs`

`get_eval_report()` → dict

`get_bpe_token_representation(str_entity_or_relation: List[str] | str) → List[List[int]] | List[int]`

Parameters

str_entity_or_relation (*corresponds to a str or a list of strings to be tokenized via BPE and shaped.*)

Return type

A list integer(s) or a list of lists containing integer(s)

`get_padded_bpe_triple_representation(triples: List[List[str]]) → Tuple[List, List, List]`

Parameters

triples

`set_model_train_mode()` → None

Switch the underlying model to training mode.

Calls `model.train()` and re-enables gradient computation for all parameters so that subsequent calls to optimisation steps work correctly after a period of inference.

`set_model_eval_mode()` → None

Switch the underlying model to evaluation mode.

Calls `model.eval()` and freezes all parameters (`requires_grad = False`) so that dropout and batch-norm layers behave deterministically during inference.

property name

sample_entity (*n*: int) → List[str]

Return *n* random entity strings without replacement.

Parameters

n (*int*) – Number of entities to sample. Must be non-negative and at most `num_entities`.

Returns

Randomly selected entity string labels.

Return type

List[str]

sample_relation (*n*: int) → List[str]

Return *n* random relation strings without replacement.

Parameters

n (*int*) – Number of relations to sample. Must be non-negative and at most `num_relations`.

Returns

Randomly selected relation string labels.

Return type

List[str]

is_seen (*entity*: str = None, *relation*: str = None) → bool

Check whether an entity or relation was present in the training set.

Exactly one of *entity* or *relation* should be provided.

Parameters

- **entity** (*str*, *optional*) – Entity string label to look up.
- **relation** (*str*, *optional*) – Relation string label to look up.

Returns

True if the given string is in the respective index mapping, False otherwise.

Return type

bool

save () → None

Persist the current model weights to the experiment directory.

The checkpoint filename encodes the current timestamp so successive calls do not overwrite each other. Ensemble models are saved with an `_ensemble_` infix in the filename.

get_entity_index (*x*: str) → int

Return the integer index for a given entity string.

Parameters

x (*str*) – Entity string label (must have been seen during training).

Returns

Corresponding row index in the entity embedding matrix.

Return type

int

get_relation_index (*x*: str) → int

Return the integer index for a given relation string.

Parameters

x (*str*) – Relation string label (must have been seen during training).

Returns

Corresponding row index in the relation embedding matrix.

Return type

int

index_triple (*head_entity: List[str], relation: List[str], tail_entity: List[str]*)
→ Tuple[torch.LongTensor, torch.LongTensor, torch.LongTensor]

Convert string triple lists to integer index tensors.

Parameters

- **head_entity** (*List[str]*) – Head entity string labels.
- **relation** (*List[str]*) – Relation string labels.
- **tail_entity** (*List[str]*) – Tail entity string labels.

Returns

idx_head_entity, idx_relation, idx_tail_entity – Each has shape (n, 1) containing the integer indices for the corresponding strings.

Return type

torch.LongTensor

add_new_entity_embeddings (*entity_name: str = None, embeddings: torch.FloatTensor = None*)
→ None

Extend the entity embedding table with a new entity at inference time.

The new entity is appended to both `entity_to_idx` / `idx_to_entity` mappings and the `entity_embeddings` weight tensor so that subsequent calls to prediction methods can reference it by name.

Parameters

- **entity_name** (*str*) – String label for the new entity. If the entity already exists in the index no modification is made.
- **embeddings** (*torch.FloatTensor*) – 1-D float tensor of length `embedding_dim` containing the pre-computed embedding for the new entity.

get_entity_embeddings (*items: List[str]*) → torch.FloatTensor

Return the embedding vectors for the given entity strings.

For standard (non-BPE) models the method looks up each string in `entity_to_idx` and returns the corresponding rows of the entity embedding matrix. For BPE models subword token embeddings are fetched and flattened into a single vector per entity.

Parameters

items (*List[str]*) – Entity string labels to retrieve.

Returns

Shape (len(items), embedding_dim).

Return type

torch.FloatTensor

get_relation_embeddings (*items: List[str]*) → torch.FloatTensor

Return the embedding vectors for the given relation strings.

Parameters

items (*List[str]*) – Relation string labels to retrieve.

Returns

Shape (len(items), embedding_dim).

Return type

`torch.FloatTensor`

construct_input_and_output (*head_entity: List[str], relation: List[str], tail_entity: List[str], labels*)
→ `Tuple[torch.LongTensor, torch.FloatTensor]`

Build an indexed triple tensor and a label tensor from string inputs.

Parameters

- **head_entity** (*List[str]*) – Head entity string labels.
- **relation** (*List[str]*) – Relation string labels.
- **tail_entity** (*List[str]*) – Tail entity string labels.
- **labels** (*list or array-like*) – Binary or soft labels (one per triple) used as training targets.

Returns

- **x** (*torch.LongTensor*) – Shape $(n, 3)$ integer-indexed triples.
- **labels** (*torch.FloatTensor*) – Shape $(n,)$ float label tensor.

parameters ()

class `dicee.abstracts.InteractiveQueryDecomposition`

Mixin that provides fuzzy-logic operators for multi-hop EPFO query answering.

The three families of operators — T-norm, T-conorm, and negation norm — are applied element-wise over entity score tensors to compose complex queries from atomic link-prediction results (e.g. 2p, 3p, 2i, ip, up).

t_norm (*tens_1: torch.Tensor, tens_2: torch.Tensor, tnorm: str = 'min'*) → `torch.Tensor`

Apply a T-norm to combine two entity score distributions.

Parameters

- **tens_1** (*torch.Tensor*) – Score tensors of identical shape, values in $[0, 1]$.
- **tens_2** (*torch.Tensor*) – Score tensors of identical shape, values in $[0, 1]$.
- **tnorm** (*str*) – Operator to use. 'min' applies the Gödel (min) T-norm; 'prod' applies the product T-norm.

Returns

Element-wise combined scores of the same shape as the inputs.

Return type

`torch.Tensor`

tensor_t_norm (*subquery_scores: torch.FloatTensor, tnorm: str = 'min'*) → `torch.FloatTensor`

Compute T-norm over $[0,1]^{n \times d}$ where n denotes the number of hops and d denotes number of entities

t_conorm (*tens_1: torch.Tensor, tens_2: torch.Tensor, tconorm: str = 'min'*) → `torch.Tensor`

Apply a T-conorm (S-norm) to combine two score distributions (union).

Parameters

- **tens_1** (*torch.Tensor*) – Score tensors of identical shape, values in $[0, 1]$.
- **tens_2** (*torch.Tensor*) – Score tensors of identical shape, values in $[0, 1]$.
- **tconorm** (*str*) – Operator to use. 'min' applies the Gödel (max) T-conorm; 'prod' applies the probabilistic sum T-conorm.

Returns

Element-wise combined scores of the same shape as the inputs.

Return type

`torch.Tensor`

negnorm (*tens_1*: `torch.Tensor`, *lambda_*: `float`, *neg_norm*: `str` = 'standard') → `torch.Tensor`

Apply a negation norm (complement) to an entity score distribution.

Parameters

- **tens_1** (`torch.Tensor`) – Input score tensor, values in $[0, 1]$.
- **lambda** (`float`) – Shape parameter used by the Sugeno and Yager negation norms. Ignored for the standard complement.
- **neg_norm** (`str`) – Which negation to apply: 'standard' ($1 - x$), 'sugeno', or 'yager'.

Returns

Complemented score tensor of the same shape as *tens_1*.

Return type

`torch.Tensor`

class `dicee.abstracts.AbstractCallback`

Bases: `abc.ABC`, `lightning.pytorch.callbacks.Callback`

Abstract base class for KGE training lifecycle callbacks.

Concrete sub-classes override one or more hook methods to perform custom actions at specific points during training (e.g. weight averaging, periodic evaluation, model checkpointing). All hooks have empty default implementations so sub-classes only need to override the hooks they care about.

Callbacks are registered by passing them to the trainer's *callbacks* list. They are also compatible with PyTorch Lightning trainers because this class extends `lightning.pytorch.callbacks.Callback`.

on_init_start (**args*, ***kwargs*)

Called when the trainer is about to be constructed.

Override to perform setup that must happen before any trainer state is initialised.

on_init_end (**args*, ***kwargs*)

Called immediately after the trainer has been constructed.

Override to perform setup that requires a fully initialised trainer.

on_fit_start (*trainer*, *model*)

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

on_train_epoch_end (*trainer*, *model*)

Called at the end of each training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

on_train_batch_end(*args, **kwargs)

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

on_fit_end(*args, **kwargs)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

class dicee.abstracts.**AbstractPPECallback**(num_epochs, path, epoch_to_start,
last_percent_to_consider)

Bases: *AbstractCallback*

Abstract base class for Post-training Parameter Ensembling (PPE) callbacks.

Sub-classes implement weight-averaging strategies (SWA, EMA, SWAG, ...) by overriding `on_train_epoch_end()` and `on_fit_end()`. Common book-keeping (epoch counter, sample counter, alpha weights) is managed here.

Parameters

- **num_epochs** (*int*) – Total number of training epochs.
- **path** (*str*) – Experiment directory where averaged checkpoints will be written.
- **epoch_to_start** (*int*) – First epoch at which the averaging procedure should begin.
- **last_percent_to_consider** (*float*) – Fraction of the total training epochs (counted from the end) whose checkpoints are included in the ensemble.

num_epochs

path

sample_counter = 0

epoch_count = 0

alphas = None

on_fit_start(trainer, model)

Called once before the first training epoch.

Parameters

- **trainer** (*AbstractTrainer* or *pl.Trainer*) – The active trainer instance.
- **model** (*BaseKGE*) – The model about to be trained.

on_fit_end(trainer, model)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

store_ensemble(param_ensemble) → None

class dicee.abstracts.**BaseInteractiveTrainKGE**

Abstract/base class for training knowledge graph embedding models interactively. This class provides methods for re-training KGE models and also Literal Embedding model.

train_triples (*h: List[str], r: List[str], t: List[str], labels: List[float], iteration=2, optimizer=None*)

train_k_vs_all (*h, r, iteration=1, lr=0.001*)

Train k vs all :param head_entity: :param relation: :param iteration: :param lr: :return:

train (*kg, lr=0.1, epoch=10, batch_size=32, neg_sample_ratio=10, num_workers=1*) → None

Retrained a pretrain model on an input KG via negative sampling.

train_literals (*train_file_path: str = None, num_epochs: int = 100, lit_lr: float = 0.001, lit_normalization_type: str = 'z-norm', batch_size: int = 1024, sampling_ratio: float = None, random_seed=1, loader_backend: str = 'pandas', freeze_entity_embeddings: bool = True, gate_residual: bool = True, device: str = None, suffle_data: bool = True*)

Trains the Literal Embeddings model using literal data.

Parameters

- **train_file_path** (*str*) – Path to the training data file.
- **num_epochs** (*int*) – Number of training epochs.
- **lit_lr** (*float*) – Learning rate for the literal model.
- **norm_type** (*str*) – Normalization type to use ('z-norm', 'min-max', or None).
- **batch_size** (*int*) – Batch size for training.
- **sampling_ratio** (*float*) – Ratio of training triples to use.
- **loader_backend** (*str*) – Backend for loading the dataset ('pandas' or 'rdflib').
- **freeze_entity_embeddings** (*bool*) – If True, freeze the entity embeddings during training.
- **gate_residual** (*bool*) – If True, use gate residual connections in the model.
- **device** (*str*) – Device to use for training ('cuda' or 'cpu'). If None, will use available GPU or CPU.
- **suffle_data** (*bool*) – If True, shuffle the dataset before training.

dicee.analyse_experiments

This script should be moved to dicee/scripts Example: python dicee/analyse_experiments.py –dir Experiments –features “model” “trainMRR” “testMRR”

Classes

Experiment

Functions

get_default_arguments()
analyse(args)

Module Contents

```
dicee.analyse_experiments.get_default_arguments()
```

```
class dicee.analyse_experiments.Experiment
```

```
    model_name = []

    callbacks = []

    embedding_dim = []

    num_params = []

    num_epochs = []

    batch_size = []

    lr = []

    byte_pair_encoding = []

    aswa = []

    path_dataset_folder = []

    full_storage_path = []

    pq = []

    train_mrr = []

    train_h1 = []

    train_h3 = []

    train_h10 = []

    val_mrr = []

    val_h1 = []

    val_h3 = []

    val_h10 = []

    test_mrr = []

    test_h1 = []

    test_h3 = []

    test_h10 = []

    runtime = []

    normalization = []

    scoring_technique = []

    save_experiment(x)
```

`to_df()`

`dicee.analyse_experiments.analyse(args)`

dicee.callbacks

Callbacks for training lifecycle events.

Provides callback classes for various training events including epoch end, model saving, weight averaging, and evaluation.

Classes

<i>AccumulateEpochLossCallback</i>	Callback to store epoch losses to a CSV file.
<i>PrintCallback</i>	Callback that prints training start/end times and total run-time.
<i>KGESaveCallback</i>	Callback that periodically saves model checkpoints during training.
<i>PseudoLabellingCallback</i>	Callback that augments the training set with pseudo-labelled triples.
<i>Eval</i>	Abstract base class for KGE training lifecycle callbacks.
<i>KronE</i>	Abstract base class for KGE training lifecycle callbacks.
<i>Perturb</i>	A callback for a three-Level Perturbation
<i>PeriodicEvalCallback</i>	Callback to periodically evaluate the model and optionally save checkpoints during training.
<i>LRScheduler</i>	Callback for managing learning rate scheduling and model snapshots.

Functions

<i>estimate_q(eps)</i>	estimate rate of convergence q from sequence esp
<i>compute_convergence(seq, i)</i>	

Module Contents

class `dicee.callbacks.AccumulateEpochLossCallback` (*path: str*)

Bases: `dicee.abstracts.AbstractCallback`

Callback to store epoch losses to a CSV file.

Parameters

path – Directory path where the loss file will be saved.

path

on_fit_end (*trainer, model*) → None

Store epoch loss history to CSV file.

Parameters

- **trainer** – The trainer instance.
- **model** – The model being trained.

class dicee.callbacks.**PrintCallback**

Bases: *dicee.abstracts.AbstractCallback*

Callback that prints training start/end times and total runtime.

start_time

on_fit_start (*trainer, pl_module*)

Called once before the first training epoch.

Parameters

- **trainer** (*AbstractTrainer* or *pl.Trainer*) – The active trainer instance.
- **model** (*BaseKGE*) – The model about to be trained.

on_fit_end (*trainer, pl_module*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

on_train_batch_end (**args, **kwargs*)

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

on_train_epoch_end (**args, **kwargs*)

Called at the end of each training epoch.

Parameters

- **trainer** (*AbstractTrainer* or *pl.Trainer*) – The active trainer instance.
- **model** (*BaseKGE*) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

class dicee.callbacks.**KGESaveCallback** (*every_x_epoch: int, max_epochs: int, path: str*)

Bases: *dicee.abstracts.AbstractCallback*

Callback that periodically saves model checkpoints during training.

Parameters

- **every_x_epoch** (*int* or *None*) – Save a checkpoint every *every_x_epoch* epochs. When *None*, the interval defaults to `max(max_epochs // 2, 1)`.
- **max_epochs** (*int*) – Total number of training epochs (used to compute the default interval when *every_x_epoch* is *None*).
- **path** (*str*) – Directory where checkpoint files will be written.

every_x_epoch

max_epochs

epoch_counter = 0

path

on_train_batch_end (**args, **kwargs*)

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

`on_fit_start (trainer, pl_module)`

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

`on_train_epoch_end (*args, **kwargs)`

Called at the end of each training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

`on_fit_end (*args, **kwargs)`

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

`on_epoch_end (model, trainer, **kwargs)`

class `dicee.callbacks.PseudoLabellingCallback (data_module, kg, batch_size)`

Bases: `dicee.abstracts.AbstractCallback`

Callback that augments the training set with pseudo-labelled triples.

At the end of each epoch the current model scores a batch of unlabelled triples and those with a predicted probability ≥ 0.90 are appended to the training dataset (semi-supervised self-training).

Parameters

- **data_module** (*object*) – Dataset module exposing a `train_set_idx` attribute and a `train_data_loader()` method.
- **kg** (*KG*) – The knowledge graph, providing `num_entities`, `num_relations`, and `unlabelled_set`.
- **batch_size** (*int*) – Number of unlabelled triples to sample and score each epoch.

data_module

kg

num_of_epochs = 0

unlabelled_size

batch_size

create_random_data()

on_epoch_end (trainer, model)

`dicee.callbacks.estimate_q (eps)`

estimate rate of convergence q from sequence esp

`dicee.callbacks.compute_convergence (seq, i)`


```
class dicee.callbacks.Eval (path, epoch_ratio: int = None)
```

Bases: `dicee.abstracts.AbstractCallback`

Abstract base class for KGE training lifecycle callbacks.

Concrete sub-classes override one or more hook methods to perform custom actions at specific points during training (e.g. weight averaging, periodic evaluation, model checkpointing). All hooks have empty default implementations so sub-classes only need to override the hooks they care about.

Callbacks are registered by passing them to the trainer's *callbacks* list. They are also compatible with PyTorch Lightning trainers because this class extends `lightning.pytorch.callbacks.Callback`.

path

reports = []

epoch_ratio

epoch_counter = 0

on_fit_start (*trainer, model*)

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

on_fit_end (*trainer, model*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

on_train_epoch_end (*trainer, model*)

Called at the end of each training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

on_train_batch_end (**args, **kwargs*)

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

```
class dicee.callbacks.KronE
```

Bases: `dicee.abstracts.AbstractCallback`

Abstract base class for KGE training lifecycle callbacks.

Concrete sub-classes override one or more hook methods to perform custom actions at specific points during training (e.g. weight averaging, periodic evaluation, model checkpointing). All hooks have empty default implementations so sub-classes only need to override the hooks they care about.

Callbacks are registered by passing them to the trainer's *callbacks* list. They are also compatible with PyTorch Lightning trainers because this class extends `lightning.pytorch.callbacks.Callback`.

f = None

static batch_kronecker_product (*a, b*)

Kronecker product of matrices a and b with leading batch dimensions. Batch dimensions are broadcast. The number of them must be the same. :type a: torch.Tensor :type b: torch.Tensor :rtype: torch.Tensor

get_kronecker_triple_representation (*indexed_triple: torch.LongTensor*)

Get kronecker embeddings

on_fit_start (*trainer, model*)

Called once before the first training epoch.

Parameters

- **trainer** (*AbstractTrainer or pl.Trainer*) – The active trainer instance.
- **model** (*BaseKGE*) – The model about to be trained.

class `dicke.callbacks.Perturb` (*level: str = 'input', ratio: float = 0.0, method: str = None, scaler: float = None, frequency=None*)

Bases: `dicke.abstracts.AbstractCallback`

A callback for a three-Level Perturbation

Input Perturbation: During training an input x is perturbed by randomly replacing its element. In the context of knowledge graph embedding models, x can denote a triple, a tuple of an entity and a relation, or a tuple of two entities. A perturbation means that a component of x is randomly replaced by an entity or a relation.

Parameter Perturbation:

Output Perturbation:

level = 'input'

ratio = 0.0

method = None

scaler = None

frequency = None

on_train_batch_start (*trainer, model, batch, batch_idx*)

Called when the train batch begins.

class `dicke.callbacks.PeriodicEvalCallback` (*experiment_path: str, max_epochs: int, eval_every_n_epoch: int = 0, eval_at_epochs: list = None, save_model_every_n_epoch: bool = True, n_epochs_eval_model: str = 'val_test'*)

Bases: `dicke.abstracts.AbstractCallback`

Callback to periodically evaluate the model and optionally save checkpoints during training.

Evaluates at regular intervals (every N epochs) or at explicitly specified epochs. Stores evaluation reports and model checkpoints.

experiment_dir

max_epochs

epoch_counter = 0

save_model_every_n_epoch = True

reports

```

n_epochs_eval_model = 'val_test'

default_eval_model = None

eval_epochs

on_fit_end(trainer, model)
    Called at the end of training. Saves final evaluation report.

on_train_epoch_end(trainer, model)
    Called at the end of each training epoch. Performs evaluation and checkpointing if scheduled.

```

class `dicee.callbacks.LRScheduler` (*adaptive_lr_config: dict, total_epochs: int, experiment_dir: str, eta_max: float = 0.1, snapshot_dir: str = 'snapshots'*)

Bases: `dicee.abstracts.AbstractCallback`

Callback for managing learning rate scheduling and model snapshots.

Supports cosine annealing (“cca”), MMCCLR (“mmcclr”), and their deferred (warmup) variants: - “deferred_cca” - “deferred_mmcclr”

At the end of each learning rate cycle, the model can optionally be saved as a snapshot.

```

total_epochs

experiment_dir

snapshot_dir

batches_per_epoch = None

total_steps = None

cycle_length = None

warmup_steps = None

lr_lambda = None

scheduler = None

step_count = 0

snapshot_loss

on_train_start(trainer, model)
    Initialize training parameters and LR scheduler at start of training.

on_train_batch_end(trainer, model, outputs, batch, batch_idx)
    Step the LR scheduler and save model snapshot if needed after each batch.

on_fit_end(trainer, model)
    Called once after the final training epoch completes.

    Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

```

dicee.config

Configuration module for DICE embeddings.

Provides the Namespace class with default configuration values for training knowledge graph embedding models.

Classes

<i>Namespace</i>	Extended Namespace with default KGE training configuration.
------------------	---

Module Contents

```
class dicee.config.Namespace (**kwargs)
```

Bases: argparse.Namespace

Extended Namespace with default KGE training configuration.

Provides sensible defaults for all training parameters while allowing easy customization through command-line arguments or direct assignment.

```
dataset_dir: str = None
```

The path of a folder containing train.txt, and/or valid.txt and/or test.txt

```
save_embeddings_as_csv: bool = False
```

Embeddings of entities and relations are stored into CSV files to facilitate easy usage.

```
storage_path: str = 'Experiments'
```

A directory named with time of execution under `storage_path` that contains related data about embeddings.

```
path_to_store_single_run: str = None
```

A single directory created that contains related data about embeddings.

```
path_single_kg = None
```

Path of a file corresponding to the input knowledge graph

```
sparql_endpoint = None
```

An endpoint of a triple store.

```
model: str = 'Keci'
```

KGE model

```
optim: str = 'Adam'
```

Optimizer

```
embedding_dim: int = 64
```

Size of continuous vector representation of an entity/relation

```
num_epochs: int = 150
```

Number of pass over the training data

```
batch_size: int = 1024
```

Mini-batch size if it is None, an automatic batch finder technique applied

```
lr: float = 0.1
```

Learning rate

add_noise_rate: float = None

The ratio of added random triples into training dataset

gpus = None

Number GPUs to be used during training

callbacks

10}}

Type

Callbacks, e.g., {"PPE"

Type

{ "last_percent_to_consider"

backend: str = 'pandas'

Backend to read, process, and index input knowledge graph. pandas, polars and rdflib available

separator: str = '\\s+'

separator for extracting head, relation and tail from a triple

trainer: str = 'torchCPUTrainer'

Trainer for knowledge graph embedding model

scoring_technique: str = 'KvsAll'

Scoring technique for knowledge graph embedding models

neg_ratio: int = 0

Negative ratio for a true triple in NegSample training_technique

weight_decay: float = 0.0

Weight decay for all trainable params

normalization: str = 'None'

LayerNorm, BatchNorm1d, or None

init_param: str = None

xavier_normal or None

gradient_accumulation_steps: int = 0

Not tested e

num_folds_for_cv: int = 0

Number of folds for CV

eval_model: str = 'train_val_test'

["None", "train", "train_val", "train_val_test", "test"]

Type

Evaluate trained model choices

save_model_at_every_epoch: int = None

Not tested

label_smoothing_rate: float = 0.0

num_core: int = 0

Number of CPUs to be used in the mini-batch loading process

random_seed: int = 0
Random Seed

sample_triples_ratio: float = None
Read some triples that are uniformly at random sampled. Ratio being between 0 and 1

read_only_few: int = None
Read only first few triples

pykeen_model_kwargs
Additional keyword arguments for pykeen models

pl_trainer_kwargs
Additional keyword arguments for the PyTorch Lightning Trainer

kernel_size: int = 3
Size of a square kernel in a convolution operation

num_of_output_channels: int = 32
Number of slices in the generated feature map by convolution.

p: int = 0
P parameter of Clifford Embeddings

q: int = 1
Q parameter of Clifford Embeddings

input_dropout_rate: float = 0.0
Dropout rate on embeddings of input triples

hidden_dropout_rate: float = 0.0
Dropout rate on hidden representations of input triples

feature_map_dropout_rate: float = 0.0
Dropout rate on a feature map generated by a convolution operation

byte_pair_encoding: bool = False
Byte pair encoding

Type
WIP

adaptive_swa: bool = False
Adaptive stochastic weight averaging

swa: bool = False
Stochastic weight averaging

swag: bool = False
Stochastic weight averaging - Gaussian

ema: bool = False
Exponential Moving Average

twa: bool = False
Trainable weight averaging

block_size: int = None
block size of LLM

continual_learning = None
Path of a pretrained model size of LLM

auto_batch_finding = False
A flag for using auto batch finding

eval_every_n_epochs: int = 0
Evaluate model every n epochs. If 0, no evaluation is applied.

save_every_n_epochs: bool = False
Save model every n epochs. If True, save model at every epoch.

eval_at_epochs: list = None
List of epoch numbers at which to evaluate the model (e.g., 1 5 10).

n_epochs_eval_model: str = 'val_test'
Evaluating link prediction performance on data splits while performing periodic evaluation.

adaptive_lr
“cca”}
Type
Adaptive learning rate parameters, e.g., “{“scheduler_name”

swa_start_epoch: int = None
Epoch at which to start applying stochastic weight averaging.

swa_c_epochs: int = 1
Number of epochs to average over for SWA, SWAG, EMA, TWA.

__iter__()

dicee.dataset_classes

Dataset classes for knowledge graph embedding training.

This package groups the various `torch.utils.data.Dataset` implementations used throughout DICE Embeddings into thematic sub-modules:

- `_bpe` – Byte-pair-encoding related datasets.
- `_negative_sampling` – Negative sampling based datasets.
- `_label_based` – Multi-label / multi-class scoring datasets.
- `_literal` – Literal (numeric) embedding dataset.
- `_factory` – `construct_dataset` / `reload_dataset` helpers.

All public names are re-exported here so that existing `from dicee.dataset_classes import ...` statements continue to work unchanged.

Classes

<i>BPE_NegativeSamplingDataset</i>	Dataset for negative sampling with byte-pair encoded triples.
<i>MultiClassClassificationDataset</i>	Dataset for autoregressive multi-class classification on sub-word units.

continues on next page

Table 7 – continued from previous page

<i>MultiLabelDataset</i>	Multi-label dataset for BPE-based KvsAll / AllvsAll training.
<i>AllvsAll</i>	Dataset for AllvsAll training (multi-label, exhaustive).
<i>KvsAll</i>	Dataset for KvsAll training (multi-label).
<i>KvsSampleDataset</i>	Dataset for KvsSample training (dynamic multi-label).
<i>OnevsAllDataset</i>	Dataset for the 1-vs-All training strategy (multi-class).
<i>LiteralDataset</i>	Dataset for loading and processing literal data for Literal Embedding models.
<i>FixedNegSampleDataset</i>	Pre-computed (fixed) negative sampling dataset.
<i>OnevsSample</i>	Dataset for 1-vs-Sample training (dynamic multi-class with negatives).
<i>TriplePredictionDataset</i>	Dataset for triple prediction with on-the-fly negative sampling.

Functions

<i>construct_dataset</i> (→ torch.utils.data.Dataset)	Build the appropriate dataset for the given training configuration.
<i>reload_dataset</i> (path, form_of_labelling, ...)	Reload training data from disk and construct a Pytorch dataset.

Package Contents

class `dicee.dataset_classes.BPE_NegativeSamplingDataset` (*train_set: torch.LongTensor, ordered_shaped_bpe_entities: torch.LongTensor, neg_ratio: int*)

Bases: `torch.utils.data.Dataset`

Dataset for negative sampling with byte-pair encoded triples.

Each sample is a BPE-encoded triple. The custom `collate_fn` constructs negatives by corrupting head or tail entities with random BPE entities.

Parameters

- **train_set** (*torch.LongTensor*) – Integer-encoded triples of shape $(N, 3)$.
- **ordered_shaped_bpe_entities** (*torch.LongTensor*) – All BPE entity representations, ordered by entity index.
- **neg_ratio** (*int*) – Number of negative samples per positive triple.

train_set

ordered_bpe_entities

num_bpe_entities

neg_ratio

num_datapoints

__len__()

__getitem__ (*idx*)

`collate_fn (batch_shaped_bpe_triples: List[Tuple[torch.Tensor, torch.Tensor]])`

```
class dicee.dataset_classes.MultiClassClassificationDataset(  
    subword_units: numpy.ndarray, block_size: int = 8)
```

Bases: `torch.utils.data.Dataset`

Dataset for autoregressive multi-class classification on sub-word units.

Splits a flat sequence of sub-word token ids into overlapping windows of size `block_size` for next-token prediction.

Parameters

- **subword_units** (`numpy.ndarray`) – 1-D array of sub-word token ids.
- **block_size** (`int`, *optional*) – Context window length (default 8).

`train_data`

`block_size = 8`

`num_of_data_points`

`collate_fn = None`

`__len__()`

`__getitem__(idx)`

```
class dicee.dataset_classes.MultiLabelDataset (train_set: torch.LongTensor,  
    train_indices_target: torch.LongTensor, target_dim: int,  
    torch_ordered_shaped_bpe_entities: torch.LongTensor)
```

Bases: `torch.utils.data.Dataset`

Multi-label dataset for BPE-based KvsAll / AllvsAll training.

Each sample is a BPE-encoded (head, relation) pair together with a binary multi-label target vector over all entities.

Parameters

- **train_set** (`torch.LongTensor`) – BPE-encoded input pairs of shape (N, 2, token_length).
- **train_indices_target** (`torch.LongTensor`) – Per-sample lists of positive target entity indices.
- **target_dim** (`int`) – Dimensionality of the target vector (number of entities).
- **torch_ordered_shaped_bpe_entities** (`torch.LongTensor`) – Ordered BPE entity representations.

`train_set`

`train_indices_target`

`target_dim`

`num_datapoints`

`torch_ordered_shaped_bpe_entities`

`collate_fn = None`

`__len__()`

`__getitem__(idx)`

`dicee.dataset_classes.construct_dataset` (* (Keyword-only parameters separator (PEP 3102)),
train_set: `numpy.ndarray` | `list`, valid_set=None, test_set=None, ordered_bpe_entities=None,
train_target_indices=None, target_dim: `int` = None, entity_to_idx: `dict`, relation_to_idx: `dict`,
form_of_labelling: `str`, scoring_technique: `str`, neg_ratio: `int`, label_smoothing_rate: `float`,
byte_pair_encoding=None, block_size: `int` = None, seed: `int` = None) → `torch.utils.data.Dataset`

Build the appropriate dataset for the given training configuration.

Parameters

- **train_set** (`numpy.ndarray` or `list`) – Raw integer-indexed triples.
- **entity_to_idx** (`dict`) – Name → index mappings.
- **relation_to_idx** (`dict`) – Name → index mappings.
- **form_of_labelling** (`str`) – 'EntityPrediction' or 'RelationPrediction'.
- **scoring_technique** (`str`) – One of 'NegSample', 'FixedNegSample', '1vsAll', '1vsSample', 'KvsAll', 'AllvsAll', 'KvsSample'.
- **neg_ratio** (`int`) – Negative sample ratio.
- **label_smoothing_rate** (`float`) – Label smoothing coefficient.

Return type

`torch.utils.data.Dataset`

`dicee.dataset_classes.reload_dataset` (`path`: `str`, `form_of_labelling`, `scoring_technique`, `neg_ratio`,
`label_smoothing_rate`)

Reload training data from disk and construct a Pytorch dataset.

class `dicee.dataset_classes.AllvsAll` (`train_set_idx`: `numpy.ndarray`, `entity_idx`s, `relation_idx`s,
`label_smoothing_rate`=0.0)

Bases: `torch.utils.data.Dataset`

Dataset for AllvsAll training (multi-label, exhaustive).

Extends the KvsAll idea: every *possible* (entity, relation) combination is included — not just those observed in the KG. Pairs without any known tail entities receive an all-zeros label vector.

Parameters

- **train_set_idx** (`numpy.ndarray`) – (N, 3) integer-indexed triples.
- **entity_idx**s (`dict`) – Entity-name → index mapping.
- **relation_idx**s (`dict`) – Relation-name → index mapping.
- **label_smoothing_rate** (`float`, *optional*) – Label smoothing coefficient (default 0.0).

`train_data` = None

`train_target` = None

`label_smoothing_rate`

`collate_fn` = None

target_dim

__len__()

__getitem__(idx)

```
class dicee.dataset_classes.KvsAll (train_set_idx: numpy.ndarray, entity_idx, relation_idx, form,  
    store=None, label_smoothing_rate: float = 0.0)
```

Bases: torch.utils.data.Dataset

Dataset for KvsAll training (multi-label).

D := {(x, y)_i | i=1^N} where

- x = (h, r) is a unique (entity, relation) pair observed in the KG,
- y ∈ [0, 1]^{**IE**} is a multi-label vector with y_j = 1 iff (h, r, e_j) ∈ KG.

Parameters

- **train_set_idx** (numpy.ndarray) – (N, 3) integer-indexed triples.
- **entity_idx** (dict) – Entity-name → index mapping.
- **relation_idx** (dict) – Relation-name → index mapping.
- **form** (str) – 'EntityPrediction' or 'RelationPrediction'.
- **label_smoothing_rate** (float, optional) – Label smoothing coefficient (default 0.0).

train_data = None

train_target = None

label_smoothing_rate

collate_fn = None

__len__()

__getitem__(idx)

```
class dicee.dataset_classes.KvsSampleDataset (train_set_idx: numpy.ndarray, entity_idx,  
    relation_idx, form, store=None, neg_ratio=None, label_smoothing_rate: float = 0.0)
```

Bases: torch.utils.data.Dataset

Dataset for KvsSample training (dynamic multi-label).

Like **KvsAll** but sub-samples the target vector at each access to keep mini-batch sizes manageable when the entity set is large.

Parameters

- **train_set_idx** (numpy.ndarray) – (N, 3) integer-indexed triples.
- **entity_idx** (dict) – Entity-name → index mapping.
- **relation_idx** (dict) – Relation-name → index mapping.
- **form** (str) – 'EntityPrediction'.
- **neg_ratio** (int) – Number of negative samples per positive target.

- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data = None

train_target = None

neg_ratio = None

num_entities

label_smoothing_rate

collate_fn = None

max_num_of_classes

__len__()

__getitem__(*idx*)

class dicee.dataset_classes.**OnevsAllDataset** (*train_set_idx: numpy.ndarray, entity_idx*s)

Bases: torch.utils.data.Dataset

Dataset for the 1-vs-All training strategy (multi-class).

Each sample is a (*head, relation*) pair with a one-hot target vector whose single active position corresponds to the true tail entity.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (*N, 3*) integer-indexed triples.
- **entity_idx**s (*dict*) – Entity-name → index mapping (used to determine the target dimension).

train_data

target_dim

collate_fn = None

__len__()

__getitem__(*idx*)

class dicee.dataset_classes.**LiteralDataset** (*file_path: str, ent_idx: dict = None, normalization_type: str = 'z-norm', sampling_ratio: float = None, loader_backend: str = 'pandas'*)

Bases: torch.utils.data.Dataset

Dataset for loading and processing literal data for Literal Embedding models.

Handles loading, normalization, and preparation of (*entity, attribute, value*) triples. Supports z-score and min-max normalization as well as optional sub-sampling for ablation studies.

Parameters

- **file_path** (*str*) – Path to the training data file (CSV / TSV / RDF).
- **ent_idx** (*dict*) – Entity-name → index mapping.
- **normalization_type** (*str, optional*) – 'z-norm', 'min-max', or None (default 'z-norm').

- **sampling_ratio** (*float or None, optional*) – Fraction of the training set to keep (default `None` = use all).
- **loader_backend** (*str, optional*) – 'pandas' or 'rdflib' (default 'pandas').

train_file_path

loader_backend = 'pandas'

normalization_type = 'z-norm'

normalization_params

sampling_ratio = `None`

entity_to_idx = `None`

num_entities

__getitem__ (*index*)

__len__ ()

static load_and_validate_literal_data (*file_path: str = None, loader_backend: str = 'pandas'*)

→ `pandas.DataFrame`

Load and validate a literal data file.

Parameters

- **file_path** (*str*) – Path to the data file.
- **loader_backend** (*str*) – 'pandas' or 'rdflib'.

Returns

Three-column `DataFrame` with columns `head`, `attribute`, `value`.

Return type

`pandas.DataFrame`

static denormalize (*preds_norm, attributes, normalization_params*) → `numpy.ndarray`

Reverse the normalization applied during training.

Parameters

- **preds_norm** (*numpy.ndarray*) – Normalized predictions.
- **attributes** (*list*) – Attribute names corresponding to each prediction.
- **normalization_params** (*dict*) – Parameters stored during training.

Returns

Denormalized predictions.

Return type

`numpy.ndarray`

class `dicee.dataset_classes.FixedNegSampleDataset` (*train_set: numpy.ndarray, num_entities: int, num_relations: int, neg_sample_ratio: int = 1, label_smoothing_rate: float = 0.0, seed: int = None*)

Bases: `torch.utils.data.Dataset`

Pre-computed (fixed) negative sampling dataset.

At construction time every positive triple is paired with one random negative (head- or tail-corrupted) using vectorized operations for efficiency. The pairs are stored so that `__getitem__` is a simple lookup.

This is useful when you want deterministic negatives across epochs (e.g., for reproducibility or debugging).

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int, optional*) – Number of negative samples per positive triple (default 1).
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

neg_sample_ratio = 1

num_entities

num_relations

label_smoothing_rate = 0.0

collate_fn = None

seed = None

train_triples

length

__len__() → int

__getitem__(*idx: int*) → Tuple[torch.Tensor, torch.Tensor]

```
class dicee.dataset_classes.OnevsSample (train_set: numpy.ndarray, num_entities: int,  
    num_relations: int, neg_sample_ratio: int = None, label_smoothing_rate: float = 0.0)
```

Bases: torch.utils.data.Dataset

Dataset for 1-vs-Sample training (dynamic multi-class with negatives).

For every positive triple (*h*, *r*, *t*) the dataset draws **neg_sample_ratio** random entities as negatives and returns a label vector that marks the true tail and the negatives.

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int*) – Number of negative samples per positive.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data

num_entities

num_relations

```

neg_sample_ratio = None

label_smoothing_rate

collate_fn = None

__len__()

__getitem__(idx)

class dicee.dataset_classes.TriplePredictionDataset (train_set: numpy.ndarray,
            num_entities: int, num_relations: int, neg_sample_ratio: int = 1, label_smoothing_rate: float = 0.0,
            seed: int = None)

Bases: torch.utils.data.Dataset

Dataset for triple prediction with on-the-fly negative sampling.

Each item is a single positive triple; the custom collate_fn generates a batch of mixed positive and negative
triples.

Parameters
    • train_set (numpy.ndarray) – (N, 3) integer-indexed triples.
    • num_entities (int) – Total number of entities.
    • num_relations (int) – Total number of relations.
    • neg_sample_ratio (int, optional) – Number of negative samples per positive triple
      (default 1).
    • label_smoothing_rate (float, optional) – Label smoothing coefficient (default 0.
      0).

label_smoothing_rate

neg_sample_ratio

seed = None

train_set

length

num_entities

num_relations

__len__()

__getitem__(idx)

collate_fn (batch: List[torch.Tensor])

```

dicee.eval_static_funcs

Static evaluation functions for KGE models.

This module provides backward compatibility by re-exporting from the new `dicee.evaluation` module.

Deprecated since version Use: `dicee.evaluation` submodules instead. This module will be removed in a future version.

Functions

<code>evaluate_ensemble_link_prediction_performance</code>	Evaluate link prediction performance of an ensemble of KGE models.
<code>evaluate_link_prediction_performance(→ Dict[str, float])</code>	Evaluate link prediction performance with head and tail prediction.
<code>evaluate_link_prediction_performance_with_head_and_tail</code>	Evaluate link prediction with BPE encoding (head and tail).
<code>evaluate_link_prediction_performance_with_head_and_tail_and_reciprocals</code>	Evaluate link prediction with BPE encoding and reciprocals.
<code>evaluate_link_prediction_performance_with_reciprocal_relations</code>	Evaluate link prediction with reciprocal relations.
<code>evaluate_lp_bpe_k_vs_all(→ Dict[str, float])</code>	Evaluate BPE link prediction with KvsAll scoring.
<code>evaluate_literal_prediction(→ Optional[pandas.DataFrame])</code>	Evaluate trained literal prediction model on a test file.

Module Contents

`dicee.eval_static_funcs.evaluate_ensemble_link_prediction_performance` (*models: List, triples, er_vocab: Dict[Tuple, List], weights: List[float] | None = None, batch_size: int = 512, weighted_averaging: bool = True, normalize_scores: bool = True*) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is True but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
```

(continues on next page)


```

...     weights=weights, weighted_averaging=True
... )
>>> print (f"MRR: {results['MRR']:.4f}")

```

`dicee.eval_static_funcs.evaluate_link_prediction_performance(model, triples, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.eval_static_funcs.evaluate_link_prediction_performance_with_bpe(model, within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.eval_static_funcs.evaluate_link_prediction_performance_with_bpe_reciprocals(model, within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.eval_static_funcs.evaluate_link_prediction_performance_with_reciprocals (
    model, triples, er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.eval_static_funcs.evaluate_lp_bpe_k_vs_all (model, triples: List[List[str]],
    er_vocab: Dict | None = None, batch_size: int | None = None,
    func_triple_to_bpe_representation: Callable | None = None,
    str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

```
dicee.eval_static_funcs.evaluate_literal_prediction (kge_model, eval_file_path: str = None,
    store_lit_preds: bool = True, eval_literals: bool = True, loader_backend: str = 'pandas',
    return_attr_error_metrics: bool = False) → pandas.DataFrame | None
```

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is `True`, otherwise `None`.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

dicee.evaluation

Evaluation module for knowledge graph embedding models.

This module provides comprehensive evaluation capabilities for KGE models, including link prediction, literal prediction, and ensemble evaluation.

Modules:

`link_prediction`: Functions for evaluating link prediction performance
`literal_prediction`: Functions for evaluating literal/attribute prediction
`ensemble`: Functions for ensemble model evaluation
`evaluator`: Main Evaluator class for integrated evaluation
`utils`: Shared utility functions for evaluation
`_filtering`: Internal helpers for filtered ranking (not exported)

Example

```
>>> from dicee.evaluation import Evaluator
>>> from dicee.evaluation.link_prediction import evaluate_link_prediction_performance
>>> from dicee.evaluation.ensemble import evaluate_ensemble_link_prediction_
→ performance
```

Submodules

dicee.evaluation.ensemble

Ensemble evaluation functions.

This module provides functions for evaluating ensemble models, including weighted averaging and score normalization.

Functions

<code>evaluate_ensemble_link_prediction_performance</code>	Evaluate link prediction performance of an ensemble of KGE models.
--	--

Module Contents

`dicee.evaluation.ensemble.evaluate_ensemble_link_prediction_performance` (*models*: List, *triples*, *er_vocab*: Dict[Tuple, List], *weights*: List[float] | None = None, *batch_size*: int = 512, *weighted_averaging*: bool = True, *normalize_scores*: bool = True) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is True but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

`dicee.evaluation.evaluator`

Main Evaluator class for KGE model evaluation.

This module provides the Evaluator class which orchestrates evaluation of knowledge graph embedding models across different datasets and scoring techniques.

Attributes

<code>VALID_SCORING_TECHNIQUES</code>

Classes

<code>Evaluator</code>

Evaluator class for KGE models in various downstream tasks.

Module Contents

`dicee.evaluation.evaluator.VALID_SCORING_TECHNIQUES`

class `dicee.evaluation.evaluator.Evaluator` (*args*, *is_continual_training*: *bool* = *False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

```

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation(dataset) → None
    Prepare vocabularies from the dataset for evaluation.
    Resolves any future objects and saves vocabularies to disk.

```

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

```

eval(dataset, trained_model, form_of_labelling: str, during_training: bool = False) → Dict | None
    Evaluate the trained model on the dataset.

```

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

```

eval_rank_of_head_and_tail_entity(*, train_set, valid_set=None, test_set=None, trained_model)
    → None

```

Evaluate with negative sampling scoring.

```

eval_rank_of_head_and_tail_byte_pair_encoded_entity(*, train_set=None, valid_set=None,
    test_set=None, ordered_bpe_entities, trained_model) → None

```

Evaluate with BPE-encoded entities and negative sampling.

```

eval_with_byte(*, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model,
    form_of_labelling) → None

```

Evaluate Byte model with generation.

```

eval_with_bpe_vs_all(*, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model,
    form_of_labelling) → None

```

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, train_set, valid_set=None, test_set=None, trained_model, form_of_labelling)
→ None

Evaluate with KvsAll or lvsAll scoring.

evaluate_lp_k_vs_all (model, triple_idx, info: str | None = None,
form_of_labelling: str | None = None) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp_with_byte (model, triples: List[List[str]], info: str | None = None) → Dict[str, float]

Evaluate BytE model with text generation.

Parameters

- **model** – BytE model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (model, triples: List[List[str]], info: str | None = None,
form_of_labelling: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (model, triple_idx, info: str) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: str) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: numpy.ndarray, *form_of_labelling*: str)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

dicee.evaluation.link_prediction

Link prediction evaluation functions.

This module provides various functions for evaluating link prediction performance of knowledge graph embedding models.

Functions

<code>evaluate_link_prediction_performance(→ Dict[str, float])</code>	Evaluate link prediction performance with head and tail prediction.
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with reciprocal relations.
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with BPE encoding and reciprocals.
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with BPE encoding (head and tail).
<code>evaluate_lp(→ Dict[str, float])</code>	Evaluate link prediction with batched processing.
<code>evaluate_bpe_lp(→ Dict[str, float])</code>	Evaluate link prediction with BPE-encoded entities.
<code>evaluate_lp_bpe_k_vs_all(→ Dict[str, float])</code>	Evaluate BPE link prediction with KvsAll scoring.

Module Contents

`dicee.evaluation.link_prediction.evaluate_link_prediction_performance` (*model*, *triples*, *er_vocab*: Dict[Tuple, List], *re_vocab*: Dict[Tuple, List]) → Dict[str, float]

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.

- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.
```

```
evaluate_link_prediction_performance_with_reciprocals (model, triples,  
er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.
```

```
evaluate_link_prediction_performance_with_bpe_reciprocals (model,  
within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_link_prediction_performance_with_bpe (  
model, within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List],  
re_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_lp(model, triple_idx, num_entities: int,
      er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List], info: str = 'Eval Starts',
      batch_size: int = 128, chunk_size: int = 1000) → Dict[str, float]
```

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.
- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_bpe_lp(model, triple_idx: List[Tuple],
      all_bpe_shaped_entities, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List],
      info: str = 'Eval Starts') → Dict[str, float]
```

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]],
      er_vocab: Dict | None = None, batch_size: int | None = None,
      func_triple_to_bpe_representation: Callable | None = None,
      str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.

- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

dicee.evaluation.literal_prediction

Literal prediction evaluation functions.

This module provides functions for evaluating literal/attribute prediction performance of knowledge graph embedding models.

Functions

<code>evaluate_literal_prediction(→ Optional[pandas.DataFrame])</code>	Evaluate trained literal prediction model on a test file.
--	---

Module Contents

`dicee.evaluation.literal_prediction.evaluate_literal_prediction(kge_model, eval_file_path: str = None, store_lit_preds: bool = True, eval_literals: bool = True, loader_backend: str = 'pandas', return_attr_error_metrics: bool = False) → pandas.DataFrame | None`

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if return_attr_error_metrics is True, otherwise None.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

dicee.evaluation.utils

Utility functions for evaluation module.

This module contains shared helper functions used across different evaluation components.

Attributes

<code>DEFAULT_HITS_RANGE</code>
<code>ALL_HITS_RANGE</code>

Functions

<code>make_iterable_verbose(→ Iterable)</code>	Wrap an iterable with tqdm progress bar if verbose is True.
<code>compute_metrics_from_ranks(→ Dict[str, float])</code>	Compute standard link prediction metrics from ranks.
<code>compute_metrics_from_ranks_simple(→ Dict[str, float])</code>	Compute link prediction metrics without scaling factor.
<code>update_hits(→ None)</code>	Update hits dictionary based on rank.
<code>create_hits_dict(→ Dict[int, List[float]])</code>	Create an initialized hits dictionary.
<code>efficient_zero_grad(→ None)</code>	Efficiently zero gradients using <code>parameter.grad = None</code> .

Module Contents

`dicee.evaluation.utils.DEFAULT_HITS_RANGE: List[int] = [1, 3, 10]`

`dicee.evaluation.utils.ALL_HITS_RANGE: List[int]`

`dicee.evaluation.utils.make_iterable_verbose(iterable_object: Iterable, verbose: bool, desc: str = 'Default', position: int | None = None, leave: bool = True) → Iterable`

Wrap an iterable with tqdm progress bar if verbose is True.

Parameters

- **iterable_object** – The iterable to potentially wrap.
- **verbose** – Whether to show progress bar.
- **desc** – Description for the progress bar.
- **position** – Position of the progress bar.

- **leave** – Whether to leave the progress bar after completion.

Returns

The original iterable or a tqdm-wrapped version.

```
dicee.evaluation.utils.compute_metrics_from_ranks (ranks: List[int], num_triples: int,
hits_dict: Dict[int, List[float]], scale_factor: int = 1) → Dict[str, float]
```

Compute standard link prediction metrics from ranks.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.
- **scale_factor** – Factor to scale the denominator (e.g., 2 for head+tail).

Returns

Dictionary containing H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.utils.compute_metrics_from_ranks_simple (ranks: List[int], num_triples: int,
hits_dict: Dict[int, List[float]]) → Dict[str, float]
```

Compute link prediction metrics without scaling factor.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.

Returns

Dictionary containing H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.utils.update_hits (hits: Dict[int, List[float]], rank: int,
hits_range: List[int] | None = None) → None
```

Update hits dictionary based on rank.

Parameters

- **hits** – Dictionary to update in-place.
- **rank** – The rank to check against hit levels.
- **hits_range** – List of hit levels to check (default: ALL_HITS_RANGE).

```
dicee.evaluation.utils.create_hits_dict (hits_range: List[int] | None = None)
→ Dict[int, List[float]]
```

Create an initialized hits dictionary.

Parameters

hits_range – List of hit levels to initialize (default: ALL_HITS_RANGE).

Returns

Dictionary with empty lists for each hit level.

```
dicee.evaluation.utils.efficient_zero_grad (model) → None
```

Efficiently zero gradients using `parameter.grad = None`.

This is more efficient than `optimizer.zero_grad()` as it avoids memory operations.

See: https://pytorch.org/tutorials/recipes/recipes/tuning_guide.html

Parameters

model – PyTorch model to zero gradients for.

Classes

<i>Evaluator</i>	Evaluator class for KGE models in various downstream tasks.
------------------	---

Functions

<i>evaluate_ensemble_link_prediction_performance</i>	Evaluate link prediction performance of an ensemble of KGE models.
<i>evaluate_bpe_lp</i> (→ Dict[str, float])	Evaluate link prediction with BPE-encoded entities.
<i>evaluate_link_prediction_performance</i> (→ Dict[str, float])	Evaluate link prediction performance with head and tail prediction.
<i>evaluate_link_prediction_performance_with_j</i>	Evaluate link prediction with BPE encoding (head and tail).
<i>evaluate_link_prediction_performance_with_j</i>	Evaluate link prediction with BPE encoding and reciprocals.
<i>evaluate_link_prediction_performance_with_</i>	Evaluate link prediction with reciprocal relations.
<i>evaluate_lp</i> (→ Dict[str, float])	Evaluate link prediction with batched processing.
<i>evaluate_lp_bpe_k_vs_all</i> (→ Dict[str, float])	Evaluate BPE link prediction with KvsAll scoring.
<i>evaluate_literal_prediction</i> (→ Optional[pandas.DataFrame])	Evaluate trained literal prediction model on a test file.

Package Contents

`dicee.evaluation.evaluate_ensemble_link_prediction_performance` (*models*: List, *triples*, *er_vocab*: Dict[Tuple, List], *weights*: List[float] | None = None, *batch_size*: int = 512, *weighted_averaging*: bool = True, *normalize_scores*: bool = True) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is `True` but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

class `dicee.evaluation.Evaluator` (*args*, *is_continual_training*: *bool* = *False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

```

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation(dataset) → None
    Prepare vocabularies from the dataset for evaluation.
    Resolves any future objects and saves vocabularies to disk.

```

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

```

eval(dataset, trained_model, form_of_labelling: str, during_training: bool = False) → Dict | None
    Evaluate the trained model on the dataset.

```

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

```

eval_rank_of_head_and_tail_entity(*, train_set, valid_set=None, test_set=None, trained_model)
    → None

```

Evaluate with negative sampling scoring.

```

eval_rank_of_head_and_tail_byte_pair_encoded_entity(*, train_set=None, valid_set=None,
    test_set=None, ordered_bpe_entities, trained_model) → None

```

Evaluate with BPE-encoded entities and negative sampling.

```

eval_with_byte(*, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model,
    form_of_labelling) → None

```

Evaluate Byte model with generation.

```

eval_with_bpe_vs_all(*, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model,
    form_of_labelling) → None

```

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, train_set, valid_set=None, test_set=None, trained_model, form_of_labelling)
→ None

Evaluate with KvsAll or lvsAll scoring.

evaluate_lp_k_vs_all (model, triple_idx, info: str | None = None,
form_of_labelling: str | None = None) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

evaluate_lp_with_byte (model, triples: List[List[str]], info: str | None = None) → Dict[str, float]

Evaluate BytE model with text generation.

Parameters

- **model** – BytE model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (model, triples: List[List[str]], info: str | None = None,
form_of_labelling: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

evaluate_lp (model, triple_idx, info: str) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: *str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: *numpy.ndarray*, *form_of_labelling*: *str*) → Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

`dicee.evaluation.evaluate_bpe_lp` (*model*, *triple_idx*: List[Tuple], *all_bpe_shaped_entities*, *er_vocab*: Dict[Tuple, List], *re_vocab*: Dict[Tuple, List], *info*: *str* = 'Eval Starts') → Dict[str, float]

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, [H@10](#), and MRR metrics.

`dicee.evaluation.evaluate_link_prediction_performance` (*model*, *triples*, *er_vocab*: Dict[Tuple, List], *re_vocab*: Dict[Tuple, List]) → Dict[str, float]

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance_with_bpe(model,
    within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List],
    re_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance_with_bpe_reciprocals(model,
    within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance_with_reciprocals(model, triples,
    er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_lp(model, triple_idx, num_entities: int, er_vocab: Dict[Tuple, List],
    re_vocab: Dict[Tuple, List], info: str = 'Eval Starts', batch_size: int = 128, chunk_size: int = 1000)
    → Dict[str, float]
```

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.
- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]],
    er_vocab: Dict | None = None, batch_size: int | None = None,
    func_triple_to_bpe_representation: Callable | None = None,
    str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

```
dicee.evaluation.evaluate_literal_prediction(kge_model, eval_file_path: str = None,
    store_lit_preds: bool = True, eval_literals: bool = True, loader_backend: str = 'pandas',
    return_attr_error_metrics: bool = False) → pandas.DataFrame | None
```

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is `True`, otherwise `None`.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

`dicee.evaluator`

Evaluator module for knowledge graph embedding models.

This module provides backward compatibility by re-exporting from the new `dicee.evaluation` module.

Deprecated since version Use: `dicee.evaluation.Evaluator` instead. This module will be removed in a future version.

Classes

Evaluator

Evaluator class for KGE models in various downstream tasks.

Module Contents

class `dicee.evaluator.Evaluator` (*args, is_continual_training: bool = False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset*, *trained_model*, *form_of_labelling*: str, *during_training*: bool = False) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').

- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set*, *valid_set=None*, *test_set=None*, *trained_model*)
→ None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set=None*, *valid_set=None*,
test_set=None, *ordered_bpe_entities*, *trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, *raw_train_set*, *raw_valid_set=None*, *raw_test_set=None*, *trained_model*,
form_of_labelling) → None

Evaluate BytE model with generation.

eval_with_bpe_vs_all (*, *raw_train_set*, *raw_valid_set=None*, *raw_test_set=None*, *trained_model*,
form_of_labelling) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, *train_set*, *valid_set=None*, *test_set=None*, *trained_model*, *form_of_labelling*)
→ None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all (*model*, *triple_idx*, *info: str | None = None*,
form_of_labelling: str | None = None) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp_with_byte (*model*, *triples: List[List[str]]*, *info: str | None = None*) → Dict[str, float]

Evaluate BytE model with text generation.

Parameters

- **model** – BytE model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (*model*, *triples: List[List[str]]*, *info: str | None = None*,
form_of_labelling: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.

- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

evaluate_lp (*model*, *triple_idx*, *info*: *str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: *str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: *numpy.ndarray*, *form_of_labelling*: *str*)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

dicee.executer

Executor module for training, retraining and evaluating KGE models.

This module provides the Execute and ContinuousExecute classes for managing the full lifecycle of knowledge graph embedding model training.

Classes

<i>Execute</i>	Executor class for training, retraining and evaluating KGE models.
<i>ContinuousExecute</i>	A subclass of Execute Class for retraining

Module Contents

class `dicee.executer.Execute` (*args*, *continuous_training: bool = False*)

Executor class for training, retraining and evaluating KGE models.

Handles the complete workflow: 1. Loading & Preprocessing & Serializing input data 2. Training & Validation & Testing 3. Storing all necessary information

args

Processed input arguments.

distributed

Whether distributed training is enabled.

rank

Process rank in distributed training.

world_size

Total number of processes.

local_rank

Local GPU rank.

trainer

Training handler instance.

trained_model

The trained model after training completes.

knowledge_graph

The loaded knowledge graph.

report

Dictionary storing training metrics and results.

evaluator

Model evaluation handler.

distributed

rank

world_size

local_rank

args

is_continual_training = False

trainer: `dicee.trainer.DICE_Trainer` | None = None

trained_model = None

knowledge_graph: *dicee.knowledge_graph.KG* | None = None

report: Dict

evaluator: *dicee.evaluator.Evaluator* | None = None

start_time: float | None = None

is_local_rank_zero() → bool

cleanup()

setup_executor() → None
 Set up storage directories for the experiment.
 Creates or reuses experiment directories based on configuration. Saves the configuration to a JSON file.

create_and_store_kg() → None
 Create knowledge graph and store as memory-mapped file.
 Only executed on local rank 0 in distributed training. Skips if memmap already exists.

load_from_memmap() → None
 Load knowledge graph from memory-mapped file.

save_trained_model() → None
 Save a knowledge graph embedding model
 (1) Send model to eval mode and cpu.
 (2) Store the memory footprint of the model.
 (3) Save the model into disk.
 (4) Update the stats of KG again ?

Parameter

rtype
 None

end(*form_of_labelling: str*) → dict
 End training
 (1) Store trained model.
 (2) Report runtimes.
 (3) Eval model if required.

Parameter

rtype
 A dict containing information about the training and/or evaluation

write_report() → None
 Report training related information in a report.json file

start() → dict
 Start training
 # (1) Loading the Data # (2) Create an evaluator object. # (3) Create a trainer object. # (4) Start the training

Parameter

rtype

A dict containing information about the training and/or evaluation

class dicee.executer.**ContinuousExecute** (*args*)

Bases: *Execute*

A subclass of Execute Class for retraining

- (1) Loading & Preprocessing & Serializing input data.
- (2) Training & Validation & Testing
- (3) Storing all necessary info

During the continual learning we can only modify * **num_epochs** * parameter. Trained model stored in the same folder as the seed model for the training. Trained model is noted with the current time.

continual_start () → dict

Start Continual Training

- (1) Initialize training.
- (2) Start continual training.
- (3) Save trained model.

Parameter

rtype

A dict containing information about the training and/or evaluation

dicee.knowledge_graph

Knowledge Graph module for data loading and preprocessing.

Provides the KG class for handling knowledge graph data including loading, preprocessing, and indexing operations.

Classes

<i>KG</i>	Knowledge Graph container and processor.
-----------	--

Module Contents

```
class dicee.knowledge_graph.KG (dataset_dir: str | None = None, byte_pair_encoding: bool = False,  
    padding: bool = False, add_noise_rate: float | None = None, sparql_endpoint: str | None = None,  
    path_single_kg: str | None = None, path_for_deserialization: str | None = None,  
    add_reciprocal: bool | None = None, eval_model: str | None = None,  
    read_only_few: int | None = None, sample_triples_ratio: float | None = None,  
    path_for_serialization: str | None = None, entity_to_idx: Dict | None = None,  
    relation_to_idx: Dict | None = None, backend: str | None = None,  
    training_technique: str | None = None, separator: str | None = None)
```

Knowledge Graph container and processor.

Handles loading, preprocessing, and indexing of knowledge graph data from various sources including files, SPARQL endpoints, and serialized formats.

dataset_dir
Path to directory containing train/valid/test files.

num_entities
Total number of unique entities.

num_relations
Total number of unique relations.

train_set
Indexed training triples as numpy array.

valid_set
Indexed validation triples (optional).

test_set
Indexed test triples (optional).

entity_to_idx
Mapping from entity strings to indices.

relation_to_idx
Mapping from relation strings to indices.

dataset_dir = None

sparql_endpoint = None

path_single_kg = None

byte_pair_encoding = False

ordered_shaped_bpe_tokens = None

add_noise_rate = None

num_entities: int | None = None

num_relations: int | None = None

path_for_deserialization = None

add_reciprocal = None

eval_model = None

read_only_few = None

sample_triples_ratio = None

path_for_serialization = None

entity_to_idx = None

relation_to_idx = None

backend = 'pandas'

training_technique = None

```

separator = None

raw_train_set = None

raw_valid_set = None

raw_test_set = None

train_set = None

valid_set = None

test_set = None

idx_entity_to_bpe_shaped: Dict

enc

num_tokens

num_bpe_entities: int | None = None

padding = False

dummy_id

max_length_subword_tokens: int | None = None

train_set_target = None

target_dim: int | None = None

train_target_indices = None

ordered_bpe_entities = None

description_of_input = None

describe() → None
    Generate a description string of the dataset statistics.

property entities_str: List[str]
    Get list of all entity strings.

property relations_str: List[str]
    Get list of all relation strings.

exists(h: str, r: str, t: str) → bool
    Check if a triple exists in the training set.

```

Parameters

- **h** – Head entity string.
- **r** – Relation string.
- **t** – Tail entity string.

Returns

True if the triple exists, False otherwise.

```

__iter__() → Iterator[Tuple[str, str, str]]
    Iterate over training triples as string tuples.

__len__() → int
    Return number of triples in the raw training set.

func_triple_to_bpe_representation(triple: List[str])

```

dicee.knowledge_graph_embeddings

Classes

<i>KGE</i>	Knowledge Graph Embedding Class for interactive usage of pre-trained models
------------	---

Module Contents

```

class dicee.knowledge_graph_embeddings.KGE (path=None, url=None, construct_ensemble=False,
      model_name=None)

    Bases:
        dicee.abstracts.BaseInteractiveKGE,
        dicee.abstracts.InteractiveQueryDecomposition,
        dicee.abstracts.BaseInteractiveTrainKGE

    Knowledge Graph Embedding Class for interactive usage of pre-trained models

    __str__()

    to(device: str) → None

    get_transductive_entity_embeddings(indices: torch.LongTensor | List[str], as_pytorch=False,
      as_numpy=False, as_list=True) → torch.FloatTensor | numpy.ndarray | List[float]

    create_vector_database(collection_name: str, distance: str, location: str = 'localhost',
      port: int = 6333)

    generate(h="", r="")

    eval_lp_performance(dataset=List[Tuple[str, str, str]], filtered=True)

    predict_missing_head_entity(relation: List[str] | str, tail_entity: List[str] | str, within=None,
      batch_size=2, topk=1, return_indices=False) → Tuple

    Given a relation and a tail entity, return top k ranked head entity.

     $\operatorname{argmax}_{\{e \in E\}} f(e, r, t)$ , where  $r \in R$ ,  $t \in E$ .

```

Parameter

relation: Union[List[str], str]
 String representation of selected relations.

tail_entity: Union[List[str], str]
 String representation of selected entities.

k: int
 Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_relations (*head_entity: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False*) → Tuple

Given a head entity and a tail entity, return top k ranked relations.

$\text{argmax}_{\{r \in R\}} f(h,r,t)$, where $h, t \in E$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_tail_entity (*head_entity: List[str] | str, relation: List[str] | str, within: List[str] = None, batch_size=2, topk=1, return_indices=False*) → torch.FloatTensor

Given a head entity and a relation, return top k ranked entities

$\text{argmax}_{\{e \in E\}} f(h,r,e)$, where $h \in E$ and $r \in R$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

Returns: Tuple

scores

predict (*, *h: List[str] | str | None = None, r: List[str] | str | None = None, t: List[str] | str | None = None, within: List[str] | None = None, logits: bool = True*) → torch.FloatTensor

Predict scores for triples or missing triple elements.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.
- **t** – Tail entity/entities. None to predict tails.
- **within** – Optional list of entities to restrict predictions to.
- **logits** – If True, return raw scores. If False, return sigmoid scores (0-1).

Returns

- Single triple (h, r, t): scalar score
- Missing element: vector of all possible scores

Return type

torch.FloatTensor of scores. Shape depends on the query type

Raises

AssertionError – If inputs are not strings or lists of strings.

Examples

```
>>> # Score a specific triple
>>> model.predict(h="Mongolia", r="isLocatedIn", t="Asia", logits=False)
tensor(0.9523)
```

```
>>> # Get scores for all possible tail entities
>>> model.predict(h="Mongolia", r="isLocatedIn", t=None)
tensor([0.21, 0.95, 0.03, ...]) # One score per entity
```

predict_topk (*, h: str | List[str] | None = None, r: str | List[str] | None = None, t: str | List[str] | None = None, topk: int = 10, within: List[str] | None = None, batch_size: int = 1024) → List[Tuple[str, float]] | List[List[Tuple[str, float]]]

Predict top-k missing items in a given triple pattern.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.
- **t** – Tail entity/entities. None to predict tails.
- **topk** – Number of top predictions to return.
- **within** – Optional list of entities to restrict predictions to.
- **batch_size** – Batch size for processing multiple queries.

Returns

List[(item, score), ...] of length topk. For batch query: List of such lists, one per query.

Return type

For single query

Raises

- **AssertionError** – If more than one of h, r, t is None.
- **AssertionError** – If the required arguments for a query type are None.

Examples

```
>>> model.predict_topk(h=["Mongolia"], r=["isLocatedIn"], topk=3)
[('Asia', 0.99), ('Europe', 0.02), ...]
```

```
>>> model.predict_topk(r=["isLocatedIn"], t=["Asia"], topk=5)
[('Mongolia', 0.85), ('China', 0.82), ...]
```


triple_score (*h: List[str] | str = None, r: List[str] | str = None, t: List[str] | str = None, logits=False*)
→ torch.FloatTensor
Predict triple score

Parameter

head_entity: List[str]
String representation of selected entities.

relation: List[str]
String representation of selected relations.

tail_entity: List[str]
String representation of selected entities.

logits: bool
If logits is True, unnormalized score returned

Returns: Tuple

pytorch tensor of triple score

return_multi_hop_query_results (*aggregated_query_for_all_entities, k: int, only_scores*)

single_hop_query_answering (*query: tuple, only_scores: bool = True, k: int = None, use_logits: bool = True*)

answer_multi_hop_query (*query_type: str | None = None, query: Tuple[str | Tuple[str, str], Ellipsis] | None = None, queries: List[Tuple[str | Tuple[str, str], Ellipsis]] | None = None, tnorm: str = 'prod', neg_norm: str = 'standard', lambda_: float = 0.0, k: int = 10, only_scores: bool = False, use_logits: bool = True*)
→ List[Tuple[str, torch.Tensor]] | List[List[Tuple[str, torch.Tensor]]]

Answer multi-hop EPFO (Existential Positive First-Order) queries.

Supports 9 query types: 1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up. See docs/guides/multi_hop_queries.md for detailed query patterns.

Parameters

- **query_type** – Query pattern name. One of: - 1p: (e, (r,)) # One-hop - 2p: (e, (r1, r2)) # Two-hop - 3p: (e, (r1, r2, r3)) # Three-hop - 2i: ((e1, (r1,)), (e2, (r2,))) # Two-way intersection - 3i: ((e1, (r1,)), (e2, (r2,)), (e3, (r3,))) # Three-way intersection - ip: (((e1, (r1,)), (e2, (r2,))), (r3,)) # Intersection + projection - pi: ((e, (r1, r2)), (r3,)) # Projection + intersection (2i meets 2p) - 2u: ((e1, (r1,)), (e2, (r2,))) # Two-way union - up: ((e, (r1, r2)), (e, (r3,))) # Union + projection
- **query** – Single query tuple matching the query_type pattern.
- **queries** – Batch of queries. If provided, query must be None.
- **tnorm** – T-norm for intersection/union. Options: “prod”, “min”.
- **neg_norm** – Negation norm. Options: “standard”, “sugeno”, “yager”.
- **lambda** – Parameter for sugeno and yager negation (0.0-1.0).
- **k** – Number of top answer entities to return.

- **only_scores** – If True, return only scores tensor. If False, return (entity, score) tuples.
- **use_logits** – If True, use raw model logits. If False, use sigmoid probabilities.

Returns

List[(entity, score), ...] of top-k answers. For batch queries: List of such lists, one per query.

Return type

For single query

Raises

- **ValueError** – If query_type is not in {1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up}.
- **AssertionError** – If query structure doesn't match query_type pattern.

Examples

```
>>> # 1p: Find entities located in Asia
>>> model.answer_multi_hop_query(
...     query_type="1p",
...     query=("Asia", ("isLocatedIn",)),
...     k=5
... )
[("Mongolia", 0.92), ("China", 0.89), ...]
```

```
>>> # 2p: Two-hop query (e.g., "capital of countries in Europe")
>>> model.answer_multi_hop_query(
...     query_type="2p",
...     query=("Europe", ("isLocatedIn", "hasCapital")),
...     k=3
... )
[("Paris", 0.85), ("Berlin", 0.82), ...]
```

```
>>> # 2i: Intersection query
>>> model.answer_multi_hop_query(
...     query_type="2i",
...     query=((("Asia", ("isLocatedIn",)), ("Mountains", ("hasGeography",))),
...     k=5
... )
[("Nepal", 0.78), ("Tibet", 0.65), ...]
```

See also

- docs/guides/multi_hop_queries.md: Complete guide with all query patterns
- tests/test_answer_multi_hop_query.py: Usage examples

find_missing_triples (confidence: float, entities: List[str] = None, relations: List[str] = None, topk: int = 10, at_most: int = sys.maxsize) → Set

Find missing triples

Iterative over a set of entities E and a set of relation R :

orall e in E and orall r in R f(e,r,x)

```

    Return (e,r,x)
otin G and f(e,r,x) > confidence
confidence: float
A threshold for an output of a sigmoid function given a triple.
topk: int
Highest ranked k item to select triples with f(e,r,x) > confidence .
at_most: int
Stop after finding at_most missing triples
{(e,r,x) | f(e,r,x) > confidence land (e,r,x)
otin G
predict_literals (entity: List[str] | str = None, attribute: List[str] | str = None,
denormalize_preds: bool = True) → numpy.ndarray
Predicts literal values for given entities and attributes.

```

Parameters

- **entity** (*Union[List[str], str]*) – Entity or list of entities to predict literals for.
- **attribute** (*Union[List[str], str]*) – Attribute or list of attributes to predict literals for.
- **denormalize_preds** (*bool*) – If True, denormalizes the predictions.

Returns

Predictions for the given entities and attributes.

Return type

numpy ndarray

dicee.models

Submodules

dicee.models.adopt

ADOPT Optimizer Implementation.

This module implements the ADOPT (Adaptive Optimization with Precise Tracking) algorithm, an advanced optimization method for training neural networks.

ADOPT Overview:

ADOPT is an adaptive learning rate optimization algorithm that combines the benefits of momentum-based methods with per-parameter learning rate adaptation. Unlike Adam, which applies momentum to raw gradients, ADOPT normalizes gradients first and then applies momentum, leading to more stable training dynamics.

Key Features: - Gradient normalization before momentum application - Adaptive per-parameter learning rates - Optional gradient clipping that grows with training steps - Support for decoupled weight decay (AdamW-style) - Multiple execution modes: single-tensor, multi-tensor (foreach), and fused (planned)

Algorithm Comparison:

Adam: $m = \beta_1 * m + (1 - \beta_1) * g$, $\theta = \theta - \alpha * m / \sqrt{v}$ ADOPT: $m = \beta_1 * m + (1 - \beta_1) * g / \sqrt{v}$, $\theta = \theta - \alpha * m$

The key difference is that ADOPT normalizes gradients before momentum, which provides better stability and can lead to improved convergence.

Classes:

- ADOPT: Main optimizer class (extends torch.optim.Optimizer)

Functions:

- adopt: Functional API for ADOPT algorithm computation
- _single_tensor_adopt: Single-tensor implementation (TorchScript compatible)
- _multi_tensor_adopt: Multi-tensor implementation using foreach operations

Performance:

- Single-tensor: Default, compatible with torch.jit.script
- Multi-tensor (foreach): 2-3x faster on GPU through vectorization
- Fused (planned): Would provide maximum performance via specialized kernels

Example:

```
>>> import torch
>>> from dicee.models.adopt import ADOPT
>>>
>>> model = torch.nn.Linear(10, 1)
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01, decouple=True)
>>>
>>> # Training loop
>>> for epoch in range(num_epochs):
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     optimizer.step()
```

References:

Original implementation: <https://github.com/iShohei220/adopt>

Notes:

This implementation is based on the original ADOPT implementation and adapted to work with the PyTorch optimizer interface and the dicee framework.

Classes

<code>ADOPT</code>	ADOPT Optimizer.
--------------------	------------------

Functions

<code>adopt(params, grads, exp_avgs, exp_avg_sqs, state_steps)</code>	Functional API that performs ADOPT algorithm computation.
---	---

Module Contents

```
class dicee.models.adopt.ADOPT (params: torch.optim.optimizer.ParamsT,  
    lr: float | torch.Tensor = 0.001, betas: Tuple[float, float] = (0.9, 0.9999), eps: float = 1e-06,  
    clip_lambda: Callable[[int], float] | None = lambda step: ..., weight_decay: float = 0.0,  
    decouple: bool = False, *, foreach: bool | None = None, maximize: bool = False,  
    capturable: bool = False, differentiable: bool = False, fused: bool | None = None)
```

Bases: torch.optim.optimizer.Optimizer

ADOPT Optimizer.

ADOPT is an adaptive learning rate optimization algorithm that combines momentum-based updates with adaptive per-parameter learning rates. It uses exponential moving averages of gradients and squared gradients, with gradient clipping for stability.

The algorithm performs the following key operations: 1. Normalizes gradients by the square root of the second moment estimate 2. Applies optional gradient clipping based on the training step 3. Updates parameters using momentum-smoothed normalized gradients 4. Supports decoupled weight decay (AdamW-style) or L2 regularization

Mathematical formulation:

$$m_t = \beta_1 * m_{t-1} + (1 - \beta_1) * \text{clip}(g_t / \sqrt{v_t}) \quad v_t = \beta_2 * v_{t-1} + (1 - \beta_2) * g_t^2 \quad \theta_t = \theta_{t-1} - \alpha * m_t$$

where:

- θ_t : parameter at step t
- g_t : gradient at step t
- m_t : first moment estimate (momentum)
- v_t : second moment estimate (variance)
- α : learning rate
- β_1, β_2 : exponential decay rates
- $\text{clip}()$: optional gradient clipping function

Reference:

Original implementation: <https://github.com/iShohei220/adopt>

Parameters

- **params** (*ParamsT*) – Iterable of parameters to optimize or dicts defining parameter groups.
- **lr** (*float or Tensor, optional*) – Learning rate. Can be a float or 1-element Tensor. Default: 1e-3

- **betas** (*Tuple[float, float], optional*) – Coefficients (β_1, β_2) for computing running averages of gradient and its square. β_1 controls momentum, β_2 controls variance. Default: (0.9, 0.9999)
- **eps** (*float, optional*) – Term added to denominator to improve numerical stability. Default: 1e-6
- **clip_lambda** (*Callable[[int], float], optional*) – Function that takes the step number and returns the gradient clipping threshold. Common choices: - lambda step: step**0.25 (default, gradually increases clipping threshold) - lambda step: 1.0 (constant clipping) - None (no clipping) Default: lambda step: step**0.25
- **weight_decay** (*float, optional*) – Weight decay coefficient (L2 penalty). Default: 0.0
- **decouple** (*bool, optional*) – If True, uses decoupled weight decay (AdamW-style), applying weight decay directly to parameters. If False, adds weight decay to gradients (L2 regularization). Default: False
- **foreach** (*bool, optional*) – If True, uses the faster foreach implementation for multi-tensor operations. Default: None (auto-select)
- **maximize** (*bool, optional*) – If True, maximizes parameters instead of minimizing. Useful for reinforcement learning. Default: False
- **capturable** (*bool, optional*) – If True, the optimizer is safe to capture in a CUDA graph. Requires learning rate as Tensor. Default: False
- **differentiable** (*bool, optional*) – If True, the optimization step can be differentiated. Useful for meta-learning. Default: False
- **fused** (*bool, optional*) – If True, uses fused kernel implementation (currently not supported). Default: None

Raises

- **ValueError** – If learning rate, epsilon, betas, or weight_decay are invalid.
- **RuntimeError** – If fused is enabled (not currently supported).
- **RuntimeError** – If lr is a Tensor with foreach=True and capturable=False.

Example

```
>>> # Basic usage
>>> optimizer = ADOPT(model.parameters(), lr=0.001)
>>> optimizer.zero_grad()
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With decoupled weight decay
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01,
↳decouple=True)
```

```
>>> # Custom gradient clipping
>>> optimizer = ADOPT(model.parameters(), clip_lambda=lambda step: max(1.0,
↳step**0.5))
```

Note

- For most use cases, the default hyperparameters work well
- Consider using `decouple=True` for better generalization (similar to AdamW)
- The `clip_lambda` function helps stabilize training in early steps

`clip_lambda`

`__setstate__` (*state*)

Restore optimizer state from a checkpoint.

This method handles backward compatibility when loading optimizer state from older versions. It ensures all required fields are present with default values and properly converts step counters to tensors if needed.

Key responsibilities: 1. Set default values for newly added hyperparameters 2. Convert old-style scalar step counters to tensor format 3. Place step tensors on appropriate devices based on capturable/fused modes

Parameters

state (*dict*) – Optimizer state dictionary (typically from `torch.load()`).

Note

- This enables loading checkpoints saved with older ADOPT versions
- Step counters are converted to appropriate device/dtype for compatibility
- Capturable and fused modes require step tensors on parameter devices

`step` (*closure=None*)

Perform a single optimization step.

This method executes one iteration of the ADOPT optimization algorithm across all parameter groups. It orchestrates the following workflow:

1. Optionally evaluates a closure to recompute the loss (useful for algorithms like LBFGS or when loss needs multiple evaluations)
2. For each parameter group: - Collects parameters with gradients and their associated state - Extracts hyperparameters (betas, learning rate, etc.) - Calls the functional `adopt()` API to perform the actual update
3. Returns the loss value if a closure was provided

The functional API (`adopt()`) handles three execution modes: - Single-tensor: Updates one parameter at a time (default, JIT-compatible) - Multi-tensor (foreach): Batches operations for better performance - Fused: Uses fused CUDA kernels (not yet implemented)

Gradient scaling support: This method is compatible with automatic mixed precision (AMP) training. It can access `grad_scale` and `found_inf` attributes for gradient unscaling and inf/nan detection when used with `GradScaler`.

Parameters

closure (*Callable, optional*) – A callable that reevaluates the model and returns the loss. The closure should: - Enable gradients (`torch.enable_grad()`) - Compute forward pass - Compute loss - Compute backward pass - Return the loss value Example: `lambda: (loss := model(x), loss.backward(), loss)[-1]` Default: None

Returns

The loss value returned by the closure, or None if no closure was provided.

Return type

Optional[Tensor]

Example

```
>>> # Standard usage
>>> loss = criterion(model(input), target)
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With closure (e.g., for line search)
>>> def closure():
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     return loss
>>> loss = optimizer.step(closure)
```

Note

- Call `zero_grad()` before computing gradients for the next step
- CUDA graph capture is checked for safety when `capturable=True`
- The method is thread-safe for different parameter groups

`dicee.models.adopt.adopt` (*params*: List[torch.Tensor], *grads*: List[torch.Tensor],
exp_avgs: List[torch.Tensor], *exp_avg_sqs*: List[torch.Tensor], *state_steps*: List[torch.Tensor],
foreach: bool | None = None, *capturable*: bool = False, *differentiable*: bool = False,
fused: bool | None = None, *grad_scale*: torch.Tensor | None = None,
found_inf: torch.Tensor | None = None, *has_complex*: bool = False, *, *beta1*: float, *beta2*: float,
lr: float | torch.Tensor, *clip_lambda*: Callable[[int], float] | None, *weight_decay*: float,
decouple: bool, *eps*: float, *maximize*: bool)

Functional API that performs ADOPT algorithm computation.

This is the main functional interface for the ADOPT optimization algorithm. It dispatches to one of three implementations based on the execution mode:

1. **Single-tensor mode** (default): Updates parameters one at a time - Compatible with `torch.jit.script` - More flexible but slower - Used when `foreach=False` or automatically for small models
2. **Multi-tensor (foreach) mode**: Batches operations across tensors - 2-3x faster on GPU through vectorization - Groups tensors by device/dtype automatically - Used when `foreach=True`
3. **Fused mode**: Uses specialized fused kernels (not yet implemented) - Would provide maximum performance - Currently raises `RuntimeError` if enabled

Algorithm overview (ADOPT):

ADOPT adapts learning rates per-parameter while using momentum on normalized gradients. The key innovation is normalizing gradients before momentum, which provides more stable training than standard Adam.

Mathematical formulation:

```
# Normalize gradient by its historical variance  $\text{normed\_g\_t} = \text{g\_t} / \sqrt{v\_t + \epsilon}$ 
# Optional gradient clipping for stability  $\text{normed\_g\_t} = \text{clip}(\text{normed\_g\_t}, \text{threshold}(t))$ 
# Momentum on normalized gradients (key difference from Adam)  $m\_t = \beta_1 * m_{t-1} + (1 - \beta_1) * \text{normed\_g\_t}$ 
# Parameter update  $\theta\_t = \theta_{t-1} - \alpha * m\_t$ 
# Update variance estimate  $v\_t = \beta_2 * v_{t-1} + (1 - \beta_2) * \text{g\_t}^2$ 
```

where:

- θ : parameters
- g : gradients
- m : first moment (momentum of normalized gradients)
- v : second moment (variance of raw gradients)
- α : learning rate
- β_1, β_2 : exponential decay rates
- ϵ : numerical stability constant
- $\text{clip}()$: gradient clipping function based on step

Automatic mode selection:

When `foreach` and `fused` are both `None` (default), the function automatically selects the best implementation based on:

- Parameter types and devices
- Whether differentiable mode is enabled
- Learning rate type (float vs Tensor)
- Capturable mode requirements

```
param params
    Parameters to optimize.

type params
    List[Tensor]

param grads
    Gradients for each parameter.

type grads
    List[Tensor]

param exp_avgs
    First moment estimates (momentum).

type exp_avgs
    List[Tensor]

param exp_avg_sqs
    Second moment estimates (variance).

type exp_avg_sqs
    List[Tensor]
```

param state_steps
Step counters (must be singleton tensors).

type state_steps
List[Tensor]

param foreach
Whether to use multi-tensor implementation. None: auto-select based on configuration (default).

type foreach
Optional[bool]

param capturable
If True, ensure CUDA graph capture safety.

type capturable
bool

param differentiable
If True, allow gradients through optimization step.

type differentiable
bool

param fused
If True, use fused kernels (not implemented).

type fused
Optional[bool]

param grad_scale
Gradient scaler for AMP training.

type grad_scale
Optional[Tensor]

param found_inf
Flag for inf/nan detection in AMP.

type found_inf
Optional[Tensor]

param has_complex
Whether any parameters are complex-valued.

type has_complex
bool

param beta1
Exponential decay rate for first moment (momentum). Typical range: 0.9-0.95.

type beta1
float

param beta2
Exponential decay rate for second moment (variance). Typical range: 0.999-0.9999 (higher than Adam).

type beta2
float

param lr
Learning rate. Can be a scalar Tensor for dynamic learning rate with capturable=True.

type lr
Union[float, Tensor]

param clip_lambda
Function that maps step number to gradient clipping threshold. None disables clipping.

type clip_lambda
Optional[Callable[[int], float]]

param weight_decay
Weight decay coefficient (L2 penalty).

type weight_decay
float

param decouple
If True, use decoupled weight decay (AdamW-style). Recommended for better generalization.

type decouple
bool

param eps
Small constant for numerical stability in normalization.

type eps
float

param maximize
If True, maximize objective instead of minimize.

type maximize
bool

raises RuntimeError
If torch.jit.script is used with foreach or fused.

raises RuntimeError
If state_steps contains non-tensor elements.

raises RuntimeError
If fused=True (not yet implemented).

raises RuntimeError
If lr is Tensor with foreach=True and capturable=False.

Example

```
>>> # Typically called by ADOPT optimizer, not directly
>>> adopt (
...     params=[p1, p2],
...     grads=[g1, g2],
...     exp_avgs=[m1, m2],
...     exp_avg_sqs=[v1, v2],
...     state_steps=[step1, step2],
...     beta1=0.9,
...     beta2=0.9999,
...     lr=0.001,
...     clip_lambda=lambda s: s**0.25,
...     weight_decay=0.01,
...     decouple=True,
```

(continues on next page)

(continued from previous page)

```
...     eps=1e-6,  
...     maximize=False,  
... )
```

Note

- For distributed training, this API is compatible with torch/distributed/optim
- The foreach mode is generally preferred for GPU training
- Complex parameters are handled transparently by viewing as real
- First optimization step only initializes variance, doesn't update parameters

See also

- ADOPT class: High-level optimizer interface
- `_single_tensor_adopt`: Single-tensor implementation details
- `_multi_tensor_adopt`: Multi-tensor implementation details

dicee.models.base_model

Classes

<i>BaseKGELightning</i>	Thin PyTorch Lightning wrapper shared by all KGE models.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>IdentityClass</i>	No-op normalisation / dropout placeholder.

Module Contents

class `dicee.models.base_model.BaseKGELightning` (*args, **kwargs)

Bases: `lightning.LightningModule`

Thin PyTorch Lightning wrapper shared by all KGE models.

Provides the standard Lightning training loop hooks (`training_step`, `on_train_epoch_end`, `configure_optimizers`) as well as a helper for reporting model size. All concrete KGE models should extend *BaseKGE* rather than this class directly.

training_step_outputs = []

mem_of_model() → Dict

Size of model in MB and number of params

training_step(batch, batch_idx=None)

Execute one optimisation step for the given mini-batch.

Handles two- and three-element batches produced by the different dataset classes (*KvsAll* / *NegSample* vs. *KvsSample*).

Parameters

- **batch** (*tuple*) – (*x*, *y*) for standard scoring, or (*x*, *y_select*, *y*) for sample-based labelling.
- **batch_idx** (*int*, *optional*) – Index of the current batch (unused, kept for Lightning API compat).

Returns

Scalar loss value for this batch.

Return type

`torch.FloatTensor`

loss_function (*yhat_batch*: `torch.FloatTensor`, *y_batch*: `torch.FloatTensor`) → `torch.FloatTensor`

Compute the loss between model predictions and targets.

Delegates to `self.loss` which is configured in `BaseKGE.__init__` based on the scoring technique (BCEWithLogitsLoss for entity/relation prediction, CrossEntropyLoss for classification).

Parameters

- **yhat_batch** (`torch.FloatTensor`) – Model output scores, shape (batch_size, *).
- **y_batch** (`torch.FloatTensor`) – Ground-truth labels of the same shape as *yhat_batch*.

Returns

Scalar loss value.

Return type

`torch.FloatTensor`

on_train_epoch_end (*args, **kwargs)

Called in the training loop at the very end of the epoch.

To access all batch outputs at the end of the epoch, you can cache step outputs as an attribute of the `LightningModule` and access them in this hook:

```
class MyLightningModule(L.LightningModule):
    def __init__(self):
        super().__init__()
        self.training_step_outputs = []

    def training_step(self):
        loss = ...
        self.training_step_outputs.append(loss)
        return loss

    def on_train_epoch_end(self):
        # do something with all training_step outputs, for example:
        epoch_mean = torch.stack(self.training_step_outputs).mean()
        self.log("training_epoch_mean", epoch_mean)
        # free up the memory
        self.training_step_outputs.clear()
```

test_epoch_end (outputs: `List[Any]`)

test_dataloader () → `None`

An iterable or collection of iterables specifying test samples.

For more information about multiple dataloaders, see this section.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

 Warning

do not assign state in `prepare_data`

- `test()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

 Note

If you don't need a test dataset and a `test_step()`, you don't need to implement this method.

`val_dataloader()` → None

An iterable or collection of iterables specifying validation samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `fit()`
- `validate()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

 Note

If you don't need a validation dataset and a `validation_step()`, you don't need to implement this method.

`predict_dataloader()` → None

An iterable or collection of iterables specifying prediction samples.

For more information about multiple dataloaders, see this section.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `predict()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Returns

A `torch.utils.data.DataLoader` or a sequence of them specifying prediction samples.

`train_dataloader()` → None

An iterable or collection of iterables specifying training samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

 Warning

do not assign state in `prepare_data`

- `fit()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

configure_optimizers (*parameters=None*)

Instantiate and return the optimiser for training.

The optimiser type is taken from `self.optimizer_name` which is set in `BaseKGE.init_params_with_sanity_checking()` from the `--optim` argument. Supported values: 'SGD', 'Adam', 'Adopt', 'AdamW', 'NAdam', 'Adagrad', 'ASGD', 'Muon'.

Parameters

parameters (*iterable, optional*) – Model parameters to optimise. Defaults to `self.parameters()` when `None`.

Returns

The configured optimiser instance.

Return type

`torch.optim.Optimizer`

class `dicee.models.base_model.BaseKGE` (*args: dict*)

Bases: `BaseKGELightning`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`


```

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

```

forward_byte_pair_encoded_k_vs_all (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*) → *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x: torch.LongTensor* | *Tuple[torch.LongTensor, torch.LongTensor]*,
y_idx: torch.LongTensor = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample()*
- (batch, 3) *tensor* → *forward_triples()*
- (batch, 2) *tensor* → *forward_k_vs_all()*
- BPE triple *tensor* → *forward_byte_pair_encoded_triple()*
- BPE pair *tensor* → *forward_byte_pair_encoded_k_vs_all()*

Parameters

- **x** (*torch.LongTensor* or *Tuple[torch.LongTensor, torch.LongTensor]*) – Either a plain index tensor or a (triple_idx, target_idx) tuple for sample-based labelling.
- **y_idx** (*torch.LongTensor*, *optional*) – Target entity indices used by *forward_k_vs_sample()*. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x: torch.LongTensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample(*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation(idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation(indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation(x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x*: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (torch.LongTensor) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (numpy.ndarray) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (numpy.ndarray) – Shape (num_relations, embedding_dim).

class dicee.models.base_model.IdentityClass (*args=None*)

Bases: torch.nn.Module

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (--normalization None). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = None

__call__ (*x*)

static forward (*x*)

dicee.models.clifford

Classes

<i>Keci</i>	Keci: Knowledge Graph Embedding via Clifford Algebra.
<i>KeciTransformer</i>	Keci with Transformer architecture.
<i>TransformerBlock</i>	A single transformer block with self-attention and MLP.
<i>TransformerSelfAttention</i>	Multi-head self-attention for the Keci Transformer.
<i>TransformerMLP</i>	MLP for the Keci Transformer.
<i>CKeci</i>	Without learning dimension scaling

continues on next page

Table 26 – continued from previous page

DeCaL

Base class for all Knowledge Graph Embedding models.

Module Contents

class `dicee.models.clifford.Keci` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Keci: Knowledge Graph Embedding via Clifford Algebra.

Embeds entities and relations as multi-vectors in the Clifford algebra $Cl_{\{p,q\}}(\mathbb{R}^d)$ and scores triples via the Clifford product. The algebra is parameterised by two non-negative integers p and q :

- $p = 0, q = 0$ — reduces to a standard bilinear (DistMult-like) model.
- $p = 0, q = 1$ — equivalent to ComplEx.
- Larger p and q capture higher-order geometric interactions.

The embedding dimension must satisfy `embedding_dim % (p + q + 1) == 0`; the resulting quotient is stored as `self.r`.

Parameters

args (*dict*) – Configuration dictionary. Recognised keys (beyond those in `BaseKGE`): p (int, default 0) and q (int, default 0).

References

Demir et al., *Clifford Embeddings — A Generalized Approach for Embedding in Normed Algebras*, ECML 2023.

name = 'Keci'

p

q

r

requires_grad_for_interactions = True

compute_sigma_pp (*hp, rp*)

Compute $\sigma_{\{pp\}} = \sum_{i=1}^{p-1} \sum_{k=i+1}^p (h_{i r k} - h_{k r i}) e_i e_k$

$\sigma_{\{pp\}}$ captures the interactions between along p bases. For instance, let p, e_1, e_2, e_3 , we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$. This can be implemented with a nested two for loops

```
results = []
for i in range(p - 1):
```

```
    for k in range(i + 1, p):
```

```
        results.append((hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i]))
```

```
sigma_pp = torch.stack(results, dim=2)
assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

```
e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3
```

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_qq (*hq, rq*)

Compute $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_j r_k - h_k r_j) e_j e_k$ σ_{qq} captures the interactions between along q bases. For instance, let $q = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$. This can be implemented with a nested two for loops

```
results = []
for j in range(q - 1):
```

```
    for k in range(j + 1, q):
```

```
        results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])
```

```
sigma_qq = torch.stack(results, dim=2)
assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

```
e2e1, e2e2, e2e3, e3e1, e3e2, e3e3
```

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_pq (**, hp, hq, rp, rq*)

```
sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j
```

```
results = []
sigma_pq = torch.zeros(b, r, p, q)
for i in range(p):
```

```
    for j in range(q):
```

```
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
```

```
print(sigma_pq.shape)
```

apply_coefficients (*hp, hq, rp, rq*)

Multiplying a base vector with its scalar coefficient

clifford_multiplication (*h0, hp, hq, r0, rp, rq*)

Compute our CL multiplication

$$h = h_0 + \sum_{i=1}^p h_i e_i + \sum_{j=p+1}^{p+q} h_j e_j$$

$$r = r_0 + \sum_{i=1}^p r_i e_i + \sum_{j=p+1}^{p+q} r_j e_j$$

$$e_i^2 = +1 \text{ for } i \leq p, e_j^2 = -1 \text{ for } p < j \leq p+q, e_i e_j = -e_j e_i \text{ for } i$$

e_j

$$h r = \sigma_0 + \sigma_p + \sigma_q + \sigma_{pp} + \sigma_q + \sigma_{pq} \text{ where}$$

(1) $\sigma_0 = h_0 r_0 + \sum_{i=1}^p (h_0 r_i - h_i r_0) e_i - \sum_{j=p+1}^{p+q} (h_j r_j) e_j$

(2) $\sigma_p = \sum_{i=1}^p (h_0 r_i + h_i r_0) e_i$

(3) $\sigma_q = \sum_{j=p+1}^{p+q} (h_0 r_j + h_j r_0) e_j$

(4) $\sigma_{pp} = \sum_{i=1}^p \sum_{k=i+1}^p (h_i r_k - h_k r_i) e_i e_k$

(5) $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_j r_k - h_k r_j) e_j e_k$

(6) $\sigma_{pq} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$

construct_cl_multivector (*x: torch.FloatTensor, r: int, p: int, q: int*)

→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Split a flat embedding vector into the three Clifford components.

Given an embedding x of dimension $d = r + r \cdot p + r \cdot q$, returns the scalar part a_0 , the p -blade part a_p , and the q -blade part a_q .

Parameters

- x (*torch.FloatTensor*) – Shape (batch_size, d).

- $\mathbf{r}$ (*int*) – Scalar block size ($\text{embedding_dim} // (\mathbf{p} + \mathbf{q} + 1)$).
- $\mathbf{p}$ (*int*) – Number of positive-signature basis elements.
- $\mathbf{q}$ (*int*) – Number of negative-signature basis elements.

Returns

- $\mathbf{a0}$ (*torch.FloatTensor*) – Shape ($\text{batch_size}, \mathbf{r}$) — scalar (grade-0) part.
- $\mathbf{ap}$ (*torch.FloatTensor*) – Shape ($\text{batch_size}, \mathbf{r}, \mathbf{p}$) — positive-blade part.
- $\mathbf{aq}$ (*torch.FloatTensor*) – Shape ($\text{batch_size}, \mathbf{r}, \mathbf{q}$) — negative-blade part.

forward_k_vs_with_explicit (x : *torch.Tensor*) $\rightarrow$ *torch.FloatTensor*

KvsAll scoring using an explicit loop over $\sigma_{pp}/\sigma_{qq}/\sigma_{pq}$ terms.

Functionally equivalent to `forward_k_vs_all()` but computes the higher-order interaction terms (σ_{pp} , σ_{qq} , σ_{pq}) with explicit nested loops rather than einsum contractions. Kept for reference and correctness verification.

Parameters

$\mathbf{x}$ (*torch.Tensor*) – Shape ($\text{batch_size}, 2$) integer tensor [$\text{head_idx}, \text{relation_idx}$].

Returns

Shape ($\text{batch_size}, \text{num_entities}$) score matrix.

Return type

torch.FloatTensor

k_vs_all_score (bpe_head_ent_emb : *torch.FloatTensor*, bpe_rel_ent_emb : *torch.FloatTensor*, E : *torch.FloatTensor*) $\rightarrow$ *torch.FloatTensor*

Compute Clifford-product scores for a head/relation batch vs. all entities.

Decomposes the head-entity and relation embeddings into Clifford multi-vectors, performs the $\text{Cl}_{\{p,q\}}$ product, and inner-products the result against the entity embedding matrix E .

Parameters

- bpe_head_ent_emb (*torch.FloatTensor*) – Head-entity embeddings, shape ($\text{batch_size}, \text{embedding_dim}$).
- bpe_rel_ent_emb (*torch.FloatTensor*) – Relation embeddings, shape ($\text{batch_size}, \text{embedding_dim}$).
- E (*torch.FloatTensor*) – All entity embeddings, shape ($\text{num_entities}, \text{embedding_dim}$).

Returns

Shape ($\text{batch_size}, \text{num_entities}$) score matrix.

Return type

torch.FloatTensor

forward_k_vs_all (x : *torch.Tensor*) $\rightarrow$ *torch.FloatTensor*

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations $\mathbb{R}^d$.
- (2) Construct head entity and relation embeddings according to $\text{Cl}_{\{p,q\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

`forward_k_vs_with_explicit` and this functions are identical Parameter ——— `x`: `torch.LongTensor` with (n,2) shape `rtype`: `torch.FloatTensor` with (n, **IEI**) shape

construct_batch_selected_cl_multivector (*x*: *torch.FloatTensor*, *r*: *int*, *p*: *int*, *q*: *int*)
→ tuple[`torch.FloatTensor`, `torch.FloatTensor`, `torch.FloatTensor`]

Split a batched, *k*-selected embedding tensor into Clifford components.

A variant of `construct_cl_multivector()` for tensors that have an extra *k* dimension (e.g. when scoring against *k* sampled targets).

Parameters

- **x** (*torch.FloatTensor*) – Shape (batch_size, *k*, *d*).
- **r** (*int*) – Scalar block size.
- **p** (*int*) – Number of positive-signature basis elements.
- **q** (*int*) – Number of negative-signature basis elements.

Returns

- **a0** (*torch.FloatTensor*) – Shape (batch_size, *k*, *r*).
- **ap** (*torch.FloatTensor*) – Shape (batch_size, *k*, *r*, *p*).
- **aq** (*torch.FloatTensor*) – Shape (batch_size, *k*, *r*, *q*).

forward_k_vs_sample (*x*: *torch.LongTensor*, *target_entity_idx*: *torch.LongTensor*) → `torch.FloatTensor`

Parameter

`x`: `torch.LongTensor` with (n,2) shape

`target_entity_idx`: `torch.LongTensor` with (n, *k*) shape *k* denotes the selected number of examples.

rtype

`torch.FloatTensor` with (n, *k*) shape

score (*h*, *r*, *t*)

forward_triples (*x*: *torch.Tensor*) → `torch.FloatTensor`

Parameter

`x`: `torch.LongTensor` with (n,3) shape

rtype

`torch.FloatTensor` with (n) shape

class `dicee.models.clifford.KeciTransformer` (*args*)

Bases: `Keci`

Keci with Transformer architecture.

Concatenates *h0*, *hp*, *hq*, *r0*, *rp*, *rq* into a single embedding vector and processes through transformer.

name = 'KeciTransformer'

use_clifford_mul

seq_len = 1

transformer

lm_head

forward_k_vs_all (*x*: *torch.Tensor*) → *torch.FloatTensor*

Kvsall training

Parameter

x: *torch.LongTensor* with (n,2) shape

rtype

torch.FloatTensor with (n, **IEI**) shape

class *dicee.models.clifford.TransformerBlock* (*n_embd*, *n_head*, *dropout=0.0*, *bias=False*)

Bases: *torch.nn.Module*

A single transformer block with self-attention and MLP.

ln_1

attn

ln_2

mlp

forward (*x*)

class *dicee.models.clifford.TransformerSelfAttention* (*n_embd*, *n_head*, *dropout=0.0*,
bias=False)

Bases: *torch.nn.Module*

Multi-head self-attention for the Keci Transformer.

c_attn

c_proj

attn_dropout

resid_dropout

n_head

n_embd

dropout = 0.0

forward (*x*)

class *dicee.models.clifford.TransformerMLP* (*n_embd*, *dropout=0.0*, *bias=False*)

Bases: *torch.nn.Module*

MLP for the Keci Transformer.

c_fc

gelu

c_proj

dropout

forward(*x*)

class dicee.models.clifford.**CKeci**(*args*)

Bases: *Keci*

Without learning dimension scaling

name = 'CKeci'

requires_grad_for_interactions = **False**

class dicee.models.clifford.**DeCaL**(*args*)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches *forward()* calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (*h*, *r*, *t*) triples.
- *forward_k_vs_all()* — score a (*h*, *r*) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by *vars(argparse.Namespace)*. Required keys: *embedding_dim*, *num_entities*, *num_relations*, *learning_rate* (or *lr*), *optim*, *scoring_technique*.

name = 'DeCaL'

entity_embeddings

relation_embeddings

p

q

r

re

forward_triples (*x: torch.Tensor*) → *torch.FloatTensor*

Parameter

x: *torch.LongTensor* with (*n*,) shape

rtype

torch.FloatTensor with (*n*) shape

cl_pqr (*a: torch.tensor*) → *torch.tensor*

Input: *tensor(batch_size, emb_dim)* —> output: *tensor* with *1+p+q+r* components with size (*batch_size*, *emb_dim/(1+p+q+r)*) each.

1) takes a tensor of size (*batch_size*, *emb_dim*), split it into *1 + p + q + r* components, hence *1+p+q+r* must be a divisor of the *emb_dim*. 2) Return a list of the *1+p+q+r* components vectors, each are tensors of size (*batch_size*, *emb_dim/(1+p+q+r)*)

compute_sigmas_single (*list_h_emb, list_r_emb, list_t_emb*)

here we compute all the sums with no others vectors interaction taken with the scalar product with t, that is,

$$s0 = h_0 r_0 t_0 s1 = \sum_{i=1}^p h_i r_i t_0 s2 = \sum_{j=p+1}^{p+q} h_j r_j t_0 s3 = \sum_{i=1}^q (h_0 r_i t_i + h_i r_0 t_i) s4 = \sum_{i=p+1}^{p+q} (h_0 r_i t_i + h_i r_0 t_i) s5 = \sum_{i=p+q+1}^{p+q+r} (h_0 r_i t_i + h_i r_0 t_i)$$

and return:

$$sigma_0 t = \sigma_0 \cdot t_0 = s0 + s1 - s2 s3, s4 \text{ and } s5$$

compute_sigmas_multivect (*list_h_emb, list_r_emb*)

Here we compute and return all the sums with vectors interaction for the same and different bases.

For same bases vectors interaction we have

$$\sigma_p p = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (h_i r_{i'} - h_{i'} r_i) (model \text{ the interactions between } e_i \text{ and } e_{i'} \text{ for } 1 \leq i, i' \leq p) \sigma_q q = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (h_j r_{j'} - h_{j'} r_j)$$

For different base vector interactions, we have

$$\sigma_p q = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) (interactions \text{ between } e_i \text{ and } e_j \text{ for } 1 \leq i \leq p \text{ and } p+1 \leq j \leq p+q) \sigma_p r = \sum_{i=1}^p \sum_{j=p+q+1}^{p+q+r} (h_i r_j - h_j r_i)$$

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q,r\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— x: torch.LongTensor with (n,) shape :rtype: torch.FloatTensor with (n, **IEI**) shape

apply_coefficients (*h0, hp, hq, hk, r0, rp, rq, rk*)

Multiplying a base vector with its scalar coefficient

construct_cl_multivector (*x: torch.FloatTensor, re: int, p: int, q: int, r: int*)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Construct a batch of multivectors $Cl_{\{p,q,r\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,d) shape

returns

- **a0** (*torch.FloatTensor*)
- **ap** (*torch.FloatTensor*)
- **aq** (*torch.FloatTensor*)
- **ar** (*torch.FloatTensor*)

compute_sigma_pp (*hp, rp*)

Compute .. math:

$$\sigma_{p,p} = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (x_i y_{i'} - x_{i'} y_i)$$

σ_{pp} captures the interactions between along p bases For instance, let p e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3, and e_2 e_3 This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_qq (*hq, rq*)

Compute

$$\sigma_{q,q}^* = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (x_j y_{j'} - x_{j'} y_j) Eq.16$$

σ_{qq} captures the interactions between along q bases For instance, let q e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3, and e_2 e_3 This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_rr (*hk, rk*)

$$\sigma_{r,r}^* = \sum_{k=p+q+1}^{p+q+r-1} \sum_{k'=k+1}^p (x_k y_{k'} - x_{k'} y_k)$$

compute_sigma_pq (*, *hp, hq, rp, rq*)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

```

print(sigma_pq.shape)
compute_sigma_pr(*, hp, hk, rp, rk)
Compute

```

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

```

results = []
sigma_pq = torch.zeros(b, r, p, q)
for i in range(p):
    for j in range(q):
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
print(sigma_pq.shape)
compute_sigma_qr(*, hq, hk, rq, rk)

```

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

```

results = []
sigma_pq = torch.zeros(b, r, p, q)
for i in range(p):
    for j in range(q):
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
print(sigma_pq.shape)

```

dicее.models.complex

Classes

<i>ConEx</i>	Convolutional ComplEx Knowledge Graph Embeddings
<i>AConEx</i>	Additive Convolutional ComplEx Knowledge Graph Embeddings
<i>ComplEx</i>	Base class for all Knowledge Graph Embedding models.

Module Contents

```

class dicее.models.complex.ConEx(args)
    Bases: dicее.models.base_model.BaseKGE
    Convolutional ComplEx Knowledge Graph Embeddings
    name = 'ConEx'
    conv2d
    fc_num_input
    fc1
    norm_fc1
    bn_conv2d
    feature_map_dropout

```

residual_convolution (*C_1*: *Tuple[torch.Tensor, torch.Tensor]*,
C_2: *Tuple[torch.Tensor, torch.Tensor]*) → *torch.FloatTensor*

Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

forward_k_vs_all (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: *torch.Tensor*, *target_entity_idx*: *torch.Tensor*)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

class `dicee.models.complex.AConEx` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Additive Convolutional ComplEx Knowledge Graph Embeddings

name = 'AConEx'

conv2d

fc_num_input

fc1

norm_fc1

bn_conv2d

feature_map_dropout

residual_convolution (*C_1*: *Tuple[torch.Tensor, torch.Tensor]*,
C_2: *Tuple[torch.Tensor, torch.Tensor]*) → *torch.FloatTensor*

Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

forward_k_vs_all (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: *torch.Tensor*, *target_entity_idx*: *torch.Tensor*)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

class `dicee.models.complex.Complex` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

```

name = 'ComplEx'

static score(head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor,
            tail_ent_emb: torch.FloatTensor)

static k_vs_all_score(emb_h: torch.FloatTensor, emb_r: torch.FloatTensor,
                    emb_E: torch.FloatTensor)

```

Parameters

- `emb_h`
- `emb_r`
- `emb_E`

forward_k_vs_all (*x*: torch.LongTensor) → torch.FloatTensor

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: torch.LongTensor, *target_entity_idx*: torch.LongTensor)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, *k*) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

dicee.models.dualE

Classes

<i>DualE</i>	Dual Quaternion Knowledge Graph Embeddings (https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657)
--------------	---

Module Contents

class dicee.models.dualE.**DualE** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Dual Quaternion Knowledge Graph Embeddings (<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

name = 'DualE'

entity_embeddings

relation_embeddings

num_ent = None

kvsall_score (*e_1_h, e_2_h, e_3_h, e_4_h, e_5_h, e_6_h, e_7_h, e_8_h, e_1_t, e_2_t, e_3_t, e_4_t, e_5_t, e_6_t, e_7_t, e_8_t, r_1, r_2, r_3, r_4, r_5, r_6, r_7, r_8*) → torch.tensor

KvsAll scoring function

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_triples (*idx_triple: torch.tensor*) → torch.tensor

Negative Sampling forward pass:

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_k_vs_all (*x*)

KvsAll forward pass

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

T (*x: torch.tensor*) → torch.tensor

Transpose function

Input: Tensor with shape (nxm) Output: Tensor with shape (mxn)

dicee.models.ensemble

Classes

EnsembleKGE

Module Contents

class dicee.models.ensemble.**EnsembleKGE** (*models: list = None, seed_model=None, pretrained_models: List = None*)

```

name

train_mode = True

args

named_children()

property example_input_array

parameters()

modules()

__iter__()

__len__()

eval()

to(device)

state_dict()
    Return the state dict of the ensemble.

load_state_dict(state_dict, strict=True)
    Load the state dict into the ensemble.

mem_of_model()

__call__(x_batch)

step()

get_embeddings()

__str__()

```

dicee.models.function_space

Classes

<i>FMult</i>	Learning Knowledge Neural Graphs
<i>GFMult</i>	Learning Knowledge Neural Graphs
<i>FMult2</i>	Learning Knowledge Neural Graphs
<i>LFMult1</i>	Embedding with trigonometric functions. We represent all entities and relations in the complex number space as:
<i>LFMult</i>	Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:

Module Contents

```

class dicee.models.function_space.FMult(args)
    Bases: dicee.models.base_model.BaseKGE
    Learning Knowledge Neural Graphs

```

```

name = 'FMult'

entity_embeddings

relation_embeddings

k

num_sample = 50

gamma

roots

weights

compute_func(weights: torch.FloatTensor, x) → torch.FloatTensor

chain_func(weights, x: torch.FloatTensor)

forward_triples(idx_triple: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

```

class dicee.models.function_space.**GFMult**(args)

Bases: *[dicee.models.base_model.BaseKGE](#)*

Learning Knowledge Neural Graphs

```

name = 'GFMult'

entity_embeddings

relation_embeddings

k

num_sample = 250

roots

weights

compute_func(weights: torch.FloatTensor, x) → torch.FloatTensor

chain_func(weights, x: torch.FloatTensor)

forward_triples(idx_triple: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

```

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

```
class dicee.models.function_space.FMult2(args)
```

Bases: *[dicee.models.base_model.BaseKGE](#)*

Learning Knowledge Neural Graphs

name = 'FMult2'

n_layers = 3

k

n = 50

score_func = 'compositional'

discrete_points

entity_embeddings

relation_embeddings

build_func(Vec)

build_chain_funcs(list_Vec)

compute_func(W, b, x) → torch.FloatTensor

function(list_W, list_b)

trapezoid(list_W, list_b)

forward_triples(idx_triple: torch.Tensor) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

```
class dicee.models.function_space.LFMult1(args)
```

Bases: *[dicee.models.base_model.BaseKGE](#)*

Embedding with trigonometric functions. We represent all entities and relations in the complex number space as: $f(x) = \sum_{k=0}^{d-1} w_k e^{kix}$. and use the three differents scoring function as in the paper to evaluate the score

name = 'LFMult1'

entity_embeddings

```

relation_embeddings

forward_triples (idx_triple)
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

tri_score (h, r, t)

vtp_score (h, r, t)

class dicee.models.function_space.LFMult (args)
    Bases: dicee.models.base_model.BaseKGE

    Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:  $f(x) = \sum_{i=0}^{d-1} a_k x^{i \% d}$  and use the three differents scoring function as in the paper to evaluate the score.
    We also consider combining with Neural Networks.

    name = 'LFMult'

    entity_embeddings

    relation_embeddings

    degree

    m

    x_values

    forward_triples (idx_triple)
        Score a batch of (head, relation, tail) index triples.

        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].

        Returns
            Shape (batch_size,) triple scores.

        Return type
            torch.FloatTensor

    construct_multi_coeff (x)

    poly_NN (x, coefh, coefr, coeft)
        Constructing a 2 layers NN to represent the embeddings.  $h = \text{sigma}(wh^T x + bh)$ ,  $r = \text{sigma}(wr^T x + br)$ ,
         $t = \text{sigma}(wt^T x + bt)$ 

    linear (x, w, b)

```

scalar_batch_NN (*a, b, c*)

element wise multiplication between a,b and c: Inputs : a, b, c ==> torch.tensor of size batch_size x m x d
Output : a tensor of size batch_size x d

tri_score (*coeff_h, coeff_r, coeff_t*)

this part implement the trilinear scoring techniques:

$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \frac{a_i b_j c_k}{1+(i+j+k)d}$$

1. generate the range for i,j and k from [0 d-1]
2. perform $\frac{a_i b_j c_k}{1+(i+j+k)d}$ in parallel for every batch
3. take the sum over each batch

vtp_score (*h, r, t*)

this part implement the vector triple product scoring techniques:

$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \frac{a_i c_j b_k - b_i c_j a_k}{(1+(i+j)d)(1+k)}$$

1. generate the range for i,j and k from [0 d-1]
2. Compute the first and second terms of the sum
3. Multiply with then denominator and take the sum
4. take the sum over each batch

comp_func (*h, r, t*)

this part implement the function composition scoring techniques: i.e. score = <hor, t>

polynomial (*coeff, x, degree*)

This function takes a matrix tensor of coefficients (coeff), a tensor vector of points x and range of integer [0,1,...d] and return a vector tensor (coeff[0][0] + coeff[0][1]x +...+ coeff[0][d]x^d,

coeff[1][0] + coeff[1][1]x +...+ coeff[1][d]x^d)

pop (*coeff, x, degree*)

This function allow us to evaluate the composition of two polynomes without for loops :) it takes a matrix tensor of coefficients (coeff), a matrix tensor of points x and range of integer [0,1,...d]

and return a tensor (coeff[0][0] + coeff[0][1]x +...+ coeff[0][d]x^d,

coeff[1][0] + coeff[1][1]x +...+ coeff[1][d]x^d)

dicee.models.literal

Classes

LiteralEmbeddings

A model for learning and predicting numerical literals using pre-trained KGE.

Module Contents

class dicee.models.literal.**LiteralEmbeddings** (*num_of_data_properties: int, embedding_dims: int, entity_embeddings: torch.tensor, dropout: float = 0.3, gate_residual=True, freeze_entity_embeddings=True*)

Bases: `torch.nn.Module`

A model for learning and predicting numerical literals using pre-trained KGE.

num_of_data_properties

Number of data properties (attributes).

Type

int

embedding_dims

Dimension of the embeddings.

Type

int

entity_embeddings

Pre-trained entity embeddings.

Type

`torch.tensor`

dropout

Dropout rate for regularization.

Type

float

gate_residual

Whether to use gated residual connections.

Type

bool

freeze_entity_embeddings

Whether to freeze the entity embeddings during training.

Type

bool

embedding_dim

num_of_data_properties

hidden_dim

gate_residual = True

freeze_entity_embeddings = True

entity_embeddings

data_property_embeddings

fc

fc_out

dropout

```

gated_residual_proj

layer_norm

forward(entity_idx, attr_idx)

```

Parameters

- **entity_idx** (*Tensor*) – Entity indices (batch).
- **attr_idx** (*Tensor*) – Attribute (Data property) indices (batch).

Returns

scalar predictions.

Return type

Tensor

property device

dicee.models.octonion

Classes

<i>OMult</i>	Base class for all Knowledge Graph Embedding models.
<i>ConvO</i>	Base class for all Knowledge Graph Embedding models.
<i>ACConvO</i>	Additive Convolutional Octonion Knowledge Graph Embeddings

Functions

<i>octonion_mul</i> (*, O_1, O_2)
<i>octonion_mul_norm</i> (*, O_1, O_2)

Module Contents

`dicee.models.octonion.octonion_mul(*, O_1, O_2)`

`dicee.models.octonion.octonion_mul_norm(*, O_1, O_2)`

class `dicee.models.octonion.OMult` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'OMult'

static octonion_normalizer (*emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4, emb_rel_e5, emb_rel_e6, emb_rel_e7*)

score (*head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor, tail_ent_emb: torch.FloatTensor*)

k_vs_all_score (*bpe_head_ent_emb, bpe_rel_ent_emb, E*)

forward_k_vs_all (*x*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples, i.e., `[score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8]`, shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch, |Entities|)

class `dicee.models.octonion.ConvO` (*args: dict*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'ConvO'

conv2d

fc_num_input

fc1

bn_conv2d

norm_fc1

feature_map_dropout

static octonion_normalizer (*emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4, emb_rel_e5, emb_rel_e6, emb_rel_e7*)

residual_convolution (*O_1, O_2*)

forward_triples (*x*: *torch.Tensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: *torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities)

class *dicee.models.octonion.AConvO* (*args*: *dict*)

Bases: *dicee.models.base_model.BaseKGE*

Additive Convolutional Octonion Knowledge Graph Embeddings

name = 'AConvO'

conv2d

fc_num_input

fc1

bn_conv2d

norm_fc1

feature_map_dropout

static *octonion_normalizer* (*emb_rel_e0*, *emb_rel_e1*, *emb_rel_e2*, *emb_rel_e3*, *emb_rel_e4*, *emb_rel_e5*, *emb_rel_e6*, *emb_rel_e7*)

residual_convolution (*O_1*, *O_2*)

forward_triples (*x*: *torch.Tensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: *torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities)

dicee.models.pykeen_models

Classes

PykeenKGE

A class for using knowledge graph embedding models implemented in Pykeen

Module Contents

class dicee.models.pykeen_models.**PykeenKGE** (*args: dict*)

Bases: *dicee.models.base_model.BaseKGE*

A class for using knowledge graph embedding models implemented in Pykeen

Notes: Pykeen_DistMult: C Pykeen_ComplEx: Pykeen_QuatE: Pykeen_MuRE: Pykeen_CP: Pykeen_HolE: Pykeen_HolE: Pykeen_HolE: Pykeen_TransD: Pykeen_TransE: Pykeen_TransF: Pykeen_TransH: Pykeen_TransR:

model_kwargs

name

model

loss_history = []

args

entity_embeddings = None

relation_embeddings = None

forward_k_vs_all (*x: torch.LongTensor*)

=> Explicit version by this we can apply bn and dropout

(1) Retrieve embeddings of heads and relations + apply Dropout & Normalization if given. h, r = self.get_head_relation_representation(x) # (2) Reshape (1). if self.last_dim > 0:

h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim, self.last_dim)

(3) Reshape all entities. if self.last_dim > 0:

t = self.entity_embeddings.weight.reshape(self.num_entities, self.embedding_dim, self.last_dim)

else:

t = self.entity_embeddings.weight

(4) Call the score_t from interactions to generate triple scores. return self.interaction.score_t(h=h, r=r, all_entities=t, slice_size=1)

forward_triples (*x: torch.LongTensor*) → torch.FloatTensor

=> Explicit version by this we can apply bn and dropout

(1) Retrieve embeddings of heads, relations and tails and apply Dropout & Normalization if given. h, r, t = self.get_triple_representation(x) # (2) Reshape (1). if self.last_dim > 0:

h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim, self.last_dim) t = t.reshape(len(x), self.embedding_dim, self.last_dim)

```
# (3) Compute the triple score return self.interaction.score(h=h, r=r, t=t, slice_size=None, slice_dim=0)

abstractmethod forward_k_vs_sample (x: torch.LongTensor, target_entity_idx)
    Score a (head, relation) batch against a sampled subset of entities.

    Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns
    Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type
    torch.FloatTensor
```

dicee.models.quaternion

Classes

<i>QMult</i>	Base class for all Knowledge Graph Embedding models.
<i>ConvQ</i>	Convolutional Quaternion Knowledge Graph Embeddings
<i>ACConvQ</i>	Additive Convolutional Quaternion Knowledge Graph Embeddings

Functions

<i>quaternion_mul_with_unit_norm</i> (*, Q_1, Q_2)
--

Module Contents

`dicee.models.quaternion.quaternion_mul_with_unit_norm(*, Q_1, Q_2)`

class `dicee.models.quaternion.QMult` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

`name = 'QMult'`

`explicit = True`

quaternion_multiplication_followed_by_inner_product (*h, r, t*)

Parameters

- **h** – shape: (**batch_dims*, dim) The head representations.
- **r** – shape: (**batch_dims*, dim) The head representations.
- **t** – shape: (**batch_dims*, dim) The tail representations.

Returns

Triple scores.

static quaternion_normalizer (*x: torch.FloatTensor*) → torch.FloatTensor

Normalize the length of relation vectors, if the forward constraint has not been applied yet.

Absolute value of a quaternion

$$|a + bi + cj + dk| = \sqrt{a^2 + b^2 + c^2 + d^2}$$

L2 norm of quaternion vector:

$$\|x\|^2 = \sum_{i=1}^d |x_i|^2 = \sum_{i=1}^d (x_i.re^2 + x_i.im_1^2 + x_i.im_2^2 + x_i.im_3^2)$$

Parameters

x – The vector.

Returns

The normalized vector.

score (*head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor, tail_ent_emb: torch.FloatTensor*)

k_vs_all_score (*bpe_head_ent_emb, bpe_rel_ent_emb, E*)

Parameters

- **bpe_head_ent_emb**
- **bpe_rel_ent_emb**
- **E**

forward_k_vs_all (*x*)

Parameters

x

forward_k_vs_sample (*x, target_entity_idx*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e., [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,| Entities|)

class dicee.models.quaternion.**ConvQ** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Convolutional Quaternion Knowledge Graph Embeddings

name = 'ConvQ'

entity_embeddings

relation_embeddings

conv2d

fc_num_input

fc1

bn_conv1

bn_conv2

feature_map_dropout

residual_convolution(*Q_1*, *Q_2*)

forward_triples(*indexed_triple*: torch.Tensor) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(*x*: torch.Tensor)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

class dicee.models.quaternion.**AConvQ**(*args*)

Bases: *dicee.models.base_model.BaseKGE*

Additive Convolutional Quaternion Knowledge Graph Embeddings

name = 'AConvQ'

entity_embeddings

relation_embeddings

conv2d

fc_num_input

fc1

bn_conv1

bn_conv2

feature_map_dropout

residual_convolution(*Q_1*, *Q_2*)

forward_triples (*indexed_triple: torch.Tensor*) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*x: torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

dicee.models.real

Classes

<i>DistMult</i>	DistMult: bilinear diagonal knowledge graph embedding.
<i>TransE</i>	TransE: translation-based knowledge graph embedding.
<i>Shallom</i>	Shallom: shallow neural model for relation prediction.
<i>Pyke</i>	Pyke: Physical Embedding Model for Knowledge Graphs.
<i>CoKEConfig</i>	Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.
<i>CoKE</i>	Contextualized Knowledge Graph Embedding (CoKE) model.

Module Contents

class dicee.models.real.**DistMult** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

DistMult: bilinear diagonal knowledge graph embedding.

Scores a triple (h, r, t) as the element-wise product of the head, relation, and tail embeddings summed over the embedding dimension:

$$f(h, r, t) = \sum_i h_i \cdot r_i \cdot t_i$$

Simple yet effective baseline; incapable of modelling asymmetric relations.

References

Yang et al., *Embedding Entities and Relations for Learning and Inference in Knowledge Bases*, ICLR 2015. <https://arxiv.org/abs/1412.6575>

name = 'DistMult'

k_vs_all_score (*emb_h: torch.FloatTensor, emb_r: torch.FloatTensor, emb_E: torch.FloatTensor*) → torch.FloatTensor

Score a head/relation batch against all entity embeddings.

Computes $(h * r) @ E^T$ after applying hidden dropout and normalisation to the element-wise product.

Parameters

- **emb_h** (*torch.FloatTensor*) – Head entity embeddings, shape (batch_size, embedding_dim).
- **emb_r** (*torch.FloatTensor*) – Relation embeddings, shape (batch_size, embedding_dim).
- **emb_E** (*torch.FloatTensor*) – All entity embeddings, shape (num_entities, embedding_dim).

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_all (*x: torch.LongTensor*) → *torch.FloatTensor*

KvsAll forward pass: score head/relation against all entities.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*) → *torch.FloatTensor*

KvsSample forward pass: score head/relation against a sampled entity subset.

Parameters

- **x** (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (*torch.LongTensor*) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

score (*h: torch.FloatTensor, r: torch.FloatTensor, t: torch.FloatTensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) embedding triples.

Parameters

- **h** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).
- **r** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).
- **t** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor


```
class dicee.models.real.TransE(args)
```

Bases: `dicee.models.base_model.BaseKGE`

TransE: translation-based knowledge graph embedding.

Models a relation r as a translation in embedding space such that $h + r \approx t$ for a true triple (h, r, t) . The score function is defined as:

$$f(h, r, t) = \text{margin} - ||h + r - t||_2$$

TransE is effective for 1-to-1 relations but struggles with reflexive, one-to-many, and many-to-one patterns.

References

Bordes et al., *Translating Embeddings for Modeling Multi-relational Data*, NeurIPS 2013. <https://proceedings.neurips.cc/paper/2013/file/1cecc7a77928ca8133fa24680a88d2f9-Paper.pdf>

```
name = 'TransE'
```

```
margin = 4
```

```
score(head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor,
      tail_ent_emb: torch.FloatTensor) → torch.FloatTensor
```

Score a batch of triples using the TransE margin-distance formula.

Parameters

- **head_ent_emb** (`torch.FloatTensor`) – Each has shape (batch_size, embedding_dim).
- **rel_ent_emb** (`torch.FloatTensor`) – Each has shape (batch_size, embedding_dim).
- **tail_ent_emb** (`torch.FloatTensor`) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) scores equal to $\text{margin} - ||h + r - t||_2$.

Return type

`torch.FloatTensor`

```
forward_k_vs_all(x: torch.Tensor) → torch.FloatTensor
```

KvsAll forward pass: score head/relation against all entities.

Computes $\text{margin} - ||h + r - e||_2$ for every entity embedding e .

Parameters

x (`torch.Tensor`) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

```
class dicee.models.real.Shallom(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Shallom: shallow neural model for relation prediction.

Represents each triple as the concatenation of head and tail entity embeddings and feeds it through a two-layer MLP to predict the relation. Designed for the `RelationPrediction` labelling form.

References

Demir et al., *A Shallow Neural Model for Relation Prediction*, ISWC 2021. <https://arxiv.org/abs/2101.09090>

`name = 'Shallom'`

`shallom`

`get_embeddings()` → Tuple[numpy.ndarray, None]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- `entity_embeddings` (numpy.ndarray) – Shape (num_entities, embedding_dim).
- `relation_embeddings` (numpy.ndarray) – Shape (num_relations, embedding_dim).

`forward_k_vs_all(x)` → torch.FloatTensor

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

`forward_triples(x)` → torch.FloatTensor

Score a batch of triples by looking up relation scores from `forward_k_vs_all`.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

`class dicee.models.real.Pyke(args)`

Bases: `dicee.models.base_model.BaseKGE`

Pyke: Physical Embedding Model for Knowledge Graphs.

Scores a triple (h, r, t) based on the average pairwise distance between head-to-relation and relation-to-tail in embedding space:

$$f(h, r, t) = \text{margin} - (||h - r||_2 + ||r - t||_2) / 2$$

The model encodes geometric proximity between entities and the relations that connect them.

`name = 'Pyke'`

`dist_func`

`margin = 1.0`

forward_triples (*x*: *torch.LongTensor*) → *torch.FloatTensor*

Score a batch of triples using the Pyke distance formula.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class *dicee.models.real.CoKEConfig*

Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.

block_size

Sequence length for transformer (3 for triples: head, relation, tail)

vocab_size

Total vocabulary size (num_entities + num_relations)

n_layer

Number of transformer layers

n_head

Number of attention heads per layer

n_embd

Embedding dimension (set to match model embedding_dim)

dropout

Dropout rate applied throughout the model

bias

Whether to use bias in linear layers

causal

Whether to use causal masking (False for bidirectional attention)

block_size: int = 3

vocab_size: int = None

n_layer: int = 6

n_head: int = 8

n_embd: int = None

dropout: float = 0.3

bias: bool = True

causal: bool = False

```
class dicee.models.real.CoKE(args, config: CoKEConfig = CoKEConfig())
```

Bases: `dicee.models.base_model.BaseKGE`

Contextualized Knowledge Graph Embedding (CoKE) model. Based on: <https://arxiv.org/pdf/1911.02168>.

CoKE uses a transformer encoder to learn contextualized representations of entities and relations. For link prediction, it predicts masked elements in (head, relation, tail) triples using bidirectional attention, similar to BERT's masked language modeling approach.

The model creates a sequence [head_emb, relation_emb, mask_emb], adds positional embeddings, and processes it through transformer layers to predict the tail entity.

```
name = 'CoKE'
```

```
config
```

```
pos_emb
```

```
mask_emb
```

```
blocks
```

```
ln_f
```

```
coke_dropout
```

```
forward_k_vs_all(x: torch.Tensor)
```

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

```
score(emb_h, emb_r, emb_t)
```

```
forward_k_vs_sample(x: torch.LongTensor, target_entity_idx: torch.LongTensor)
```

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

`dicee.models.static_funcs`

Functions

```
quaternion_mul(→ Tuple[torch.Tensor, torch.Tensor, ...]) Perform quaternion multiplication
```

Module Contents

```
dicee.models.static_funcs.quaternion_mul(* , Q_1, Q_2)
    → Tuple[torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor]
    Perform quaternion multiplication :param Q_1: :param Q_2: :return:
```

`dicee.models.transformers`

Full definition of a GPT Language Model, all of it in this single file. References: 1) the official GPT-2 TensorFlow implementation released by OpenAI: <https://github.com/openai/gpt-2/blob/master/src/model.py> 2) huggingface/transformers PyTorch implementation: https://github.com/huggingface/transformers/blob/main/src/transformers/models/gpt2/modeling_gpt2.py

Classes

<i>ByteE</i>	Base class for all Knowledge Graph Embedding models.
<i>LayerNorm</i>	LayerNorm but with an optional bias. PyTorch doesn't support simply bias=False
<i>SelfAttention</i>	Base class for all neural network modules.
<i>MLP</i>	Base class for all neural network modules.
<i>Block</i>	Base class for all neural network modules.
<i>GPTConfig</i>	
<i>GPT</i>	Base class for all neural network modules.

Module Contents

```
class dicee.models.transformers.ByteE(*args, **kwargs)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

```
name = 'ByteE'
```

```
config
```

```
temperature = 0.5
```

```
topk = 2
```

```
transformer
```

lm_head

loss_function (*yhat_batch*, *y_batch*)

Compute the loss between model predictions and targets.

Delegates to `self.loss` which is configured in `BaseKGE.__init__` based on the scoring technique (BCEWithLogitsLoss for entity/relation prediction, CrossEntropyLoss for classification).

Parameters

- **yhat_batch** (*torch.FloatTensor*) – Model output scores, shape (batch_size, *).
- **y_batch** (*torch.FloatTensor*) – Ground-truth labels of the same shape as *yhat_batch*.

Returns

Scalar loss value.

Return type

torch.FloatTensor

forward (*x*: *torch.LongTensor*)

Parameters

x (*B by T tensor*)

generate (*idx*, *max_new_tokens*, *temperature=1.0*, *top_k=None*)

Take a conditioning sequence of indices *idx* (*LongTensor* of shape (b,t)) and complete the sequence *max_new_tokens* times, feeding the predictions back into the model each time. Most likely you'll want to make sure to be in `model.eval()` mode of operation for this.

training_step (*batch*, *batch_idx=None*)

Execute one optimisation step for the given mini-batch.

Handles two- and three-element batches produced by the different dataset classes (*KvsAll* / *NegSample* vs. *KvsSample*).

Parameters

- **batch** (*tuple*) – (*x*, *y*) for standard scoring, or (*x*, *y_select*, *y*) for sample-based labelling.
- **batch_idx** (*int*, *optional*) – Index of the current batch (unused, kept for Lightning API compat).

Returns

Scalar loss value for this batch.

Return type

torch.FloatTensor

class `dicee.models.transformers.LayerNorm` (*ndim*, *bias*)

Bases: `torch.nn.Module`

LayerNorm but with an optional bias. PyTorch doesn't support simply `bias=False`

weight

bias

forward (*input*)

```
class dicee.models.transformers.SelfAttention(config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training` (*bool*) – Boolean represents whether this module is in training or evaluation mode.

`c_attn`

`c_proj`

`attn_dropout`

`resid_dropout`

`n_head`

`n_embd`

`dropout`

`causal`

`flash = True`

`forward` (*x*)

```
class dicee.models.transformers.MLP (config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

c_fc

gelu

c_proj

dropout

forward (*x*)

```
class dicee.models.transformers.Block (config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:


```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`ln_1`

`attn`

`ln_2`

`mlp`

`forward(x)`

```
class dicee.models.transformers.GPTConfig
```

```
    block_size: int = 1024
```

```
    vocab_size: int = 50304
```

```
    n_layer: int = 12
```

```
    n_head: int = 12
```

```
    n_embd: int = 768
```

```
    dropout: float = 0.0
```

```
    bias: bool = False
```

```
    causal: bool = True
```

```
class dicee.models.transformers.GPT(config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

config

transformer

lm_head

get_num_params (*non_embedding=True*)

Return the number of parameters in the model. For non-embedding count (default), the position embeddings get subtracted. The token embeddings would too, except due to the parameter sharing these params are actually used as weights in the final layer, so we include them.

forward (*idx, targets=None*)

crop_block_size (*block_size*)

classmethod from_pretrained (*model_type, override_args=None*)

configure_optimizers (*weight_decay, learning_rate, betas, device_type*)

estimate_mfu (*fwdbwd_per_iter, dt*)

estimate model flops utilization (MFU) in units of A100 bfloat16 peak FLOPS

Classes

<i>ADOPT</i>	ADOPT Optimizer.
<i>BaseKGELightning</i>	Thin PyTorch Lightning wrapper shared by all KGE models.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>IdentityClass</i>	No-op normalisation / dropout placeholder.
<i>Block</i>	Base class for all neural network modules.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>DistMult</i>	DistMult: bilinear diagonal knowledge graph embedding.
<i>TransE</i>	TransE: translation-based knowledge graph embedding.
<i>Shallom</i>	Shallom: shallow neural model for relation prediction.
<i>Pyke</i>	Pyke: Physical Embedding Model for Knowledge Graphs.
<i>CoKEConfig</i>	Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.
<i>CoKE</i>	Contextualized Knowledge Graph Embedding (CoKE) model.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>ConEx</i>	Convolutional ComplEx Knowledge Graph Embeddings
<i>AConEx</i>	Additive Convolutional ComplEx Knowledge Graph Embeddings
<i>Complex</i>	Base class for all Knowledge Graph Embedding models.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>IdentityClass</i>	No-op normalisation / dropout placeholder.
<i>QMult</i>	Base class for all Knowledge Graph Embedding models.
<i>ConvQ</i>	Convolutional Quaternion Knowledge Graph Embeddings
<i>AConvQ</i>	Additive Convolutional Quaternion Knowledge Graph Embeddings
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>IdentityClass</i>	No-op normalisation / dropout placeholder.
<i>OMult</i>	Base class for all Knowledge Graph Embedding models.
<i>ConvO</i>	Base class for all Knowledge Graph Embedding models.
<i>AConvO</i>	Additive Convolutional Octonion Knowledge Graph Embeddings
<i>Keci</i>	Keci: Knowledge Graph Embedding via Clifford Algebra.
<i>CKeci</i>	Without learning dimension scaling
<i>DeCaL</i>	Base class for all Knowledge Graph Embedding models.
<i>KeciTransformer</i>	Keci with Transformer architecture.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>PykeenKGE</i>	A class for using knowledge graph embedding models implemented in Pykeen
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>FMult</i>	Learning Knowledge Neural Graphs
<i>GFMult</i>	Learning Knowledge Neural Graphs
<i>FMult2</i>	Learning Knowledge Neural Graphs
<i>LFMult1</i>	Embedding with trigonometric functions. We represent all entities and relations in the complex number space as:
<i>LFMult</i>	Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:

continues on next page

Table 40 – continued from previous page

<i>DualE</i>	Dual Quaternion Knowledge Graph Embeddings (https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657)
--------------	---

Functions

<code>quaternion_mul</code>	($\rightarrow$ Tuple[torch.Tensor, torch.Tensor, ...])	Perform quaternion multiplication
<code>quaternion_mul_with_unit_norm</code>	(*, Q_1, Q_2)	
<code>octonion_mul</code>	(*, O_1, O_2)	
<code>octonion_mul_norm</code>	(*, O_1, O_2)	

Package Contents

```
class dicee.models.ADOPT (params: torch.optim.optimizer.ParamsT, lr: float | torch.Tensor = 0.001,
    betas: Tuple[float, float] = (0.9, 0.9999), eps: float = 1e-06,
    clip_lambda: Callable[[int], float] | None = lambda step: ..., weight_decay: float = 0.0,
    decouple: bool = False, *, foreach: bool | None = None, maximize: bool = False,
    capturable: bool = False, differentiable: bool = False, fused: bool | None = None)
```

Bases: torch.optim.optimizer.Optimizer

ADOPT Optimizer.

ADOPT is an adaptive learning rate optimization algorithm that combines momentum-based updates with adaptive per-parameter learning rates. It uses exponential moving averages of gradients and squared gradients, with gradient clipping for stability.

The algorithm performs the following key operations: 1. Normalizes gradients by the square root of the second moment estimate 2. Applies optional gradient clipping based on the training step 3. Updates parameters using momentum-smoothed normalized gradients 4. Supports decoupled weight decay (AdamW-style) or L2 regularization

Mathematical formulation:

$$m_t = \beta_1 * m_{\{t-1\}} + (1 - \beta_1) * \text{clip}(g_t / \sqrt{v_t}) \quad v_t = \beta_2 * v_{\{t-1\}} + (1 - \beta_2) * g_t^2 \quad \theta_t = \theta_{\{t-1\}} - \alpha * m_t$$

where:

- θ_t : parameter at step t
- g_t : gradient at step t
- m_t : first moment estimate (momentum)
- v_t : second moment estimate (variance)
- α : learning rate
- β_1, β_2 : exponential decay rates
- `clip()`: optional gradient clipping function

Reference:

Original implementation: <https://github.com/iShohei220/adapt>

Parameters

- **params** (*ParamsT*) – Iterable of parameters to optimize or dicts defining parameter groups.

- **lr** (*float or Tensor, optional*) – Learning rate. Can be a float or 1-element Tensor. Default: 1e-3
- **betas** (*Tuple[float, float], optional*) – Coefficients (β_1, β_2) for computing running averages of gradient and its square. β_1 controls momentum, β_2 controls variance. Default: (0.9, 0.9999)
- **eps** (*float, optional*) – Term added to denominator to improve numerical stability. Default: 1e-6
- **clip_lambda** (*Callable[[int], float], optional*) – Function that takes the step number and returns the gradient clipping threshold. Common choices: - lambda step: step**0.25 (default, gradually increases clipping threshold) - lambda step: 1.0 (constant clipping) - None (no clipping) Default: lambda step: step**0.25
- **weight_decay** (*float, optional*) – Weight decay coefficient (L2 penalty). Default: 0.0
- **decouple** (*bool, optional*) – If True, uses decoupled weight decay (AdamW-style), applying weight decay directly to parameters. If False, adds weight decay to gradients (L2 regularization). Default: False
- **foreach** (*bool, optional*) – If True, uses the faster foreach implementation for multi-tensor operations. Default: None (auto-select)
- **maximize** (*bool, optional*) – If True, maximizes parameters instead of minimizing. Useful for reinforcement learning. Default: False
- **capturable** (*bool, optional*) – If True, the optimizer is safe to capture in a CUDA graph. Requires learning rate as Tensor. Default: False
- **differentiable** (*bool, optional*) – If True, the optimization step can be differentiated. Useful for meta-learning. Default: False
- **fused** (*bool, optional*) – If True, uses fused kernel implementation (currently not supported). Default: None

Raises

- **ValueError** – If learning rate, epsilon, betas, or weight_decay are invalid.
- **RuntimeError** – If fused is enabled (not currently supported).
- **RuntimeError** – If lr is a Tensor with foreach=True and capturable=False.

Example

```
>>> # Basic usage
>>> optimizer = ADOPT(model.parameters(), lr=0.001)
>>> optimizer.zero_grad()
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With decoupled weight decay
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01,
↳decouple=True)
```

```
>>> # Custom gradient clipping
>>> optimizer = ADOPT(model.parameters(), clip_lambda=lambda step: max(1.0,
↳step**0.5))
```

Note

- For most use cases, the default hyperparameters work well
- Consider using `decouple=True` for better generalization (similar to AdamW)
- The `clip_lambda` function helps stabilize training in early steps

`clip_lambda`

`__setstate__` (*state*)

Restore optimizer state from a checkpoint.

This method handles backward compatibility when loading optimizer state from older versions. It ensures all required fields are present with default values and properly converts step counters to tensors if needed.

Key responsibilities: 1. Set default values for newly added hyperparameters 2. Convert old-style scalar step counters to tensor format 3. Place step tensors on appropriate devices based on capturable/fused modes

Parameters

state (*dict*) – Optimizer state dictionary (typically from `torch.load()`).

Note

- This enables loading checkpoints saved with older ADOPT versions
- Step counters are converted to appropriate device/dtype for compatibility
- Capturable and fused modes require step tensors on parameter devices

`step` (*closure=None*)

Perform a single optimization step.

This method executes one iteration of the ADOPT optimization algorithm across all parameter groups. It orchestrates the following workflow:

1. Optionally evaluates a closure to recompute the loss (useful for algorithms like LBFGS or when loss needs multiple evaluations)
2. For each parameter group: - Collects parameters with gradients and their associated state - Extracts hyperparameters (betas, learning rate, etc.) - Calls the functional `adopt()` API to perform the actual update
3. Returns the loss value if a closure was provided

The functional API (`adopt()`) handles three execution modes: - Single-tensor: Updates one parameter at a time (default, JIT-compatible) - Multi-tensor (foreach): Batches operations for better performance - Fused: Uses fused CUDA kernels (not yet implemented)

Gradient scaling support: This method is compatible with automatic mixed precision (AMP) training. It can access `grad_scale` and `found_inf` attributes for gradient unscaling and inf/nan detection when used with `GradScaler`.

Parameters

closure (*Callable, optional*) – A callable that reevaluates the model and returns the loss. The closure should: - Enable gradients (`torch.enable_grad()`) - Compute forward pass - Compute loss - Compute backward pass - Return the loss value Example: `lambda: (loss := model(x), loss.backward(), loss)[-1]` Default: `None`

Returns

The loss value returned by the closure, or None if no closure was provided.

Return type

Optional[Tensor]

Example

```
>>> # Standard usage
>>> loss = criterion(model(input), target)
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With closure (e.g., for line search)
>>> def closure():
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     return loss
>>> loss = optimizer.step(closure)
```

Note

- Call `zero_grad()` before computing gradients for the next step
- CUDA graph capture is checked for safety when `capturable=True`
- The method is thread-safe for different parameter groups

```
class dicee.models.BaseKGELightning(*args, **kwargs)
```

Bases: `lightning.LightningModule`

Thin PyTorch Lightning wrapper shared by all KGE models.

Provides the standard Lightning training loop hooks (`training_step`, `on_train_epoch_end`, `configure_optimizers`) as well as a helper for reporting model size. All concrete KGE models should extend `BaseKGE` rather than this class directly.

```
training_step_outputs = []
```

```
mem_of_model() → Dict
```

Size of model in MB and number of params

```
training_step(batch, batch_idx=None)
```

Execute one optimisation step for the given mini-batch.

Handles two- and three-element batches produced by the different dataset classes (`KvsAll` / `NegSample` vs. `KvsSample`).

Parameters

- **batch** (*tuple*) – (`x`, `y`) for standard scoring, or (`x`, `y_select`, `y`) for sample-based labelling.

- **batch_idx** (*int, optional*) – Index of the current batch (unused, kept for Lightning API compat).

Returns

Scalar loss value for this batch.

Return type

`torch.FloatTensor`

loss_function (*yhat_batch: torch.FloatTensor, y_batch: torch.FloatTensor*) → `torch.FloatTensor`

Compute the loss between model predictions and targets.

Delegates to `self.loss` which is configured in `BaseKGE.__init__` based on the scoring technique (BCEWithLogitsLoss for entity/relation prediction, CrossEntropyLoss for classification).

Parameters

- **yhat_batch** (*torch.FloatTensor*) – Model output scores, shape (batch_size, *).
- **y_batch** (*torch.FloatTensor*) – Ground-truth labels of the same shape as *yhat_batch*.

Returns

Scalar loss value.

Return type

`torch.FloatTensor`

on_train_epoch_end (*args, **kwargs)

Called in the training loop at the very end of the epoch.

To access all batch outputs at the end of the epoch, you can cache step outputs as an attribute of the `LightningModule` and access them in this hook:

```
class MyLightningModule(L.LightningModule):
    def __init__(self):
        super().__init__()
        self.training_step_outputs = []

    def training_step(self):
        loss = ...
        self.training_step_outputs.append(loss)
        return loss

    def on_train_epoch_end(self):
        # do something with all training_step outputs, for example:
        epoch_mean = torch.stack(self.training_step_outputs).mean()
        self.log("training_epoch_mean", epoch_mean)
        # free up the memory
        self.training_step_outputs.clear()
```

test_epoch_end (outputs: List[Any])

test_dataloader () → None

An iterable or collection of iterables specifying test samples.

For more information about multiple dataloaders, see this section.

For data processing use the following pattern:

- download in `prepare_data()`

- process and split in `setup()`

However, the above are only necessary for distributed processing.

Warning

do not assign state in `prepare_data`

- `test()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Note

If you don't need a test dataset and a `test_step()`, you don't need to implement this method.

`val_dataloader()` → None

An iterable or collection of iterables specifying validation samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `fit()`
- `validate()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Note

If you don't need a validation dataset and a `validation_step()`, you don't need to implement this method.

`predict_dataloader()` → None

An iterable or collection of iterables specifying prediction samples.

For more information about multiple dataloaders, see this section.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `predict()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Returns

A `torch.utils.data.DataLoader` or a sequence of them specifying prediction samples.

`train_dataloader()` → None

An iterable or collection of iterables specifying training samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

Warning

do not assign state in `prepare_data`

- `fit()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

configure_optimizers (*parameters=None*)

Instantiate and return the optimiser for training.

The optimiser type is taken from `self.optimizer_name` which is set in `BaseKGE.init_params_with_sanity_checking()` from the `--optim` argument. Supported values: 'SGD', 'Adam', 'Adopt', 'AdamW', 'NAdam', 'Adagrad', 'ASGD', 'Muon'.

Parameters

parameters (*iterable, optional*) – Model parameters to optimise. Defaults to `self.parameters()` when None.

Returns

The configured optimiser instance.

Return type

`torch.optim.Optimizer`

class `dicee.models.BaseKGE` (*args: dict*)

Bases: `BaseKGELightning`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

`loss`

`selected_optimizer = None`

`normalizer_class = None`

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

`hidden_dropout`

`loss_history = []`

`byte_pair_encoding`

`max_length_subword_tokens`

`block_size`

`forward_byte_pair_encoded_k_vs_all` (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

`forward_byte_pair_encoded_triple` (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*) → *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x*: `torch.LongTensor` | `Tuple[torch.LongTensor, torch.LongTensor]`,
 y_idx: `torch.LongTensor` = None) → `torch.FloatTensor`

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- `Tuple (x, y_idx)` → `forward_k_vs_sample()`
- `(batch, 3)` tensor → `forward_triples()`
- `(batch, 2)` tensor → `forward_k_vs_all()`
- BPE triple tensor → `forward_byte_pair_encoded_triple()`
- BPE pair tensor → `forward_byte_pair_encoded_k_vs_all()`

Parameters

- **x** (`torch.LongTensor` or `Tuple[torch.LongTensor, torch.LongTensor]`) – Either a plain index tensor or a `(triple_idx, target_idx)` tuple for sample-based labelling.
- **y_idx** (`torch.LongTensor`, optional) – Target entity indices used by `forward_k_vs_sample()`. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.LongTensor`) → `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

- **x** (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings() → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

```
class dicee.models.IdentityClass (args=None)
```

Bases: torch.nn.Module

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (`--normalization None`). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = None

__call__(x)

static forward(x)

```
class dicee.models.Block (config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

ln_1

attn

ln_2

mlp

forward(*x*)

class `dicee.models.BaseKGE` (*args: dict*)

Bases: `BaseKGELightning`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (*h*, *r*, *t*) triples.
- `forward_k_vs_all()` — score a (*h*, *r*) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None


```

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

```

forward_byte_pair_encoded_k_vs_all (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*) → *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple*[*torch.LongTensor*, *torch.LongTensor*],
y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample()*
- (batch, 3) tensor → *forward_triples()*
- (batch, 2) tensor → *forward_k_vs_all()*
- BPE triple tensor → *forward_byte_pair_encoded_triple()*
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all()*

Parameters

- **x** (*torch.LongTensor* or *Tuple*[*torch.LongTensor*, *torch.LongTensor*]) – Either a plain index tensor or a (*triple_idx*, *target_idx*) tuple for sample-based labelling.
- **y_idx** (*torch.LongTensor*, optional) – Target entity indices used by *forward_k_vs_sample()*. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.LongTensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample(*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation(idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation(indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation(x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x*: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (torch.LongTensor) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (numpy.ndarray) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (numpy.ndarray) – Shape (num_relations, embedding_dim).

class dicee.models.DistMult (*args*)

Bases: `dicee.models.base_model.BaseKGE`

DistMult: bilinear diagonal knowledge graph embedding.

Scores a triple (h, r, t) as the element-wise product of the head, relation, and tail embeddings summed over the embedding dimension:

$$f(h, r, t) = \sum_i h_i \cdot r_i \cdot t_i$$

Simple yet effective baseline; incapable of modelling asymmetric relations.

References

Yang et al., *Embedding Entities and Relations for Learning and Inference in Knowledge Bases*, ICLR 2015. <https://arxiv.org/abs/1412.6575>

name = 'DistMult'

k_vs_all_score (*emb_h*: torch.FloatTensor, *emb_r*: torch.FloatTensor, *emb_E*: torch.FloatTensor)

→ torch.FloatTensor

Score a head/relation batch against all entity embeddings.

Computes $(h * r) @ E^T$ after applying hidden dropout and normalisation to the element-wise product.

Parameters

- **emb_h** (torch.FloatTensor) – Head entity embeddings, shape (batch_size, embedding_dim).

- **emb_r** (*torch.FloatTensor*) – Relation embeddings, shape (batch_size, embedding_dim).
- **emb_E** (*torch.FloatTensor*) – All entity embeddings, shape (num_entities, embedding_dim).

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_all (*x: torch.LongTensor*) → *torch.FloatTensor*

KvsAll forward pass: score head/relation against all entities.

Parameters

• **x** (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*) → *torch.FloatTensor*

KvsSample forward pass: score head/relation against a sampled entity subset.

Parameters

- **x** (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (*torch.LongTensor*) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

score (*h: torch.FloatTensor, r: torch.FloatTensor, t: torch.FloatTensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) embedding triples.

Parameters

- **h** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).
- **r** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).
- **t** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class *dicee.models.TransE* (*args*)

Bases: *dicee.models.base_model.BaseKGE*

TransE: translation-based knowledge graph embedding.

Models a relation r as a translation in embedding space such that $h + r \approx t$ for a true triple (h, r, t) . The score function is defined as:

$$f(h, r, t) = \text{margin} - ||h + r - t||_2$$

TransE is effective for 1-to-1 relations but struggles with reflexive, one-to-many, and many-to-one patterns.

References

Bordes et al., *Translating Embeddings for Modeling Multi-relational Data*, NeurIPS 2013. <https://proceedings.neurips.cc/paper/2013/file/1cecc7a77928ca8133fa24680a88d2f9-Paper.pdf>

name = 'TransE'

margin = 4

score (*head_ent_emb*: torch.FloatTensor, *rel_ent_emb*: torch.FloatTensor, *tail_ent_emb*: torch.FloatTensor) → torch.FloatTensor

Score a batch of triples using the TransE margin-distance formula.

Parameters

- **head_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **rel_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **tail_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) scores equal to $\text{margin} - ||h + r - t||_2$.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: torch.Tensor) → torch.FloatTensor

KvsAll forward pass: score head/relation against all entities.

Computes $\text{margin} - ||h + r - e||_2$ for every entity embedding e .

Parameters

x (torch.Tensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

class dicee.models.Shallom (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Shallom: shallow neural model for relation prediction.

Represents each triple as the concatenation of head and tail entity embeddings and feeds it through a two-layer MLP to predict the relation. Designed for the RelationPrediction labelling form.

References

Demir et al., *A Shallow Neural Model for Relation Prediction*, ISWC 2021. <https://arxiv.org/abs/2101.09090>

```
name = 'Shallom'
```

```
shallom
```

```
get_embeddings() → Tuple[numpy.ndarray, None]
```

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

```
forward_k_vs_all(x) → torch.FloatTensor
```

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

```
forward_triples(x) → torch.FloatTensor
```

Score a batch of triples by looking up relation scores from `forward_k_vs_all`.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

```
class dicee.models.Pyke(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Pyke: Physical Embedding Model for Knowledge Graphs.

Scores a triple (h, r, t) based on the average pairwise distance between head-to-relation and relation-to-tail in embedding space:

$$f(h, r, t) = \text{margin} - (||h - r||_2 + ||r - t||_2) / 2$$

The model encodes geometric proximity between entities and the relations that connect them.

```
name = 'Pyke'
```

```
dist_func
```

```
margin = 1.0
```

forward_triples (*x*: *torch.LongTensor*) → *torch.FloatTensor*

Score a batch of triples using the Pyke distance formula.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class *dicee.models.CoKEConfig*

Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.

block_size

Sequence length for transformer (3 for triples: head, relation, tail)

vocab_size

Total vocabulary size (num_entities + num_relations)

n_layer

Number of transformer layers

n_head

Number of attention heads per layer

n_embd

Embedding dimension (set to match model embedding_dim)

dropout

Dropout rate applied throughout the model

bias

Whether to use bias in linear layers

causal

Whether to use causal masking (False for bidirectional attention)

block_size: int = 3

vocab_size: int = None

n_layer: int = 6

n_head: int = 8

n_embd: int = None

dropout: float = 0.3

bias: bool = True

causal: bool = False


```
class dicee.models.CoKE(args, config: CoKEConfig = CoKEConfig())
```

Bases: `dicee.models.base_model.BaseKGE`

Contextualized Knowledge Graph Embedding (CoKE) model. Based on: <https://arxiv.org/pdf/1911.02168>.

CoKE uses a transformer encoder to learn contextualized representations of entities and relations. For link prediction, it predicts masked elements in (head, relation, tail) triples using bidirectional attention, similar to BERT's masked language modeling approach.

The model creates a sequence [head_emb, relation_emb, mask_emb], adds positional embeddings, and processes it through transformer layers to predict the tail entity.

```
name = 'CoKE'
```

```
config
```

```
pos_emb
```

```
mask_emb
```

```
blocks
```

```
ln_f
```

```
coke_dropout
```

```
forward_k_vs_all (x: torch.Tensor)
```

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

```
score (emb_h, emb_r, emb_t)
```

```
forward_k_vs_sample (x: torch.LongTensor, target_entity_idx: torch.LongTensor)
```

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

```
class dicee.models.BaseKGE(args: dict)
```

Bases: `BaseKGELightning`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.

- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

loss

`selected_optimizer = None`

`normalizer_class = None`

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

`hidden_dropout`

`loss_history = []`

`byte_pair_encoding`

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*) → *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from *self.args* with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the *args* dict. Falls back to sensible defaults for any missing key so that minimal *args* dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple[torch.LongTensor, torch.LongTensor]*,
y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample()*
- (batch, 3) tensor → *forward_triples()*
- (batch, 2) tensor → *forward_k_vs_all()*
- BPE triple tensor → *forward_byte_pair_encoded_triple()*
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all()*

Parameters

- **x** (*torch.LongTensor* or *Tuple[torch.LongTensor, torch.LongTensor]*) – Either a plain index tensor or a (*triple_idx*, *target_idx*) tuple for sample-based labelling.
- **y_idx** (*torch.LongTensor*, *optional*) – Target entity indices used by *forward_k_vs_sample()*. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x: torch.LongTensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises *ValueError* to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by *KvsSample* and *1vsSample* datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (*idx_hrt*)

→ *Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]*

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (*indexed_triple*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) where *T* is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of *T* subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

```

class dicee.models.ConEx(args)
    Bases: dicee.models.base_model.BaseKGE
    Convolutional ComplEx Knowledge Graph Embeddings
    name = 'ConEx'
    conv2d
    fc_num_input
    fc1
    norm_fc1
    bn_conv2d
    feature_map_dropout
    residual_convolution(C_1: Tuple[torch.Tensor, torch.Tensor],
        C_2: Tuple[torch.Tensor, torch.Tensor]) → torch.FloatTensor
        Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors
        that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds
        complex-valued embeddings :return:
    forward_k_vs_all(x: torch.Tensor) → torch.FloatTensor
        Score a (head, relation) batch against every entity.
        Sub-classes must override this method. The default implementation raises ValueError to make missing
        overrides obvious at runtime.
        Returns
            Shape (batch_size, num_entities) score matrix.
        Return type
            torch.FloatTensor
    forward_triples(x: torch.Tensor) → torch.FloatTensor
        Score a batch of (head, relation, tail) index triples.
        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].
        Returns
            Shape (batch_size,) triple scores.
        Return type
            torch.FloatTensor
    forward_k_vs_sample(x: torch.Tensor, target_entity_idx: torch.Tensor)
        Score a (head, relation) batch against a sampled subset of entities.
        Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must over-
        ride this method.
        Returns
            Shape (batch_size, k) score matrix where k is the number of sampled target entities.
        Return type
            torch.FloatTensor

```

```

class dicee.models.AConEx (args)
    Bases: dicee.models.base_model.BaseKGE
    Additive Convolutional ComplEx Knowledge Graph Embeddings
    name = 'AConEx'

    conv2d

    fc_num_input

    fc1

    norm_fc1

    bn_conv2d

    feature_map_dropout

    residual_convolution (C_1: Tuple[torch.Tensor, torch.Tensor],
                          C_2: Tuple[torch.Tensor, torch.Tensor]) → torch.FloatTensor
        Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors
        that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds
        complex-valued embeddings :return:

    forward_k_vs_all (x: torch.Tensor) → torch.FloatTensor
        Score a (head, relation) batch against every entity.

        Sub-classes must override this method. The default implementation raises ValueError to make missing
        overrides obvious at runtime.

        Returns
            Shape (batch_size, num_entities) score matrix.

        Return type
            torch.FloatTensor

    forward_triples (x: torch.Tensor) → torch.FloatTensor
        Score a batch of (head, relation, tail) index triples.

        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].

        Returns
            Shape (batch_size,) triple scores.

        Return type
            torch.FloatTensor

    forward_k_vs_sample (x: torch.Tensor, target_entity_idx: torch.Tensor)
        Score a (head, relation) batch against a sampled subset of entities.

        Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must over-
        ride this method.

        Returns
            Shape (batch_size, k) score matrix where k is the number of sampled target entities.

        Return type
            torch.FloatTensor

```

```
class dicee.models.Complex(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (*h*, *r*, *t*) triples.
- `forward_k_vs_all()` — score a (*h*, *r*) batch against every entity.

Parameters

args (*dict*) — Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

```
name = 'Complex'
```

```
static score(head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor,  
             tail_ent_emb: torch.FloatTensor)
```

```
static k_vs_all_score(emb_h: torch.FloatTensor, emb_r: torch.FloatTensor,  
                     emb_E: torch.FloatTensor)
```

Parameters

- `emb_h`
- `emb_r`
- `emb_E`

```
forward_k_vs_all(x: torch.LongTensor) → torch.FloatTensor
```

Score a (*head*, *relation*) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (*batch_size*, *num_entities*) score matrix.

Return type

`torch.FloatTensor`

```
forward_k_vs_sample(x: torch.LongTensor, target_entity_idx: torch.LongTensor)
```

Score a (*head*, *relation*) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (*batch_size*, *k*) score matrix where *k* is the number of sampled target entities.

Return type

`torch.FloatTensor`


```
class dicee.models.BaseKGE (args: dict)
```

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (*h*, *r*, *t*) triples.
- *forward_k_vs_all()* — score a (*h*, *r*) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

```

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all(x: torch.LongTensor) → torch.FloatTensor
    KvsAll scoring for BPE-encoded head entities and relations.

    Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

    Returns
        Shape (batch_size, num_bpe_entities) score matrix.

    Return type
        torch.FloatTensor

forward_byte_pair_encoded_triple(x: Tuple[torch.LongTensor, torch.LongTensor])
    → torch.FloatTensor
    NegSample scoring for BPE-encoded (head, relation, tail) triples.

    Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3, T) BPE token indices.

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

init_params_with_sanity_checking() → None
    Populate model hyper-parameters from self.args with safe defaults.

    Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the args dict. Falls back to sensible defaults for any missing key so that minimal args dicts (e.g. for unit tests) are still valid.

forward(x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
        y_idx: torch.LongTensor = None) → torch.FloatTensor
    Route the forward pass to the appropriate scoring method.

    Inspects the shape and type of x to decide which low-level scorer to call:
    • Tuple (x, y_idx) → forward_k_vs_sample()

```

- (batch, 3) tensor $\rightarrow$ `forward_triples()`
- (batch, 2) tensor $\rightarrow$ `forward_k_vs_all()`
- BPE triple tensor $\rightarrow$ `forward_byte_pair_encoded_triple()`
- BPE pair tensor $\rightarrow$ `forward_byte_pair_encoded_k_vs_all()`

Parameters

- **x** (`torch.LongTensor` or `Tuple[torch.LongTensor, torch.LongTensor]`) – Either a plain index tensor or a (triple_idx, target_idx) tuple for sample-based labelling.
- **y_idx** (`torch.LongTensor`, optional) – Target entity indices used by `forward_k_vs_sample()`. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.LongTensor`) $\rightarrow$ `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

- **x** (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

`torch.FloatTensor`

get_triple_representation (*idx_hrt*)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (*indexed_triple*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, *T*) where *T* is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, *T*, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of *T* subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, *T*) where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, *T*, embedding_dim), L2-normalised over the (*T*, *D*) dimensions.

Return type

`torch.FloatTensor`

`get_embeddings()` → `Tuple[numpy.ndarray, numpy.ndarray]`

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (`numpy.ndarray`) – Shape (`num_entities`, `embedding_dim`).
- **relation_embeddings** (`numpy.ndarray`) – Shape (`num_relations`, `embedding_dim`).

```
class dicee.models.IdentityClass (args=None)
```

Bases: `torch.nn.Module`

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (`--normalization None`). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = `None`

__call__ (`x`)

static forward (`x`)

```
dicee.models.quaternion_mul (*, Q_1, Q_2)
```

→ `Tuple[torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor]`

Perform quaternion multiplication :param Q_1: :param Q_2: :return:

```
dicee.models.quaternion_mul_with_unit_norm (*, Q_1, Q_2)
```

```
class dicee.models.QMult (args)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (`h`, `r`, `t`) triples.
- `forward_k_vs_all()` — score a (`h`, `r`) batch against every entity.

Parameters

args (`dict`) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = `'QMult'`

explicit = `True`

quaternion_multiplication_followed_by_inner_product (`h`, `r`, `t`)

Parameters

- **h** – shape: (`*batch_dims`, `dim`) The head representations.
- **r** – shape: (`*batch_dims`, `dim`) The head representations.

- **t** – shape: (**batch_dims*, dim) The tail representations.

Returns

Triple scores.

static quaternion_normalizer (*x*: torch.FloatTensor) → torch.FloatTensor

Normalize the length of relation vectors, if the forward constraint has not been applied yet.

Absolute value of a quaternion

$$|a + bi + cj + dk| = \sqrt{a^2 + b^2 + c^2 + d^2}$$

L2 norm of quaternion vector:

$$\|x\|^2 = \sum_{i=1}^d |x_i|^2 = \sum_{i=1}^d (x_i.re^2 + x_i.im_1^2 + x_i.im_2^2 + x_i.im_3^2)$$

Parameters

x – The vector.

Returns

The normalized vector.

score (*head_ent_emb*: torch.FloatTensor, *rel_ent_emb*: torch.FloatTensor, *tail_ent_emb*: torch.FloatTensor)

k_vs_all_score (*bpe_head_ent_emb*, *bpe_rel_ent_emb*, *E*)

Parameters

- **bpe_head_ent_emb**
- **bpe_rel_ent_emb**
- **E**

forward_k_vs_all (*x*)

Parameters

x

forward_k_vs_sample (*x*, *target_entity_idx*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e., [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,l Entities!)

class dicee.models.ConvQ (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Convolutional Quaternion Knowledge Graph Embeddings

name = 'ConvQ'

entity_embeddings

relation_embeddings

conv2d

fc_num_input

```

    fc1

    bn_conv1

    bn_conv2

    feature_map_dropout

    residual_convolution(Q_1, Q_2)

    forward_triples(indexed_triple: torch.Tensor) → torch.Tensor
        Score a batch of (head, relation, tail) index triples.

        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].

        Returns
            Shape (batch_size,) triple scores.

        Return type
            torch.FloatTensor

    forward_k_vs_all(x: torch.Tensor)
        Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] =>
        [0.0,0.1,...,0.8], shape=> (1, Entities) Given a batch of head entities and relations => shape (size of batch,|
        Entities)

class dicee.models.AConvQ(args)
    Bases: dicee.models.base_model.BaseKGE

    Additive Convolutional Quaternion Knowledge Graph Embeddings

    name = 'AConvQ'

    entity_embeddings

    relation_embeddings

    conv2d

    fc_num_input

    fc1

    bn_conv1

    bn_conv2

    feature_map_dropout

    residual_convolution(Q_1, Q_2)

    forward_triples(indexed_triple: torch.Tensor) → torch.Tensor
        Score a batch of (head, relation, tail) index triples.

        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].

        Returns
            Shape (batch_size,) triple scores.

```

Return type

torch.FloatTensor

forward_k_vs_all (*x*: torch.Tensor)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **|Entities|**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

class dicee.models.**BaseKGE** (*args*: dict)

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (dict) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

`hidden_dropout`

`loss_history = []`

`byte_pair_encoding`

`max_length_subword_tokens`

`block_size`

`forward_byte_pair_encoded_k_vs_all` (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

`forward_byte_pair_encoded_triple` (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*) → *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

`init_params_with_sanity_checking`() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x*: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
y_idx: torch.LongTensor = None) → torch.FloatTensor

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- Tuple (*x*, y_idx) → *forward_k_vs_sample()*
- (batch, 3) tensor → *forward_triples()*
- (batch, 2) tensor → *forward_k_vs_all()*
- BPE triple tensor → *forward_byte_pair_encoded_triple()*
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all()*

Parameters

- **x** (torch.LongTensor or Tuple[torch.LongTensor, torch.LongTensor]) – Either a plain index tensor or a (triple_idx, target_idx) tuple for sample-based labelling.
- **y_idx** (torch.LongTensor, optional) – Target entity indices used by *forward_k_vs_sample()*. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x*: torch.LongTensor) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

- **x** (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (torch.LongTensor) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings() → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (numpy.ndarray) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (numpy.ndarray) – Shape (num_relations, embedding_dim).

```
class dicee.models.IdentityClass (args=None)
```

Bases: torch.nn.Module

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (`--normalization None`). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = None

__call__(x)

static forward(x)

```
dicee.models.octonion_mul(*, O_1, O_2)
```

```
dicee.models.octonion_mul_norm(*, O_1, O_2)
```

```
class dicee.models.OMult (args)
```

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (dict) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'OMult'

static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4, emb_rel_e5, emb_rel_e6, emb_rel_e7)

score(head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor, tail_ent_emb: torch.FloatTensor)

k_vs_all_score(bpe_head_ent_emb, bpe_rel_ent_emb, E)

forward_k_vs_all(*x*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e., [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,| Entities|)

class dicee.models.ConvO(*args: dict*)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'ConvO'

conv2d

fc_num_input

fc1

bn_conv2d

norm_fc1

feature_map_dropout

static octonion_normalizer(*emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4, emb_rel_e5, emb_rel_e6, emb_rel_e7*)

residual_convolution(*O_1, O_2*)

forward_triples(*x: torch.Tensor*) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(*x: torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,| Entities|)

```

class dicee.models.AConvO(args: dict)
    Bases: dicee.models.base_model.BaseKGE
    Additive Convolutional Octonion Knowledge Graph Embeddings
    name = 'AConvO'
    conv2d
    fc_num_input
    fc1
    bn_conv2d
    norm_fc1
    feature_map_dropout
    static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                               emb_rel_e5, emb_rel_e6, emb_rel_e7)
    residual_convolution(O_1, O_2)
    forward_triples(x: torch.Tensor) → torch.Tensor
        Score a batch of (head, relation, tail) index triples.
        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].
        Returns
            Shape (batch_size,) triple scores.
        Return type
            torch.FloatTensor
    forward_k_vs_all(x: torch.Tensor)
        Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] =>
        [0.0,0.1,...,0.8], shape=> (1, Entities) Given a batch of head entities and relations => shape (size of batch,|
        Entities)
class dicee.models.Keci(args)
    Bases: dicee.models.base_model.BaseKGE
    Keci: Knowledge Graph Embedding via Clifford Algebra.
    Embeds entities and relations as multi-vectors in the Clifford algebra  $Cl_{\{p,q\}}(R^d)$  and scores triples via the
    Clifford product. The algebra is parameterised by two non-negative integers  $p$  and  $q$ :
    •  $p = 0, q = 0$  — reduces to a standard bilinear (DistMult-like) model.
    •  $p = 0, q = 1$  — equivalent to ComplEx.
    • Larger  $p$  and  $q$  capture higher-order geometric interactions.
    The embedding dimension must satisfy  $embedding\_dim \% (p + q + 1) == 0$ ; the resulting quotient is stored
    as self.r.
    Parameters
        args (dict) – Configuration dictionary. Recognised keys (beyond those in BaseKGE):  $p$  (int,
        default 0) and  $q$  (int, default 0).

```

References

Demir et al., *Clifford Embeddings — A Generalized Approach for Embedding in Normed Algebras*, ECML 2023.

```
name = 'Keci'
```

```
p
```

```
q
```

```
r
```

```
requires_grad_for_interactions = True
```

```
compute_sigma_pp(hp, rp)
```

Compute $\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{k=i+1}^p (h_{ir_k} - h_{kr_i}) e_i e_k$

σ_{pp} captures the interactions between along p bases For instance, let p e_1, e_2, e_3 , we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

```
results = [] for i in range(p - 1):
```

```
    for k in range(i + 1, p):
```

```
        results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])
```

```
sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

```
compute_sigma_qq(hq, rq)
```

Compute $\sigma_{qq} = \sum_{j=1}^{q-1} \sum_{k=j+1}^q (h_{jr_k} - h_{kr_j}) e_j e_k$ σ_{qq} captures the interactions between along q bases For instance, let q e_1, e_2, e_3 , we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

```
results = [] for j in range(q - 1):
```

```
    for k in range(j + 1, q):
```

```
        results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])
```

```
sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

```
compute_sigma_pq(*, hp, hq, rp, rq)
```

$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_{ir_j} - h_{jr_i}) e_i e_j$

```
results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):
```

```
    for j in range(q):
```

```
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
```

```
print(sigma_pq.shape)
```

apply_coefficients (*hp, hq, rp, rq*)

Multiplying a base vector with its scalar coefficient

clifford_multiplication (*h0, hp, hq, r0, rp, rq*)

Compute our CL multiplication

$$h = h_0 + \sum_{i=1}^p h_i e_i + \sum_{j=p+1}^{p+q} h_j e_j \quad r = r_0 + \sum_{i=1}^p r_i e_i + \sum_{j=p+1}^{p+q} r_j e_j$$

$$e_i^2 = +1 \text{ for } i \leq p, e_j^2 = -1 \text{ for } p < j \leq p+q, e_i e_j = -e_j e_i \text{ for } i < j$$

$$h r = \sigma_0 + \sigma_p + \sigma_q + \sigma_{pp} + \sigma_q + \sigma_{pq} \text{ where}$$

$$(1) \sigma_0 = h_0 r_0 + \sum_{i=1}^p (h_0 r_i - h_i r_0) e_i - \sum_{j=p+1}^{p+q} (h_j r_j) e_j$$

$$(2) \sigma_p = \sum_{i=1}^p (h_0 r_i + h_i r_0) e_i$$

$$(3) \sigma_q = \sum_{j=p+1}^{p+q} (h_0 r_j + h_j r_0) e_j$$

$$(4) \sigma_{pp} = \sum_{i=1}^p \sum_{k=i+1}^p (h_i r_k - h_k r_i) e_i e_k$$

$$(5) \sigma_{qq} = \sum_{j=1}^{p+q} \sum_{k=j+1}^{p+q} (h_j r_k - h_k r_j) e_j e_k$$

$$(6) \sigma_{pq} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

construct_cl_multivector (*x: torch.FloatTensor, r: int, p: int, q: int*)

→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Split a flat embedding vector into the three Clifford components.

Given an embedding *x* of dimension $d = r + r \cdot p + r \cdot q$, returns the scalar part *a0*, the *p*-blade part *ap*, and the *q*-blade part *aq*.

Parameters

- **x** (*torch.FloatTensor*) – Shape (batch_size, d).
- **r** (*int*) – Scalar block size (embedding_dim // (p + q + 1)).
- **p** (*int*) – Number of positive-signature basis elements.
- **q** (*int*) – Number of negative-signature basis elements.

Returns

- **a0** (*torch.FloatTensor*) – Shape (batch_size, r) — scalar (grade-0) part.
- **ap** (*torch.FloatTensor*) – Shape (batch_size, r, p) — positive-blade part.
- **aq** (*torch.FloatTensor*) – Shape (batch_size, r, q) — negative-blade part.

forward_k_vs_with_explicit (*x: torch.Tensor*) → torch.FloatTensor

KvsAll scoring using an explicit loop over $\sigma_{pp}/\sigma_{qq}/\sigma_{pq}$ terms.

Functionally equivalent to `forward_k_vs_all()` but computes the higher-order interaction terms (σ_{pp} , σ_{qq} , σ_{pq}) with explicit nested loops rather than einsum contractions. Kept for reference and correctness verification.

Parameters

x (*torch.Tensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

k_vs_all_score (*bpe_head_ent_emb*: torch.FloatTensor, *bpe_rel_ent_emb*: torch.FloatTensor, *E*: torch.FloatTensor) → torch.FloatTensor

Compute Clifford-product scores for a head/relation batch vs. all entities.

Decomposes the head-entity and relation embeddings into Clifford multi-vectors, performs the $Cl_{\{p,q\}}$ product, and inner-products the result against the entity embedding matrix *E*.

Parameters

- **bpe_head_ent_emb** (torch.FloatTensor) – Head-entity embeddings, shape (batch_size, embedding_dim).
- **bpe_rel_ent_emb** (torch.FloatTensor) – Relation embeddings, shape (batch_size, embedding_dim).
- **E** (torch.FloatTensor) – All entity embeddings, shape (num_entities, embedding_dim).

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: torch.Tensor) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations $\mathbb{R}^d$.
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— *x*: torch.LongTensor with (n,2) shape :rtype: torch.FloatTensor with (n, **IEI**) shape

construct_batch_selected_cl_multivector (*x*: torch.FloatTensor, *r*: int, *p*: int, *q*: int) → tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Split a batched, *k*-selected embedding tensor into Clifford components.

A variant of `construct_cl_multivector()` for tensors that have an extra *k* dimension (e.g. when scoring against *k* sampled targets).

Parameters

- **x** (torch.FloatTensor) – Shape (batch_size, *k*, *d*).
- **r** (int) – Scalar block size.
- **p** (int) – Number of positive-signature basis elements.
- **q** (int) – Number of negative-signature basis elements.

Returns

- **a0** (torch.FloatTensor) – Shape (batch_size, *k*, *r*).
- **ap** (torch.FloatTensor) – Shape (batch_size, *k*, *r*, *p*).
- **aq** (torch.FloatTensor) – Shape (batch_size, *k*, *r*, *q*).

forward_k_vs_sample (*x*: *torch.LongTensor*, *target_entity_idx*: *torch.LongTensor*) → *torch.FloatTensor*

Parameter

x: *torch.LongTensor* with (n,2) shape

target_entity_idx: *torch.LongTensor* with (n, k) shape k denotes the selected number of examples.

rtype

torch.FloatTensor with (n, k) shape

score (*h*, *r*, *t*)

forward_triples (*x*: *torch.Tensor*) → *torch.FloatTensor*

Parameter

x: *torch.LongTensor* with (n,3) shape

rtype

torch.FloatTensor with (n) shape

class *dicee.models.CKeci* (*args*)

Bases: *Keci*

Without learning dimension scaling

name = 'CKeci'

requires_grad_for_interactions = False

class *dicee.models.DeCaL* (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches *forward()* calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (*h*, *r*, *t*) triples.
- *forward_k_vs_all()* — score a (*h*, *r*) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by *vars(argparse.Namespace)*. Required keys: *embedding_dim*, *num_entities*, *num_relations*, *learning_rate* (or *lr*), *optim*, *scoring_technique*.

name = 'DeCaL'

entity_embeddings

relation_embeddings

p

q

r

re

forward_triples (*x*: torch.Tensor) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,) shape

rtype

torch.FloatTensor with (n) shape

cl_pqr (*a*: torch.tensor) → torch.tensor

Input: tensor(batch_size, emb_dim) → output: tensor with 1+p+q+r components with size (batch_size, emb_dim/(1+p+q+r)) each.

1) takes a tensor of size (batch_size, emb_dim), split it into 1 + p + q + r components, hence 1+p+q+r must be a divisor of the emb_dim. 2) Return a list of the 1+p+q+r components vectors, each are tensors of size (batch_size, emb_dim/(1+p+q+r))

compute_sigmas_single (*list_h_emb, list_r_emb, list_t_emb*)

here we compute all the sums with no others vectors interaction taken with the scalar product with t, that is,

$$s0 = h_0 r_0 t_0 s1 = \sum_{i=1}^p h_i r_i t_0 s2 = \sum_{j=p+1}^{p+q} h_j r_j t_0 s3 = \sum_{i=1}^q (h_0 r_i t_i + h_i r_0 t_i) s4 = \sum_{i=p+1}^{p+q} (h_0 r_i t_i + h_i r_0 t_i) s5 = \sum_{i=p+q+1}^{p+q+r} (h_0 r_i t_i + h_i r_0 t_i)$$

and return:

$$sigma_0 t = \sigma_0 \cdot t_0 = s0 + s1 - s2 s3, s4 \text{ and } s5$$

compute_sigmas_multivect (*list_h_emb, list_r_emb*)

Here we compute and return all the sums with vectors interaction for the same and different bases.

For same bases vectors interaction we have

$$\sigma_p p = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (h_i r_{i'} - h_{i'} r_i) (\text{model the interactions between } e_i \text{ and } e_{i'} \text{ for } 1 \leq i, i' \leq p) \sigma_q q = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (h_j r_{j'} - h_{j'} r_j)$$

For different base vector interactions, we have

$$\sigma_p q = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) (\text{interactions between } e_i \text{ and } e_j \text{ for } 1 \leq i \leq p \text{ and } p+1 \leq j \leq p+q) \sigma_p r = \sum_{i=1}^p$$

forward_k_vs_all (*x*: torch.Tensor) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q,r\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functons are identical Parameter ——— *x*: torch.LongTensor with (n,) shape :rtype: torch.FloatTensor with (n, **IEI**) shape

apply_coefficients (*h0, hp, hq, hk, r0, rp, rq, rk*)

Multiplying a base vector with its scalar coefficient

construct_cl_multivector (*x: torch.FloatTensor, re: int, p: int, q: int, r: int*)
 → tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Construct a batch of multivectors $\text{Cl}_{\{p,q,r\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,d) shape

returns

- **a0** (*torch.FloatTensor*)
- **ap** (*torch.FloatTensor*)
- **aq** (*torch.FloatTensor*)
- **ar** (*torch.FloatTensor*)

compute_sigma_pp (*hp, rp*)

Compute .. math:

$$\backslash \sigma_{\{p,p\}}^{*} = \backslash \sum_{i=1}^{p-1} \backslash \sum_{i'=i+1}^p (x_{i'} y_{i'} - x_i y_i)$$

$\sigma_{\{pp\}}$ captures the interactions between along p bases For instance, let p e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3 , and e_2 e_3 This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_qq (*hq, rq*)

Compute

$$\sigma_{q,q}^{*} = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (x_j y_{j'} - x_{j'} y_j) Eq.16$$

$\sigma_{\{q\}}$ captures the interactions between along q bases For instance, let q e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3 , and e_2 e_3 This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_rr(hk, rk)

$$\sigma_{r,r}^* = \sum_{k=p+q+1}^{p+q+r-1} \sum_{k'=k+1}^p (x_k y_{k'} - x_{k'} y_k)$$

compute_sigma_pq(*, hp, hq, rp, rq)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

compute_sigma_pr(*, hp, hk, rp, rk)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

compute_sigma_qr(*, hq, hk, rq, rk)

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

class dicee.models.**KeciTransformer**(args)

Bases: [Keci](#)

Keci with Transformer architecture.

Concatenates h0, hp, hq, r0, rp, rq into a single embedding vector and processes through transformer.

name = 'KeciTransformer'

use_clifford_mul

seq_len = 1

transformer

lm_head

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

Kvsall training

Parameter

x: torch.LongTensor with (n,2) shape

rtype

torch.FloatTensor with (n, **IEI**) shape

class dicee.models.**BaseKGE** (*args: dict*)

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

```

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all(x: torch.LongTensor) → torch.FloatTensor

```

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

```

forward_byte_pair_encoded_triple(x: Tuple[torch.LongTensor, torch.LongTensor])
→ torch.FloatTensor

```

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x*: `torch.LongTensor` | `Tuple[torch.LongTensor, torch.LongTensor]`,
 y_idx: `torch.LongTensor` = None) → `torch.FloatTensor`

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- `Tuple (x, y_idx)` → `forward_k_vs_sample()`
- `(batch, 3)` tensor → `forward_triples()`
- `(batch, 2)` tensor → `forward_k_vs_all()`
- BPE triple tensor → `forward_byte_pair_encoded_triple()`
- BPE pair tensor → `forward_byte_pair_encoded_k_vs_all()`

Parameters

- **x** (`torch.LongTensor` or `Tuple[torch.LongTensor, torch.LongTensor]`) – Either a plain index tensor or a `(triple_idx, target_idx)` tuple for sample-based labelling.
- **y_idx** (`torch.LongTensor`, optional) – Target entity indices used by `forward_k_vs_sample()`. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.LongTensor`) → `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

- **x** (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings() → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

class dicee.models.**PykeenKGE** (*args: dict*)

Bases: *dicee.models.base_model.BaseKGE*

A class for using knowledge graph embedding models implemented in Pykeen

Notes: Pykeen_DistMult: C Pykeen_ComplEx: Pykeen_QuatE: Pykeen_MuRE: Pykeen_CP: Pykeen_HolE: Pykeen_HolE: Pykeen_HolE: Pykeen_TransD: Pykeen_TransE: Pykeen_TransF: Pykeen_TransH: Pykeen_TransR:

model_kwargs

name

model

loss_history = []

args

entity_embeddings = None

relation_embeddings = None

forward_k_vs_all (*x: torch.LongTensor*)

=> Explicit version by this we can apply bn and dropout

(1) Retrieve embeddings of heads and relations + apply Dropout & Normalization if given. h, r = self.get_head_relation_representation(x) # (2) Reshape (1). if self.last_dim > 0:

h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim, self.last_dim)

(3) Reshape all entities. if self.last_dim > 0:

t = self.entity_embeddings.weight.reshape(self.num_entities, self.embedding_dim, self.last_dim)

else:

t = self.entity_embeddings.weight

```

# (4) Call the score_t from interactions to generate triple scores. return self.interaction.score_t(h=h, r=r,
all_entities=t, slice_size=1)

forward_triples (x: torch.LongTensor) → torch.FloatTensor
# => Explicit version by this we can apply bn and dropout

# (1) Retrieve embeddings of heads, relations and tails and apply Dropout & Normalization if given. h, r, t =
self.get_triple_representation(x) # (2) Reshape (1). if self.last_dim > 0:
    h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim,
    self.last_dim) t = t.reshape(len(x), self.embedding_dim, self.last_dim)

# (3) Compute the triple score return self.interaction.score(h=h, r=r, t=t, slice_size=None, slice_dim=0)

abstractmethod forward_k_vs_sample (x: torch.LongTensor, target_entity_idx)
Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must over-
ride this method.

Returns
    Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type
    torch.FloatTensor

class dicee.models.BaseKGE (args: dict)
Bases: BaseKGELightning

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from BaseKGELightning and adds the embedding tables, normalisation /
dropout layers, and the routing logic that dispatches forward() calls to the appropriate scoring method.

Sub-classes must implement at minimum:



- forward_triples() — score a batch of (h, r, t) triples.
- forward_k_vs_all() — score a (h, r) batch against every entity.

Parameters
    args (dict) – Flat configuration dictionary produced by vars(argparse.Namespace). Re-
    quired keys: embedding_dim, num_entities, num_relations, learning_rate (or lr),
    optim, scoring_technique.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

```

```

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

```

forward_byte_pair_encoded_k_vs_all (x : *torch.LongTensor*) $\rightarrow$ *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

$\mathbf{x}$ (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple*[*torch.LongTensor*, *torch.LongTensor*])
→ *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from *self.args* with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the *args* dict. Falls back to sensible defaults for any missing key so that minimal *args* dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple*[*torch.LongTensor*, *torch.LongTensor*],
y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample*()
- (batch, 3) tensor → *forward_triples*()
- (batch, 2) tensor → *forward_k_vs_all*()
- BPE triple tensor → *forward_byte_pair_encoded_triple*()
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all*()

Parameters

• **x** (*torch.LongTensor* or *Tuple*[*torch.LongTensor*, *torch.LongTensor*]) – Either a plain index tensor or a (triple_idx, target_idx) tuple for sample-based labelling.

• **y_idx** (*torch.LongTensor*, optional) – Target entity indices used by *forward_k_vs_sample*(). Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.LongTensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

class dicee.models.FMult (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Learning Knowledge Neural Graphs

name = 'FMult'

entity_embeddings

relation_embeddings

k

num_sample = 50

gamma

roots

weights

compute_func (*weights: torch.FloatTensor, x*) → torch.FloatTensor

chain_func (*weights, x: torch.FloatTensor*)

forward_triples (*idx_triple: torch.Tensor*) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

```
class dicee.models.GFMult(args)
```

Bases: *dicee.models.base_model.BaseKGE*

Learning Knowledge Neural Graphs

name = 'GFMult'

entity_embeddings

relation_embeddings

k

num_sample = 250

roots

weights

compute_func (*weights: torch.FloatTensor, x*) → torch.FloatTensor

chain_func (*weights, x: torch.FloatTensor*)

forward_triples (*idx_triple: torch.Tensor*) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

```
class dicee.models.FMult2(args)
```

Bases: *dicee.models.base_model.BaseKGE*

Learning Knowledge Neural Graphs

name = 'FMult2'

n_layers = 3

k

n = 50


```

score_func = 'compositional'

discrete_points

entity_embeddings

relation_embeddings

build_func (Vec)

build_chain_funcs (list_Vec)

compute_func (W, b, x) → torch.FloatTensor

function (list_W, list_b)

trapezoid (list_W, list_b)

forward_triples (idx_triple: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

```

class dicee.models.LFMult1 (args)

Bases: *dicee.models.base_model.BaseKGE*

Embedding with trigonometric functions. We represent all entities and relations in the complex number space as: $f(x) = \sum_{k=0}^{d-1} w_k e^{kix}$. and use the three different scoring function as in the paper to evaluate the score

```

name = 'LFMult1'

entity_embeddings

relation_embeddings

forward_triples (idx_triple)
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

tri_score (h, r, t)

vtp_score (h, r, t)

```

```
class dicee.models.LFMult (args)
```

Bases: *dicee.models.base_model.BaseKGE*

Embedding with polynomial functions. We represent all entities and relations in the polynomial space as: $f(x) = \sum_{i=0}^{d-1} a_k x^i$ and use the three different scoring function as in the paper to evaluate the score. We also consider combining with Neural Networks.

name = 'LFMult'

entity_embeddings

relation_embeddings

degree

m

x_values

forward_triples(idx_triple)

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

construct_multi_coeff(x)

poly_NN(x, coefh, coefr, coeft)

Constructing a 2 layers NN to represent the embeddings. $h = \sigma(w_h^T x + b_h)$, $r = \sigma(w_r^T x + b_r)$, $t = \sigma(w_t^T x + b_t)$

linear(x, w, b)

scalar_batch_NN(a, b, c)

element wise multiplication between a,b and c: Inputs : a, b, c ==> torch.tensor of size batch_size x m x d
Output : a tensor of size batch_size x d

tri_score(coeff_h, coeff_r, coeff_t)

this part implement the trilinear scoring techniques:

$score(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \frac{a_i b_j c_k}{1+(i+j+k)d}$

1. generate the range for i,j and k from [0 d-1]
2. perform $\frac{a_i b_j c_k}{1+(i+j+k)d}$ in parallel for every batch
3. take the sum over each batch

vtp_score(h, r, t)

this part implement the vector triple product scoring techniques:

$score(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \frac{a_i c_j b_k - b_i c_j a_k}{(1+(i+j)d)(1+k)}$

1. generate the range for i,j and k from [0 d-1]

2. Compute the first and second terms of the sum
3. Multiply with then denominator and take the sum
4. take the sum over each batch

comp_func (*h, r, t*)

this part implement the function composition scoring techniques: i.e. score = <hor, t>

polynomial (*coeff, x, degree*)

This function takes a matrix tensor of coefficients (*coeff*), a tensor vector of points *x* and range of integer $[0, 1, \dots, d]$ and return a vector tensor ($\text{coeff}[0][0] + \text{coeff}[0][1]x + \dots + \text{coeff}[0][d]x^d$,

$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d$)

pop (*coeff, x, degree*)

This function allow us to evaluate the composition of two polynomes without for loops :) it takes a matrix tensor of coefficients (*coeff*), a matrix tensor of points *x* and range of integer $[0, 1, \dots, d]$

and return a tensor ($\text{coeff}[0][0] + \text{coeff}[0][1]x + \dots + \text{coeff}[0][d]x^d$,

$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d$)

class `dicee.models.DualE` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Dual Quaternion Knowledge Graph Embeddings (<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

name = 'DualE'

entity_embeddings

relation_embeddings

num_ent = None

kvsall_score (*e_1_h, e_2_h, e_3_h, e_4_h, e_5_h, e_6_h, e_7_h, e_8_h, e_1_t, e_2_t, e_3_t, e_4_t, e_5_t, e_6_t, e_7_t, e_8_t, r_1, r_2, r_3, r_4, r_5, r_6, r_7, r_8*) → torch.tensor

KvsAll scoring function

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_triples (*idx_triple: torch.tensor*) → torch.tensor

Negative Sampling forward pass:

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_k_vs_all (*x*)
KvsAll forward pass

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

T (*x: torch.tensor*) → torch.tensor
Transpose function

Input: Tensor with shape (nxm) Output: Tensor with shape (mxn)

dicee.query_generator

Classes

QueryGenerator

Module Contents

```
class dicee.query_generator.QueryGenerator(train_path: str, val_path: str, test_path: str,
    ent2id: Dict = None, rel2id: Dict = None, seed: int = 1, gen_valid: bool = False,
    gen_test: bool = True)

    train_path
    val_path
    test_path
    gen_valid = False
    gen_test = True
    seed = 1
    max_ans_num = 1000000.0
    mode
    ent2id = None
    rel2id: Dict = None
    ent_in: Dict
    ent_out: Dict
```

query_name_to_struct

list2tuple (*list_data*)

tuple2list (*x: List | Tuple*) → List | Tuple
Convert a nested tuple to a nested list.

set_global_seed (*seed: int*)
Set seed

construct_graph (*paths: List[str]*) → Tuple[Dict, Dict]
Construct graph from triples Returns dicts with incoming and outgoing edges

fill_query (*query_structure: List[str | List], ent_in: Dict, ent_out: Dict, answer: int*) → bool
Private method for fill_query logic.

achieve_answer (*query: List[str | List], ent_in: Dict, ent_out: Dict*) → set
Private method for achieve_answer logic. @TODO: Document the code

write_links (*ent_out, small_ent_out*)

ground_queries (*query_structure: List[str | List], ent_in: Dict, ent_out: Dict, small_ent_in: Dict, small_ent_out: Dict, gen_num: int, query_name: str*)
Generating queries and achieving answers

unmap (*query_type, queries, tp_answers, fp_answers, fn_answers*)

unmap_query (*query_structure, query, id2ent, id2rel*)

generate_queries (*query_struct: List, gen_num: int, query_type: str*)
Passing incoming and outgoing edges to ground queries depending on mode [train valid or text] and getting queries and answers in return @ TODO: create a class for each single query struct

save_queries (*query_type: str, gen_num: int, save_path: str*)

abstractmethod load_queries (*path*)

get_queries (*query_type: str, gen_num: int*)

static save_queries_and_answers (*path: str, data: List[Tuple[str, Tuple[collections.defaultdict]]]*)
→ None
Save Queries into Disk

static load_queries_and_answers (*path: str*) → List[Tuple[str, Tuple[collections.defaultdict]]]
Load Queries from Disk to Memory

dicee.read_preprocess_save_load_kg

Submodules

dicee.read_preprocess_save_load_kg.preprocess

Classes

Module Contents

class `dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG(kg)`

Preprocess the data in memory

kg

start () → None

Preprocess train, valid and test datasets stored in knowledge graph instance

Parameter

rtype

None

preprocess_with_byte_pair_encoding ()

preprocess_with_byte_pair_encoding_with_padding () → None

Preprocess with byte pair encoding and add padding

preprocess_with_pandas () → None

Preprocess with pandas: add reciprocal triples, construct vocabulary, and index datasets

preprocess_with_polars () → None

Preprocess with polars: add reciprocal triples and create indexed datasets

sequential_vocabulary_construction () → None

(1) Read input data into memory

(2) Remove triples with a condition

(3) **Serialize vocabularies in a pandas dataframe where**

=> the index is integer and => a single column is string (e.g. URI)

dicee.read_preprocess_save_load_kg.read_from_disk

Classes

ReadFromDisk

Read the data from disk into memory

Module Contents

class `dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk(kg)`

Read the data from disk into memory

kg

start () → None

Read a knowledge graph from disk into memory

Data will be available at the `train_set`, `test_set`, `valid_set` attributes.

Parameter

None

rtype

None

`add_noisy_triples_into_training()`

`dicee.read_preprocess_save_load_kg.save_load_disk`

Classes

LoadSaveToDisk

Module Contents

`class dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk(kg)`

`kg`

`save()`

`load()`

`dicee.read_preprocess_save_load_kg.util`

Functions

<code>polars_dataframe_indexer(→ polars.DataFrame)</code>	Replaces 'subject', 'relation', and 'object' columns in the input Polars DataFrame with their corresponding index values
<code>pandas_dataframe_indexer(→ pandas.DataFrame)</code>	Replaces 'subject', 'relation', and 'object' columns in the input Pandas DataFrame with their corresponding index values
<code>apply_reciprocal_or_noise(add_reciprocal, eval_model)</code>	Add reciprocal triples if conditions are met
<code>timeit(func)</code>	
<code>read_with_polars(→ polars.DataFrame)</code>	Load and Preprocess via Polars
<code>read_with_pandas(data_path[, read_only_few, ...])</code>	Load and Preprocess via Pandas
<code>read_from_disk(→ Tuple[polars.DataFrame, pandas.DataFrame])</code>	
<code>count_triples(→ int)</code>	Returns the total number of triples in the triple store.
<code>fetch_worker(endpoint, offsets, chunk_size, ...)</code>	Worker process: fetch assigned chunks and save to disk with per-worker tqdm.
<code>read_from_triple_store_with_polars(endpoint[, ...])</code>	Main function to read all triples in parallel, save as Parquet, and load into Polars dataframe.
<code>read_from_triple_store_with_pandas([endpoint])</code>	Read triples from triple store into pandas dataframe
<code>get_er_vocab(data[, file_path])</code>	
<code>get_re_vocab(data[, file_path])</code>	
<code>get_ee_vocab(data[, file_path])</code>	
<code>create_constraints(triples[, file_path])</code>	

continues on next page

Table 46 – continued from previous page

<code>load_with_pandas(→ None)</code>	Deserialize data
<code>save_numpy_ndarray(*, data, file_path)</code>	
<code>load_numpy_ndarray(*, file_path)</code>	
<code>save_pickle(*, data[, file_path])</code>	
<code>load_pickle(*[, file_path])</code>	
<code>create_reciprocal_triples(x)</code>	Add inverse triples into dask dataframe
<code>dataset_sanity_checking(→ None)</code>	

Module Contents

```
dicee.read_preprocess_save_load_kg.util.polars_dataframe_indexer(
    df_polars: polars.DataFrame, idx_entity: polars.DataFrame, idx_relation: polars.DataFrame)
    → polars.DataFrame
```

Replaces ‘subject’, ‘relation’, and ‘object’ columns in the input Polars DataFrame with their corresponding index values from the entity and relation index DataFrames.

This function processes the DataFrame in three main steps: 1. Replace the ‘relation’ values with the corresponding index from *idx_relation*. 2. Replace the ‘subject’ values with the corresponding index from *idx_entity*. 3. Replace the ‘object’ values with the corresponding index from *idx_entity*.

Parameters:

df_polars

[polars.DataFrame] The input Polars DataFrame containing columns: ‘subject’, ‘relation’, and ‘object’.

idx_entity

[polars.DataFrame] A Polars DataFrame that contains the mapping between entity names and their corresponding indices. Must have columns: ‘entity’ and ‘index’.

idx_relation

[polars.DataFrame] A Polars DataFrame that contains the mapping between relation names and their corresponding indices. Must have columns: ‘relation’ and ‘index’.

Returns:

polars.DataFrame

A DataFrame with the ‘subject’, ‘relation’, and ‘object’ columns replaced by their corresponding indices.

Example Usage:

```
>>> df_polars = pl.DataFrame({
    "subject": ["Alice", "Bob", "Charlie"],
    "relation": ["knows", "works_with", "lives_in"],
    "object": ["Dave", "Eve", "Frank"]
})
>>> idx_entity = pl.DataFrame({
    "entity": ["Alice", "Bob", "Charlie", "Dave", "Eve", "Frank"],
    "index": [0, 1, 2, 3, 4, 5]
})
>>> idx_relation = pl.DataFrame({
    "relation": ["knows", "works_with", "lives_in"],
    "index": [0, 1, 2]
```

(continues on next page)


```

    })
>>> polars_dataframe_indexer(df_polars, idx_entity, idx_relation)

```

Steps:

1. Join the input DataFrame *df_polars* on the 'relation' column with *idx_relation* to replace the relations with their indices.
2. Join on 'subject' to replace it with the corresponding entity index using a left join on *idx_entity*.
3. Join on 'object' to replace it with the corresponding entity index using a left join on *idx_entity*.
4. Select only the 'subject', 'relation', and 'object' columns to return the final result.

```

dicee.read_preprocess_save_load_kg.util.pandas_dataframe_indexer(
    df_pandas: pandas.DataFrame, idx_entity: pandas.DataFrame, idx_relation: pandas.DataFrame)
    → pandas.DataFrame

```

Replaces 'subject', 'relation', and 'object' columns in the input Pandas DataFrame with their corresponding index values from the entity and relation index DataFrames.

Parameters:**df_pandas**

[pd.DataFrame] The input Pandas DataFrame containing columns: 'subject', 'relation', and 'object'.

idx_entity

[pd.DataFrame] A Pandas DataFrame that contains the mapping between entity names and their corresponding indices. Must have columns: 'entity' and 'index'.

idx_relation

[pd.DataFrame] A Pandas DataFrame that contains the mapping between relation names and their corresponding indices. Must have columns: 'relation' and 'index'.

Returns:**pd.DataFrame**

A DataFrame with the 'subject', 'relation', and 'object' columns replaced by their corresponding indices.

```

dicee.read_preprocess_save_load_kg.util.apply_reciprocal_or_noise(add_reciprocal: bool,
    eval_model: str, df: object = None, info: str = None)

```

Add reciprocal triples if conditions are met

```

dicee.read_preprocess_save_load_kg.util.timeit(func)

```

```

dicee.read_preprocess_save_load_kg.util.read_with_polars(data_path,
    read_only_few: int = None, sample_triples_ratio: float = None, separator: str = None)
    → polars.DataFrame

```

Load and Preprocess via Polars

```

dicee.read_preprocess_save_load_kg.util.read_with_pandas(data_path,
    read_only_few: int = None, sample_triples_ratio: float = None, separator: str = None)

```

Load and Preprocess via Pandas

`dicee.read_preprocess_save_load_kg.util.read_from_disk(data_path: str, read_only_few: int = None, sample_triples_ratio: float = None, backend: str = None, separator: str = None) → Tuple[polars.DataFrame, pandas.DataFrame]`

`dicee.read_preprocess_save_load_kg.util.count_triples(endpoint: str) → int`

Returns the total number of triples in the triple store.

`dicee.read_preprocess_save_load_kg.util.fetch_worker(endpoint: str, offsets: list[int], chunk_size: int, output_dir: str, worker_id: int)`

Worker process: fetch assigned chunks and save to disk with per-worker tqdm.

`dicee.read_preprocess_save_load_kg.util.read_from_triple_store_with_polars(endpoint: str, chunk_size: int = 500000, output_dir: str = 'triples_parquet')`

Main function to read all triples in parallel, save as Parquet, and load into Polars dataframe.

`dicee.read_preprocess_save_load_kg.util.read_from_triple_store_with_pandas(endpoint: str = None)`

Read triples from triple store into pandas dataframe

`dicee.read_preprocess_save_load_kg.util.get_er_vocab(data, file_path: str = None)`

`dicee.read_preprocess_save_load_kg.util.get_re_vocab(data, file_path: str = None)`

`dicee.read_preprocess_save_load_kg.util.get_ee_vocab(data, file_path: str = None)`

`dicee.read_preprocess_save_load_kg.util.create_constraints(triples, file_path: str = None)`

(1) Extract domains and ranges of relations

(2) Store a mapping from relations to entities that are outside of the domain and range. Create constrained entities based on the range of relations :param triples: :return: Tuple[dict, dict]

`dicee.read_preprocess_save_load_kg.util.load_with_pandas(self) → None`

Deserialize data

`dicee.read_preprocess_save_load_kg.util.save_numpy_ndarray(*, data: numpy.ndarray, file_path: str)`

`dicee.read_preprocess_save_load_kg.util.load_numpy_ndarray(*, file_path: str)`

`dicee.read_preprocess_save_load_kg.util.save_pickle(*, data: object, file_path=str)`

`dicee.read_preprocess_save_load_kg.util.load_pickle(*, file_path=str)`

`dicee.read_preprocess_save_load_kg.util.create_recipriocal_triples(x)`

Add inverse triples into dask dataframe :param x: :return:

`dicee.read_preprocess_save_load_kg.util.dataset_sanity_checking(train_set: numpy.ndarray, num_entities: int, num_relations: int) → None`

Parameters

- **train_set**
- **num_entities**
- **num_relations**

Returns

Classes

<i>PreprocessKG</i>	Preprocess the data in memory
<i>LoadSaveToDisk</i>	
<i>ReadFromDisk</i>	Read the data from disk into memory

Package Contents

class dicee.read_preprocess_save_load_kg.**PreprocessKG** (*kg*)

Preprocess the data in memory

kg

start () → None

Preprocess train, valid and test datasets stored in knowledge graph instance

Parameter

rtype

None

preprocess_with_byte_pair_encoding ()

preprocess_with_byte_pair_encoding_with_padding () → None

Preprocess with byte pair encoding and add padding

preprocess_with_pandas () → None

Preprocess with pandas: add reciprocal triples, construct vocabulary, and index datasets

preprocess_with_polars () → None

Preprocess with polars: add reciprocal triples and create indexed datasets

sequential_vocabulary_construction () → None

(1) Read input data into memory

(2) Remove triples with a condition

(3) **Serialize vocabularies in a pandas dataframe where**

=> the index is integer and => a single column is string (e.g. URI)

class dicee.read_preprocess_save_load_kg.**LoadSaveToDisk** (*kg*)

kg

save ()

load ()

class dicee.read_preprocess_save_load_kg.**ReadFromDisk** (*kg*)

Read the data from disk into memory

kg

start () → None

Read a knowledge graph from disk into memory

Data will be available at the train_set, test_set, valid_set attributes.

Parameter

None

rtype

None

`add_noisy_triples_into_training()`

dicee.sanity_checkers

Functions

<code>is_sparql_endpoint_alive([sparql_endpoint])</code>	
<code>validate_knowledge_graph(args)</code>	Validating the source of knowledge graph
<code>sanity_checking_with_arguments(args)</code>	
<code>sanity_check_callback_args(args)</code>	Perform sanity checks on callback-related arguments.

Module Contents

`dicee.sanity_checkers.is_sparql_endpoint_alive(sparql_endpoint: str = None)`

`dicee.sanity_checkers.validate_knowledge_graph(args)`

Validating the source of knowledge graph

`dicee.sanity_checkers.sanity_checking_with_arguments(args)`

`dicee.sanity_checkers.sanity_check_callback_args(args)`

Perform sanity checks on callback-related arguments.

dicee.scripts

Submodules

dicee.scripts.index_serve

```
$ docker pull qdrant/qdrant && docker run -p 6333:6333 -p 6334:6334 -v $(pwd)/qdrant_storage:/qdrant/storage:z  
qdrant/qdrant $ dicee_vector_db -index -serve -path CountryEmbeddings -collection "countries_vdb"
```

Attributes

<code>app</code>
<code>neural_searcher</code>

Classes

<code>NeuralSearcher</code>
<code>StringListRequest</code>

Functions

```
get_default_arguments()
index(args)
root()
search_embeddings(q)
retrieve_embeddings(q)
search_embeddings_batch(request)
serve(args)
main()
```

Module Contents

```
dicee.scripts.index_serve.get_default_arguments ()

dicee.scripts.index_serve.index (args)

dicee.scripts.index_serve.app

dicee.scripts.index_serve.neural_searcher = None

class dicee.scripts.index_serve.NeuralSearcher (args)

    collection_name

    entity_to_idx = None

    qdrant_client

    topk = 5

    retrieve_embedding (entity: str = None, entities: List[str] = None) → List

    search (entity: str)

async dicee.scripts.index_serve.root ()

async dicee.scripts.index_serve.search_embeddings (q: str)

async dicee.scripts.index_serve.retrieve_embeddings (q: str)

class dicee.scripts.index_serve.StringListRequest
    Bases: pydantic.BaseModel

    queries: List[str]

    reducer: str | None = None

async dicee.scripts.index_serve.search_embeddings_batch (request: StringListRequest)

dicee.scripts.index_serve.serve (args)

dicee.scripts.index_serve.main ()
```

dicee.scripts.run

Functions

<code>get_default_arguments([description])</code>	Extends lightning Trainer's arguments with ours
<code>main()</code>	

Module Contents

`dicee.scripts.run.get_default_arguments (description=None)`

Extends lightning Trainer's arguments with ours

`dicee.scripts.run.main()`

dicee.static_funcs

Static utility functions for DICE embeddings.

This module provides utility functions for model initialization, data loading, serialization, and various helper operations.

Attributes

<code>MODEL_REGISTRY</code>

Functions

<code>create_recipriocal_triples(→ pandas.DataFrame)</code>	Add inverse triples to a DataFrame.
<code>get_er_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build entity-relation to tail vocabulary.
<code>get_re_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build relation-entity (tail) to head vocabulary.
<code>get_ee_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build entity-entity to relation vocabulary.
<code>timeit(→ Callable)</code>	Decorator to measure and print execution time and memory usage.
<code>setup_distributed_training(→ Dict[str, Union[bool, int]])</code>	Resolve distributed-training state and initialize custom DDP when needed.
<code>save_pickle(→ None)</code>	Save data to a pickle file.
<code>load_pickle(→ object)</code>	Load data from a pickle file.
<code>load_term_mapping(→ Union[dict, polars.DataFrame])</code>	Load term-to-index mapping from pickle or CSV file.
<code>select_model(args[, is_continual_training, storage_path])</code>	
<code>load_model(→ Tuple[object, Tuple[dict, dict]])</code>	Load weights and initialize pytorch module from namespace arguments
<code>load_model_ensemble(...)</code>	Construct Ensemble Of weights and initialize pytorch module from namespace arguments
<code>save_numpy_ndarray(*, data, file_path)</code>	
<code>numpy_data_type_changer(→ numpy.ndarray)</code>	Detect most efficient data type for a given triples
<code>save_checkpoint_model(→ None)</code>	Store Pytorch model into disk
<code>store(→ None)</code>	

continues on next page

Table 54 – continued from previous page

<code>add_noisy_triples(→ pandas.DataFrame)</code>	Add randomly constructed triples
<code>read_or_load_kg(args, cls)</code>	
<code>intialize_model(...)</code>	Initialize a knowledge graph embedding model.
<code>load_json(→ Dict)</code>	Load JSON file into a dictionary.
<code>save_embeddings(→ None)</code>	Save embeddings to a CSV file.
<code>vocab_to_parquet(vocab_to_idx, name, ...)</code>	
<code>create_experiment_folder(→ str)</code>	Create a timestamped experiment folder.
<code>continual_training_setup_executor(→ None)</code>	
<code>exponential_function(→ torch.FloatTensor)</code>	
<code>load_numpy(→ numpy.ndarray)</code>	
<code>evaluate(entity_to_idx, scores, easy_answers, hard_answers)</code>	# @TODO: CD: Renamed this function
<code>download_file(url[, destination_folder])</code>	
<code>download_files_from_url(→ None)</code>	
<code>download_pretrained_model(→ str)</code>	
<code>write_csv_from_model_parallel(path)</code>	Create
<code>from_pretrained_model_write_embeddings_int(→ None)</code>	

Module Contents

`dicee.static_funcs.MODEL_REGISTRY: Dict[str, Tuple[Type, str]]`

`dicee.static_funcs.create_recipriocal_triples(df: pandas.DataFrame) → pandas.DataFrame`

Add inverse triples to a DataFrame.

For each triple (s, p, o), creates an inverse triple (o, p_inverse, s).

Parameters

df – DataFrame with ‘subject’, ‘relation’, and ‘object’ columns.

Returns

DataFrame with original and inverse triples concatenated.

`dicee.static_funcs.get_er_vocab(data: numpy.ndarray, file_path: str | None = None)`

→ Dict[Tuple[int, int], List[int]]

Build entity-relation to tail vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (head, relation) pairs to list of tail entities.

`dicee.static_funcs.get_re_vocab(data: numpy.ndarray, file_path: str | None = None)`

→ Dict[Tuple[int, int], List[int]]

Build relation-entity (tail) to head vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (relation, tail) pairs to list of head entities.

```
dicee.static_funcs.get_ee_vocab(data: numpy.ndarray, file_path: str | None = None)
→ Dict[Tuple[int, int], List[int]]
```

Build entity-entity to relation vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (head, tail) pairs to list of relations.

```
dicee.static_funcs.timeit(func: Callable) → Callable
```

Decorator to measure and print execution time and memory usage.

Parameters

func – Function to be timed.

Returns

Wrapped function that prints timing information.

```
dicee.static_funcs.setup_distributed_training(args) → Dict[str, bool | int]
```

Resolve distributed-training state and initialize custom DDP when needed.

```
dicee.static_funcs.save_pickle(*, data: object | None = None, file_path: str) → None
```

Save data to a pickle file.

Note: Consider using more portable formats (JSON, Parquet) for new code.

Parameters

- **data** – Object to serialize. If None, nothing is saved.
- **file_path** – Path where the pickle file will be saved.

```
dicee.static_funcs.load_pickle(file_path: str) → object
```

Load data from a pickle file.

Note: Consider using more portable formats (JSON, Parquet) for new code.

Parameters

file_path – Path to the pickle file.

Returns

Deserialized object from the pickle file.

```
dicee.static_funcs.load_term_mapping(file_path: str) → dict | polars.DataFrame
```

Load term-to-index mapping from pickle or CSV file.

Attempts to load from pickle first, falls back to CSV if not found.

Parameters

file_path – Base path without extension.

Returns

Dictionary or Polars DataFrame containing the mapping.

```
dicee.static_funcs.select_model(args: dict, is_continual_training: bool = None,
storage_path: str = None)
```


`dicee.static_funcs.load_model (path_of_experiment_folder: str, model_name='model.pt', verbose=0)`
→ Tuple[object, Tuple[dict, dict]]

Load weights and initialize pytorch module from namespace arguments

`dicee.static_funcs.load_model_ensemble (path_of_experiment_folder: str)`
→ Tuple[`dicee.models.base_model.BaseKGE`, Tuple[pandas.DataFrame, pandas.DataFrame]]

Construct Ensemble Of weights and initialize pytorch module from namespace arguments

- (1) Detect models under given path
- (2) Accumulate parameters of detected models
- (3) Normalize parameters
- (4) Insert (3) into model.

`dicee.static_funcs.save_numpy_ndarray (*, data: numpy.ndarray, file_path: str)`

`dicee.static_funcs.numpy_data_type_changer (train_set: numpy.ndarray, num: int)`
→ numpy.ndarray

Detect most efficient data type for a given triples :param train_set: :param num: :return:

`dicee.static_funcs.save_checkpoint_model (model, path: str) → None`

Store Pytorch model into disk

`dicee.static_funcs.store (trained_model, model_name: str = 'model', full_storage_path: str = None, save_embeddings_as_csv=False) → None`

`dicee.static_funcs.add_noisy_triples (train_set: pandas.DataFrame, add_noise_rate: float)`
→ pandas.DataFrame

Add randomly constructed triples :param train_set: :param add_noise_rate: :return:

`dicee.static_funcs.read_or_load_kg (args, cls)`

`dicee.static_funcs.initialize_model (args: Dict, verbose: int = 0)`
→ Tuple[`dicee.models.base_model.BaseKGE`, str]

Initialize a knowledge graph embedding model.

Parameters

- **args** – Dictionary containing model configuration including ‘model’ key.
- **verbose** – Verbosity level. If > 0, prints initialization message.

Returns

Tuple of (initialized model, form of labelling string).

Raises

ValueError – If the model name is not recognized.

`dicee.static_funcs.load_json (path: str) → Dict`

Load JSON file into a dictionary.

Parameters

path – Path to the JSON file.

Returns

Dictionary containing the JSON data.

Raises

- **FileNotFoundError** – If the file does not exist.

- `json.JSONDecodeError` – If the file contains invalid JSON.

`dicee.static_funcs.save_embeddings(embeddings: numpy.ndarray, indexes: List, path: str) → None`

Save embeddings to a CSV file.

Parameters

- **embeddings** – NumPy array of embeddings with shape (n_items, embedding_dim).
- **indexes** – List of index labels (entity/relation names).
- **path** – Output file path.

`dicee.static_funcs.vocab_to_parquet(vocab_to_idx, name, path_for_serialization, print_into)`

`dicee.static_funcs.create_experiment_folder(folder_name: str = 'Experiments') → str`

Create a timestamped experiment folder.

Parameters

folder_name – Base directory name for experiments.

Returns

Full path to the created experiment folder.

`dicee.static_funcs.continual_training_setup_executor(executor) → None`

`dicee.static_funcs.exponential_function(x: numpy.ndarray, lam: float, ascending_order=True) → torch.FloatTensor`

`dicee.static_funcs.load_numpy(path) → numpy.ndarray`

`dicee.static_funcs.evaluate(entity_to_idx, scores, easy_answers, hard_answers)`

@TODO: CD: Renamed this function Evaluate multi hop query answering on different query types

`dicee.static_funcs.download_file(url, destination_folder='.')`

`dicee.static_funcs.download_files_from_url(base_url: str, destination_folder='.') → None`

Parameters

- **base_url** (e.g. <https://files.dice-research.org/projects/DiceEmbeddings/KINSHIP-Keci-dim128-epoch256-KvsAll>)
- **destination_folder** (e.g. `"KINSHIP-Keci-dim128-epoch256-KvsAll"`)

`dicee.static_funcs.download_pretrained_model(url: str) → str`

`dicee.static_funcs.write_csv_from_model_parallel(path: str)`

Create

`dicee.static_funcs.from_pretrained_model_write_embeddings_into_csv(path: str) → None`

`dicee.static_funcs_training`

Training-related static functions.

This module provides backward compatibility by re-exporting evaluation functions from the new `dicee.evaluation` module, along with training utilities.

Deprecated since version Evaluation: functions have moved to `dicee.evaluation`. Use that module for new code. This module will continue to export training utilities.

Functions

<code>evaluate_bpe_lp(→ Dict[str, float])</code>	Evaluate link prediction with BPE-encoded entities.
<code>evaluate_lp(→ Dict[str, float])</code>	Evaluate link prediction with batched processing.
<code>efficient_zero_grad(→ None)</code>	Efficiently zero gradients using <code>parameter.grad = None</code> .
<code>make_iterable_verbose(→ Iterable)</code>	Wrap an iterable with tqdm progress bar if verbose is True.

Module Contents

`dicee.static_funcs_training.evaluate_bpe_lp(model, triple_idx: List[Tuple],
all_bpe_shaped_entities, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List],
info: str = 'Eval Starts') → Dict[str, float]`

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.static_funcs_training.evaluate_lp(model, triple_idx, num_entities: int,
er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List], info: str = 'Eval Starts',
batch_size: int = 128, chunk_size: int = 1000) → Dict[str, float]`

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.
- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.static_funcs_training.efficient_zero_grad(model) → None`

Efficiently zero gradients using `parameter.grad = None`.

This is more efficient than `optimizer.zero_grad()` as it avoids memory operations.

See: https://pytorch.org/tutorials/recipes/recipes/tuning_guide.html

Parameters

model – PyTorch model to zero gradients for.

`dicee.static_funcs_training.make_iterable_verbose(iterable_object: Iterable, verbose: bool, desc: str = 'Default', position: int | None = None, leave: bool = True) → Iterable`

Wrap an iterable with `tqdm` progress bar if `verbose` is `True`.

Parameters

- **iterable_object** – The iterable to potentially wrap.
- **verbose** – Whether to show progress bar.
- **desc** – Description for the progress bar.
- **position** – Position of the progress bar.
- **leave** – Whether to leave the progress bar after completion.

Returns

The original iterable or a `tqdm`-wrapped version.

`dicee.static_preprocess_funcs`

Attributes

`enable_log`

Functions

<code>timeit(func)</code>	
<code>preprocesses_input_args(args)</code>	Sanity Checking in input arguments
<code>create_constraints(→ Tuple[dict, dict, dict, dict])</code>	
<code>get_er_vocab(data)</code>	
<code>get_re_vocab(data)</code>	
<code>get_ee_vocab(data)</code>	
<code>mapping_from_first_two_cols_to_third(train_se</code>	

Module Contents

`dicee.static_preprocess_funcs.enable_log = False`

`dicee.static_preprocess_funcs.timeit(func)`

`dicee.static_preprocess_funcs.preprocesses_input_args(args)`
Sanity Checking in input arguments

`dicee.static_preprocess_funcs.create_constraints(triples: numpy.ndarray)`
→ `Tuple[dict, dict, dict, dict]`

- (1) Extract domains and ranges of relations
- (2) Store a mapping from relations to entities that are outside of the domain and range. Create constraints entities based on the range of relations :param triples: :return:

```
dicee.static_preprocess_funcs.get_er_vocab(data)

dicee.static_preprocess_funcs.get_re_vocab(data)

dicee.static_preprocess_funcs.get_ee_vocab(data)

dicee.static_preprocess_funcs.mapping_from_first_two_cols_to_third(train_set_idx)
```

dicee.trainer

Submodules

dicee.trainer.dice_trainer

DICE Trainer module for knowledge graph embedding training.

Provides the DICE_Trainer class which supports multiple training backends including PyTorch Lightning, DDP, and custom CPU/GPU trainers.

Classes

<code>DICE_Trainer</code>	DICE_Trainer implement
---------------------------	------------------------

Functions

<code>load_term_mapping(→ polars.DataFrame)</code>	Load term-to-index mapping from CSV file.
<code>initialize_trainer(...)</code>	Initialize the appropriate trainer based on configuration.
<code>get_callbacks(→ List)</code>	Create list of callbacks based on configuration.

Module Contents

```
dicee.trainer.dice_trainer.load_term_mapping(file_path: str) → polars.DataFrame
```

Load term-to-index mapping from CSV file.

Parameters

file_path – Base path without extension.

Returns

Polars DataFrame containing the mapping.

```
dicee.trainer.dice_trainer.initialize_trainer(args, callbacks: List)
```

→ `dicee.trainer.torch_trainer.TorchTrainer` | `dicee.trainer.model_parallelism.TensorParallel` | `dicee.trainer.torch_trainer_ddp`

Initialize the appropriate trainer based on configuration.

Parameters

- **args** – Configuration arguments containing trainer type.
- **callbacks** – List of training callbacks.

Returns

Initialized trainer instance.

Raises

AssertionError – If trainer is None after initialization.

`dicee.trainer.dice_trainer.get_callbacks(args) → List`

Create list of callbacks based on configuration.

Parameters

args – Configuration arguments.

Returns

List of callback instances.

```
class dicee.trainer.dice_trainer.DICE_Trainer (args, is_continual_training: bool, storage_path,
        evaluator=None)
```

DICE_Trainer implement

- 1- Pytorch Lightning trainer (<https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html>)
- 2- Multi-GPU Trainer(<https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html>)
- 3- CPU Trainer

args

is_continual_training:bool

storage_path:str

evaluator:

report:dict

report

args

trainer = None

is_continual_training

storage_path

evaluator = None

form_of_labelling = None

continual_start (knowledge_graph)

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

initialize_trainer (*callbacks: List*)

→ lightning.Trainer | *dicee.trainer.model_parallelism.TensorParallel* | *dicee.trainer.torch_trainer.TorchTrainer* | *dicee.t*

Initialize Trainer from input arguments

`initialize_or_load_model()`

`init_data_loader(dataset: torch.utils.data.Dataset) → torch.utils.data.DataLoader`

`init_dataset() → torch.utils.data.Dataset`

`start(knowledge_graph: dicee.knowledge_graph.KG | numpy.memmap)`
→ Tuple[dicee.models.base_model.BaseKGE, str]

Start the training

(1) Initialize Trainer

(2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

`k_fold_cross_validation(dataset) → Tuple[dicee.models.base_model.BaseKGE, str]`

Perform K-fold Cross-Validation

1. Obtain K train and test splits.

2. **For each split,**

2.1 initialize trainer and model 2.2. Train model with configuration provided in args. 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.

3. Report the mean and average MRR .

Parameters

- `self`
- `dataset`

Returns

model

`dicee.trainer.model_parallelism`

Classes

TensorParallel

Abstract base class for KGE model trainers.

Functions

extract_input_outputs(z[, device])

find_good_batch_size(train_loader,
tp_ensemble_model)

forward_backward_update_loss(→ float)

Module Contents

`dicee.trainer.model_parallelism.extract_input_outputs(z: list, device=None)`

`dicee.trainer.model_parallelism.find_good_batch_size(train_loader, tp_ensemble_model)`

`dicee.trainer.model_parallelism.forward_backward_update_loss` (*z: Tuple, ensemble_model*)
→ float

class `dicee.trainer.model_parallelism.TensorParallel` (*args, callbacks*)

Bases: `dicee.abstracts.AbstractTrainer`

Abstract base class for KGE model trainers.

Provides the callback dispatch mechanism shared by all concrete trainer implementations (TorchTrainer, TorchD-
DPTrainer, etc.). Sub-classes call the `on_*` hooks at the appropriate points in the training loop so that any registered `AbstractCallback` can react.

Parameters

- **args** (*argparse.Namespace or similar*) – Processed configuration object. Must expose at least `random_seed` (int).
- **callbacks** (*list of AbstractCallback*) – Ordered list of callback instances to invoke at each lifecycle hook.

fit (**args, **kwargs*)

Train model

`dicee.trainer.torch_trainer`

Classes

TorchTrainer

TorchTrainer for using single GPU or multi CPUs on a single node

Module Contents

class `dicee.trainer.torch_trainer.TorchTrainer` (*args, callbacks*)

Bases: `dicee.abstracts.AbstractTrainer`

TorchTrainer for using single GPU or multi CPUs on a single node

Arguments

`callbacks`: list of Abstract callback instances

`loss_function` = None

`optimizer` = None

`model` = None

`train_dataloaders` = None

`training_step` = None

`process`

fit (**args, train_dataloaders, **kwargs*) → None

Training starts

Arguments

kwargs: Tuple
empty dictionary

Return type
batch loss (float)

forward_backward_update (*x_batch: torch.Tensor, y_batch: torch.Tensor*) → torch.Tensor

Compute forward, loss, backward, and parameter update

Arguments

Return type
batch loss (float)

extract_input_outputs_set_device (*batch: list*) → Tuple

Construct inputs and outputs from a batch of inputs with outputs From a batch of inputs and put

Arguments

Return type
(tuple) mini-batch on select device

dicee.trainer.torch_trainer_ddp

Classes

TorchDDPTrainer

A Trainer based on torch.nn.parallel.DistributedDataParallel

NodeTrainer

Functions

make_iterable_verbose(→ Iterable)

Module Contents

`dicee.trainer.torch_trainer_ddp.make_iterable_verbose` (*iterable_object, verbose, desc='Default', position=None, leave=True*) → Iterable

class `dicee.trainer.torch_trainer_ddp.TorchDDPTrainer` (*args, callbacks*)

Bases: *dicee.abstracts.AbstractTrainer*

A Trainer based on torch.nn.parallel.DistributedDataParallel

Arguments

entity_idx
mapping.

relation_idx
mapping.

form
?

store
?

label_smoothing_rate
Using hard targets (0,1) drives weights to infinity. An outlier produces enormous gradients.

Return type
torch.utils.data.Dataset

fit (*args, **kwargs)
Train model

```
class dicee.trainer.torch_trainer_ddp.NodeTrainer(trainer, model: torch.nn.Module,
    train_dataset_loader: torch.utils.data.DataLoader, callbacks, num_epochs: int)
```

trainer

local_rank

global_rank

optimizer

train_dataset_loader

loss_func

callbacks

model

num_epochs

loss_history = []

ctx

scaler

extract_input_outputs (z: list)

train()
Training loop for DDP

Classes

DICE_Trainer

DICE_Trainer implement

Package Contents

```
class dicee.trainer.DICE_Trainer(args, is_continual_training: bool, storage_path, evaluator=None)
```

DICE_Trainer implement

- 1- Pytorch Lightning trainer (<https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html>)
- 2- Multi-GPU Trainer(<https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html>)
- 3- CPU Trainer

```

    args
    is_continual_training:bool
    storage_path:str
    evaluator:
    report:dict

report

args

trainer = None

is_continual_training

storage_path

evaluator = None

form_of_labelling = None

continual_start (knowledge_graph)

```

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

```

initialize_trainer (callbacks: List)
    → lightning.Trainer | dicke.trainer.model_parallelism.TensorParallel | dicke.trainer.torch_trainer.TorchTrainer | dicke.t
    Initialize Trainer from input arguments

```

```

initialize_or_load_model ()

```

```

init_dataloader (dataset: torch.utils.data.Dataset) → torch.utils.data.DataLoader

```

```

init_dataset () → torch.utils.data.Dataset

```

```

start (knowledge_graph: dicke.knowledge_graph.KG | numpy.memmap)
    → Tuple[dicke.models.base_model.BaseKGE, str]

```

Start the training

- (1) Initialize Trainer
- (2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

k_fold_cross_validation (*dataset*) → Tuple[*dicke.models.base_model.BaseKGE*, str]

Perform K-fold Cross-Validation

1. Obtain K train and test splits.
2. **For each split,**
 - 2.1 initialize trainer and model
 - 2.2. Train model with configuration provided in args.
 - 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.
3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

dicke.weight_averaging

Classes

<i>ASWA</i>	Adaptive stochastic weight averaging
<i>SWA</i>	Stochastic Weight Averaging callback.
<i>SWAG</i>	Stochastic Weight Averaging - Gaussian (SWAG).
<i>EMA</i>	Exponential Moving Average (EMA) callback.
<i>TWA</i>	Train with Weight Averaging (TWA) using subspace projection + averaging.

Module Contents

class *dicke.weight_averaging.ASWA* (*num_epochs*, *path*)

Bases: *dicke.abstracts.AbstractCallback*

Adaptive stochastic weight averaging ASWE keeps track of the validation performance and update s the ensemble model accordingly.

path

num_epochs

initial_eval_setting = None

epoch_count = 0

alphas = []

val_aswa = -1

on_fit_end (*trainer*, *model*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

static compute_mrr (*trainer*, *model*) → float

`get_aswa_state_dict (model)`

`decide (running_model_state_dict, ensemble_state_dict, val_running_model, mrr_updated_ensemble_model)`

Perform Hard Update, software or rejection

Parameters

- `running_model_state_dict`
- `ensemble_state_dict`
- `val_running_model`
- `mrr_updated_ensemble_model`

`on_train_epoch_end (trainer, model)`

Called at the end of each training epoch.

Parameters

- `trainer` (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- `model` (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

`class dicee.weight_averaging.SWA (swa_start_epoch: int = 1, swa_c_epochs: int = 1, lr_init: float = 0.1, swa_lr: float = 0.05, max_epochs: int = None)`

Bases: `dicee.abstracts.AbstractCallback`

Stochastic Weight Averaging callback.

Initialize SWA callback.

swa_start_epoch: int

The epoch at which to start SWA.

swa_c_epochs: int

The number of epochs to use for SWA.

lr_init: float

The initial learning rate.

swa_lr: float

The learning rate to use during SWA.

max_epochs: int

The maximum number of epochs. `args.num_epochs`

`swa_start_epoch`

`swa_c_epochs = 1`

`swa_lr = 0.05`

`lr_init = 0.1`

`max_epochs = None`

`swa_model = None`

`swa_n = 0`

```

current_epoch = -1

static moving_average (swa_model, running_model, alpha)
    Update SWA model with moving average of current model. Math: # SWA update:  $\theta_{\text{swa}} \leftarrow (1 - \alpha) * \theta_{\text{swa}} + \alpha * \theta$  #  $\alpha = 1 / (n + 1)$ , where  $n$  = number of models already averaged #  $\alpha$  is tracked via self.swa_n in code

on_train_epoch_start (trainer, model)
    Update learning rate according to SWA schedule.

on_train_epoch_end (trainer, model)
    Apply SWA averaging if conditions are met.

on_fit_end (trainer, model)
    Replace main model with SWA model at the end of training.

class dicee.weight_averaging.SWAG (swa_start_epoch, swa_c_epochs: int = 1, lr_init: float = 0.1, swa_lr: float = 0.05, max_epochs: int = None, max_num_models: int = 20, var_clamp: float = 1e-30)
    Bases: dicee.abstracts.AbstractCallback

    Stochastic Weight Averaging - Gaussian (SWAG). Parameters

    swa_start_epoch
        [int] Epoch at which to start collecting weights.

    swa_c_epochs
        [int] Interval of epochs between updates.

    lr_init
        [float] Initial LR.

    swa_lr
        [float] LR in SWA / GSWA phase.

    max_epochs
        [int] Total number of epochs.

    max_num_models
        [int] Number of models to keep for low-rank covariance approx.

    var_clamp
        [float] Clamp low variance for stability.

swa_start_epoch

swa_c_epochs = 1

swa_lr = 0.05

lr_init = 0.1

max_epochs = None

max_num_models = 20

var_clamp = 1e-30

mean = None

sq_mean = None

```

```

deviations = []

gswa_n = 0

current_epoch = -1

get_mean_and_var()
    Return mean + variance (diagonal part).

sample(base_model, scale=0.5)
    Sample new model from SWAG posterior distribution.

    Math: # From SWAG, posterior is approximated as:  $\theta \sim N(\text{mean}, \Sigma)$  # where  $\Sigma \approx \text{diag}(\text{var}) + (1/(K-1)) * D D^T$ 
    # - mean = running average of weights # - var = elementwise variance ( $\text{sq\_mean} - \text{mean}^2$ ) # - D = [dev_1, dev_2, ..., dev_K], deviations from mean (low-rank approx) # - K = number of collected models

    # Sampling step: # 1.  $\theta_{\text{diag}} = \text{mean} + \text{scale} * \text{std} \odot \varepsilon$ , where  $\varepsilon \sim N(0, I)$  # 2.  $\theta_{\text{lowrank}} = \theta_{\text{diag}} + (D z) / \sqrt{K-1}$ , where  $z \sim N(0, I_K)$  # Final sample =  $\theta_{\text{lowrank}}$ 

on_train_epoch_start(trainer, model)
    Update LR schedule (same as SWA).

on_train_epoch_end(trainer, model)
    Collect Gaussian stats at the end of epochs after swa_start.

on_fit_end(trainer, model)
    Set model weights to the collected SWAG mean at the end of training.

class dicee.weight_averaging.EMA(ema_start_epoch: int, decay: float = 0.999,
    max_epochs: int = None, ema_c_epochs: int = 1)
    Bases: dicee.abstracts.AbstractCallback

    Exponential Moving Average (EMA) callback.

    Parameters
    • ema_start_epoch (int) – Epoch to start EMA.
    • decay (float) – EMA decay rate (typical: 0.99 - 0.9999) Math:  $\theta_{\text{ema}} \leftarrow \text{decay} * \theta_{\text{ema}} + (1 - \text{decay}) * \theta$ 
    • max_epochs (int) – Maximum number of epochs.

ema_start_epoch

decay = 0.999

max_epochs = None

ema_c_epochs = 1

ema_model = None

current_epoch = -1

static ema_update(ema_model, running_model, decay: float)
    Update EMA model with exponential moving average of current model. Math: # EMA update:  $\theta_{\text{ema}} \leftarrow (1 - \alpha) * \theta_{\text{ema}} + \alpha * \theta$  #  $\alpha = 1 - \text{decay}$ , where decay is the EMA smoothing factor (typical 0.99 - 0.999) #  $\alpha$  controls how much of the current model  $\theta$  contributes to the EMA # decay is fixed in code
    -> can be extended to scheduled

```

```

on_train_epoch_start (trainer, model)
    Track current epoch.

on_train_epoch_end (trainer, model)
    Update EMA if past start epoch.

on_fit_end (trainer, model)
    Replace main model with EMA model at the end of training.

class dicee.weight_averaging.TWA (twa_start_epoch: int, lr_init: float, num_samples: int = 5,
    reg_lambda: float = 0.0, max_epochs: int = None, twa_c_epochs: int = 1)
    Bases: dicee.abstracts.AbstractCallback

    Train with Weight Averaging (TWA) using subspace projection + averaging.

Parameters

    twa_start_epoch
        [int] Epoch to start TWA.

    lr_init
        [float] Learning rate used for  $\beta$  updates.

    num_samples
        [int] Number of sampled weight snapshots to build projection subspace.

    reg_lambda
        [float] Regularization coefficient for  $\beta$  updates.

    max_epochs
        [int] Total number of training epochs.

    twa_c_epochs
        [int] Interval of epochs between TWA updates.

twa_start_epoch

num_samples = 5

reg_lambda = 0.0

max_epochs = None

lr_init

twa_c_epochs = 1

current_epoch = -1

weight_samples = []

twa_model = None

base_weights = None

P = None

beta = None

sample_weights (model)
    Collect sampled weights from the current model and maintain rolling buffer.

```


build_projection (*weight_samples*, *k=None*)

Build projection subspace from collected weight samples. :param weight_samples: list of flat weight tensors [(D,), ...] :param k: number of basis vectors to keep. Defaults to min(N, D).

Returns

(D,) base weight vector (average) P: (D, k) projection matrix with top-k basis directions

Return type

mean_w

on_train_epoch_start (*trainer*, *model*)

Track epoch.

on_train_epoch_end (*trainer*, *model*)

Main TWA logic: build subspace and update in β space.

Math: # TWA weight update: # $w_{twa} = \text{mean_w} + P * \beta$ # $\text{mean_w} = (1/n) * \sum_i w_i$ (SWA baseline)
$\beta \leftarrow (1 - \eta * \lambda) * \beta - \eta * P^T * g$ # g = gradient of training loss w.r.t. full model weights
η = learning rate, λ = ridge regularization # P = orthonormal basis spanning sampled checkpoints
{ w_i }

on_fit_end (*trainer*, *model*)

Replace with TWA model at the end.

14.2 Attributes

`__version__`

14.3 Classes

<i>Execute</i>	Executor class for training, retraining and evaluating KGE models.
<i>KGE</i>	Knowledge Graph Embedding Class for interactive usage of pre-trained models
<i>QueryGenerator</i>	
<i>DICE_Trainer</i>	DICE_Trainer implement
<i>Evaluator</i>	Evaluator class for KGE models in various downstream tasks.

14.4 Package Contents

class `dicee.Execute` (*args*, *continuous_training: bool = False*)

Executor class for training, retraining and evaluating KGE models.

Handles the complete workflow: 1. Loading & Preprocessing & Serializing input data 2. Training & Validation & Testing 3. Storing all necessary information

args

Processed input arguments.

distributed

Whether distributed training is enabled.

rank
Process rank in distributed training.

world_size
Total number of processes.

local_rank
Local GPU rank.

trainer
Training handler instance.

trained_model
The trained model after training completes.

knowledge_graph
The loaded knowledge graph.

report
Dictionary storing training metrics and results.

evaluator
Model evaluation handler.

distributed

rank

world_size

local_rank

args

is_continual_training = False

trainer: *dicce.trainer.DICE_Trainer* | None = None

trained_model = None

knowledge_graph: *dicce.knowledge_graph.KG* | None = None

report: Dict

evaluator: *dicce.evaluator.Evaluator* | None = None

start_time: float | None = None

is_local_rank_zero() → bool

cleanup()

setup_executor() → None
Set up storage directories for the experiment.

Creates or reuses experiment directories based on configuration. Saves the configuration to a JSON file.

create_and_store_kg() → None
Create knowledge graph and store as memory-mapped file.

Only executed on local rank 0 in distributed training. Skips if memmap already exists.

load_from_memmap() → None

Load knowledge graph from memory-mapped file.

save_trained_model() → None

Save a knowledge graph embedding model

- (1) Send model to eval mode and cpu.
- (2) Store the memory footprint of the model.
- (3) Save the model into disk.
- (4) Update the stats of KG again ?

Parameter

rtype

None

end(*form_of_labelling: str*) → dict

End training

- (1) Store trained model.
- (2) Report runtimes.
- (3) Eval model if required.

Parameter

rtype

A dict containing information about the training and/or evaluation

write_report() → None

Report training related information in a report.json file

start() → dict

Start training

(1) Loading the Data # (2) Create an evaluator object. # (3) Create a trainer object. # (4) Start the training

Parameter

rtype

A dict containing information about the training and/or evaluation

class `dicee.KGE` (*path=None, url=None, construct_ensemble=False, model_name=None*)

Bases: `dicee.abstracts.BaseInteractiveKGE`, `dicee.abstracts.InteractiveQueryDecomposition`, `dicee.abstracts.BaseInteractiveTrainKGE`

Knowledge Graph Embedding Class for interactive usage of pre-trained models

__str__()

to (*device: str*) → None

get_transductive_entity_embeddings (*indices: torch.LongTensor | List[str], as_pytorch=False, as_numpy=False, as_list=True*) → `torch.FloatTensor` | `numpy.ndarray` | `List[float]`

create_vector_database (*collection_name: str, distance: str, location: str = 'localhost',
port: int = 6333*)

generate (*h="", r=""*)

eval_lp_performance (*dataset=List[Tuple[str, str, str]], filtered=True*)

predict_missing_head_entity (*relation: List[str] | str, tail_entity: List[str] | str, within=None,
batch_size=2, topk=1, return_indices=False*) → Tuple

Given a relation and a tail entity, return top k ranked head entity.

$\text{argmax}_{\{e \in E\}} f(e, r, t)$, where $r \in R, t \in E$.

Parameter

relation: Union[List[str], str]

String representation of selected relations.

tail_entity: Union[List[str], str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_relations (*head_entity: List[str] | str, tail_entity: List[str] | str, within=None,
batch_size=2, topk=1, return_indices=False*) → Tuple

Given a head entity and a tail entity, return top k ranked relations.

$\text{argmax}_{\{r \in R\}} f(h, r, t)$, where $h, t \in E$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_tail_entity (*head_entity: List[str] | str, relation: List[str] | str,
within: List[str] = None, batch_size=2, topk=1, return_indices=False*) → torch.FloatTensor

Given a head entity and a relation, return top k ranked entities

$\text{argmax}_{\{e \in E\}} f(h, r, e)$, where $h \in E$ and $r \in R$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

Returns: Tuple

scores

predict (*, h: List[str] | str | None = None, r: List[str] | str | None = None, t: List[str] | str | None = None, within: List[str] | None = None, logits: bool = True) → torch.FloatTensor

Predict scores for triples or missing triple elements.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.
- **t** – Tail entity/entities. None to predict tails.
- **within** – Optional list of entities to restrict predictions to.
- **logits** – If True, return raw scores. If False, return sigmoid scores (0-1).

Returns

- Single triple (h, r, t): scalar score
- Missing element: vector of all possible scores

Return type

torch.FloatTensor of scores. Shape depends on the query type

Raises

AssertionError – If inputs are not strings or lists of strings.

Examples

```
>>> # Score a specific triple
>>> model.predict(h="Mongolia", r="isLocatedIn", t="Asia", logits=False)
tensor(0.9523)
```

```
>>> # Get scores for all possible tail entities
>>> model.predict(h="Mongolia", r="isLocatedIn", t=None)
tensor([0.21, 0.95, 0.03, ...]) # One score per entity
```

predict_topk (*, h: str | List[str] | None = None, r: str | List[str] | None = None, t: str | List[str] | None = None, topk: int = 10, within: List[str] | None = None, batch_size: int = 1024) → List[Tuple[str, float]] | List[List[Tuple[str, float]]]

Predict top-k missing items in a given triple pattern.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.

- **t** – Tail entity/entities. None to predict tails.
- **topk** – Number of top predictions to return.
- **within** – Optional list of entities to restrict predictions to.
- **batch_size** – Batch size for processing multiple queries.

Returns

List[(item, score), ...] of length topk. For batch query: List of such lists, one per query.

Return type

For single query

Raises

- **AssertionError** – If more than one of h, r, t is None.
- **AssertionError** – If the required arguments for a query type are None.

Examples

```
>>> model.predict_topk(h=["Mongolia"], r=["isLocatedIn"], topk=3)
[('Asia', 0.99), ('Europe', 0.02), ...]
```

```
>>> model.predict_topk(r=["isLocatedIn"], t=["Asia"], topk=5)
[('Mongolia', 0.85), ('China', 0.82), ...]
```

triple_score (*h: List[str] | str = None, r: List[str] | str = None, t: List[str] | str = None, logits=False*)

→ torch.FloatTensor

Predict triple score

Parameter

head_entity: List[str]

String representation of selected entities.

relation: List[str]

String representation of selected relations.

tail_entity: List[str]

String representation of selected entities.

logits: bool

If logits is True, unnormalized score returned

Returns: Tuple

pytorch tensor of triple score

return_multi_hop_query_results (*aggregated_query_for_all_entities, k: int, only_scores*)

single_hop_query_answering (*query: tuple, only_scores: bool = True, k: int = None, use_logits: bool = True*)

```

answer_multi_hop_query (query_type: str | None = None,
    query: Tuple[str | Tuple[str, str], Ellipsis] | None = None,
    queries: List[Tuple[str | Tuple[str, str], Ellipsis]] | None = None, tnorm: str = 'prod',
    neg_norm: str = 'standard', lambda_: float = 0.0, k: int = 10, only_scores: bool = False,
    use_logits: bool = True)
    → List[Tuple[str, torch.Tensor]] | List[List[Tuple[str, torch.Tensor]]]

```

Answer multi-hop EPFO (Existential Positive First-Order) queries.

Supports 9 query types: 1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up. See docs/guides/multi_hop_queries.md for detailed query patterns.

Parameters

- **query_type** – Query pattern name. One of: - 1p: (e, (r,)) # One-hop - 2p: (e, (r1, r2)) # Two-hop - 3p: (e, (r1, r2, r3)) # Three-hop - 2i: ((e1, (r1,)), (e2, (r2,))) # Two-way intersection - 3i: ((e1, (r1,)), (e2, (r2,)), (e3, (r3,))) # Three-way intersection - ip: (((e1, (r1,)), (e2, (r2,))), (r3,)) # Intersection + projection - pi: ((e, (r1, r2)), (r3,)) # Projection + intersection (2i meets 2p) - 2u: ((e1, (r1,)), (e2, (r2,))) # Two-way union - up: ((e, (r1, r2)), (e, (r3,))) # Union + projection
- **query** – Single query tuple matching the query_type pattern.
- **queries** – Batch of queries. If provided, query must be None.
- **tnorm** – T-norm for intersection/union. Options: “prod”, “min”.
- **neg_norm** – Negation norm. Options: “standard”, “sugeno”, “yager”.
- **lambda** – Parameter for sugeno and yager negation (0.0-1.0).
- **k** – Number of top answer entities to return.
- **only_scores** – If True, return only scores tensor. If False, return (entity, score) tuples.
- **use_logits** – If True, use raw model logits. If False, use sigmoid probabilities.

Returns

List[(entity, score), ...] of top-k answers. For batch queries: List of such lists, one per query.

Return type

For single query

Raises

- **ValueError** – If query_type is not in {1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up}.
- **AssertionError** – If query structure doesn't match query_type pattern.

Examples

```

>>> # 1p: Find entities located in Asia
>>> model.answer_multi_hop_query(
...     query_type="1p",
...     query=("Asia", ("isLocatedIn",)),
...     k=5
... )
[("Mongolia", 0.92), ("China", 0.89), ...]

```

```

>>> # 2p: Two-hop query (e.g., "capital of countries in Europe")
>>> model.answer_multi_hop_query(
...     query_type="2p",

```

(continues on next page)

(continued from previous page)

```
...     query=("Europe", ("isLocatedIn", "hasCapital")),
...     k=3
... )
[("Paris", 0.85), ("Berlin", 0.82), ...]
```

```
>>> # 2i: Intersection query
>>> model.answer_multi_hop_query(
...     query_type="2i",
...     query=((("Asia", ("isLocatedIn",)), ("Mountains", ("hasGeography",))),
...     k=5
... )
[("Nepal", 0.78), ("Tibet", 0.65), ...]
```

See also

- docs/guides/multi_hop_queries.md: Complete guide with all query patterns
- tests/test_answer_multi_hop_query.py: Usage examples

find_missing_triples (*confidence: float, entities: List[str] = None, relations: List[str] = None, topk: int = 10, at_most: int = sys.maxsize*) → Set

Find missing triples

Iterative over a set of entities E and a set of relation R :

forall e in E and forall r in R f(e,r,x)

Return (e,r,x)

otin G and f(e,r,x) > confidence

confidence: float

A threshold for an output of a sigmoid function given a triple.

topk: int

Highest ranked k item to select triples with f(e,r,x) > confidence .

at_most: int

Stop after finding at_most missing triples

{(e,r,x) | f(e,r,x) > confidence land (e,r,x)}

otin G

predict_literals (*entity: List[str] | str = None, attribute: List[str] | str = None, denormalize_preds: bool = True*) → numpy.ndarray

Predicts literal values for given entities and attributes.

Parameters

- **entity** (*Union[List[str], str]*) – Entity or list of entities to predict literals for.
- **attribute** (*Union[List[str], str]*) – Attribute or list of attributes to predict literals for.
- **denormalize_preds** (*bool*) – If True, denormalizes the predictions.

Returns

Predictions for the given entities and attributes.

Return type

numpy ndarray

```
class dicee.QueryGenerator (train_path: str, val_path: str, test_path: str, ent2id: Dict = None,  
    rel2id: Dict = None, seed: int = 1, gen_valid: bool = False, gen_test: bool = True)  
  
    train_path  
  
    val_path  
  
    test_path  
  
    gen_valid = False  
  
    gen_test = True  
  
    seed = 1  
  
    max_ans_num = 1000000.0  
  
    mode  
  
    ent2id = None  
  
    rel2id: Dict = None  
  
    ent_in: Dict  
  
    ent_out: Dict  
  
    query_name_to_struct  
  
    list2tuple (list_data)  
  
    tuple2list (x: List | Tuple) → List | Tuple  
        Convert a nested tuple to a nested list.  
  
    set_global_seed (seed: int)  
        Set seed  
  
    construct_graph (paths: List[str]) → Tuple[Dict, Dict]  
        Construct graph from triples Returns dicts with incoming and outgoing edges  
  
    fill_query (query_structure: List[str | List], ent_in: Dict, ent_out: Dict, answer: int) → bool  
        Private method for fill_query logic.  
  
    achieve_answer (query: List[str | List], ent_in: Dict, ent_out: Dict) → set  
        Private method for achieve_answer logic. @TODO: Document the code  
  
    write_links (ent_out, small_ent_out)  
  
    ground_queries (query_structure: List[str | List], ent_in: Dict, ent_out: Dict, small_ent_in: Dict,  
        small_ent_out: Dict, gen_num: int, query_name: str)  
        Generating queries and achieving answers  
  
    unmap (query_type, queries, tp_answers, fp_answers, fn_answers)
```

```

unmap_query (query_structure, query, id2ent, id2rel)

generate_queries (query_struct: List, gen_num: int, query_type: str)
    Passing incoming and outgoing edges to ground queries depending on mode [train valid or text] and getting
    queries and answers in return @ TODO: create a class for each single query struct

save_queries (query_type: str, gen_num: int, save_path: str)

abstractmethod load_queries (path)

get_queries (query_type: str, gen_num: int)

static save_queries_and_answers (path: str, data: List[Tuple[str, Tuple[collections.defaultdict]]])
    → None
    Save Queries into Disk

static load_queries_and_answers (path: str) → List[Tuple[str, Tuple[collections.defaultdict]]]
    Load Queries from Disk to Memory

class dicee.DICE_Trainer (args, is_continual_training: bool, storage_path, evaluator=None)

    DICE_Trainer implement
    1- Pytorch Lightning trainer (https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html)
    2- Multi-GPU Trainer(https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html)
    3- CPU Trainer

    args
    is_continual_training:bool
    storage_path:str
    evaluator:
    report:dict

report

args

trainer = None

is_continual_training

storage_path

evaluator = None

form_of_labelling = None

continual_start (knowledge_graph)

    (1) Initialize training.
    (2) Load model
    (3) Load trainer (3) Fit model

```

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

initialize_trainer (*callbacks: List*)

→ `lightning.Trainer` | `diccee.trainer.model_parallelism.TensorParallel` | `diccee.trainer.torch_trainer.TorchTrainer` | `diccee.`

Initialize Trainer from input arguments

initialize_or_load_model ()

init_dataloader (*dataset: torch.utils.data.Dataset*) → `torch.utils.data.DataLoader`

init_dataset () → `torch.utils.data.Dataset`

start (*knowledge_graph: diccee.knowledge_graph.KG* | *numpy.memmap*)

→ `Tuple[diccee.models.base_model.BaseKGE, str]`

Start the training

- (1) Initialize Trainer
- (2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

k_fold_cross_validation (*dataset*) → `Tuple[diccee.models.base_model.BaseKGE, str]`

Perform K-fold Cross-Validation

1. Obtain K train and test splits.
2. **For each split,**
 - 2.1 initialize trainer and model
 - 2.2. Train model with configuration provided in args.
 - 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.
3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

class diccee.Evaluator (*args, is_continual_training: bool = False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset*, *trained_model*, *form_of_labelling*: str, *during_training*: bool = False) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).

- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set*, *valid_set=None*, *test_set=None*, *trained_model*)
→ None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set=None*, *valid_set=None*,
test_set=None, *ordered_bpe_entities*, *trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, *raw_train_set*, *raw_valid_set=None*, *raw_test_set=None*, *trained_model*,
form_of_labelling) → None

Evaluate Byte model with generation.

eval_with_bpe_vs_all (*, *raw_train_set*, *raw_valid_set=None*, *raw_test_set=None*, *trained_model*,
form_of_labelling) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, *train_set*, *valid_set=None*, *test_set=None*, *trained_model*, *form_of_labelling*)
→ None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all (*model*, *triple_idx*, *info: str | None = None*,
form_of_labelling: str | None = None) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – 'EntityPrediction' or 'RelationPrediction'.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

evaluate_lp_with_byte (*model*, *triples: List[List[str]]*, *info: str | None = None*) → Dict[str, float]

Evaluate Byte model with text generation.

Parameters

- **model** – Byte model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (*model*, *triples: List[List[str]]*, *info: str | None = None*,
form_of_labelling: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (*model*, *triple_idx*, *info*: *str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: *str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: *numpy.ndarray*, *form_of_labelling*: *str*)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

```
dicee.__version__ = '0.3.3'
```

Python Module Index

d

- `dicee`, 12
- `dicee.__main__`, 12
- `dicee.abstracts`, 12
- `dicee.analyse_experiments`, 20
- `dicee.callbacks`, 22
- `dicee.config`, 28
- `dicee.dataset_classes`, 31
- `dicee.eval_static_funcs`, 39
- `dicee.evaluation`, 43
 - `dicee.evaluation.ensemble`, 43
 - `dicee.evaluation.evaluator`, 44
 - `dicee.evaluation.link_prediction`, 48
 - `dicee.evaluation.literal_prediction`, 51
 - `dicee.evaluation.utils`, 52
- `dicee.evaluator`, 61
- `dicee.executer`, 64
- `dicee.knowledge_graph`, 67
- `dicee.knowledge_graph_embeddings`, 70
- `dicee.models`, 75
 - `dicee.models.adopt`, 75
 - `dicee.models.base_model`, 84
 - `dicee.models.clifford`, 92
 - `dicee.models.complex`, 101
 - `dicee.models.dualE`, 104
 - `dicee.models.ensemble`, 105
 - `dicee.models.function_space`, 106
 - `dicee.models.literal`, 110
 - `dicee.models.octonion`, 112
 - `dicee.models.pykeen_models`, 115
 - `dicee.models.quaternion`, 116
 - `dicee.models.real`, 119
 - `dicee.models.static_funcs`, 124
 - `dicee.models.transformers`, 125
- `dicee.query_generator`, 196
- `dicee.read_preprocess_save_load_kg`, 197
- `dicee.read_preprocess_save_load_kg.preprocess`, 197
- `dicee.read_preprocess_save_load_kg.read_from_disk`, 198
- `dicee.read_preprocess_save_load_kg.save_load_disk`, 199
- `dicee.read_preprocess_save_load_kg.util`, 199
- `dicee.sanity_checkers`, 204
- `dicee.scripts`, 204
 - `dicee.scripts.index_serve`, 204
 - `dicee.scripts.run`, 206
- `dicee.static_funcs`, 206
- `dicee.static_funcs_training`, 210
- `dicee.static_preprocess_funcs`, 212
- `dicee.trainer`, 213
 - `dicee.trainer.dice_trainer`, 213
 - `dicee.trainer.model_parallelism`, 215
 - `dicee.trainer.torch_trainer`, 216
 - `dicee.trainer.torch_trainer_ddp`, 217
- `dicee.weight_averaging`, 220

Index

Non-alphabetical

`__call__()` (*dicee.models.base_model.IdentityClass method*), 92
`__call__()` (*dicee.models.ensemble.EnsembleKGE method*), 106
`__call__()` (*dicee.models.IdentityClass method*), 143, 165, 172
`__getitem__()` (*dicee.dataset_classes.AllvsAll method*), 35
`__getitem__()` (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 32
`__getitem__()` (*dicee.dataset_classes.FixedNegSampleDataset method*), 38
`__getitem__()` (*dicee.dataset_classes.KvsAll method*), 35
`__getitem__()` (*dicee.dataset_classes.KvsSampleDataset method*), 36
`__getitem__()` (*dicee.dataset_classes.LiteralDataset method*), 37
`__getitem__()` (*dicee.dataset_classes.MultiClassClassificationDataset method*), 33
`__getitem__()` (*dicee.dataset_classes.MultiLabelDataset method*), 34
`__getitem__()` (*dicee.dataset_classes.OnevsAllDataset method*), 36
`__getitem__()` (*dicee.dataset_classes.OnevsSample method*), 39
`__getitem__()` (*dicee.dataset_classes.TriplePredictionDataset method*), 39
`__iter__()` (*dicee.config.Namespace method*), 31
`__iter__()` (*dicee.knowledge_graph.KG method*), 69
`__iter__()` (*dicee.models.ensemble.EnsembleKGE method*), 106
`__len__()` (*dicee.dataset_classes.AllvsAll method*), 35
`__len__()` (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 32
`__len__()` (*dicee.dataset_classes.FixedNegSampleDataset method*), 38
`__len__()` (*dicee.dataset_classes.KvsAll method*), 35
`__len__()` (*dicee.dataset_classes.KvsSampleDataset method*), 36
`__len__()` (*dicee.dataset_classes.LiteralDataset method*), 37
`__len__()` (*dicee.dataset_classes.MultiClassClassificationDataset method*), 33
`__len__()` (*dicee.dataset_classes.MultiLabelDataset method*), 33
`__len__()` (*dicee.dataset_classes.OnevsAllDataset method*), 36
`__len__()` (*dicee.dataset_classes.OnevsSample method*), 39
`__len__()` (*dicee.dataset_classes.TriplePredictionDataset method*), 39
`__len__()` (*dicee.knowledge_graph.KG method*), 70
`__len__()` (*dicee.models.ensemble.EnsembleKGE method*), 106
`__setstate__()` (*dicee.models.ADOPT method*), 134
`__setstate__()` (*dicee.models.adopt.ADOPT method*), 79
`__str__()` (*dicee.KGE method*), 227
`__str__()` (*dicee.knowledge_graph_embeddings.KGE method*), 70
`__str__()` (*dicee.models.ensemble.EnsembleKGE method*), 106
`__version__` (in module *dicee*), 238

A

`AbstractCallback` (class in *dicee.abstracts*), 18
`AbstractPPECallback` (class in *dicee.abstracts*), 19
`AbstractTrainer` (class in *dicee.abstracts*), 12
`AccumulateEpochLossCallback` (class in *dicee.callbacks*), 22
`achieve_answer()` (*dicee.query_generator.QueryGenerator method*), 197
`achieve_answer()` (*dicee.QueryGenerator method*), 233
`AConEx` (class in *dicee.models*), 158
`AConEx` (class in *dicee.models.complex*), 102
`AConvO` (class in *dicee.models*), 173
`AConvO` (class in *dicee.models.octonion*), 114
`AConvQ` (class in *dicee.models*), 167
`AConvQ` (class in *dicee.models.quaternion*), 118
`adaptive_lr` (*dicee.config.Namespace attribute*), 31
`adaptive_swa` (*dicee.config.Namespace attribute*), 30
`add_new_entity_embeddings()` (*dicee.abstracts.BaseInteractiveKGE method*), 16
`add_noise_rate` (*dicee.config.Namespace attribute*), 28
`add_noise_rate` (*dicee.knowledge_graph.KG attribute*), 68
`add_noisy_triples()` (in module *dicee.static_funcs*), 209
`add_noisy_triples_into_training()` (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk method*), 199
`add_noisy_triples_into_training()` (*dicee.read_preprocess_save_load_kg.ReadFromDisk method*), 204
`add_reciprocal` (*dicee.knowledge_graph.KG attribute*), 68
`ADOPT` (class in *dicee.models*), 132
`ADOPT` (class in *dicee.models.adopt*), 77
`adopt()` (in module *dicee.models.adopt*), 80
`ALL_HITS_RANGE` (in module *dicee.evaluation.utils*), 52
`AllvsAll` (class in *dicee.dataset_classes*), 34

- alphas (*dicee.abstracts.AbstractPPECallback* attribute), 19
- alphas (*dicee.weight_averaging.ASWA* attribute), 220
- analyse() (in module *dicee.analyse_experiments*), 22
- answer_multi_hop_query() (*dicee.KGE* method), 230
- answer_multi_hop_query() (*dicee.knowledge_graph_embeddings.KGE* method), 73
- app (in module *dicee.scripts.index_serve*), 205
- apply_coefficients() (*dicee.models.clifford.DeCaL* method), 99
- apply_coefficients() (*dicee.models.clifford.Keci* method), 94
- apply_coefficients() (*dicee.models.DeCaL* method), 179
- apply_coefficients() (*dicee.models.Keci* method), 175
- apply_reciprocal_or_noise() (in module *dicee.read_preprocess_save_load_kg.util*), 201
- apply_semantic_constraint (*dicee.abstracts.BaseInteractiveKGE* attribute), 14
- apply_unit_norm (*dicee.models.base_model.BaseKGE* attribute), 88
- apply_unit_norm (*dicee.models.BaseKGE* attribute), 139, 144, 154, 161, 168, 182, 187
- args (*dicee.DICE_Trainer* attribute), 234
- args (*dicee.evaluation.Evaluator* attribute), 55, 56
- args (*dicee.evaluation.evaluator.Evaluator* attribute), 45, 46
- args (*dicee.Evaluator* attribute), 236
- args (*dicee.evaluator.Evaluator* attribute), 62
- args (*dicee.Execute* attribute), 225, 226
- args (*dicee.executer.Execute* attribute), 65
- args (*dicee.models.base_model.BaseKGE* attribute), 88
- args (*dicee.models.base_model.IdentityClass* attribute), 92
- args (*dicee.models.BaseKGE* attribute), 139, 144, 154, 161, 168, 182, 187
- args (*dicee.models.ensemble.EnsembleKGE* attribute), 106
- args (*dicee.models.IdentityClass* attribute), 143, 165, 172
- args (*dicee.models.pykeen_models.PykeenKGE* attribute), 115
- args (*dicee.models.PykeenKGE* attribute), 186
- args (*dicee.trainer.DICE_Trainer* attribute), 219
- args (*dicee.trainer.dice_trainer.DICE_Trainer* attribute), 214
- ASWA (class in *dicee.weight_averaging*), 220
- aswa (*dicee.analyse_experiments.Experiment* attribute), 21
- attn (*dicee.models.Block* attribute), 144
- attn (*dicee.models.clifford.TransformerBlock* attribute), 97
- attn (*dicee.models.transformers.Block* attribute), 129
- attn_dropout (*dicee.models.clifford.TransformerSelfAttention* attribute), 97
- attn_dropout (*dicee.models.transformers.SelfAttention* attribute), 127
- attributes (*dicee.abstracts.AbstractTrainer* attribute), 13
- auto_batch_finding (*dicee.config.Namespace* attribute), 31

B

- backend (*dicee.config.Namespace* attribute), 29
- backend (*dicee.knowledge_graph.KG* attribute), 68
- base_weights (*dicee.weight_averaging.TWA* attribute), 224
- BaseInteractiveKGE (class in *dicee.abstracts*), 13
- BaseInteractiveTrainKGE (class in *dicee.abstracts*), 19
- BaseKGE (class in *dicee.models*), 139, 144, 153, 160, 168, 182, 187
- BaseKGE (class in *dicee.models.base_model*), 88
- BaseKGELightning (class in *dicee.models*), 135
- BaseKGELightning (class in *dicee.models.base_model*), 84
- batch_kronecker_product() (*dicee.callbacks.KronE* static method), 25
- batch_size (*dicee.analyse_experiments.Experiment* attribute), 21
- batch_size (*dicee.callbacks.PseudoLabellingCallback* attribute), 24
- batch_size (*dicee.config.Namespace* attribute), 28
- batches_per_epoch (*dicee.callbacks.LRScheduler* attribute), 27
- beta (*dicee.weight_averaging.TWA* attribute), 224
- bias (*dicee.models.CoKEConfig* attribute), 152
- bias (*dicee.models.real.CoKEConfig* attribute), 123
- bias (*dicee.models.transformers.GPTConfig* attribute), 129
- bias (*dicee.models.transformers.LayerNorm* attribute), 126
- Block (class in *dicee.models*), 143
- Block (class in *dicee.models.transformers*), 128
- block_size (*dicee.config.Namespace* attribute), 30
- block_size (*dicee.dataset_classes.MultiClassClassificationDataset* attribute), 33
- block_size (*dicee.models.base_model.BaseKGE* attribute), 89
- block_size (*dicee.models.BaseKGE* attribute), 140, 145, 155, 162, 169, 183, 188
- block_size (*dicee.models.CoKEConfig* attribute), 152

block_size (*dicee.models.real.CoKEConfig* attribute), 123
 block_size (*dicee.models.transformers.GPTConfig* attribute), 129
 blocks (*dicee.models.CoKE* attribute), 153
 blocks (*dicee.models.real.CoKE* attribute), 124
 bn_conv1 (*dicee.models.AConvQ* attribute), 167
 bn_conv1 (*dicee.models.ConvQ* attribute), 167
 bn_conv1 (*dicee.models.quaternion.AConvQ* attribute), 118
 bn_conv1 (*dicee.models.quaternion.ConvQ* attribute), 118
 bn_conv2 (*dicee.models.AConvQ* attribute), 167
 bn_conv2 (*dicee.models.ConvQ* attribute), 167
 bn_conv2 (*dicee.models.quaternion.AConvQ* attribute), 118
 bn_conv2 (*dicee.models.quaternion.ConvQ* attribute), 118
 bn_conv2d (*dicee.models.AConEx* attribute), 159
 bn_conv2d (*dicee.models.AConvO* attribute), 174
 bn_conv2d (*dicee.models.complex.AConEx* attribute), 102
 bn_conv2d (*dicee.models.complex.ConEx* attribute), 101
 bn_conv2d (*dicee.models.ConEx* attribute), 158
 bn_conv2d (*dicee.models.ConvO* attribute), 173
 bn_conv2d (*dicee.models.octonion.AConvO* attribute), 114
 bn_conv2d (*dicee.models.octonion.ConvO* attribute), 113
 BPE_NegativeSamplingDataset (class in *dicee.dataset_classes*), 32
 build_chain_funcs () (*dicee.models.FMult2* method), 193
 build_chain_funcs () (*dicee.models.function_space.FMult2* method), 108
 build_func () (*dicee.models.FMult2* method), 193
 build_func () (*dicee.models.function_space.FMult2* method), 108
 build_projection () (*dicee.weight_averaging.TWA* method), 224
 Byte (class in *dicee.models.transformers*), 125
 byte_pair_encoding (*dicee.analyse_experiments.Experiment* attribute), 21
 byte_pair_encoding (*dicee.config.Namespace* attribute), 30
 byte_pair_encoding (*dicee.knowledge_graph.KG* attribute), 68
 byte_pair_encoding (*dicee.models.base_model.BaseKGE* attribute), 89
 byte_pair_encoding (*dicee.models.BaseKGE* attribute), 140, 145, 154, 162, 169, 183, 188

C

c_attn (*dicee.models.clifford.TransformerSelfAttention* attribute), 97
 c_attn (*dicee.models.transformers.SelfAttention* attribute), 127
 c_fc (*dicee.models.clifford.TransformerMLP* attribute), 97
 c_fc (*dicee.models.transformers.MLP* attribute), 128
 c_proj (*dicee.models.clifford.TransformerMLP* attribute), 97
 c_proj (*dicee.models.clifford.TransformerSelfAttention* attribute), 97
 c_proj (*dicee.models.transformers.MLP* attribute), 128
 c_proj (*dicee.models.transformers.SelfAttention* attribute), 127
 callbacks (*dicee.abstracts.AbstractTrainer* attribute), 13
 callbacks (*dicee.analyse_experiments.Experiment* attribute), 21
 callbacks (*dicee.config.Namespace* attribute), 29
 callbacks (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 218
 causal (*dicee.models.CoKEConfig* attribute), 152
 causal (*dicee.models.real.CoKEConfig* attribute), 123
 causal (*dicee.models.transformers.GPTConfig* attribute), 129
 causal (*dicee.models.transformers.SelfAttention* attribute), 127
 chain_func () (*dicee.models.FMult* method), 191
 chain_func () (*dicee.models.function_space.FMult* method), 107
 chain_func () (*dicee.models.function_space.GFMult* method), 107
 chain_func () (*dicee.models.GFMult* method), 192
 CKeci (class in *dicee.models*), 178
 CKeci (class in *dicee.models.clifford*), 98
 cl_pqr () (*dicee.models.clifford.DeCaL* method), 98
 cl_pqr () (*dicee.models.DeCaL* method), 179
 cleanup () (*dicee.Execute* method), 226
 cleanup () (*dicee.executer.Execute* method), 66
 clifford_multiplication () (*dicee.models.clifford.Keci* method), 94
 clifford_multiplication () (*dicee.models.Keci* method), 176
 clip_lambda (*dicee.models.ADOPT* attribute), 134
 clip_lambda (*dicee.models.adopt.ADOPT* attribute), 79
 CoKE (class in *dicee.models*), 152
 CoKE (class in *dicee.models.real*), 123
 coke_dropout (*dicee.models.CoKE* attribute), 153

coke_dropout (*dicee.models.real.CoKE attribute*), 124
 CoKEConfig (*class in dicee.models*), 152
 CoKEConfig (*class in dicee.models.real*), 123
 collate_fn (*dicee.dataset_classes.AllvsAll attribute*), 34
 collate_fn (*dicee.dataset_classes.FixedNegSampleDataset attribute*), 38
 collate_fn (*dicee.dataset_classes.KvsAll attribute*), 35
 collate_fn (*dicee.dataset_classes.KvsSampleDataset attribute*), 36
 collate_fn (*dicee.dataset_classes.MultiClassClassificationDataset attribute*), 33
 collate_fn (*dicee.dataset_classes.MultiLabelDataset attribute*), 33
 collate_fn (*dicee.dataset_classes.OnevsAllDataset attribute*), 36
 collate_fn (*dicee.dataset_classes.OnevsSample attribute*), 39
 collate_fn () (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 32
 collate_fn () (*dicee.dataset_classes.TriplePredictionDataset method*), 39
 collection_name (*dicee.scripts.index_serve.NeuralSearcher attribute*), 205
 comp_func () (*dicee.models.function_space.LFMult method*), 110
 comp_func () (*dicee.models.LFMult method*), 195
 ComplEx (*class in dicee.models*), 159
 ComplEx (*class in dicee.models.complex*), 103
 compute_convergence () (*in module dicee.callbacks*), 24
 compute_func () (*dicee.models.FMult method*), 191
 compute_func () (*dicee.models.FMult2 method*), 193
 compute_func () (*dicee.models.function_space.FMult method*), 107
 compute_func () (*dicee.models.function_space.FMult2 method*), 108
 compute_func () (*dicee.models.function_space.GFMult method*), 107
 compute_func () (*dicee.models.GFMult method*), 192
 compute_metrics_from_ranks () (*in module dicee.evaluation.utils*), 53
 compute_metrics_from_ranks_simple () (*in module dicee.evaluation.utils*), 53
 compute_mrr () (*dicee.weight_averaging.ASWA static method*), 220
 compute_sigma_pp () (*dicee.models.clifford.DeCaL method*), 99
 compute_sigma_pp () (*dicee.models.clifford.Keci method*), 93
 compute_sigma_pp () (*dicee.models.DeCaL method*), 180
 compute_sigma_pp () (*dicee.models.Keci method*), 175
 compute_sigma_pq () (*dicee.models.clifford.DeCaL method*), 100
 compute_sigma_pq () (*dicee.models.clifford.Keci method*), 94
 compute_sigma_pq () (*dicee.models.DeCaL method*), 181
 compute_sigma_pq () (*dicee.models.Keci method*), 175
 compute_sigma_pr () (*dicee.models.clifford.DeCaL method*), 101
 compute_sigma_pr () (*dicee.models.DeCaL method*), 181
 compute_sigma_qq () (*dicee.models.clifford.DeCaL method*), 100
 compute_sigma_qq () (*dicee.models.clifford.Keci method*), 93
 compute_sigma_qq () (*dicee.models.DeCaL method*), 180
 compute_sigma_qq () (*dicee.models.Keci method*), 175
 compute_sigma_qr () (*dicee.models.clifford.DeCaL method*), 101
 compute_sigma_qr () (*dicee.models.DeCaL method*), 181
 compute_sigma_rr () (*dicee.models.clifford.DeCaL method*), 100
 compute_sigma_rr () (*dicee.models.DeCaL method*), 181
 compute_sigmas_multivect () (*dicee.models.clifford.DeCaL method*), 99
 compute_sigmas_multivect () (*dicee.models.DeCaL method*), 179
 compute_sigmas_single () (*dicee.models.clifford.DeCaL method*), 99
 compute_sigmas_single () (*dicee.models.DeCaL method*), 179
 ConEx (*class in dicee.models*), 157
 ConEx (*class in dicee.models.complex*), 101
 config (*dicee.models.CoKE attribute*), 153
 config (*dicee.models.real.CoKE attribute*), 124
 config (*dicee.models.transformers.BytE attribute*), 125
 config (*dicee.models.transformers.GPT attribute*), 130
 configs (*dicee.abstracts.BaseInteractiveKGE attribute*), 14
 configure_optimizers () (*dicee.models.base_model.BaseKGELightning method*), 88
 configure_optimizers () (*dicee.models.BaseKGELightning method*), 138
 configure_optimizers () (*dicee.models.transformers.GPT method*), 130
 construct_batch_selected_cl_multivector () (*dicee.models.clifford.Keci method*), 96
 construct_batch_selected_cl_multivector () (*dicee.models.Keci method*), 177
 construct_cl_multivector () (*dicee.models.clifford.DeCaL method*), 99
 construct_cl_multivector () (*dicee.models.clifford.Keci method*), 94
 construct_cl_multivector () (*dicee.models.DeCaL method*), 180
 construct_cl_multivector () (*dicee.models.Keci method*), 176
 construct_dataset () (*in module dicee.dataset_classes*), 34
 construct_ensemble (*dicee.abstracts.BaseInteractiveKGE attribute*), 14

`construct_graph()` (*dicee.query_generator.QueryGenerator* method), 197
`construct_graph()` (*dicee.QueryGenerator* method), 233
`construct_input_and_output()` (*dicee.abstracts.BaseInteractiveKGE* method), 17
`construct_multi_coeff()` (*dicee.models.function_space.LFMulti* method), 109
`construct_multi_coeff()` (*dicee.models.LFMulti* method), 194
`continual_learning` (*dicee.config.Namespace* attribute), 30
`continual_start()` (*dicee.DICE_Trainer* method), 234
`continual_start()` (*dicee.executer.ContinuousExecute* method), 67
`continual_start()` (*dicee.trainer.DICE_Trainer* method), 219
`continual_start()` (*dicee.trainer.dice_trainer.DICE_Trainer* method), 214
`continual_training_setup_executor()` (in module *dicee.static_funcs*), 210
`ContinuousExecute` (class in *dicee.executer*), 67
`conv2d` (*dicee.models.AConEx* attribute), 159
`conv2d` (*dicee.models.AConvO* attribute), 174
`conv2d` (*dicee.models.AConvQ* attribute), 167
`conv2d` (*dicee.models.complex.AConEx* attribute), 102
`conv2d` (*dicee.models.complex.ConEx* attribute), 101
`conv2d` (*dicee.models.ConEx* attribute), 158
`conv2d` (*dicee.models.ConvO* attribute), 173
`conv2d` (*dicee.models.ConvQ* attribute), 166
`conv2d` (*dicee.models.octonion.AConvO* attribute), 114
`conv2d` (*dicee.models.octonion.ConvO* attribute), 113
`conv2d` (*dicee.models.quaternion.AConvQ* attribute), 118
`conv2d` (*dicee.models.quaternion.ConvQ* attribute), 118
`ConvO` (class in *dicee.models*), 173
`ConvO` (class in *dicee.models.octonion*), 113
`ConvQ` (class in *dicee.models*), 166
`ConvQ` (class in *dicee.models.quaternion*), 117
`count_triples()` (in module *dicee.read_preprocess_save_load_kg.util*), 202
`create_and_store_kg()` (*dicee.Execute* method), 226
`create_and_store_kg()` (*dicee.executer.Execute* method), 66
`create_constraints()` (in module *dicee.read_preprocess_save_load_kg.util*), 202
`create_constraints()` (in module *dicee.static_preprocess_funcs*), 212
`create_experiment_folder()` (in module *dicee.static_funcs*), 210
`create_hits_dict()` (in module *dicee.evaluation.utils*), 53
`create_random_data()` (*dicee.callbacks.PseudoLabellingCallback* method), 24
`create_recipriocal_triples()` (in module *dicee.read_preprocess_save_load_kg.util*), 202
`create_recipriocal_triples()` (in module *dicee.static_funcs*), 207
`create_vector_database()` (*dicee.KGE* method), 227
`create_vector_database()` (*dicee.knowledge_graph_embeddings.KGE* method), 70
`crop_block_size()` (*dicee.models.transformers.GPT* method), 130
`ctx` (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 218
`current_epoch` (*dicee.weight_averaging.EMA* attribute), 223
`current_epoch` (*dicee.weight_averaging.SWA* attribute), 221
`current_epoch` (*dicee.weight_averaging.SWAG* attribute), 223
`current_epoch` (*dicee.weight_averaging.TWA* attribute), 224
`cycle_length` (*dicee.callbacks.LRScheduler* attribute), 27

D

`data_module` (*dicee.callbacks.PseudoLabellingCallback* attribute), 24
`data_property_embeddings` (*dicee.models.literal.LiteralEmbeddings* attribute), 111
`dataset_dir` (*dicee.config.Namespace* attribute), 28
`dataset_dir` (*dicee.knowledge_graph.KG* attribute), 67, 68
`dataset_sanity_checking()` (in module *dicee.read_preprocess_save_load_kg.util*), 202
`DeCaL` (class in *dicee.models*), 178
`DeCaL` (class in *dicee.models.clifford*), 98
`decay` (*dicee.weight_averaging.EMA* attribute), 223
`decide()` (*dicee.weight_averaging.ASWA* method), 221
`default_eval_model` (*dicee.callbacks.PeriodicEvalCallback* attribute), 27
`DEFAULT_HITS_RANGE` (in module *dicee.evaluation.utils*), 52
`degree` (*dicee.models.function_space.LFMulti* attribute), 109
`degree` (*dicee.models.LFMulti* attribute), 194
`denormalize()` (*dicee.dataset_classes.LiteralDataset* static method), 37
`describe()` (*dicee.knowledge_graph.KG* method), 69
`description_of_input` (*dicee.knowledge_graph.KG* attribute), 69
`deviations` (*dicee.weight_averaging.SWAG* attribute), 222
`device` (*dicee.models.literal.LiteralEmbeddings* property), 112

- DICE_Trainer (*class in dicee*), 234
- DICE_Trainer (*class in dicee.trainer*), 218
- DICE_Trainer (*class in dicee.trainer.dice_trainer*), 214
- dicee
 - module, 12
- dicee.__main__
 - module, 12
- dicee.abstracts
 - module, 12
- dicee.analyse_experiments
 - module, 20
- dicee.callbacks
 - module, 22
- dicee.config
 - module, 28
- dicee.dataset_classes
 - module, 31
- dicee.eval_static_funcs
 - module, 39
- dicee.evaluation
 - module, 43
- dicee.evaluation.ensemble
 - module, 43
- dicee.evaluation.evaluator
 - module, 44
- dicee.evaluation.link_prediction
 - module, 48
- dicee.evaluation.literal_prediction
 - module, 51
- dicee.evaluation.utils
 - module, 52
- dicee.evaluator
 - module, 61
- dicee.executer
 - module, 64
- dicee.knowledge_graph
 - module, 67
- dicee.knowledge_graph_embeddings
 - module, 70
- dicee.models
 - module, 75
- dicee.models.adopt
 - module, 75
- dicee.models.base_model
 - module, 84
- dicee.models.clifford
 - module, 92
- dicee.models.complex
 - module, 101
- dicee.models.dualE
 - module, 104
- dicee.models.ensemble
 - module, 105
- dicee.models.function_space
 - module, 106
- dicee.models.literal
 - module, 110
- dicee.models.octonion
 - module, 112
- dicee.models.pykeen_models
 - module, 115
- dicee.models.quaternion
 - module, 116
- dicee.models.real
 - module, 119
- dicee.models.static_funcs
 - module, 124
- dicee.models.transformers

- module, 125
- dicee.query_generator
 - module, 196
- dicee.read_preprocess_save_load_kg
 - module, 197
- dicee.read_preprocess_save_load_kg.preprocess
 - module, 197
- dicee.read_preprocess_save_load_kg.read_from_disk
 - module, 198
- dicee.read_preprocess_save_load_kg.save_load_disk
 - module, 199
- dicee.read_preprocess_save_load_kg.util
 - module, 199
- dicee.sanity_checkers
 - module, 204
- dicee.scripts
 - module, 204
- dicee.scripts.index_serve
 - module, 204
- dicee.scripts.run
 - module, 206
- dicee.static_funcs
 - module, 206
- dicee.static_funcs_training
 - module, 210
- dicee.static_preprocess_funcs
 - module, 212
- dicee.trainer
 - module, 213
- dicee.trainer.dice_trainer
 - module, 213
- dicee.trainer.model_parallelism
 - module, 215
- dicee.trainer.torch_trainer
 - module, 216
- dicee.trainer.torch_trainer_ddp
 - module, 217
- dicee.weight_averaging
 - module, 220
- discrete_points (*dicee.models.FMult2 attribute*), 193
- discrete_points (*dicee.models.function_space.FMult2 attribute*), 108
- dist_func (*dicee.models.Pyke attribute*), 151
- dist_func (*dicee.models.real.Pyke attribute*), 122
- DistMult (*class in dicee.models*), 148
- DistMult (*class in dicee.models.real*), 119
- distributed (*dicee.Execute attribute*), 225, 226
- distributed (*dicee.executer.Execute attribute*), 65
- domain_constraints_per_rel (*dicee.evaluation.Evaluator attribute*), 56
- domain_constraints_per_rel (*dicee.evaluation.evaluator.Evaluator attribute*), 46
- domain_constraints_per_rel (*dicee.Evaluator attribute*), 236
- domain_constraints_per_rel (*dicee.evaluator.Evaluator attribute*), 62
- download_file() (*in module dicee.static_funcs*), 210
- download_files_from_url() (*in module dicee.static_funcs*), 210
- download_pretrained_model() (*in module dicee.static_funcs*), 210
- dropout (*dicee.models.clifford.TransformerMLP attribute*), 97
- dropout (*dicee.models.clifford.TransformerSelfAttention attribute*), 97
- dropout (*dicee.models.CoKEConfig attribute*), 152
- dropout (*dicee.models.literal.LiteralEmbeddings attribute*), 111
- dropout (*dicee.models.real.CoKEConfig attribute*), 123
- dropout (*dicee.models.transformers.GPTConfig attribute*), 129
- dropout (*dicee.models.transformers.MLP attribute*), 128
- dropout (*dicee.models.transformers.SelfAttention attribute*), 127
- DualE (*class in dicee.models*), 195
- DualE (*class in dicee.models.dualE*), 104
- dummy_eval() (*dicee.evaluation.Evaluator method*), 57
- dummy_eval() (*dicee.evaluation.evaluator.Evaluator method*), 47
- dummy_eval() (*dicee.Evaluator method*), 238
- dummy_eval() (*dicee.evaluator.Evaluator method*), 64

dummy_id (*dicee.knowledge_graph.KG attribute*), 69
 during_training (*dicee.evaluation.Evaluator attribute*), 55, 56
 during_training (*dicee.evaluation.evaluator.Evaluator attribute*), 45, 46
 during_training (*dicee.Evaluator attribute*), 236
 during_training (*dicee.evaluator.Evaluator attribute*), 62

E

ee_vocab (*dicee.evaluation.Evaluator attribute*), 55
 ee_vocab (*dicee.evaluation.evaluator.Evaluator attribute*), 45
 ee_vocab (*dicee.Evaluator attribute*), 235, 236
 ee_vocab (*dicee.evaluator.Evaluator attribute*), 61, 62
 efficient_zero_grad() (*in module dicee.evaluation.utils*), 53
 efficient_zero_grad() (*in module dicee.static_funcs_training*), 211
 EMA (*class in dicee.weight_averaging*), 223
 ema (*dicee.config.Namespace attribute*), 30
 ema_c_epochs (*dicee.weight_averaging.EMA attribute*), 223
 ema_model (*dicee.weight_averaging.EMA attribute*), 223
 ema_start_epoch (*dicee.weight_averaging.EMA attribute*), 223
 ema_update() (*dicee.weight_averaging.EMA static method*), 223
 embedding_dim (*dicee.analyse_experiments.Experiment attribute*), 21
 embedding_dim (*dicee.config.Namespace attribute*), 28
 embedding_dim (*dicee.models.base_model.BaseKGE attribute*), 88
 embedding_dim (*dicee.models.BaseKGE attribute*), 139, 144, 154, 161, 168, 182, 187
 embedding_dim (*dicee.models.literal.LiteralEmbeddings attribute*), 111
 embedding_dims (*dicee.models.literal.LiteralEmbeddings attribute*), 111
 enable_log (*in module dicee.static_preprocess_funcs*), 212
 enc (*dicee.knowledge_graph.KG attribute*), 69
 end() (*dicee.Execute method*), 227
 end() (*dicee.executer.Execute method*), 66
 EnsembleKGE (*class in dicee.models.ensemble*), 105
 ent2id (*dicee.query_generator.QueryGenerator attribute*), 196
 ent2id (*dicee.QueryGenerator attribute*), 233
 ent_in (*dicee.query_generator.QueryGenerator attribute*), 196
 ent_in (*dicee.QueryGenerator attribute*), 233
 ent_out (*dicee.query_generator.QueryGenerator attribute*), 196
 ent_out (*dicee.QueryGenerator attribute*), 233
 entities_str (*dicee.knowledge_graph.KG property*), 69
 entity_embeddings (*dicee.models.AConvQ attribute*), 167
 entity_embeddings (*dicee.models.clifford.DeCaL attribute*), 98
 entity_embeddings (*dicee.models.ConvQ attribute*), 166
 entity_embeddings (*dicee.models.DeCaL attribute*), 178
 entity_embeddings (*dicee.models.DualE attribute*), 195
 entity_embeddings (*dicee.models.dualE.DualE attribute*), 104
 entity_embeddings (*dicee.models.FMult attribute*), 191
 entity_embeddings (*dicee.models.FMult2 attribute*), 193
 entity_embeddings (*dicee.models.function_space.FMult attribute*), 107
 entity_embeddings (*dicee.models.function_space.FMult2 attribute*), 108
 entity_embeddings (*dicee.models.function_space.GFMult attribute*), 107
 entity_embeddings (*dicee.models.function_space.LFMult attribute*), 109
 entity_embeddings (*dicee.models.function_space.LFMult1 attribute*), 108
 entity_embeddings (*dicee.models.GFMult attribute*), 192
 entity_embeddings (*dicee.models.LFMult attribute*), 194
 entity_embeddings (*dicee.models.LFMult1 attribute*), 193
 entity_embeddings (*dicee.models.literal.LiteralEmbeddings attribute*), 111
 entity_embeddings (*dicee.models.pykeen_models.PykeenKGE attribute*), 115
 entity_embeddings (*dicee.models.PykeenKGE attribute*), 186
 entity_embeddings (*dicee.models.quaternion.AConvQ attribute*), 118
 entity_embeddings (*dicee.models.quaternion.ConvQ attribute*), 117
 entity_to_idx (*dicee.dataset_classes.LiteralDataset attribute*), 37
 entity_to_idx (*dicee.knowledge_graph.KG attribute*), 68
 entity_to_idx (*dicee.scripts.index_serve.NeuralSearcher attribute*), 205
 epoch_count (*dicee.abstracts.AbstractPPECallback attribute*), 19
 epoch_count (*dicee.weight_averaging.ASWA attribute*), 220
 epoch_counter (*dicee.callbacks.Eval attribute*), 25
 epoch_counter (*dicee.callbacks.KGESaveCallback attribute*), 23
 epoch_counter (*dicee.callbacks.PeriodicEvalCallback attribute*), 26
 epoch_ratio (*dicee.callbacks.Eval attribute*), 25

er_vocab (*dicee.evaluation.Evaluator attribute*), 55
 er_vocab (*dicee.evaluation.evaluator.Evaluator attribute*), 45
 er_vocab (*dicee.Evaluator attribute*), 235, 236
 er_vocab (*dicee.evaluator.Evaluator attribute*), 61, 62
 estimate_mfu () (*dicee.models.transformers.GPT method*), 130
 estimate_q () (*in module dicee.callbacks*), 24
 Eval (*class in dicee.callbacks*), 24
 eval () (*dicee.evaluation.Evaluator method*), 56
 eval () (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval () (*dicee.Evaluator method*), 236
 eval () (*dicee.evaluator.Evaluator method*), 62
 eval () (*dicee.models.ensemble.EnsembleKGE method*), 106
 eval_at_epochs (*dicee.config.Namespace attribute*), 31
 eval_epochs (*dicee.callbacks.PeriodicEvalCallback attribute*), 27
 eval_every_n_epochs (*dicee.config.Namespace attribute*), 31
 eval_lp_performance () (*dicee.KGE method*), 228
 eval_lp_performance () (*dicee.knowledge_graph_embeddings.KGE method*), 70
 eval_model (*dicee.config.Namespace attribute*), 29
 eval_model (*dicee.knowledge_graph.KG attribute*), 68
 eval_rank_of_head_and_tail_byte_pair_encoded_entity () (*dicee.evaluation.Evaluator method*), 56
 eval_rank_of_head_and_tail_byte_pair_encoded_entity () (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_rank_of_head_and_tail_byte_pair_encoded_entity () (*dicee.Evaluator method*), 237
 eval_rank_of_head_and_tail_byte_pair_encoded_entity () (*dicee.evaluator.Evaluator method*), 63
 eval_rank_of_head_and_tail_entity () (*dicee.evaluation.Evaluator method*), 56
 eval_rank_of_head_and_tail_entity () (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_rank_of_head_and_tail_entity () (*dicee.Evaluator method*), 237
 eval_rank_of_head_and_tail_entity () (*dicee.evaluator.Evaluator method*), 63
 eval_with_bpe_vs_all () (*dicee.evaluation.Evaluator method*), 56
 eval_with_bpe_vs_all () (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_with_bpe_vs_all () (*dicee.Evaluator method*), 237
 eval_with_bpe_vs_all () (*dicee.evaluator.Evaluator method*), 63
 eval_with_byte () (*dicee.evaluation.Evaluator method*), 56
 eval_with_byte () (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_with_byte () (*dicee.Evaluator method*), 237
 eval_with_byte () (*dicee.evaluator.Evaluator method*), 63
 eval_with_data () (*dicee.evaluation.Evaluator method*), 58
 eval_with_data () (*dicee.evaluation.evaluator.Evaluator method*), 48
 eval_with_data () (*dicee.Evaluator method*), 238
 eval_with_data () (*dicee.evaluator.Evaluator method*), 64
 eval_with_vs_all () (*dicee.evaluation.Evaluator method*), 56
 eval_with_vs_all () (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_with_vs_all () (*dicee.Evaluator method*), 237
 eval_with_vs_all () (*dicee.evaluator.Evaluator method*), 63
 evaluate () (*in module dicee.static_funcs*), 210
 evaluate_bpe_lp () (*in module dicee.evaluation*), 58
 evaluate_bpe_lp () (*in module dicee.evaluation.link_prediction*), 50
 evaluate_bpe_lp () (*in module dicee.static_funcs_training*), 211
 evaluate_ensemble_link_prediction_performance () (*in module dicee.eval_static_funcs*), 40
 evaluate_ensemble_link_prediction_performance () (*in module dicee.evaluation*), 54
 evaluate_ensemble_link_prediction_performance () (*in module dicee.evaluation.ensemble*), 44
 evaluate_link_prediction_performance () (*in module dicee.eval_static_funcs*), 41
 evaluate_link_prediction_performance () (*in module dicee.evaluation*), 58
 evaluate_link_prediction_performance () (*in module dicee.evaluation.link_prediction*), 48
 evaluate_link_prediction_performance_with_bpe () (*in module dicee.eval_static_funcs*), 41
 evaluate_link_prediction_performance_with_bpe () (*in module dicee.evaluation*), 59
 evaluate_link_prediction_performance_with_bpe () (*in module dicee.evaluation.link_prediction*), 49
 evaluate_link_prediction_performance_with_bpe_reciprocals () (*in module dicee.eval_static_funcs*), 41
 evaluate_link_prediction_performance_with_bpe_reciprocals () (*in module dicee.evaluation*), 59
 evaluate_link_prediction_performance_with_bpe_reciprocals () (*in module dicee.evaluation.link_prediction*), 49
 evaluate_link_prediction_performance_with_reciprocals () (*in module dicee.eval_static_funcs*), 41
 evaluate_link_prediction_performance_with_reciprocals () (*in module dicee.evaluation*), 59
 evaluate_link_prediction_performance_with_reciprocals () (*in module dicee.evaluation.link_prediction*), 49
 evaluate_literal_prediction () (*in module dicee.eval_static_funcs*), 42
 evaluate_literal_prediction () (*in module dicee.evaluation*), 60
 evaluate_literal_prediction () (*in module dicee.evaluation.literal_prediction*), 51
 evaluate_lp () (*dicee.evaluation.Evaluator method*), 57
 evaluate_lp () (*dicee.evaluation.evaluator.Evaluator method*), 47
 evaluate_lp () (*dicee.Evaluator method*), 238

[evaluate_lp\(\)](#) (*dicee.evaluator.Evaluator method*), 64
[evaluate_lp\(\)](#) (*in module dicee.evaluation*), 59
[evaluate_lp\(\)](#) (*in module dicee.evaluation.link_prediction*), 49
[evaluate_lp\(\)](#) (*in module dicee.static_funcs_training*), 211
[evaluate_lp_bpe_k_vs_all\(\)](#) (*dicee.evaluation.Evaluator method*), 57
[evaluate_lp_bpe_k_vs_all\(\)](#) (*dicee.evaluation.evaluator.Evaluator method*), 47
[evaluate_lp_bpe_k_vs_all\(\)](#) (*dicee.Evaluator method*), 237
[evaluate_lp_bpe_k_vs_all\(\)](#) (*dicee.evaluator.Evaluator method*), 63
[evaluate_lp_bpe_k_vs_all\(\)](#) (*in module dicee.eval_static_funcs*), 42
[evaluate_lp_bpe_k_vs_all\(\)](#) (*in module dicee.evaluation*), 60
[evaluate_lp_bpe_k_vs_all\(\)](#) (*in module dicee.evaluation.link_prediction*), 50
[evaluate_lp_k_vs_all\(\)](#) (*dicee.evaluation.Evaluator method*), 57
[evaluate_lp_k_vs_all\(\)](#) (*dicee.evaluation.evaluator.Evaluator method*), 47
[evaluate_lp_k_vs_all\(\)](#) (*dicee.Evaluator method*), 237
[evaluate_lp_k_vs_all\(\)](#) (*dicee.evaluator.Evaluator method*), 63
[evaluate_lp_with_byte\(\)](#) (*dicee.evaluation.Evaluator method*), 57
[evaluate_lp_with_byte\(\)](#) (*dicee.evaluation.evaluator.Evaluator method*), 47
[evaluate_lp_with_byte\(\)](#) (*dicee.Evaluator method*), 237
[evaluate_lp_with_byte\(\)](#) (*dicee.evaluator.Evaluator method*), 63
[Evaluator](#) (*class in dicee*), 235
[Evaluator](#) (*class in dicee.evaluation*), 55
[Evaluator](#) (*class in dicee.evaluation.evaluator*), 45
[Evaluator](#) (*class in dicee.evaluator*), 61
[evaluator](#) (*dicee.DICE_Trainer attribute*), 234
[evaluator](#) (*dicee.Execute attribute*), 226
[evaluator](#) (*dicee.executer.Execute attribute*), 65, 66
[evaluator](#) (*dicee.trainer.DICE_Trainer attribute*), 219
[evaluator](#) (*dicee.trainer.dice_trainer.DICE_Trainer attribute*), 214
[every_x_epoch](#) (*dicee.callbacks.KGESaveCallback attribute*), 23
[example_input_array](#) (*dicee.models.ensemble.EnsembleKGE property*), 106
[Execute](#) (*class in dicee*), 225
[Execute](#) (*class in dicee.executer*), 65
[exists\(\)](#) (*dicee.knowledge_graph.KG method*), 69
[Experiment](#) (*class in dicee.analyse_experiments*), 21
[experiment_dir](#) (*dicee.callbacks.LRScheduler attribute*), 27
[experiment_dir](#) (*dicee.callbacks.PeriodicEvalCallback attribute*), 26
[explicit](#) (*dicee.models.QMult attribute*), 165
[explicit](#) (*dicee.models.quaternion.QMult attribute*), 116
[exponential_function\(\)](#) (*in module dicee.static_funcs*), 210
[extract_input_outputs\(\)](#) (*dicee.trainer.torch_trainer_ddp.NodeTrainer method*), 218
[extract_input_outputs\(\)](#) (*in module dicee.trainer.model_parallelism*), 215
[extract_input_outputs_set_device\(\)](#) (*dicee.trainer.torch_trainer.TorchTrainer method*), 217

F

[f](#) (*dicee.callbacks.KronE attribute*), 25
[fc](#) (*dicee.models.literal.LiteralEmbeddings attribute*), 111
[fc1](#) (*dicee.models.AConEx attribute*), 159
[fc1](#) (*dicee.models.AConvO attribute*), 174
[fc1](#) (*dicee.models.AConvQ attribute*), 167
[fc1](#) (*dicee.models.complex.AConEx attribute*), 102
[fc1](#) (*dicee.models.complex.ConEx attribute*), 101
[fc1](#) (*dicee.models.ConEx attribute*), 158
[fc1](#) (*dicee.models.ConvO attribute*), 173
[fc1](#) (*dicee.models.ConvQ attribute*), 166
[fc1](#) (*dicee.models.octonion.AConvO attribute*), 114
[fc1](#) (*dicee.models.octonion.ConvO attribute*), 113
[fc1](#) (*dicee.models.quaternion.AConvQ attribute*), 118
[fc1](#) (*dicee.models.quaternion.ConvQ attribute*), 118
[fc_num_input](#) (*dicee.models.AConEx attribute*), 159
[fc_num_input](#) (*dicee.models.AConvO attribute*), 174
[fc_num_input](#) (*dicee.models.AConvQ attribute*), 167
[fc_num_input](#) (*dicee.models.complex.AConEx attribute*), 102
[fc_num_input](#) (*dicee.models.complex.ConEx attribute*), 101
[fc_num_input](#) (*dicee.models.ConEx attribute*), 158
[fc_num_input](#) (*dicee.models.ConvO attribute*), 173
[fc_num_input](#) (*dicee.models.ConvQ attribute*), 166
[fc_num_input](#) (*dicee.models.octonion.AConvO attribute*), 114

`fc_num_input` (*dicее.models.octonion.ConvO attribute*), 113
`fc_num_input` (*dicее.models.quaternion.AConvQ attribute*), 118
`fc_num_input` (*dicее.models.quaternion.ConvQ attribute*), 118
`fc_out` (*dicее.models.literal.LiteralEmbeddings attribute*), 111
`feature_map_dropout` (*dicее.models.AConEx attribute*), 159
`feature_map_dropout` (*dicее.models.AConvO attribute*), 174
`feature_map_dropout` (*dicее.models.AConvQ attribute*), 167
`feature_map_dropout` (*dicее.models.complex.AConEx attribute*), 102
`feature_map_dropout` (*dicее.models.complex.ConEx attribute*), 101
`feature_map_dropout` (*dicее.models.ConEx attribute*), 158
`feature_map_dropout` (*dicее.models.ConvO attribute*), 173
`feature_map_dropout` (*dicее.models.ConvQ attribute*), 167
`feature_map_dropout` (*dicее.models.octonion.AConvO attribute*), 114
`feature_map_dropout` (*dicее.models.octonion.ConvO attribute*), 113
`feature_map_dropout` (*dicее.models.quaternion.AConvQ attribute*), 118
`feature_map_dropout` (*dicее.models.quaternion.ConvQ attribute*), 118
`feature_map_dropout_rate` (*dicее.config.Namespace attribute*), 30
`feature_map_dropout_rate` (*dicее.models.base_model.BaseKGE attribute*), 89
`feature_map_dropout_rate` (*dicее.models.BaseKGE attribute*), 139, 145, 154, 161, 168, 182, 188
`fetch_worker()` (in module *dicее.read_preprocess_save_load_kg.util*), 202
`fill_query()` (*dicее.query_generator.QueryGenerator method*), 197
`fill_query()` (*dicее.QueryGenerator method*), 233
`find_good_batch_size()` (in module *dicее.trainer.model_parallelism*), 215
`find_missing_triples()` (*dicее.KGE method*), 232
`find_missing_triples()` (*dicее.knowledge_graph_embeddings.KGE method*), 74
`fit()` (*dicее.trainer.model_parallelism.TensorParallel method*), 216
`fit()` (*dicее.trainer.torch_trainer_ddp.TorchDDPTrainer method*), 218
`fit()` (*dicее.trainer.torch_trainer.TorchTrainer method*), 216
`FixedNegSampleDataset` (class in *dicее.dataset_classes*), 37
`flash` (*dicее.models.transformers.SelfAttention attribute*), 127
`FMult` (class in *dicее.models*), 191
`FMult` (class in *dicее.models.function_space*), 106
`FMult2` (class in *dicее.models*), 192
`FMult2` (class in *dicее.models.function_space*), 108
`form_of_labelling` (*dicее.DICE_Trainer attribute*), 234
`form_of_labelling` (*dicее.trainer.DICE_Trainer attribute*), 219
`form_of_labelling` (*dicее.trainer.dice_trainer.DICE_Trainer attribute*), 214
`forward()` (*dicее.models.base_model.BaseKGE method*), 90
`forward()` (*dicее.models.base_model.IdentityClass static method*), 92
`forward()` (*dicее.models.BaseKGE method*), 141, 146, 155, 162, 169, 184, 189
`forward()` (*dicее.models.Block method*), 144
`forward()` (*dicее.models.clifford.TransformerBlock method*), 97
`forward()` (*dicее.models.clifford.TransformerMLP method*), 98
`forward()` (*dicее.models.clifford.TransformerSelfAttention method*), 97
`forward()` (*dicее.models.IdentityClass static method*), 143, 165, 172
`forward()` (*dicее.models.literal.LiteralEmbeddings method*), 112
`forward()` (*dicее.models.transformers.Block method*), 129
`forward()` (*dicее.models.transformers.BytE method*), 126
`forward()` (*dicее.models.transformers.GPT method*), 130
`forward()` (*dicее.models.transformers.LayerNorm method*), 126
`forward()` (*dicее.models.transformers.MLP method*), 128
`forward()` (*dicее.models.transformers.SelfAttention method*), 127
`forward_backward_update()` (*dicее.trainer.torch_trainer.TorchTrainer method*), 217
`forward_backward_update_loss()` (in module *dicее.trainer.model_parallelism*), 215
`forward_byte_pair_encoded_k_vs_all()` (*dicее.models.base_model.BaseKGE method*), 89
`forward_byte_pair_encoded_k_vs_all()` (*dicее.models.BaseKGE method*), 140, 145, 155, 162, 169, 183, 188
`forward_byte_pair_encoded_triple()` (*dicее.models.base_model.BaseKGE method*), 89
`forward_byte_pair_encoded_triple()` (*dicее.models.BaseKGE method*), 140, 145, 155, 162, 169, 183, 188
`forward_k_vs_all()` (*dicее.models.AConEx method*), 159
`forward_k_vs_all()` (*dicее.models.AConvO method*), 174
`forward_k_vs_all()` (*dicее.models.AConvQ method*), 168
`forward_k_vs_all()` (*dicее.models.base_model.BaseKGE method*), 90
`forward_k_vs_all()` (*dicее.models.BaseKGE method*), 141, 146, 156, 163, 170, 184, 190
`forward_k_vs_all()` (*dicее.models.clifford.DeCaL method*), 99
`forward_k_vs_all()` (*dicее.models.clifford.Keci method*), 95
`forward_k_vs_all()` (*dicее.models.clifford.KeciTransformer method*), 97
`forward_k_vs_all()` (*dicее.models.CoKE method*), 153
`forward_k_vs_all()` (*dicее.models.ComplEx method*), 160

`forward_k_vs_all()` (*dicee.models.complex.AConEx method*), 103
`forward_k_vs_all()` (*dicee.models.complex.ComplEx method*), 104
`forward_k_vs_all()` (*dicee.models.complex.ConEx method*), 102
`forward_k_vs_all()` (*dicee.models.ConEx method*), 158
`forward_k_vs_all()` (*dicee.models.ConvO method*), 173
`forward_k_vs_all()` (*dicee.models.ConvQ method*), 167
`forward_k_vs_all()` (*dicee.models.DeCaL method*), 179
`forward_k_vs_all()` (*dicee.models.DistMult method*), 149
`forward_k_vs_all()` (*dicee.models.DualE method*), 196
`forward_k_vs_all()` (*dicee.models.dualE.DualE method*), 105
`forward_k_vs_all()` (*dicee.models.Keci method*), 177
`forward_k_vs_all()` (*dicee.models.KeciTransformer method*), 182
`forward_k_vs_all()` (*dicee.models.octonion.AConvO method*), 114
`forward_k_vs_all()` (*dicee.models.octonion.ConvO method*), 114
`forward_k_vs_all()` (*dicee.models.octonion.OMult method*), 113
`forward_k_vs_all()` (*dicee.models.OMult method*), 172
`forward_k_vs_all()` (*dicee.models.pykeen_models.PykeenKGE method*), 115
`forward_k_vs_all()` (*dicee.models.PykeenKGE method*), 186
`forward_k_vs_all()` (*dicee.models.QMult method*), 166
`forward_k_vs_all()` (*dicee.models.quaternion.AConvQ method*), 119
`forward_k_vs_all()` (*dicee.models.quaternion.ConvQ method*), 118
`forward_k_vs_all()` (*dicee.models.quaternion.QMult method*), 117
`forward_k_vs_all()` (*dicee.models.real.CoKE method*), 124
`forward_k_vs_all()` (*dicee.models.real.DistMult method*), 120
`forward_k_vs_all()` (*dicee.models.real.Shallom method*), 122
`forward_k_vs_all()` (*dicee.models.real.TransE method*), 121
`forward_k_vs_all()` (*dicee.models.Shallom method*), 151
`forward_k_vs_all()` (*dicee.models.TransE method*), 150
`forward_k_vs_sample()` (*dicee.models.AConEx method*), 159
`forward_k_vs_sample()` (*dicee.models.base_model.BaseKGE method*), 91
`forward_k_vs_sample()` (*dicee.models.BaseKGE method*), 141, 147, 156, 163, 170, 184, 190
`forward_k_vs_sample()` (*dicee.models.clifford.Keci method*), 96
`forward_k_vs_sample()` (*dicee.models.CoKE method*), 153
`forward_k_vs_sample()` (*dicee.models.ComplEx method*), 160
`forward_k_vs_sample()` (*dicee.models.complex.AConEx method*), 103
`forward_k_vs_sample()` (*dicee.models.complex.ComplEx method*), 104
`forward_k_vs_sample()` (*dicee.models.complex.ConEx method*), 102
`forward_k_vs_sample()` (*dicee.models.ConEx method*), 158
`forward_k_vs_sample()` (*dicee.models.DistMult method*), 149
`forward_k_vs_sample()` (*dicee.models.Keci method*), 177
`forward_k_vs_sample()` (*dicee.models.pykeen_models.PykeenKGE method*), 116
`forward_k_vs_sample()` (*dicee.models.PykeenKGE method*), 187
`forward_k_vs_sample()` (*dicee.models.QMult method*), 166
`forward_k_vs_sample()` (*dicee.models.quaternion.QMult method*), 117
`forward_k_vs_sample()` (*dicee.models.real.CoKE method*), 124
`forward_k_vs_sample()` (*dicee.models.real.DistMult method*), 120
`forward_k_vs_with_explicit()` (*dicee.models.clifford.Keci method*), 95
`forward_k_vs_with_explicit()` (*dicee.models.Keci method*), 176
`forward_triples()` (*dicee.models.AConEx method*), 159
`forward_triples()` (*dicee.models.AConvO method*), 174
`forward_triples()` (*dicee.models.AConvQ method*), 167
`forward_triples()` (*dicee.models.base_model.BaseKGE method*), 90
`forward_triples()` (*dicee.models.BaseKGE method*), 141, 146, 156, 163, 170, 184, 189
`forward_triples()` (*dicee.models.clifford.DeCaL method*), 98
`forward_triples()` (*dicee.models.clifford.Keci method*), 96
`forward_triples()` (*dicee.models.complex.AConEx method*), 103
`forward_triples()` (*dicee.models.complex.ConEx method*), 102
`forward_triples()` (*dicee.models.ConEx method*), 158
`forward_triples()` (*dicee.models.ConvO method*), 173
`forward_triples()` (*dicee.models.ConvQ method*), 167
`forward_triples()` (*dicee.models.DeCaL method*), 179
`forward_triples()` (*dicee.models.DualE method*), 195
`forward_triples()` (*dicee.models.dualE.DualE method*), 105
`forward_triples()` (*dicee.models.FMult method*), 191
`forward_triples()` (*dicee.models.FMult2 method*), 193
`forward_triples()` (*dicee.models.function_space.FMult method*), 107
`forward_triples()` (*dicee.models.function_space.FMult2 method*), 108
`forward_triples()` (*dicee.models.function_space.GFMult method*), 107

`forward_triples()` (*dicee.models.function_space.LFMult method*), 109
`forward_triples()` (*dicee.models.function_space.LFMult1 method*), 109
`forward_triples()` (*dicee.models.GFMult method*), 192
`forward_triples()` (*dicee.models.Keci method*), 178
`forward_triples()` (*dicee.models.LFMult method*), 194
`forward_triples()` (*dicee.models.LFMult1 method*), 193
`forward_triples()` (*dicee.models.octonion.AConvO method*), 114
`forward_triples()` (*dicee.models.octonion.ConvO method*), 113
`forward_triples()` (*dicee.models.Pyke method*), 151
`forward_triples()` (*dicee.models.pykeen_models.PykeenKGE method*), 115
`forward_triples()` (*dicee.models.PykeenKGE method*), 187
`forward_triples()` (*dicee.models.quaternion.AConvQ method*), 118
`forward_triples()` (*dicee.models.quaternion.ConvQ method*), 118
`forward_triples()` (*dicee.models.real.Pyke method*), 122
`forward_triples()` (*dicee.models.real.Shallom method*), 122
`forward_triples()` (*dicee.models.Shallom method*), 151
`freeze_entity_embeddings` (*dicee.models.literal.LiteralEmbeddings attribute*), 111
`frequency` (*dicee.callbacks.Perturb attribute*), 26
`from_pretrained()` (*dicee.models.transformers.GPT class method*), 130
`from_pretrained_model_write_embeddings_into_csv()` (*in module dicee.static_funcs*), 210
`full_storage_path` (*dicee.analyse_experiments.Experiment attribute*), 21
`func_triple_to_bpe_representation` (*dicee.evaluation.Evaluator attribute*), 56
`func_triple_to_bpe_representation` (*dicee.evaluation.evaluator.Evaluator attribute*), 46
`func_triple_to_bpe_representation` (*dicee.Evaluator attribute*), 236
`func_triple_to_bpe_representation` (*dicee.evaluator.Evaluator attribute*), 62
`func_triple_to_bpe_representation()` (*dicee.knowledge_graph.KG method*), 70
`function()` (*dicee.models.FMult2 method*), 193
`function()` (*dicee.models.function_space.FMult2 method*), 108

G

`gamma` (*dicee.models.FMult attribute*), 191
`gamma` (*dicee.models.function_space.FMult attribute*), 107
`gate_residual` (*dicee.models.literal.LiteralEmbeddings attribute*), 111
`gated_residual_proj` (*dicee.models.literal.LiteralEmbeddings attribute*), 111
`gelu` (*dicee.models.clifford.TransformerMLP attribute*), 97
`gelu` (*dicee.models.transformers.MLP attribute*), 128
`gen_test` (*dicee.query_generator.QueryGenerator attribute*), 196
`gen_test` (*dicee.QueryGenerator attribute*), 233
`gen_valid` (*dicee.query_generator.QueryGenerator attribute*), 196
`gen_valid` (*dicee.QueryGenerator attribute*), 233
`generate()` (*dicee.KGE method*), 228
`generate()` (*dicee.knowledge_graph_embeddings.KGE method*), 70
`generate()` (*dicee.models.transformers.BytE method*), 126
`generate_queries()` (*dicee.query_generator.QueryGenerator method*), 197
`generate_queries()` (*dicee.QueryGenerator method*), 234
`get_aswa_state_dict()` (*dicee.weight_averaging.ASWA method*), 220
`get_bpe_head_and_relation_representation()` (*dicee.models.base_model.BaseKGE method*), 92
`get_bpe_head_and_relation_representation()` (*dicee.models.BaseKGE method*), 142, 148, 157, 164, 171, 185, 191
`get_bpe_token_representation()` (*dicee.abstracts.BaseInteractiveKGE method*), 14
`get_callbacks()` (*in module dicee.trainer.dice_trainer*), 214
`get_default_arguments()` (*in module dicee.analyse_experiments*), 21
`get_default_arguments()` (*in module dicee.scripts.index_serve*), 205
`get_default_arguments()` (*in module dicee.scripts.run*), 206
`get_ee_vocab()` (*in module dicee.read_preprocess_save_load_kg.util*), 202
`get_ee_vocab()` (*in module dicee.static_funcs*), 208
`get_ee_vocab()` (*in module dicee.static_preprocess_funcs*), 213
`get_embeddings()` (*dicee.models.base_model.BaseKGE method*), 92
`get_embeddings()` (*dicee.models.BaseKGE method*), 143, 148, 157, 165, 172, 186, 191
`get_embeddings()` (*dicee.models.ensemble.EnsembleKGE method*), 106
`get_embeddings()` (*dicee.models.real.Shallom method*), 122
`get_embeddings()` (*dicee.models.Shallom method*), 151
`get_entity_embeddings()` (*dicee.abstracts.BaseInteractiveKGE method*), 16
`get_entity_index()` (*dicee.abstracts.BaseInteractiveKGE method*), 15
`get_er_vocab()` (*in module dicee.read_preprocess_save_load_kg.util*), 202
`get_er_vocab()` (*in module dicee.static_funcs*), 207
`get_er_vocab()` (*in module dicee.static_preprocess_funcs*), 213
`get_eval_report()` (*dicee.abstracts.BaseInteractiveKGE method*), 14

`get_head_relation_representation()` (*dicee.models.base_model.BaseKGE method*), 91
`get_head_relation_representation()` (*dicee.models.BaseKGE method*), 142, 147, 157, 164, 171, 185, 190
`get_kronecker_triple_representation()` (*dicee.callbacks.KronE method*), 26
`get_mean_and_var()` (*dicee.weight_averaging.SWAG method*), 223
`get_num_params()` (*dicee.models.transformers.GPT method*), 130
`get_padded_bpe_triple_representation()` (*dicee.abstracts.BaseInteractiveKGE method*), 14
`get_queries()` (*dicee.query_generator.QueryGenerator method*), 197
`get_queries()` (*dicee.QueryGenerator method*), 234
`get_re_vocab()` (*in module dicee.read_preprocess_save_load_kg.util*), 202
`get_re_vocab()` (*in module dicee.static_funcs*), 207
`get_re_vocab()` (*in module dicee.static_preprocess_funcs*), 213
`get_relation_embeddings()` (*dicee.abstracts.BaseInteractiveKGE method*), 16
`get_relation_index()` (*dicee.abstracts.BaseInteractiveKGE method*), 15
`get_sentence_representation()` (*dicee.models.base_model.BaseKGE method*), 91
`get_sentence_representation()` (*dicee.models.BaseKGE method*), 142, 147, 157, 164, 171, 185, 190
`get_transductive_entity_embeddings()` (*dicee.KGE method*), 227
`get_transductive_entity_embeddings()` (*dicee.knowledge_graph_embeddings.KGE method*), 70
`get_triple_representation()` (*dicee.models.base_model.BaseKGE method*), 91
`get_triple_representation()` (*dicee.models.BaseKGE method*), 142, 147, 156, 163, 171, 185, 190
`GFMult` (*class in dicee.models*), 192
`GFMult` (*class in dicee.models.function_space*), 107
`global_rank` (*dicee.abstracts.AbstractTrainer attribute*), 13
`global_rank` (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 218
`GPT` (*class in dicee.models.transformers*), 129
`GPTConfig` (*class in dicee.models.transformers*), 129
`gpus` (*dicee.config.Namespace attribute*), 29
`gradient_accumulation_steps` (*dicee.config.Namespace attribute*), 29
`ground_queries()` (*dicee.query_generator.QueryGenerator method*), 197
`ground_queries()` (*dicee.QueryGenerator method*), 233
`gswa_n` (*dicee.weight_averaging.SWAG attribute*), 223

H

`hidden_dim` (*dicee.models.literal.LiteralEmbeddings attribute*), 111
`hidden_dropout` (*dicee.models.base_model.BaseKGE attribute*), 89
`hidden_dropout` (*dicee.models.BaseKGE attribute*), 140, 145, 154, 162, 169, 183, 188
`hidden_dropout_rate` (*dicee.config.Namespace attribute*), 30
`hidden_dropout_rate` (*dicee.models.base_model.BaseKGE attribute*), 88
`hidden_dropout_rate` (*dicee.models.BaseKGE attribute*), 139, 144, 154, 161, 168, 182, 187
`hidden_normalizer` (*dicee.models.base_model.BaseKGE attribute*), 89
`hidden_normalizer` (*dicee.models.BaseKGE attribute*), 140, 145, 154, 161, 169, 183, 188

I

`IdentityClass` (*class in dicee.models*), 143, 165, 172
`IdentityClass` (*class in dicee.models.base_model*), 92
`idx_entity_to_bpe_shaped` (*dicee.knowledge_graph.KG attribute*), 69
`index()` (*in module dicee.scripts.index_serve*), 205
`index_triple()` (*dicee.abstracts.BaseInteractiveKGE method*), 16
`init_data_loader()` (*dicee.DICE_Trainer method*), 235
`init_data_loader()` (*dicee.trainer.DICE_Trainer method*), 219
`init_data_loader()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 215
`init_dataset()` (*dicee.DICE_Trainer method*), 235
`init_dataset()` (*dicee.trainer.DICE_Trainer method*), 219
`init_dataset()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 215
`init_param` (*dicee.config.Namespace attribute*), 29
`init_params_with_sanity_checking()` (*dicee.models.base_model.BaseKGE method*), 90
`init_params_with_sanity_checking()` (*dicee.models.BaseKGE method*), 140, 146, 155, 162, 169, 183, 189
`initial_eval_setting` (*dicee.weight_averaging.ASWA attribute*), 220
`initialize_or_load_model()` (*dicee.DICE_Trainer method*), 235
`initialize_or_load_model()` (*dicee.trainer.DICE_Trainer method*), 219
`initialize_or_load_model()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 214
`initialize_trainer()` (*dicee.DICE_Trainer method*), 235
`initialize_trainer()` (*dicee.trainer.DICE_Trainer method*), 219
`initialize_trainer()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 214
`initialize_trainer()` (*in module dicee.trainer.dice_trainer*), 213
`input_dp_ent_real` (*dicee.models.base_model.BaseKGE attribute*), 89
`input_dp_ent_real` (*dicee.models.BaseKGE attribute*), 140, 145, 154, 162, 169, 183, 188
`input_dp_rel_real` (*dicee.models.base_model.BaseKGE attribute*), 89

dicee.models.BaseKGE attribute), 140, 145, 154, 162, 169, 183, 188
dicee.config.Namespace attribute), 30
dicee.models.base_model.BaseKGE attribute), 88
dicee.models.BaseKGE attribute), 139, 144, 154, 161, 168, 182, 187
InteractiveQueryDecomposition (*class in dicee.abstracts*), 17
intialize_model() (*in module dicee.static_funcs*), 209
is_continual_training (*dicee.DICE_Trainer attribute*), 234
is_continual_training (*dicee.evaluation.Evaluator attribute*), 56
is_continual_training (*dicee.evaluation.evaluator.Evaluator attribute*), 46
is_continual_training (*dicee.Evaluator attribute*), 236
is_continual_training (*dicee.evaluator.Evaluator attribute*), 62
is_continual_training (*dicee.Execute attribute*), 226
is_continual_training (*dicee.executer.Execute attribute*), 65
is_continual_training (*dicee.trainer.DICE_Trainer attribute*), 219
is_continual_training (*dicee.trainer.dice_trainer.DICE_Trainer attribute*), 214
is_global_zero (*dicee.abstracts.AbstractTrainer attribute*), 13
is_local_rank_zero() (*dicee.Execute method*), 226
is_local_rank_zero() (*dicee.executer.Execute method*), 66
is_seen() (*dicee.abstracts.BaseInteractiveKGE method*), 15
is_sparql_endpoint_alive() (*in module dicee.sanity_checkers*), 204

K

k (*dicee.models.FMult attribute*), 191
k (*dicee.models.FMult2 attribute*), 192
k (*dicee.models.function_space.FMult attribute*), 107
k (*dicee.models.function_space.FMult2 attribute*), 108
k (*dicee.models.function_space.GFMult attribute*), 107
k (*dicee.models.GFMult attribute*), 192
k_fold_cross_validation() (*dicee.DICE_Trainer method*), 235
k_fold_cross_validation() (*dicee.trainer.DICE_Trainer method*), 219
k_fold_cross_validation() (*dicee.trainer.dice_trainer.DICE_Trainer method*), 215
k_vs_all_score() (*dicee.models.clifford.Keci method*), 95
k_vs_all_score() (*dicee.models.ComplEx static method*), 160
k_vs_all_score() (*dicee.models.complex.ComplEx static method*), 104
k_vs_all_score() (*dicee.models.DistMult method*), 148
k_vs_all_score() (*dicee.models.Keci method*), 177
k_vs_all_score() (*dicee.models.octonion.OMult method*), 113
k_vs_all_score() (*dicee.models.OMult method*), 172
k_vs_all_score() (*dicee.models.QMult method*), 166
k_vs_all_score() (*dicee.models.quaternion.QMult method*), 117
k_vs_all_score() (*dicee.models.real.DistMult method*), 119
Keci (*class in dicee.models*), 174
Keci (*class in dicee.models.clifford*), 93
KeciTransformer (*class in dicee.models*), 181
KeciTransformer (*class in dicee.models.clifford*), 96
kernel_size (*dicee.config.Namespace attribute*), 30
kernel_size (*dicee.models.base_model.BaseKGE attribute*), 89
kernel_size (*dicee.models.BaseKGE attribute*), 139, 145, 154, 161, 168, 182, 188
KG (*class in dicee.knowledge_graph*), 67
kg (*dicee.callbacks.PseudoLabellingCallback attribute*), 24
kg (*dicee.read_preprocess_save_load_kg.LoadSaveToDisk attribute*), 203
kg (*dicee.read_preprocess_save_load_kg.PreprocessKG attribute*), 203
kg (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG attribute*), 198
kg (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk attribute*), 198
kg (*dicee.read_preprocess_save_load_kg.ReadFromDisk attribute*), 203
kg (*dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk attribute*), 199
KGE (*class in dicee*), 227
KGE (*class in dicee.knowledge_graph_embeddings*), 70
KGESaveCallback (*class in dicee.callbacks*), 23
knowledge_graph (*dicee.Execute attribute*), 226
knowledge_graph (*dicee.executer.Execute attribute*), 65
KronE (*class in dicee.callbacks*), 25
KvsAll (*class in dicee.dataset_classes*), 35
kvsall_score() (*dicee.models.DualE method*), 195
kvsall_score() (*dicee.models.dualE.DualE method*), 105
KvsSampleDataset (*class in dicee.dataset_classes*), 35

L

label_smoothing_rate (*dicee.config.Namespace* attribute), 29
 label_smoothing_rate (*dicee.dataset_classes.AllvsAll* attribute), 34
 label_smoothing_rate (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 38
 label_smoothing_rate (*dicee.dataset_classes.KvsAll* attribute), 35
 label_smoothing_rate (*dicee.dataset_classes.KvsSampleDataset* attribute), 36
 label_smoothing_rate (*dicee.dataset_classes.OnevsSample* attribute), 39
 label_smoothing_rate (*dicee.dataset_classes.TriplePredictionDataset* attribute), 39
 layer_norm (*dicee.models.literal.LiteralEmbeddings* attribute), 112
 LayerNorm (*class in dicee.models.transformers*), 126
 learning_rate (*dicee.models.base_model.BaseKGE* attribute), 88
 learning_rate (*dicee.models.BaseKGE* attribute), 139, 144, 154, 161, 168, 182, 187
 length (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 38
 length (*dicee.dataset_classes.TriplePredictionDataset* attribute), 39
 level (*dicee.callbacks.Perturb* attribute), 26
 LFMult (*class in dicee.models*), 193
 LFMult (*class in dicee.models.function_space*), 109
 LFMult1 (*class in dicee.models*), 193
 LFMult1 (*class in dicee.models.function_space*), 108
 linear() (*dicee.models.function_space.LFMult* method), 109
 linear() (*dicee.models.LFMult* method), 194
 list2tuple() (*dicee.query_generator.QueryGenerator* method), 197
 list2tuple() (*dicee.QueryGenerator* method), 233
 LiteralDataset (*class in dicee.dataset_classes*), 36
 LiteralEmbeddings (*class in dicee.models.literal*), 110
 lm_head (*dicee.models.clifford.KeciTransformer* attribute), 97
 lm_head (*dicee.models.KeciTransformer* attribute), 182
 lm_head (*dicee.models.transformers.BytE* attribute), 125
 lm_head (*dicee.models.transformers.GPT* attribute), 130
 ln_1 (*dicee.models.Block* attribute), 144
 ln_1 (*dicee.models.clifford.TransformerBlock* attribute), 97
 ln_1 (*dicee.models.transformers.Block* attribute), 129
 ln_2 (*dicee.models.Block* attribute), 144
 ln_2 (*dicee.models.clifford.TransformerBlock* attribute), 97
 ln_2 (*dicee.models.transformers.Block* attribute), 129
 ln_f (*dicee.models.CoKE* attribute), 153
 ln_f (*dicee.models.real.CoKE* attribute), 124
 load() (*dicee.read_preprocess_save_load_kg.LoadSaveToDisk* method), 203
 load() (*dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk* method), 199
 load_and_validate_literal_data() (*dicee.dataset_classes.LiteralDataset* static method), 37
 load_from_memmap() (*dicee.Execute* method), 226
 load_from_memmap() (*dicee.executer.Execute* method), 66
 load_json() (*in module dicee.static_funcs*), 209
 load_model() (*in module dicee.static_funcs*), 208
 load_model_ensemble() (*in module dicee.static_funcs*), 209
 load_numpy() (*in module dicee.static_funcs*), 210
 load_numpy_ndarray() (*in module dicee.read_preprocess_save_load_kg.util*), 202
 load_pickle() (*in module dicee.read_preprocess_save_load_kg.util*), 202
 load_pickle() (*in module dicee.static_funcs*), 208
 load_queries() (*dicee.query_generator.QueryGenerator* method), 197
 load_queries() (*dicee.QueryGenerator* method), 234
 load_queries_and_answers() (*dicee.query_generator.QueryGenerator* static method), 197
 load_queries_and_answers() (*dicee.QueryGenerator* static method), 234
 load_state_dict() (*dicee.models.ensemble.EnsembleKGE* method), 106
 load_term_mapping() (*in module dicee.static_funcs*), 208
 load_term_mapping() (*in module dicee.trainer.dice_trainer*), 213
 load_with_pandas() (*in module dicee.read_preprocess_save_load_kg.util*), 202
 loader_backend (*dicee.dataset_classes.LiteralDataset* attribute), 37
 LoadSaveToDisk (*class in dicee.read_preprocess_save_load_kg*), 203
 LoadSaveToDisk (*class in dicee.read_preprocess_save_load_kg.save_load_disk*), 199
 local_rank (*dicee.abstracts.AbstractTrainer* attribute), 13
 local_rank (*dicee.Execute* attribute), 226
 local_rank (*dicee.executer.Execute* attribute), 65
 local_rank (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 218
 loss (*dicee.models.base_model.BaseKGE* attribute), 89
 loss (*dicee.models.BaseKGE* attribute), 140, 145, 154, 161, 168, 182, 188
 loss_func (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 218
 loss_function (*dicee.trainer.torch_trainer.TorchTrainer* attribute), 216

- `loss_function()` (*dicее.models.base_model.BaseKGELightning method*), 85
- `loss_function()` (*dicее.models.BaseKGELightning method*), 136
- `loss_function()` (*dicее.models.transformers.BytE method*), 126
- `loss_history` (*dicее.models.base_model.BaseKGE attribute*), 89
- `loss_history` (*dicее.models.BaseKGE attribute*), 140, 145, 154, 162, 169, 183, 188
- `loss_history` (*dicее.models.pykeen_models.PykeenKGE attribute*), 115
- `loss_history` (*dicее.models.PykeenKGE attribute*), 186
- `loss_history` (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 218
- `lr` (*dicее.analyse_experiments.Experiment attribute*), 21
- `lr` (*dicее.config.Namespace attribute*), 28
- `lr_init` (*dicее.weight_averaging.SWA attribute*), 221
- `lr_init` (*dicее.weight_averaging.SWAG attribute*), 222
- `lr_init` (*dicее.weight_averaging.TWA attribute*), 224
- `lr_lambda` (*dicее.callbacks.LRScheduler attribute*), 27
- `LRScheduler` (*class in dicее.callbacks*), 27

M

- `m` (*dicее.models.function_space.LFMult attribute*), 109
- `m` (*dicее.models.LFMult attribute*), 194
- `main()` (*in module dicее.scripts.index_serve*), 205
- `main()` (*in module dicее.scripts.run*), 206
- `make_iterable_verbose()` (*in module dicее.evaluation.utils*), 52
- `make_iterable_verbose()` (*in module dicее.static_funcs_training*), 212
- `make_iterable_verbose()` (*in module dicее.trainer.torch_trainer_ddp*), 217
- `mapping_from_first_two_cols_to_third()` (*in module dicее.static_preprocess_funcs*), 213
- `margin` (*dicее.models.Pyke attribute*), 151
- `margin` (*dicее.models.real.Pyke attribute*), 122
- `margin` (*dicее.models.real.TransE attribute*), 121
- `margin` (*dicее.models.TransE attribute*), 150
- `mask_emb` (*dicее.models.CoKE attribute*), 153
- `mask_emb` (*dicее.models.real.CoKE attribute*), 124
- `max_ans_num` (*dicее.query_generator.QueryGenerator attribute*), 196
- `max_ans_num` (*dicее.QueryGenerator attribute*), 233
- `max_epochs` (*dicее.callbacks.KGESaveCallback attribute*), 23
- `max_epochs` (*dicее.callbacks.PeriodicEvalCallback attribute*), 26
- `max_epochs` (*dicее.weight_averaging.EMA attribute*), 223
- `max_epochs` (*dicее.weight_averaging.SWA attribute*), 221
- `max_epochs` (*dicее.weight_averaging.SWAG attribute*), 222
- `max_epochs` (*dicее.weight_averaging.TWA attribute*), 224
- `max_length_subword_tokens` (*dicее.knowledge_graph.KG attribute*), 69
- `max_length_subword_tokens` (*dicее.models.base_model.BaseKGE attribute*), 89
- `max_length_subword_tokens` (*dicее.models.BaseKGE attribute*), 140, 145, 154, 162, 169, 183, 188
- `max_num_models` (*dicее.weight_averaging.SWAG attribute*), 222
- `max_num_of_classes` (*dicее.dataset_classes.KvsSampleDataset attribute*), 36
- `mean` (*dicее.weight_averaging.SWAG attribute*), 222
- `mem_of_model()` (*dicее.models.base_model.BaseKGELightning method*), 84
- `mem_of_model()` (*dicее.models.BaseKGELightning method*), 135
- `mem_of_model()` (*dicее.models.ensemble.EnsembleKGE method*), 106
- `method` (*dicее.callbacks.Perturb attribute*), 26
- `MLP` (*class in dicее.models.transformers*), 127
- `mlp` (*dicее.models.Block attribute*), 144
- `mlp` (*dicее.models.clifford.TransformerBlock attribute*), 97
- `mlp` (*dicее.models.transformers.Block attribute*), 129
- `mode` (*dicее.query_generator.QueryGenerator attribute*), 196
- `mode` (*dicее.QueryGenerator attribute*), 233
- `model` (*dicее.config.Namespace attribute*), 28
- `model` (*dicее.models.pykeen_models.PykeenKGE attribute*), 115
- `model` (*dicее.models.PykeenKGE attribute*), 186
- `model` (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 218
- `model` (*dicее.trainer.torch_trainer.TorchTrainer attribute*), 216
- `model_kwargs` (*dicее.models.pykeen_models.PykeenKGE attribute*), 115
- `model_kwargs` (*dicее.models.PykeenKGE attribute*), 186
- `model_name` (*dicее.analyse_experiments.Experiment attribute*), 21
- `MODEL_REGISTRY` (*in module dicее.static_funcs*), 207
- `module`
 - `dicее`, 12
 - `dicее.__main__`, 12

- `dicee.abstracts`, 12
- `dicee.analyse_experiments`, 20
- `dicee.callbacks`, 22
- `dicee.config`, 28
- `dicee.dataset_classes`, 31
- `dicee.eval_static_funcs`, 39
- `dicee.evaluation`, 43
 - `dicee.evaluation.ensemble`, 43
 - `dicee.evaluation.evaluator`, 44
 - `dicee.evaluation.link_prediction`, 48
 - `dicee.evaluation.literal_prediction`, 51
 - `dicee.evaluation.utils`, 52
- `dicee.evaluator`, 61
- `dicee.executer`, 64
- `dicee.knowledge_graph`, 67
- `dicee.knowledge_graph_embeddings`, 70
- `dicee.models`, 75
 - `dicee.models.adopt`, 75
 - `dicee.models.base_model`, 84
 - `dicee.models.clifford`, 92
 - `dicee.models.complex`, 101
 - `dicee.models.dualE`, 104
 - `dicee.models.ensemble`, 105
 - `dicee.models.function_space`, 106
 - `dicee.models.literal`, 110
 - `dicee.models.octonion`, 112
 - `dicee.models.pykeen_models`, 115
 - `dicee.models.quaternion`, 116
 - `dicee.models.real`, 119
 - `dicee.models.static_funcs`, 124
 - `dicee.models.transformers`, 125
- `dicee.query_generator`, 196
- `dicee.read_preprocess_save_load_kg`, 197
 - `dicee.read_preprocess_save_load_kg.preprocess`, 197
 - `dicee.read_preprocess_save_load_kg.read_from_disk`, 198
 - `dicee.read_preprocess_save_load_kg.save_load_disk`, 199
 - `dicee.read_preprocess_save_load_kg.util`, 199
- `dicee.sanity_checkers`, 204
- `dicee.scripts`, 204
 - `dicee.scripts.index_serve`, 204
 - `dicee.scripts.run`, 206
- `dicee.static_funcs`, 206
 - `dicee.static_funcs_training`, 210
 - `dicee.static_preprocess_funcs`, 212
- `dicee.trainer`, 213
 - `dicee.trainer.dice_trainer`, 213
 - `dicee.trainer.model_parallelism`, 215
 - `dicee.trainer.torch_trainer`, 216
 - `dicee.trainer.torch_trainer_ddp`, 217
 - `dicee.weight_averaging`, 220
- `modules()` (*`dicee.models.ensemble.EnsembleKGE` method*), 106
- `moving_average()` (*`dicee.weight_averaging.SWA` static method*), 222
- `MultiClassClassificationDataset` (*class in `dicee.dataset_classes`*), 33
- `MultiLabelDataset` (*class in `dicee.dataset_classes`*), 33

N

- `n` (*`dicee.models.FMult2` attribute*), 192
- `n` (*`dicee.models.function_space.FMult2` attribute*), 108
- `n_embd` (*`dicee.models.clifford.TransformerSelfAttention` attribute*), 97
- `n_embd` (*`dicee.models.CoKEConfig` attribute*), 152
- `n_embd` (*`dicee.models.real.CoKEConfig` attribute*), 123
- `n_embd` (*`dicee.models.transformers.GPTConfig` attribute*), 129
- `n_embd` (*`dicee.models.transformers.SelfAttention` attribute*), 127
- `n_epochs_eval_model` (*`dicee.callbacks.PeriodicEvalCallback` attribute*), 26
- `n_epochs_eval_model` (*`dicee.config.Namespace` attribute*), 31
- `n_head` (*`dicee.models.clifford.TransformerSelfAttention` attribute*), 97
- `n_head` (*`dicee.models.CoKEConfig` attribute*), 152

`n_head` (*dicee.models.real.CoKEConfig* attribute), 123
`n_head` (*dicee.models.transformers.GPTConfig* attribute), 129
`n_head` (*dicee.models.transformers.SelfAttention* attribute), 127
`n_layer` (*dicee.models.CoKEConfig* attribute), 152
`n_layer` (*dicee.models.real.CoKEConfig* attribute), 123
`n_layer` (*dicee.models.transformers.GPTConfig* attribute), 129
`n_layers` (*dicee.models.FMult2* attribute), 192
`n_layers` (*dicee.models.function_space.FMult2* attribute), 108
`name` (*dicee.abstracts.BaseInteractiveKGE* property), 14
`name` (*dicee.models.AConEx* attribute), 159
`name` (*dicee.models.AConvO* attribute), 174
`name` (*dicee.models.AConvQ* attribute), 167
`name` (*dicee.models.CKeci* attribute), 178
`name` (*dicee.models.clifford.CKeci* attribute), 98
`name` (*dicee.models.clifford.DeCaL* attribute), 98
`name` (*dicee.models.clifford.Keci* attribute), 93
`name` (*dicee.models.clifford.KeciTransformer* attribute), 96
`name` (*dicee.models.CoKE* attribute), 153
`name` (*dicee.models.ComplEx* attribute), 160
`name` (*dicee.models.complex.AConEx* attribute), 102
`name` (*dicee.models.complex.ComplEx* attribute), 103
`name` (*dicee.models.complex.ConEx* attribute), 101
`name` (*dicee.models.ConEx* attribute), 158
`name` (*dicee.models.ConvO* attribute), 173
`name` (*dicee.models.ConvQ* attribute), 166
`name` (*dicee.models.DeCaL* attribute), 178
`name` (*dicee.models.DistMult* attribute), 148
`name` (*dicee.models.DualE* attribute), 195
`name` (*dicee.models.dualE.DualE* attribute), 104
`name` (*dicee.models.ensemble.EnsembleKGE* attribute), 105
`name` (*dicee.models.FMult* attribute), 191
`name` (*dicee.models.FMult2* attribute), 192
`name` (*dicee.models.function_space.FMult* attribute), 106
`name` (*dicee.models.function_space.FMult2* attribute), 108
`name` (*dicee.models.function_space.GFMult* attribute), 107
`name` (*dicee.models.function_space.LFMult* attribute), 109
`name` (*dicee.models.function_space.LFMult1* attribute), 108
`name` (*dicee.models.GFMult* attribute), 192
`name` (*dicee.models.Keci* attribute), 175
`name` (*dicee.models.KeciTransformer* attribute), 181
`name` (*dicee.models.LFMult* attribute), 194
`name` (*dicee.models.LFMult1* attribute), 193
`name` (*dicee.models.octonion.AConvO* attribute), 114
`name` (*dicee.models.octonion.ConvO* attribute), 113
`name` (*dicee.models.octonion.OMult* attribute), 113
`name` (*dicee.models.OMult* attribute), 172
`name` (*dicee.models.Pyke* attribute), 151
`name` (*dicee.models.pykeen_models.PykeenKGE* attribute), 115
`name` (*dicee.models.PykeenKGE* attribute), 186
`name` (*dicee.models.QMult* attribute), 165
`name` (*dicee.models.quaternion.AConvQ* attribute), 118
`name` (*dicee.models.quaternion.ConvQ* attribute), 117
`name` (*dicee.models.quaternion.QMult* attribute), 116
`name` (*dicee.models.real.CoKE* attribute), 124
`name` (*dicee.models.real.DistMult* attribute), 119
`name` (*dicee.models.real.Pyke* attribute), 122
`name` (*dicee.models.real.Shallom* attribute), 122
`name` (*dicee.models.real.TransE* attribute), 121
`name` (*dicee.models.Shallom* attribute), 151
`name` (*dicee.models.TransE* attribute), 150
`name` (*dicee.models.transformers.BytE* attribute), 125
`named_children()` (*dicee.models.ensemble.EnsembleKGE* method), 106
`Namespace` (class in *dicee.config*), 28
`neg_ratio` (*dicee.config.Namespace* attribute), 29
`neg_ratio` (*dicee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 32
`neg_ratio` (*dicee.dataset_classes.KvsSampleDataset* attribute), 36
`neg_sample_ratio` (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 38
`neg_sample_ratio` (*dicee.dataset_classes.OnevsSample* attribute), 38

neg_sample_ratio (*dicdee.dataset_classes.TriplePredictionDataset* attribute), 39
 negnorm() (*dicdee.abstracts.InteractiveQueryDecomposition* method), 18
 neural_searcher (in module *dicdee.scripts.index_serve*), 205
 NeuralSearcher (class in *dicdee.scripts.index_serve*), 205
 NodeTrainer (class in *dicdee.trainer.torch_trainer_ddp*), 218
 norm_fc1 (*dicdee.models.AConEx* attribute), 159
 norm_fc1 (*dicdee.models.AConvO* attribute), 174
 norm_fc1 (*dicdee.models.complex.AConEx* attribute), 102
 norm_fc1 (*dicdee.models.complex.ConEx* attribute), 101
 norm_fc1 (*dicdee.models.ConEx* attribute), 158
 norm_fc1 (*dicdee.models.ConvO* attribute), 173
 norm_fc1 (*dicdee.models.octonion.AConvO* attribute), 114
 norm_fc1 (*dicdee.models.octonion.ConvO* attribute), 113
 normalization (*dicdee.analyse_experiments.Experiment* attribute), 21
 normalization (*dicdee.config.Namespace* attribute), 29
 normalization_params (*dicdee.dataset_classes.LiteralDataset* attribute), 37
 normalization_type (*dicdee.dataset_classes.LiteralDataset* attribute), 37
 normalize_head_entity_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 89
 normalize_head_entity_embeddings (*dicdee.models.BaseKGE* attribute), 140, 145, 154, 161, 168, 183, 188
 normalize_relation_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 89
 normalize_relation_embeddings (*dicdee.models.BaseKGE* attribute), 140, 145, 154, 161, 169, 183, 188
 normalize_tail_entity_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 89
 normalize_tail_entity_embeddings (*dicdee.models.BaseKGE* attribute), 140, 145, 154, 161, 169, 183, 188
 normalizer_class (*dicdee.models.base_model.BaseKGE* attribute), 89
 normalizer_class (*dicdee.models.BaseKGE* attribute), 140, 145, 154, 161, 168, 183, 188
 num_bpe_entities (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 32
 num_bpe_entities (*dicdee.knowledge_graph.KG* attribute), 69
 num_core (*dicdee.config.Namespace* attribute), 29
 num_datapoints (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 32
 num_datapoints (*dicdee.dataset_classes.MultiLabelDataset* attribute), 33
 num_ent (*dicdee.models.DualE* attribute), 195
 num_ent (*dicdee.models.dualE.DualE* attribute), 105
 num_entities (*dicdee.dataset_classes.FixedNegSampleDataset* attribute), 38
 num_entities (*dicdee.dataset_classes.KvsSampleDataset* attribute), 36
 num_entities (*dicdee.dataset_classes.LiteralDataset* attribute), 37
 num_entities (*dicdee.dataset_classes.OnevsSample* attribute), 38
 num_entities (*dicdee.dataset_classes.TriplePredictionDataset* attribute), 39
 num_entities (*dicdee.evaluation.Evaluator* attribute), 55, 56
 num_entities (*dicdee.evaluation.evaluator.Evaluator* attribute), 45, 46
 num_entities (*dicdee.Evaluator* attribute), 235, 236
 num_entities (*dicdee.evaluator.Evaluator* attribute), 61, 62
 num_entities (*dicdee.knowledge_graph.KG* attribute), 68
 num_entities (*dicdee.models.base_model.BaseKGE* attribute), 88
 num_entities (*dicdee.models.BaseKGE* attribute), 139, 144, 154, 161, 168, 182, 187
 num_epochs (*dicdee.abstracts.AbstractPPECallback* attribute), 19
 num_epochs (*dicdee.analyse_experiments.Experiment* attribute), 21
 num_epochs (*dicdee.config.Namespace* attribute), 28
 num_epochs (*dicdee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 218
 num_epochs (*dicdee.weight_averaging.ASWA* attribute), 220
 num_folds_for_cv (*dicdee.config.Namespace* attribute), 29
 num_of_data_points (*dicdee.dataset_classes.MultiClassClassificationDataset* attribute), 33
 num_of_data_properties (*dicdee.models.literal.LiteralEmbeddings* attribute), 111
 num_of_epochs (*dicdee.callbacks.PseudoLabellingCallback* attribute), 24
 num_of_output_channels (*dicdee.config.Namespace* attribute), 30
 num_of_output_channels (*dicdee.models.base_model.BaseKGE* attribute), 89
 num_of_output_channels (*dicdee.models.BaseKGE* attribute), 139, 145, 154, 161, 168, 182, 188
 num_params (*dicdee.analyse_experiments.Experiment* attribute), 21
 num_relations (*dicdee.dataset_classes.FixedNegSampleDataset* attribute), 38
 num_relations (*dicdee.dataset_classes.OnevsSample* attribute), 38
 num_relations (*dicdee.dataset_classes.TriplePredictionDataset* attribute), 39
 num_relations (*dicdee.evaluation.Evaluator* attribute), 55, 56
 num_relations (*dicdee.evaluation.evaluator.Evaluator* attribute), 45, 46
 num_relations (*dicdee.Evaluator* attribute), 236
 num_relations (*dicdee.evaluator.Evaluator* attribute), 61, 62
 num_relations (*dicdee.knowledge_graph.KG* attribute), 68
 num_relations (*dicdee.models.base_model.BaseKGE* attribute), 88
 num_relations (*dicdee.models.BaseKGE* attribute), 139, 144, 154, 161, 168, 182, 187
 num_sample (*dicdee.models.FMult* attribute), 191

num_sample (*dicee.models.function_space.FMult attribute*), 107
 num_sample (*dicee.models.function_space.GFMult attribute*), 107
 num_sample (*dicee.models.GFMult attribute*), 192
 num_samples (*dicee.weight_averaging.TWA attribute*), 224
 num_tokens (*dicee.knowledge_graph.KG attribute*), 69
 num_tokens (*dicee.models.base_model.BaseKGE attribute*), 88
 num_tokens (*dicee.models.BaseKGE attribute*), 139, 144, 154, 161, 168, 182, 187
 numpy_data_type_changer () (*in module dicee.static_funcs*), 209

O

octonion_mul () (*in module dicee.models*), 172
 octonion_mul () (*in module dicee.models.octonion*), 112
 octonion_mul_norm () (*in module dicee.models*), 172
 octonion_mul_norm () (*in module dicee.models.octonion*), 112
 octonion_normalizer () (*dicee.models.AConvO static method*), 174
 octonion_normalizer () (*dicee.models.ConvO static method*), 173
 octonion_normalizer () (*dicee.models.octonion.AConvO static method*), 114
 octonion_normalizer () (*dicee.models.octonion.ConvO static method*), 113
 octonion_normalizer () (*dicee.models.octonion.OMult static method*), 113
 octonion_normalizer () (*dicee.models.OMult static method*), 172
 OMult (*class in dicee.models*), 172
 OMult (*class in dicee.models.octonion*), 112
 on_epoch_end () (*dicee.callbacks.KGESaveCallback method*), 24
 on_epoch_end () (*dicee.callbacks.PseudoLabellingCallback method*), 24
 on_fit_end () (*dicee.abstracts.AbstractCallback method*), 19
 on_fit_end () (*dicee.abstracts.AbstractPPECallback method*), 19
 on_fit_end () (*dicee.abstracts.AbstractTrainer method*), 13
 on_fit_end () (*dicee.callbacks.AccumulateEpochLossCallback method*), 22
 on_fit_end () (*dicee.callbacks.Eval method*), 25
 on_fit_end () (*dicee.callbacks.KGESaveCallback method*), 24
 on_fit_end () (*dicee.callbacks.LRScheduler method*), 27
 on_fit_end () (*dicee.callbacks.PeriodicEvalCallback method*), 27
 on_fit_end () (*dicee.callbacks.PrintCallback method*), 23
 on_fit_end () (*dicee.weight_averaging.ASWA method*), 220
 on_fit_end () (*dicee.weight_averaging.EMA method*), 224
 on_fit_end () (*dicee.weight_averaging.SWA method*), 222
 on_fit_end () (*dicee.weight_averaging.SWAG method*), 223
 on_fit_end () (*dicee.weight_averaging.TWA method*), 225
 on_fit_start () (*dicee.abstracts.AbstractCallback method*), 18
 on_fit_start () (*dicee.abstracts.AbstractPPECallback method*), 19
 on_fit_start () (*dicee.abstracts.AbstractTrainer method*), 13
 on_fit_start () (*dicee.callbacks.Eval method*), 25
 on_fit_start () (*dicee.callbacks.KGESaveCallback method*), 23
 on_fit_start () (*dicee.callbacks.KronE method*), 26
 on_fit_start () (*dicee.callbacks.PrintCallback method*), 23
 on_init_end () (*dicee.abstracts.AbstractCallback method*), 18
 on_init_start () (*dicee.abstracts.AbstractCallback method*), 18
 on_train_batch_end () (*dicee.abstracts.AbstractCallback method*), 19
 on_train_batch_end () (*dicee.abstracts.AbstractTrainer method*), 13
 on_train_batch_end () (*dicee.callbacks.Eval method*), 25
 on_train_batch_end () (*dicee.callbacks.KGESaveCallback method*), 23
 on_train_batch_end () (*dicee.callbacks.LRScheduler method*), 27
 on_train_batch_end () (*dicee.callbacks.PrintCallback method*), 23
 on_train_batch_start () (*dicee.callbacks.Perturb method*), 26
 on_train_epoch_end () (*dicee.abstracts.AbstractCallback method*), 18
 on_train_epoch_end () (*dicee.abstracts.AbstractTrainer method*), 13
 on_train_epoch_end () (*dicee.callbacks.Eval method*), 25
 on_train_epoch_end () (*dicee.callbacks.KGESaveCallback method*), 24
 on_train_epoch_end () (*dicee.callbacks.PeriodicEvalCallback method*), 27
 on_train_epoch_end () (*dicee.callbacks.PrintCallback method*), 23
 on_train_epoch_end () (*dicee.models.base_model.BaseKGELightning method*), 85
 on_train_epoch_end () (*dicee.models.BaseKGELightning method*), 136
 on_train_epoch_end () (*dicee.weight_averaging.ASWA method*), 221
 on_train_epoch_end () (*dicee.weight_averaging.EMA method*), 224
 on_train_epoch_end () (*dicee.weight_averaging.SWA method*), 222
 on_train_epoch_end () (*dicee.weight_averaging.SWAG method*), 223
 on_train_epoch_end () (*dicee.weight_averaging.TWA method*), 225

on_train_epoch_start() (*dicee.abstracts.AbstractTrainer method*), 13
 on_train_epoch_start() (*dicee.weight_averaging.EMA method*), 223
 on_train_epoch_start() (*dicee.weight_averaging.SWA method*), 222
 on_train_epoch_start() (*dicee.weight_averaging.SWAG method*), 223
 on_train_epoch_start() (*dicee.weight_averaging.TWA method*), 225
 on_train_start() (*dicee.callbacks.LRScheduler method*), 27
 OnevsAllDataset (*class in dicee.dataset_classes*), 36
 OnevsSample (*class in dicee.dataset_classes*), 38
 optim (*dicee.config.Namespace attribute*), 28
 optimizer (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 218
 optimizer (*dicee.trainer.torch_trainer.TorchTrainer attribute*), 216
 optimizer_name (*dicee.models.base_model.BaseKGE attribute*), 88
 optimizer_name (*dicee.models.BaseKGE attribute*), 139, 144, 154, 161, 168, 182, 188
 ordered_bpe_entities (*dicee.dataset_classes.BPE_NegativeSamplingDataset attribute*), 32
 ordered_bpe_entities (*dicee.knowledge_graph.KG attribute*), 69
 ordered_shaped_bpe_tokens (*dicee.knowledge_graph.KG attribute*), 68

P

p (*dicee.config.Namespace attribute*), 30
 p (*dicee.models.clifford.DeCaL attribute*), 98
 p (*dicee.models.clifford.Keci attribute*), 93
 p (*dicee.models.DeCaL attribute*), 178
 p (*dicee.models.Keci attribute*), 175
 P (*dicee.weight_averaging.TWA attribute*), 224
 padding (*dicee.knowledge_graph.KG attribute*), 69
 pandas_dataframe_indexer() (*in module dicee.read_preprocess_save_load_kg.util*), 201
 param_init (*dicee.models.base_model.BaseKGE attribute*), 89
 param_init (*dicee.models.BaseKGE attribute*), 140, 145, 154, 161, 169, 183, 188
 parameters() (*dicee.abstracts.BaseInteractiveKGE method*), 17
 parameters() (*dicee.models.ensemble.EnsembleKGE method*), 106
 path (*dicee.abstracts.AbstractPPECallback attribute*), 19
 path (*dicee.callbacks.AccumulateEpochLossCallback attribute*), 22
 path (*dicee.callbacks.Eval attribute*), 25
 path (*dicee.callbacks.KGESaveCallback attribute*), 23
 path (*dicee.weight_averaging.ASWA attribute*), 220
 path_dataset_folder (*dicee.analyse_experiments.Experiment attribute*), 21
 path_for_deserialization (*dicee.knowledge_graph.KG attribute*), 68
 path_for_serialization (*dicee.knowledge_graph.KG attribute*), 68
 path_single_kg (*dicee.config.Namespace attribute*), 28
 path_single_kg (*dicee.knowledge_graph.KG attribute*), 68
 path_to_store_single_run (*dicee.config.Namespace attribute*), 28
 PeriodicEvalCallback (*class in dicee.callbacks*), 26
 Perturb (*class in dicee.callbacks*), 26
 pl_trainer_kwargs (*dicee.config.Namespace attribute*), 30
 polars_dataframe_indexer() (*in module dicee.read_preprocess_save_load_kg.util*), 200
 poly_NN() (*dicee.models.function_space.LFMult method*), 109
 poly_NN() (*dicee.models.LFMult method*), 194
 polynomial() (*dicee.models.function_space.LFMult method*), 110
 polynomial() (*dicee.models.LFMult method*), 195
 pop() (*dicee.models.function_space.LFMult method*), 110
 pop() (*dicee.models.LFMult method*), 195
 pos_emb (*dicee.models.CoKE attribute*), 153
 pos_emb (*dicee.models.real.CoKE attribute*), 124
 pq (*dicee.analyse_experiments.Experiment attribute*), 21
 predict() (*dicee.KGE method*), 229
 predict() (*dicee.knowledge_graph_embeddings.KGE method*), 71
 predict_data_loader() (*dicee.models.base_model.BaseKGELighting method*), 87
 predict_data_loader() (*dicee.models.BaseKGELighting method*), 137
 predict_literals() (*dicee.KGE method*), 232
 predict_literals() (*dicee.knowledge_graph_embeddings.KGE method*), 75
 predict_missing_head_entity() (*dicee.KGE method*), 228
 predict_missing_head_entity() (*dicee.knowledge_graph_embeddings.KGE method*), 70
 predict_missing_relations() (*dicee.KGE method*), 228
 predict_missing_relations() (*dicee.knowledge_graph_embeddings.KGE method*), 71
 predict_missing_tail_entity() (*dicee.KGE method*), 228
 predict_missing_tail_entity() (*dicee.knowledge_graph_embeddings.KGE method*), 71
 predict_topk() (*dicee.KGE method*), 229

predict_topk() (*dicee.knowledge_graph_embeddings.KGE method*), 72
 preprocess_with_byte_pair_encoding() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 203
 preprocess_with_byte_pair_encoding() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 198
 preprocess_with_byte_pair_encoding_with_padding() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 203
 preprocess_with_byte_pair_encoding_with_padding() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 198
 preprocess_with_pandas() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 203
 preprocess_with_pandas() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 198
 preprocess_with_polars() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 203
 preprocess_with_polars() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 198
 preprocesses_input_args() (*in module dicee.static_preprocess_funcs*), 212
 PreprocessKG (*class in dicee.read_preprocess_save_load_kg*), 203
 PreprocessKG (*class in dicee.read_preprocess_save_load_kg.preprocess*), 198
 PrintCallback (*class in dicee.callbacks*), 22
 process (*dicee.trainer.torch_trainer.TorchTrainer attribute*), 216
 PseudoLabellingCallback (*class in dicee.callbacks*), 24
 Pyke (*class in dicee.models*), 151
 Pyke (*class in dicee.models.real*), 122
 pykeen_model_kwargs (*dicee.config.Namespace attribute*), 30
 PykeenKGE (*class in dicee.models*), 186
 PykeenKGE (*class in dicee.models.pykeen_models*), 115

Q

q (*dicee.config.Namespace attribute*), 30
 q (*dicee.models.clifford.DeCaL attribute*), 98
 q (*dicee.models.clifford.Keci attribute*), 93
 q (*dicee.models.DeCaL attribute*), 178
 q (*dicee.models.Keci attribute*), 175
 qdrant_client (*dicee.scripts.index_serve.NeuralSearcher attribute*), 205
 QMult (*class in dicee.models*), 165
 QMult (*class in dicee.models.quaternion*), 116
 quaternion_mul() (*in module dicee.models*), 165
 quaternion_mul() (*in module dicee.models.static_funcs*), 125
 quaternion_mul_with_unit_norm() (*in module dicee.models*), 165
 quaternion_mul_with_unit_norm() (*in module dicee.models.quaternion*), 116
 quaternion_multiplication_followed_by_inner_product() (*dicee.models.QMult method*), 165
 quaternion_multiplication_followed_by_inner_product() (*dicee.models.quaternion.QMult method*), 116
 quaternion_normalizer() (*dicee.models.QMult static method*), 166
 quaternion_normalizer() (*dicee.models.quaternion.QMult static method*), 117
 queries (*dicee.scripts.index_serve.StringListRequest attribute*), 205
 query_name_to_struct (*dicee.query_generator.QueryGenerator attribute*), 196
 query_name_to_struct (*dicee.QueryGenerator attribute*), 233
 QueryGenerator (*class in dicee*), 233
 QueryGenerator (*class in dicee.query_generator*), 196

R

r (*dicee.models.clifford.DeCaL attribute*), 98
 r (*dicee.models.clifford.Keci attribute*), 93
 r (*dicee.models.DeCaL attribute*), 178
 r (*dicee.models.Keci attribute*), 175
 random_seed (*dicee.config.Namespace attribute*), 29
 range_constraints_per_rel (*dicee.evaluation.Evaluator attribute*), 56
 range_constraints_per_rel (*dicee.evaluation.evaluator.Evaluator attribute*), 46
 range_constraints_per_rel (*dicee.Evaluator attribute*), 236
 range_constraints_per_rel (*dicee.evaluator.Evaluator attribute*), 62
 rank (*dicee.Execute attribute*), 225, 226
 rank (*dicee.executer.Execute attribute*), 65
 ratio (*dicee.callbacks.Perturb attribute*), 26
 raw_test_set (*dicee.knowledge_graph.KG attribute*), 69
 raw_train_set (*dicee.knowledge_graph.KG attribute*), 69
 raw_valid_set (*dicee.knowledge_graph.KG attribute*), 69
 re (*dicee.models.clifford.DeCaL attribute*), 98
 re (*dicee.models.DeCaL attribute*), 179
 re_vocab (*dicee.evaluation.Evaluator attribute*), 55
 re_vocab (*dicee.evaluation.evaluator.Evaluator attribute*), 45
 re_vocab (*dicee.Evaluator attribute*), 235, 236
 re_vocab (*dicee.evaluator.Evaluator attribute*), 61, 62
 read_from_disk() (*in module dicee.read_preprocess_save_load_kg.util*), 201

`read_from_triple_store_with_pandas()` (in module `dicee.read_preprocess_save_load_kg.util`), 202
`read_from_triple_store_with_polars()` (in module `dicee.read_preprocess_save_load_kg.util`), 202
`read_only_few` (`dicee.config.Namespace` attribute), 30
`read_only_few` (`dicee.knowledge_graph.KG` attribute), 68
`read_or_load_kg()` (in module `dicee.static_funcs`), 209
`read_with_pandas()` (in module `dicee.read_preprocess_save_load_kg.util`), 201
`read_with_polars()` (in module `dicee.read_preprocess_save_load_kg.util`), 201
`ReadFromDisk` (class in `dicee.read_preprocess_save_load_kg`), 203
`ReadFromDisk` (class in `dicee.read_preprocess_save_load_kg.read_from_disk`), 198
`reducer` (`dicee.scripts.index_serve.StringListRequest` attribute), 205
`reg_lambda` (`dicee.weight_averaging.TWA` attribute), 224
`rel2id` (`dicee.query_generator.QueryGenerator` attribute), 196
`rel2id` (`dicee.QueryGenerator` attribute), 233
`relation_embeddings` (`dicee.models.AConvQ` attribute), 167
`relation_embeddings` (`dicee.models.clifford.DeCaL` attribute), 98
`relation_embeddings` (`dicee.models.ConvQ` attribute), 166
`relation_embeddings` (`dicee.models.DeCaL` attribute), 178
`relation_embeddings` (`dicee.models.DualE` attribute), 195
`relation_embeddings` (`dicee.models.dualE.DualE` attribute), 104
`relation_embeddings` (`dicee.models.FMult` attribute), 191
`relation_embeddings` (`dicee.models.FMult2` attribute), 193
`relation_embeddings` (`dicee.models.function_space.FMult` attribute), 107
`relation_embeddings` (`dicee.models.function_space.FMult2` attribute), 108
`relation_embeddings` (`dicee.models.function_space.GFMult` attribute), 107
`relation_embeddings` (`dicee.models.function_space.LFMult` attribute), 109
`relation_embeddings` (`dicee.models.function_space.LFMult1` attribute), 108
`relation_embeddings` (`dicee.models.GFMult` attribute), 192
`relation_embeddings` (`dicee.models.LFMult` attribute), 194
`relation_embeddings` (`dicee.models.LFMult1` attribute), 193
`relation_embeddings` (`dicee.models.pykeen_models.PykeenKGE` attribute), 115
`relation_embeddings` (`dicee.models.PykeenKGE` attribute), 186
`relation_embeddings` (`dicee.models.quaternion.AConvQ` attribute), 118
`relation_embeddings` (`dicee.models.quaternion.ConvQ` attribute), 117
`relation_to_idx` (`dicee.knowledge_graph.KG` attribute), 68
`relations_str` (`dicee.knowledge_graph.KG` property), 69
`reload_dataset()` (in module `dicee.dataset_classes`), 34
`report` (`dicee.DICE_Trainer` attribute), 234
`report` (`dicee.evaluation.Evaluator` attribute), 55, 56
`report` (`dicee.evaluation.evaluator.Evaluator` attribute), 45, 46
`report` (`dicee.Evaluator` attribute), 236
`report` (`dicee.evaluator.Evaluator` attribute), 62
`report` (`dicee.Execute` attribute), 226
`report` (`dicee.executer.Execute` attribute), 65, 66
`report` (`dicee.trainer.DICE_Trainer` attribute), 219
`report` (`dicee.trainer.dice_trainer.DICE_Trainer` attribute), 214
`reports` (`dicee.callbacks.Eval` attribute), 25
`reports` (`dicee.callbacks.PeriodicEvalCallback` attribute), 26
`requires_grad_for_interactions` (`dicee.models.CKeci` attribute), 178
`requires_grad_for_interactions` (`dicee.models.clifford.CKeci` attribute), 98
`requires_grad_for_interactions` (`dicee.models.clifford.Keci` attribute), 93
`requires_grad_for_interactions` (`dicee.models.Keci` attribute), 175
`resid_dropout` (`dicee.models.clifford.TransformerSelfAttention` attribute), 97
`resid_dropout` (`dicee.models.transformers.SelfAttention` attribute), 127
`residual_convolution()` (`dicee.models.AConEx` method), 159
`residual_convolution()` (`dicee.models.AConvO` method), 174
`residual_convolution()` (`dicee.models.AConvQ` method), 167
`residual_convolution()` (`dicee.models.complex.AConEx` method), 102
`residual_convolution()` (`dicee.models.complex.ConEx` method), 101
`residual_convolution()` (`dicee.models.ConEx` method), 158
`residual_convolution()` (`dicee.models.ConvO` method), 173
`residual_convolution()` (`dicee.models.ConvQ` method), 167
`residual_convolution()` (`dicee.models.octonion.AConvO` method), 114
`residual_convolution()` (`dicee.models.octonion.ConvO` method), 113
`residual_convolution()` (`dicee.models.quaternion.AConvQ` method), 118
`residual_convolution()` (`dicee.models.quaternion.ConvQ` method), 118
`retrieve_embedding()` (`dicee.scripts.index_serve.NeuralSearcher` method), 205
`retrieve_embeddings()` (in module `dicee.scripts.index_serve`), 205
`return_multi_hop_query_results()` (`dicee.KGE` method), 230

`return_multi_hop_query_results()` (*dicee.knowledge_graph_embeddings.KGE method*), 73
`root()` (in module *dicee.scripts.index_serve*), 205
`roots` (*dicee.models.FMult attribute*), 191
`roots` (*dicee.models.function_space.FMult attribute*), 107
`roots` (*dicee.models.function_space.GFMult attribute*), 107
`roots` (*dicee.models.GFMult attribute*), 192
`runtime` (*dicee.analyse_experiments.Experiment attribute*), 21

S

`sample()` (*dicee.weight_averaging.SWAG method*), 223
`sample_counter` (*dicee.abstracts.AbstractPPECallback attribute*), 19
`sample_entity()` (*dicee.abstracts.BaseInteractiveKGE method*), 14
`sample_relation()` (*dicee.abstracts.BaseInteractiveKGE method*), 15
`sample_triples_ratio` (*dicee.config.Namespace attribute*), 30
`sample_triples_ratio` (*dicee.knowledge_graph.KG attribute*), 68
`sample_weights()` (*dicee.weight_averaging.TWA method*), 224
`sampling_ratio` (*dicee.dataset_classes.LiteralDataset attribute*), 37
`sanity_check_callback_args()` (in module *dicee.sanity_checkers*), 204
`sanity_checking_with_arguments()` (in module *dicee.sanity_checkers*), 204
`save()` (*dicee.abstracts.BaseInteractiveKGE method*), 15
`save()` (*dicee.read_preprocess_save_load_kg.LoadSaveToDisk method*), 203
`save()` (*dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk method*), 199
`save_checkpoint()` (*dicee.abstracts.AbstractTrainer static method*), 13
`save_checkpoint_model()` (in module *dicee.static_funcs*), 209
`save_embeddings()` (in module *dicee.static_funcs*), 210
`save_embeddings_as_csv` (*dicee.config.Namespace attribute*), 28
`save_every_n_epochs` (*dicee.config.Namespace attribute*), 31
`save_experiment()` (*dicee.analyse_experiments.Experiment method*), 21
`save_model_at_every_epoch` (*dicee.config.Namespace attribute*), 29
`save_model_every_n_epoch` (*dicee.callbacks.PeriodicEvalCallback attribute*), 26
`save_numpy_ndarray()` (in module *dicee.read_preprocess_save_load_kg.util*), 202
`save_numpy_ndarray()` (in module *dicee.static_funcs*), 209
`save_pickle()` (in module *dicee.read_preprocess_save_load_kg.util*), 202
`save_pickle()` (in module *dicee.static_funcs*), 208
`save_queries()` (*dicee.query_generator.QueryGenerator method*), 197
`save_queries()` (*dicee.QueryGenerator method*), 234
`save_queries_and_answers()` (*dicee.query_generator.QueryGenerator static method*), 197
`save_queries_and_answers()` (*dicee.QueryGenerator static method*), 234
`save_trained_model()` (*dicee.Execute method*), 227
`save_trained_model()` (*dicee.executer.Execute method*), 66
`scalar_batch_NN()` (*dicee.models.function_space.LFMult method*), 109
`scalar_batch_NN()` (*dicee.models.LFMult method*), 194
`scaler` (*dicee.callbacks.Perturb attribute*), 26
`scaler` (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 218
`scheduler` (*dicee.callbacks.LRScheduler attribute*), 27
`score()` (*dicee.models.clifford.Keci method*), 96
`score()` (*dicee.models.CoKE method*), 153
`score()` (*dicee.models.ComplEx static method*), 160
`score()` (*dicee.models.complex.ComplEx static method*), 104
`score()` (*dicee.models.DistMult method*), 149
`score()` (*dicee.models.Keci method*), 178
`score()` (*dicee.models.octonion.OMult method*), 113
`score()` (*dicee.models.OMult method*), 172
`score()` (*dicee.models.QMult method*), 166
`score()` (*dicee.models.quaternion.QMult method*), 117
`score()` (*dicee.models.real.CoKE method*), 124
`score()` (*dicee.models.real.DistMult method*), 120
`score()` (*dicee.models.real.TransE method*), 121
`score()` (*dicee.models.TransE method*), 150
`score_func` (*dicee.models.FMult2 attribute*), 192
`score_func` (*dicee.models.function_space.FMult2 attribute*), 108
`scoring_technique` (*dicee.analyse_experiments.Experiment attribute*), 21
`scoring_technique` (*dicee.config.Namespace attribute*), 29
`search()` (*dicee.scripts.index_serve.NeuralSearcher method*), 205
`search_embeddings()` (in module *dicee.scripts.index_serve*), 205
`search_embeddings_batch()` (in module *dicee.scripts.index_serve*), 205
`seed` (*dicee.dataset_classes.FixedNegSampleDataset attribute*), 38

seed (*dicee.dataset_classes.TriplePredictionDataset* attribute), 39
 seed (*dicee.query_generator.QueryGenerator* attribute), 196
 seed (*dicee.QueryGenerator* attribute), 233
 select_model() (in module *dicee.static_funcs*), 208
 selected_optimizer (*dicee.models.base_model.BaseKGE* attribute), 89
 selected_optimizer (*dicee.models.BaseKGE* attribute), 140, 145, 154, 161, 168, 183, 188
 SelfAttention (class in *dicee.models.transformers*), 126
 separator (*dicee.config.Namespace* attribute), 29
 separator (*dicee.knowledge_graph.KG* attribute), 68
 seq_len (*dicee.models.clifford.KeciTransformer* attribute), 96
 seq_len (*dicee.models.KeciTransformer* attribute), 181
 sequential_vocabulary_construction() (*dicee.read_preprocess_save_load_kg.PreprocessKG* method), 203
 sequential_vocabulary_construction() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 198
 serve() (in module *dicee.scripts.index_serve*), 205
 set_global_seed() (*dicee.query_generator.QueryGenerator* method), 197
 set_global_seed() (*dicee.QueryGenerator* method), 233
 set_model_eval_mode() (*dicee.abstracts.BaseInteractiveKGE* method), 14
 set_model_train_mode() (*dicee.abstracts.BaseInteractiveKGE* method), 14
 setup_distributed_training() (in module *dicee.static_funcs*), 208
 setup_executor() (*dicee.Execute* method), 226
 setup_executor() (*dicee.executor.Execute* method), 66
 Shallom (class in *dicee.models*), 150
 Shallom (class in *dicee.models.real*), 121
 shallom (*dicee.models.real.Shallom* attribute), 122
 shallom (*dicee.models.Shallom* attribute), 151
 single_hop_query_answering() (*dicee.KGE* method), 230
 single_hop_query_answering() (*dicee.knowledge_graph_embeddings.KGE* method), 73
 snapshot_dir (*dicee.callbacks.LRScheduler* attribute), 27
 snapshot_loss (*dicee.callbacks.LRScheduler* attribute), 27
 sparql_endpoint (*dicee.config.Namespace* attribute), 28
 sparql_endpoint (*dicee.knowledge_graph.KG* attribute), 68
 sq_mean (*dicee.weight_averaging.SWAG* attribute), 222
 start() (*dicee.DICE_Trainer* method), 235
 start() (*dicee.Execute* method), 227
 start() (*dicee.executor.Execute* method), 66
 start() (*dicee.read_preprocess_save_load_kg.PreprocessKG* method), 203
 start() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 198
 start() (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk* method), 198
 start() (*dicee.read_preprocess_save_load_kg.ReadFromDisk* method), 203
 start() (*dicee.trainer.DICE_Trainer* method), 219
 start() (*dicee.trainer.dice_trainer.DICE_Trainer* method), 215
 start_time (*dicee.callbacks.PrintCallback* attribute), 23
 start_time (*dicee.Execute* attribute), 226
 start_time (*dicee.executor.Execute* attribute), 66
 state_dict() (*dicee.models.ensemble.EnsembleKGE* method), 106
 step() (*dicee.models.ADOPT* method), 134
 step() (*dicee.models.adapt.ADOPT* method), 79
 step() (*dicee.models.ensemble.EnsembleKGE* method), 106
 step_count (*dicee.callbacks.LRScheduler* attribute), 27
 storage_path (*dicee.config.Namespace* attribute), 28
 storage_path (*dicee.DICE_Trainer* attribute), 234
 storage_path (*dicee.trainer.DICE_Trainer* attribute), 219
 storage_path (*dicee.trainer.dice_trainer.DICE_Trainer* attribute), 214
 store() (in module *dicee.static_funcs*), 209
 store_ensemble() (*dicee.abstracts.AbstractPPECallback* method), 19
 strategy (*dicee.abstracts.AbstractTrainer* attribute), 13
 StringListRequest (class in *dicee.scripts.index_serve*), 205
 SWA (class in *dicee.weight_averaging*), 221
 swa (*dicee.config.Namespace* attribute), 30
 swa_c_epochs (*dicee.config.Namespace* attribute), 31
 swa_c_epochs (*dicee.weight_averaging.SWA* attribute), 221
 swa_c_epochs (*dicee.weight_averaging.SWAG* attribute), 222
 swa_lr (*dicee.weight_averaging.SWA* attribute), 221
 swa_lr (*dicee.weight_averaging.SWAG* attribute), 222
 swa_model (*dicee.weight_averaging.SWA* attribute), 221
 swa_n (*dicee.weight_averaging.SWA* attribute), 221
 swa_start_epoch (*dicee.config.Namespace* attribute), 31
 swa_start_epoch (*dicee.weight_averaging.SWA* attribute), 221

swa_start_epoch (*dicce.weight_averaging.SWAG attribute*), 222
 SWAG (*class in dicce.weight_averaging*), 222
 swag (*dicce.config.Namespace attribute*), 30

T

T () (*dicce.models.DualE method*), 196
 T () (*dicce.models.dualE.DualE method*), 105
 t_conorm () (*dicce.abstracts.InteractiveQueryDecomposition method*), 17
 t_norm () (*dicce.abstracts.InteractiveQueryDecomposition method*), 17
 target_dim (*dicce.dataset_classes.AllvsAll attribute*), 34
 target_dim (*dicce.dataset_classes.MultiLabelDataset attribute*), 33
 target_dim (*dicce.dataset_classes.OnevsAllDataset attribute*), 36
 target_dim (*dicce.knowledge_graph.KG attribute*), 69
 temperature (*dicce.models.transformers.BytE attribute*), 125
 tensor_t_norm () (*dicce.abstracts.InteractiveQueryDecomposition method*), 17
 TensorParallel (*class in dicce.trainer.model_parallelism*), 216
 test_data_loader () (*dicce.models.base_model.BaseKGELightning method*), 85
 test_data_loader () (*dicce.models.BaseKGELightning method*), 136
 test_epoch_end () (*dicce.models.base_model.BaseKGELightning method*), 85
 test_epoch_end () (*dicce.models.BaseKGELightning method*), 136
 test_h1 (*dicce.analyse_experiments.Experiment attribute*), 21
 test_h3 (*dicce.analyse_experiments.Experiment attribute*), 21
 test_h10 (*dicce.analyse_experiments.Experiment attribute*), 21
 test_mrr (*dicce.analyse_experiments.Experiment attribute*), 21
 test_path (*dicce.query_generator.QueryGenerator attribute*), 196
 test_path (*dicce.QueryGenerator attribute*), 233
 test_set (*dicce.knowledge_graph.KG attribute*), 68, 69
 timeit () (*in module dicce.read_preprocess_save_load_kg.util*), 201
 timeit () (*in module dicce.static_funcs*), 208
 timeit () (*in module dicce.static_preprocess_funcs*), 212
 to () (*dicce.KGE method*), 227
 to () (*dicce.knowledge_graph_embeddings.KGE method*), 70
 to () (*dicce.models.ensemble.EnsembleKGE method*), 106
 to_df () (*dicce.analyse_experiments.Experiment method*), 21
 topk (*dicce.models.transformers.BytE attribute*), 125
 topk (*dicce.scripts.index_serve.NeuralSearcher attribute*), 205
 torch_ordered_shaped_bpe_entities (*dicce.dataset_classes.MultiLabelDataset attribute*), 33
 TorchDDPTrainer (*class in dicce.trainer.torch_trainer_ddp*), 217
 TorchTrainer (*class in dicce.trainer.torch_trainer*), 216
 total_epochs (*dicce.callbacks.LRScheduler attribute*), 27
 total_steps (*dicce.callbacks.LRScheduler attribute*), 27
 train () (*dicce.abstracts.BaseInteractiveTrainKGE method*), 20
 train () (*dicce.trainer.torch_trainer_ddp.NodeTrainer method*), 218
 train_data (*dicce.dataset_classes.AllvsAll attribute*), 34
 train_data (*dicce.dataset_classes.KvsAll attribute*), 35
 train_data (*dicce.dataset_classes.KvsSampleDataset attribute*), 36
 train_data (*dicce.dataset_classes.MultiClassClassificationDataset attribute*), 33
 train_data (*dicce.dataset_classes.OnevsAllDataset attribute*), 36
 train_data (*dicce.dataset_classes.OnevsSample attribute*), 38
 train_data_loader () (*dicce.models.base_model.BaseKGELightning method*), 87
 train_data_loader () (*dicce.models.BaseKGELightning method*), 138
 train_data_loaders (*dicce.trainer.torch_trainer.TorchTrainer attribute*), 216
 train_dataset_loader (*dicce.trainer.torch_trainer_ddp.NodeTrainer attribute*), 218
 train_file_path (*dicce.dataset_classes.LiteralDataset attribute*), 37
 train_h1 (*dicce.analyse_experiments.Experiment attribute*), 21
 train_h3 (*dicce.analyse_experiments.Experiment attribute*), 21
 train_h10 (*dicce.analyse_experiments.Experiment attribute*), 21
 train_indices_target (*dicce.dataset_classes.MultiLabelDataset attribute*), 33
 train_k_vs_all () (*dicce.abstracts.BaseInteractiveTrainKGE method*), 20
 train_literals () (*dicce.abstracts.BaseInteractiveTrainKGE method*), 20
 train_mode (*dicce.models.ensemble.EnsembleKGE attribute*), 106
 train_mrr (*dicce.analyse_experiments.Experiment attribute*), 21
 train_path (*dicce.query_generator.QueryGenerator attribute*), 196
 train_path (*dicce.QueryGenerator attribute*), 233
 train_set (*dicce.dataset_classes.BPE_NegativeSamplingDataset attribute*), 32
 train_set (*dicce.dataset_classes.MultiLabelDataset attribute*), 33
 train_set (*dicce.dataset_classes.TriplePredictionDataset attribute*), 39

train_set (*dicee.knowledge_graph.KG attribute*), 68, 69
 train_set_target (*dicee.knowledge_graph.KG attribute*), 69
 train_target (*dicee.dataset_classes.AllvsAll attribute*), 34
 train_target (*dicee.dataset_classes.KvsAll attribute*), 35
 train_target (*dicee.dataset_classes.KvsSampleDataset attribute*), 36
 train_target_indices (*dicee.knowledge_graph.KG attribute*), 69
 train_triples (*dicee.dataset_classes.FixedNegSampleDataset attribute*), 38
 train_triples() (*dicee.abstracts.BaseInteractiveTrainKGE method*), 19
 trained_model (*dicee.Execute attribute*), 226
 trained_model (*dicee.executer.Execute attribute*), 65
 trainer (*dicee.config.Namespace attribute*), 29
 trainer (*dicee.DICE_Trainer attribute*), 234
 trainer (*dicee.Execute attribute*), 226
 trainer (*dicee.executer.Execute attribute*), 65
 trainer (*dicee.trainer.DICE_Trainer attribute*), 219
 trainer (*dicee.trainer.dice_trainer.DICE_Trainer attribute*), 214
 trainer (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 218
 training_step (*dicee.trainer.torch_trainer.TorchTrainer attribute*), 216
 training_step() (*dicee.models.base_model.BaseKGELightning method*), 84
 training_step() (*dicee.models.BaseKGELightning method*), 135
 training_step() (*dicee.models.transformers.BytE method*), 126
 training_step_outputs (*dicee.models.base_model.BaseKGELightning attribute*), 84
 training_step_outputs (*dicee.models.BaseKGELightning attribute*), 135
 training_technique (*dicee.knowledge_graph.KG attribute*), 68
 TransE (*class in dicee.models*), 149
 TransE (*class in dicee.models.real*), 120
 transformer (*dicee.models.clifford.KeciTransformer attribute*), 96
 transformer (*dicee.models.KeciTransformer attribute*), 181
 transformer (*dicee.models.transformers.BytE attribute*), 125
 transformer (*dicee.models.transformers.GPT attribute*), 130
 TransformerBlock (*class in dicee.models.clifford*), 97
 TransformerMLP (*class in dicee.models.clifford*), 97
 TransformerSelfAttention (*class in dicee.models.clifford*), 97
 trapezoid() (*dicee.models.FMult2 method*), 193
 trapezoid() (*dicee.models.function_space.FMult2 method*), 108
 tri_score() (*dicee.models.function_space.LFMult method*), 110
 tri_score() (*dicee.models.function_space.LFMult1 method*), 109
 tri_score() (*dicee.models.LFMult method*), 194
 tri_score() (*dicee.models.LFMult1 method*), 193
 triple_score() (*dicee.KGE method*), 230
 triple_score() (*dicee.knowledge_graph_embeddings.KGE method*), 72
 TriplePredictionDataset (*class in dicee.dataset_classes*), 39
 tuple2list() (*dicee.query_generator.QueryGenerator method*), 197
 tuple2list() (*dicee.QueryGenerator method*), 233
 TWA (*class in dicee.weight_averaging*), 224
 twa (*dicee.config.Namespace attribute*), 30
 twa_c_epochs (*dicee.weight_averaging.TWA attribute*), 224
 twa_model (*dicee.weight_averaging.TWA attribute*), 224
 twa_start_epoch (*dicee.weight_averaging.TWA attribute*), 224

U

unlabelled_size (*dicee.callbacks.PseudoLabellingCallback attribute*), 24
 unmap() (*dicee.query_generator.QueryGenerator method*), 197
 unmap() (*dicee.QueryGenerator method*), 233
 unmap_query() (*dicee.query_generator.QueryGenerator method*), 197
 unmap_query() (*dicee.QueryGenerator method*), 233
 update_hits() (*in module dicee.evaluation.utils*), 53
 use_clifford_mul (*dicee.models.clifford.KeciTransformer attribute*), 96
 use_clifford_mul (*dicee.models.KeciTransformer attribute*), 181

V

val_aswa (*dicee.weight_averaging.ASWA attribute*), 220
 val_dataloader() (*dicee.models.base_model.BaseKGELightning method*), 86
 val_dataloader() (*dicee.models.BaseKGELightning method*), 137
 val_h1 (*dicee.analyse_experiments.Experiment attribute*), 21
 val_h3 (*dicee.analyse_experiments.Experiment attribute*), 21
 val_h10 (*dicee.analyse_experiments.Experiment attribute*), 21

- `val_mrr` (*dicee.analyse_experiments.Experiment* attribute), 21
- `val_path` (*dicee.query_generator.QueryGenerator* attribute), 196
- `val_path` (*dicee.QueryGenerator* attribute), 233
- `VALID_SCORING_TECHNIQUES` (in module *dicee.evaluation.evaluator*), 45
- `valid_set` (*dicee.knowledge_graph.KG* attribute), 68, 69
- `validate_knowledge_graph()` (in module *dicee.sanity_checkers*), 204
- `var_clamp` (*dicee.weight_averaging.SWAG* attribute), 222
- `vocab_preparation()` (*dicee.evaluation.Evaluator* method), 56
- `vocab_preparation()` (*dicee.evaluation.evaluator.Evaluator* method), 46
- `vocab_preparation()` (*dicee.Evaluator* method), 236
- `vocab_preparation()` (*dicee.evaluator.Evaluator* method), 62
- `vocab_size` (*dicee.models.CoKEConfig* attribute), 152
- `vocab_size` (*dicee.models.real.CoKEConfig* attribute), 123
- `vocab_size` (*dicee.models.transformers.GPTConfig* attribute), 129
- `vocab_to_parquet()` (in module *dicee.static_funcs*), 210
- `vtp_score()` (*dicee.models.function_space.LFMult* method), 110
- `vtp_score()` (*dicee.models.function_space.LFMultI* method), 109
- `vtp_score()` (*dicee.models.LFMult* method), 194
- `vtp_score()` (*dicee.models.LFMultI* method), 193

W

- `warmup_steps` (*dicee.callbacks.LRScheduler* attribute), 27
- `weight` (*dicee.models.transformers.LayerNorm* attribute), 126
- `weight_decay` (*dicee.config.Namespace* attribute), 29
- `weight_decay` (*dicee.models.base_model.BaseKGE* attribute), 89
- `weight_decay` (*dicee.models.BaseKGE* attribute), 139, 145, 154, 161, 168, 182, 188
- `weight_samples` (*dicee.weight_averaging.TWA* attribute), 224
- `weights` (*dicee.models.FMult* attribute), 191
- `weights` (*dicee.models.function_space.FMult* attribute), 107
- `weights` (*dicee.models.function_space.GFMult* attribute), 107
- `weights` (*dicee.models.GFMult* attribute), 192
- `world_size` (*dicee.Execute* attribute), 226
- `world_size` (*dicee.executer.Execute* attribute), 65
- `write_csv_from_model_parallel()` (in module *dicee.static_funcs*), 210
- `write_links()` (*dicee.query_generator.QueryGenerator* method), 197
- `write_links()` (*dicee.QueryGenerator* method), 233
- `write_report()` (*dicee.Execute* method), 227
- `write_report()` (*dicee.executer.Execute* method), 66

X

- `x_values` (*dicee.models.function_space.LFMult* attribute), 109
- `x_values` (*dicee.models.LFMult* attribute), 194