
DICE Embeddings

Release 0.1.3.2

Caglar Demir

Aug 18, 2026

Contents:

1	Dicee Manual	2
2	Installation	3
2.1	Installation from Source	3
3	Download Knowledge Graphs	3
4	Knowledge Graph Embedding Models	3
5	How to Train	3
6	Creating an Embedding Vector Database	5
6.1	Learning Embeddings	5
6.2	Loading Embeddings into Qdrant Vector Database	6
6.3	Launching Webservice	6
7	Answering Complex Queries	6
8	Predicting Missing Links	8
9	Downloading Pretrained Models	8
10	How to Deploy	8
11	Docker	8
12	Coverage Report	8
13	How to cite	10
14	dicee	12
14.1	Submodules	12
14.2	Attributes	244
14.3	Classes	244
14.4	Package Contents	245
	Python Module Index	259

DICE Embeddings¹: Hardware-agnostic Framework for Large-scale Knowledge Graph Embeddings:

1 Dicee Manual

Version: dicee 0.3.2

GitHub repository: <https://github.com/dice-group/dice-embeddings>

Publisher and maintainer: Caglar Demir²

Contact: caglar.demir@upb.de

License: OSI Approved :: MIT License

Dicee is a hardware-agnostic framework for large-scale knowledge graph embeddings.

Knowledge graph embedding research has mainly focused on learning continuous representations of knowledge graphs towards the link prediction problem. Recently developed frameworks can be effectively applied in a wide range of research-related applications. Yet, using these frameworks in real-world applications becomes more challenging as the size of the knowledge graph grows

We developed the DICE Embeddings framework (dicee) to compute embeddings for large-scale knowledge graphs in a hardware-agnostic manner. To achieve this goal, we rely on

1. **Pandas**³ & Co. to use parallelism at preprocessing a large knowledge graph,
2. **PyTorch**⁴ & Co. to learn knowledge graph embeddings via multi-CPUs, GPUs, TPUs or computing cluster, and
3. **Huggingface**⁵ to ease the deployment of pre-trained models.

Why Pandas⁶ & Co. ? A large knowledge graph can be read and preprocessed (e.g. removing literals) by pandas, modin, or polars in parallel. Through polars, a knowledge graph having more than 1 billion triples can be read in parallel fashion. Importantly, using these frameworks allow us to perform all necessary computations on a single CPU as well as a cluster of computers.

Why PyTorch⁷ & Co. ? PyTorch is one of the most popular machine learning frameworks available at the time of writing. PytorchLightning facilitates scaling the training procedure of PyTorch without boilerplate. In our framework, we combine **PyTorch**⁸ & **PytorchLightning**⁹. Users can choose the trainer class (e.g., DDP by Pytorch) to train large knowledge graph embedding models with billions of parameters. PytorchLightning allows us to use state-of-the-art model parallelism techniques (e.g. Fully Sharded Training, FairScale, or DeepSpeed) without extra effort. With our framework, practitioners can directly use PytorchLightning for model parallelism to train gigantic embedding models.

Why Huggingface¹⁰? Seamlessly deploy and share pre-trained embedding models through the Huggingface ecosystem.

¹ <https://github.com/dice-group/dice-embeddings>

² <https://github.com/Demirrr>

³ <https://pandas.pydata.org/>

⁴ <https://pytorch.org/>

⁵ <https://huggingface.co/>

⁶ <https://pandas.pydata.org/>

⁷ <https://pytorch.org/>

⁸ <https://pytorch.org/>

⁹ <https://www.pytorchlightning.ai/>

¹⁰ <https://huggingface.co/>

2 Installation

2.1 Installation from Source

```
git clone https://github.com/dice-group/dice-embeddings.git
conda create -n dice python=3.10.13 --no-default-packages && conda activate dice &&
↪ cd dice-embeddings &&
pip3 install -e .
```

or

```
pip install dicee
```

3 Download Knowledge Graphs

```
wget https://files.dice-research.org/datasets/dice-embeddings/KGs.zip --no-check-
↪ certificate && unzip KGs.zip
```

To test the Installation

```
python -m pytest -p no:warnings -x # Runs >114 tests leading to > 15 mins
python -m pytest -p no:warnings --lf # run only the last failed test
python -m pytest -p no:warnings --ff # to run the failures first and then the rest of
↪ the tests.
```

4 Knowledge Graph Embedding Models

1. TransE, DistMult, ComplEx, ConEx, QMult, OMult, ConvO, ConvQ, Keci
2. All 44 models available in <https://github.com/pykeen/pykeen#models>

For more, please refer to examples.

5 How to Train

To Train a KGE model (KECI) and evaluate it on the train, validation, and test sets of the UMLS benchmark dataset.

```
from dicee.executer import Execute
from dicee.config import Namespace
args = Namespace()
args.model = 'Keci'
args.scoring_technique = "KvsAll" # 1vsAll, or AllvsAll, or NegSample
args.dataset_dir = "KGs/UMLS"
args.path_to_store_single_run = "Keci_UMLS"
args.num_epochs = 100
args.embedding_dim = 32
args.batch_size = 1024
reports = Execute(args).start()
print(reports["Train"]["MRR"]) # => 0.9912
print(reports["Test"]["MRR"]) # => 0.8155
# See the Keci_UMLS folder embeddings and all other files
```

where the data is in the following form

```
$ head -3 KGs/UMLS/train.txt
acquired_abnormality    location_of             experimental_model_of_disease
anatomical_abnormality  manifestation_of        physiologic_function
alga    isa      entity
```

A KGE model can also be trained from the command line

```
dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
```

dicee automatically detects available GPUs and trains a model with distributed data parallels technique. Under the hood, dicee uses lightning as a default trainer.

```
# Train a model by only using the GPU-0
CUDA_VISIBLE_DEVICES=0 dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model
↪ "train_val_test"
# Train a model by only using GPU-1
CUDA_VISIBLE_DEVICES=1 dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model
↪ "train_val_test"
NCCL_P2P_DISABLE=1 CUDA_VISIBLE_DEVICES=0,1 python dicee/scripts/run.py --trainer PL -
↪ --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
```

Under the hood, dicee executes run.py script and uses lightning as a default trainer

```
# Two equivalent executions
# (1)
dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
↪ 9753123402351737}
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
↪ 'MRR': 0.8072362996241839}
# Evaluate Keci on Test set: Evaluate Keci on Test set
# {'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
↪ 'MRR': 0.8064032293278861}

# (2)
CUDA_VISIBLE_DEVICES=0,1 python dicee/scripts/run.py --trainer PL --dataset_dir "KGs/
↪ UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
↪ 9753123402351737}
# Evaluate Keci on Train set: Evaluate Keci on Train set
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
↪ 'MRR': 0.8072362996241839}
# Evaluate Keci on Test set: Evaluate Keci on Test set
# {'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
↪ 'MRR': 0.8064032293278861}
```

Similarly, models can be easily trained with torchrun

```
torchrun --standalone --nnodes=1 --nproc_per_node=gpu dicee/scripts/run.py --trainer_
→torchDDP --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
→9753123402351737}
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
→'MRR': 0.8072499937521418}
# Evaluate Keci on Test set: Evaluate Keci on Test set
{'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
→'MRR': 0.8064032293278861}
```

You can also train a model in multi-node multi-gpu setting.

```
torchrun --nnodes 2 --nproc_per_node=gpu --node_rank 0 --rdzv_id 455 --rdzv_backend_
→c10d --rdzv_endpoint=nebula dicee/scripts/run.py --trainer torchDDP --dataset_dir_
→KGs/UMLS
torchrun --nnodes 2 --nproc_per_node=gpu --node_rank 1 --rdzv_id 455 --rdzv_backend_
→c10d --rdzv_endpoint=nebula dicee/scripts/run.py --trainer torchDDP --dataset_dir_
→KGs/UMLS
```

Train a KGE model by providing the path of a single file and store all parameters under newly created directory called KeciFamilyRun.

```
dicee --path_single_kg "KGs/Family/family-benchmark_rich_background.owl" --model Keci_
→--path_to_store_single_run KeciFamilyRun --backend rdflib
```

where the data is in the following form

```
$ head -3 KGs/Family/train.txt
_:1 <http://www.w3.org/1999/02/22-rdf-syntax-ns#type> <http://www.w3.org/2002/07/owl
→#Ontology> .
<http://www.benchmark.org/family#hasChild> <http://www.w3.org/1999/02/22-rdf-syntax-ns
→#type> <http://www.w3.org/2002/07/owl#ObjectProperty> .
<http://www.benchmark.org/family#hasParent> <http://www.w3.org/1999/02/22-rdf-syntax-
→ns#type> <http://www.w3.org/2002/07/owl#ObjectProperty> .
```

Apart from n-triples or standard link prediction dataset formats, we support ["owl", "nt", "turtle", "rdf/xml", "n3"]*. Moreover, a KGE model can be also trained by providing an endpoint of a triple store.

```
dicee --sparql_endpoint "http://localhost:3030/mutagenesis/" --model Keci
```

For more, please refer to examples.

6 Creating an Embedding Vector Database

6.1 Learning Embeddings

```
# Train an embedding model
dicee --dataset_dir KGs/Countries-S1 --path_to_store_single_run CountryEmbeddings --
→model Keci --p 0 --q 1 --embedding_dim 32 --adaptive_swa
```

6.2 Loading Embeddings into Qdrant Vector Database

```
# Ensure that Qdrant available
# docker pull qdrant/qdrant && docker run -p 6333:6333 -p 6334:6334 -v $(pwd)/
↪qdrant_storage:/qdrant/storage:z qdrant/qdrant
diceeindex --path_model "CountryEmbeddings" --collection_name "dummy" --location
↪"localhost"
```

6.3 Launching Webservice

```
diceeserve --path_model "CountryEmbeddings" --collection_name "dummy" --collection_
↪location "localhost"
```

Retrieve and Search

Get embedding of germany

```
curl -X 'GET' 'http://0.0.0.0:8000/api/get?q=germany' -H 'accept: application/json'
```

Get most similar things to europe

```
curl -X 'GET' 'http://0.0.0.0:8000/api/search?q=europe' -H 'accept: application/json'
{"result":[{"hit":"europe","score":1.0},
{"hit":"northern_europe","score":0.67126536},
{"hit":"western_europe","score":0.6010134},
{"hit":"puerto_rico","score":0.5051694},
{"hit":"southern_europe","score":0.4829831}]}
```

7 Answering Complex Queries

```
# pip install dicee
# wget https://files.dice-research.org/datasets/dice-embeddings/KGs.zip --no-check-
↪certificate & unzip KGs.zip
from dicee.executor import Execute
from dicee.config import Namespace
from dicee.knowledge_graph_embeddings import KGE
# (1) Train a KGE model
args = Namespace()
args.model = 'Keci'
args.p=0
args.q=1
args.optim = 'Adam'
args.scoring_technique = "AllvsAll"
args.path_single_kg = "KGs/Family/family-benchmark_rich_background.owl"
args.backend = "rdflib"
args.num_epochs = 200
args.batch_size = 1024
args.lr = 0.1
args.embedding_dim = 512
result = Execute(args).start()
# (2) Load the pre-trained model
```

(continues on next page)

(continued from previous page)

```
pre_trained_kge = KGE(path=result['path_experiment_folder'])
# (3) Single-hop query answering
# Query: ?E : \exist E.hasSibling(E, F9M167)
# Question: Who are the siblings of F9M167?
# Answer: [F9M157, F9F141], as (F9M167, hasSibling, F9M157) and (F9M167, hasSibling, ↵
↵F9F141)
predictions = pre_trained_kge.answer_multi_hop_query(query_type="1p",
                                                        query=('http://www.benchmark.org/
↵family#F9M167',
                                                                ('http://www.benchmark.
↵org/family#hasSibling',)),
                                                        tnorm="min", k=3)
top_entities = [topk_entity for topk_entity, query_score in predictions]
assert "http://www.benchmark.org/family#F9F141" in top_entities
assert "http://www.benchmark.org/family#F9M157" in top_entities
# (2) Two-hop query answering
# Query: ?D : \exist E.Married(D, E) \land hasSibling(E, F9M167)
# Question: To whom a sibling of F9M167 is married to?
# Answer: [F9F158, F9M142] as (F9M157 #married F9F158) and (F9F141 #married F9M142)
predictions = pre_trained_kge.answer_multi_hop_query(query_type="2p",
                                                        query=("http://www.benchmark.org/
↵family#F9M167",
                                                                ("http://www.benchmark.
↵org/family#hasSibling",
                                                                "http://www.benchmark.
↵org/family#married")),
                                                        tnorm="min", k=3)
top_entities = [topk_entity for topk_entity, query_score in predictions]
assert "http://www.benchmark.org/family#F9M142" in top_entities
assert "http://www.benchmark.org/family#F9F158" in top_entities
# (3) Three-hop query answering
# Query: ?T : \exist D.type(D,T) \land Married(D,E) \land hasSibling(E, F9M167)
# Question: What are the type of people who are married to a sibling of F9M167?
# (3) Answer: [Person, Male, Father] since F9M157 is [Brother Father Grandfather ↵
↵Male] and F9M142 is [Male Grandfather Father]
predictions = pre_trained_kge.answer_multi_hop_query(query_type="3p", query=("http://
↵www.benchmark.org/family#F9M167",
                                                                ("http://
↵www.benchmark.org/family#hasSibling",
                                                                "http://
↵www.benchmark.org/family#married",
                                                                "http://
↵www.w3.org/1999/02/22-rdf-syntax-ns#type")),
                                                        tnorm="min", k=5)
top_entities = [topk_entity for topk_entity, query_score in predictions]
print(top_entities)
assert "http://www.benchmark.org/family#Person" in top_entities
assert "http://www.benchmark.org/family#Father" in top_entities
assert "http://www.benchmark.org/family#Male" in top_entities
```

For more, please refer to examples/multi_hop_query_answering.

8 Predicting Missing Links

```
from dicee import KGE
# (1) Train a knowledge graph embedding model..
# (2) Load a pretrained model
pre_trained_kge = KGE(path='../')
# (3) Predict missing links through head entity rankings
pre_trained_kge.predict_topk(h=[".."],r=[".."],topk=10)
# (4) Predict missing links through relation rankings
pre_trained_kge.predict_topk(h=[".."],t=[".."],topk=10)
# (5) Predict missing links through tail entity rankings
pre_trained_kge.predict_topk(r=[".."],t=[".."],topk=10)
```

9 Downloading Pretrained Models

```
from dicee import KGE
# (1) Load a pretrained ConEx on DBpedia
model = KGE(url="https://files.dice-research.org/projects/DiceEmbeddings/KINSHIP-Keci-
↳dim128-epoch256-KvsAll")
```

- For more please look at dice-research.org/projects/DiceEmbeddings/¹¹

10 How to Deploy

```
from dicee import KGE
KGE(path='../').deploy(share=True, top_k=10)
```

11 Docker

To build the Docker image:

```
docker build -t dice-embeddings .
```

To test the Docker image:

```
docker run --rm -v ~/.local/share/dicee/KGs:/dicee/KGs dice-embeddings ./main.py --
↳model AConEx --embedding_dim 16
```

12 Coverage Report

The coverage report is generated using `coverage.py`¹²:

Name	Stmts	Miss	Cover	Missing
-----	-----	-----	-----	-----
dicee/___init___ .py	7	0	100%	
dicee/abstracts.py	338	115	66%	112-113, 115

(continues on next page)

¹¹ <https://files.dice-research.org/projects/DiceEmbeddings/>

¹² <https://coverage.readthedocs.io/en/7.6.0/>

(continued from previous page)

↪131, 154-155, 160, 173, 197, 240-254, 290, 303-306, 309-313, 353-364, 379-387, 402, ↪				
↪413-417, 427-428, 434-436, 442-445, 448-453, 576-596, 602-606, 610-612, 631, 658-696				
dicee/callbacks.py	248	103	58%	50-55, ↪
↪67-73, 76, 88-93, 98-103, 106-109, 116-133, 138-142, 146-147, 247, 281-285, 291-292, ↪				
↪ 310-316, 319, 324-325, 337-343, 349-358, 363-365, 410, 421-434, 438-473, 485-491				
dicee/config.py	97	2	98%	146-147
dicee/dataset_classes.py	430	146	66%	16, 44, ↪
↪57, 89-98, 104, 111-118, 121, 124, 127-151, 207-213, 216, 219-221, 324, 335-338, ↪				
↪354, 420-421, 439, 562-581, 583, 587-599, 606-615, 618, 622-636, 780-787, 790-794, ↪				
↪845, 866-878, 902-915, 937, 941-954, 964-967, 973, 985, 987, 989, 1012-1022				
dicee/eval_static_funcs.py	256	100	61%	104, 109, ↪
↪ 114, 261-356, 363-414, 442, 465-468				
dicee/evaluator.py	267	48	82%	48, 53, ↪
↪58, 77, 82-83, 86, 102, 119, 130, 134, 139, 173-184, 191-202, 310, 340-358, 452, ↪				
↪462, 480-485				
dicee/executer.py	134	16	88%	53-57, ↪
↪166-176, 235-236, 283				
dicee/knowledge_graph.py	82	10	88%	84, 94- ↪
↪95, 124, 128, 132-134, 137-138, 140				
dicee/knowledge_graph_embeddings.py	654	415	37%	25, 28- ↪
↪29, 37-50, 55-88, 91-125, 129-137, 171, 173-229, 261, 265, 276-277, 301-303, 311, ↪				
↪339-362, 493, 497-519, 523-547, 580, 656, 665, 710-716, 748, 806-1171, 1202-1263, ↪				
↪1267-1295, 1326, 1332				
dicee/models/___init___py	9	0	100%	
dicee/models/adopt.py	187	172	8%	50-86, ↪
↪99-110, 129-185, 195-242, 266-322, 346-448, 484-517				
dicee/models/base_model.py	240	35	85%	30-35, ↪
↪64, 66, 92, 99-116, 171, 204, 244, 250, 259, 262, 266, 273, 277, 279, 294, 307-308, ↪				
↪362, 365, 438, 450				
dicee/models/clifford.py	470	278	41%	10, 12, ↪
↪16, 24-25, 52-56, 79-87, 101-103, 108-109, 140-160, 184, 191, 195-256, 273-277, 289, ↪				
↪ 292, 297, 302, 346-361, 377-444, 464-470, 483, 486, 491, 496, 525-531, 544, 547, ↪				
↪552, 557, 567-576, 592-593, 613-685, 696-699, 724-749, 773-806, 842-846, 859, 869, ↪				
↪872, 877, 882, 887, 891, 895, 904-905, 935, 942, 947, 975-979, 1007-1016, 1026-1034, ↪				
↪ 1052-1054, 1072-1074, 1090-1092				
dicee/models/complex.py	162	25	85%	86-109, ↪
↪273-287				
dicee/models/dualE.py	59	10	83%	93-102, ↪
↪142-156				
dicee/models/ensemble.py	89	67	25%	7-29, 31, ↪
↪ 34, 37, 40, 43, 46, 49, 52-54, 56-58, 64-68, 71-90, 93-94, 97-112, 131				
dicee/models/function_space.py	262	221	16%	10-23, ↪
↪27-36, 39-48, 52-69, 76-87, 90-99, 102-111, 115-127, 135-157, 160-166, 169-186, 189- ↪				
↪195, 198-206, 209, 214-235, 244-247, 251-255, 259-268, 272-293, 302-308, 312-329, ↪				
↪333-336, 345-353, 356, 367-373, 393-407, 425-439, 444-454, 462-466, 475-479				
dicee/models/literal.py	33	1	97%	82
dicee/models/octonion.py	227	83	63%	21-44, ↪
↪320-329, 334-345, 348-370, 374-416, 426-474				
dicee/models/pykeen_models.py	55	5	91%	77-80, ↪
↪135				
dicee/models/quaternion.py	192	69	64%	7-21, 30- ↪
↪55, 68-72, 107, 185, 328-342, 345-364, 368-389, 399-426				

(continues on next page)

(continued from previous page)

dicee/models/real.py	61	12	80%	37-42, ↩
↪70-73, 91, 107-110				
dicee/models/static_funcs.py	10	0	100%	
dicee/models/transformers.py	234	189	19%	20-39, ↩
↪42, 56-71, 80-98, 101-112, 119-121, 124, 130-147, 151-176, 182-186, 189-193, 199-				
↪203, 206-208, 225-252, 261-264, 267-272, 275-300, 306-311, 315-368, 372-394, 400-410				
dicee/query_generator.py	374	346	7%	17-51, ↩
↪55, 61-64, 68-69, 77-91, 99-146, 154-187, 191-205, 211-268, 273-302, 306-442, 452-				
↪471, 479-502, 509-513, 518, 523-529				
dicee/read_preprocess_save_load_kg/___init___py	3	0	100%	
dicee/read_preprocess_save_load_kg/preprocess.py	243	40	84%	33, 39, ↩
↪76, 100-125, 131, 136-149, 175, 205, 380-381				
dicee/read_preprocess_save_load_kg/read_from_disk.py	36	11	69%	34, 38-
↪40, 47, 55, 58-72				
dicee/read_preprocess_save_load_kg/save_load_disk.py	53	21	60%	29-30, ↩
↪38, 47-68				
dicee/read_preprocess_save_load_kg/util.py	236	125	47%	159, 173-
↪175, 179-180, 198-204, 207-209, 214-216, 230, 244-247, 252-260, 265-271, 276-281, ↩				
↪286-291, 303-324, 330-386, 390-394, 398-399, 403, 407-408, 436, 441, 448-449				
dicee/sanity_checkers.py	47	19	60%	8-12, 21-
↪31, 46, 51, 58, 69-79				
dicee/static_funcs.py	483	194	60%	42, 52, ↩
↪58-63, 85, 92-96, 109-119, 129-131, 136, 143, 167, 172, 184, 190, 198, 202, 229-233,				
↪295, 303-309, 320-330, 341-361, 389, 413-414, 419-420, 437-438, 440-441, 443-444, ↩				
↪452, 470-474, 491-494, 498-503, 507-511, 515-516, 522-524, 539-553, 558-561, 566-				
↪569, 578-629, 634-646, 663-680, 683-691, 695-713, 724				
dicee/static_funcs_training.py	155	66	57%	7-10, ↩
↪222-319, 327-328				
dicee/static_preprocess_funcs.py	98	43	56%	17-25, ↩
↪50, 57, 59, 70, 83-107, 112-115, 120-123, 128-131				
dicee/trainer/___init___py	1	0	100%	
dicee/trainer/dice_trainer.py	151	18	88%	22, 30-
↪31, 33-35, 97, 104, 109-114, 152, 237, 280-283				
dicee/trainer/model_parallelism.py	99	87	12%	10-25, ↩
↪30-116, 121-132, 136, 141-197				
dicee/trainer/torch_trainer.py	77	6	92%	31, 102, ↩
↪168, 179-181				
dicee/trainer/torch_trainer_ddp.py	89	71	20%	11-14, ↩
↪43, 47-67, 78-94, 113-122, 126-136, 151-158, 168-191				

TOTAL	6948	3169	54%	

13 How to cite

Currently, we are working on our manuscript describing our framework. If you really like our work and want to cite it now, feel free to chose one :)

```
# Keci
@inproceedings{demir2023clifford,
  title={Clifford Embeddings--A Generalized Approach for Embedding in Normed Algebras}
  ↪,
```

(continues on next page)

```

    author={Demir, Caglar and Ngonga Ngomo, Axel-Cyrille},
    booktitle={Joint European Conference on Machine Learning and Knowledge Discovery in
↪Databases},
    pages={567--582},
    year={2023},
    organization={Springer}
}
# LitCQD
@inproceedings{demir2023litcq,
    title={LitCQD: Multi-Hop Reasoning in Incomplete Knowledge Graphs with Numeric
↪Literals},
    author={Demir, Caglar and Wiebesiek, Michel and Lu, Renzhong and Ngonga Ngomo, Axel-
↪Cyrille and Heindorf, Stefan},
    booktitle={Joint European Conference on Machine Learning and Knowledge Discovery in
↪Databases},
    pages={617--633},
    year={2023},
    organization={Springer}
}
# DICE Embedding Framework
@article{demir2022hardware,
    title={Hardware-agnostic computation for large-scale knowledge graph embeddings},
    author={Demir, Caglar and Ngomo, Axel-Cyrille Ngonga},
    journal={Software Impacts},
    year={2022},
    publisher={Elsevier}
}
# KronE
@inproceedings{demir2022kronecker,
    title={Kronecker decomposition for knowledge graph embeddings},
    author={Demir, Caglar and Lienen, Julian and Ngonga Ngomo, Axel-Cyrille},
    booktitle={Proceedings of the 33rd ACM Conference on Hypertext and Social Media},
    pages={1--10},
    year={2022}
}
# QMult, OMult, ConvQ, ConvO
@InProceedings{pmlr-v157-demir21a,
    title = {Convolutional Hypercomplex Embeddings for Link Prediction},
    author = {Demir, Caglar and Moussallem, Diego and Heindorf, Stefan and Ngonga
↪Ngomo, Axel-Cyrille},
    booktitle = {Proceedings of The 13th Asian Conference on Machine Learning},
    pages = {656--671},
    year = {2021},
    editor = {Balasubramanian, Vineeth N. and Tsang, Ivor},
    volume = {157},
    series = {Proceedings of Machine Learning Research},
    month = {17--19 Nov},
    publisher = {PMLR},
    pdf = {https://proceedings.mlr.press/v157/demir21a/demir21a.pdf},
    url = {https://proceedings.mlr.press/v157/demir21a.html},
}
# ConEx

```

```
@inproceedings{demir2021convolutional,
  title={Convolutional Complex Knowledge Graph Embeddings},
  author={Caglar Demir and Axel-Cyrille Ngonga Ngomo},
  booktitle={Eighteenth Extended Semantic Web Conference - Research Track},
  year={2021},
  url={https://openreview.net/forum?id=6T45-4TFqaX}}
# Shallom
@inproceedings{demir2021shallow,
  title={A shallow neural model for relation prediction},
  author={Demir, Caglar and Moussallem, Diego and Ngomo, Axel-Cyrille Ngonga},
  booktitle={2021 IEEE 15th International Conference on Semantic Computing (ICSC)},
  pages={179--182},
  year={2021},
  organization={IEEE}
```

For any questions or wishes, please contact: caglar.demir@upb.de

14 dicee

DICE Embeddings - Knowledge Graph Embedding Library.

A library for training and using knowledge graph embedding models with support for various scoring techniques and training strategies.

Submodules:

evaluation: Model evaluation functions and Evaluator class models: KGE model implementations trainer: Training orchestration scripts: Utility scripts

14.1 Submodules

`dicee.__main__`

`dicee.abstracts`

Attributes

logger

Classes

<i>AbstractTrainer</i>	Abstract base class for KGE model trainers.
<i>BaseInteractiveKGE</i>	Base class for interactive, post-training use of KGE models.
<i>InteractiveQueryDecomposition</i>	Mixin that provides fuzzy-logic operators for multi-hop EPFO query answering.
<i>AbstractCallback</i>	Abstract base class for KGE training lifecycle callbacks.
<i>AbstractPPECallback</i>	Abstract base class for Post-training Parameter Ensembling (PPE) callbacks.
<i>BaseInteractiveTrainKGE</i>	Abstract/base class for training knowledge graph embedding models interactively.

Module Contents

`dicee.abstracts.logger`

class `dicee.abstracts.AbstractTrainer` (*args, callbacks*)

Abstract base class for KGE model trainers.

Provides the callback dispatch mechanism shared by all concrete trainer implementations (TorchTrainer, TorchD-
DPTrainer, etc.). Sub-classes call the `on_*` hooks at the appropriate points in the training loop so that any registered *AbstractCallback* can react.

Parameters

- **args** (*argparse.Namespace or similar*) – Processed configuration object. Must expose at least `random_seed` (int).
- **callbacks** (*list of AbstractCallback*) – Ordered list of callback instances to invoke at each lifecycle hook.

attributes

callbacks

is_global_zero = `True`

global_rank = `0`

local_rank = `0`

strategy = `None`

on_fit_start (**args, **kwargs*)

Dispatch `on_fit_start` to all registered callbacks.

Called once before the first training epoch begins.

on_fit_end (**args, **kwargs*)

Dispatch `on_fit_end` to all registered callbacks.

Called once after the last training epoch completes.

on_train_epoch_start (**args, **kwargs*)

Dispatch `on_train_epoch_start` to all registered callbacks.

Called at the beginning of every epoch.

on_train_epoch_end (**args, **kwargs*)

Dispatch `on_train_epoch_end` to all registered callbacks.

Called at the end of every epoch after the loss has been accumulated.

on_train_batch_end (**args, **kwargs*)

Dispatch `on_train_batch_end` to all registered callbacks.

Called after each mini-batch gradient update.

static save_checkpoint (*full_path: str, model*) → `None`

Persist model weights to disk.

Parameters

- **full_path** (*str*) – Absolute or relative file path (including filename) where the `state_dict` will be written, e.g. `'Experiments/run1/model.pt'`.

- **model** (*torch.nn.Module*) – The model whose `state_dict` is to be saved.

```
class dicee.abstracts.BaseInteractiveKGE (path: str = None, url: str = None,  

construct_ensemble: bool = False, model_name: str = None,  

apply_semantic_constraint: bool = False)
```

Base class for interactive, post-training use of KGE models.

Loads a pre-trained model from disk (or a remote URL) together with its entity/relation index mappings and exposes the prediction API used by *KGE*.

Parameters

- **path** (*str, optional*) – Path to the experiment directory produced by `Execute`. Must contain `model.pt`, `configuration.json`, `entity_to_idx.csv` and `relation_to_idx.csv`.
- **url** (*str, optional*) – Remote URL of a pre-trained model to download. Mutually exclusive with *path*.
- **construct_ensemble** (*bool, optional*) – When `True`, load all checkpoint files in *path* and average their weights to form an ensemble model. Defaults to `False`.
- **model_name** (*str, optional*) – Filename (without extension) of the checkpoint to load when multiple `.pt` files exist in *path*.
- **apply_semantic_constraint** (*bool, optional*) – Reserved for future use. Defaults to `False`.

```
construct_ensemble = False
```

```
apply_semantic_constraint = False
```

```
configs
```

```
get_eval_report () → dict
```

```
get_bpe_token_representation (str_entity_or_relation: List[str] | str) → List[List[int]] | List[int]
```

Parameters

str_entity_or_relation (*corresponds to a str or a list of strings to be tokenized via BPE and shaped.*)

Return type

A list integer(s) or a list of lists containing integer(s)

```
get_padded_bpe_triple_representation (triples: List[List[str]]) → Tuple[List, List, List]
```

Parameters

triples

```
set_model_train_mode () → None
```

Switch the underlying model to training mode.

Calls `model.train()` and re-enables gradient computation for all parameters so that subsequent calls to optimisation steps work correctly after a period of inference.

```
set_model_eval_mode () → None
```

Switch the underlying model to evaluation mode.

Calls `model.eval()` and freezes all parameters (`requires_grad = False`) so that dropout and batch-norm layers behave deterministically during inference.

property name

sample_entity (*n*: int) → List[str]

Return *n* random entity strings without replacement.

Parameters

n (*int*) – Number of entities to sample. Must be non-negative and at most `num_entities`.

Returns

Randomly selected entity string labels.

Return type

List[str]

sample_relation (*n*: int) → List[str]

Return *n* random relation strings without replacement.

Parameters

n (*int*) – Number of relations to sample. Must be non-negative and at most `num_relations`.

Returns

Randomly selected relation string labels.

Return type

List[str]

is_seen (*entity*: str = None, *relation*: str = None) → bool

Check whether an entity or relation was present in the training set.

Exactly one of *entity* or *relation* should be provided.

Parameters

- **entity** (*str*, *optional*) – Entity string label to look up.
- **relation** (*str*, *optional*) – Relation string label to look up.

Returns

True if the given string is in the respective index mapping, False otherwise.

Return type

bool

save () → None

Persist the current model weights to the experiment directory.

The checkpoint filename encodes the current timestamp so successive calls do not overwrite each other. Ensemble models are saved with an `_ensemble_` infix in the filename.

get_entity_index (*x*: str) → int

Return the integer index for a given entity string.

Parameters

x (*str*) – Entity string label (must have been seen during training).

Returns

Corresponding row index in the entity embedding matrix.

Return type

int

get_relation_index (*x: str*) → int

Return the integer index for a given relation string.

Parameters

x (*str*) – Relation string label (must have been seen during training).

Returns

Corresponding row index in the relation embedding matrix.

Return type

int

index_triple (*head_entity: List[str], relation: List[str], tail_entity: List[str]*)
→ Tuple[torch.LongTensor, torch.LongTensor, torch.LongTensor]

Convert string triple lists to integer index tensors.

Parameters

- **head_entity** (*List[str]*) – Head entity string labels.
- **relation** (*List[str]*) – Relation string labels.
- **tail_entity** (*List[str]*) – Tail entity string labels.

Returns

idx_head_entity, idx_relation, idx_tail_entity – Each has shape (n, 1) containing the integer indices for the corresponding strings.

Return type

torch.LongTensor

add_new_entity_embeddings (*entity_name: str = None, embeddings: torch.FloatTensor = None*)
→ None

Extend the entity embedding table with a new entity at inference time.

The new entity is appended to both `entity_to_idx` / `idx_to_entity` mappings and the `entity_embeddings` weight tensor so that subsequent calls to prediction methods can reference it by name.

Parameters

- **entity_name** (*str*) – String label for the new entity. If the entity already exists in the index no modification is made.
- **embeddings** (*torch.FloatTensor*) – 1-D float tensor of length `embedding_dim` containing the pre-computed embedding for the new entity.

get_entity_embeddings (*items: List[str]*) → torch.FloatTensor

Return the embedding vectors for the given entity strings.

For standard (non-BPE) models the method looks up each string in `entity_to_idx` and returns the corresponding rows of the entity embedding matrix. For BPE models subword token embeddings are fetched and flattened into a single vector per entity.

Parameters

items (*List[str]*) – Entity string labels to retrieve.

Returns

Shape (len(items), embedding_dim).

Return type

torch.FloatTensor

get_relation_embeddings (*items: List[str]*) → torch.FloatTensor

Return the embedding vectors for the given relation strings.

Parameters

items (*List[str]*) – Relation string labels to retrieve.

Returns

Shape (len(items), embedding_dim).

Return type

torch.FloatTensor

construct_input_and_output (*head_entity: List[str], relation: List[str], tail_entity: List[str], labels*)
→ Tuple[torch.LongTensor, torch.FloatTensor]

Build an indexed triple tensor and a label tensor from string inputs.

Parameters

- **head_entity** (*List[str]*) – Head entity string labels.
- **relation** (*List[str]*) – Relation string labels.
- **tail_entity** (*List[str]*) – Tail entity string labels.
- **labels** (*list or array-like*) – Binary or soft labels (one per triple) used as training targets.

Returns

- **x** (*torch.LongTensor*) – Shape (n, 3) integer-indexed triples.
- **labels** (*torch.FloatTensor*) – Shape (n,) float label tensor.

parameters ()

class dicee.abstracts.**InteractiveQueryDecomposition**

Mixin that provides fuzzy-logic operators for multi-hop EPFO query answering.

The three families of operators — T-norm, T-conorm, and negation norm — are applied element-wise over entity score tensors to compose complex queries from atomic link-prediction results (e.g. 2p, 3p, 2i, ip, up).

t_norm (*tens_1: torch.Tensor, tens_2: torch.Tensor, tnorm: str = 'min'*) → torch.Tensor

Apply a T-norm to combine two entity score distributions.

Parameters

- **tens_1** (*torch.Tensor*) – Score tensors of identical shape, values in [0, 1].
- **tens_2** (*torch.Tensor*) – Score tensors of identical shape, values in [0, 1].
- **tnorm** (*str*) – Operator to use. 'min' applies the Gödel (min) T-norm; 'prod' applies the product T-norm.

Returns

Element-wise combined scores of the same shape as the inputs.

Return type

torch.Tensor

tensor_t_norm (*subquery_scores: torch.FloatTensor, tnorm: str = 'min'*) → torch.FloatTensor

Compute T-norm over $[0,1]^{n \times d}$ where n denotes the number of hops and d denotes number of entities

t_conorm (*tens_1*: torch.Tensor, *tens_2*: torch.Tensor, *tconorm*: str = 'min') → torch.Tensor

Apply a T-conorm (S-norm) to combine two score distributions (union).

Parameters

- **tens_1** (torch.Tensor) – Score tensors of identical shape, values in [0, 1].
- **tens_2** (torch.Tensor) – Score tensors of identical shape, values in [0, 1].
- **tconorm** (str) – Operator to use. 'min' applies the Gödel (max) T-conorm; 'prod' applies the probabilistic sum T-conorm.

Returns

Element-wise combined scores of the same shape as the inputs.

Return type

torch.Tensor

negnorm (*tens_1*: torch.Tensor, *lambda_*: float, *neg_norm*: str = 'standard') → torch.Tensor

Apply a negation norm (complement) to an entity score distribution.

Parameters

- **tens_1** (torch.Tensor) – Input score tensor, values in [0, 1].
- **lambda** (float) – Shape parameter used by the Sugeno and Yager negation norms. Ignored for the standard complement.
- **neg_norm** (str) – Which negation to apply: 'standard' (1 - x), 'sugeno', or 'yager'.

Returns

Complemented score tensor of the same shape as *tens_1*.

Return type

torch.Tensor

class dicee.abstracts.**AbstractCallback**

Bases: abc.ABC, lightning.pytorch.callbacks.Callback

Abstract base class for KGE training lifecycle callbacks.

Concrete sub-classes override one or more hook methods to perform custom actions at specific points during training (e.g. weight averaging, periodic evaluation, model checkpointing). All hooks have empty default implementations so sub-classes only need to override the hooks they care about.

Callbacks are registered by passing them to the trainer's *callbacks* list. They are also compatible with PyTorch Lightning trainers because this class extends `lightning.pytorch.callbacks.Callback`.

on_init_start (*args, **kwargs)

Called when the trainer is about to be constructed.

Override to perform setup that must happen before any trainer state is initialised.

on_init_end (*args, **kwargs)

Called immediately after the trainer has been constructed.

Override to perform setup that requires a fully initialised trainer.

on_fit_start (trainer, model)

Called once before the first training epoch.

Parameters

- **trainer** (AbstractTrainer or pl.Trainer) – The active trainer instance.

- **model** (`BaseKGE`) – The model about to be trained.

`on_train_epoch_end(trainer, model)`

Called at the end of each training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

`on_train_batch_end(*args, **kwargs)`

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

`on_fit_end(*args, **kwargs)`

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

class `dicee.abstracts.AbstractPPECallback(num_epochs, path, epoch_to_start, last_percent_to_consider)`

Bases: `AbstractCallback`

Abstract base class for Post-training Parameter Ensembling (PPE) callbacks.

Sub-classes implement weight-averaging strategies (SWA, EMA, SWAG, ...) by overriding `on_train_epoch_end()` and `on_fit_end()`. Common book-keeping (epoch counter, sample counter, alpha weights) is managed here.

Parameters

- **num_epochs** (`int`) – Total number of training epochs.
- **path** (`str`) – Experiment directory where averaged checkpoints will be written.
- **epoch_to_start** (`int`) – First epoch at which the averaging procedure should begin.
- **last_percent_to_consider** (`float`) – Fraction of the total training epochs (counted from the end) whose checkpoints are included in the ensemble.

`num_epochs`

`path`

`sample_counter = 0`

`epoch_count = 0`

`alphas = None`

`on_fit_start(trainer, model)`

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

on_fit_end (*trainer, model*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

store_ensemble (*param_ensemble*) → None

class dicee.abstracts.**BaseInteractiveTrainKGE**

Abstract/base class for training knowledge graph embedding models interactively. This class provides methods for re-training KGE models and also Literal Embedding model.

train_triples (*h: List[str], r: List[str], t: List[str], labels: List[float], iteration=2, optimizer=None*)

train_k_vs_all (*h, r, iteration=1, lr=0.001*)

Train k vs all :param head_entity: :param relation: :param iteration: :param lr: :return:

train (*kg, lr=0.1, epoch=10, batch_size=32, neg_sample_ratio=10, num_workers=1*) → None

Retrained a pretrain model on an input KG via negative sampling.

train_literals (*train_file_path: str = None, num_epochs: int = 100, lit_lr: float = 0.001, lit_normalization_type: str = 'z-norm', batch_size: int = 1024, sampling_ratio: float = None, random_seed=1, loader_backend: str = 'pandas', freeze_entity_embeddings: bool = True, gate_residual: bool = True, device: str = None, shuffle_data: bool = True*)

Trains the Literal Embeddings model using literal data.

Parameters

- **train_file_path** (*str*) – Path to the training data file.
- **num_epochs** (*int*) – Number of training epochs.
- **lit_lr** (*float*) – Learning rate for the literal model.
- **norm_type** (*str*) – Normalization type to use ('z-norm', 'min-max', or None).
- **batch_size** (*int*) – Batch size for training.
- **sampling_ratio** (*float*) – Ratio of training triples to use.
- **loader_backend** (*str*) – Backend for loading the dataset ('pandas' or 'rdflib').
- **freeze_entity_embeddings** (*bool*) – If True, freeze the entity embeddings during training.
- **gate_residual** (*bool*) – If True, use gate residual connections in the model.
- **device** (*str*) – Device to use for training ('cuda' or 'cpu'). If None, will use available GPU or CPU.
- **shuffle_data** (*bool*) – If True, shuffle the dataset before training.

dicee.analyse_experiments

This script should be moved to dicee/scripts Example: `python dicee/analyse_experiments.py --dir Experiments --features "model" "trainMRR" "testMRR"`

Attributes

logger

Classes

Experiment

Functions

get_default_arguments()
analyse(args)

Module Contents

```
dicee.analyse_experiments.logger

dicee.analyse_experiments.get_default_arguments()

class dicee.analyse_experiments.Experiment

    model_name = []

    callbacks = []

    embedding_dim = []

    num_params = []

    num_epochs = []

    batch_size = []

    lr = []

    byte_pair_encoding = []

    aswa = []

    path_dataset_folder = []

    full_storage_path = []

    pq = []

    train_mrr = []

    train_h1 = []

    train_h3 = []

    train_h10 = []

    val_mrr = []

    val_h1 = []
```

```

val_h3 = []

val_h10 = []

test_mrr = []

test_h1 = []

test_h3 = []

test_h10 = []

runtime = []

normalization = []

scoring_technique = []

save_experiment(x)

to_df()

dicee.analyse_experiments.analyse(args)

```

dicee.callbacks

Callbacks for training lifecycle events.

Provides callback classes for various training events including epoch end, model saving, weight averaging, and evaluation.

Attributes

<i>logger</i>

Classes

<i>AccumulateEpochLossCallback</i>	Callback to store epoch losses to a CSV file.
<i>PrintCallback</i>	Callback that prints training start/end times and total run-time.
<i>KGESaveCallback</i>	Callback that periodically saves model checkpoints during training.
<i>PseudoLabellingCallback</i>	Callback that augments the training set with pseudo-labelled triples.
<i>Eval</i>	Abstract base class for KGE training lifecycle callbacks.
<i>KronE</i>	Abstract base class for KGE training lifecycle callbacks.
<i>Perturb</i>	A callback for a three-Level Perturbation
<i>PeriodicEvalCallback</i>	Callback to periodically evaluate the model and optionally save checkpoints during training.
<i>LRScheduler</i>	Callback for managing learning rate scheduling and model snapshots.

Functions

<code>estimate_q(eps)</code>	estimate rate of convergence q from sequence esp
<code>compute_convergence(seq, i)</code>	

Module Contents

`dicee.callbacks.logger`

class `dicee.callbacks.AccumulateEpochLossCallback` (*path: str*)

Bases: `dicee.abstracts.AbstractCallback`

Callback to store epoch losses to a CSV file.

Parameters

path – Directory path where the loss file will be saved.

path

on_fit_end (*trainer, model*) → None

Store epoch loss history to CSV file.

Parameters

- **trainer** – The trainer instance.
- **model** – The model being trained.

class `dicee.callbacks.PrintCallback`

Bases: `dicee.abstracts.AbstractCallback`

Callback that prints training start/end times and total runtime.

start_time

on_fit_start (*trainer, pl_module*)

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

on_fit_end (*trainer, pl_module*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

on_train_batch_end (**args, **kwargs*)

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

on_train_epoch_end (**args, **kwargs*)

Called at the end of each training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.

- **model** (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

class `dicee.callbacks.KGESaveCallback` (*every_x_epoch: int, max_epochs: int, path: str*)

Bases: `dicee.abstracts.AbstractCallback`

Callback that periodically saves model checkpoints during training.

Parameters

- **every_x_epoch** (*int or None*) – Save a checkpoint every *every_x_epoch* epochs. When *None*, the interval defaults to `max(max_epochs // 2, 1)`.
- **max_epochs** (*int*) – Total number of training epochs (used to compute the default interval when *every_x_epoch* is *None*).
- **path** (*str*) – Directory where checkpoint files will be written.

every_x_epoch

max_epochs

epoch_counter = 0

path

on_train_batch_end (**args, **kwargs*)

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

on_fit_start (*trainer, pl_module*)

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

on_train_epoch_end (**args, **kwargs*)

Called at the end of each training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

on_fit_end (**args, **kwargs*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

on_epoch_end (*model, trainer, **kwargs*)

class `dicee.callbacks.PseudoLabellingCallback` (*data_module, kg, batch_size*)

Bases: `dicee.abstracts.AbstractCallback`

Callback that augments the training set with pseudo-labelled triples.

At the end of each epoch the current model scores a batch of unlabelled triples and those with a predicted probability ≥ 0.90 are appended to the training dataset (semi-supervised self-training).

Parameters

- **data_module** (*object*) – Dataset module exposing a `train_set_idx` attribute and a `train_data_loader()` method.
- **kg** (*KG*) – The knowledge graph, providing `num_entities`, `num_relations`, and `unlabelled_set`.
- **batch_size** (*int*) – Number of unlabelled triples to sample and score each epoch.

data_module

kg

num_of_epochs = 0

unlabelled_size

batch_size

create_random_data()

on_epoch_end(*trainer, model*)

`dicee.callbacks.estimate_q`(*eps*)

estimate rate of convergence *q* from sequence *esp*

`dicee.callbacks.compute_convergence`(*seq, i*)

class `dicee.callbacks.Eval`(*path, epoch_ratio: int = None*)

Bases: `dicee.abstracts.AbstractCallback`

Abstract base class for KGE training lifecycle callbacks.

Concrete sub-classes override one or more hook methods to perform custom actions at specific points during training (e.g. weight averaging, periodic evaluation, model checkpointing). All hooks have empty default implementations so sub-classes only need to override the hooks they care about.

Callbacks are registered by passing them to the trainer's *callbacks* list. They are also compatible with PyTorch Lightning trainers because this class extends `lightning.pytorch.callbacks.Callback`.

path

reports = []

epoch_ratio

epoch_counter = 0

on_fit_start(*trainer, model*)

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

on_fit_end(*trainer, model*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

`on_train_epoch_end(trainer, model)`

Called at the end of each training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model being trained. `model.loss_history` contains the per-epoch average losses accumulated so far.

`on_train_batch_end(*args, **kwargs)`

Called after each mini-batch gradient update.

Override to inspect or modify the model at a finer granularity than epoch-level hooks.

class `dicee.callbacks.KronE`

Bases: `dicee.abstracts.AbstractCallback`

Abstract base class for KGE training lifecycle callbacks.

Concrete sub-classes override one or more hook methods to perform custom actions at specific points during training (e.g. weight averaging, periodic evaluation, model checkpointing). All hooks have empty default implementations so sub-classes only need to override the hooks they care about.

Callbacks are registered by passing them to the trainer's `callbacks` list. They are also compatible with PyTorch Lightning trainers because this class extends `lightning.pytorch.callbacks.Callback`.

f = `None`

static `batch_kronecker_product(a, b)`

Kronecker product of matrices a and b with leading batch dimensions. Batch dimensions are broadcast. The number of them must match: type a: `torch.Tensor` :type b: `torch.Tensor` :rtype: `torch.Tensor`

get_kronecker_triple_representation (`indexed_triple: torch.LongTensor`)

Get kronecker embeddings

on_fit_start (`trainer, model`)

Called once before the first training epoch.

Parameters

- **trainer** (`AbstractTrainer` or `pl.Trainer`) – The active trainer instance.
- **model** (`BaseKGE`) – The model about to be trained.

class `dicee.callbacks.Perturb` (`level: str = 'input', ratio: float = 0.0, method: str = None, scaler: float = None, frequency=None`)

Bases: `dicee.abstracts.AbstractCallback`

A callback for a three-Level Perturbation

Input Perturbation: During training an input x is perturbed by randomly replacing its element. In the context of knowledge graph embedding models, x can denote a triple, a tuple of an entity and a relation, or a tuple of two entities. A perturbation means that a component of x is randomly replaced by an entity or a relation.

Parameter Perturbation:

Output Perturbation:

level = `'input'`

ratio = `0.0`

```

method = None

scaler = None

frequency = None

on_train_batch_start (trainer, model, batch, batch_idx)
    Called when the train batch begins.

class dicee.callbacks.PeriodicEvalCallback (experiment_path: str, max_epochs: int,
    eval_every_n_epoch: int = 0, eval_at_epochs: list = None,
    save_model_every_n_epoch: bool = True, n_epochs_eval_model: str = 'val_test')
    Bases: dicee.abstracts.AbstractCallback

    Callback to periodically evaluate the model and optionally save checkpoints during training.

    Evaluates at regular intervals (every N epochs) or at explicitly specified epochs. Stores evaluation reports and model
    checkpoints.

    experiment_dir

    max_epochs

    epoch_counter = 0

    save_model_every_n_epoch = True

    reports

    n_epochs_eval_model = 'val_test'

    default_eval_model = None

    eval_epochs

    on_fit_end (trainer, model)
        Called at the end of training. Saves final evaluation report.

    on_train_epoch_end (trainer, model)
        Called at the end of each training epoch. Performs evaluation and checkpointing if scheduled.

class dicee.callbacks.LRScheduler (adaptive_lr_config: dict, total_epochs: int, experiment_dir: str,
    eta_max: float = 0.1, snapshot_dir: str = 'snapshots')
    Bases: dicee.abstracts.AbstractCallback

    Callback for managing learning rate scheduling and model snapshots.

    Supports cosine annealing (“cca”), MMCCLR (“mmcclr”), and their deferred (warmup) variants: - “deferred_cca”
    - “deferred_mmcclr”

    At the end of each learning rate cycle, the model can optionally be saved as a snapshot.

    total_epochs

    experiment_dir

    snapshot_dir

    batches_per_epoch = None

    total_steps = None

```

cycle_length = None

warmup_steps = None

lr_lambda = None

scheduler = None

step_count = 0

snapshot_loss

on_train_start (*trainer, model*)

Initialize training parameters and LR scheduler at start of training.

on_train_batch_end (*trainer, model, outputs, batch, batch_idx*)

Step the LR scheduler and save model snapshot if needed after each batch.

on_fit_end (*trainer, model*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

dicee.config

Configuration module for DICE embeddings.

Provides the Namespace class with default configuration values for training knowledge graph embedding models.

Classes

<i>Namespace</i>	Extended Namespace with default KGE training configuration.
------------------	---

Module Contents

class `dicee.config.Namespace` (***kwargs*)

Bases: `argparse.Namespace`

Extended Namespace with default KGE training configuration.

Provides sensible defaults for all training parameters while allowing easy customization through command-line arguments or direct assignment.

dataset_dir: `str` | `None` = `None`

The path of a folder containing train.txt, and/or valid.txt and/or test.txt

save_embeddings_as_csv: `bool` = `False`

Embeddings of entities and relations are stored into CSV files to facilitate easy usage.

storage_path: `str` = `'Experiments'`

A directory named with time of execution under `storage_path` that contains related data about embeddings.

path_to_store_single_run: `str` | `None` = `None`

A single directory created that contains related data about embeddings.

path_single_kg = None
 Path of a file corresponding to the input knowledge graph

sparql_endpoint = None
 An endpoint of a triple store.

model: str = 'Keci'
 KGE model

optim: str = 'Adam'
 Optimizer

embedding_dim: int = 64
 Size of continuous vector representation of an entity/relation

num_epochs: int = 150
 Number of pass over the training data

batch_size: int = 1024
 Mini-batch size if it is None, an automatic batch finder technique applied

lr: float = 0.1
 Learning rate

add_noise_rate: float | None = None
 The ratio of added random triples into training dataset

gpus = None
 Number GPUs to be used during training

callbacks
 10}}

Type
 Callbacks, e.g., {"PPE"}

Type
 {"last_percent_to_consider"}

backend: str = 'pandas'
 Backend to read, process, and index input knowledge graph. pandas, polars and rdflib available

separator: str = '\\s+'
 separator for extracting head, relation and tail from a triple

trainer: str = 'torchCPUTrainer'
 'torchCPUTrainer' (CPU/single GPU), 'PL' (PyTorch Lightning multi-GPU), 'torchDDP' (native DDP), 'TP' (Tensor Parallelism - implements 'Multiple Run Ensemble Learning with Low-Dimensional Knowledge Graph Embeddings')

Type
 Trainer for knowledge graph embedding model. Options

scoring_technique: str = 'KvsAll'
 Scoring technique for knowledge graph embedding models

neg_ratio: int = 0
 Negative ratio for a true triple in NegSample training_technique

weight_decay: float = 0.0
 Weight decay for all trainable params

normalization: str = 'None'
 LayerNorm, BatchNorm1d, or None

init_param: str | None = None
 xavier_normal or None

gradient_accumulation_steps: int = 0
 Not tested e

num_folds_for_cv: int = 0
 Number of folds for CV

eval_model: str = 'train_val_test'
 ["None", "train", "train_val", "train_val_test", "test"]

Type
 Evaluate trained model choices

save_model_at_every_epoch: int | None = None
 Not tested

label_smoothing_rate: float = 0.0

num_core: int = 0
 Number of CPUs to be used in the mini-batch loading process

random_seed: int = 0
 Random Seed

log_level: str = 'INFO'
 DEBUG, INFO, WARNING, ERROR, or CRITICAL. Dataset info, timing, and checkpoint messages are logged at INFO; set to WARNING or higher to silence them.

Type
 Logging verbosity

sample_triples_ratio: float | None = None
 Read some triples that are uniformly at random sampled. Ratio being between 0 and 1

read_only_few: int | None = None
 Read only first few triples

pykeen_model_kwargs
 Additional keyword arguments for pykeen models

pl_trainer_kwargs
 Additional keyword arguments for the PyTorch Lightning Trainer

kernel_size: int = 3
 Size of a square kernel in a convolution operation

num_of_output_channels: int = 32
 Number of slices in the generated feature map by convolution.

p: int = 0
 P parameter of Clifford Embeddings

q: int = 1
 Q parameter of Clifford Embeddings

input_dropout_rate: float = 0.0
 Dropout rate on embeddings of input triples

hidden_dropout_rate: float = 0.0
 Dropout rate on hidden representations of input triples

feature_map_dropout_rate: float = 0.0
 Dropout rate on a feature map generated by a convolution operation

byte_pair_encoding: bool = False
 Byte pair encoding

Type
 WIP

adaptive_swa: bool = False
 Adaptive stochastic weight averaging

swa: bool = False
 Stochastic weight averaging

swag: bool = False
 Stochastic weight averaging - Gaussian

ema: bool = False
 Exponential Moving Average

twa: bool = False
 Trainable weight averaging

block_size: int | None = None
 block size of LLM

continual_learning: str | None = None
 Path of a pretrained model size of LLM

auto_batch_finding: bool = False
 A flag for using auto batch finding

eval_every_n_epochs: int = 0
 Evaluate model every n epochs. If 0, no evaluation is applied.

save_every_n_epochs: bool = False
 Save model every n epochs. If True, save model at every epoch.

eval_at_epochs: list | None = None
 List of epoch numbers at which to evaluate the model (e.g., 1 5 10).

n_epochs_eval_model: str = 'val_test'
 Evaluating link prediction performance on data splits while performing periodic evaluation.

adaptive_lr
 “cca”}}

Type
 Adaptive learning rate parameters, e.g., ‘{“scheduler_name”

```

swa_start_epoch: int | None = None
    Epoch at which to start applying stochastic weight averaging.

swa_c_epochs: int = 1
    Number of epochs to average over for SWA, SWAG, EMA, TWA.

__iter__()

```

dicee.dataset_classes

Dataset classes for knowledge graph embedding training.

This package groups the various `torch.utils.data.Dataset` implementations used throughout DICE Embeddings into thematic sub-modules:

- `_bpe` – Byte-pair-encoding related datasets.
- `_negative_sampling` – Negative sampling based datasets.
- `_label_based` – Multi-label / multi-class scoring datasets.
- `_literal` – Literal (numeric) embedding dataset.
- `_factory` – `construct_dataset` / `reload_dataset` helpers.

All public names are re-exported here so that existing `from dicee.dataset_classes import ...` statements continue to work unchanged.

Classes

<i>BPE_NegativeSamplingDataset</i>	Dataset for negative sampling with byte-pair encoded triples.
<i>MultiClassClassificationDataset</i>	Dataset for autoregressive multi-class classification on sub-word units.
<i>MultiLabelDataset</i>	Multi-label dataset for BPE-based KvsAll / AllvsAll training.
<i>AllvsAll</i>	Dataset for AllvsAll training (multi-label, exhaustive).
<i>FSDP1vsSampleDataset</i>	Positive-triple dataset for FSDP 1vsSample training with true-negative sampling.
<i>KvsAll</i>	Dataset for KvsAll training (multi-label).
<i>KvsSampleDataset</i>	Dataset for KvsSample training (dynamic multi-label).
<i>OnevsAllDataset</i>	Dataset for the 1-vs-All training strategy (multi-class).
<i>LiteralDataset</i>	Dataset for loading and processing literal data for Literal Embedding models.
<i>FixedNegSampleDataset</i>	Pre-computed (fixed) negative sampling dataset.
<i>OnevsSample</i>	Dataset for 1-vs-Sample training (dynamic multi-class with negatives).
<i>TriplePredictionDataset</i>	Dataset for triple prediction with on-the-fly negative sampling.

Functions

<i>construct_dataset(→ torch.utils.data.Dataset)</i>	Build the appropriate dataset for the given training configuration.
--	---

continues on next page

Table 11 – continued from previous page

<code>reload_dataset(path, form_of_labelling, ...)</code>	Reload training data from disk and construct a Pytorch dataset.
---	---

Package Contents

class `dicee.dataset_classes.BPE_NegativeSamplingDataset` (*train_set: torch.LongTensor, ordered_shaped_bpe_entities: torch.LongTensor, neg_ratio: int*)

Bases: `torch.utils.data.Dataset`

Dataset for negative sampling with byte-pair encoded triples.

Each sample is a BPE-encoded triple. The custom `collate_fn` constructs negatives by corrupting head or tail entities with random BPE entities.

Parameters

- **train_set** (*torch.LongTensor*) – Integer-encoded triples of shape $(N, 3)$.
- **ordered_shaped_bpe_entities** (*torch.LongTensor*) – All BPE entity representations, ordered by entity index.
- **neg_ratio** (*int*) – Number of negative samples per positive triple.

train_set

ordered_bpe_entities

num_bpe_entities

neg_ratio

num_datapoints

__len__()

__getitem__ (*idx*)

collate_fn (*batch_shaped_bpe_triples: List[Tuple[torch.Tensor, torch.Tensor]]*)

class `dicee.dataset_classes.MultiClassClassificationDataset` (*subword_units: numpy.ndarray, block_size: int = 8*)

Bases: `torch.utils.data.Dataset`

Dataset for autoregressive multi-class classification on sub-word units.

Splits a flat sequence of sub-word token ids into overlapping windows of size `block_size` for next-token prediction.

Parameters

- **subword_units** (*numpy.ndarray*) – 1-D array of sub-word token ids.
- **block_size** (*int, optional*) – Context window length (default 8).

train_data

block_size = 8

num_of_data_points

```
collate_fn = None
```

```
__len__()
```

```
__getitem__(idx)
```

```
class dicee.dataset_classes.MultiLabelDataset (train_set: torch.LongTensor,  
        train_indices_target: torch.LongTensor, target_dim: int,  
        torch_ordered_shaped_bpe_entities: torch.LongTensor)
```

Bases: torch.utils.data.Dataset

Multi-label dataset for BPE-based KvsAll / AllvsAll training.

Each sample is a BPE-encoded (head, relation) pair together with a binary multi-label target vector over all entities.

Parameters

- **train_set** (*torch.LongTensor*) – BPE-encoded input pairs of shape (N, 2, token_length).
- **train_indices_target** (*torch.LongTensor*) – Per-sample lists of positive target entity indices.
- **target_dim** (*int*) – Dimensionality of the target vector (number of entities).
- **torch_ordered_shaped_bpe_entities** (*torch.LongTensor*) – Ordered BPE entity representations.

```
train_set
```

```
train_indices_target
```

```
target_dim
```

```
num_datapoints
```

```
torch_ordered_shaped_bpe_entities
```

```
collate_fn = None
```

```
__len__()
```

```
__getitem__(idx)
```

```
dicee.dataset_classes.construct_dataset (* (Keyword-only parameters separator (PEP 3102)),  
        train_set: numpy.ndarray | list, valid_set=None, test_set=None, ordered_bpe_entities=None,  
        train_target_indices=None, target_dim: int = None, entity_to_idx: dict, relation_to_idx: dict,  
        form_of_labelling: str, scoring_technique: str, neg_ratio: int, label_smoothing_rate: float,  
        byte_pair_encoding=None, block_size: int = None, seed: int = None, sort_train_set: bool = True)  
    → torch.utils.data.Dataset
```

Build the appropriate dataset for the given training configuration.

Parameters

- **train_set** (*numpy.ndarray* or *list*) – Raw integer-indexed triples.
- **entity_to_idx** (*dict*) – Name → index mappings.
- **relation_to_idx** (*dict*) – Name → index mappings.
- **form_of_labelling** (*str*) – 'EntityPrediction' or 'RelationPrediction'.

- **scoring_technique** (*str*) – One of 'NegSample', 'FixedNegSample', '1vsAll', '1vsSample', 'KvsAll', 'AllvsAll', 'KvsSample'.
- **neg_ratio** (*int*) – Negative sample ratio.
- **label_smoothing_rate** (*float*) – Label smoothing coefficient.

Return type

`torch.utils.data.Dataset`

`dicee.dataset_classes.reload_dataset` (*path: str, form_of_labelling, scoring_technique, neg_ratio, label_smoothing_rate*)

Reload training data from disk and construct a Pytorch dataset.

class `dicee.dataset_classes.AllvsAll` (*train_set_idx: numpy.ndarray, entity_idxes, relation_idxes, label_smoothing_rate=0.0*)

Bases: `torch.utils.data.Dataset`

Dataset for AllvsAll training (multi-label, exhaustive).

Extends the KvsAll idea: every *possible* (*entity, relation*) combination is included — not just those observed in the KG. Pairs without any known tail entities receive an all-zeros label vector.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **entity_idxes** (*dict*) – Entity-name → index mapping.
- **relation_idxes** (*dict*) – Relation-name → index mapping.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data = None

train_target = None

label_smoothing_rate

collate_fn = None

target_dim

__len__ ()

__getitem__ (*idx*)

class `dicee.dataset_classes.FSDP1vsSampleDataset` (*train_set_idx: numpy.ndarray, entity_idxes, relation_idxes, form, neg_ratio=None, label_smoothing_rate: float = 0.0*)

Bases: `torch.utils.data.Dataset`

Positive-triple dataset for FSDP 1vsSample training with true-negative sampling.

Each dataset item is a single positive triple (*h, r, t*). The *collate_fn* builds the full (*source, target_idx, labels*) batch in a *DataLoader* worker, sampling true negatives via index remapping so they never coincide with any known positive tail for that (*h, r*) pair.

Fixed batch width follows KvsSample:

`max_num_of_classes = max_positives_per_pair + neg_ratio`

Each sample contributes 1 positive and (`max_num_of_classes - 1`) negatives.

```

train_data

num_entities

num_relations

neg_ratio = None

label_smoothing_rate = 0.0

max_num_of_classes

num_negatives

collate_fn

__len__()

__getitem__(idx)

class dicee.dataset_classes.KvsAll (train_set_idx: numpy.ndarray, entity_idx, relation_idx, form,
    store=None, label_smoothing_rate: float = 0.0)
    Bases: torch.utils.data.Dataset
    Dataset for KvsAll training (multi-label).
     $D := \{(x, y)_i\}_{i=1}^N$  where
        •  $x = (h, r)$  is a unique (entity, relation) pair observed in the KG,
        •  $y \in [0, 1]^{|E|}$  is a multi-label vector with  $y_j = 1$  iff  $(h, r, e_j) \in KG$ .

    Parameters
        • train_set_idx (numpy.ndarray) – (N, 3) integer-indexed triples.
        • entity_idx (dict) – Entity-name  $\rightarrow$  index mapping.
        • relation_idx (dict) – Relation-name  $\rightarrow$  index mapping.
        • form (str) – 'EntityPrediction' or 'RelationPrediction'.
        • label_smoothing_rate (float, optional) – Label smoothing coefficient (default 0.0).

train_data = None

train_target = None

label_smoothing_rate

collate_fn = None

__len__()

__getitem__(idx)

```

```
class dicee.dataset_classes.KvsSampleDataset (train_set_idx: numpy.ndarray, entity_idxes,  

relation_idxes, form, store=None, neg_ratio=None, label_smoothing_rate: float = 0.0)
```

Bases: torch.utils.data.Dataset

Dataset for KvsSample training (dynamic multi-label).

Like `KvsAll` but sub-samples the target vector at each access to keep mini-batch sizes manageable when the entity set is large.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **entity_idxes** (*dict*) – Entity-name → index mapping.
- **relation_idxes** (*dict*) – Relation-name → index mapping.
- **form** (*str*) – 'EntityPrediction'.
- **neg_ratio** (*int*) – Number of negative samples per positive target.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data = None

train_target = None

neg_ratio = None

num_entities

label_smoothing_rate

collate_fn = None

max_num_of_classes

__len__()

__getitem__ (*idx*)

```
class dicee.dataset_classes.OnevsAllDataset (train_set_idx: numpy.ndarray, entity_idxes)
```

Bases: torch.utils.data.Dataset

Dataset for the 1-vs-All training strategy (multi-class).

Each sample is a (head, relation) pair with a one-hot target vector whose single active position corresponds to the true tail entity.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **entity_idxes** (*dict*) – Entity-name → index mapping (used to determine the target dimension).

train_data

target_dim

collate_fn = None

__len__()

`__getitem__` (*idx*)

class dicee.dataset_classes.**LiteralDataset** (*file_path: str, ent_idx: dict = None, normalization_type: str = 'z-norm', sampling_ratio: float = None, loader_backend: str = 'pandas'*)

Bases: `torch.utils.data.Dataset`

Dataset for loading and processing literal data for Literal Embedding models.

Handles loading, normalization, and preparation of (entity, attribute, value) triples. Supports z-score and min-max normalization as well as optional sub-sampling for ablation studies.

Parameters

- **file_path** (*str*) – Path to the training data file (CSV / TSV / RDF).
- **ent_idx** (*dict*) – Entity-name → index mapping.
- **normalization_type** (*str, optional*) – 'z-norm', 'min-max', or None (default 'z-norm').
- **sampling_ratio** (*float or None, optional*) – Fraction of the training set to keep (default None = use all).
- **loader_backend** (*str, optional*) – 'pandas' or 'rdflib' (default 'pandas').

`train_file_path`

`loader_backend = 'pandas'`

`normalization_type = 'z-norm'`

`normalization_params`

`sampling_ratio = None`

`entity_to_idx = None`

`num_entities`

`__getitem__` (*index*)

`__len__` ()

static `load_and_validate_literal_data` (*file_path: str = None, loader_backend: str = 'pandas'*)
→ `pandas.DataFrame`

Load and validate a literal data file.

Parameters

- **file_path** (*str*) – Path to the data file.
- **loader_backend** (*str*) – 'pandas' or 'rdflib'.

Returns

Three-column `DataFrame` with columns head, attribute, value.

Return type

`pandas.DataFrame`

static `denormalize` (*preds_norm, attributes, normalization_params*) → `numpy.ndarray`

Reverse the normalization applied during training.

Parameters

- **preds_norm** (*numpy.ndarray*) – Normalized predictions.
- **attributes** (*list*) – Attribute names corresponding to each prediction.
- **normalization_params** (*dict*) – Parameters stored during training.

Returns

Denormalized predictions.

Return type

`numpy.ndarray`

```
class dicee.dataset_classes.FixedNegSampleDataset (train_set: numpy.ndarray, num_entities: int,  
num_relations: int, neg_sample_ratio: int = 1, label_smoothing_rate: float = 0.0, seed: int = None)
```

Bases: `torch.utils.data.Dataset`

Pre-computed (fixed) negative sampling dataset.

At construction time every positive triple is paired with one random negative (head- or tail-corrupted) using vectorized operations for efficiency. The pairs are stored so that `__getitem__` is a simple lookup.

This is useful when you want deterministic negatives across epochs (e.g., for reproducibility or debugging).

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int, optional*) – Number of negative samples per positive triple (default 1).
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

```
neg_sample_ratio = 1
```

```
num_entities
```

```
num_relations
```

```
label_smoothing_rate = 0.0
```

```
collate_fn = None
```

```
seed = None
```

```
train_triples
```

```
length
```

```
__len__() → int
```

```
__getitem__(idx: int) → Tuple[torch.Tensor, torch.Tensor]
```

```
class dicee.dataset_classes.OnevsSample (train_set: numpy.ndarray, num_entities: int,  
num_relations: int, neg_sample_ratio: int = None, label_smoothing_rate: float = 0.0)
```

Bases: `torch.utils.data.Dataset`

Dataset for 1-vs-Sample training (dynamic multi-class with negatives).

For every positive triple (h, r, t) the dataset draws `neg_sample_ratio` random entities as negatives and returns a label vector that marks the true tail and the negatives.

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int*) – Number of negative samples per positive.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data

num_entities

num_relations

neg_sample_ratio = None

label_smoothing_rate

collate_fn = None

__len__()

__getitem__(*idx*)

```
class dicee.dataset_classes.TriplePredictionDataset (train_set: numpy.ndarray,  
    num_entities: int, num_relations: int, neg_sample_ratio: int = 1, label_smoothing_rate: float = 0.0,  
    seed: int = None, sort_train_set: bool = True)
```

Bases: `torch.utils.data.Dataset`

Dataset for triple prediction with on-the-fly negative sampling.

Each item is a single positive triple; the custom `collate_fn` generates a batch of mixed positive and negative triples.

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int, optional*) – Number of negative samples per positive triple (default 1).
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

label_smoothing_rate

neg_sample_ratio

seed = None

length

num_entities

num_relations


```

__len__()

__getitem__(idx)

collate_fn(batch: List[torch.Tensor])

```

dicee.eval_static_funcs

Static evaluation functions for KGE models.

This module provides backward compatibility by re-exporting from the new `dicee.evaluation` module.

Deprecated since version Use: `dicee.evaluation` submodules instead. This module will be removed in a future version.

Functions

<code>evaluate_ensemble_link_prediction_performance</code>	Evaluate link prediction performance of an ensemble of KGE models.
<code>evaluate_link_prediction_performance(→ Dict[str, float])</code>	Evaluate link prediction performance with head and tail prediction.
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with BPE encoding (head and tail).
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with BPE encoding and reciprocals.
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with reciprocal relations.
<code>evaluate_lp_bpe_k_vs_all(→ Dict[str, float])</code>	Evaluate BPE link prediction with KvsAll scoring.
<code>evaluate_literal_prediction(→ Optional[pandas.DataFrame])</code>	Evaluate trained literal prediction model on a test file.

Module Contents

`dicee.eval_static_funcs.evaluate_ensemble_link_prediction_performance` (*models: List, triples, er_vocab: Dict[Tuple, List], weights: List[float] | None = None, batch_size: int = 512, weighted_averaging: bool = True, normalize_scores: bool = True*) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is `True` but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

`dicee.eval_static_funcs.evaluate_link_prediction_performance(model, triples, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

`dicee.eval_static_funcs.evaluate_link_prediction_performance_with_bpe(model, within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

`dicee.eval_static_funcs.evaluate_link_prediction_performance_with_bpe_reciprocals(model, within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.eval_static_funcs.evaluate_link_prediction_performance_with_reciprocals(  
    model, triples, er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.eval_static_funcs.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]],  
    er_vocab: Dict | None = None, batch_size: int | None = None,  
    func_triple_to_bpe_representation: Callable | None = None,  
    str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

```
dicee.eval_static_funcs.evaluate_literal_prediction(kge_model, eval_file_path: str = None,  
    store_lit_preds: bool = True, eval_literals: bool = True, loader_backend: str = 'pandas',  
    return_attr_error_metrics: bool = False) → pandas.DataFrame | None
```

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is True, otherwise None.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

dicee.evaluation

Evaluation module for knowledge graph embedding models.

This module provides comprehensive evaluation capabilities for KGE models, including link prediction, literal prediction, and ensemble evaluation.

Modules:

link_prediction: Functions for evaluating link prediction performance
literal_prediction: Functions for evaluating literal/attribute prediction
ensemble: Functions for ensemble model evaluation
evaluator: Main Evaluator class for integrated evaluation
utils: Shared utility functions for evaluation
_filtering: Internal helpers for filtered ranking (not exported)

Example

```
>>> from dicee.evaluation import Evaluator
>>> from dicee.evaluation.link_prediction import evaluate_link_prediction_performance
>>> from dicee.evaluation.ensemble import evaluate_ensemble_link_prediction_
↪ performance
```

Submodules

`dicee.evaluation.ensemble`

Ensemble evaluation functions.

This module provides functions for evaluating ensemble models, including weighted averaging and score normalization.

Functions

<code>evaluate_ensemble_link_prediction_performance</code>	Evaluate link prediction performance of an ensemble of KGE models.
--	--

Module Contents

`dicee.evaluation.ensemble.evaluate_ensemble_link_prediction_performance` (*models*: List, *triples*, *er_vocab*: Dict[Tuple, List], *weights*: List[float] | None = None, *batch_size*: int = 512, *weighted_averaging*: bool = True, *normalize_scores*: bool = True) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is True but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

dicee.evaluation.evaluator

Main Evaluator class for KGE model evaluation.

This module provides the Evaluator class which orchestrates evaluation of knowledge graph embedding models across different datasets and scoring techniques.

Attributes

<code>logger</code> <code>VALID_SCORING_TECHNIQUES</code>
--

Classes

<code>Evaluator</code>	Evaluator class for KGE models in various downstream tasks.
------------------------	---

Module Contents

`dicee.evaluation.evaluator.logger`

`dicee.evaluation.evaluator.VALID_SCORING_TECHNIQUES`

class `dicee.evaluation.evaluator.Evaluator` (*args*, *is_continual_training*: *bool* = *False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset*, *trained_model*, *form_of_labelling*: str, *during_training*: bool = False) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set*, *valid_set*=None, *test_set*=None, *trained_model*) → None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set*=None, *valid_set*=None, *test_set*=None, *ordered_bpe_entities*, *trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte(**, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate ByteE model with generation.

eval_with_bpe_vs_all(**, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all(**, train_set, valid_set=None, test_set=None, trained_model, form_of_labelling*) → None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all(*model, triple_idx, info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp_with_byte(*model, triples: List[List[str]], info: str | None = None*) → Dict[str, float]

Evaluate ByteE model with text generation.

Parameters

- **model** – ByteE model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all(*model, triples: List[List[str]], info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp(*model, triple_idx, info: str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: str) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: numpy.ndarray, *form_of_labelling*: str)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

dicee.evaluation.link_prediction

Link prediction evaluation functions.

This module provides various functions for evaluating link prediction performance of knowledge graph embedding models.

Attributes

logger

Functions

<i>evaluate_link_prediction_performance</i> (→ Dict[str, float])	Evaluate link prediction performance with head and tail prediction.
<i>evaluate_link_prediction_performance_with_</i>	Evaluate link prediction with reciprocal relations.
<i>evaluate_link_prediction_performance_with_</i>	Evaluate link prediction with BPE encoding and reciprocals.
<i>evaluate_link_prediction_performance_with_</i>	Evaluate link prediction with BPE encoding (head and tail).
<i>evaluate_lp</i> (→ Dict[str, float])	Evaluate link prediction with batched processing.

continues on next page

Table 17 – continued from previous page

<code>evaluate_bpe_lp(→ Dict[str, float])</code>	Evaluate link prediction with BPE-encoded entities.
<code>evaluate_lp_bpe_k_vs_all(→ Dict[str, float])</code>	Evaluate BPE link prediction with KvsAll scoring.

Module Contents

`dicee.evaluation.link_prediction.logger`

`dicee.evaluation.link_prediction.evaluate_link_prediction_performance(model, triples, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.link_prediction.`

`evaluate_link_prediction_performance_with_reciprocals(model, triples, er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.link_prediction.`

`evaluate_link_prediction_performance_with_bpe_reciprocals(model, within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_link_prediction_performance_with_bpe(
    model, within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List],
    re_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_lp(model, triple_idx, num_entities: int,
    er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List], info: str = 'Eval Starts',
    batch_size: int = 128, chunk_size: int = 1000) → Dict[str, float]
```

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.
- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_bpe_lp(model, triple_idx: List[Tuple],
    all_bpe_shaped_entities, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List],
    info: str = 'Eval Starts') → Dict[str, float]
```

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]],
er_vocab: Dict | None = None, batch_size: int | None = None,
func_triple_to_bpe_representation: Callable | None = None,
str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

dicee.evaluation.literal_prediction

Literal prediction evaluation functions.

This module provides functions for evaluating literal/attribute prediction performance of knowledge graph embedding models.

Attributes

logger

Functions

<code>evaluate_literal_prediction(→ Optional[pandas.DataFrame])</code>	Evaluate trained literal prediction model on a test file.
--	---

Module Contents

`dicee.evaluation.literal_prediction.logger`

```
dicee.evaluation.literal_prediction.evaluate_literal_prediction(kge_model,
eval_file_path: str = None, store_lit_preds: bool = True, eval_literals: bool = True,
loader_backend: str = 'pandas', return_attr_error_metrics: bool = False)
→ pandas.DataFrame | None
```

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is True, otherwise None.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

dicee.evaluation.utils

Utility functions for evaluation module.

This module contains shared helper functions used across different evaluation components.

Attributes

`DEFAULT_HITS_RANGE`
`ALL_HITS_RANGE`

Functions

<code>make_iterable_verbose(→ Iterable)</code>	Wrap an iterable with tqdm progress bar if verbose is True.
<code>compute_metrics_from_ranks(→ Dict[str, float])</code>	Compute standard link prediction metrics from ranks.
<code>compute_metrics_from_ranks_simple(→ Dict[str, float])</code>	Compute link prediction metrics without scaling factor.
<code>update_hits(→ None)</code>	Update hits dictionary based on rank.

continues on next page

Table 21 – continued from previous page

<code>create_hits_dict(→ Dict[int, List[float]])</code>	Create an initialized hits dictionary.
<code>efficient_zero_grad(→ None)</code>	Efficiently zero gradients using <code>parameter.grad = None</code> .

Module Contents

`dicee.evaluation.utils.DEFAULT_HITS_RANGE: List[int] = [1, 3, 10]`

`dicee.evaluation.utils.ALL_HITS_RANGE: List[int]`

`dicee.evaluation.utils.make_iterable_verbose (iterable_object: Iterable, verbose: bool, desc: str = 'Default', position: int | None = None, leave: bool = True) → Iterable`

Wrap an iterable with tqdm progress bar if verbose is True.

Parameters

- **iterable_object** – The iterable to potentially wrap.
- **verbose** – Whether to show progress bar.
- **desc** – Description for the progress bar.
- **position** – Position of the progress bar.
- **leave** – Whether to leave the progress bar after completion.

Returns

The original iterable or a tqdm-wrapped version.

`dicee.evaluation.utils.compute_metrics_from_ranks (ranks: List[int], num_triples: int, hits_dict: Dict[int, List[float]], scale_factor: int = 1) → Dict[str, float]`

Compute standard link prediction metrics from ranks.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.
- **scale_factor** – Factor to scale the denominator (e.g., 2 for head+tail).

Returns

Dictionary containing H@1, H@3, [H@10](#), and MRR metrics.

`dicee.evaluation.utils.compute_metrics_from_ranks_simple (ranks: List[int], num_triples: int, hits_dict: Dict[int, List[float]]) → Dict[str, float]`

Compute link prediction metrics without scaling factor.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.

Returns

Dictionary containing H@1, H@3, [H@10](#), and MRR metrics.

`dicee.evaluation.utils.update_hits(hits: Dict[int, List[float]], rank: int, hits_range: List[int] | None = None) → None`

Update hits dictionary based on rank.

Parameters

- **hits** – Dictionary to update in-place.
- **rank** – The rank to check against hit levels.
- **hits_range** – List of hit levels to check (default: ALL_HITS_RANGE).

`dicee.evaluation.utils.create_hits_dict(hits_range: List[int] | None = None) → Dict[int, List[float]]`

Create an initialized hits dictionary.

Parameters

hits_range – List of hit levels to initialize (default: ALL_HITS_RANGE).

Returns

Dictionary with empty lists for each hit level.

`dicee.evaluation.utils.efficient_zero_grad(model) → None`

Efficiently zero gradients using `parameter.grad = None`.

This is more efficient than `optimizer.zero_grad()` as it avoids memory operations.

See: https://pytorch.org/tutorials/recipes/recipes/tuning_guide.html

Parameters

model – PyTorch model to zero gradients for.

Classes

<i>Evaluator</i>	Evaluator class for KGE models in various downstream tasks.
------------------	---

Functions

<i>evaluate_ensemble_link_prediction_performance</i>	Evaluate link prediction performance of an ensemble of KGE models.
<i>evaluate_bpe_lp</i> (→ Dict[str, float])	Evaluate link prediction with BPE-encoded entities.
<i>evaluate_link_prediction_performance</i> (→ Dict[str, float])	Evaluate link prediction performance with head and tail prediction.
<i>evaluate_link_prediction_performance_with_j</i>	Evaluate link prediction with BPE encoding (head and tail).
<i>evaluate_link_prediction_performance_with_j</i>	Evaluate link prediction with BPE encoding and reciprocals.
<i>evaluate_link_prediction_performance_with_</i>	Evaluate link prediction with reciprocal relations.
<i>evaluate_lp</i> (→ Dict[str, float])	Evaluate link prediction with batched processing.
<i>evaluate_lp_bpe_k_vs_all</i> (→ Dict[str, float])	Evaluate BPE link prediction with KvsAll scoring.
<i>evaluate_literal_prediction</i> (→ Optional[pandas.DataFrame])	Evaluate trained literal prediction model on a test file.

Package Contents

`dicee.evaluation.evaluate_ensemble_link_prediction_performance` (*models: List, triples, er_vocab: Dict[Tuple, List], weights: List[float] | None = None, batch_size: int = 512, weighted_averaging: bool = True, normalize_scores: bool = True*) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is `True`. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If `True`, use weighted averaging of predictions. If `False`, use simple mean.
- **normalize_scores** – If `True`, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is `True` but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

class `dicee.evaluation.Evaluator` (*args, is_continual_training: bool = False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset*, *trained_model*, *form_of_labelling*: str, *during_training*: bool = False) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).

- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, train_set, valid_set=None, test_set=None, trained_model) → None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, train_set=None, valid_set=None, test_set=None, ordered_bpe_entities, trained_model) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling) → None

Evaluate Byte model with generation.

eval_with_bpe_vs_all (*, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, train_set, valid_set=None, test_set=None, trained_model, form_of_labelling) → None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all (model, triple_idx, info: str | None = None, form_of_labelling: str | None = None) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – 'EntityPrediction' or 'RelationPrediction'.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

evaluate_lp_with_byte (model, triples: List[List[str]], info: str | None = None) → Dict[str, float]

Evaluate Byte model with text generation.

Parameters

- **model** – Byte model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (model, triples: List[List[str]], info: str | None = None, form_of_labelling: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (*model*, *triple_idx*, *info*: *str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: *str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: *numpy.ndarray*, *form_of_labelling*: *str*)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

`dicee.evaluation.evaluate_bpe_lp` (*model*, *triple_idx*: *List[Tuple]*, *all_bpe_shaped_entities*,
er_vocab: *Dict[Tuple, List]*, *re_vocab*: *Dict[Tuple, List]*, *info*: *str* = 'Eval Starts') → Dict[str, float]

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.

- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance(model, triples,
er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance_with_bpe(model,
within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List],
re_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance_with_bpe_reciprocals(model,
within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.evaluate_link_prediction_performance_with_reciprocals(model, triples, er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.evaluate_lp(model, triple_idx, num_entities: int, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List], info: str = 'Eval Starts', batch_size: int = 128, chunk_size: int = 1000) → Dict[str, float]`

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.
- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]], er_vocab: Dict | None = None, batch_size: int | None = None, func_triple_to_bpe_representation: Callable | None = None, str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]`

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If `batch_size` is not provided.

`dicee.evaluation.evaluate_literal_prediction(kge_model, eval_file_path: str = None, store_lit_preds: bool = True, eval_literals: bool = True, loader_backend: str = 'pandas', return_attr_error_metrics: bool = False) → pandas.DataFrame | None`

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is True, otherwise None.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

`dicee.evaluator`

Evaluator module for knowledge graph embedding models.

This module provides backward compatibility by re-exporting from the new `dicee.evaluation` module.

Deprecated since version Use: `dicee.evaluation.Evaluator` instead. This module will be removed in a future version.

Classes

<i>Evaluator</i>	Evaluator class for KGE models in various downstream tasks.
------------------	---

Module Contents

class `dicee.evaluator.Evaluator` (*args*, *is_continual_training*: *bool* = *False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset, trained_model, form_of_labelling: str, during_training: bool = False*) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set, valid_set=None, test_set=None, trained_model*)
→ None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set=None, valid_set=None, test_set=None, ordered_bpe_entities, trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, *raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate Byte model with generation.

eval_with_bpe_vs_all (*, *raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, *train_set, valid_set=None, test_set=None, trained_model, form_of_labelling*)
→ None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all (*model, triple_idx, info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.

- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp_with_byte (*model, triples: List[List[str]], info: str | None = None*) → Dict[str, float]

Evaluate ByteE model with text generation.

Parameters

- **model** – ByteE model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (*model, triples: List[List[str]], info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (*model, triple_idx, info: str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model, form_of_labelling: str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset, trained_model, triple_idx: numpy.ndarray, form_of_labelling: str*)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

dicee.executer

Executor module for training, retraining and evaluating KGE models.

This module provides the Execute and ContinuousExecute classes for managing the full lifecycle of knowledge graph embedding model training.

Attributes

<i>logger</i>

Classes

<i>Execute</i>	Executor class for training, retraining and evaluating KGE models.
<i>ContinuousExecute</i>	A subclass of Execute Class for retraining

Module Contents

`dicee.executer.logger`

class `dicee.executer.Execute` (*args, continuous_training: bool = False*)

Executor class for training, retraining and evaluating KGE models.

Handles the complete workflow: 1. Loading & Preprocessing & Serializing input data 2. Training & Validation & Testing 3. Storing all necessary information

args

Processed input arguments.

distributed

Whether distributed training is enabled.

rank

Process rank in distributed training.

world_size
Total number of processes.

local_rank
Local GPU rank.

trainer
Training handler instance.

trained_model
The trained model after training completes.

knowledge_graph
The loaded knowledge graph.

report
Dictionary storing training metrics and results.

evaluator
Model evaluation handler.

distributed

rank

world_size

local_rank

args

is_continual_training = False

trainer: *dicее.trainer.DICE_Trainer* | None = None

trained_model = None

knowledge_graph: *dicее.knowledge_graph.KG* | None = None

report: Dict

evaluator: *dicее.evaluator.Evaluator* | None = None

start_time: float | None = None

is_local_rank_zero() → bool

is_global_rank_zero() → bool

cleanup()

setup_executor() → None
Set up storage directories for the experiment.

Creates or reuses experiment directories based on configuration. Saves the configuration to a JSON file.

create_and_store_kg() → None
Create knowledge graph and store as memory-mapped file.

Only executed on global rank 0 in distributed training, so exactly one process writes the shared memmap/vocab files regardless of how many nodes are involved. Skips if memmap already exists.

load_from_memmap () → None

Load knowledge graph from memory-mapped file.

save_trained_model () → None

Save a knowledge graph embedding model

- (1) Send model to eval mode and cpu.
- (2) Store the memory footprint of the model.
- (3) Save the model into disk.
- (4) Update the stats of KG again ?

Parameter

rtype

None

end (*form_of_labelling: str*) → dict

End training

- (1) Store trained model.
- (2) Report runtimes.
- (3) Eval model if required.

Parameter

rtype

A dict containing information about the training and/or evaluation

write_report () → None

Report training related information in a report.json file

start () → dict

Start training

(1) Loading the Data # (2) Create an evaluator object. # (3) Create a trainer object. # (4) Start the training

Parameter

rtype

A dict containing information about the training and/or evaluation

class `dicee.executer.ContinuousExecute` (*args*)

Bases: `Execute`

A subclass of Execute Class for retraining

- (1) Loading & Preprocessing & Serializing input data.
- (2) Training & Validation & Testing
- (3) Storing all necessary info

During the continual learning we can only modify * **num_epochs** * parameter. Trained model stored in the same folder as the seed model for the training. Trained model is noted with the current time.

continual_start() → dict

Start Continual Training

- (1) Initialize training.
- (2) Start continual training.
- (3) Save trained model.

Parameter

rtype

A dict containing information about the training and/or evaluation

dicee.knowledge_graph

Knowledge Graph module for data loading and preprocessing.

Provides the KG class for handling knowledge graph data including loading, preprocessing, and indexing operations.

Classes

<i>KG</i>	Knowledge Graph container and processor.
-----------	--

Module Contents

```
class dicee.knowledge_graph.KG (dataset_dir: str | None = None, byte_pair_encoding: bool = False,
    padding: bool = False, add_noise_rate: float | None = None, sparql_endpoint: str | None = None,
    path_single_kg: str | None = None, path_for_deserialization: str | None = None,
    add_reciprocal: bool | None = None, eval_model: str | None = None,
    read_only_few: int | None = None, sample_triples_ratio: float | None = None,
    path_for_serialization: str | None = None, entity_to_idx: Dict | None = None,
    relation_to_idx: Dict | None = None, backend: str | None = None,
    training_technique: str | None = None, separator: str | None = None)
```

Knowledge Graph container and processor.

Handles loading, preprocessing, and indexing of knowledge graph data from various sources including files, SPARQL endpoints, and serialized formats.

dataset_dir

Path to directory containing train/valid/test files.

num_entities

Total number of unique entities.

num_relations

Total number of unique relations.

train_set

Indexed training triples as numpy array.

valid_set

Indexed validation triples (optional).

test_set

Indexed test triples (optional).

```

entity_to_idx
    Mapping from entity strings to indices.
relation_to_idx
    Mapping from relation strings to indices.
dataset_dir = None
sparql_endpoint = None
path_single_kg = None
byte_pair_encoding = False
ordered_shaped_bpe_tokens = None
add_noise_rate = None
num_entities: int | None = None
num_relations: int | None = None
path_for_deserialization = None
add_reciprocal = None
eval_model = None
read_only_few = None
sample_triples_ratio = None
path_for_serialization = None
entity_to_idx = None
relation_to_idx = None
backend = 'pandas'
training_technique = None
separator = None
raw_train_set = None
raw_valid_set = None
raw_test_set = None
train_set = None
valid_set = None
test_set = None
idx_entity_to_bpe_shaped: Dict
enc
num_tokens

```

```

num_bpe_entities: int | None = None

padding = False

dummy_id

max_length_subword_tokens: int | None = None

train_set_target = None

target_dim: int | None = None

train_target_indices = None

ordered_bpe_entities = None

description_of_input = None

describe() → None
    Generate a description string of the dataset statistics.

property_entities_str: List[str]
    Get list of all entity strings.

property_relations_str: List[str]
    Get list of all relation strings.

exists(h: str, r: str, t: str) → bool
    Check if a triple exists in the training set.

Parameters
    • h – Head entity string.
    • r – Relation string.
    • t – Tail entity string.

Returns
    True if the triple exists, False otherwise.

__iter__() → Iterator[Tuple[str, str, str]]
    Iterate over training triples as string tuples.

__len__() → int
    Return number of triples in the raw training set.

func_triple_to_bpe_representation(triple: List[str])

```

dicee.knowledge_graph_embeddings

Attributes

logger

Classes

<i>KGE</i>	Knowledge Graph Embedding Class for interactive usage of pre-trained models
------------	---

Module Contents

`dicee.knowledge_graph_embeddings.logger`

class `dicee.knowledge_graph_embeddings.KGE` (*path=None, url=None, construct_ensemble=False, model_name=None*)

Bases: `dicee.abstracts.BaseInteractiveKGE`, `dicee.abstracts.InteractiveQueryDecomposition`, `dicee.abstracts.BaseInteractiveTrainKGE`

Knowledge Graph Embedding Class for interactive usage of pre-trained models

`__str__` ()

`to` (*device: str*) $\rightarrow$ None

`get_transductive_entity_embeddings` (*indices: torch.LongTensor | List[str], as_pytorch=False, as_numpy=False, as_list=True*) $\rightarrow$ `torch.FloatTensor | numpy.ndarray | List[float]`

`create_vector_database` (*collection_name: str, distance: str, location: str = 'localhost', port: int = 6333*)

`generate` (*h="", r=""*)

`eval_lp_performance` (*dataset=List[Tuple[str, str, str]], filtered=True*)

`predict_missing_head_entity` (*relation: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False*) $\rightarrow$ Tuple

Given a relation and a tail entity, return top k ranked head entity.

$\text{argmax}_{\{e \in E\}} f(e, r, t)$, where $r \in R, t \in E$.

Parameter

relation: Union[List[str], str]

String representation of selected relations.

tail_entity: Union[List[str], str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_relations (*head_entity: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False*) $\rightarrow$ Tuple

Given a head entity and a tail entity, return top k ranked relations.

$\text{argmax}_{\{r \in R\}} f(h, r, t)$, where $h, t \in E$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_tail_entity (*head_entity: List[str] | str, relation: List[str] | str,*
within: List[str] = None, batch_size=2, topk=1, return_indices=False) $\rightarrow$ torch.FloatTensor

Given a head entity and a relation, return top k ranked entities

$\text{argmax}_{\{e \in E\}} f(h, r, e)$, where h in E and r in R .

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

Returns: Tuple

scores

predict (*, *h: List[str] | str | None = None, r: List[str] | str | None = None, t: List[str] | str | None = None,*
within: List[str] | None = None, logits: bool = True) $\rightarrow$ torch.FloatTensor

Predict scores for triples or missing triple elements.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.
- **t** – Tail entity/entities. None to predict tails.
- **within** – Optional list of entities to restrict predictions to.
- **logits** – If True, return raw scores. If False, return sigmoid scores (0-1).

Returns

- Single triple (h, r, t): scalar score
- Missing element: vector of all possible scores

Return type

torch.FloatTensor of scores. Shape depends on the query type

Raises

- **TypeError** – If inputs are not strings or lists of strings.
- **ValueError** – If a required argument for the query type is missing.

Examples

```
>>> # Score a specific triple
>>> model.predict(h="Mongolia", r="isLocatedIn", t="Asia", logits=False)
tensor(0.9523)
```

```
>>> # Get scores for all possible tail entities
>>> model.predict(h="Mongolia", r="isLocatedIn", t=None)
tensor([0.21, 0.95, 0.03, ...]) # One score per entity
```

predict_topk (*, *h*: str | List[str] | None = None, *r*: str | List[str] | None = None, *t*: str | List[str] | None = None, *topk*: int = 10, *within*: List[str] | None = None, *batch_size*: int = 1024) → List[Tuple[str, float]] | List[List[Tuple[str, float]]]

Predict top-k missing items in a given triple pattern.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.
- **t** – Tail entity/entities. None to predict tails.
- **topk** – Number of top predictions to return.
- **within** – Optional list of entities to restrict predictions to.
- **batch_size** – Batch size for processing multiple queries.

Returns

List[(item, score), ...] of length topk. For batch query: List of such lists, one per query.

Return type

For single query

Raises

- **TypeError** – If h, r, or t is not a str or list of str.
- **ValueError** – If the required arguments for a query type are None.

Examples

```
>>> model.predict_topk(h=["Mongolia"], r=["isLocatedIn"], topk=3)
[('Asia', 0.99), ('Europe', 0.02), ...]
```

```
>>> model.predict_topk(r=["isLocatedIn"], t=["Asia"], topk=5)
[('Mongolia', 0.85), ('China', 0.82), ...]
```

triple_score (*h*: List[str] | str = None, *r*: List[str] | str = None, *t*: List[str] | str = None, *logits*=False) → torch.FloatTensor

Predict triple score

Parameter

head_entity: List[str]

String representation of selected entities.

relation: List[str]

String representation of selected relations.

tail_entity: List[str]

String representation of selected entities.

logits: bool

If logits is True, unnormalized score returned

Returns: Tuple

pytorch tensor of triple score

return_multi_hop_query_results (*aggregated_query_for_all_entities, k: int, only_scores*)

single_hop_query_answering (*query: tuple, only_scores: bool = True, k: int = None, use_logits: bool = True*)

answer_multi_hop_query (*query_type: str | None = None, query: Tuple[str | Tuple[str, str], Ellipsis] | None = None, queries: List[Tuple[str | Tuple[str, str], Ellipsis]] | None = None, tnorm: str = 'prod', neg_norm: str = 'standard', lambda_: float = 0.0, k: int = 10, only_scores: bool = False, use_logits: bool = True*)
→ List[Tuple[str, torch.Tensor]] | List[List[Tuple[str, torch.Tensor]]]

Answer multi-hop EPFO (Existential Positive First-Order) queries.

Supports 9 query types: 1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up. See docs/guides/multi_hop_queries.md for detailed query patterns.

Parameters

- **query_type** – Query pattern name. One of: - 1p: (e, (r,)) # One-hop - 2p: (e, (r1, r2)) # Two-hop - 3p: (e, (r1, r2, r3)) # Three-hop - 2i: ((e1, (r1,)), (e2, (r2,))) # Two-way intersection - 3i: ((e1, (r1,)), (e2, (r2,)), (e3, (r3,))) # Three-way intersection - ip: (((e1, (r1,)), (e2, (r2,))), (r3,)) # Intersection + projection - pi: ((e, (r1, r2)), (r3,)) # Projection + intersection (2i meets 2p) - 2u: ((e1, (r1,)), (e2, (r2,))) # Two-way union - up: ((e, (r1, r2)), (e, (r3,))) # Union + projection
- **query** – Single query tuple matching the query_type pattern.
- **queries** – Batch of queries. If provided, query must be None.
- **tnorm** – T-norm for intersection/union. Options: “prod”, “min”.
- **neg_norm** – Negation norm. Options: “standard”, “sugeno”, “yager”.
- **lambda** – Parameter for sugeno and yager negation (0.0-1.0).
- **k** – Number of top answer entities to return.
- **only_scores** – If True, return only scores tensor. If False, return (entity, score) tuples.
- **use_logits** – If True, use raw model logits. If False, use sigmoid probabilities.

Returns

List[(entity, score), ...] of top-k answers. For batch queries: List of such lists, one per query.

Return type

For single query

Raises

- **ValueError** – If query_type is not in {1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up}.
- **AssertionError** – If query structure doesn't match query_type pattern.

Examples

```
>>> # 1p: Find entities located in Asia
>>> model.answer_multi_hop_query(
...     query_type="1p",
...     query=("Asia", ("isLocatedIn",)),
...     k=5
... )
[("Mongolia", 0.92), ("China", 0.89), ...]
```

```
>>> # 2p: Two-hop query (e.g., "capital of countries in Europe")
>>> model.answer_multi_hop_query(
...     query_type="2p",
...     query=("Europe", ("isLocatedIn", "hasCapital")),
...     k=3
... )
[("Paris", 0.85), ("Berlin", 0.82), ...]
```

```
>>> # 2i: Intersection query
>>> model.answer_multi_hop_query(
...     query_type="2i",
...     query=((("Asia", ("isLocatedIn",)), ("Mountains", ("hasGeography",))),
...     k=5
... )
[("Nepal", 0.78), ("Tibet", 0.65), ...]
```

See also

- docs/guides/multi_hop_queries.md: Complete guide with all query patterns
- tests/test_answer_multi_hop_query.py: Usage examples

find_missing_triples (confidence: float, entities: List[str] = None, relations: List[str] = None, topk: int = 10, at_most: int = sys.maxsize) → Set

Find missing triples

Iterative over a set of entities E and a set of relation R :

orall e in E and orall r in R f(e,r,x)

Return (e,r,x)

otin G and f(e,r,x) > confidence

confidence: float

A threshold for an output of a sigmoid function given a triple.

topk: int

Highest ranked k item to select triples with $f(e,r,x) > \text{confidence}$.

at_most: int

Stop after finding at_most missing triples

$\{(e,r,x) \mid f(e,r,x) > \text{confidence} \text{ and } (e,r,x)$

otin G

predict_literals (*entity*: List[str] | str = None, *attribute*: List[str] | str = None, *denormalize_preds*: bool = True) → numpy.ndarray

Predicts literal values for given entities and attributes.

Parameters

- **entity** (*Union*[List[str], str]) – Entity or list of entities to predict literals for.
- **attribute** (*Union*[List[str], str]) – Attribute or list of attributes to predict literals for.
- **denormalize_preds** (*bool*) – If True, denormalizes the predictions.

Returns

Predictions for the given entities and attributes.

Return type

numpy ndarray

dicee.models

Submodules

dicee.models.adopt

ADOPT Optimizer Implementation.

This module implements the ADOPT (Adaptive Optimization with Precise Tracking) algorithm, an advanced optimization method for training neural networks.

ADOPT Overview:

ADOPT is an adaptive learning rate optimization algorithm that combines the benefits of momentum-based methods with per-parameter learning rate adaptation. Unlike Adam, which applies momentum to raw gradients, ADOPT normalizes gradients first and then applies momentum, leading to more stable training dynamics.

Key Features: - Gradient normalization before momentum application - Adaptive per-parameter learning rates - Optional gradient clipping that grows with training steps - Support for decoupled weight decay (AdamW-style) - Multiple execution modes: single-tensor, multi-tensor (foreach), and fused (planned)

Algorithm Comparison:

Adam: $m = \beta_1 * m + (1 - \beta_1) * g$, $\theta = \theta - \alpha * m / \sqrt{v}$ ADOPT: $m = \beta_1 * m + (1 - \beta_1) * g / \sqrt{v}$, $\theta = \theta - \alpha * m$

The key difference is that ADOPT normalizes gradients before momentum, which provides better stability and can lead to improved convergence.

Classes:

- ADOPT: Main optimizer class (extends torch.optim.Optimizer)

Functions:

- adopt: Functional API for ADOPT algorithm computation
- _single_tensor_adopt: Single-tensor implementation (TorchScript compatible)
- _multi_tensor_adopt: Multi-tensor implementation using foreach operations

Performance:

- Single-tensor: Default, compatible with torch.jit.script
- Multi-tensor (foreach): 2-3x faster on GPU through vectorization
- Fused (planned): Would provide maximum performance via specialized kernels

Example:

```
>>> import torch
>>> from dicee.models.adopt import ADOPT
>>>
>>> model = torch.nn.Linear(10, 1)
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01, decouple=True)
>>>
>>> # Training loop
>>> for epoch in range(num_epochs):
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     optimizer.step()
```

References:

Original implementation: <https://github.com/iShohei220/adopt>

Notes:

This implementation is based on the original ADOPT implementation and adapted to work with the PyTorch optimizer interface and the dicee framework.

Classes

ADOPT

ADOPT Optimizer.

Functions

`adopt(params, grads, exp_avgs, exp_avg_sqs, state_steps)`

Functional API that performs ADOPT algorithm computation.

Module Contents

```
class dicee.models.adopt.ADOPT(params: torch.optim.optimizer.ParamsT,  
    lr: float | torch.Tensor = 0.001, betas: Tuple[float, float] = (0.9, 0.9999), eps: float = 1e-06,  
    clip_lambda: Callable[[int], float] | None = lambda step: ..., weight_decay: float = 0.0,  
    decouple: bool = False, *, foreach: bool | None = None, maximize: bool = False,  
    capturable: bool = False, differentiable: bool = False, fused: bool | None = None)
```

Bases: `torch.optim.optimizer.Optimizer`

ADOPT Optimizer.

ADOPT is an adaptive learning rate optimization algorithm that combines momentum-based updates with adaptive per-parameter learning rates. It uses exponential moving averages of gradients and squared gradients, with gradient clipping for stability.

The algorithm performs the following key operations: 1. Normalizes gradients by the square root of the second moment estimate 2. Applies optional gradient clipping based on the training step 3. Updates parameters using momentum-smoothed normalized gradients 4. Supports decoupled weight decay (AdamW-style) or L2 regularization

Mathematical formulation:

$$m_t = \beta_1 * m_{t-1} + (1 - \beta_1) * \text{clip}(g_t / \sqrt{v_t}) \quad v_t = \beta_2 * v_{t-1} + (1 - \beta_2) * g_t^2 \quad \theta_t = \theta_{t-1} - \alpha * m_t$$

where:

- θ_t : parameter at step t
- g_t : gradient at step t
- m_t : first moment estimate (momentum)
- v_t : second moment estimate (variance)
- α : learning rate
- β_1, β_2 : exponential decay rates
- `clip()`: optional gradient clipping function

Reference:

Original implementation: <https://github.com/iShohei220/adopt>

Parameters

- **params** (*ParamsT*) – Iterable of parameters to optimize or dicts defining parameter groups.
- **lr** (*float or Tensor, optional*) – Learning rate. Can be a float or 1-element Tensor. Default: `1e-3`
- **betas** (*Tuple[float, float], optional*) – Coefficients (β_1, β_2) for computing running averages of gradient and its square. β_1 controls momentum, β_2 controls variance. Default: `(0.9, 0.9999)`

- **eps** (*float, optional*) – Term added to denominator to improve numerical stability. Default: 1e-6
- **clip_lambda** (*Callable[[int], float], optional*) – Function that takes the step number and returns the gradient clipping threshold. Common choices: - lambda step: step**0.25 (default, gradually increases clipping threshold) - lambda step: 1.0 (constant clipping) - None (no clipping) Default: lambda step: step**0.25
- **weight_decay** (*float, optional*) – Weight decay coefficient (L2 penalty). Default: 0.0
- **decouple** (*bool, optional*) – If True, uses decoupled weight decay (AdamW-style), applying weight decay directly to parameters. If False, adds weight decay to gradients (L2 regularization). Default: False
- **foreach** (*bool, optional*) – If True, uses the faster foreach implementation for multi-tensor operations. Default: None (auto-select)
- **maximize** (*bool, optional*) – If True, maximizes parameters instead of minimizing. Useful for reinforcement learning. Default: False
- **capturable** (*bool, optional*) – If True, the optimizer is safe to capture in a CUDA graph. Requires learning rate as Tensor. Default: False
- **differentiable** (*bool, optional*) – If True, the optimization step can be differentiated. Useful for meta-learning. Default: False
- **fused** (*bool, optional*) – If True, uses fused kernel implementation (currently not supported). Default: None

Raises

- **ValueError** – If learning rate, epsilon, betas, or weight_decay are invalid.
- **RuntimeError** – If fused is enabled (not currently supported).
- **RuntimeError** – If lr is a Tensor with foreach=True and capturable=False.

Example

```
>>> # Basic usage
>>> optimizer = ADOPT(model.parameters(), lr=0.001)
>>> optimizer.zero_grad()
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With decoupled weight decay
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01,
↳ decouple=True)
```

```
>>> # Custom gradient clipping
>>> optimizer = ADOPT(model.parameters(), clip_lambda=lambda step: max(1.0,
↳ step**0.5))
```

Note

- For most use cases, the default hyperparameters work well
- Consider using decouple=True for better generalization (similar to AdamW)

- The `clip_lambda` function helps stabilize training in early steps

`clip_lambda`

`__setstate__` (*state*)

Restore optimizer state from a checkpoint.

This method handles backward compatibility when loading optimizer state from older versions. It ensures all required fields are present with default values and properly converts step counters to tensors if needed.

Key responsibilities: 1. Set default values for newly added hyperparameters 2. Convert old-style scalar step counters to tensor format 3. Place step tensors on appropriate devices based on capturable/fused modes

Parameters

state (*dict*) – Optimizer state dictionary (typically from `torch.load()`).

Note

- This enables loading checkpoints saved with older ADOPT versions
- Step counters are converted to appropriate device/dtype for compatibility
- Capturable and fused modes require step tensors on parameter devices

`step` (*closure=None*)

Perform a single optimization step.

This method executes one iteration of the ADOPT optimization algorithm across all parameter groups. It orchestrates the following workflow:

1. Optionally evaluates a closure to recompute the loss (useful for algorithms like LBFGS or when loss needs multiple evaluations)
2. For each parameter group: - Collects parameters with gradients and their associated state - Extracts hyperparameters (betas, learning rate, etc.) - Calls the functional `adopt()` API to perform the actual update
3. Returns the loss value if a closure was provided

The functional API (`adopt()`) handles three execution modes: - Single-tensor: Updates one parameter at a time (default, JIT-compatible) - Multi-tensor (foreach): Batches operations for better performance - Fused: Uses fused CUDA kernels (not yet implemented)

Gradient scaling support: This method is compatible with automatic mixed precision (AMP) training. It can access `grad_scale` and `found_inf` attributes for gradient unscaling and inf/nan detection when used with `GradScaler`.

Parameters

closure (*Callable, optional*) – A callable that reevaluates the model and returns the loss. The closure should: - Enable gradients (`torch.enable_grad()`) - Compute forward pass - Compute loss - Compute backward pass - Return the loss value Example: `lambda: (loss := model(x), loss.backward(), loss)[-1]` Default: `None`

Returns

The loss value returned by the closure, or `None` if no closure was provided.

Return type
Optional[Tensor]

Example

```
>>> # Standard usage
>>> loss = criterion(model(input), target)
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With closure (e.g., for line search)
>>> def closure():
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     return loss
>>> loss = optimizer.step(closure)
```

Note

- Call `zero_grad()` before computing gradients for the next step
- CUDA graph capture is checked for safety when `capturable=True`
- The method is thread-safe for different parameter groups

`dicee.models.adapt`.**adopt** (*params: List[torch.Tensor], grads: List[torch.Tensor], exp_avgs: List[torch.Tensor], exp_avg_sqs: List[torch.Tensor], state_steps: List[torch.Tensor], foreach: bool | None = None, capturable: bool = False, differentiable: bool = False, fused: bool | None = None, grad_scale: torch.Tensor | None = None, found_inf: torch.Tensor | None = None, has_complex: bool = False, *, beta1: float, beta2: float, lr: float | torch.Tensor, clip_lambda: Callable[[int], float] | None, weight_decay: float, decouple: bool, eps: float, maximize: bool*)

Functional API that performs ADOPT algorithm computation.

This is the main functional interface for the ADOPT optimization algorithm. It dispatches to one of three implementations based on the execution mode:

1. **Single-tensor mode** (default): Updates parameters one at a time - Compatible with `torch.jit.script` - More flexible but slower - Used when `foreach=False` or automatically for small models
2. **Multi-tensor (foreach) mode**: Batches operations across tensors - 2-3x faster on GPU through vectorization - Groups tensors by device/dtype automatically - Used when `foreach=True`
3. **Fused mode**: Uses specialized fused kernels (not yet implemented) - Would provide maximum performance - Currently raises `RuntimeError` if enabled

Algorithm overview (ADOPT):

ADOPT adapts learning rates per-parameter while using momentum on normalized gradients. The key innovation is normalizing gradients before momentum, which provides more stable training than standard Adam.

Mathematical formulation:

Normalize gradient by its historical variance $\text{normed_g_t} = \text{g_t} / \sqrt{(\text{v_t} + \epsilon)}$

```

# Optional gradient clipping for stability normed_g_t = clip(normed_g_t, threshold(t))
# Momentum on normalized gradients (key difference from Adam) m_t =  $\beta_1 * m_{t-1} + (1 - \beta_1) * \text{normed\_g\_t}$ 
# Parameter update  $\theta_t = \theta_{t-1} - \alpha * m_t$ 
# Update variance estimate  $v_t = \beta_2 * v_{t-1} + (1 - \beta_2) * g_t^2$ 

```

where:

- θ : parameters
- g : gradients
- m : first moment (momentum of normalized gradients)
- v : second moment (variance of raw gradients)
- α : learning rate
- β_1, β_2 : exponential decay rates
- ϵ : numerical stability constant
- `clip()`: gradient clipping function based on step

Automatic mode selection:

When `foreach` and `fused` are both `None` (default), the function automatically selects the best implementation based on: - Parameter types and devices - Whether differentiable mode is enabled - Learning rate type (float vs Tensor) - Capturable mode requirements

param params
Parameters to optimize.

type params
`List[Tensor]`

param grads
Gradients for each parameter.

type grads
`List[Tensor]`

param exp_avgs
First moment estimates (momentum).

type exp_avgs
`List[Tensor]`

param exp_avg_sqs
Second moment estimates (variance).

type exp_avg_sqs
`List[Tensor]`

param state_steps
Step counters (must be singleton tensors).

type state_steps
`List[Tensor]`

param foreach
Whether to use multi-tensor implementation. `None`: auto-select based on configuration (default).

type foreach
Optional[bool]

param capturable
If True, ensure CUDA graph capture safety.

type capturable
bool

param differentiable
If True, allow gradients through optimization step.

type differentiable
bool

param fused
If True, use fused kernels (not implemented).

type fused
Optional[bool]

param grad_scale
Gradient scaler for AMP training.

type grad_scale
Optional[Tensor]

param found_inf
Flag for inf/nan detection in AMP.

type found_inf
Optional[Tensor]

param has_complex
Whether any parameters are complex-valued.

type has_complex
bool

param beta1
Exponential decay rate for first moment (momentum). Typical range: 0.9-0.95.

type beta1
float

param beta2
Exponential decay rate for second moment (variance). Typical range: 0.999-0.9999 (higher than Adam).

type beta2
float

param lr
Learning rate. Can be a scalar Tensor for dynamic learning rate with capturable=True.

type lr
Union[float, Tensor]

param clip_lambda
Function that maps step number to gradient clipping threshold. None disables clipping.

type clip_lambda
Optional[Callable[[int], float]]

param weight_decay
Weight decay coefficient (L2 penalty).

type weight_decay
float

param decouple
If True, use decoupled weight decay (AdamW-style). Recommended for better generalization.

type decouple
bool

param eps
Small constant for numerical stability in normalization.

type eps
float

param maximize
If True, maximize objective instead of minimize.

type maximize
bool

raises RuntimeError
If torch.jit.script is used with foreach or fused.

raises RuntimeError
If state_steps contains non-tensor elements.

raises RuntimeError
If fused=True (not yet implemented).

raises RuntimeError
If lr is Tensor with foreach=True and capturable=False.

Example

```
>>> # Typically called by ADOPT optimizer, not directly
>>> adopt (
...     params=[p1, p2],
...     grads=[g1, g2],
...     exp_avgs=[m1, m2],
...     exp_avg_sqs=[v1, v2],
...     state_steps=[step1, step2],
...     beta1=0.9,
...     beta2=0.9999,
...     lr=0.001,
...     clip_lambda=lambda s: s**0.25,
...     weight_decay=0.01,
...     decouple=True,
...     eps=1e-6,
...     maximize=False,
... )
```

 **Note**

- For distributed training, this API is compatible with torch/distributed/optim
- The foreach mode is generally preferred for GPU training
- Complex parameters are handled transparently by viewing as real
- First optimization step only initializes variance, doesn't update parameters

➡ See also

- ADOPT class: High-level optimizer interface
- `_single_tensor_adopt`: Single-tensor implementation details
- `_multi_tensor_adopt`: Multi-tensor implementation details

dicee.models.base_model

Attributes

logger

Classes

<i>BaseKGELightning</i>	Thin PyTorch Lightning wrapper shared by all KGE models.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>IdentityClass</i>	No-op normalisation / dropout placeholder.

Module Contents

`dicee.models.base_model.logger`

class `dicee.models.base_model.BaseKGELightning` (*args, **kwargs)

Bases: `lightning.LightningModule`

Thin PyTorch Lightning wrapper shared by all KGE models.

Provides the standard Lightning training loop hooks (`training_step`, `on_train_epoch_end`, `configure_optimizers`) as well as a helper for reporting model size. All concrete KGE models should extend *BaseKGE* rather than this class directly.

training_step_outputs = []

mem_of_model() → Dict

Size of model in MB and number of params

training_step(batch, batch_idx=None)

Execute one optimisation step for the given mini-batch.

Handles two- and three-element batches produced by the different dataset classes (`KvsAll` / `NegSample` vs. `KvsSample`).

Parameters

- **batch** (*tuple*) – (*x*, *y*) for standard scoring, or (*x*, *y_select*, *y*) for sample-based labelling.
- **batch_idx** (*int*, *optional*) – Index of the current batch (unused, kept for Lightning API compat).

Returns

Scalar loss value for this batch.

Return type

`torch.FloatTensor`

loss_function (*yhat_batch*: `torch.FloatTensor`, *y_batch*: `torch.FloatTensor`) → `torch.FloatTensor`

Compute the loss between model predictions and targets.

Delegates to `self.loss` which is configured in `BaseKGE.__init__` based on the scoring technique (BCEWithLogitsLoss for entity/relation prediction, CrossEntropyLoss for classification).

Parameters

- **yhat_batch** (`torch.FloatTensor`) – Model output scores, shape (batch_size, *).
- **y_batch** (`torch.FloatTensor`) – Ground-truth labels of the same shape as *yhat_batch*.

Returns

Scalar loss value.

Return type

`torch.FloatTensor`

on_train_epoch_end (*args, **kwargs)

Called in the training loop at the very end of the epoch.

To access all batch outputs at the end of the epoch, you can cache step outputs as an attribute of the `LightningModule` and access them in this hook:

```
class MyLightningModule(L.LightningModule):
    def __init__(self):
        super().__init__()
        self.training_step_outputs = []

    def training_step(self):
        loss = ...
        self.training_step_outputs.append(loss)
        return loss

    def on_train_epoch_end(self):
        # do something with all training_step outputs, for example:
        epoch_mean = torch.stack(self.training_step_outputs).mean()
        self.log("training_epoch_mean", epoch_mean)
        # free up the memory
        self.training_step_outputs.clear()
```

test_epoch_end (outputs: `List[Any]`)

test_dataloader () → `None`

An iterable or collection of iterables specifying test samples.

For more information about multiple dataloaders, see this section.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

 Warning

do not assign state in `prepare_data`

- `test()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

 Note

If you don't need a test dataset and a `test_step()`, you don't need to implement this method.

`val_dataloader()` → None

An iterable or collection of iterables specifying validation samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `fit()`
- `validate()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

 Note

If you don't need a validation dataset and a `validation_step()`, you don't need to implement this method.

`predict_dataloader()` → None

An iterable or collection of iterables specifying prediction samples.

For more information about multiple dataloaders, see this section.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `predict()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Returns

A `torch.utils.data.DataLoader` or a sequence of them specifying prediction samples.

`train_dataloader()` → None

An iterable or collection of iterables specifying training samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs`** to a positive integer.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

 Warning

do not assign state in `prepare_data`

- `fit()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

configure_optimizers (*parameters=None*)

Instantiate and return the optimiser for training.

The optimiser type is taken from `self.optimizer_name` which is set in `BaseKGE.init_params_with_sanity_checking()` from the `--optim` argument. Supported values: 'SGD', 'Adam', 'Adopt', 'AdamW', 'NAdam', 'Adagrad', 'ASGD', 'Muon'.

Parameters

parameters (*iterable, optional*) – Model parameters to optimise. Defaults to `self.parameters()` when `None`.

Returns

The configured optimiser instance.

Return type

`torch.optim.Optimizer`

class `dicee.models.base_model.BaseKGE` (*args: dict*)

Bases: `BaseKGELightning`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

```

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

defer_large_embeddings

```

init_entity_embeddings (*embedding_dim*: *int* | *None* = *None*) → *None*

Create (or re-create) the entity embedding table.

This is the single place that honours *defer_large_embeddings*: when entity rows are sharded across FSDP ranks the table must stay *None* until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning `self.entity_embeddings` directly, so the deferral is never silently undone.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

init_relation_embeddings (*embedding_dim*: *int* | *None* = *None*) → *None*

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

forward_byte_pair_encoded_k_vs_all (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*)

→ *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from *self.args* with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the *args* dict. Falls back to sensible defaults for any missing key so that minimal *args* dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple[torch.LongTensor, torch.LongTensor]*,

y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample()*
- (batch, 3) tensor → *forward_triples()*
- (batch, 2) tensor → *forward_k_vs_all()*
- BPE triple tensor → *forward_byte_pair_encoded_triple()*
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all()*

Parameters

• **x** (*torch.LongTensor* or *Tuple[torch.LongTensor, torch.LongTensor]*) – Either a plain index tensor or a (*triple_idx*, *target_idx*) tuple for sample-based labelling.

• **y_idx** (*torch.LongTensor*, optional) – Target entity indices used by *forward_k_vs_sample()*. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x*: torch.LongTensor) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises ValueError to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

class dicee.models.base_model.IdentityClass (*args=None*)

Bases: torch.nn.Module

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (--normalization None). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = None

```
__call__(x)

static forward(x)
```

dicee.models.clifford

Attributes

<code>logger</code>

Classes

<code>Keci</code>	Keci: Knowledge Graph Embedding via Clifford Algebra.
<code>KeciTransformer</code>	Keci with Transformer architecture.
<code>TransformerBlock</code>	A single transformer block with self-attention and MLP.
<code>TransformerSelfAttention</code>	Multi-head self-attention for the Keci Transformer.
<code>TransformerMLP</code>	MLP for the Keci Transformer.
<code>CKeci</code>	Without learning dimension scaling
<code>DeCaL</code>	Base class for all Knowledge Graph Embedding models.

Module Contents

`dicee.models.clifford.logger`

class `dicee.models.clifford.Keci` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Keci: Knowledge Graph Embedding via Clifford Algebra.

Embeds entities and relations as multi-vectors in the Clifford algebra $Cl_{\{p,q\}}(R^d)$ and scores triples via the Clifford product. The algebra is parameterised by two non-negative integers p and q :

- $p = 0, q = 0$ — reduces to a standard bilinear (DistMult-like) model.
- $p = 0, q = 1$ — equivalent to ComplEx.
- Larger p and q capture higher-order geometric interactions.

The embedding dimension must satisfy `embedding_dim % (p + q + 1) == 0`; the resulting quotient is stored as `self.r`.

Parameters

args (*dict*) – Configuration dictionary. Recognised keys (beyond those in `BaseKGE`): p (int, default 0) and q (int, default 0).

References

Demir et al., *Clifford Embeddings — A Generalized Approach for Embedding in Normed Algebras*, ECML 2023.

`name = 'Keci'`

p

q

r

requires_grad_for_interactions = True

compute_sigma_pp (*hp, rp*)

Compute $\sigma_{pp} = \sum_{i=1}^p \sum_{k=i+1}^p (h_{i,r,k} - h_{k,r,i}) e_i e_k$

σ_{pp} captures the interactions between along p bases. For instance, let $p = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$. This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_qq (*hq, rq*)

Compute $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_{j,r,k} - h_{k,r,j}) e_j e_k$ captures the interactions between along q bases. For instance, let $q = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$. This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_pq (**, hp, hq, rp, rq*)

$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_{i,r,j} - h_{j,r,i}) e_i e_j$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

apply_coefficients (*hp, hq, rp, rq*)

Multiplying a base vector with its scalar coefficient

clifford_multiplication (*h0, hp, hq, r0, rp, rq*)

Compute our CL multiplication

$h = h_0 + \sum_{i=1}^p h_{i,r} e_i + \sum_{j=p+1}^{p+q} h_{j,r} e_j$
 $r = r_0 + \sum_{i=1}^p r_{i,r} e_i + \sum_{j=p+1}^{p+q} r_{j,r} e_j$

$e_i^2 = +1$ for $i \leq p$ $e_j^2 = -1$ for $p < j \leq p+q$ $e_i e_j = -e_j e_i$ for i

$e_i e_j$

$h\ r = \sigma_0 + \sigma_p + \sigma_q + \sigma_{\{pp\}} + \sigma_{\{q\}} + \sigma_{\{pq\}}$ where

(1) $\sigma_0 = h_0\ r_0 + \sum_{i=1}^p (h_0\ r_i - h_i\ r_0) e_i - \sum_{j=p+1}^{p+q} (h_j\ r_j) e_j$

(2) $\sigma_p = \sum_{i=1}^p (h_0\ r_i + h_i\ r_0) e_i$

(3) $\sigma_q = \sum_{j=p+1}^{p+q} (h_0\ r_j + h_j\ r_0) e_j$

(4) $\sigma_{\{pp\}} = \sum_{i=1}^{p-1} \sum_{k=i+1}^p (h_i\ r_k - h_k\ r_i) e_i e_k$

(5) $\sigma_{\{qq\}} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_j\ r_k - h_k\ r_j) e_j e_k$

(6) $\sigma_{\{pq\}} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i\ r_j - h_j\ r_i) e_i e_j$

construct_cl_multivector (*x*: torch.FloatTensor, *r*: int, *p*: int, *q*: int)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Split a flat embedding vector into the three Clifford components.

Given an embedding *x* of dimension $d = r + r \cdot p + r \cdot q$, returns the scalar part *a*₀, the *p*-blade part *a*_p, and the *q*-blade part *a*_q.

Parameters

- **x** (torch.FloatTensor) – Shape (batch_size, d).
- **r** (int) – Scalar block size (embedding_dim // (p + q + 1)).
- **p** (int) – Number of positive-signature basis elements.
- **q** (int) – Number of negative-signature basis elements.

Returns

- **a**₀ (torch.FloatTensor) – Shape (batch_size, r) — scalar (grade-0) part.
- **a**_p (torch.FloatTensor) – Shape (batch_size, r, p) — positive-blade part.
- **a**_q (torch.FloatTensor) – Shape (batch_size, r, q) — negative-blade part.

forward_k_vs_with_explicit (*x*: torch.Tensor) → torch.FloatTensor

KvsAll scoring using an explicit loop over $\sigma_{pp}/\sigma_{qq}/\sigma_{pq}$ terms.

Functionally equivalent to `forward_k_vs_all()` but computes the higher-order interaction terms (σ_{pp} , σ_{qq} , σ_{pq}) with explicit nested loops rather than einsum contractions. Kept for reference and correctness verification.

Parameters

x (torch.Tensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

k_vs_all_score (*bpe_head_ent_emb*: torch.FloatTensor, *bpe_rel_ent_emb*: torch.FloatTensor, *E*: torch.FloatTensor) → torch.FloatTensor

Compute Clifford-product scores for a head/relation batch vs. all entities.

Decomposes the head-entity and relation embeddings into Clifford multi-vectors, performs the $Cl_{\{p,q\}}$ product, and inner-products the result against the entity embedding matrix *E*.

Parameters

- **bpe_head_ent_emb** (*torch.FloatTensor*) – Head-entity embeddings, shape (batch_size, embedding_dim).
- **bpe_rel_ent_emb** (*torch.FloatTensor*) – Relation embeddings, shape (batch_size, embedding_dim).
- **E** (*torch.FloatTensor*) – All entity embeddings, shape (num_entities, embedding_dim).

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_all (*x: torch.Tensor*) → *torch.FloatTensor*

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations $\mathbb{R}^d$.
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— *x*: *torch.LongTensor* with (n,2) shape :rtype: *torch.FloatTensor* with (n, **IEI**) shape

construct_batch_selected_cl_multivector (*x: torch.FloatTensor, r: int, p: int, q: int*)
→ tuple[*torch.FloatTensor, torch.FloatTensor, torch.FloatTensor*]

Split a batched, *k*-selected embedding tensor into Clifford components.

A variant of *construct_cl_multivector()* for tensors that have an extra *k* dimension (e.g. when scoring against *k* sampled targets).

Parameters

- **x** (*torch.FloatTensor*) – Shape (batch_size, *k*, *d*).
- **r** (*int*) – Scalar block size.
- **p** (*int*) – Number of positive-signature basis elements.
- **q** (*int*) – Number of negative-signature basis elements.

Returns

- **a0** (*torch.FloatTensor*) – Shape (batch_size, *k*, *r*).
- **ap** (*torch.FloatTensor*) – Shape (batch_size, *k*, *r*, *p*).
- **aq** (*torch.FloatTensor*) – Shape (batch_size, *k*, *r*, *q*).

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*) → *torch.FloatTensor*

Parameter

x: *torch.LongTensor* with (n,2) shape

target_entity_idx: *torch.LongTensor* with (n, *k*) shape *k* denotes the selected number of examples.

rtype

torch.FloatTensor with (n, *k*) shape

score (*h, r, t*)

forward_triples (*x: torch.Tensor*) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,3) shape

rtype

torch.FloatTensor with (n) shape

class dicee.models.clifford.**KeciTransformer** (*args*)

Bases: *Keci*

Keci with Transformer architecture.

Concatenates h0, hp, hq, r0, rp, rq into a single embedding vector and processes through transformer.

name = 'KeciTransformer'

use_clifford_mul

seq_len = 1

transformer

lm_head

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

Kvsall training

Parameter

x: torch.LongTensor with (n,2) shape

rtype

torch.FloatTensor with (n, **IEI**) shape

class dicee.models.clifford.**TransformerBlock** (*n_embd, n_head, dropout=0.0, bias=False*)

Bases: torch.nn.Module

A single transformer block with self-attention and MLP.

ln_1

attn

ln_2

mlp

forward (*x*)

class dicee.models.clifford.**TransformerSelfAttention** (*n_embd, n_head, dropout=0.0, bias=False*)

Bases: torch.nn.Module

Multi-head self-attention for the Keci Transformer.

c_attn

```

    c_proj

    attn_dropout

    resid_dropout

    n_head

    n_embd

    dropout = 0.0

    forward(x)

```

class dicee.models.clifford.**TransformerMLP** (*n_embd*, *dropout*=0.0, *bias*=False)

Bases: torch.nn.Module

MLP for the Keci Transformer.

```

    c_fc

    gelu

    c_proj

    dropout

    forward(x)

```

class dicee.models.clifford.**CKeci** (*args*)

Bases: *Keci*

Without learning dimension scaling

```

    name = 'CKeci'

    requires_grad_for_interactions = False

```

class dicee.models.clifford.**DeCaL** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches *forward()* calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by *vars(argparse.Namespace)*. Required keys: *embedding_dim*, *num_entities*, *num_relations*, *learning_rate* (or *lr*), *optim*, *scoring_technique*.

```

    name = 'DeCaL'

    p

```

q

r

re

forward_triples (x : *torch.Tensor*) $\rightarrow$ torch.FloatTensor

Parameter

x: torch.LongTensor with (n,) shape

rtype

torch.FloatTensor with (n) shape

cl_pqr (a : *torch.tensor*) $\rightarrow$ torch.tensor

Input: tensor(batch_size, emb_dim) $\rightarrow$ output: tensor with 1+p+q+r components with size (batch_size, emb_dim/(1+p+q+r)) each.

1) takes a tensor of size (batch_size, emb_dim), split it into 1 + p + q + r components, hence 1+p+q+r must be a divisor of the emb_dim. 2) Return a list of the 1+p+q+r components vectors, each are tensors of size (batch_size, emb_dim/(1+p+q+r))

compute_sigmas_single (*list_h_emb, list_r_emb, list_t_emb*)

here we compute all the sums with no others vectors interaction taken with the scalar product with t, that is,

$$s0 = h_0 r_0 t_0 s1 = \sum_{i=1}^p h_i r_i t_0 s2 = \sum_{j=p+1}^{p+q} h_j r_j t_0 s3 = \sum_{i=1}^q (h_0 r_i t_i + h_i r_0 t_i) s4 = \sum_{i=p+1}^{p+q} (h_0 r_i t_i + h_i r_0 t_i) s5 = \sum_{i=p+q+1}^{p+q+r} (h_0 r_i t_i + h_i r_0 t_i)$$

and return:

$$\sigma_{pqr} t = \sigma_0 \cdot t_0 = s0 + s1 - s2s3, s4 \text{ and } s5$$

compute_sigmas_multivect (*list_h_emb, list_r_emb*)

Here we compute and return all the sums with vectors interaction for the same and different bases.

For same bases vectors interaction we have

$$\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (h_i r_{i'} - h_{i'} r_i) (\text{model the interactions between } e_i \text{ and } e_{i'} \text{ for } 1 \leq i, i' \leq p) \sigma_{pq} = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (h_j r_{j'} - h_{j'} r_j)$$

For different base vector interactions, we have

$$\sigma_{pq} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) (\text{interactions between } e_i \text{ and } e_j \text{ for } 1 \leq i \leq p \text{ and } p+1 \leq j \leq p+q) \sigma_{pr} = \sum_{i=1}^p \sum_{j=p+q+1}^{p+q+r} (h_i r_j - h_j r_i)$$

forward_k_vs_all (x : *torch.Tensor*) $\rightarrow$ torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations
- (2) Construct head entity and relation embeddings according to $\text{Cl}_{\{p,q,r\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— x: torch.LongTensor with (n,) shape :rtype: torch.FloatTensor with (n, **IE**) shape

apply_coefficients (h0, hp, hq, hk, r0, rp, rq, rk)

Multiplying a base vector with its scalar coefficient

construct_cl_multivector (x: torch.FloatTensor, re: int, p: int, q: int, r: int)

→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Construct a batch of multivectors $Cl_{\{p,q,r\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,d) shape

returns

- **a0** (torch.FloatTensor)
- **ap** (torch.FloatTensor)
- **aq** (torch.FloatTensor)
- **ar** (torch.FloatTensor)

compute_sigma_pp (hp, rp)

Compute .. math:

$$\sigma_{\{p,p\}}^* = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (x_{iy_{i'}} - x_{i'y_i})$$

$\sigma_{\{pp\}}$ captures the interactions between along p bases For instance, let p e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3 , and e_2 e_3 This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_qq (hq, rq)

Compute

$$\sigma_{q,q}^* = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (x_j y_{j'} - x_{j'} y_j) E q.16$$

$\sigma_{\{q\}}$ captures the interactions between along q bases For instance, let q e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3 , and e_2 e_3 This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_rr(hk, rk)

$$\sigma_{r,r}^* = \sum_{k=p+q+1}^{p+q+r-1} \sum_{k'=k+1}^p (x_k y_{k'} - x_{k'} y_k)$$

compute_sigma_pq(*, hp, hq, rp, rq)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

compute_sigma_pr(*, hp, hk, rp, rk)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

compute_sigma_qr(*, hq, hk, rq, rk)

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

dicce.models.complex

Classes

<i>ConEx</i>	Convolutional ComplEx Knowledge Graph Embeddings
<i>AConEx</i>	Additive Convolutional ComplEx Knowledge Graph Embeddings
<i>ComplEx</i>	Base class for all Knowledge Graph Embedding models.
<i>RotatE</i>	RotatE: knowledge graph embedding by relational rotation in complex space.

Module Contents

class `dicee.models.complex.ConEx` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Convolutional ComplEx Knowledge Graph Embeddings

name = 'ConEx'

conv2d

fc_num_input

fc1

norm_fc1

bn_conv2d

feature_map_dropout

residual_convolution (*C_1*: `Tuple[torch.Tensor, torch.Tensor]`,
C_2: `Tuple[torch.Tensor, torch.Tensor]`) → `torch.FloatTensor`

Compute residual score of two complex-valued embeddings. :param *C_1*: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param *C_2*: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

forward_k_vs_all (*x*: `torch.Tensor`) → `torch.FloatTensor`

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.Tensor`) → `torch.FloatTensor`

Score a batch of (head, relation, tail) index triples.

Parameters

x (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*x*: *torch.Tensor*, *target_entity_idx*: *torch.Tensor*)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

`torch.FloatTensor`

class `dicdee.models.complex.AConEx` (*args*)

Bases: `dicdee.models.base_model.BaseKGE`

Additive Convolutional ComplEx Knowledge Graph Embeddings

name = 'AConEx'

conv2d

fc_num_input

fc1

norm_fc1

bn_conv2d

feature_map_dropout

residual_convolution (*C_1*: *Tuple[torch.Tensor, torch.Tensor]*,
C_2: *Tuple[torch.Tensor, torch.Tensor]*) → `torch.FloatTensor`

Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

forward_k_vs_all (*x*: *torch.Tensor*) → `torch.FloatTensor`

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_triples (*x*: *torch.Tensor*) → `torch.FloatTensor`

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*x*: *torch.Tensor*, *target_entity_idx*: *torch.Tensor*)

Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

class dicee.models.complex.**Complex** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from BaseKGELightning and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches forward() calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- forward_triples() — score a batch of (h, r, t) triples.
- forward_k_vs_all() — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by vars(argparse.Namespace). Required keys: embedding_dim, num_entities, num_relations, learning_rate (or lr), optim, scoring_technique.

name = 'Complex'

static score (*head_ent_emb*: *torch.FloatTensor*, *rel_ent_emb*: *torch.FloatTensor*, *tail_ent_emb*: *torch.FloatTensor*)

static k_vs_all_score (*emb_h*: *torch.FloatTensor*, *emb_r*: *torch.FloatTensor*, *emb_E*: *torch.FloatTensor*)

Parameters

- emb_h
- emb_r
- emb_E

forward_k_vs_all (*x*: *torch.LongTensor*) → torch.FloatTensor

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises ValueError to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: torch.LongTensor, *target_entity_idx*: torch.LongTensor)

Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

class dicee.models.complex.RotatE(*args*)

Bases: `dicee.models.base_model.BaseKGE`

RotatE: knowledge graph embedding by relational rotation in complex space.

Represents each entity as a complex vector in $\mathbb{C}^{(d/2)}$, obtained by splitting its d -dimensional embedding into a real half and an imaginary half. Each relation is represented by $d/2$ phase angles θ that define a unit-modulus rotation $r_i = e^{i\theta_i}$; since a phase angle needs only one real number, the relation embedding table is reinitialised at half of the entity embedding size. A true triple (h, r, t) should satisfy $h \circ r \approx t$ under element-wise complex multiplication, giving the score:

$$f(h, r, t) = \text{margin} - ||h \circ r - t||_2$$

Unlike TransE, RotatE can model symmetric, antisymmetric, inverse, and composition relation patterns.

References

Sun et al., *RotatE: Knowledge Graph Embedding by Relational Rotation in Complex Space*, ICLR 2019. <https://arxiv.org/abs/1902.10197>

name = 'RotatE'

margin = 6.0

half_dim

relation_embeddings

score (*head_ent_emb*: torch.FloatTensor, *rel_ent_emb*: torch.FloatTensor, *tail_ent_emb*: torch.FloatTensor) → torch.FloatTensor

Score a batch of triples using the RotatE margin-distance formula.

Parameters

- **head_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **tail_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **rel_ent_emb** (torch.FloatTensor) – Shape (batch_size, $d/2$) relation phase angles θ .

Returns

Shape (batch_size,) scores equal to $\text{margin} - ||h \circ r - t||_2$.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: *torch.Tensor*) → *torch.FloatTensor*

KvsAll forward pass: score head/relation against all entities.

Computes $\text{margin} = ||h \circ r - e||_2$ for every entity embedding *e*.

Parameters

x (*torch.Tensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: *torch.Tensor*, *target_entity_idx*: *torch.Tensor*) → *torch.FloatTensor*

KvsSample forward pass: score head/relation against a sampled entity subset.

Computes $\text{margin} = ||h \circ r - e||_2$ for each of the *k* sampled entities *e*.

Parameters

- **x** (*torch.Tensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

- **target_entity_idx** (*torch.Tensor*) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

dicee.models.dualE

Classes

DualE

Dual Quaternion Knowledge Graph Embeddings
(<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

Module Contents

class `dicee.models.dualE.DualE` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Dual Quaternion Knowledge Graph Embeddings (<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

name = 'DualE'

entity_embeddings

relation_embeddings

num_ent = None

kvsall_score (*e_1_h, e_2_h, e_3_h, e_4_h, e_5_h, e_6_h, e_7_h, e_8_h, e_1_t, e_2_t, e_3_t, e_4_t, e_5_t, e_6_t, e_7_t, e_8_t, r_1, r_2, r_3, r_4, r_5, r_6, r_7, r_8*) → torch.tensor

KvsAll scoring function

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_triples (*idx_triple: torch.tensor*) → torch.tensor

Negative Sampling forward pass:

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_k_vs_all (*x*)

KvsAll forward pass

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

T (*x: torch.tensor*) → torch.tensor

Transpose function

Input: Tensor with shape (nxm) Output: Tensor with shape (mxn)

dicее.models.ensemble

Classes

EnsembleKGE

Module Contents

```
class dicее.models.ensemble.EnsembleKGE (models: list = None, seed_model=None,
    pretrained_models: List = None)
```

```
    name
```

```
    train_mode = True
```

```

args

named_children()

property example_input_array

parameters()

modules()

__iter__()

__len__()

eval()

to(device)

state_dict()
    Return the state dict of the ensemble.

load_state_dict(state_dict, strict=True)
    Load the state dict into the ensemble.

mem_of_model()

__call__(x_batch)

step()

get_embeddings()

__str__()

```

dicее.models.fsdp_models

Row-wise sharded entity embeddings for multi-GPU training.

Entity embeddings are partitioned row-wise across ranks:

- `dist.all_to_all_single` routes index requests to the owning rank
- A custom autograd Function propagates gradients back through the `all_to_all`
- `_LocalSparseAdam` updates only the rows that received a non-zero gradient

Public API:

`model.setup_fsdп_training(device, lr)` — called by trainer before loop `model.gather_entity_embeddings_on_rank_zero()` — called by trainer after loop `model.save_local_shard_checkpoint(path)` — per-rank, called periodically during loop `model.load_local_shard_checkpoint(path)` — per-rank, called by trainer on resume `create_fsdп_sharded_model_class(ModelClass)` — factory used by `static_funcs.py`

Classes

<i>RowWiseShardedEmbedding</i>	Entity embedding table sharded row-wise across ranks.
<i>FSDPShardedEntityModel</i>	Mixin that defers entity embedding allocation and wires up row-wise sharding.

Functions

`create_fsdp_sharded_model_class(...)`Return a row-wise sharded variant of *model_class*.

Module Contents

```
class dicee.models.fsdp_models.RowWiseShardedEmbedding (num_embeddings: int,  
    embedding_dim: int, rank: int, world_size: int, device: torch.device,  
    weight_dtype: torch.dtype = torch.bfloat16)
```

Bases: `torch.nn.Module`

Entity embedding table sharded row-wise across ranks.

Each rank owns rows [start_row, end_row). `forward()` looks like `nn.Embedding` so every existing scoring function works unchanged.

num_embeddings

embedding_dim

rank

world_size

device

shard_size

start_row

end_row

local_rows

weight

forward (*entity_ids: torch.LongTensor*) → `torch.FloatTensor`

```
class dicee.models.fsdp_models.FSDPShardedEntityModel (args)
```

Bases: `dicee.models.base_model.BaseKGE`

Mixin that defers entity embedding allocation and wires up row-wise sharding.

Lifecycle

1. `__init__()` — `entity_embeddings` is `None`
2. `setup_fsdp_training(dev, lr)` — creates `RowWiseShardedEmbedding`
3. `gather_entity_embeddings_on_rank_zero()` — called by trainer after last epoch

```
setup_fsdp_training (device: torch.device, lr: float, optimizer_cls=None,  
    optimizer_kwargs: dict = None, adam_device: torch.device = None) → None
```

Create the per-rank embedding shard and its local sparse Adam.

adam_device — where `exp_avg` / `exp_avg_sq` are stored:

GPU (default) : pure GPU kernels, no PCIe, ~42.8 GiB extra GPU RAM per rank. CPU : lower GPU footprint, PCIe transfer each step.

Falls back to the embedding device (GPU) when not specified.

gather_entity_embeddings_on_rank_zero () → torch.nn.Embedding | None

Gather the full entity table on rank 0; return None on other ranks.

save_local_shard_checkpoint (*path: str*) → None

Save this rank's entity embedding shard + its local Adam state to *path*.

Cheap and $O(\text{local_rows})$ — unlike `gather_entity_embeddings_on_rank_zero()`, no collective communication is involved, so this is safe to call periodically even when the full table is too large to gather.

load_local_shard_checkpoint (*path: str*) → None

Load this rank's entity embedding shard + local Adam state from *path*.

The checkpoint must have been saved by a run with the same `world_size`: shard boundaries (`start_row/local_rows`) are derived from `world_size` and `num_entities`, so a mismatch means this rank does not own the same rows as when the checkpoint was written.

get_embeddings ()

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (`num_entities`, `embedding_dim`).
- **relation_embeddings** (*numpy.ndarray*) – Shape (`num_relations`, `embedding_dim`).

forward_k_vs_all (**args, **kwargs*)

Score a (`head`, `relation`) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (`batch_size`, `num_entities`) score matrix.

Return type

torch.FloatTensor

`dicee.models.fsdps_models.create_fsdps_sharded_model_class` (
 model_class: Type[dicee.models.base_model.BaseKGE]
 → *Type[dicee.models.base_model.BaseKGE]*

Return a row-wise sharded variant of *model_class*.

The returned class inherits all scoring functions from *model_class* unchanged. Only `entity_embeddings` is replaced at training time.

dicee.models.function_space

Attributes

logger

Classes

FMult

Learning Knowledge Neural Graphs

continues on next page

Table 42 – continued from previous page

<i>GFMult</i>	Learning Knowledge Neural Graphs
<i>FMult2</i>	Learning Knowledge Neural Graphs
<i>LFMult1</i>	Embedding with trigonometric functions. We represent all entities and relations in the complex number space as:
<i>LFMult</i>	Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:

Module Contents

`dicee.models.function_space.logger`

class `dicee.models.function_space.FMult` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Learning Knowledge Neural Graphs

name = 'FMult'

entity_embeddings

relation_embeddings

k

num_sample = 50

gamma

roots

weights

compute_func (*weights: torch.FloatTensor, x*) → `torch.FloatTensor`

chain_func (*weights, x: torch.FloatTensor*)

forward_triples (*idx_triple: torch.Tensor*) → `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

x (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

class `dicee.models.function_space.GFMult` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Learning Knowledge Neural Graphs

name = 'GFMult'

entity_embeddings

```

relation_embeddings

k

num_sample = 250

roots

weights

compute_func(weights: torch.FloatTensor, x) → torch.FloatTensor

chain_func(weights, x: torch.FloatTensor)

forward_triples(idx_triple: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

```

```

class dicee.models.function_space.FMult2(args)
    Bases: dicee.models.base_model.BaseKGE
    Learning Knowledge Neural Graphs

    name = 'FMult2'

    n_layers = 3

    k

    n = 50

    score_func = 'compositional'

    discrete_points

    entity_embeddings

    relation_embeddings

    build_func(Vec)

    build_chain_funcs(list_Vec)

    compute_func(W, b, x) → torch.FloatTensor

    function(list_W, list_b)

    trapezoid(list_W, list_b)

```

forward_triples (*idx_triple: torch.Tensor*) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

```
class dicee.models.function_space.LFMult1(args)
```

Bases: *dicee.models.base_model.BaseKGE*

Embedding with trigonometric functions. We represent all entities and relations in the complex number space as: $f(x) = \sum_{k=0}^{d-1} w_k e^{kix}$. and use the three differents scoring function as in the paper to evaluate the score

name = 'LFMult1'

entity_embeddings

relation_embeddings

forward_triples (*idx_triple*)

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

tri_score (*h, r, t*)

vtp_score (*h, r, t*)

```
class dicee.models.function_space.LFMult(args)
```

Bases: *dicee.models.base_model.BaseKGE*

Embedding with polynomial functions. We represent all entities and relations in the polynomial space as: $f(x) = \sum_{i=0}^{d-1} a_k x^{i\%d}$ and use the three differents scoring function as in the paper to evaluate the score. We also consider combining with Neural Networks.

name = 'LFMult'

entity_embeddings

relation_embeddings

degree

m

x_values

forward_triples (*idx_triple*)

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

construct_multi_coeff (*x*)

poly_NN (*x, coefh, coefr, coeft*)

Constructing a 2 layers NN to represent the embeddings. $h = \text{sigma}(wh^T x + bh)$, $r = \text{sigma}(wr^T x + br)$, $t = \text{sigma}(wt^T x + bt)$

linear (*x, w, b*)

scalar_batch_NN (*a, b, c*)

element wise multiplication between a,b and c: Inputs : a, b, c ==> torch.tensor of size batch_size x m x d
Output : a tensor of size batch_size x d

tri_score (*coeff_h, coeff_r, coeff_t*)

this part implement the trilinear scoring techniques:

$\text{score}(h,r,t) = \int_{\{0\}^1} h(x)r(x)t(x) dx = \sum_{\{i,j,k=0\}^{d-1}} \text{dfrac}\{a_i*b_j*c_k\}\{1+(i+j+k)\%d\}$

1. generate the range for i,j and k from [0 d-1]
2. perform $\text{dfrac}\{a_i*b_j*c_k\}\{1+(i+j+k)\%d\}$ in parallel for every batch
3. take the sum over each batch

vtp_score (*h, r, t*)

this part implement the vector triple product scoring techniques:

$\text{score}(h,r,t) = \int_{\{0\}^1} h(x)r(x)t(x) dx = \sum_{\{i,j,k=0\}^{d-1}} \text{dfrac}\{a_i*c_j*b_k - b_i*c_j*a_k\}\{(1+(i+j)\%d)(1+k)\}$

1. generate the range for i,j and k from [0 d-1]
2. Compute the first and second terms of the sum
3. Multiply with then denominator and take the sum
4. take the sum over each batch

comp_func (*h, r, t*)

this part implement the function composition scoring techniques: i.e. score = <hor, t>

polynomial (*coeff, x, degree*)

This function takes a matrix tensor of coefficients (coeff), a tensor vector of points x and range of integer [0,1,...d] and return a vector tensor (coeff[0][0] + coeff[0][1]x +...+ coeff[0][d]x^d,

coeff[1][0] + coeff[1][1]x +...+ coeff[1][d]x^d)

pop (*coeff, x, degree*)

This function allow us to evaluate the composition of two polynomes without for loops :) it takes a matrix tensor of coefficients (coeff), a matrix tensor of points x and range of integer [0,1,...d]

and return a tensor ($\text{coeff}[0][0] + \text{coeff}[0][1]x + \dots + \text{coeff}[0][d]x^d$,
 $\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d$)

dicee.models.literal

Classes

<i>LiteralEmbeddings</i>	A model for learning and predicting numerical literals using pre-trained KGE.
--------------------------	---

Module Contents

```
class dicee.models.literal.LiteralEmbeddings (num_of_data_properties: int, embedding_dims: int,
        entity_embeddings: torch.tensor, dropout: float = 0.3, gate_residual=True,
        freeze_entity_embeddings=True)
```

Bases: torch.nn.Module

A model for learning and predicting numerical literals using pre-trained KGE.

num_of_data_properties

Number of data properties (attributes).

Type

int

embedding_dims

Dimension of the embeddings.

Type

int

entity_embeddings

Pre-trained entity embeddings.

Type

torch.tensor

dropout

Dropout rate for regularization.

Type

float

gate_residual

Whether to use gated residual connections.

Type

bool

freeze_entity_embeddings

Whether to freeze the entity embeddings during training.

Type

bool

embedding_dim

```

num_of_data_properties

hidden_dim

gate_residual = True

freeze_entity_embeddings = True

entity_embeddings

data_property_embeddings

fc

fc_out

dropout

gated_residual_proj

layer_norm

forward(entity_idx, attr_idx)

```

Parameters

- **entity_idx** (*Tensor*) – Entity indices (batch).
- **attr_idx** (*Tensor*) – Attribute (Data property) indices (batch).

Returns

scalar predictions.

Return type

Tensor

property device

dicee.models.octonion

Classes

<i>OMult</i>	Base class for all Knowledge Graph Embedding models.
<i>ConvO</i>	Base class for all Knowledge Graph Embedding models.
<i>AConvO</i>	Additive Convolutional Octonion Knowledge Graph Embeddings

Functions

<i>octonion_mul</i> (*, O_1, O_2)
<i>octonion_mul_norm</i> (*, O_1, O_2)

Module Contents

`dicee.models.octonion.octonion_mul(*, O_1, O_2)`

`dicee.models.octonion.octonion_mul_norm(*, O_1, O_2)`

class `dicee.models.octonion.OMult(args)`

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'OMult'

static `octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4, emb_rel_e5, emb_rel_e6, emb_rel_e7)`

score (*head_ent_emb*: `torch.FloatTensor`, *rel_ent_emb*: `torch.FloatTensor`, *tail_ent_emb*: `torch.FloatTensor`)

k_vs_all_score (*bpe_head_ent_emb*, *bpe_rel_ent_emb*, *E*)

forward_k_vs_all (*x*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples, i.e., `[score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8]`, shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch, |Entities|)

class `dicee.models.octonion.ConvO(args: dict)`

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'ConvO'

```

conv2d

fc_num_input

fc1

bn_conv2d

norm_fc1

feature_map_dropout

static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                           emb_rel_e5, emb_rel_e6, emb_rel_e7)

residual_convolution(O_1, O_2)

forward_triples(x: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

forward_k_vs_all(x: torch.Tensor)
    Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] =>
    [0.0,0.1,...,0.8], shape=> (1, Entities) Given a batch of head entities and relations => shape (size of batch,|
    Entities)

class dicee.models.octonion.AConvO(args: dict)
    Bases: dicee.models.base_model.BaseKGE
    Additive Convolutional Octonion Knowledge Graph Embeddings

    name = 'AConvO'

    conv2d

    fc_num_input

    fc1

    bn_conv2d

    norm_fc1

    feature_map_dropout

    static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                               emb_rel_e5, emb_rel_e6, emb_rel_e7)

    residual_convolution(O_1, O_2)

```


forward_triples (*x: torch.Tensor*) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*x: torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

dicee.models.pykeen_models

Attributes

logger

Classes

PykeenKGE

A class for using knowledge graph embedding models implemented in Pykeen

Module Contents

dicee.models.pykeen_models.**logger**

class dicee.models.pykeen_models.**PykeenKGE** (*args: dict*)

Bases: *dicee.models.base_model.BaseKGE*

A class for using knowledge graph embedding models implemented in Pykeen

Notes: Pykeen_DistMult: C Pykeen_ComplEx: Pykeen_QuatE: Pykeen_MuRE: Pykeen_CP: Pykeen_HolE: Pykeen_HolE: Pykeen_HolE: Pykeen_TransD: Pykeen_TransE: Pykeen_TransF: Pykeen_TransH: Pykeen_TransR:

model_kwargs

name

model

loss_history = []

args

entity_embeddings = None

relation_embeddings = None

```

forward_k_vs_all (x: torch.LongTensor)
    # => Explicit version by this we can apply bn and dropout

    # (1) Retrieve embeddings of heads and relations + apply Dropout & Normalization if given. h, r =
    self.get_head_relation_representation(x) # (2) Reshape (1). if self.last_dim > 0:
        h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim,
        self.last_dim)

    # (3) Reshape all entities. if self.last_dim > 0:
        t = self.entity_embeddings.weight.reshape(self.num_entities, self.embedding_dim, self.last_dim)

    else:
        t = self.entity_embeddings.weight

    # (4) Call the score_t from interactions to generate triple scores. return self.interaction.score_t(h=h, r=r,
    all_entities=t, slice_size=1)

forward_triples (x: torch.LongTensor) → torch.FloatTensor
    # => Explicit version by this we can apply bn and dropout

    # (1) Retrieve embeddings of heads, relations and tails and apply Dropout & Normalization if given. h, r, t =
    self.get_triple_representation(x) # (2) Reshape (1). if self.last_dim > 0:
        h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim,
        self.last_dim) t = t.reshape(len(x), self.embedding_dim, self.last_dim)

    # (3) Compute the triple score return self.interaction.score(h=h, r=r, t=t, slice_size=None, slice_dim=0)

abstractmethod forward_k_vs_sample (x: torch.LongTensor, target_entity_idx)
    Score a (head, relation) batch against a sampled subset of entities.

    Used by KvsSample and lvsSample datasets. Sub-classes that support sample-based labelling must over-
    ride this method.

    Returns
        Shape (batch_size, k) score matrix where k is the number of sampled target entities.

    Return type
        torch.FloatTensor

```

dicee.models.quaternion

Classes

<i>QMult</i>	Base class for all Knowledge Graph Embedding models.
<i>ConvQ</i>	Convolutional Quaternion Knowledge Graph Embed-
	dings
<i>AConvQ</i>	Additive Convolutional Quaternion Knowledge Graph
	Embeddings

Functions

<i>quaternion_mul_with_unit_norm</i> (*, Q_1, Q_2)
--

Module Contents

`dicee.models.quaternion.quaternion_mul_with_unit_norm(*, Q_1, Q_2)`

class `dicee.models.quaternion.QMult` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (*h*, *r*, *t*) triples.
- `forward_k_vs_all()` — score a (*h*, *r*) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'QMult'

explicit = True

quaternion_multiplication_followed_by_inner_product (*h*, *r*, *t*)

Parameters

- **h** – shape: (**batch_dims*, *dim*) The head representations.
- **r** – shape: (**batch_dims*, *dim*) The head representations.
- **t** – shape: (**batch_dims*, *dim*) The tail representations.

Returns

Triple scores.

static quaternion_normalizer (*x*: `torch.FloatTensor`) → `torch.FloatTensor`

Normalize the length of relation vectors, if the forward constraint has not been applied yet.

Absolute value of a quaternion

$$|a + bi + cj + dk| = \sqrt{a^2 + b^2 + c^2 + d^2}$$

L2 norm of quaternion vector:

$$\|x\|^2 = \sum_{i=1}^d |x_i|^2 = \sum_{i=1}^d (x_i.re^2 + x_i.im_1^2 + x_i.im_2^2 + x_i.im_3^2)$$

Parameters

x – The vector.

Returns

The normalized vector.

score (*head_ent_emb*: `torch.FloatTensor`, *rel_ent_emb*: `torch.FloatTensor`, *tail_ent_emb*: `torch.FloatTensor`)

`k_vs_all_score(bpe_head_ent_emb, bpe_rel_ent_emb, E)`

Parameters

- `bpe_head_ent_emb`
- `bpe_rel_ent_emb`
- `E`

`forward_k_vs_all(x)`

Parameters

`x`

`forward_k_vs_sample(x, target_entity_idx)`

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e.,
[score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **|Entities|**) Given a batch of head entities and
relations => shape (size of batch,| Entities|)

`class dicee.models.quaternion.ConvQ(args)`

Bases: `dicee.models.base_model.BaseKGE`

Convolutional Quaternion Knowledge Graph Embeddings

`name = 'ConvQ'`

`entity_embeddings`

`relation_embeddings`

`conv2d`

`fc_num_input`

`fc1`

`bn_conv1`

`bn_conv2`

`feature_map_dropout`

`residual_convolution(Q_1, Q_2)`

`forward_triples(indexed_triple: torch.Tensor) → torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

`x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
[head_idx, relation_idx, tail_idx].`

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

`forward_k_vs_all(x: torch.Tensor)`

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] =>
[0.0,0.1,...,0.8], shape=> (1, **|Entities|**) Given a batch of head entities and relations => shape (size of batch,|
Entities|)

```

class dicee.models.quaternion.AConvQ(args)
    Bases: dicee.models.base_model.BaseKGE
    Additive Convolutional Quaternion Knowledge Graph Embeddings
    name = 'AConvQ'
    entity_embeddings
    relation_embeddings
    conv2d
    fc_num_input
    fc1
    bn_conv1
    bn_conv2
    feature_map_dropout
    residual_convolution(Q_1, Q_2)
    forward_triples(indexed_triple: torch.Tensor) → torch.Tensor
        Score a batch of (head, relation, tail) index triples.

        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].

        Returns
            Shape (batch_size,) triple scores.

        Return type
            torch.FloatTensor

    forward_k_vs_all(x: torch.Tensor)
        Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] =>
        [0.0,0.1,...,0.8], shape=> (1, |Entities|) Given a batch of head entities and relations => shape (size of batch,|
        Entities)

```

dicee.models.real

Classes

<i>DistMult</i>	DistMult: bilinear diagonal knowledge graph embedding.
<i>TransE</i>	TransE: translation-based knowledge graph embedding.
<i>TransH</i>	TransH: translation-based knowledge graph embedding on relation-specific hyperplanes.
<i>MuRE</i>	MuRE: multi-relational graph embedding with relation-specific diagonal scaling, translation, and entity biases.
<i>Shallom</i>	Shallom: shallow neural model for relation prediction.
<i>Pyke</i>	Pyke: Physical Embedding Model for Knowledge Graphs.
<i>CoKEConfig</i>	Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.

continues on next page

Table 50 – continued from previous page

CoKE

Contextualized Knowledge Graph Embedding (CoKE) model.

Module Contents

class `dicee.models.real.DistMult` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

DistMult: bilinear diagonal knowledge graph embedding.

Scores a triple (h , r , t) as the element-wise product of the head, relation, and tail embeddings summed over the embedding dimension:

$$f(h, r, t) = \sum_i h_i \cdot r_i \cdot t_i$$

Simple yet effective baseline; incapable of modelling asymmetric relations.

References

Yang et al., *Embedding Entities and Relations for Learning and Inference in Knowledge Bases*, ICLR 2015. <https://arxiv.org/abs/1412.6575>

name = 'DistMult'

k_vs_all_score (*emb_h*: `torch.FloatTensor`, *emb_r*: `torch.FloatTensor`, *emb_E*: `torch.FloatTensor`)
→ `torch.FloatTensor`

Score a head/relation batch against all entity embeddings.

Computes $(h * r) @ E^T$ after applying hidden dropout and normalisation to the element-wise product.

Parameters

- **emb_h** (`torch.FloatTensor`) – Head entity embeddings, shape (batch_size, embedding_dim).
- **emb_r** (`torch.FloatTensor`) – Relation embeddings, shape (batch_size, embedding_dim).
- **emb_E** (`torch.FloatTensor`) – All entity embeddings, shape (num_entities, embedding_dim).

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_all (*x*: `torch.LongTensor`) → `torch.FloatTensor`

KvsAll forward pass: score head/relation against all entities.

Parameters

x (`torch.LongTensor`) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*x*: torch.LongTensor, *target_entity_idx*: torch.LongTensor) → torch.FloatTensor

KvsSample forward pass: score head/relation against a sampled entity subset.

Parameters

- **x** (torch.LongTensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (torch.LongTensor) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

score (*h*: torch.FloatTensor, *r*: torch.FloatTensor, *t*: torch.FloatTensor) → torch.FloatTensor

Score a batch of (head, relation, tail) embedding triples.

Parameters

- **h** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **r** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **t** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class dicee.models.real.**TransE** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

TransE: translation-based knowledge graph embedding.

Models a relation *r* as a translation in embedding space such that $h + r \approx t$ for a true triple (*h*, *r*, *t*). The score function is defined as:

$$f(h, r, t) = \text{margin} - ||h + r - t||_2$$

TransE is effective for 1-to-1 relations but struggles with reflexive, one-to-many, and many-to-one patterns.

References

Bordes et al., *Translating Embeddings for Modeling Multi-relational Data*, NeurIPS 2013. <https://proceedings.neurips.cc/paper/2013/file/1cecc7a77928ca8133fa24680a88d2f9-Paper.pdf>

name = 'TransE'

margin = 4

score (*head_ent_emb*: torch.FloatTensor, *rel_ent_emb*: torch.FloatTensor, *tail_ent_emb*: torch.FloatTensor) → torch.FloatTensor

Score a batch of triples using the TransE margin-distance formula.

Parameters

- **head_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).

- **rel_ent_emb** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).
- **tail_ent_emb** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) scores equal to $\text{margin} - ||h + r - t||_2$.

Return type

torch.FloatTensor

forward_k_vs_all (*x: torch.Tensor*) → *torch.FloatTensor*

KvsAll forward pass: score head/relation against all entities.

Computes $\text{margin} - ||h + r - e||_2$ for every entity embedding *e*.

Parameters

x (*torch.Tensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

class dicee.models.real.**TransH** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

TransH: translation-based knowledge graph embedding on relation-specific hyperplanes.

Addresses TransE’s inability to model one-to-many, many-to-one, and many-to-many relations by letting each relation *r* define its own hyperplane. An entity is first projected onto that hyperplane before the translation is applied, so the same entity can be positioned differently depending on the relation. Concretely, for a triple (*h*, *r*, *t*) with unit-norm hyperplane normal *r_w* and in-hyperplane translation vector *r_d*:

```
h_r = h - (r_w^T h) * r_w
t_r = t - (r_w^T t) * r_w
f(h, r, t) = -||h_r + r_d - t_r||_2^2
```

r_w reuses the inherited *relation_embeddings* table and is normalised to unit norm before every use, since only a unit-norm normal yields a true orthogonal projection onto the hyperplane. *r_d* is a second, relation-specific embedding table analogous to TransE’s relation vector.

References

Wang et al., *Knowledge Graph Embedding by Translating on Hyperplanes*, AAAI 2014. <https://aaai.org/papers/8870-knowledge-graph-embedding-by-translating-on-hyperplanes/>

name = 'TransH'

relation_normal_translations

forward_triples (*x: torch.LongTensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: torch.LongTensor) → torch.FloatTensor

KvsAll forward pass: score head/relation against all entities.

Every candidate tail entity must be projected onto the batch row's relation-specific hyperplane, and the hyperplane normal differs per batch row, so unlike TransE this materialises a full (batch_size, num_entities, embedding_dim) tensor of projected tail candidates – $O(B * E * D)$ memory, versus TransE's $O(B * E)$.

Parameters

x (torch.LongTensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: torch.LongTensor, *target_entity_idx*: torch.LongTensor) → torch.FloatTensor

KvsSample forward pass: score head/relation against a sampled entity subset.

Same hyperplane-projection logic as `forward_k_vs_all()`, but restricted to the *k* sampled tail candidates instead of every entity, so it costs $O(B * k * D)$ memory instead of $O(B * E * D)$.

Parameters

- **x** (torch.LongTensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (torch.LongTensor) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

class dicee.models.real.**MuRE** (*args*)

Bases: `dicee.models.base_model.BaseKGE`

MuRE: multi-relational graph embedding with relation-specific diagonal scaling, translation, and entity biases.

Scores a triple (*h*, *r*, *t*) as:

$$f(h, r, t) = -||R_r \odot h + t_r - t||_2 + b_h + b_t$$

R_r is a relation-specific diagonal matrix applied to the head via an element-wise (Hadamard) product; it reuses the inherited `relation_embeddings` table directly, shape (num_relations, embedding_dim). t_r is a second, relation-specific translation vector (own embedding table, same shape convention: (num_relations, embedding_dim)) analogous to TransE's relation vector. b_h and b_t are learnable scalar biases indexed per entity (head and tail respectively), not per relation.

References

Balažević et al., *Multi-relational Poincaré Graph Embeddings*, NeurIPS 2019. <https://arxiv.org/abs/1905.09791>

`name = 'MuRE'`

`relation_translations`

`b_h`

`b_t`

`forward_triples` (x : `torch.LongTensor`) $\rightarrow$ `torch.FloatTensor`

Score a batch of (head, relation, tail) index triples.

Parameters

$\mathbf{x}$ (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

`forward_k_vs_all` (x : `torch.LongTensor`) $\rightarrow$ `torch.FloatTensor`

KvsAll forward pass: score head/relation against all entities.

Computes $-||\mathbf{R}_r \odot \mathbf{h} + \mathbf{t}_r - \mathbf{e}||_2 + \mathbf{b}_h + \mathbf{b}_e$ for every entity embedding e .

Parameters

$\mathbf{x}$ (`torch.LongTensor`) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

`forward_k_vs_sample` (x : `torch.LongTensor`, $target_entity_idx$: `torch.LongTensor`) $\rightarrow$ `torch.FloatTensor`

KvsSample forward pass: score head/relation against a sampled entity subset.

Parameters

- $\mathbf{x}$ (`torch.LongTensor`) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- `target_entity_idx` (`torch.LongTensor`) – Shape (batch_size, k) indices of the k target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

`torch.FloatTensor`

`class dicee.models.real.Shallom` ($args$)

Bases: `dicee.models.base_model.BaseKGE`

Shallom: shallow neural model for relation prediction.

Represents each triple as the concatenation of head and tail entity embeddings and feeds it through a two-layer MLP to predict the relation. Designed for the `RelationPrediction` labelling form.

References

Demir et al., *A Shallow Neural Model for Relation Prediction*, ISWC 2021. <https://arxiv.org/abs/2101.09090>

```
name = 'Shallom'
```

```
shallom
```

```
get_embeddings() → Tuple[numpy.ndarray, None]
```

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

```
forward_k_vs_all(x) → torch.FloatTensor
```

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

```
forward_triples(x) → torch.FloatTensor
```

Score a batch of triples by looking up relation scores from `forward_k_vs_all`.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

```
class dicee.models.real.Pyke(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Pyke: Physical Embedding Model for Knowledge Graphs.

Scores a triple (h, r, t) based on the average pairwise distance between head-to-relation and relation-to-tail in embedding space:

$$f(h, r, t) = \text{margin} - (||h - r||_2 + ||r - t||_2) / 2$$

The model encodes geometric proximity between entities and the relations that connect them.

```
name = 'Pyke'
```

```
dist_func
```

```
margin = 1.0
```

forward_triples (x : *torch.LongTensor*) $\rightarrow$ *torch.FloatTensor*

Score a batch of triples using the Pyke distance formula.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class *dicee.models.real.CoKEConfig*

Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.

block_size

Sequence length for transformer (3 for triples: head, relation, tail)

vocab_size

Total vocabulary size (num_entities + num_relations)

n_layer

Number of transformer layers

n_head

Number of attention heads per layer

n_embd

Embedding dimension (set to match model embedding_dim)

dropout

Dropout rate applied throughout the model

bias

Whether to use bias in linear layers

causal

Whether to use causal masking (False for bidirectional attention)

block_size: int = 3

vocab_size: int = None

n_layer: int = 6

n_head: int = 8

n_embd: int = None

dropout: float = 0.3

bias: bool = True

causal: bool = False

```
class dicee.models.real.CoKE(args, config: CoKEConfig = CoKEConfig())
```

Bases: `dicee.models.base_model.BaseKGE`

Contextualized Knowledge Graph Embedding (CoKE) model. Based on: <https://arxiv.org/pdf/1911.02168>.

CoKE uses a transformer encoder to learn contextualized representations of entities and relations. For link prediction, it predicts masked elements in (head, relation, tail) triples using bidirectional attention, similar to BERT's masked language modeling approach.

The model creates a sequence [head_emb, relation_emb, mask_emb], adds positional embeddings, and processes it through transformer layers to predict the tail entity.

```
name = 'CoKE'
```

```
config
```

```
pos_emb
```

```
mask_emb
```

```
blocks
```

```
ln_f
```

```
coke_dropout
```

```
forward_k_vs_all(x: torch.Tensor)
```

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

```
score(emb_h, emb_r, emb_t)
```

```
forward_k_vs_sample(x: torch.LongTensor, target_entity_idx: torch.LongTensor)
```

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

`dicee.models.static_funcs`

Functions

```
quaternion_mul( $\rightarrow$  Tuple[torch.Tensor, torch.Tensor, ...]) Perform quaternion multiplication
```

Module Contents

```
dicee.models.static_funcs.quaternion_mul(* , Q_1, Q_2)  
→ Tuple[torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor]  
Perform quaternion multiplication :param Q_1: :param Q_2: :return:
```

`dicee.models.transformers`

Full definition of a GPT Language Model, all of it in this single file. References: 1) the official GPT-2 TensorFlow implementation released by OpenAI: <https://github.com/openai/gpt-2/blob/master/src/model.py> 2) huggingface/transformers PyTorch implementation: https://github.com/huggingface/transformers/blob/main/src/transformers/models/gpt2/modeling_gpt2.py

Attributes

<code>logger</code>

Classes

<i>Byte</i>	Base class for all Knowledge Graph Embedding models.
<i>LayerNorm</i>	LayerNorm but with an optional bias. PyTorch doesn't support simply bias=False
<i>SelfAttention</i>	Base class for all neural network modules.
<i>MLP</i>	Base class for all neural network modules.
<i>Block</i>	Base class for all neural network modules.
<i>GPTConfig</i>	
<i>GPT</i>	Base class for all neural network modules.

Module Contents

```
dicee.models.transformers.logger
```

```
class dicee.models.transformers.Byte(*args, **kwargs)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

```
name = 'Byte'
```

config

temperature = 0.5

topk = 2

transformer

lm_head

loss_function(*yhat_batch*, *y_batch*)

Compute the loss between model predictions and targets.

Delegates to `self.loss` which is configured in `BaseKGE.__init__` based on the scoring technique (BCEWithLogitsLoss for entity/relation prediction, CrossEntropyLoss for classification).

Parameters

- **yhat_batch** (*torch.FloatTensor*) – Model output scores, shape (batch_size, *).
- **y_batch** (*torch.FloatTensor*) – Ground-truth labels of the same shape as *yhat_batch*.

Returns

Scalar loss value.

Return type

torch.FloatTensor

forward(*x*: *torch.LongTensor*)

Parameters

x (*B by T tensor*)

generate(*idx*, *max_new_tokens*, *temperature*=1.0, *top_k*=None)

Take a conditioning sequence of indices *idx* (*LongTensor* of shape (b,t)) and complete the sequence *max_new_tokens* times, feeding the predictions back into the model each time. Most likely you'll want to make sure to be in `model.eval()` mode of operation for this.

training_step(*batch*, *batch_idx*=None)

Execute one optimisation step for the given mini-batch.

Handles two- and three-element batches produced by the different dataset classes (*KvsAll* / *NegSample* vs. *KvsSample*).

Parameters

- **batch** (*tuple*) – (*x*, *y*) for standard scoring, or (*x*, *y_select*, *y*) for sample-based labelling.
- **batch_idx** (*int*, *optional*) – Index of the current batch (unused, kept for Lightning API compat).

Returns

Scalar loss value for this batch.

Return type

torch.FloatTensor

class `dicee.models.transformers.LayerNorm` (*ndim*, *bias*)

Bases: `torch.nn.Module`

LayerNorm but with an optional bias. PyTorch doesn't support simply `bias=False`

weight

bias

forward(*input*)

```
class dicee.models.transformers.SelfAttention(config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

c_attn

c_proj

attn_dropout

resid_dropout

n_head

n_embd

dropout

causal


```
flash = True
```

```
forward(x)
```

```
class dicee.models.transformers.MLP (config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

c_fc

gelu

c_proj

dropout

forward(*x*)

```
class dicee.models.transformers.Block (config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`ln_1`

`attn`

`ln_2`

`mlp`

`forward(x)`

```
class dicee.models.transformers.GPTConfig
```

`block_size: int = 1024`

`vocab_size: int = 50304`

`n_layer: int = 12`

`n_head: int = 12`

`n_embd: int = 768`

`dropout: float = 0.0`

`bias: bool = False`

`causal: bool = True`

```
class dicee.models.transformers.GPT(config)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

config

transformer

lm_head

get_num_params (*non_embedding=True*)

Return the number of parameters in the model. For non-embedding count (default), the position embeddings get subtracted. The token embeddings would too, except due to the parameter sharing these params are actually used as weights in the final layer, so we include them.

forward (*idx, targets=None*)

crop_block_size (*block_size*)

classmethod from_pretrained (*model_type, override_args=None*)

configure_optimizers (*weight_decay, learning_rate, betas, device_type*)

estimate_mfu (*fwdbwd_per_iter, dt*)

estimate model flops utilization (MFU) in units of A100 bfloat16 peak FLOPS

Attributes

<code>logger</code>
<code>logger</code>
<code>logger</code>

Classes

<code>ADOPT</code>	ADOPT Optimizer.
<code>BaseKGELightning</code>	Thin PyTorch Lightning wrapper shared by all KGE models.
<code>BaseKGE</code>	Base class for all Knowledge Graph Embedding models.
<code>IdentityClass</code>	No-op normalisation / dropout placeholder.
<code>Block</code>	Base class for all neural network modules.
<code>BaseKGE</code>	Base class for all Knowledge Graph Embedding models.
<code>DistMult</code>	DistMult: bilinear diagonal knowledge graph embedding.
<code>TransE</code>	TransE: translation-based knowledge graph embedding.
<code>TransH</code>	TransH: translation-based knowledge graph embedding on relation-specific hyperplanes.
<code>MuRE</code>	MuRE: multi-relational graph embedding with relation-specific diagonal scaling, translation, and entity biases.
<code>Shallom</code>	Shallom: shallow neural model for relation prediction.
<code>Pyke</code>	Pyke: Physical Embedding Model for Knowledge Graphs.
<code>CoKEConfig</code>	Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.
<code>CoKE</code>	Contextualized Knowledge Graph Embedding (CoKE) model.
<code>BaseKGE</code>	Base class for all Knowledge Graph Embedding models.
<code>ConEx</code>	Convolutional ComplEx Knowledge Graph Embeddings
<code>AConEx</code>	Additive Convolutional ComplEx Knowledge Graph Embeddings
<code>Complex</code>	Base class for all Knowledge Graph Embedding models.
<code>RotatE</code>	RotatE: knowledge graph embedding by relational rotation in complex space.
<code>BaseKGE</code>	Base class for all Knowledge Graph Embedding models.
<code>IdentityClass</code>	No-op normalisation / dropout placeholder.
<code>QMult</code>	Base class for all Knowledge Graph Embedding models.
<code>ConvQ</code>	Convolutional Quaternion Knowledge Graph Embeddings
<code>AConvQ</code>	Additive Convolutional Quaternion Knowledge Graph Embeddings
<code>BaseKGE</code>	Base class for all Knowledge Graph Embedding models.
<code>IdentityClass</code>	No-op normalisation / dropout placeholder.
<code>OMult</code>	Base class for all Knowledge Graph Embedding models.
<code>ConvO</code>	Base class for all Knowledge Graph Embedding models.
<code>AConvO</code>	Additive Convolutional Octonion Knowledge Graph Embeddings
<code>Keci</code>	Keci: Knowledge Graph Embedding via Clifford Algebra.
<code>CKeci</code>	Without learning dimension scaling
<code>DeCaL</code>	Base class for all Knowledge Graph Embedding models.

continues on next page

Table 55 – continued from previous page

<i>KeciTransformer</i>	Keci with Transformer architecture.
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>PykeenKGE</i>	A class for using knowledge graph embedding models implemented in Pykeen
<i>BaseKGE</i>	Base class for all Knowledge Graph Embedding models.
<i>FMult</i>	Learning Knowledge Neural Graphs
<i>GFMult</i>	Learning Knowledge Neural Graphs
<i>FMult2</i>	Learning Knowledge Neural Graphs
<i>LFMult1</i>	Embedding with trigonometric functions. We represent all entities and relations in the complex number space as:
<i>LFMult</i>	Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:
<i>DualE</i>	Dual Quaternion Knowledge Graph Embeddings (https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657)

Functions

<i>quaternion_mul</i> ($\rightarrow$ Tuple[torch.Tensor, torch.Tensor, ...])	Perform quaternion multiplication
<i>quaternion_mul_with_unit_norm</i> (*, Q_1, Q_2)	
<i>octonion_mul</i> (*, O_1, O_2)	
<i>octonion_mul_norm</i> (*, O_1, O_2)	

Package Contents

class `dicee.models.ADOPT` (*params: torch.optim.optimizer.ParamsT, lr: float | torch.Tensor = 0.001, betas: Tuple[float, float] = (0.9, 0.9999), eps: float = 1e-06, clip_lambda: Callable[[int], float] | None = lambda step: ..., weight_decay: float = 0.0, decouple: bool = False, *, foreach: bool | None = None, maximize: bool = False, capturable: bool = False, differentiable: bool = False, fused: bool | None = None*)

Bases: `torch.optim.optimizer.Optimizer`

ADOPT Optimizer.

ADOPT is an adaptive learning rate optimization algorithm that combines momentum-based updates with adaptive per-parameter learning rates. It uses exponential moving averages of gradients and squared gradients, with gradient clipping for stability.

The algorithm performs the following key operations: 1. Normalizes gradients by the square root of the second moment estimate 2. Applies optional gradient clipping based on the training step 3. Updates parameters using momentum-smoothed normalized gradients 4. Supports decoupled weight decay (AdamW-style) or L2 regularization

Mathematical formulation:

$$m_t = \beta_1 * m_{\{t-1\}} + (1 - \beta_1) * \text{clip}(g_t / \sqrt{v_t}) \quad v_t = \beta_2 * v_{\{t-1\}} + (1 - \beta_2) * g_t^2 \quad \theta_t = \theta_{\{t-1\}} - \alpha * m_t$$

where:

- θ_t : parameter at step t
- g_t : gradient at step t
- m_t : first moment estimate (momentum)

- v_t : second moment estimate (variance)
- α : learning rate
- β_1, β_2 : exponential decay rates
- `clip()`: optional gradient clipping function

Reference:

Original implementation: <https://github.com/iShohei220/adopt>

Parameters

- **params** (*ParamsT*) – Iterable of parameters to optimize or dicts defining parameter groups.
- **lr** (*float or Tensor, optional*) – Learning rate. Can be a float or 1-element Tensor. Default: 1e-3
- **betas** (*Tuple[float, float], optional*) – Coefficients (β_1, β_2) for computing running averages of gradient and its square. β_1 controls momentum, β_2 controls variance. Default: (0.9, 0.9999)
- **eps** (*float, optional*) – Term added to denominator to improve numerical stability. Default: 1e-6
- **clip_lambda** (*Callable[[int], float], optional*) – Function that takes the step number and returns the gradient clipping threshold. Common choices: - lambda step: $\text{step} \times 0.25$ (default, gradually increases clipping threshold) - lambda step: 1.0 (constant clipping) - None (no clipping) Default: lambda step: $\text{step} \times 0.25$
- **weight_decay** (*float, optional*) – Weight decay coefficient (L2 penalty). Default: 0.0
- **decouple** (*bool, optional*) – If True, uses decoupled weight decay (AdamW-style), applying weight decay directly to parameters. If False, adds weight decay to gradients (L2 regularization). Default: False
- **foreach** (*bool, optional*) – If True, uses the faster foreach implementation for multi-tensor operations. Default: None (auto-select)
- **maximize** (*bool, optional*) – If True, maximizes parameters instead of minimizing. Useful for reinforcement learning. Default: False
- **capturable** (*bool, optional*) – If True, the optimizer is safe to capture in a CUDA graph. Requires learning rate as Tensor. Default: False
- **differentiable** (*bool, optional*) – If True, the optimization step can be differentiated. Useful for meta-learning. Default: False
- **fused** (*bool, optional*) – If True, uses fused kernel implementation (currently not supported). Default: None

Raises

- **ValueError** – If learning rate, epsilon, betas, or weight_decay are invalid.
- **RuntimeError** – If fused is enabled (not currently supported).
- **RuntimeError** – If lr is a Tensor with foreach=True and capturable=False.

Example

```
>>> # Basic usage
>>> optimizer = ADOPT(model.parameters(), lr=0.001)
>>> optimizer.zero_grad()
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With decoupled weight decay
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01,
↳decouple=True)
```

```
>>> # Custom gradient clipping
>>> optimizer = ADOPT(model.parameters(), clip_lambda=lambda step: max(1.0,
↳step**0.5))
```

Note

- For most use cases, the default hyperparameters work well
- Consider using `decouple=True` for better generalization (similar to AdamW)
- The `clip_lambda` function helps stabilize training in early steps

`clip_lambda`

`__setstate__(state)`

Restore optimizer state from a checkpoint.

This method handles backward compatibility when loading optimizer state from older versions. It ensures all required fields are present with default values and properly converts step counters to tensors if needed.

Key responsibilities: 1. Set default values for newly added hyperparameters 2. Convert old-style scalar step counters to tensor format 3. Place step tensors on appropriate devices based on capturable/fused modes

Parameters

state (*dict*) – Optimizer state dictionary (typically from `torch.load()`).

Note

- This enables loading checkpoints saved with older ADOPT versions
- Step counters are converted to appropriate device/dtype for compatibility
- Capturable and fused modes require step tensors on parameter devices

`step(closure=None)`

Perform a single optimization step.

This method executes one iteration of the ADOPT optimization algorithm across all parameter groups. It orchestrates the following workflow:

1. Optionally evaluates a closure to recompute the loss (useful for algorithms like LBFGS or when loss needs multiple evaluations)

2. For each parameter group: - Collects parameters with gradients and their associated state - Extracts hyperparameters (betas, learning rate, etc.) - Calls the functional adopt() API to perform the actual update
3. Returns the loss value if a closure was provided

The functional API (adopt()) handles three execution modes: - Single-tensor: Updates one parameter at a time (default, JIT-compatible) - Multi-tensor (foreach): Batches operations for better performance - Fused: Uses fused CUDA kernels (not yet implemented)

Gradient scaling support: This method is compatible with automatic mixed precision (AMP) training. It can access grad_scale and found_inf attributes for gradient unscaling and inf/nan detection when used with GradScaler.

Parameters

closure (*Callable, optional*) – A callable that reevaluates the model and returns the loss. The closure should: - Enable gradients (torch.enable_grad()) - Compute forward pass - Compute loss - Compute backward pass - Return the loss value Example: lambda: (loss := model(x), loss.backward(), loss)[-1] Default: None

Returns

The loss value returned by the closure, or None if no closure was provided.

Return type

Optional[Tensor]

Example

```
>>> # Standard usage
>>> loss = criterion(model(input), target)
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With closure (e.g., for line search)
>>> def closure():
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     return loss
>>> loss = optimizer.step(closure)
```

Note

- Call zero_grad() before computing gradients for the next step
- CUDA graph capture is checked for safety when capturable=True
- The method is thread-safe for different parameter groups

dicee.models.logger

class dicee.models.BaseKGELightning(*args, **kwargs)

Bases: lightning.LightningModule

Thin PyTorch Lightning wrapper shared by all KGE models.

Provides the standard Lightning training loop hooks (`training_step`, `on_train_epoch_end`, `configure_optimizers`) as well as a helper for reporting model size. All concrete KGE models should extend `BaseKGE` rather than this class directly.

`training_step_outputs = []`

`mem_of_model()` → Dict

Size of model in MB and number of params

`training_step(batch, batch_idx=None)`

Execute one optimisation step for the given mini-batch.

Handles two- and three-element batches produced by the different dataset classes (`KvsAll` / `NegSample` vs. `KvsSample`).

Parameters

- **`batch`** (*tuple*) – (`x`, `y`) for standard scoring, or (`x`, `y_select`, `y`) for sample-based labelling.
- **`batch_idx`** (*int*, *optional*) – Index of the current batch (unused, kept for Lightning API compat).

Returns

Scalar loss value for this batch.

Return type

`torch.FloatTensor`

`loss_function(yhat_batch: torch.FloatTensor, y_batch: torch.FloatTensor)` → `torch.FloatTensor`

Compute the loss between model predictions and targets.

Delegates to `self.loss` which is configured in `BaseKGE.__init__` based on the scoring technique (`BCEWithLogitsLoss` for entity/relation prediction, `CrossEntropyLoss` for classification).

Parameters

- **`yhat_batch`** (*torch.FloatTensor*) – Model output scores, shape `(batch_size, *)`.
- **`y_batch`** (*torch.FloatTensor*) – Ground-truth labels of the same shape as `yhat_batch`.

Returns

Scalar loss value.

Return type

`torch.FloatTensor`

`on_train_epoch_end(*args, **kwargs)`

Called in the training loop at the very end of the epoch.

To access all batch outputs at the end of the epoch, you can cache step outputs as an attribute of the `LightningModule` and access them in this hook:

```
class MyLightningModule(L.LightningModule):
    def __init__(self):
        super().__init__()
        self.training_step_outputs = []

    def training_step(self):
        loss = ...
```

(continues on next page)

(continued from previous page)

```
self.training_step_outputs.append(loss)
return loss

def on_train_epoch_end(self):
    # do something with all training_step outputs, for example:
    epoch_mean = torch.stack(self.training_step_outputs).mean()
    self.log("training_epoch_mean", epoch_mean)
    # free up the memory
    self.training_step_outputs.clear()
```

test_epoch_end(*outputs: List[Any]*)

test_dataloader() → None

An iterable or collection of iterables specifying test samples.

For more information about multiple dataloaders, see this section.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

Warning

do not assign state in `prepare_data`

- `test()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Note

If you don't need a test dataset and a `test_step()`, you don't need to implement this method.

val_dataloader() → None

An iterable or collection of iterables specifying validation samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:`~lightning.pytorch.trainer.trainer.Trainer.reload\_dataloaders\_every\_n\_epochs`** to a positive integer.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `fit()`
- `validate()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Note

If you don't need a validation dataset and a `validation_step()`, you don't need to implement this method.

`predict_dataloader()` → None

An iterable or collection of iterables specifying prediction samples.

For more information about multiple dataloaders, see this section.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `predict()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Returns

A `torch.utils.data.DataLoader` or a sequence of them specifying prediction samples.

`train_dataloader()` → None

An iterable or collection of iterables specifying training samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **`:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs``** to a positive integer.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

Warning

do not assign state in `prepare_data`

- `fit()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

`configure_optimizers` (*parameters=None*)

Instantiate and return the optimiser for training.

The optimiser type is taken from `self.optimizer_name` which is set in `BaseKGE.init_params_with_sanity_checking()` from the `--optim` argument. Supported values: 'SGD', 'Adam', 'Adopt', 'AdamW', 'NAdam', 'Adagrad', 'ASGD', 'Muon'.

Parameters

`parameters` (*iterable, optional*) – Model parameters to optimise. Defaults to `self.parameters()` when `None`.

Returns

The configured optimiser instance.

Return type

`torch.optim.Optimizer`

`class` `dicee.models.BaseKGE` (*args: dict*)

Bases: `BaseKGELightning`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (`h`, `r`, `t`) triples.
- `forward_k_vs_all()` — score a (`h`, `r`) batch against every entity.

Parameters

`args` (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

`args`

`embedding_dim` = `None`

`num_entities` = `None`

```

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

defer_large_embeddings

init_entity_embeddings (embedding_dim: int | None = None) → None
    Create (or re-create) the entity embedding table.

```

This is the single place that honours `defer_large_embeddings`: when entity rows are sharded across FSDP ranks the table must stay `None` until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning `self.entity_embeddings` directly, so the deferral is never silently undone.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

init_relation_embeddings (*embedding_dim: int | None = None*) → None

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

forward_byte_pair_encoded_k_vs_all (*x: torch.LongTensor*) → torch.FloatTensor

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x: Tuple[torch.LongTensor, torch.LongTensor]*)
→ torch.FloatTensor

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from *self.args* with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the *args* dict. Falls back to sensible defaults for any missing key so that minimal *args* dicts (e.g. for unit tests) are still valid.

forward (*x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],*
y_idx: torch.LongTensor = None) → torch.FloatTensor

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- Tuple (*x*, *y_idx*) → *forward_k_vs_sample()*
- (batch, 3) tensor → *forward_triples()*
- (batch, 2) tensor → *forward_k_vs_all()*
- BPE triple tensor → *forward_byte_pair_encoded_triple()*

- BPE pair tensor $\rightarrow$ `forward_byte_pair_encoded_k_vs_all()`

Parameters

- **x** (`torch.LongTensor` or `Tuple[torch.LongTensor, torch.LongTensor]`) – Either a plain index tensor or a (`triple_idx`, `target_idx`) tuple for sample-based labelling.
- **y_idx** (`torch.LongTensor`, optional) – Target entity indices used by `forward_k_vs_sample()`. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.LongTensor`) $\rightarrow$ `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

- **x** (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

`torch.FloatTensor`

get_triple_representation (idx_hrt)

$\rightarrow$ `Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]`

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

- **idx_hrt** (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (*indexed_triple*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).

- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

class dicee.models.IdentityClass (*args=None*)

Bases: torch.nn.Module

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (`--normalization None`). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = None

__call__(*x*)

static forward(*x*)

class dicee.models.Block (*config*)

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

ln_1

attn

ln_2

`mlp`

`forward(x)`

`class dicee.models.BaseKGE (args: dict)`

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

loss

`selected_optimizer = None`

`normalizer_class = None`

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

`hidden_dropout`

`loss_history = []`

`byte_pair_encoding`

`max_length_subword_tokens`

`block_size`

`defer_large_embeddings`

`init_entity_embeddings (embedding_dim: int | None = None) → None`

Create (or re-create) the entity embedding table.

This is the single place that honours `defer_large_embeddings`: when entity rows are sharded across FSDP ranks the table must stay `None` until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning `self.entity_embeddings` directly, so the deferral is never silently undone.

Parameters

`embedding_dim` (*Optional[int]*) – Width of the table. Defaults to `embedding_dim`.

`init_relation_embeddings (embedding_dim: int | None = None) → None`

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

`embedding_dim` (*Optional[int]*) – Width of the table. Defaults to `embedding_dim`.

`forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor) → torch.FloatTensor`

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

`x` (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

`torch.FloatTensor`

`forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])`

→ `torch.FloatTensor`

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple*[*torch.LongTensor*, *torch.LongTensor*],
y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample()*
- (batch, 3) tensor → *forward_triples()*
- (batch, 2) tensor → *forward_k_vs_all()*
- BPE triple tensor → *forward_byte_pair_encoded_triple()*
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all()*

Parameters

- **x** (*torch.LongTensor* or *Tuple*[*torch.LongTensor*, *torch.LongTensor*]) – Either a plain index tensor or a (*triple_idx*, *target_idx*) tuple for sample-based labelling.
- **y_idx** (*torch.LongTensor*, optional) – Target entity indices used by *forward_k_vs_sample()*. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.LongTensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample(*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation(idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation(indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation(x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x*: torch.LongTensor)
→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (torch.LongTensor) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (numpy.ndarray) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (numpy.ndarray) – Shape (num_relations, embedding_dim).

class dicee.models.DistMult (*args*)

Bases: `dicee.models.base_model.BaseKGE`

DistMult: bilinear diagonal knowledge graph embedding.

Scores a triple (h, r, t) as the element-wise product of the head, relation, and tail embeddings summed over the embedding dimension:

$$f(h, r, t) = \sum_i h_i \cdot r_i \cdot t_i$$

Simple yet effective baseline; incapable of modelling asymmetric relations.

References

Yang et al., *Embedding Entities and Relations for Learning and Inference in Knowledge Bases*, ICLR 2015. <https://arxiv.org/abs/1412.6575>

name = 'DistMult'

k_vs_all_score (*emb_h*: torch.FloatTensor, *emb_r*: torch.FloatTensor, *emb_E*: torch.FloatTensor)
→ torch.FloatTensor

Score a head/relation batch against all entity embeddings.

Computes $(h * r) @ E^T$ after applying hidden dropout and normalisation to the element-wise product.

Parameters

- **emb_h** (torch.FloatTensor) – Head entity embeddings, shape (batch_size, embedding_dim).

- **emb_r** (*torch.FloatTensor*) – Relation embeddings, shape (batch_size, embedding_dim).
- **emb_E** (*torch.FloatTensor*) – All entity embeddings, shape (num_entities, embedding_dim).

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_all (*x: torch.LongTensor*) → *torch.FloatTensor*

KvsAll forward pass: score head/relation against all entities.

Parameters

• **x** (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*) → *torch.FloatTensor*

KvsSample forward pass: score head/relation against a sampled entity subset.

Parameters

- **x** (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (*torch.LongTensor*) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

score (*h: torch.FloatTensor, r: torch.FloatTensor, t: torch.FloatTensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) embedding triples.

Parameters

- **h** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).
- **r** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).
- **t** (*torch.FloatTensor*) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class *dicee.models.TransE* (*args*)

Bases: *dicee.models.base_model.BaseKGE*

TransE: translation-based knowledge graph embedding.

Models a relation r as a translation in embedding space such that $h + r \approx t$ for a true triple (h, r, t) . The score function is defined as:

$$f(h, r, t) = \text{margin} - ||h + r - t||_2$$

TransE is effective for 1-to-1 relations but struggles with reflexive, one-to-many, and many-to-one patterns.

References

Bordes et al., *Translating Embeddings for Modeling Multi-relational Data*, NeurIPS 2013. <https://proceedings.neurips.cc/paper/2013/file/1cecc7a77928ca8133fa24680a88d2f9-Paper.pdf>

name = 'TransE'

margin = 4

score (*head_ent_emb*: torch.FloatTensor, *rel_ent_emb*: torch.FloatTensor, *tail_ent_emb*: torch.FloatTensor) → torch.FloatTensor

Score a batch of triples using the TransE margin-distance formula.

Parameters

- **head_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **rel_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **tail_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).

Returns

Shape (batch_size,) scores equal to $\text{margin} - ||h + r - t||_2$.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: torch.Tensor) → torch.FloatTensor

KvsAll forward pass: score head/relation against all entities.

Computes $\text{margin} - ||h + r - e||_2$ for every entity embedding e .

Parameters

x (torch.Tensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

class dicee.models.TransH(*args*)

Bases: `dicee.models.base_model.BaseKGE`

TransH: translation-based knowledge graph embedding on relation-specific hyperplanes.

Addresses TransE's inability to model one-to-many, many-to-one, and many-to-many relations by letting each relation r define its own hyperplane. An entity is first projected onto that hyperplane before the translation is applied, so the same entity can be positioned differently depending on the relation. Concretely, for a triple (h, r, t) with unit-norm hyperplane normal r_w and in-hyperplane translation vector r_d :


```

h_r = h - (r_w^T h) * r_w
t_r = t - (r_w^T t) * r_w
f(h, r, t) = -||h_r + r_d - t_r||_2^2

```

`r_w` reuses the inherited `relation_embeddings` table and is normalised to unit norm before every use, since only a unit-norm normal yields a true orthogonal projection onto the hyperplane. `r_d` is a second, relation-specific embedding table analogous to TransE's relation vector.

References

Wang et al., *Knowledge Graph Embedding by Translating on Hyperplanes*, AAAI 2014. <https://aaai.org/papers/8870-knowledge-graph-embedding-by-translating-on-hyperplanes/>

name = 'TransH'

relation_normal_translations

forward_triples (*x*: *torch.LongTensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll forward pass: score head/relation against all entities.

Every candidate tail entity must be projected onto the batch row's relation-specific hyperplane, and the hyperplane normal differs per batch row, so unlike TransE this materialises a full (batch_size, num_entities, embedding_dim) tensor of projected tail candidates – O(B * E * D) memory, versus TransE's O(B * E).

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: *torch.LongTensor*, *target_entity_idx*: *torch.LongTensor*) → *torch.FloatTensor*

KvsSample forward pass: score head/relation against a sampled entity subset.

Same hyperplane-projection logic as `forward_k_vs_all()`, but restricted to the *k* sampled tail candidates instead of every entity, so it costs O(B * k * D) memory instead of O(B * E * D).

Parameters

- **x** (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (*torch.LongTensor*) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

class dicee.models.**MuRE**(args)

Bases: `dicee.models.base_model.BaseKGE`

MuRE: multi-relational graph embedding with relation-specific diagonal scaling, translation, and entity biases.

Scores a triple (h, r, t) as:

$$f(h, r, t) = -||R_r \odot h + t_r - t||_2 + b_h + b_t$$

R_r is a relation-specific diagonal matrix applied to the head via an element-wise (Hadamard) product; it reuses the inherited `relation_embeddings` table directly, shape (num_relations, embedding_dim). t_r is a second, relation-specific translation vector (own embedding table, same shape convention: (num_relations, embedding_dim)) analogous to TransE's relation vector. b_h and b_t are learnable scalar biases indexed per entity (head and tail respectively), not per relation.

References

Balažević et al., *Multi-relational Poincaré Graph Embeddings*, NeurIPS 2019. <https://arxiv.org/abs/1905.09791>

name = 'MuRE'

relation_translations

b_h

b_t

forward_triples(x: torch.LongTensor) → torch.FloatTensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(x: torch.LongTensor) → torch.FloatTensor

KvsAll forward pass: score head/relation against all entities.

Computes $-||R_r \odot h + t_r - e||_2 + b_h + b_e$ for every entity embedding e .

Parameters

x (torch.LongTensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: torch.LongTensor, *target_entity_idx*: torch.LongTensor) → torch.FloatTensor

KvsSample forward pass: score head/relation against a sampled entity subset.

Parameters

- **x** (torch.LongTensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (torch.LongTensor) – Shape (batch_size, k) indices of the *k* target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

class dicee.models.Shallom (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Shallom: shallow neural model for relation prediction.

Represents each triple as the concatenation of head and tail entity embeddings and feeds it through a two-layer MLP to predict the relation. Designed for the RelationPrediction labelling form.

References

Demir et al., *A Shallow Neural Model for Relation Prediction*, ISWC 2021. <https://arxiv.org/abs/2101.09090>

name = 'Shallom'

shallom

get_embeddings () → Tuple[numpy.ndarray, None]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (numpy.ndarray) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (numpy.ndarray) – Shape (num_relations, embedding_dim).

forward_k_vs_all (*x*) → torch.FloatTensor

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises ValueError to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_triples (*x*) → torch.FloatTensor

Score a batch of triples by looking up relation scores from forward_k_vs_all.

Parameters

- **x** (torch.LongTensor) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class dicee.models.**Pyke**(args)Bases: *dicee.models.base_model.BaseKGE*

Pyke: Physical Embedding Model for Knowledge Graphs.

Scores a triple (h, r, t) based on the average pairwise distance between head-to-relation and relation-to-tail in embedding space:

$$f(h, r, t) = \text{margin} - (||h - r||_2 + ||r - t||_2) / 2$$

The model encodes geometric proximity between entities and the relations that connect them.

name = 'Pyke'**dist_func****margin** = 1.0**forward_triples**(x: torch.LongTensor) → torch.FloatTensor

Score a batch of triples using the Pyke distance formula.

Parameters**x** (torch.LongTensor) – Shape (batch_size, 3) integer tensor [head_idx, relation_idx, tail_idx].**Returns**

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

class dicee.models.**CoKEConfig**

Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.

block_size

Sequence length for transformer (3 for triples: head, relation, tail)

vocab_size

Total vocabulary size (num_entities + num_relations)

n_layer

Number of transformer layers

n_head

Number of attention heads per layer

n_embd

Embedding dimension (set to match model embedding_dim)

dropout

Dropout rate applied throughout the model

bias

Whether to use bias in linear layers

causal

Whether to use causal masking (False for bidirectional attention)

```

block_size: int = 3

vocab_size: int = None

n_layer: int = 6

n_head: int = 8

n_embd: int = None

dropout: float = 0.3

bias: bool = True

causal: bool = False

```

```

class dicee.models.CoKE(args, config: CoKEConfig = CoKEConfig())

```

Bases: `dicee.models.base_model.BaseKGE`

Contextualized Knowledge Graph Embedding (CoKE) model. Based on: <https://arxiv.org/pdf/1911.02168>.

CoKE uses a transformer encoder to learn contextualized representations of entities and relations. For link prediction, it predicts masked elements in (head, relation, tail) triples using bidirectional attention, similar to BERT's masked language modeling approach.

The model creates a sequence [head_emb, relation_emb, mask_emb], adds positional embeddings, and processes it through transformer layers to predict the tail entity.

```

name = 'CoKE'

```

```

config

```

```

pos_emb

```

```

mask_emb

```

```

blocks

```

```

ln_f

```

```

coke_dropout

```

```

forward_k_vs_all(x: torch.Tensor)

```

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

```

score(emb_h, emb_r, emb_t)

```

```

forward_k_vs_sample(x: torch.LongTensor, target_entity_idx: torch.LongTensor)

```

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

class dicee.models.**BaseKGE** (*args: dict*)

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

`hidden_dropout`

`loss_history = []`

`byte_pair_encoding`

`max_length_subword_tokens`

`block_size`

`defer_large_embeddings`

`init_entity_embeddings` (*embedding_dim*: *int* | *None* = *None*) → *None*

Create (or re-create) the entity embedding table.

This is the single place that honours `defer_large_embeddings`: when entity rows are sharded across FSDP ranks the table must stay *None* until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning `self.entity_embeddings` directly, so the deferral is never silently undone.

Parameters

`embedding_dim` (*Optional[int]*) – Width of the table. Defaults to `embedding_dim`.

`init_relation_embeddings` (*embedding_dim*: *int* | *None* = *None*) → *None*

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

`embedding_dim` (*Optional[int]*) – Width of the table. Defaults to `embedding_dim`.

`forward_byte_pair_encoded_k_vs_all` (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and *T* is `max_length_subword_tokens`.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple*[*torch.LongTensor*, *torch.LongTensor*])
→ *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from *self.args* with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the *args* dict. Falls back to sensible defaults for any missing key so that minimal *args* dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple*[*torch.LongTensor*, *torch.LongTensor*],
y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample*()
- (batch, 3) tensor → *forward_triples*()
- (batch, 2) tensor → *forward_k_vs_all*()
- BPE triple tensor → *forward_byte_pair_encoded_triple*()
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all*()

Parameters

• **x** (*torch.LongTensor* or *Tuple*[*torch.LongTensor*, *torch.LongTensor*]) – Either a plain index tensor or a (triple_idx, target_idx) tuple for sample-based labelling.

• **y_idx** (*torch.LongTensor*, optional) – Target entity indices used by *forward_k_vs_sample*(). Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.LongTensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample(*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation(idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation(indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation(x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[*torch.FloatTensor*, *torch.FloatTensor*]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[*numpy.ndarray*, *numpy.ndarray*]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

class dicee.models.ConEx (args)

Bases: *dicee.models.base_model.BaseKGE*

Convolutional ComplEx Knowledge Graph Embeddings

name = 'ConEx'

conv2d

fc_num_input

fc1

norm_fc1

bn_conv2d

feature_map_dropout

residual_convolution (C_1: Tuple[*torch.Tensor*, *torch.Tensor*],
C_2: Tuple[*torch.Tensor*, *torch.Tensor*]) → *torch.FloatTensor*

Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

forward_k_vs_all (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: *torch.Tensor*, *target_entity_idx*: *torch.Tensor*)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

class `dicee.models.AConEx` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Additive Convolutional ComplEx Knowledge Graph Embeddings

name = 'AConEx'

conv2d

fc_num_input

fc1

norm_fc1

bn_conv2d

feature_map_dropout

residual_convolution (*C_1*: *Tuple[torch.Tensor, torch.Tensor]*,
C_2: *Tuple[torch.Tensor, torch.Tensor]*) → *torch.FloatTensor*

Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

forward_k_vs_all (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_triples (*x*: *torch.Tensor*) → *torch.FloatTensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: *torch.Tensor*, *target_entity_idx*: *torch.Tensor*)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

torch.FloatTensor

class `dicee.models.Complex` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

name = 'Complex'

static score (*head_ent_emb*: *torch.FloatTensor*, *rel_ent_emb*: *torch.FloatTensor*,
tail_ent_emb: *torch.FloatTensor*)

```
static k_vs_all_score (emb_h: torch.FloatTensor, emb_r: torch.FloatTensor,
                      emb_E: torch.FloatTensor)
```

Parameters

- `emb_h`
- `emb_r`
- `emb_E`

```
forward_k_vs_all (x: torch.LongTensor) → torch.FloatTensor
```

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

```
forward_k_vs_sample (x: torch.LongTensor, target_entity_idx: torch.LongTensor)
```

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

`torch.FloatTensor`

```
class dicee.models.RotatE (args)
```

Bases: `dicee.models.base_model.BaseKGE`

RotatE: knowledge graph embedding by relational rotation in complex space.

Represents each entity as a complex vector in $\mathbb{C}^{(d/2)}$, obtained by splitting its d -dimensional embedding into a real half and an imaginary half. Each relation is represented by $d/2$ phase angles θ that define a unit-modulus rotation $r_i = e^{i\theta_i}$; since a phase angle needs only one real number, the relation embedding table is reinitialised at half of the entity embedding size. A true triple (h, r, t) should satisfy $h \circ r \approx t$ under element-wise complex multiplication, giving the score:

$$f(h, r, t) = \text{margin} - ||h \circ r - t||_2$$

Unlike TransE, RotatE can model symmetric, antisymmetric, inverse, and composition relation patterns.

References

Sun et al., *RotatE: Knowledge Graph Embedding by Relational Rotation in Complex Space*, ICLR 2019. <https://arxiv.org/abs/1902.10197>

```
name = 'RotatE'
```

```
margin = 6.0
```

```
half_dim
```

relation_embeddings

score (*head_ent_emb*: torch.FloatTensor, *rel_ent_emb*: torch.FloatTensor,
tail_ent_emb: torch.FloatTensor) → torch.FloatTensor

Score a batch of triples using the RotatE margin-distance formula.

Parameters

- **head_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **tail_ent_emb** (torch.FloatTensor) – Each has shape (batch_size, embedding_dim).
- **rel_ent_emb** (torch.FloatTensor) – Shape (batch_size, d/2) relation phase angles θ .

Returns

Shape (batch_size,) scores equal to $\text{margin} - ||h \circ r - t||_2$.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: torch.Tensor) → torch.FloatTensor

KvsAll forward pass: score head/relation against all entities.

Computes $\text{margin} - ||h \circ r - e||_2$ for every entity embedding e .

Parameters

x (torch.Tensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*x*: torch.Tensor, *target_entity_idx*: torch.Tensor) → torch.FloatTensor

KvsSample forward pass: score head/relation against a sampled entity subset.

Computes $\text{margin} - ||h \circ r - e||_2$ for each of the k sampled entities e .

Parameters

- **x** (torch.Tensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].
- **target_entity_idx** (torch.Tensor) – Shape (batch_size, k) indices of the k target entities per sample.

Returns

Shape (batch_size, k) score matrix.

Return type

torch.FloatTensor

class dicee.models.**BaseKGE** (*args*: dict)

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

loss

`selected_optimizer = None`

`normalizer_class = None`

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

`hidden_dropout`

`loss_history = []`

`byte_pair_encoding`

`max_length_subword_tokens`

`block_size`

`defer_large_embeddings`

`init_entity_embeddings(embedding_dim: int | None = None) → None`

Create (or re-create) the entity embedding table.

This is the single place that honours `defer_large_embeddings`: when entity rows are sharded across FSDP ranks the table must stay `None` until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning `self.entity_embeddings` directly, so the deferral is never silently undone.

Parameters

`embedding_dim` (*Optional[int]*) – Width of the table. Defaults to `embedding_dim`.

`init_relation_embeddings(embedding_dim: int | None = None) → None`

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

`embedding_dim` (*Optional[int]*) – Width of the table. Defaults to `embedding_dim`.

`forward_byte_pair_encoded_k_vs_all(x: torch.LongTensor) → torch.FloatTensor`

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

`x` (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and *T* is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

`torch.FloatTensor`

`forward_byte_pair_encoded_triple(x: Tuple[torch.LongTensor, torch.LongTensor])`

→ `torch.FloatTensor`

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

`x` (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

init_params_with_sanity_checking() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x*: `torch.LongTensor` | `Tuple[torch.LongTensor, torch.LongTensor]`,
 y_idx: `torch.LongTensor` = None) → `torch.FloatTensor`

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- `Tuple (x, y_idx)` → `forward_k_vs_sample()`
- `(batch, 3) tensor` → `forward_triples()`
- `(batch, 2) tensor` → `forward_k_vs_all()`
- BPE triple tensor → `forward_byte_pair_encoded_triple()`
- BPE pair tensor → `forward_byte_pair_encoded_k_vs_all()`

Parameters

- **x** (`torch.LongTensor` or `Tuple[torch.LongTensor, torch.LongTensor]`) – Either a plain index tensor or a `(triple_idx, target_idx)` tuple for sample-based labelling.
- **y_idx** (`torch.LongTensor`, optional) – Target entity indices used by `forward_k_vs_sample()`. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.LongTensor`) → `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

- **x** (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings() → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

class dicee.models.IdentityClass (*args=None*)

Bases: torch.nn.Module

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (`--normalization None`). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = None

__call__(*x*)

static forward(*x*)

dicee.models.quaternion_mul(*, *Q_1, Q_2*)

→ Tuple[torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor]

Perform quaternion multiplication :param Q_1: :param Q_2: :return:

dicee.models.quaternion_mul_with_unit_norm(*, *Q_1, Q_2*)

class dicee.models.QMult (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

`name = 'QMult'`

`explicit = True`

`quaternion_multiplication_followed_by_inner_product (h, r, t)`

Parameters

- **h** – shape: (*batch_dims*, dim) The head representations.
- **r** – shape: (*batch_dims*, dim) The head representations.
- **t** – shape: (*batch_dims*, dim) The tail representations.

Returns

Triple scores.

static `quaternion_normalizer (x: torch.FloatTensor) → torch.FloatTensor`

Normalize the length of relation vectors, if the forward constraint has not been applied yet.

Absolute value of a quaternion

$$|a + bi + cj + dk| = \sqrt{a^2 + b^2 + c^2 + d^2}$$

L2 norm of quaternion vector:

$$\|x\|^2 = \sum_{i=1}^d |x_i|^2 = \sum_{i=1}^d (x_i.re^2 + x_i.im_1^2 + x_i.im_2^2 + x_i.im_3^2)$$

Parameters

x – The vector.

Returns

The normalized vector.

score (*head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor, tail_ent_emb: torch.FloatTensor*)

k_vs_all_score (*bpe_head_ent_emb, bpe_rel_ent_emb, E*)

Parameters

- **bpe_head_ent_emb**
- **bpe_rel_ent_emb**
- **E**

forward_k_vs_all (*x*)

Parameters

x

forward_k_vs_sample (*x, target_entity_idx*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e., [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,| Entities|)

class `dicee.models.ConvQ (args)`

Bases: `dicee.models.base_model.BaseKGE`

Convolutional Quaternion Knowledge Graph Embeddings

```

name = 'ConvQ'

entity_embeddings

relation_embeddings

conv2d

fc_num_input

fc1

bn_conv1

bn_conv2

feature_map_dropout

residual_convolution( $Q_1, Q_2$ )

forward_triples(indexed_triple: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
         $\mathbf{x}$  (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

forward_k_vs_all(x: torch.Tensor)
    Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] =>
    [0.0,0.1,...,0.8], shape=> (1, Entities) Given a batch of head entities and relations => shape (size of batch,|
    Entities)

class dicee.models.AConvQ(args)
    Bases: dicee.models.base_model.BaseKGE

    Additive Convolutional Quaternion Knowledge Graph Embeddings

    name = 'AConvQ'

    entity_embeddings

    relation_embeddings

    conv2d

    fc_num_input

    fc1

    bn_conv1

    bn_conv2

    feature_map_dropout

```

residual_convolution(*Q_1*, *Q_2*)

forward_triples(*indexed_triple*: torch.Tensor) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all(*x*: torch.Tensor)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch, Entities)

class dicee.models.**BaseKGE**(*args*: dict)

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (dict) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

```

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

defer_large_embeddings

```

init_entity_embeddings (*embedding_dim: int | None = None*) → None

Create (or re-create) the entity embedding table.

This is the single place that honours *defer_large_embeddings*: when entity rows are sharded across FSDP ranks the table must stay *None* until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning *self.entity_embeddings* directly, so the deferral is never silently undone.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

init_relation_embeddings (*embedding_dim: int | None = None*) → None

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

forward_byte_pair_encoded_k_vs_all (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*)

→ *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from *self.args* with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the *args* dict. Falls back to sensible defaults for any missing key so that minimal *args* dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple[torch.LongTensor, torch.LongTensor]*,

y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample*()
- (batch, 3) tensor → *forward_triples*()
- (batch, 2) tensor → *forward_k_vs_all*()
- BPE triple tensor → *forward_byte_pair_encoded_triple*()
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all*()

Parameters

• **x** (*torch.LongTensor* or *Tuple[torch.LongTensor, torch.LongTensor]*) – Either a plain index tensor or a (*triple_idx*, *target_idx*) tuple for sample-based labelling.

• **y_idx** (*torch.LongTensor*, optional) – Target entity indices used by *forward_k_vs_sample*(). Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.LongTensor`) → `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

x (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

`torch.FloatTensor`

get_triple_representation (idx_hrt)

→ Tuple[`torch.FloatTensor`, `torch.FloatTensor`, `torch.FloatTensor`]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

`torch.FloatTensor`

get_head_relation_representation (indexed_triple)

→ Tuple[`torch.FloatTensor`, `torch.FloatTensor`]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

class dicee.models.IdentityClass (*args=None*)

Bases: torch.nn.Module

No-op normalisation / dropout placeholder.

Used whenever no normalisation layer is requested (--normalization None). All inputs are returned unchanged so that the rest of the model code does not need conditional checks around normalisation calls.

args = None

```

__call__(x)

static forward(x)

```

`dicee.models.octonion_mul(*, O_1, O_2)`

`dicee.models.octonion_mul_norm(*, O_1, O_2)`

class `dicee.models.OMult` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

```

name = 'OMult'

static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                           emb_rel_e5, emb_rel_e6, emb_rel_e7)

score(head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor,
      tail_ent_emb: torch.FloatTensor)

k_vs_all_score(bpe_head_ent_emb, bpe_rel_ent_emb, E)

forward_k_vs_all(x)

```

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples, i.e., `[score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8]`, shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch, |Entities|)

class `dicee.models.ConvO` (*args: dict*)

Bases: `dicee.models.base_model.BaseKGE`

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from `BaseKGELightning` and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

```

name = 'ConvO'

conv2d

fc_num_input

fc1

bn_conv2d

norm_fc1

feature_map_dropout

static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                           emb_rel_e5, emb_rel_e6, emb_rel_e7)

residual_convolution(O_1, O_2)

forward_triples(x: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

forward_k_vs_all(x: torch.Tensor)
    Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] =>
    [0.0,0.1,...,0.8], shape=> (1, Entities) Given a batch of head entities and relations => shape (size of batch,|
    Entities)

class dicee.models.AConvO(args: dict)
    Bases: dicee.models.base_model.BaseKGE

    Additive Convolutional Octonion Knowledge Graph Embeddings

    name = 'AConvO'

    conv2d

    fc_num_input

    fc1

    bn_conv2d

    norm_fc1

    feature_map_dropout

    static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                               emb_rel_e5, emb_rel_e6, emb_rel_e7)

    residual_convolution(O_1, O_2)

```

forward_triples (*x*: *torch.Tensor*) → *torch.Tensor*

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: *torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities)

class *dicee.models.Keci* (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Keci: Knowledge Graph Embedding via Clifford Algebra.

Embeds entities and relations as multi-vectors in the Clifford algebra $Cl_{\{p,q\}}(R^d)$ and scores triples via the Clifford product. The algebra is parameterised by two non-negative integers *p* and *q*:

- *p* = 0, *q* = 0 — reduces to a standard bilinear (DistMult-like) model.
- *p* = 0, *q* = 1 — equivalent to ComplEx.
- Larger *p* and *q* capture higher-order geometric interactions.

The embedding dimension must satisfy `embedding_dim % (p + q + 1) == 0`; the resulting quotient is stored as `self.r`.

Parameters

args (*dict*) – Configuration dictionary. Recognised keys (beyond those in *BaseKGE*): *p* (int, default 0) and *q* (int, default 0).

References

Demir et al., *Clifford Embeddings — A Generalized Approach for Embedding in Normed Algebras*, ECML 2023.

name = 'Keci'

p

q

r

requires_grad_for_interactions = True

compute_sigma_pp (*hp*, *rp*)

Compute $\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{k=i+1}^p (h_i r_k - h_k r_i) e_i e_k$

σ_{pp} captures the interactions between along *p* bases For instance, let *p* *e*₁, *e*₂, *e*₃, we compute interactions between *e*₁ *e*₂, *e*₁ *e*₃, and *e*₂ *e*₃ This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

```
sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

```
e2e1, e2e2, e2e3, e3e1, e3e2, e3e3
```

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

```
compute_sigma_qq (hq, rq)
```

Compute $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_{j,r,k} - h_{k,r,j}) e_j e_k \sigma_{qj}$ captures the interactions between along q bases. For instance, let $q = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$. This can be implemented with a nested two for loops

```
results = [] for j in range(q - 1):
```

```
    for k in range(j + 1, q):
```

```
        results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])
```

```
sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

```
e2e1, e2e2, e2e3, e3e1, e3e2, e3e3
```

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

```
compute_sigma_pq (*, hp, hq, rp, rq)
```

```
sum_{i=1}^p sum_{j=p+1}^{p+q} (h_{i,r,j} - h_{j,r,i}) e_i e_j
```

```
results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):
```

```
    for j in range(q):
```

```
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
```

```
print(sigma_pq.shape)
```

```
apply_coefficients (hp, hq, rp, rq)
```

Multiplying a base vector with its scalar coefficient

```
clifford_multiplication (h0, hp, hq, r0, rp, rq)
```

Compute our CL multiplication

$$h = h_0 + \sum_{i=1}^p h_i e_i + \sum_{j=p+1}^{p+q} h_j e_j \quad r = r_0 + \sum_{i=1}^p r_i e_i + \sum_{j=p+1}^{p+q} r_j e_j$$

$$e_i^2 = +1 \text{ for } i \leq p, e_j^2 = -1 \text{ for } p < j \leq p+q, e_i e_j = -e_j e_i \text{ for } i$$

e_j

$h = \sigma_0 + \sigma_p + \sigma_q + \sigma_{pp} + \sigma_q + \sigma_{pq}$ where

(1) $\sigma_0 = h_0 r_0 + \sum_{i=1}^p (h_0 r_i - h_i r_0) e_i - \sum_{j=p+1}^{p+q} (h_j r_j - h_r j) e_j$

(2) $\sigma_p = \sum_{i=1}^p (h_0 r_i + h_i r_0) e_i$

(3) $\sigma_q = \sum_{j=p+1}^{p+q} (h_0 r_j + h_j r_0) e_j$

(4) $\sigma_{pp} = \sum_{i=1}^p \sum_{k=i+1}^p (h_i r_k - h_k r_i) e_i e_k$

(5) $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_j r_k - h_k r_j) e_j e_k$

(6) $\sigma_{pq} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$

construct_cl_multivector (*x*: torch.FloatTensor, *r*: int, *p*: int, *q*: int)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Split a flat embedding vector into the three Clifford components.

Given an embedding *x* of dimension $d = r + r*p + r*q$, returns the scalar part *a0*, the *p*-blade part *ap*, and the *q*-blade part *aq*.

Parameters

- **x** (torch.FloatTensor) – Shape (batch_size, d).
- **r** (int) – Scalar block size (embedding_dim // (p + q + 1)).
- **p** (int) – Number of positive-signature basis elements.
- **q** (int) – Number of negative-signature basis elements.

Returns

- **a0** (torch.FloatTensor) – Shape (batch_size, r) — scalar (grade-0) part.
- **ap** (torch.FloatTensor) – Shape (batch_size, r, p) — positive-blade part.
- **aq** (torch.FloatTensor) – Shape (batch_size, r, q) — negative-blade part.

forward_k_vs_with_explicit (*x*: torch.Tensor) → torch.FloatTensor

KvsAll scoring using an explicit loop over sigma_pp/qq/pq terms.

Functionally equivalent to `forward_k_vs_all()` but computes the higher-order interaction terms (sigma_pp, sigma_qq, sigma_pq) with explicit nested loops rather than einsum contractions. Kept for reference and correctness verification.

Parameters

x (torch.Tensor) – Shape (batch_size, 2) integer tensor [head_idx, relation_idx].

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

k_vs_all_score (*bpe_head_ent_emb*: torch.FloatTensor, *bpe_rel_ent_emb*: torch.FloatTensor, *E*: torch.FloatTensor) → torch.FloatTensor

Compute Clifford-product scores for a head/relation batch vs. all entities.

Decomposes the head-entity and relation embeddings into Clifford multi-vectors, performs the $Cl_{\{p,q\}}$ product, and inner-products the result against the entity embedding matrix *E*.

Parameters

- **bpe_head_ent_emb** (torch.FloatTensor) – Head-entity embeddings, shape (batch_size, embedding_dim).
- **bpe_rel_ent_emb** (torch.FloatTensor) – Relation embeddings, shape (batch_size, embedding_dim).
- **E** (torch.FloatTensor) – All entity embeddings, shape (num_entities, embedding_dim).

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_all (*x*: torch.Tensor) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations $\mathbb{R}^d$.
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— *x*: torch.LongTensor with (n,2) shape :rtype: torch.FloatTensor with (n, **IEI**) shape

construct_batch_selected_cl_multivector (*x*: torch.FloatTensor, *r*: int, *p*: int, *q*: int)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Split a batched, *k*-selected embedding tensor into Clifford components.

A variant of `construct_cl_multivector()` for tensors that have an extra *k* dimension (e.g. when scoring against *k* sampled targets).

Parameters

- **x** (torch.FloatTensor) – Shape (batch_size, *k*, *d*).
- **r** (int) – Scalar block size.
- **p** (int) – Number of positive-signature basis elements.
- **q** (int) – Number of negative-signature basis elements.

Returns

- **a0** (torch.FloatTensor) – Shape (batch_size, *k*, *r*).
- **ap** (torch.FloatTensor) – Shape (batch_size, *k*, *r*, *p*).
- **aq** (torch.FloatTensor) – Shape (batch_size, *k*, *r*, *q*).

forward_k_vs_sample (*x*: torch.LongTensor, *target_entity_idx*: torch.LongTensor) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,2) shape

target_entity_idx: torch.LongTensor with (n, *k*) shape *k* denotes the selected number of examples.

rtype

torch.FloatTensor with (n, *k*) shape

score (*h*, *r*, *t*)

forward_triples (*x*: torch.Tensor) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,3) shape

rtype

torch.FloatTensor with (n) shape


```
class dicee.models.CKeci (args)
```

Bases: *Keci*

Without learning dimension scaling

```
name = 'CKeci'
```

```
requires_grad_for_interactions = False
```

```
class dicee.models.DeCaL (args)
```

Bases: *dicee.models.base_model.BaseKGE*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- *forward_triples()* — score a batch of (h, r, t) triples.
- *forward_k_vs_all()* — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

```
name = 'DeCaL'
```

```
p
```

```
q
```

```
r
```

```
re
```

```
forward_triples (x: torch.Tensor) → torch.FloatTensor
```

Parameter

x: torch.LongTensor with (n,) shape

rtype

torch.FloatTensor with (n) shape

```
cl_pqr (a: torch.tensor) → torch.tensor
```

Input: tensor(batch_size, emb_dim) → output: tensor with 1+p+q+r components with size (batch_size, emb_dim/(1+p+q+r)) each.

1) takes a tensor of size (batch_size, emb_dim), split it into 1 + p + q + r components, hence 1+p+q+r must be a divisor of the emb_dim. 2) Return a list of the 1+p+q+r components vectors, each are tensors of size (batch_size, emb_dim/(1+p+q+r))

```
compute_sigmas_single (list_h_emb, list_r_emb, list_t_emb)
```

here we compute all the sums with no others vectors interaction taken with the scalar product with t, that is,

$$s_0 = h_0 r_0 t_0 s_1 = \sum_{i=1}^p h_i r_i t_0 s_2 = \sum_{j=p+1}^{p+q} h_j r_j t_0 s_3 = \sum_{i=1}^q (h_0 r_i t_i + h_i r_0 t_i) s_4 = \sum_{i=p+1}^{p+q} (h_0 r_i t_i + h_i r_0 t_i) s_5 = \sum_{i=p+q+1}^{p+q+r} (h_0 r_i t_i + h_i r_0 t_i)$$

and return:

$$\sigma_0 t = \sigma_0 \cdot t_0 = s_0 + s_1 - s_2 s_3, s_4 \text{ and } s_5$$

compute_sigmas_multivect (*list_h_emb, list_r_emb*)

Here we compute and return all the sums with vectors interaction for the same and different bases.

For same bases vectors interaction we have

$$\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (h_i r_{i'} - h_{i'} r_i) (\text{model the interactions between } e_i \text{ and } e_{i'} \text{ for } 1 \leq i, i' \leq p) \sigma_{qq} = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (h_j r_{j'} - h_{j'} r_j)$$

For different base vector interactions, we have

$$\sigma_{pq} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) (\text{interactions between } e_i \text{ and } e_j \text{ for } 1 \leq i \leq p \text{ and } p+1 \leq j \leq p+q) \sigma_{pr} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i)$$

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations
- (2) Construct head entity and relation embeddings according to $\text{Cl}_{\{p,q,r\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— *x*: torch.LongTensor with (n,) shape :rtype: torch.FloatTensor with (n, **IE**) shape

apply_coefficients (*h0, hp, hq, hk, r0, rp, rq, rk*)

Multiplying a base vector with its scalar coefficient

construct_cl_multivector (*x: torch.FloatTensor, re: int, p: int, q: int, r: int*)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Construct a batch of multivectors $\text{Cl}_{\{p,q,r\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,d) shape

returns

- **a0** (*torch.FloatTensor*)
- **ap** (*torch.FloatTensor*)
- **aq** (*torch.FloatTensor*)
- **ar** (*torch.FloatTensor*)

compute_sigma_pp (*hp, rp*)

Compute .. math:

$$\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (x_{iy_{i'}} - x_{i'y_i})$$

$\text{sigma_}\{pp\}$ captures the interactions between along p bases For instance, let $p \in \{e_1, e_2, e_3\}$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

```
results = [] for i in range(p - 1):
    for k in range(i + 1, p):
        results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3,$

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_qq (hq, rq)

Compute

$$\sigma_{q,q}^* = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (x_j y_{j'} - x_{j'} y_j) Eq.16$$

$\text{sigma_}\{q\}$ captures the interactions between along q bases For instance, let $q \in \{e_1, e_2, e_3\}$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

```
results = [] for j in range(q - 1):
    for k in range(j + 1, q):
        results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3,$

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_rr (hk, rk)

$$\sigma_{r,r}^* = \sum_{k=p+q+1}^{p+q+r-1} \sum_{k'=k+1}^p (x_k y_{k'} - x_{k'} y_k)$$

compute_sigma_pq ($*, hp, hq, rp, rq$)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

```
results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):
```

```
    for j in range(q):
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
```

```
print(sigma_pq.shape)
```

compute_sigma_pr(* , hp, hk, rp, rk)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

compute_sigma_qr(* , hq, hk, rq, rk)

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

class dicee.models.**KeciTransformer**(args)

Bases: *Keci*

Keci with Transformer architecture.

Concatenates h0, hp, hq, r0, rp, rq into a single embedding vector and processes through transformer.

name = 'KeciTransformer'

use_clifford_mul

seq_len = 1

transformer

lm_head

forward_k_vs_all(x: torch.Tensor) → torch.FloatTensor

Kvsall training

Parameter

x: torch.LongTensor with (n,2) shape

rtype

torch.FloatTensor with (n, **IEI**) shape

class dicee.models.**BaseKGE**(args: dict)

Bases: *BaseKGELightning*

Base class for all Knowledge Graph Embedding models.

Inherits the Lightning training loop from *BaseKGELightning* and adds the embedding tables, normalisation / dropout layers, and the routing logic that dispatches `forward()` calls to the appropriate scoring method.

Sub-classes must implement at minimum:

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

args (*dict*) – Flat configuration dictionary produced by `vars(argparse.Namespace)`. Required keys: `embedding_dim`, `num_entities`, `num_relations`, `learning_rate` (or `lr`), `optim`, `scoring_technique`.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

loss

`selected_optimizer = None`

`normalizer_class = None`

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

`hidden_dropout`

`loss_history = []`

byte_pair_encoding

max_length_subword_tokens

block_size

defer_large_embeddings

init_entity_embeddings (*embedding_dim*: int | None = None) → None

Create (or re-create) the entity embedding table.

This is the single place that honours *defer_large_embeddings*: when entity rows are sharded across FSDP ranks the table must stay None until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning `self.entity_embeddings` directly, so the deferral is never silently undone.

Parameters

embedding_dim (*Optional*[int]) – Width of the table. Defaults to *embedding_dim*.

init_relation_embeddings (*embedding_dim*: int | None = None) → None

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

embedding_dim (*Optional*[int]) – Width of the table. Defaults to *embedding_dim*.

forward_byte_pair_encoded_k_vs_all (*x*: torch.LongTensor) → torch.FloatTensor

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: Tuple[torch.LongTensor, torch.LongTensor])
→ torch.FloatTensor

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from `self.args` with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the `args` dict. Falls back to sensible defaults for any missing key so that minimal `args` dicts (e.g. for unit tests) are still valid.

forward (*x*: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
 y_idx: torch.LongTensor = None) → torch.FloatTensor

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- Tuple (*x*, *y_idx*) → `forward_k_vs_sample()`
- (batch, 3) tensor → `forward_triples()`
- (batch, 2) tensor → `forward_k_vs_all()`
- BPE triple tensor → `forward_byte_pair_encoded_triple()`
- BPE pair tensor → `forward_byte_pair_encoded_k_vs_all()`

Parameters

- **x** (torch.LongTensor or Tuple[torch.LongTensor, torch.LongTensor]) – Either a plain index tensor or a (triple_idx, target_idx) tuple for sample-based labelling.
- **y_idx** (torch.LongTensor, optional) – Target entity indices used by `forward_k_vs_sample()`. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

torch.FloatTensor

forward_triples (*x*: torch.LongTensor) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

- **x** (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

torch.FloatTensor

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where k is the number of sampled target entities.

Return type

torch.FloatTensor

get_triple_representation (idx_hrt)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (torch.LongTensor) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_head_relation_representation (indexed_triple)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (torch.LongTensor) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (torch.LongTensor) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (x: torch.LongTensor)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings() → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

dicee.models.logger

class dicee.models.PykeenKGE (*args: dict*)

Bases: *dicee.models.base_model.BaseKGE*

A class for using knowledge graph embedding models implemented in Pykeen

Notes: Pykeen_DistMult: C Pykeen_ComplEx: Pykeen_QuatE: Pykeen_MuRE: Pykeen_CP: Pykeen_HolE: Pykeen_HolE: Pykeen_HolE: Pykeen_TransD: Pykeen_TransE: Pykeen_TransF: Pykeen_TransH: Pykeen_TransR:

model_kwargs

name

model

loss_history = []

args

entity_embeddings = None

relation_embeddings = None

forward_k_vs_all (*x: torch.LongTensor*)

=> Explicit version by this we can apply bn and dropout

(1) Retrieve embeddings of heads and relations + apply Dropout & Normalization if given. h, r = self.get_head_relation_representation(x) # (2) Reshape (1). if self.last_dim > 0:

h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim, self.last_dim)

(3) Reshape all entities. if self.last_dim > 0:

t = self.entity_embeddings.weight.reshape(self.num_entities, self.embedding_dim, self.last_dim)

else:

t = self.entity_embeddings.weight

```
# (4) Call the score_t from interactions to generate triple scores. return self.interaction.score_t(h=h, r=r,
all_entities=t, slice_size=1)
```

```
forward_triples (x: torch.LongTensor) → torch.FloatTensor
```

```
# => Explicit version by this we can apply bn and dropout
```

```
# (1) Retrieve embeddings of heads, relations and tails and apply Dropout & Normalization if given. h, r, t =
self.get_triple_representation(x) # (2) Reshape (1). if self.last_dim > 0:
```

```
    h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim,
    self.last_dim) t = t.reshape(len(x), self.embedding_dim, self.last_dim)
```

```
# (3) Compute the triple score return self.interaction.score(h=h, r=r, t=t, slice_size=None, slice_dim=0)
```

```
abstractmethod forward_k_vs_sample (x: torch.LongTensor, target_entity_idx)
```

```
Score a (head, relation) batch against a sampled subset of entities.
```

```
Used by KvsSample and 1vsSample datasets. Sub-classes that support sample-based labelling must over-
ride this method.
```

Returns

```
Shape (batch_size, k) score matrix where k is the number of sampled target entities.
```

Return type

```
torch.FloatTensor
```

```
class dicee.models.BaseKGE (args: dict)
```

```
Bases: BaseKGELightning
```

```
Base class for all Knowledge Graph Embedding models.
```

```
Inherits the Lightning training loop from BaseKGELightning and adds the embedding tables, normalisation /
dropout layers, and the routing logic that dispatches forward() calls to the appropriate scoring method.
```

```
Sub-classes must implement at minimum:
```

- `forward_triples()` — score a batch of (h, r, t) triples.
- `forward_k_vs_all()` — score a (h, r) batch against every entity.

Parameters

```
args (dict) – Flat configuration dictionary produced by vars(argparse.Namespace). Re-
quired keys: embedding_dim, num_entities, num_relations, learning_rate (or lr),
optim, scoring_technique.
```

```
args
```

```
embedding_dim = None
```

```
num_entities = None
```

```
num_relations = None
```

```
num_tokens = None
```

```
learning_rate = None
```

```
apply_unit_norm = None
```

```
input_dropout_rate = None
```

```

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

defer_large_embeddings

```

init_entity_embeddings (*embedding_dim*: *int* | *None* = *None*) → *None*

Create (or re-create) the entity embedding table.

This is the single place that honours *defer_large_embeddings*: when entity rows are sharded across FSDP ranks the table must stay *None* until the trainer allocates the sharded adapter. Subclasses that need an entity table of a non-default width must call this method instead of assigning `self.entity_embeddings` directly, so the deferral is never silently undone.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

init_relation_embeddings (*embedding_dim*: *int* | *None* = *None*) → *None*

Create (or re-create) the relation embedding table.

Relation tables are never deferred - they are small enough to be replicated on every rank.

Parameters

embedding_dim (*Optional[int]*) – Width of the table. Defaults to *embedding_dim*.

forward_byte_pair_encoded_k_vs_all (*x*: *torch.LongTensor*) → *torch.FloatTensor*

KvsAll scoring for BPE-encoded head entities and relations.

Retrieves subword-unit embeddings for the head entity and relation, reduces them to fixed-size vectors via a linear projection, then scores against all BPE entity embeddings.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) BPE token indices where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

Shape (batch_size, num_bpe_entities) score matrix.

Return type

torch.FloatTensor

forward_byte_pair_encoded_triple (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*)

→ *torch.FloatTensor*

NegSample scoring for BPE-encoded (head, relation, tail) triples.

Retrieves subword-unit embeddings for all three elements and reduces them to fixed-size vectors via a linear projection before computing the triple score.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) BPE token indices.

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

init_params_with_sanity_checking() → None

Populate model hyper-parameters from *self.args* with safe defaults.

Reads embedding dimension, learning rate, dropout rates, normalisation strategy, optimizer name, and parameter initialisation scheme from the *args* dict. Falls back to sensible defaults for any missing key so that minimal *args* dicts (e.g. for unit tests) are still valid.

forward (*x*: *torch.LongTensor* | *Tuple[torch.LongTensor, torch.LongTensor]*,

y_idx: *torch.LongTensor* = None) → *torch.FloatTensor*

Route the forward pass to the appropriate scoring method.

Inspects the shape and type of *x* to decide which low-level scorer to call:

- *Tuple* (*x*, *y_idx*) → *forward_k_vs_sample()*
- (batch, 3) tensor → *forward_triples()*
- (batch, 2) tensor → *forward_k_vs_all()*
- BPE triple tensor → *forward_byte_pair_encoded_triple()*
- BPE pair tensor → *forward_byte_pair_encoded_k_vs_all()*

Parameters

• **x** (*torch.LongTensor* or *Tuple[torch.LongTensor, torch.LongTensor]*) – Either a plain index tensor or a (triple_idx, target_idx) tuple for sample-based labelling.

• **y_idx** (*torch.LongTensor*, optional) – Target entity indices used by *forward_k_vs_sample()*. Ignored when *x* is a plain tensor.

Returns

Score tensor whose shape depends on the selected scorer.

Return type

`torch.FloatTensor`

forward_triples (*x*: `torch.LongTensor`) → `torch.Tensor`

Score a batch of (head, relation, tail) index triples.

Parameters

x (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

`torch.FloatTensor`

forward_k_vs_all (*args, **kwargs)

Score a (head, relation) batch against every entity.

Sub-classes must override this method. The default implementation raises `ValueError` to make missing overrides obvious at runtime.

Returns

Shape (batch_size, num_entities) score matrix.

Return type

`torch.FloatTensor`

forward_k_vs_sample (*args, **kwargs)

Score a (head, relation) batch against a sampled subset of entities.

Used by `KvsSample` and `1vsSample` datasets. Sub-classes that support sample-based labelling must override this method.

Returns

Shape (batch_size, k) score matrix where *k* is the number of sampled target entities.

Return type

`torch.FloatTensor`

get_triple_representation (idx_hrt)

→ Tuple[`torch.FloatTensor`, `torch.FloatTensor`, `torch.FloatTensor`]

Retrieve and normalise embedding vectors for a triple index batch.

Parameters

idx_hrt (`torch.LongTensor`) – Shape (batch_size, 3) integer tensor with columns [head_idx, relation_idx, tail_idx].

Returns

head_ent_emb, rel_ent_emb, tail_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

`torch.FloatTensor`

get_head_relation_representation (indexed_triple)

→ Tuple[`torch.FloatTensor`, `torch.FloatTensor`]

Retrieve and normalise embedding vectors for head entities and relations.

Parameters

indexed_triple (*torch.LongTensor*) – Shape (batch_size, 2) integer tensor with columns [head_idx, relation_idx].

Returns

head_ent_emb, rel_ent_emb – Each has shape (batch_size, embedding_dim) after applying the configured dropout and normalisation.

Return type

torch.FloatTensor

get_sentence_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Retrieve BPE subword-unit embeddings for a batch of triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3, T) where T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb, tail_emb – Each has shape (batch_size, T, embedding_dim).

Return type

torch.FloatTensor

get_bpe_head_and_relation_representation (*x: torch.LongTensor*)

→ Tuple[torch.FloatTensor, torch.FloatTensor]

Retrieve unit-normalised BPE embeddings for head entities and relations.

Each entity/relation is represented as a sequence of T subword tokens. Their token embeddings are L2-normalised across the sequence dimension so that the resulting matrix has unit Frobenius norm.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 2, T) where dim 1 indexes [head, relation] and T is max_length_subword_tokens.

Returns

head_ent_emb, rel_emb – Each has shape (batch_size, T, embedding_dim), L2-normalised over the (T, D) dimensions.

Return type

torch.FloatTensor

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

Return the entity and relation embedding matrices as numpy arrays.

Returns

- **entity_embeddings** (*numpy.ndarray*) – Shape (num_entities, embedding_dim).
- **relation_embeddings** (*numpy.ndarray*) – Shape (num_relations, embedding_dim).

dicee.models.logger

class dicee.models.FMult (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Learning Knowledge Neural Graphs

name = 'FMult'

```

entity_embeddings

relation_embeddings

k

num_sample = 50

gamma

roots

weights

compute_func(weights: torch.FloatTensor, x) → torch.FloatTensor

chain_func(weights, x: torch.FloatTensor)

forward_triples(idx_triple: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

```

class dicee.models.**GFMult**(args)

Bases: *dicee.models.base_model.BaseKGE*

Learning Knowledge Neural Graphs

```

name = 'GFMult'

entity_embeddings

relation_embeddings

k

num_sample = 250

roots

weights

compute_func(weights: torch.FloatTensor, x) → torch.FloatTensor

chain_func(weights, x: torch.FloatTensor)

forward_triples(idx_triple: torch.Tensor) → torch.Tensor
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

```

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

```
class dicee.models.FMult2(args)
```

Bases: *[dicee.models.base_model.BaseKGE](#)*

Learning Knowledge Neural Graphs

name = 'FMult2'

n_layers = 3

k

n = 50

score_func = 'compositional'

discrete_points

entity_embeddings

relation_embeddings

build_func(Vec)

build_chain_funcs(list_Vec)

compute_func(W, b, x) → torch.FloatTensor

function(list_W, list_b)

trapezoid(list_W, list_b)

forward_triples(idx_triple: torch.Tensor) → torch.Tensor

Score a batch of (head, relation, tail) index triples.

Parameters

x (*torch.LongTensor*) – Shape (batch_size, 3) integer tensor where each row is [head_idx, relation_idx, tail_idx].

Returns

Shape (batch_size,) triple scores.

Return type

torch.FloatTensor

```
class dicee.models.LFMult1(args)
```

Bases: *[dicee.models.base_model.BaseKGE](#)*

Embedding with trigonometric functions. We represent all entities and relations in the complex number space as: $f(x) = \sum_{k=0}^{k=d-1} w_k e^{kix}$. and use the three differents scoring function as in the paper to evaluate the score

name = 'LFMult1'

entity_embeddings


```

relation_embeddings

forward_triples (idx_triple)
    Score a batch of (head, relation, tail) index triples.

    Parameters
        x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
        [head_idx, relation_idx, tail_idx].

    Returns
        Shape (batch_size,) triple scores.

    Return type
        torch.FloatTensor

tri_score (h, r, t)

vtp_score (h, r, t)

class dicee.models.LFMult (args)
    Bases: dicee.models.base_model.BaseKGE

    Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:  $f(x) = \sum_{i=0}^{d-1} a_k x^{i \bmod d}$  and use the three differents scoring function as in the paper to evaluate the score.
    We also consider combining with Neural Networks.

    name = 'LFMult'

    entity_embeddings

    relation_embeddings

    degree

    m

    x_values

    forward_triples (idx_triple)
        Score a batch of (head, relation, tail) index triples.

        Parameters
            x (torch.LongTensor) – Shape (batch_size, 3) integer tensor where each row is
            [head_idx, relation_idx, tail_idx].

        Returns
            Shape (batch_size,) triple scores.

        Return type
            torch.FloatTensor

    construct_multi_coeff (x)

    poly_NN (x, coefh, coefr, coeft)
        Constructing a 2 layers NN to represent the embeddings.  $h = \text{sigma}(w_h^T x + b_h)$ ,  $r = \text{sigma}(w_r^T x + b_r)$ ,
         $t = \text{sigma}(w_t^T x + b_t)$ 

    linear (x, w, b)

```

scalar_batch_NN (*a, b, c*)

element wise multiplication between a,b and c: Inputs : a, b, c ==> torch.tensor of size batch_size x m x d
Output : a tensor of size batch_size x d

tri_score (*coeff_h, coeff_r, coeff_t*)

this part implement the trilinear scoring techniques:

$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \frac{a_i b_j c_k}{1+(i+j+k)d}$$

1. generate the range for i,j and k from [0 d-1]
2. perform $\frac{a_i b_j c_k}{1+(i+j+k)d}$ in parallel for every batch
3. take the sum over each batch

vtp_score (*h, r, t*)

this part implement the vector triple product scoring techniques:

$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \frac{a_i c_j b_k - b_i c_j a_k}{(1+(i+j)d)(1+k)}$$

1. generate the range for i,j and k from [0 d-1]
2. Compute the first and second terms of the sum
3. Multiply with then denominator and take the sum
4. take the sum over each batch

comp_func (*h, r, t*)

this part implement the function composition scoring techniques: i.e. score = <hor, t>

polynomial (*coeff, x, degree*)

This function takes a matrix tensor of coefficients (coeff), a tensor vector of points x and range of integer [0,1,...d] and return a vector tensor (coeff[0][0] + coeff[0][1]x +...+ coeff[0][d]x^d,

$$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d)$$

pop (*coeff, x, degree*)

This function allow us to evaluate the composition of two polynomes without for loops :) it takes a matrix tensor of coefficients (coeff), a matrix tensor of points x and range of integer [0,1,...d]

and return a tensor (coeff[0][0] + coeff[0][1]x +...+ coeff[0][d]x^d,

$$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d)$$

class dicee.models.DualE (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Dual Quaternion Knowledge Graph Embeddings (<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

name = 'DualE'

entity_embeddings

relation_embeddings

num_ent = None

kvsall_score (*e_1_h, e_2_h, e_3_h, e_4_h, e_5_h, e_6_h, e_7_h, e_8_h, e_1_t, e_2_t, e_3_t, e_4_t, e_5_t, e_6_t, e_7_t, e_8_t, r_1, r_2, r_3, r_4, r_5, r_6, r_7, r_8*) → torch.tensor

KvsAll scoring function

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_triples (*idx_triple: torch.tensor*) → torch.tensor

Negative Sampling forward pass:

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_k_vs_all (*x*)

KvsAll forward pass

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

T (*x: torch.tensor*) → torch.tensor

Transpose function

Input: Tensor with shape (nxm) Output: Tensor with shape (mxn)

dicee.query_generator

Attributes

logger

Classes

QueryGenerator

Module Contents

`dicee.query_generator.logger`

class `dicee.query_generator.QueryGenerator` (*train_path: str, val_path: str, test_path: str, ent2id: Dict = None, rel2id: Dict = None, seed: int = 1, gen_valid: bool = False, gen_test: bool = True*)

```

train_path

val_path

test_path

gen_valid = False

gen_test = True

seed = 1

max_ans_num = 1000000.0

mode

ent2id = None

rel2id: Dict = None

ent_in: Dict

ent_out: Dict

query_name_to_struct

list2tuple(list_data)

tuple2list(x: List | Tuple) → List | Tuple
    Convert a nested tuple to a nested list.

set_global_seed(seed: int)
    Set seed

construct_graph(paths: List[str]) → Tuple[Dict, Dict]
    Construct graph from triples Returns dicts with incoming and outgoing edges

fill_query(query_structure: List[str | List], ent_in: Dict, ent_out: Dict, answer: int) → bool
    Private method for fill_query logic.

achieve_answer(query: List[str | List], ent_in: Dict, ent_out: Dict) → set
    Private method for achieve_answer logic. @TODO: Document the code

write_links(ent_out, small_ent_out)

ground_queries(query_structure: List[str | List], ent_in: Dict, ent_out: Dict, small_ent_in: Dict,
               small_ent_out: Dict, gen_num: int, query_name: str)
    Generating queries and achieving answers

unmap(query_type, queries, tp_answers, fp_answers, fn_answers)

unmap_query(query_structure, query, id2ent, id2rel)

generate_queries(query_struct: List, gen_num: int, query_type: str)
    Passing incoming and outgoing edges to ground queries depending on mode [train valid or text] and getting
    queries and answers in return @ TODO: create a class for each single query struct

save_queries(query_type: str, gen_num: int, save_path: str)

abstractmethod load_queries(path)

```

```

get_queries (query_type: str, gen_num: int)

static save_queries_and_answers (path: str, data: List[Tuple[str, Tuple[collections.defaultdict]]])
    → None
    Save Queries into Disk

static load_queries_and_answers (path: str) → List[Tuple[str, Tuple[collections.defaultdict]]]
    Load Queries from Disk to Memory

```

dicee.read_preprocess_save_load_kg

Submodules

dicee.read_preprocess_save_load_kg.preprocess

Attributes

logger

Classes

PreprocessKG

Preprocess the data in memory

Module Contents

dicee.read_preprocess_save_load_kg.preprocess.**logger**

class dicee.read_preprocess_save_load_kg.preprocess.**PreprocessKG** (*kg*)

Preprocess the data in memory

kg

start () → None

Preprocess train, valid and test datasets stored in knowledge graph instance

Parameter

rtype

None

preprocess_with_byte_pair_encoding ()

preprocess_with_byte_pair_encoding_with_padding () → None

Preprocess with byte pair encoding and add padding

preprocess_with_pandas () → None

Preprocess with pandas: add reciprocal triples, construct vocabulary, and index datasets

preprocess_with_polars () → None

Preprocess with polars: add reciprocal triples and create indexed datasets

`sequential_vocabulary_construction()` → None

- (1) Read input data into memory
- (2) Remove triples with a condition
- (3) **Serialize vocabularies in a pandas dataframe where**
=> the index is integer and => a single column is string (e.g. URI)

`dicee.read_preprocess_save_load_kg.read_from_disk`

Attributes

logger

Classes

ReadFromDisk

Read the data from disk into memory

Module Contents

`dicee.read_preprocess_save_load_kg.read_from_disk.logger`

class `dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk(kg)`

Read the data from disk into memory

kg

start() → None

Read a knowledge graph from disk into memory

Data will be available at the `train_set`, `test_set`, `valid_set` attributes.

Parameter

None

rtype

None

add_noisy_triples_into_training()

`dicee.read_preprocess_save_load_kg.save_load_disk`

Attributes

logger

Classes

Module Contents

`dicee.read_preprocess_save_load_kg.save_load_disk.logger`

class `dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk(kg)`

`kg`

`save()`

`load()`

`dicee.read_preprocess_save_load_kg.util`

Attributes

logger

Functions

<code>polars_dataframe_indexer(→ polars.DataFrame)</code>	Replaces 'subject', 'relation', and 'object' columns in the input Polars DataFrame with their corresponding index values
<code>pandas_dataframe_indexer(→ pandas.DataFrame)</code>	Replaces 'subject', 'relation', and 'object' columns in the input Pandas DataFrame with their corresponding index values
<code>apply_reciprocal_or_noise(add_reciprocal, eval_model)</code>	Add reciprocal triples if conditions are met
<code>timeit(func)</code>	
<code>read_with_polars(→ polars.DataFrame)</code>	Load and Preprocess via Polars
<code>read_with_pandas(data_path[, read_only_few, ...])</code>	Load and Preprocess via Pandas
<code>read_from_disk(→ Tuple[polars.DataFrame, pandas.DataFrame])</code>	
<code>count_triples(→ int)</code>	Returns the total number of triples in the triple store.
<code>fetch_worker(endpoint, offsets, chunk_size, ...)</code>	Worker process: fetch assigned chunks and save to disk with per-worker tqdm.
<code>read_from_triple_store_with_polars(endpoint[, ...])</code>	Main function to read all triples in parallel, save as Parquet, and load into Polars dataframe.
<code>read_from_triple_store_with_pandas([endpoint]</code>	Read triples from triple store into pandas dataframe
<code>get_er_vocab(data[, file_path])</code>	
<code>get_re_vocab(data[, file_path])</code>	
<code>get_ee_vocab(data[, file_path])</code>	
<code>create_constraints(triples[, file_path])</code>	
<code>load_with_pandas(→ None)</code>	Deserialize data
<code>save_numpy_ndarray(*, data, file_path)</code>	
<code>load_numpy_ndarray(*, file_path)</code>	
<code>save_pickle(*, data[, file_path])</code>	
<code>load_pickle(*[, file_path])</code>	

continues on next page

Table 66 – continued from previous page

<code>create_recipriocal_triples(x)</code>	Add inverse triples into dask dataframe
<code>dataset_sanity_checking(→ None)</code>	

Module Contents

`dicee.read_preprocess_save_load_kg.util.logger`

`dicee.read_preprocess_save_load_kg.util.polars_dataframe_indexer(`
 `df_polars: polars.DataFrame, idx_entity: polars.DataFrame, idx_relation: polars.DataFrame)`
 `→ polars.DataFrame`

Replaces ‘subject’, ‘relation’, and ‘object’ columns in the input Polars DataFrame with their corresponding index values from the entity and relation index DataFrames.

This function processes the DataFrame in three main steps: 1. Replace the ‘relation’ values with the corresponding index from `idx_relation`. 2. Replace the ‘subject’ values with the corresponding index from `idx_entity`. 3. Replace the ‘object’ values with the corresponding index from `idx_entity`.

Parameters:

df_polars

[polars.DataFrame] The input Polars DataFrame containing columns: ‘subject’, ‘relation’, and ‘object’.

idx_entity

[polars.DataFrame] A Polars DataFrame that contains the mapping between entity names and their corresponding indices. Must have columns: ‘entity’ and ‘index’.

idx_relation

[polars.DataFrame] A Polars DataFrame that contains the mapping between relation names and their corresponding indices. Must have columns: ‘relation’ and ‘index’.

Returns:

polars.DataFrame

A DataFrame with the ‘subject’, ‘relation’, and ‘object’ columns replaced by their corresponding indices.

Example Usage:

```
>>> df_polars = pl.DataFrame({
    "subject": ["Alice", "Bob", "Charlie"],
    "relation": ["knows", "works_with", "lives_in"],
    "object": ["Dave", "Eve", "Frank"]
})
>>> idx_entity = pl.DataFrame({
    "entity": ["Alice", "Bob", "Charlie", "Dave", "Eve", "Frank"],
    "index": [0, 1, 2, 3, 4, 5]
})
>>> idx_relation = pl.DataFrame({
    "relation": ["knows", "works_with", "lives_in"],
    "index": [0, 1, 2]
})
>>> polars_dataframe_indexer(df_polars, idx_entity, idx_relation)
```


Steps:

1. Join the input DataFrame *df_polars* on the 'relation' column with *idx_relation* to replace the relations with their indices.
2. Join on 'subject' to replace it with the corresponding entity index using a left join on *idx_entity*.
3. Join on 'object' to replace it with the corresponding entity index using a left join on *idx_entity*.
4. Select only the 'subject', 'relation', and 'object' columns to return the final result.

```
dicee.read_preprocess_save_load_kg.util.pandas_dataframe_indexer(  
    df_pandas: pandas.DataFrame, idx_entity: pandas.DataFrame, idx_relation: pandas.DataFrame)  
    → pandas.DataFrame
```

Replaces 'subject', 'relation', and 'object' columns in the input Pandas DataFrame with their corresponding index values from the entity and relation index DataFrames.

Parameters:

df_pandas

[pd.DataFrame] The input Pandas DataFrame containing columns: 'subject', 'relation', and 'object'.

idx_entity

[pd.DataFrame] A Pandas DataFrame that contains the mapping between entity names and their corresponding indices. Must have columns: 'entity' and 'index'.

idx_relation

[pd.DataFrame] A Pandas DataFrame that contains the mapping between relation names and their corresponding indices. Must have columns: 'relation' and 'index'.

Returns:

pd.DataFrame

A DataFrame with the 'subject', 'relation', and 'object' columns replaced by their corresponding indices.

```
dicee.read_preprocess_save_load_kg.util.apply_reciprocal_or_noise(add_reciprocal: bool,  
    eval_model: str, df: object = None, info: str = None)
```

Add reciprocal triples if conditions are met

```
dicee.read_preprocess_save_load_kg.util.timeit(func)
```

```
dicee.read_preprocess_save_load_kg.util.read_with_polars(data_path,  
    read_only_few: int = None, sample_triples_ratio: float = None, separator: str = None)  
    → polars.DataFrame
```

Load and Preprocess via Polars

```
dicee.read_preprocess_save_load_kg.util.read_with_pandas(data_path,  
    read_only_few: int = None, sample_triples_ratio: float = None, separator: str = None)
```

Load and Preprocess via Pandas

```
dicee.read_preprocess_save_load_kg.util.read_from_disk(data_path: str,  
    read_only_few: int = None, sample_triples_ratio: float = None, backend: str = None,  
    separator: str = None) → Tuple[polars.DataFrame, pandas.DataFrame]
```

```
dicee.read_preprocess_save_load_kg.util.count_triples(endpoint: str) → int
```

Returns the total number of triples in the triple store.

```
dicee.read_preprocess_save_load_kg.util.fetch_worker(endpoint: str, offsets: list[int],
chunk_size: int, output_dir: str, worker_id: int)
```

Worker process: fetch assigned chunks and save to disk with per-worker tqdm.

```
dicee.read_preprocess_save_load_kg.util.read_from_triple_store_with_polars(
endpoint: str, chunk_size: int = 500000, output_dir: str = 'triples_parquet')
```

Main function to read all triples in parallel, save as Parquet, and load into Polars dataframe.

```
dicee.read_preprocess_save_load_kg.util.read_from_triple_store_with_pandas(
endpoint: str = None)
```

Read triples from triple store into pandas dataframe

```
dicee.read_preprocess_save_load_kg.util.get_er_vocab(data, file_path: str = None)
```

```
dicee.read_preprocess_save_load_kg.util.get_re_vocab(data, file_path: str = None)
```

```
dicee.read_preprocess_save_load_kg.util.get_ee_vocab(data, file_path: str = None)
```

```
dicee.read_preprocess_save_load_kg.util.create_constraints(triples, file_path: str = None)
```

(1) Extract domains and ranges of relations

(2) Store a mapping from relations to entities that are outside of the domain and range. Crete constrained entities based on the range of relations :param triples: :return: Tuple[dict, dict]

```
dicee.read_preprocess_save_load_kg.util.load_with_pandas(self) → None
```

Deserialize data

```
dicee.read_preprocess_save_load_kg.util.save_numpy_ndarray(*, data: numpy.ndarray,
file_path: str)
```

```
dicee.read_preprocess_save_load_kg.util.load_numpy_ndarray(*, file_path: str)
```

```
dicee.read_preprocess_save_load_kg.util.save_pickle(*, data: object, file_path=str)
```

```
dicee.read_preprocess_save_load_kg.util.load_pickle(*, file_path=str)
```

```
dicee.read_preprocess_save_load_kg.util.create_recipriocal_triples(x)
```

Add inverse triples into dask dataframe :param x: :return:

```
dicee.read_preprocess_save_load_kg.util.dataset_sanity_checking(
train_set: numpy.ndarray, num_entities: int, num_relations: int) → None
```

Parameters

- **train_set**
- **num_entities**
- **num_relations**

Returns

Classes

<i>PreprocessKG</i>	Preprocess the data in memory
<i>LoadSaveToDisk</i>	
<i>ReadFromDisk</i>	Read the data from disk into memory

Package Contents

class `dicee.read_preprocess_save_load_kg.PreprocessKG(kg)`

Preprocess the data in memory

kg

start() → None

Preprocess train, valid and test datasets stored in knowledge graph instance

Parameter

rtype

None

preprocess_with_byte_pair_encoding()

preprocess_with_byte_pair_encoding_with_padding() → None

Preprocess with byte pair encoding and add padding

preprocess_with_pandas() → None

Preprocess with pandas: add reciprocal triples, construct vocabulary, and index datasets

preprocess_with_polars() → None

Preprocess with polars: add reciprocal triples and create indexed datasets

sequential_vocabulary_construction() → None

(1) Read input data into memory

(2) Remove triples with a condition

(3) **Serialize vocabularies in a pandas dataframe where**

=> the index is integer and => a single column is string (e.g. URI)

class `dicee.read_preprocess_save_load_kg.LoadSaveToDisk(kg)`

kg

save()

load()

class `dicee.read_preprocess_save_load_kg.ReadFromDisk(kg)`

Read the data from disk into memory

kg

start() → None

Read a knowledge graph from disk into memory

Data will be available at the train_set, test_set, valid_set attributes.

Parameter

None

rtype

None

add_noisy_triples_into_training()

dicee.sanity_checkers

Attributes

<code>logger</code>

Functions

<code>is_sparql_endpoint_alive([sparql_endpoint])</code>	
<code>validate_knowledge_graph(args)</code>	Validating the source of knowledge graph
<code>sanity_checking_with_arguments(args)</code>	
<code>sanity_check_callback_args(args)</code>	Perform sanity checks on callback-related arguments.

Module Contents

`dicee.sanity_checkers.logger`

`dicee.sanity_checkers.is_sparql_endpoint_alive(sparql_endpoint: str = None)`

`dicee.sanity_checkers.validate_knowledge_graph(args)`

Validating the source of knowledge graph

`dicee.sanity_checkers.sanity_checking_with_arguments(args)`

`dicee.sanity_checkers.sanity_check_callback_args(args)`

Perform sanity checks on callback-related arguments.

dicee.scripts

Submodules

dicee.scripts.index_serve

\$ docker pull qdrant/qdrant && docker run -p 6333:6333 -p 6334:6334 -v \$(pwd)/qdrant_storage:/qdrant/storage:z
qdrant/qdrant \$ dicee_vector_db -index -serve -path CountryEmbeddings -collection "countries_vdb"

Attributes

<code>logger</code>
<code>app</code>
<code>neural_searcher</code>

Classes

<code>NeuralSearcher</code>
<code>StringListRequest</code>

Functions

```
get_default_arguments()
index(args)
root()
search_embeddings(q)
retrieve_embeddings(q)
search_embeddings_batch(request)
serve(args)
main()
```

Module Contents

```
dicee.scripts.index_serve.logger

dicee.scripts.index_serve.get_default_arguments()

dicee.scripts.index_serve.index(args)

dicee.scripts.index_serve.app

dicee.scripts.index_serve.neural_searcher = None

class dicee.scripts.index_serve.NeuralSearcher(args)

    collection_name

    entity_to_idx = None

    qdrant_client

    topk = 5

    retrieve_embedding(entity: str = None, entities: List[str] = None) → List

    search(entity: str)

async dicee.scripts.index_serve.root()

async dicee.scripts.index_serve.search_embeddings(q: str)

async dicee.scripts.index_serve.retrieve_embeddings(q: str)

class dicee.scripts.index_serve.StringListRequest
    Bases: pydantic.BaseModel

    queries: List[str]

    reducer: str | None = None

async dicee.scripts.index_serve.search_embeddings_batch(request: StringListRequest)

dicee.scripts.index_serve.serve(args)

dicee.scripts.index_serve.main()
```

dicee.scripts.run

Functions

<code>get_default_arguments([description])</code>	Extends lightning Trainer's arguments with ours
<code>main()</code>	

Module Contents

`dicee.scripts.run.get_default_arguments` (*description=None*)

Extends lightning Trainer's arguments with ours

`dicee.scripts.run.main()`

dicee.static_funcs

Static utility functions for DICE embeddings.

This module provides utility functions for model initialization, data loading, serialization, and various helper operations.

Attributes

<code>logger</code>	
<code>MODEL_REGISTRY</code>	

Functions

<code>create_recipriocal_triples(→ pandas.DataFrame)</code>	Add inverse triples to a DataFrame.
<code>get_er_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build entity-relation to tail vocabulary.
<code>get_re_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build relation-entity (tail) to head vocabulary.
<code>get_ee_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build entity-entity to relation vocabulary.
<code>timeit(→ Callable)</code>	Decorator to measure and print execution time and memory usage.
<code>setup_distributed_training(→ Dict[str, Union[bool, int]])</code>	Resolve distributed-training state and initialize custom DDP when needed.
<code>save_pickle(→ None)</code>	Save data to a pickle file.
<code>load_pickle(→ object)</code>	Load data from a pickle file.
<code>load_term_mapping(→ Union[dict, pandas.DataFrame])</code>	Load term-to-index mapping from pickle or CSV file.
<code>select_model(args[, is_continual_training, storage_path])</code>	
<code>load_model(→ Tuple[object, Tuple[dict, dict]])</code>	Load weights and initialize pytorch module from namespace arguments
<code>load_model_ensemble(...)</code>	Construct Ensemble Of weights and initialize pytorch module from namespace arguments
<code>save_numpy_ndarray(*, data, file_path)</code>	
<code>numpy_data_type_changer(→ numpy.ndarray)</code>	Detect most efficient data type for a given triples
<code>save_checkpoint_model(→ None)</code>	Store Pytorch model into disk

continues on next page

Table 75 – continued from previous page

<code>store(→ None)</code>	
<code>add_noisy_triples(→ pandas.DataFrame)</code>	Add randomly constructed triples
<code>read_or_load_kg(args, cls)</code>	
<code>intialize_model(...)</code>	Initialize a knowledge graph embedding model.
<code>load_json(→ Dict)</code>	Load JSON file into a dictionary.
<code>save_embeddings(→ None)</code>	Save embeddings to a CSV file.
<code>vocab_to_parquet(vocab_to_idx, name, ...)</code>	
<code>create_experiment_folder(→ str)</code>	Create a timestamped experiment folder.
<code>continual_training_setup_executor(→ None)</code>	
<code>exponential_function(→ torch.FloatTensor)</code>	
<code>load_numpy(→ numpy.ndarray)</code>	
<code>evaluate(entity_to_idx, scores, easy_answers, hard_answers)</code>	# @TODO: CD: Renamed this function
<code>download_file(url[, destination_folder])</code>	
<code>download_files_from_url(→ None)</code>	
<code>download_pretrained_model(→ str)</code>	
<code>write_csv_from_model_parallel(path)</code>	Create
<code>from_pretrained_model_write_embeddings_int(→ None)</code>	

Module Contents

`dicee.static_funcs.logger`

`dicee.static_funcs.MODEL_REGISTRY: Dict[str, Tuple[Type, str]]`

`dicee.static_funcs.create_recipriocal_triples(df: pandas.DataFrame) → pandas.DataFrame`
Add inverse triples to a DataFrame.

For each triple (s, p, o), creates an inverse triple (o, p_inverse, s).

Parameters

df – DataFrame with ‘subject’, ‘relation’, and ‘object’ columns.

Returns

DataFrame with original and inverse triples concatenated.

`dicee.static_funcs.get_er_vocab(data: numpy.ndarray, file_path: str | None = None)`
→ Dict[Tuple[int, int], List[int]]

Build entity-relation to tail vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (head, relation) pairs to list of tail entities.

`dicee.static_funcs.get_re_vocab(data: numpy.ndarray, file_path: str | None = None)`
→ Dict[Tuple[int, int], List[int]]

Build relation-entity (tail) to head vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).

- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (relation, tail) pairs to list of head entities.

```
dicee.static_funcs.get_ee_vocab(data: numpy.ndarray, file_path: str | None = None)
→ Dict[Tuple[int, int], List[int]]
```

Build entity-entity to relation vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (head, tail) pairs to list of relations.

```
dicee.static_funcs.timeit(func: Callable) → Callable
```

Decorator to measure and print execution time and memory usage.

Parameters

func – Function to be timed.

Returns

Wrapped function that prints timing information.

```
dicee.static_funcs.setup_distributed_training(args) → Dict[str, bool | int]
```

Resolve distributed-training state and initialize custom DDP when needed.

```
dicee.static_funcs.save_pickle(*, data: object | None = None, file_path: str) → None
```

Save data to a pickle file.

Note: Consider using more portable formats (JSON, Parquet) for new code.

Parameters

- **data** – Object to serialize. If None, nothing is saved.
- **file_path** – Path where the pickle file will be saved.

```
dicee.static_funcs.load_pickle(file_path: str) → object
```

Load data from a pickle file.

Note: Consider using more portable formats (JSON, Parquet) for new code.

Parameters

file_path – Path to the pickle file.

Returns

Deserialized object from the pickle file.

```
dicee.static_funcs.load_term_mapping(file_path: str) → dict | polars.DataFrame
```

Load term-to-index mapping from pickle or CSV file.

Attempts to load from pickle first, falls back to CSV if not found.

Parameters

file_path – Base path without extension.

Returns

Dictionary or Polars DataFrame containing the mapping.

`dicee.static_funcs.select_model` (*args: dict, is_continual_training: bool = None, storage_path: str = None*)

`dicee.static_funcs.load_model` (*path_of_experiment_folder: str, model_name='model.pt', verbose=0*)
→ Tuple[object, Tuple[dict, dict]]

Load weights and initialize pytorch module from namespace arguments

`dicee.static_funcs.load_model_ensemble` (*path_of_experiment_folder: str*)
→ Tuple[`dicee.models.base_model.BaseKGE`, Tuple[pandas.DataFrame, pandas.DataFrame]]

Construct Ensemble Of weights and initialize pytorch module from namespace arguments

- (1) Detect models under given path
- (2) Accumulate parameters of detected models
- (3) Normalize parameters
- (4) Insert (3) into model.

`dicee.static_funcs.save_numpy_ndarray` (*, *data: numpy.ndarray, file_path: str*)

`dicee.static_funcs.numpy_data_type_changer` (*train_set: numpy.ndarray, num: int*)
→ numpy.ndarray

Detect most efficient data type for a given triples :param train_set: :param num: :return:

`dicee.static_funcs.save_checkpoint_model` (*model, path: str*) → None

Store Pytorch model into disk

`dicee.static_funcs.store` (*trained_model, model_name: str = 'model', full_storage_path: str = None, save_embeddings_as_csv=False*) → None

`dicee.static_funcs.add_noisy_triples` (*train_set: pandas.DataFrame, add_noise_rate: float*)
→ pandas.DataFrame

Add randomly constructed triples :param train_set: :param add_noise_rate: :return:

`dicee.static_funcs.read_or_load_kg` (*args, cls*)

`dicee.static_funcs.intialize_model` (*args: Dict, verbose: int = 0, for_inference: bool = False*)
→ Tuple[`dicee.models.base_model.BaseKGE`, str]

Initialize a knowledge graph embedding model.

Parameters

- **args** – Dictionary containing model configuration including ‘model’ key.
- **verbose** – Verbosity level. If > 0, prints initialization message.
- **for_inference** – If True, never build the FSDP row-wise-sharded shell, even if `args[“trainer”] == “torchFSDP”`. A completed run’s `model.pt` (written by `TorchFSDP-Trainer._materialize_model()`) already has a plain, full `entity_embeddings.weight` — loading it into the sharded shell (which defers `entity_embeddings` until `setup_fsdp_training()`) fails with “Unexpected key(s): `entity_embeddings.weight`”. Callers loading a finished model for inference (`load_model`, `load_model_ensemble`) must set this; callers building a fresh model to train/resume with (`select_model`) must not.

Returns

Tuple of (initialized model, form of labelling string).

Raises

ValueError – If the model name is not recognized.

`dicee.static_funcs.load_json(path: str) → Dict`

Load JSON file into a dictionary.

Parameters

path – Path to the JSON file.

Returns

Dictionary containing the JSON data.

Raises

- **FileNotFoundError** – If the file does not exist.
- **json.JSONDecodeError** – If the file contains invalid JSON.

`dicee.static_funcs.save_embeddings(embeddings: numpy.ndarray, indexes: List, path: str) → None`

Save embeddings to a CSV file.

Parameters

- **embeddings** – NumPy array of embeddings with shape (n_items, embedding_dim).
- **indexes** – List of index labels (entity/relation names).
- **path** – Output file path.

`dicee.static_funcs.vocab_to_parquet(vocab_to_idx, name, path_for_serialization, print_into)`

`dicee.static_funcs.create_experiment_folder(folder_name: str = 'Experiments') → str`

Create a timestamped experiment folder.

Parameters

folder_name – Base directory name for experiments.

Returns

Full path to the created experiment folder.

`dicee.static_funcs.continual_training_setup_executor(executor) → None`

`dicee.static_funcs.exponential_function(x: numpy.ndarray, lam: float, ascending_order=True)`
→ torch.FloatTensor

`dicee.static_funcs.load_numpy(path) → numpy.ndarray`

`dicee.static_funcs.evaluate(entity_to_idx, scores, easy_answers, hard_answers)`

@TODO: CD: Renamed this function Evaluate multi hop query answering on different query types

`dicee.static_funcs.download_file(url, destination_folder='.')`

`dicee.static_funcs.download_files_from_url(base_url: str, destination_folder='.') → None`

Parameters

- **base_url** (e.g. <https://files.dice-research.org/projects/DiceEmbeddings/KINSHIP-Keci-dim128-epoch256-KvsAll>)
- **destination_folder** (e.g. `"KINSHIP-Keci-dim128-epoch256-KvsAll"`)

`dicee.static_funcs.download_pretrained_model(url: str) → str`

`dicee.static_funcs.write_csv_from_model_parallel(path: str)`

Create

`dicee.static_funcs.from_pretrained_model_write_embeddings_into_csv(path: str) → None`

dicee.static_funcs_training

Training-related static functions.

This module provides backward compatibility by re-exporting evaluation functions from the new `dicee.evaluation` module, along with training utilities.

Deprecated since version Evaluation: functions have moved to `dicee.evaluation`. Use that module for new code. This module will continue to export training utilities.

Functions

<code>evaluate_bpe_lp(→ Dict[str, float])</code>	Evaluate link prediction with BPE-encoded entities.
<code>evaluate_lp(→ Dict[str, float])</code>	Evaluate link prediction with batched processing.
<code>efficient_zero_grad(→ None)</code>	Efficiently zero gradients using <code>parameter.grad = None</code> .
<code>make_iterable_verbose(→ Iterable)</code>	Wrap an iterable with <code>tqdm</code> progress bar if <code>verbose</code> is <code>True</code> .

Module Contents

```
dicee.static_funcs_training.evaluate_bpe_lp(model, triple_idx: List[Tuple],
    all_bpe_shaped_entities, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List],
    info: str = 'Eval Starts') → Dict[str, float]
```

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.static_funcs_training.evaluate_lp(model, triple_idx, num_entities: int,
    er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List], info: str = 'Eval Starts',
    batch_size: int = 128, chunk_size: int = 1000) → Dict[str, float]
```

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.

- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.static_funcs_training.efficient_zero_grad(model) → None`

Efficiently zero gradients using `parameter.grad = None`.

This is more efficient than `optimizer.zero_grad()` as it avoids memory operations.

See: https://pytorch.org/tutorials/recipes/recipes/tuning_guide.html

Parameters

model – PyTorch model to zero gradients for.

`dicee.static_funcs_training.make_iterable_verbose(iterable_object: Iterable, verbose: bool, desc: str = 'Default', position: int | None = None, leave: bool = True) → Iterable`

Wrap an iterable with tqdm progress bar if verbose is True.

Parameters

- **iterable_object** – The iterable to potentially wrap.
- **verbose** – Whether to show progress bar.
- **desc** – Description for the progress bar.
- **position** – Position of the progress bar.
- **leave** – Whether to leave the progress bar after completion.

Returns

The original iterable or a tqdm-wrapped version.

`dicee.static_preprocess_funcs`

Attributes

<code>logger</code>
<code>enable_log</code>

Functions

<code>timeit(func)</code>	
<code>preprocesses_input_args(args)</code>	Sanity Checking in input arguments
<code>create_constraints(→ Tuple[dict, dict, dict])</code>	
<code>get_er_vocab(data)</code>	
<code>get_re_vocab(data)</code>	
<code>get_ee_vocab(data)</code>	
<code>mapping_from_first_two_cols_to_third(train_se</code>	

Module Contents

`dicee.static_preprocess_funcs.logger`

`dicee.static_preprocess_funcs.enable_log = False`

`dicee.static_preprocess_funcs.timeit (func)`

`dicee.static_preprocess_funcs.preprocesses_input_args (args)`

Sanity Checking in input arguments

`dicee.static_preprocess_funcs.create_constraints (triples: numpy.ndarray)`
→ Tuple[dict, dict, dict, dict]

(1) Extract domains and ranges of relations

(2) Store a mapping from relations to entities that are outside of the domain and range. Create constraints entities based on the range of relations :param triples: :return:

`dicee.static_preprocess_funcs.get_er_vocab (data)`

`dicee.static_preprocess_funcs.get_re_vocab (data)`

`dicee.static_preprocess_funcs.get_ee_vocab (data)`

`dicee.static_preprocess_funcs.mapping_from_first_two_cols_to_third (train_set_idx)`

`dicee.trainer`

Submodules

`dicee.trainer.auto_batch_finder`

Attributes

logger

Functions

<code>find_good_batch_size(→ Tuple[int, Optional[float]])</code>	Find a batch size that uses GPU memory efficiently.
--	---

Module Contents

`dicee.trainer.auto_batch_finder.logger`

`dicee.trainer.auto_batch_finder.find_good_batch_size (`
`train_loader: torch.utils.data.DataLoader, training_step_fn: Callable, device)`
→ Tuple[int, float | None]

Find a batch size that uses GPU memory efficiently.

Progressively increases batch size (with tunable delta) until GPU memory exceeds 90% or a CUDA OOM error is raised, then backs off to the last safe batch size. Only supported on CUDA devices; returns the initial batch size unchanged on CPU.

Parameters

- **train_loader** – DataLoader wrapping the training dataset.
- **training_step_fn** – Callable[[batch], float] — runs one forward+backward pass and returns the scalar loss. Each trainer passes its own implementation so this function stays trainer-agnostic.
- **device** – torch.device (or anything accepted by torch.device()) used to monitor GPU memory.

Returns

(batch_size, runtime) where runtime is the wall-clock seconds of the last successful batch, or None when batch finding was skipped.

dicee.trainer.dice_trainer

DICE Trainer module for knowledge graph embedding training.

Provides the DICE_Trainer class which supports multiple training backends including PyTorch Lightning, DDP, and custom CPU/GPU trainers.

Attributes

<i>logger</i>

Classes

<i>DICE_Trainer</i>	DICE_Trainer implement
---------------------	------------------------

Functions

<i>load_term_mapping</i> (→ polars.DataFrame)	Load term-to-index mapping from CSV file.
<i>initialize_trainer</i> (...)	Initialize the appropriate trainer based on configuration.
<i>get_callbacks</i> (→ List)	Create list of callbacks based on configuration.

Module Contents

`dicee.trainer.dice_trainer.logger`

`dicee.trainer.dice_trainer.load_term_mapping` (*file_path: str*) → polars.DataFrame
Load term-to-index mapping from CSV file.

Parameters

file_path – Base path without extension.

Returns

Polars DataFrame containing the mapping.

`dicee.trainer.dice_trainer.initialize_trainer` (*args, callbacks: List*)

→ *dicee.trainer.torch_trainer.TorchTrainer* | *dicee.trainer.model_parallelism.TensorParallel* | *dicee.trainer.torch_trainer_ddp*

Initialize the appropriate trainer based on configuration.

Parameters

- **args** – Configuration arguments containing trainer type.
- **callbacks** – List of training callbacks.

Returns

Initialized trainer instance.

Raises

AssertionError – If trainer is None after initialization.

`dicee.trainer.dice_trainer.get_callbacks(args)` → List

Create list of callbacks based on configuration.

Parameters

args – Configuration arguments.

Returns

List of callback instances.

class `dicee.trainer.dice_trainer.DICE_Trainer` (*args, is_continual_training: bool, storage_path, evaluator=None*)

DICE_Trainer implement

- 1- Pytorch Lightning trainer (<https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html>)
- 2- Multi-GPU Trainer(<https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html>)
- 3- CPU Trainer

args

is_continual_training:bool

storage_path:str

evaluator:

report:dict

report

args

trainer = None

is_continual_training

storage_path

evaluator = None

form_of_labelling = None

continual_start (*knowledge_graph*)

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

initialize_trainer (*callbacks: List*)

→ *lightning.Trainer* | *dicее.trainer.model_parallelism.TensorParallel* | *dicее.trainer.torch_trainer.TorchTrainer* | *dicее.t*

Initialize Trainer from input arguments

initialize_or_load_model ()

init_dataloader (*dataset: torch.utils.data.Dataset*) → *torch.utils.data.DataLoader*

init_dataset () → *torch.utils.data.Dataset*

start (*knowledge_graph: dicее.knowledge_graph.KG* | *numpy.memmap*)

→ *Tuple[dicее.models.base_model.BaseKGE, str]*

Start the training

- (1) Initialize Trainer
- (2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

k_fold_cross_validation (*dataset*) → *Tuple[dicее.models.base_model.BaseKGE, str]*

Perform K-fold Cross-Validation

1. Obtain K train and test splits.
2. **For each split,**
 - 2.1 initialize trainer and model
 - 2.2. Train model with configuration provided in args.
 - 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.
3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

dicее.trainer.model_parallelism

Tensor Parallelism trainer for ensemble knowledge graph embeddings.

This module implements the tensor parallelism training strategy described in:

“Multiple Run Ensemble Learning with Low-Dimensional Knowledge Graph Embeddings”

<https://arxiv.org/abs/2104.05003>

The TensorParallel trainer creates an ensemble of models, each trained on a separate GPU, allowing efficient utilization of multi-GPU systems through model parallelism.

Attributes

<code>logger</code>

Classes

<code>TensorParallel</code>

Abstract base class for KGE model trainers.

Functions

<code>extract_input_outputs(z[, device])</code>
<code>forward_backward_update_loss(→ float)</code>

Module Contents

`dicee.trainer.model_parallelism.logger`

`dicee.trainer.model_parallelism.extract_input_outputs` (*z*: list, *device*=None)

`dicee.trainer.model_parallelism.forward_backward_update_loss` (*z*: Tuple, *ensemble_model*)
→ float

class `dicee.trainer.model_parallelism.TensorParallel` (*args*, *callbacks*)

Bases: `dicee.abstracts.AbstractTrainer`

Abstract base class for KGE model trainers.

Provides the callback dispatch mechanism shared by all concrete trainer implementations (TorchTrainer, TorchD-
DPTrainer, etc.). Sub-classes call the `on_*` hooks at the appropriate points in the training loop so that any registered
AbstractCallback can react.

Parameters

- **args** (*argparse.Namespace* or similar) – Processed configuration object. Must expose at least `random_seed` (int).
- **callbacks** (list of AbstractCallback) – Ordered list of callback instances to invoke at each lifecycle hook.

fit (**args*, ***kwargs*)

Train model

`dicee.trainer.torch_trainer`

Attributes

<code>logger</code>

Classes

<i>TorchTrainer</i>	TorchTrainer for using single GPU or multi CPUs on a single node
---------------------	--

Module Contents

`dicee.trainer.torch_trainer.logger`

class `dicee.trainer.torch_trainer.TorchTrainer` (*args*, *callbacks*)

Bases: `dicee.abstracts.AbstractTrainer`

TorchTrainer for using single GPU or multi CPUs on a single node

Arguments

callbacks: list of Abstract callback instances

loss_function = None

optimizer = None

model = None

train_dataloaders = None

training_step = None

process

fit (**args*, *train_dataloaders*, ***kwargs*) → None

Training starts

Arguments

kwargs:Tuple

empty dictionary

Return type

batch loss (float)

forward_backward_update (*x_batch: torch.Tensor*, *y_batch: torch.Tensor*) → `torch.Tensor`

Compute forward, loss, backward, and parameter update

Arguments

Return type

batch loss (float)

extract_input_outputs_set_device (*batch: list*) → `Tuple`

Construct inputs and outputs from a batch of inputs with outputs From a batch of inputs and put

Arguments

Return type

(tuple) mini-batch on select device

dicee.trainer.torch_trainer_ddp

Attributes

<code>logger</code>

Classes

<code>TorchDDPTrainer</code>	Abstract base class for KGE model trainers.
<code>NodeTrainer</code>	

Functions

<code>make_iterable_verbose(→ Iterable)</code>
--

Module Contents

`dicee.trainer.torch_trainer_ddp.logger`

`dicee.trainer.torch_trainer_ddp.make_iterable_verbose(iterable_object, verbose, desc='Default', position=None, leave=True) → Iterable`

class `dicee.trainer.torch_trainer_ddp.TorchDDPTrainer(args, callbacks)`

Bases: `dicee.abstracts.AbstractTrainer`

Abstract base class for KGE model trainers.

Provides the callback dispatch mechanism shared by all concrete trainer implementations (TorchTrainer, TorchDDPTrainer, etc.). Sub-classes call the `on_*` hooks at the appropriate points in the training loop so that any registered `AbstractCallback` can react.

Parameters

- **args** (*argparse.Namespace or similar*) – Processed configuration object. Must expose at least `random_seed` (int).
- **callbacks** (*list of AbstractCallback*) – Ordered list of callback instances to invoke at each lifecycle hook.

fit (**args, **kwargs*)

class `dicee.trainer.torch_trainer_ddp.NodeTrainer(trainer, model: torch.nn.Module, train_dataset_loader: torch.utils.data.DataLoader, callbacks, num_epochs: int)`

trainer

local_rank

global_rank

optimizer

train_dataset_loader

```

loss_func

callbacks

model

num_epochs

loss_history = []

ctx

scaler

extract_input_outputs (z: list)

train()

```

dicee.trainer.torch_trainer_fsdp

Multi-GPU trainer: row-wise sharded entity embeddings + FSDP dense parameters.

Entity embeddings are sharded row-wise across ranks via `dist.all_to_all_single`. Each rank owns a contiguous slice of entity rows; forward and backward use `all_to_all` to route index requests and gradient returns to the owning rank. A per-rank `_LocalSparseAdam` updates only the rows that received a non-zero gradient each batch.

Dense parameters (relation embeddings, Clifford coefficients, ...) are wrapped with FSDP and updated by the standard dense optimizer.

Attributes

<code>logger</code>
<code>OptimizedModule</code>

Classes

<code>TorchFSDPTrainer</code>	Single-node multi-GPU trainer: row-wise sharded entity embeddings + FSDP.
-------------------------------	---

Module Contents

`dicee.trainer.torch_trainer_fsdp.logger`

`dicee.trainer.torch_trainer_fsdp.OptimizedModule = None`

class `dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer` (*args*, *callbacks*)

Bases: `dicee.abstracts.AbstractTrainer`

Single-node multi-GPU trainer: row-wise sharded entity embeddings + FSDP.

Entity embeddings

Sharded row-wise across GPUs. `all_to_all` routes each index request to its owning rank; embeddings and gradients travel back the same way. `_LocalSparseAdam` updates only the rows accessed in each batch — $O(\text{active_rows})$ not $O(\text{shard_size})$. Each entity receives gradient from every rank's batches.

Dense parameters (relation embeddings, Clifford coefficients, ...)

Wrapped with FSDP; updated by the standard dense optimizer step().

Usage

Launched with torchrun:

```
torchrun -nproc_per_node=4 -u dicee --trainer torchFSDP ...
```

`local_rank`

`global_rank`

`is_global_zero`

`device`

`model = None`

`raw_model = None`

`optimizer = None`

`loss_func = None`

`train_dataset_loader = None`

`ptdtype`

`ctx`

`scaler`

`sharding_strategy`

`gradient_clip_val`

`use_compile`

`num_workers`

`prefetch_factor`

`entity_optim_device`

`checkpoint_every_n_epochs`

`fit(*args, **kwargs)`

`extract_input_outputs(z)`

Classes

DICE_Trainer

DICE_Trainer implement

Package Contents

class `diccee.trainer.DICE_Trainer` (*args*, *is_continual_training*: *bool*, *storage_path*, *evaluator*=*None*)

DICE_Trainer implement

- 1- Pytorch Lightning trainer (<https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html>)
- 2- Multi-GPU Trainer(<https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html>)
- 3- CPU Trainer

args

is_continual_training:*bool*

storage_path:*str*

evaluator:

report:*dict*

report

args

trainer = *None*

is_continual_training

storage_path

evaluator = *None*

form_of_labelling = *None*

continual_start (*knowledge_graph*)

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

initialize_trainer (*callbacks*: *List*)

→ `lightning.Trainer` | `diccee.trainer.model_parallelism.TensorParallel` | `diccee.trainer.torch_trainer.TorchTrainer` | `diccee.trainer.cpu_trainer.CPUTrainer`

Initialize Trainer from input arguments

initialize_or_load_model ()

init_dataloader (*dataset*: `torch.utils.data.Dataset`) → `torch.utils.data.DataLoader`

init_dataset () → `torch.utils.data.Dataset`

start (*knowledge_graph*: *dicee.knowledge_graph.KG* | *numpy.memmap*)
→ Tuple[*dicee.models.base_model.BaseKGE*, str]

Start the training

- (1) Initialize Trainer
- (2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

k_fold_cross_validation (*dataset*) → Tuple[*dicee.models.base_model.BaseKGE*, str]

Perform K-fold Cross-Validation

1. Obtain K train and test splits.
2. **For each split,**
 - 2.1 initialize trainer and model
 - 2.2. Train model with configuration provided in args.
 - 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.
3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

dicee.weight_averaging

Classes

<i>ASWA</i>	Adaptive stochastic weight averaging
<i>SWA</i>	Stochastic Weight Averaging callback.
<i>SWAG</i>	Stochastic Weight Averaging - Gaussian (SWAG).
<i>EMA</i>	Exponential Moving Average (EMA) callback.
<i>TWA</i>	Train with Weight Averaging (TWA) using subspace projection + averaging.

Module Contents

class *dicee.weight_averaging.ASWA* (*num_epochs*, *path*)

Bases: *dicee.abstracts.AbstractCallback*

Adaptive stochastic weight averaging ASWE keeps track of the validation performance and update s the ensemble model accordingly.

path

num_epochs

initial_eval_setting = None

epoch_count = 0

alphas = []

val_aswa = -1

on_fit_end(*trainer*, *model*)

Called once after the final training epoch completes.

Override to perform post-training actions such as saving the final model state, computing evaluation metrics, or cleaning up resources.

static compute_mrr(*trainer*, *model*) → float

get_aswa_state_dict(*model*)

decide(*running_model_state_dict*, *ensemble_state_dict*, *val_running_model*,
mrr_updated_ensemble_model)

Perform Hard Update, software or rejection

Parameters

- **running_model_state_dict**
- **ensemble_state_dict**
- **val_running_model**
- **mrr_updated_ensemble_model**

on_train_epoch_end(*trainer*, *model*)

Called at the end of each training epoch.

Parameters

- **trainer** (*AbstractTrainer* or *pl.Trainer*) – The active trainer instance.
- **model** (*BaseKGE*) – The model being trained. *model.loss_history* contains the per-epoch average losses accumulated so far.

class *dicee.weight_averaging.SWA*(*swa_start_epoch*, *swa_c_epochs*: *int* = 1, *lr_init*: *float* = 0.1,
swa_lr: *float* = 0.05, *max_epochs*: *int* = None)

Bases: *dicee.abstracts.AbstractCallback*

Stochastic Weight Averaging callback.

Initialize SWA callback.

swa_start_epoch: **int**

The epoch at which to start SWA.

swa_c_epochs: **int**

The number of epochs to use for SWA.

lr_init: **float**

The initial learning rate.

swa_lr: **float**

The learning rate to use during SWA.

max_epochs: **int**

The maximum number of epochs. *args.num_epochs*

swa_start_epoch

swa_c_epochs = 1


```

swa_lr = 0.05

lr_init = 0.1

max_epochs = None

swa_model = None

swa_n = 0

current_epoch = -1

static moving_average (swa_model, running_model, alpha)
    Update SWA model with moving average of current model. Math: # SWA update:  $\theta_{\text{swa}} \leftarrow (1 - \alpha) * \theta_{\text{swa}} + \alpha * \theta$  #  $\alpha = 1 / (n + 1)$ , where  $n$  = number of models already averaged #  $\alpha$  is tracked via self.swa_n in code

on_train_epoch_start (trainer, model)
    Update learning rate according to SWA schedule.

on_train_epoch_end (trainer, model)
    Apply SWA averaging if conditions are met.

on_fit_end (trainer, model)
    Replace main model with SWA model at the end of training.

class dicee.weight_averaging.SWAG (swa_start_epoch: int = 1, swa_c_epochs: int = 1, lr_init: float = 0.1,
    swa_lr: float = 0.05, max_epochs: int = None, max_num_models: int = 20,
    var_clamp: float = 1e-30)
    Bases: dicee.abstracts.AbstractCallback

    Stochastic Weight Averaging - Gaussian (SWAG). Parameters

    swa_start_epoch
        [int] Epoch at which to start collecting weights.

    swa_c_epochs
        [int] Interval of epochs between updates.

    lr_init
        [float] Initial LR.

    swa_lr
        [float] LR in SWA / GSWA phase.

    max_epochs
        [int] Total number of epochs.

    max_num_models
        [int] Number of models to keep for low-rank covariance approx.

    var_clamp
        [float] Clamp low variance for stability.

swa_start_epoch

swa_c_epochs = 1

swa_lr = 0.05

lr_init = 0.1

```

```

max_epochs = None

max_num_models = 20

var_clamp = 1e-30

mean = None

sq_mean = None

deviations = []

gswa_n = 0

current_epoch = -1

get_mean_and_var()
    Return mean + variance (diagonal part).

sample(base_model, scale=0.5)
    Sample new model from SWAG posterior distribution.

    Math: # From SWAG, posterior is approximated as: #  $\theta \sim N(\text{mean}, \Sigma)$  # where  $\Sigma \approx \text{diag}(\text{var}) + (1/(K-1)) * D D^T$  # - mean = running average of weights # - var = elementwise variance (sq_mean - mean^2) # - D = [dev_1, dev_2, ..., dev_K], deviations from mean (low-rank approx) # - K = number of collected models

    # Sampling step: # 1.  $\theta_{\text{diag}} = \text{mean} + \text{scale} * \text{std} \odot \epsilon$ , where  $\epsilon \sim N(0, I)$  # 2.  $\theta_{\text{lowrank}} = \theta_{\text{diag}} + (D z) / \sqrt{K-1}$ , where  $z \sim N(0, I_K)$  # Final sample =  $\theta_{\text{lowrank}}$ 

on_train_epoch_start(trainer, model)
    Update LR schedule (same as SWA).

on_train_epoch_end(trainer, model)
    Collect Gaussian stats at the end of epochs after swa_start.

on_fit_end(trainer, model)
    Set model weights to the collected SWAG mean at the end of training.

```

class dicee.weight_averaging.EMA(ema_start_epoch: int, decay: float = 0.999, max_epochs: int = None, ema_c_epochs: int = 1)

Bases: `dicee.abstracts.AbstractCallback`

Exponential Moving Average (EMA) callback.

Parameters

- **ema_start_epoch** (int) – Epoch to start EMA.
- **decay** (float) – EMA decay rate (typical: 0.99 - 0.9999) Math: $\theta_{\text{ema}} \leftarrow \text{decay} * \theta_{\text{ema}} + (1 - \text{decay}) * \theta$
- **max_epochs** (int) – Maximum number of epochs.

```

ema_start_epoch

decay = 0.999

max_epochs = None

ema_c_epochs = 1

```

```

ema_model = None

current_epoch = -1

static ema_update(ema_model, running_model, decay: float)
    Update EMA model with exponential moving average of current model. Math:  $\theta_{ema} \leftarrow (1 - \alpha) * \theta_{ema} + \alpha * \theta$  #  $\alpha = 1 - \text{decay}$ , where decay is the EMA smoothing factor (typical 0.99 - 0.999) #  $\alpha$  controls how much of the current model  $\theta$  contributes to the EMA # decay is fixed in code
    -> can be extended to scheduled

on_train_epoch_start(trainer, model)
    Track current epoch.

on_train_epoch_end(trainer, model)
    Update EMA if past start epoch.

on_fit_end(trainer, model)
    Replace main model with EMA model at the end of training.

class dicee.weight_averaging.TWA(twa_start_epoch: int, lr_init: float, num_samples: int = 5,
    reg_lambda: float = 0.0, max_epochs: int = None, twa_c_epochs: int = 1)
    Bases: dicee.abstracts.AbstractCallback

    Train with Weight Averaging (TWA) using subspace projection + averaging.

    Parameters

        twa_start_epoch
            [int] Epoch to start TWA.

        lr_init
            [float] Learning rate used for  $\beta$  updates.

        num_samples
            [int] Number of sampled weight snapshots to build projection subspace.

        reg_lambda
            [float] Regularization coefficient for  $\beta$  updates.

        max_epochs
            [int] Total number of training epochs.

        twa_c_epochs
            [int] Interval of epochs between TWA updates.

twa_start_epoch

num_samples = 5

reg_lambda = 0.0

max_epochs = None

lr_init

twa_c_epochs = 1

current_epoch = -1

weight_samples = []

```

```

twamodel = None

base_weights = None

P = None

beta = None

sample_weights(model)
    Collect sampled weights from the current model and maintain rolling buffer.

build_projection(weight_samples, k=None)
    Build projection subspace from collected weight samples. :param weight_samples: list of flat weight tensors [(D,), ...] :param k: number of basis vectors to keep. Defaults to min(N, D).

    Returns
        (D,) base weight vector (average) P: (D, k) projection matrix with top-k basis directions

    Return type
        mean_w

on_train_epoch_start(trainer, model)
    Track epoch.

on_train_epoch_end(trainer, model)
    Main TWA logic: build subspace and update in  $\beta$  space.

    # Math: # TWA weight update: # w_twa = mean_w + P * beta # mean_w = (1/n) * sum_i w_i (SWA baseline)
    # beta <- (1 - eta * lambda) * beta - eta * PT * g # g = gradient of training loss w.r.t. full model weights
    # eta = learning rate, lambda = ridge regularization # P = orthonormal basis spanning sampled checkpoints {w_i}

on_fit_end(trainer, model)
    Replace with TWA model at the end.

```

14.2 Attributes

<code>__version__</code>	
--------------------------	--

14.3 Classes

<i>Execute</i>	Executor class for training, retraining and evaluating KGE models.
<i>KGE</i>	Knowledge Graph Embedding Class for interactive usage of pre-trained models
<i>QueryGenerator</i>	
<i>DICE_Trainer</i>	DICE_Trainer implement
<i>Evaluator</i>	Evaluator class for KGE models in various downstream tasks.

14.4 Package Contents

class `dicee.Execute` (*args*, *continuous_training: bool = False*)

Executor class for training, retraining and evaluating KGE models.

Handles the complete workflow: 1. Loading & Preprocessing & Serializing input data 2. Training & Validation & Testing 3. Storing all necessary information

args

Processed input arguments.

distributed

Whether distributed training is enabled.

rank

Process rank in distributed training.

world_size

Total number of processes.

local_rank

Local GPU rank.

trainer

Training handler instance.

trained_model

The trained model after training completes.

knowledge_graph

The loaded knowledge graph.

report

Dictionary storing training metrics and results.

evaluator

Model evaluation handler.

distributed

rank

world_size

local_rank

args

is_continual_training = False

trainer: `dicee.trainer.DICE_Trainer` | None = None

trained_model = None

knowledge_graph: `dicee.knowledge_graph.KG` | None = None

report: Dict

evaluator: `dicee.evaluator.Evaluator` | None = None

start_time: float | None = None

is_local_rank_zero() → bool

is_global_rank_zero() → bool

cleanup()

setup_executor() → None

Set up storage directories for the experiment.

Creates or reuses experiment directories based on configuration. Saves the configuration to a JSON file.

create_and_store_kg() → None

Create knowledge graph and store as memory-mapped file.

Only executed on global rank 0 in distributed training, so exactly one process writes the shared memmap/vocab files regardless of how many nodes are involved. Skips if memmap already exists.

load_from_memmap() → None

Load knowledge graph from memory-mapped file.

save_trained_model() → None

Save a knowledge graph embedding model

(1) Send model to eval mode and cpu.

(2) Store the memory footprint of the model.

(3) Save the model into disk.

(4) Update the stats of KG again ?

Parameter

rtype

None

end(*form_of_labelling: str*) → dict

End training

(1) Store trained model.

(2) Report runtimes.

(3) Eval model if required.

Parameter

rtype

A dict containing information about the training and/or evaluation

write_report() → None

Report training related information in a report.json file

start() → dict

Start training

(1) Loading the Data # (2) Create an evaluator object. # (3) Create a trainer object. # (4) Start the training

Parameter

rtype

A dict containing information about the training and/or evaluation

```
class dicee.KGE (path=None, url=None, construct_ensemble=False, model_name=None)
```

Bases: `dicee.abstracts.BaseInteractiveKGE`, `dicee.abstracts.InteractiveQueryDecomposition`, `dicee.abstracts.BaseInteractiveTrainKGE`

Knowledge Graph Embedding Class for interactive usage of pre-trained models

`__str__()`

`to(device: str) → None`

`get_transductive_entity_embeddings(indices: torch.LongTensor | List[str], as_pytorch=False, as_numpy=False, as_list=True) → torch.FloatTensor | numpy.ndarray | List[float]`

`create_vector_database(collection_name: str, distance: str, location: str = 'localhost', port: int = 6333)`

`generate(h="", r="")`

`eval_lp_performance(dataset=List[Tuple[str, str, str]], filtered=True)`

`predict_missing_head_entity(relation: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False) → Tuple`

Given a relation and a tail entity, return top k ranked head entity.

$\operatorname{argmax}_{\{e \in E\}} f(e, r, t)$, where $r \in R$, $t \in E$.

Parameter

relation: Union[List[str], str]

String representation of selected relations.

tail_entity: Union[List[str], str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

`predict_missing_relations(head_entity: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False) → Tuple`

Given a head entity and a tail entity, return top k ranked relations.

$\operatorname{argmax}_{\{r \in R\}} f(h, r, t)$, where $h, t \in E$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_tail_entity (*head_entity: List[str] | str, relation: List[str] | str,*
within: List[str] = None, batch_size=2, topk=1, return_indices=False) → torch.FloatTensor

Given a head entity and a relation, return top k ranked entities

$\text{argmax}_{\{e \in E\}} f(h, r, e)$, where h in E and r in R .

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

Returns: Tuple

scores

predict (*, *h: List[str] | str | None = None, r: List[str] | str | None = None, t: List[str] | str | None = None,*
within: List[str] | None = None, logits: bool = True) → torch.FloatTensor

Predict scores for triples or missing triple elements.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.
- **t** – Tail entity/entities. None to predict tails.
- **within** – Optional list of entities to restrict predictions to.
- **logits** – If True, return raw scores. If False, return sigmoid scores (0-1).

Returns

- Single triple (h, r, t): scalar score
- Missing element: vector of all possible scores

Return type

torch.FloatTensor of scores. Shape depends on the query type

Raises

- **TypeError** – If inputs are not strings or lists of strings.
- **ValueError** – If a required argument for the query type is missing.

Examples

```
>>> # Score a specific triple
>>> model.predict(h="Mongolia", r="isLocatedIn", t="Asia", logits=False)
tensor(0.9523)
```

```
>>> # Get scores for all possible tail entities
>>> model.predict(h="Mongolia", r="isLocatedIn", t=None)
tensor([0.21, 0.95, 0.03, ...]) # One score per entity
```

predict_topk (*, *h*: str | List[str] | None = None, *r*: str | List[str] | None = None, *t*: str | List[str] | None = None, *topk*: int = 10, *within*: List[str] | None = None, *batch_size*: int = 1024) → List[Tuple[str, float]] | List[List[Tuple[str, float]]]

Predict top-k missing items in a given triple pattern.

Parameters

- **h** – Head entity/entities. None to predict heads.
- **r** – Relation/relations. None to predict relations.
- **t** – Tail entity/entities. None to predict tails.
- **topk** – Number of top predictions to return.
- **within** – Optional list of entities to restrict predictions to.
- **batch_size** – Batch size for processing multiple queries.

Returns

List[(item, score), ...] of length topk. For batch query: List of such lists, one per query.

Return type

For single query

Raises

- **TypeError** – If h, r, or t is not a str or list of str.
- **ValueError** – If the required arguments for a query type are None.

Examples

```
>>> model.predict_topk(h=["Mongolia"], r=["isLocatedIn"], topk=3)
[('Asia', 0.99), ('Europe', 0.02), ...]
```

```
>>> model.predict_topk(r=["isLocatedIn"], t=["Asia"], topk=5)
[('Mongolia', 0.85), ('China', 0.82), ...]
```

triple_score (*h*: List[str] | str = None, *r*: List[str] | str = None, *t*: List[str] | str = None, *logits*=False) → torch.FloatTensor

Predict triple score

Parameter

head_entity: List[str]

String representation of selected entities.

relation: List[str]

String representation of selected relations.

tail_entity: List[str]

String representation of selected entities.

logits: bool

If logits is True, unnormalized score returned

Returns: Tuple

pytorch tensor of triple score

return_multi_hop_query_results (*aggregated_query_for_all_entities*, *k*: int, *only_scores*)

single_hop_query_answering (*query*: tuple, *only_scores*: bool = True, *k*: int = None, *use_logits*: bool = True)

answer_multi_hop_query (*query_type*: str | None = None, *query*: Tuple[str | Tuple[str, str], Ellipsis] | None = None, *queries*: List[Tuple[str | Tuple[str, str], Ellipsis]] | None = None, *tnorm*: str = 'prod', *neg_norm*: str = 'standard', *lambda_*: float = 0.0, *k*: int = 10, *only_scores*: bool = False, *use_logits*: bool = True)
→ List[Tuple[str, torch.Tensor]] | List[List[Tuple[str, torch.Tensor]]]

Answer multi-hop EPFO (Existential Positive First-Order) queries.

Supports 9 query types: 1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up. See docs/guides/multi_hop_queries.md for detailed query patterns.

Parameters

- **query_type** – Query pattern name. One of: - 1p: (e, (r,)) # One-hop - 2p: (e, (r1, r2)) # Two-hop - 3p: (e, (r1, r2, r3)) # Three-hop - 2i: ((e1, (r1,)), (e2, (r2,))) # Two-way intersection - 3i: ((e1, (r1,)), (e2, (r2,)), (e3, (r3,))) # Three-way intersection - ip: (((e1, (r1,)), (e2, (r2,))), (r3,)) # Intersection + projection - pi: ((e, (r1, r2)), (r3,)) # Projection + intersection (2i meets 2p) - 2u: ((e1, (r1,)), (e2, (r2,))) # Two-way union - up: ((e, (r1, r2)), (e, (r3,))) # Union + projection
- **query** – Single query tuple matching the query_type pattern.
- **queries** – Batch of queries. If provided, query must be None.
- **tnorm** – T-norm for intersection/union. Options: “prod”, “min”.
- **neg_norm** – Negation norm. Options: “standard”, “sugeno”, “yager”.
- **lambda** – Parameter for sugeno and yager negation (0.0-1.0).
- **k** – Number of top answer entities to return.
- **only_scores** – If True, return only scores tensor. If False, return (entity, score) tuples.
- **use_logits** – If True, use raw model logits. If False, use sigmoid probabilities.

Returns

List[(entity, score), ...] of top-k answers. For batch queries: List of such lists, one per query.

Return type

For single query

Raises

- **ValueError** – If query_type is not in { 1p, 2p, 3p, 2i, 3i, ip, pi, 2u, up }.

- **AssertionError** – If query structure doesn't match query_type pattern.

Examples

```
>>> # 1p: Find entities located in Asia
>>> model.answer_multi_hop_query(
...     query_type="1p",
...     query=("Asia", ("isLocatedIn",)),
...     k=5
... )
[("Mongolia", 0.92), ("China", 0.89), ...]
```

```
>>> # 2p: Two-hop query (e.g., "capital of countries in Europe")
>>> model.answer_multi_hop_query(
...     query_type="2p",
...     query=("Europe", ("isLocatedIn", "hasCapital")),
...     k=3
... )
[("Paris", 0.85), ("Berlin", 0.82), ...]
```

```
>>> # 2i: Intersection query
>>> model.answer_multi_hop_query(
...     query_type="2i",
...     query=((("Asia", ("isLocatedIn",)), ("Mountains", ("hasGeography",))),
...     k=5
... )
[("Nepal", 0.78), ("Tibet", 0.65), ...]
```

➡ See also

- docs/guides/multi_hop_queries.md: Complete guide with all query patterns
- tests/test_answer_multi_hop_query.py: Usage examples

find_missing_triples (confidence: float, entities: List[str] = None, relations: List[str] = None, topk: int = 10, at_most: int = sys.maxsize) → Set

Find missing triples

Iterative over a set of entities E and a set of relation R :

orall e in E and orall r in R f(e,r,x)

Return (e,r,x)

otin G and f(e,r,x) > confidence

confidence: float

A threshold for an output of a sigmoid function given a triple.

topk: int

Highest ranked k item to select triples with f(e,r,x) > confidence .

at_most: int

Stop after finding at_most missing triples

$\{(e,r,x) \mid f(e,r,x) > \text{confidence_land}(e,r,x)\}$

otin G

predict_literals (*entity*: *List[str] | str = None*, *attribute*: *List[str] | str = None*,
denormalize_preds: *bool = True*) $\rightarrow$ *numpy.ndarray*

Predicts literal values for given entities and attributes.

Parameters

- **entity** (*Union[List[str], str]*) – Entity or list of entities to predict literals for.
- **attribute** (*Union[List[str], str]*) – Attribute or list of attributes to predict literals for.
- **denormalize_preds** (*bool*) – If True, denormalizes the predictions.

Returns

Predictions for the given entities and attributes.

Return type

numpy.ndarray

class *dicee.QueryGenerator* (*train_path*: *str*, *val_path*: *str*, *test_path*: *str*, *ent2id*: *Dict = None*,
rel2id: *Dict = None*, *seed*: *int = 1*, *gen_valid*: *bool = False*, *gen_test*: *bool = True*)

train_path

val_path

test_path

gen_valid = *False*

gen_test = *True*

seed = *1*

max_ans_num = *1000000.0*

mode

ent2id = *None*

rel2id: *Dict* = *None*

ent_in: *Dict*

ent_out: *Dict*

query_name_to_struct

list2tuple (*list_data*)

tuple2list (*x*: *List | Tuple*) $\rightarrow$ *List | Tuple*

Convert a nested tuple to a nested list.

set_global_seed (*seed*: *int*)

Set seed

```

construct_graph (paths: List[str]) → Tuple[Dict, Dict]
    Construct graph from triples Returns dicts with incoming and outgoing edges

fill_query (query_structure: List[str | List], ent_in: Dict, ent_out: Dict, answer: int) → bool
    Private method for fill_query logic.

achieve_answer (query: List[str | List], ent_in: Dict, ent_out: Dict) → set
    Private method for achieve_answer logic. @TODO: Document the code

write_links (ent_out, small_ent_out)

ground_queries (query_structure: List[str | List], ent_in: Dict, ent_out: Dict, small_ent_in: Dict,
    small_ent_out: Dict, gen_num: int, query_name: str)
    Generating queries and achieving answers

unmap (query_type, queries, tp_answers, fp_answers, fn_answers)

unmap_query (query_structure, query, id2ent, id2rel)

generate_queries (query_struct: List, gen_num: int, query_type: str)
    Passing incoming and outgoing edges to ground queries depending on mode [train valid or text] and getting
    queries and answers in return @ TODO: create a class for each single query struct

save_queries (query_type: str, gen_num: int, save_path: str)

abstractmethod load_queries (path)

get_queries (query_type: str, gen_num: int)

static save_queries_and_answers (path: str, data: List[Tuple[str, Tuple[collections.defaultdict]]])
    → None
    Save Queries into Disk

static load_queries_and_answers (path: str) → List[Tuple[str, Tuple[collections.defaultdict]]]
    Load Queries from Disk to Memory

class dicee.DICE_Trainer (args, is_continual_training: bool, storage_path, evaluator=None)

    DICE_Trainer implement
    1- Pytorch Lightning trainer (https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html)
    2- Multi-GPU Trainer(https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html)
    3- CPU Trainer

    args
    is_continual_training:bool
    storage_path:str
    evaluator:
    report:dict

    report

    args

    trainer = None

    is_continual_training

```

storage_path

evaluator = None

form_of_labelling = None

continual_start (*knowledge_graph*)

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

initialize_trainer (*callbacks: List*)

→ `lightning.Trainer` | `diccee.trainer.model_parallelism.TensorParallel` | `diccee.trainer.torch_trainer.TorchTrainer` | `diccee.`

Initialize Trainer from input arguments

initialize_or_load_model ()

init_dataloader (*dataset: torch.utils.data.Dataset*) → `torch.utils.data.DataLoader`

init_dataset () → `torch.utils.data.Dataset`

start (*knowledge_graph: diccee.knowledge_graph.KG* | *numpy.memmap*)

→ `Tuple[diccee.models.base_model.BaseKGE, str]`

Start the training

- (1) Initialize Trainer
- (2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

k_fold_cross_validation (*dataset*) → `Tuple[diccee.models.base_model.BaseKGE, str]`

Perform K-fold Cross-Validation

1. Obtain K train and test splits.
2. **For each split,**
 - 2.1 initialize trainer and model
 - 2.2. Train model with configuration provided in args.
 - 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.
3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

```
class dicee.Evaluator(args, is_continual_training: bool = False)
```

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset*, *trained_model*, *form_of_labelling*: str, *during_training*: bool = False) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set*, *valid_set*=None, *test_set*=None, *trained_model*)
→ None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set*=None, *valid_set*=None, *test_set*=None, *ordered_bpe_entities*, *trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, *raw_train_set*, *raw_valid_set*=None, *raw_test_set*=None, *trained_model*, *form_of_labelling*) → None

Evaluate Byte model with generation.

eval_with_bpe_vs_all (*, *raw_train_set*, *raw_valid_set*=None, *raw_test_set*=None, *trained_model*, *form_of_labelling*) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, *train_set*, *valid_set*=None, *test_set*=None, *trained_model*, *form_of_labelling*)
→ None

Evaluate with KvsAll or lvsAll scoring.

evaluate_lp_k_vs_all (*model*, *triple_idx*, *info*: str | None = None, *form_of_labelling*: str | None = None) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – 'EntityPrediction' or 'RelationPrediction'.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

evaluate_lp_with_byte (*model*, *triples*: List[List[str]], *info*: str | None = None) → Dict[str, float]

Evaluate ByteE model with text generation.

Parameters

- **model** – ByteE model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (*model*, *triples*: List[List[str]], *info*: str | None = None, *form_of_labelling*: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (*model*, *triple_idx*, *info*: str) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: str) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: numpy.ndarray, *form_of_labelling*: str) → Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

```
dicee.__version__ = '0.3.3'
```

Python Module Index

d

- `dicee`, 12
- `dicee.__main__`, 12
- `dicee.abstracts`, 12
- `dicee.analyse_experiments`, 20
- `dicee.callbacks`, 22
- `dicee.config`, 28
- `dicee.dataset_classes`, 32
- `dicee.eval_static_funcs`, 41
- `dicee.evaluation`, 44
 - `dicee.evaluation.ensemble`, 45
 - `dicee.evaluation.evaluator`, 46
 - `dicee.evaluation.link_prediction`, 49
 - `dicee.evaluation.literal_prediction`, 52
 - `dicee.evaluation.utils`, 53
- `dicee.evaluator`, 62
- `dicee.executer`, 66
- `dicee.knowledge_graph`, 69
- `dicee.knowledge_graph_embeddings`, 71
- `dicee.models`, 77
 - `dicee.models.adopt`, 77
 - `dicee.models.base_model`, 86
 - `dicee.models.clifford`, 95
 - `dicee.models.complex`, 103
 - `dicee.models.dualE`, 108
 - `dicee.models.ensemble`, 109
 - `dicee.models.fsdps_models`, 110
 - `dicee.models.function_space`, 112
 - `dicee.models.literal`, 117
 - `dicee.models.octonion`, 118
 - `dicee.models.pykeen_models`, 121
 - `dicee.models.quaternion`, 122
 - `dicee.models.real`, 125
 - `dicee.models.static_funcs`, 133
 - `dicee.models.transformers`, 134
- `dicee.query_generator`, 211
- `dicee.read_preprocess_save_load_kg`, 213
- `dicee.read_preprocess_save_load_kg.preprocess`, 213
- `dicee.read_preprocess_save_load_kg.read_from_disk`, 214
- `dicee.read_preprocess_save_load_kg.save_load_disk`, 214
- `dicee.read_preprocess_save_load_kg.util`, 215
- `dicee.sanity_checkers`, 220
- `dicee.scripts`, 220
 - `dicee.scripts.index_serve`, 220
 - `dicee.scripts.run`, 222
- `dicee.static_funcs`, 222
 - `dicee.static_funcs_training`, 227
 - `dicee.static_preprocess_funcs`, 228
- `dicee.trainer`, 229
 - `dicee.trainer.auto_batch_finder`, 229
 - `dicee.trainer.dice_trainer`, 230
 - `dicee.trainer.model_parallelism`, 232
 - `dicee.trainer.torch_trainer`, 233
 - `dicee.trainer.torch_trainer_ddp`, 235
 - `dicee.trainer.torch_trainer_fsdp`, 236
 - `dicee.weight_averaging`, 239

Index

Non-alphabetical

`__call__()` (*dicee.models.base_model.IdentityClass method*), 94
`__call__()` (*dicee.models.ensemble.EnsembleKGE method*), 110
`__call__()` (*dicee.models.IdentityClass method*), 153, 179, 186
`__getitem__()` (*dicee.dataset_classes.AllvsAll method*), 35
`__getitem__()` (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 33
`__getitem__()` (*dicee.dataset_classes.FixedNegSampleDataset method*), 39
`__getitem__()` (*dicee.dataset_classes.FSDP1vsSampleDataset method*), 36
`__getitem__()` (*dicee.dataset_classes.KvsAll method*), 36
`__getitem__()` (*dicee.dataset_classes.KvsSampleDataset method*), 37
`__getitem__()` (*dicee.dataset_classes.LiteralDataset method*), 38
`__getitem__()` (*dicee.dataset_classes.MultiClassClassificationDataset method*), 34
`__getitem__()` (*dicee.dataset_classes.MultiLabelDataset method*), 34
`__getitem__()` (*dicee.dataset_classes.OnevsAllDataset method*), 37
`__getitem__()` (*dicee.dataset_classes.OnevsSample method*), 40
`__getitem__()` (*dicee.dataset_classes.TriplePredictionDataset method*), 41
`__iter__()` (*dicee.config.Namespace method*), 32
`__iter__()` (*dicee.knowledge_graph.KG method*), 71
`__iter__()` (*dicee.models.ensemble.EnsembleKGE method*), 110
`__len__()` (*dicee.dataset_classes.AllvsAll method*), 35
`__len__()` (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 33
`__len__()` (*dicee.dataset_classes.FixedNegSampleDataset method*), 39
`__len__()` (*dicee.dataset_classes.FSDP1vsSampleDataset method*), 36
`__len__()` (*dicee.dataset_classes.KvsAll method*), 36
`__len__()` (*dicee.dataset_classes.KvsSampleDataset method*), 37
`__len__()` (*dicee.dataset_classes.LiteralDataset method*), 38
`__len__()` (*dicee.dataset_classes.MultiClassClassificationDataset method*), 34
`__len__()` (*dicee.dataset_classes.MultiLabelDataset method*), 34
`__len__()` (*dicee.dataset_classes.OnevsAllDataset method*), 37
`__len__()` (*dicee.dataset_classes.OnevsSample method*), 40
`__len__()` (*dicee.dataset_classes.TriplePredictionDataset method*), 40
`__len__()` (*dicee.knowledge_graph.KG method*), 71
`__len__()` (*dicee.models.ensemble.EnsembleKGE method*), 110
`__setstate__()` (*dicee.models.ADOPT method*), 143
`__setstate__()` (*dicee.models.adopt.ADOPT method*), 81
`__str__()` (*dicee.KGE method*), 247
`__str__()` (*dicee.knowledge_graph_embeddings.KGE method*), 72
`__str__()` (*dicee.models.ensemble.EnsembleKGE method*), 110
`__version__` (*in module dicee*), 258

A

`AbstractCallback` (*class in dicee.abstracts*), 18
`AbstractPPECallback` (*class in dicee.abstracts*), 19
`AbstractTrainer` (*class in dicee.abstracts*), 13
`AccumulateEpochLossCallback` (*class in dicee.callbacks*), 23
`achieve_answer()` (*dicee.query_generator.QueryGenerator method*), 212
`achieve_answer()` (*dicee.QueryGenerator method*), 253
`AConEx` (*class in dicee.models*), 171
`AConEx` (*class in dicee.models.complex*), 105
`AConvO` (*class in dicee.models*), 188
`AConvO` (*class in dicee.models.octonion*), 120
`AConvQ` (*class in dicee.models*), 181
`AConvQ` (*class in dicee.models.quaternion*), 124
`adaptive_lr` (*dicee.config.Namespace attribute*), 31
`adaptive_swa` (*dicee.config.Namespace attribute*), 31
`add_new_entity_embeddings()` (*dicee.abstracts.BaseInteractiveKGE method*), 16
`add_noise_rate` (*dicee.config.Namespace attribute*), 29
`add_noise_rate` (*dicee.knowledge_graph.KG attribute*), 70
`add_noisy_triples()` (*in module dicee.static_funcs*), 225
`add_noisy_triples_into_training()` (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk method*), 214
`add_noisy_triples_into_training()` (*dicee.read_preprocess_save_load_kg.ReadFromDisk method*), 219
`add_reciprocal` (*dicee.knowledge_graph.KG attribute*), 70
`ADOPT` (*class in dicee.models*), 141
`ADOPT` (*class in dicee.models.adopt*), 79
`adopt()` (*in module dicee.models.adopt*), 82

ALL_HITS_RANGE (in module *dicee.evaluation.utils*), 54
 AllvsAll (class in *dicee.dataset_classes*), 35
 alphas (*dicee.abstracts.AbstractPPECallback* attribute), 19
 alphas (*dicee.weight_averaging.ASWA* attribute), 239
 analyse () (in module *dicee.analyse_experiments*), 22
 answer_multi_hop_query () (*dicee.KGE* method), 250
 answer_multi_hop_query () (*dicee.knowledge_graph_embeddings.KGE* method), 75
 app (in module *dicee.scripts.index_serve*), 221
 apply_coefficients () (*dicee.models.clifford.DeCaL* method), 102
 apply_coefficients () (*dicee.models.clifford.Keci* method), 96
 apply_coefficients () (*dicee.models.DeCaL* method), 194
 apply_coefficients () (*dicee.models.Keci* method), 190
 apply_reciprocal_or_noise () (in module *dicee.read_preprocess_save_load_kg.util*), 217
 apply_semantic_constraint (*dicee.abstracts.BaseInteractiveKGE* attribute), 14
 apply_unit_norm (*dicee.models.base_model.BaseKGE* attribute), 90
 apply_unit_norm (*dicee.models.BaseKGE* attribute), 149, 154, 166, 175, 182, 197, 202
 args (*dicee.DICE_Trainer* attribute), 253
 args (*dicee.evaluation.Evaluator* attribute), 57
 args (*dicee.evaluation.evaluator.Evaluator* attribute), 46, 47
 args (*dicee.Evaluator* attribute), 255
 args (*dicee.evaluator.Evaluator* attribute), 63, 64
 args (*dicee.Execute* attribute), 245
 args (*dicee.executer.Execute* attribute), 66, 67
 args (*dicee.models.base_model.BaseKGE* attribute), 90
 args (*dicee.models.base_model.IdentityClass* attribute), 94
 args (*dicee.models.BaseKGE* attribute), 148, 154, 166, 175, 182, 197, 202
 args (*dicee.models.ensemble.EnsembleKGE* attribute), 109
 args (*dicee.models.IdentityClass* attribute), 153, 179, 186
 args (*dicee.models.pykeen_models.PykeenKGE* attribute), 121
 args (*dicee.models.PykeenKGE* attribute), 201
 args (*dicee.trainer.DICE_Trainer* attribute), 238
 args (*dicee.trainer.dice_trainer.DICE_Trainer* attribute), 231
 ASWA (class in *dicee.weight_averaging*), 239
 aswa (*dicee.analyse_experiments.Experiment* attribute), 21
 attn (*dicee.models.Block* attribute), 153
 attn (*dicee.models.clifford.TransformerBlock* attribute), 99
 attn (*dicee.models.transformers.Block* attribute), 138
 attn_dropout (*dicee.models.clifford.TransformerSelfAttention* attribute), 100
 attn_dropout (*dicee.models.transformers.SelfAttention* attribute), 136
 attributes (*dicee.abstracts.AbstractTrainer* attribute), 13
 auto_batch_finding (*dicee.config.Namespace* attribute), 31

B

b_h (*dicee.models.MuRE* attribute), 162
 b_h (*dicee.models.real.MuRE* attribute), 130
 b_t (*dicee.models.MuRE* attribute), 162
 b_t (*dicee.models.real.MuRE* attribute), 130
 backend (*dicee.config.Namespace* attribute), 29
 backend (*dicee.knowledge_graph.KG* attribute), 70
 base_weights (*dicee.weight_averaging.TWA* attribute), 244
 BaseInteractiveKGE (class in *dicee.abstracts*), 14
 BaseInteractiveTrainKGE (class in *dicee.abstracts*), 20
 BaseKGE (class in *dicee.models*), 148, 154, 166, 174, 182, 196, 202
 BaseKGE (class in *dicee.models.base_model*), 90
 BaseKGELightning (class in *dicee.models*), 144
 BaseKGELightning (class in *dicee.models.base_model*), 86
 batch_kronecker_product () (*dicee.callbacks.KronE* static method), 26
 batch_size (*dicee.analyse_experiments.Experiment* attribute), 21
 batch_size (*dicee.callbacks.PseudoLabellingCallback* attribute), 25
 batch_size (*dicee.config.Namespace* attribute), 29
 batches_per_epoch (*dicee.callbacks.LRScheduler* attribute), 27
 beta (*dicee.weight_averaging.TWA* attribute), 244
 bias (*dicee.models.CoKEConfig* attribute), 164, 165
 bias (*dicee.models.real.CoKEConfig* attribute), 132
 bias (*dicee.models.transformers.GPTConfig* attribute), 138
 bias (*dicee.models.transformers.LayerNorm* attribute), 136
 Block (class in *dicee.models*), 153

Block (class in *dicее.models.transformers*), 137
 block_size (*dicее.config.Namespace* attribute), 31
 block_size (*dicее.dataset_classes.MultiClassClassificationDataset* attribute), 33
 block_size (*dicее.models.base_model.BaseKGE* attribute), 91
 block_size (*dicее.models.BaseKGE* attribute), 149, 155, 167, 176, 183, 198, 203
 block_size (*dicее.models.CoKEConfig* attribute), 164
 block_size (*dicее.models.real.CoKEConfig* attribute), 132
 block_size (*dicее.models.transformers.GPTConfig* attribute), 138
 blocks (*dicее.models.CoKE* attribute), 165
 blocks (*dicее.models.real.CoKE* attribute), 133
 bn_conv1 (*dicее.models.AConvQ* attribute), 181
 bn_conv1 (*dicее.models.ConvQ* attribute), 181
 bn_conv1 (*dicее.models.quaternion.AConvQ* attribute), 125
 bn_conv1 (*dicее.models.quaternion.ConvQ* attribute), 124
 bn_conv2 (*dicее.models.AConvQ* attribute), 181
 bn_conv2 (*dicее.models.ConvQ* attribute), 181
 bn_conv2 (*dicее.models.quaternion.AConvQ* attribute), 125
 bn_conv2 (*dicее.models.quaternion.ConvQ* attribute), 124
 bn_conv2d (*dicее.models.AConEx* attribute), 171
 bn_conv2d (*dicее.models.AConvO* attribute), 188
 bn_conv2d (*dicее.models.complex.AConEx* attribute), 105
 bn_conv2d (*dicее.models.complex.ConEx* attribute), 104
 bn_conv2d (*dicее.models.ConEx* attribute), 170
 bn_conv2d (*dicее.models.ConvO* attribute), 188
 bn_conv2d (*dicее.models.octonion.AConvO* attribute), 120
 bn_conv2d (*dicее.models.octonion.ConvO* attribute), 120
 BPE_NegativeSamplingDataset (class in *dicее.dataset_classes*), 33
 build_chain_funcs () (*dicее.models.FMult2* method), 208
 build_chain_funcs () (*dicее.models.function_space.FMult2* method), 114
 build_func () (*dicее.models.FMult2* method), 208
 build_func () (*dicее.models.function_space.FMult2* method), 114
 build_projection () (*dicее.weight_averaging.TWA* method), 244
 Byte (class in *dicее.models.transformers*), 134
 byte_pair_encoding (*dicее.analyse_experiments.Experiment* attribute), 21
 byte_pair_encoding (*dicее.config.Namespace* attribute), 31
 byte_pair_encoding (*dicее.knowledge_graph.KG* attribute), 70
 byte_pair_encoding (*dicее.models.base_model.BaseKGE* attribute), 91
 byte_pair_encoding (*dicее.models.BaseKGE* attribute), 149, 155, 167, 176, 183, 197, 203

C

c_attn (*dicее.models.clifford.TransformerSelfAttention* attribute), 99
 c_attn (*dicее.models.transformers.SelfAttention* attribute), 136
 c_fc (*dicее.models.clifford.TransformerMLP* attribute), 100
 c_fc (*dicее.models.transformers.MLP* attribute), 137
 c_proj (*dicее.models.clifford.TransformerMLP* attribute), 100
 c_proj (*dicее.models.clifford.TransformerSelfAttention* attribute), 99
 c_proj (*dicее.models.transformers.MLP* attribute), 137
 c_proj (*dicее.models.transformers.SelfAttention* attribute), 136
 callbacks (*dicее.abstracts.AbstractTrainer* attribute), 13
 callbacks (*dicее.analyse_experiments.Experiment* attribute), 21
 callbacks (*dicее.config.Namespace* attribute), 29
 callbacks (*dicее.trainer.torch_trainer_ddp.NodeTrainer* attribute), 236
 causal (*dicее.models.CoKEConfig* attribute), 164, 165
 causal (*dicее.models.real.CoKEConfig* attribute), 132
 causal (*dicее.models.transformers.GPTConfig* attribute), 138
 causal (*dicее.models.transformers.SelfAttention* attribute), 136
 chain_func () (*dicее.models.FMult* method), 207
 chain_func () (*dicее.models.function_space.FMult* method), 113
 chain_func () (*dicее.models.function_space.GFMult* method), 114
 chain_func () (*dicее.models.GFMult* method), 207
 checkpoint_every_n_epochs (*dicее.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237
 CKeci (class in *dicее.models*), 192
 CKeci (class in *dicее.models.clifford*), 100
 cl_pqr () (*dicее.models.clifford.DeCaL* method), 101
 cl_pqr () (*dicее.models.DeCaL* method), 193
 cleanup () (*dicее.Execute* method), 246
 cleanup () (*dicее.executer.Execute* method), 67

`clifford_multiplication()` (*dicee.models.clifford.Keci method*), 96
`clifford_multiplication()` (*dicee.models.Keci method*), 190
`clip_lambda` (*dicee.models.ADOPT attribute*), 143
`clip_lambda` (*dicee.models.adapt.ADOPT attribute*), 81
`CoKE` (*class in dicee.models*), 165
`CoKE` (*class in dicee.models.real*), 132
`coke_dropout` (*dicee.models.CoKE attribute*), 165
`coke_dropout` (*dicee.models.real.CoKE attribute*), 133
`CoKEConfig` (*class in dicee.models*), 164
`CoKEConfig` (*class in dicee.models.real*), 132
`collate_fn` (*dicee.dataset_classes.AllvsAll attribute*), 35
`collate_fn` (*dicee.dataset_classes.FixedNegSampleDataset attribute*), 39
`collate_fn` (*dicee.dataset_classes.FSDP1vsSampleDataset attribute*), 36
`collate_fn` (*dicee.dataset_classes.KvsAll attribute*), 36
`collate_fn` (*dicee.dataset_classes.KvsSampleDataset attribute*), 37
`collate_fn` (*dicee.dataset_classes.MultiClassClassificationDataset attribute*), 33
`collate_fn` (*dicee.dataset_classes.MultiLabelDataset attribute*), 34
`collate_fn` (*dicee.dataset_classes.OnevsAllDataset attribute*), 37
`collate_fn` (*dicee.dataset_classes.OnevsSample attribute*), 40
`collate_fn()` (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 33
`collate_fn()` (*dicee.dataset_classes.TriplePredictionDataset method*), 41
`collection_name` (*dicee.scripts.index_serve.NeuralSearcher attribute*), 221
`comp_func()` (*dicee.models.function_space.LFMult method*), 116
`comp_func()` (*dicee.models.LFMult method*), 210
`ComplEx` (*class in dicee.models*), 172
`ComplEx` (*class in dicee.models.complex*), 106
`compute_convergence()` (*in module dicee.callbacks*), 25
`compute_func()` (*dicee.models.FMult method*), 207
`compute_func()` (*dicee.models.FMult2 method*), 208
`compute_func()` (*dicee.models.function_space.FMult method*), 113
`compute_func()` (*dicee.models.function_space.FMult2 method*), 114
`compute_func()` (*dicee.models.function_space.GFMult method*), 114
`compute_func()` (*dicee.models.GFMult method*), 207
`compute_metrics_from_ranks()` (*in module dicee.evaluation.utils*), 54
`compute_metrics_from_ranks_simple()` (*in module dicee.evaluation.utils*), 54
`compute_mrr()` (*dicee.weight_averaging.ASWA static method*), 240
`compute_sigma_pp()` (*dicee.models.clifford.DeCaL method*), 102
`compute_sigma_pp()` (*dicee.models.clifford.Keci method*), 96
`compute_sigma_pp()` (*dicee.models.DeCaL method*), 194
`compute_sigma_pp()` (*dicee.models.Keci method*), 189
`compute_sigma_pq()` (*dicee.models.clifford.DeCaL method*), 103
`compute_sigma_pq()` (*dicee.models.clifford.Keci method*), 96
`compute_sigma_pq()` (*dicee.models.DeCaL method*), 195
`compute_sigma_pq()` (*dicee.models.Keci method*), 190
`compute_sigma_pr()` (*dicee.models.clifford.DeCaL method*), 103
`compute_sigma_pr()` (*dicee.models.DeCaL method*), 195
`compute_sigma_qq()` (*dicee.models.clifford.DeCaL method*), 102
`compute_sigma_qq()` (*dicee.models.clifford.Keci method*), 96
`compute_sigma_qq()` (*dicee.models.DeCaL method*), 195
`compute_sigma_qq()` (*dicee.models.Keci method*), 190
`compute_sigma_qr()` (*dicee.models.clifford.DeCaL method*), 103
`compute_sigma_qr()` (*dicee.models.DeCaL method*), 196
`compute_sigma_rr()` (*dicee.models.clifford.DeCaL method*), 103
`compute_sigma_rr()` (*dicee.models.DeCaL method*), 195
`compute_sigmas_multivect()` (*dicee.models.clifford.DeCaL method*), 101
`compute_sigmas_multivect()` (*dicee.models.DeCaL method*), 194
`compute_sigmas_single()` (*dicee.models.clifford.DeCaL method*), 101
`compute_sigmas_single()` (*dicee.models.DeCaL method*), 193
`ConEx` (*class in dicee.models*), 170
`ConEx` (*class in dicee.models.complex*), 104
`config` (*dicee.models.CoKE attribute*), 165
`config` (*dicee.models.real.CoKE attribute*), 133
`config` (*dicee.models.transformers.BytE attribute*), 134
`config` (*dicee.models.transformers.GPT attribute*), 139
`configs` (*dicee.abstracts.BaseInteractiveKGE attribute*), 14
`configure_optimizers()` (*dicee.models.base_model.BaseKGELightning method*), 90
`configure_optimizers()` (*dicee.models.BaseKGELightning method*), 148
`configure_optimizers()` (*dicee.models.transformers.GPT method*), 139

`construct_batch_selected_cl_multivector()` (*dicee.models.clifford.Keci method*), 98
`construct_batch_selected_cl_multivector()` (*dicee.models.Keci method*), 192
`construct_cl_multivector()` (*dicee.models.clifford.DeCaL method*), 102
`construct_cl_multivector()` (*dicee.models.clifford.Keci method*), 97
`construct_cl_multivector()` (*dicee.models.DeCaL method*), 194
`construct_cl_multivector()` (*dicee.models.Keci method*), 190
`construct_dataset()` (*in module dicee.dataset_classes*), 34
`construct_ensemble` (*dicee.abstracts.BaseInteractiveKGE attribute*), 14
`construct_graph()` (*dicee.query_generator.QueryGenerator method*), 212
`construct_graph()` (*dicee.QueryGenerator method*), 252
`construct_input_and_output()` (*dicee.abstracts.BaseInteractiveKGE method*), 17
`construct_multi_coeff()` (*dicee.models.function_space.LFMMult method*), 116
`construct_multi_coeff()` (*dicee.models.LFMMult method*), 209
`continual_learning` (*dicee.config.Namespace attribute*), 31
`continual_start()` (*dicee.DICE_Trainer method*), 254
`continual_start()` (*dicee.executer.ContinuousExecute method*), 68
`continual_start()` (*dicee.trainer.DICE_Trainer method*), 238
`continual_start()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 231
`continual_training_setup_executor()` (*in module dicee.static_funcs*), 226
`ContinuousExecute` (*class in dicee.executer*), 68
`conv2d` (*dicee.models.AConEx attribute*), 171
`conv2d` (*dicee.models.AConvO attribute*), 188
`conv2d` (*dicee.models.AConvQ attribute*), 181
`conv2d` (*dicee.models.complex.AConEx attribute*), 105
`conv2d` (*dicee.models.complex.ConEx attribute*), 104
`conv2d` (*dicee.models.ConEx attribute*), 170
`conv2d` (*dicee.models.ConvO attribute*), 188
`conv2d` (*dicee.models.ConvQ attribute*), 181
`conv2d` (*dicee.models.octonion.AConvO attribute*), 120
`conv2d` (*dicee.models.octonion.ConvO attribute*), 119
`conv2d` (*dicee.models.quaternion.AConvQ attribute*), 125
`conv2d` (*dicee.models.quaternion.ConvQ attribute*), 124
`ConvO` (*class in dicee.models*), 187
`ConvO` (*class in dicee.models.octonion*), 119
`ConvQ` (*class in dicee.models*), 180
`ConvQ` (*class in dicee.models.quaternion*), 124
`count_triples()` (*in module dicee.read_preprocess_save_load_kg.util*), 217
`create_and_store_kg()` (*dicee.Execute method*), 246
`create_and_store_kg()` (*dicee.executer.Execute method*), 67
`create_constraints()` (*in module dicee.read_preprocess_save_load_kg.util*), 218
`create_constraints()` (*in module dicee.static_preprocess_funcs*), 229
`create_experiment_folder()` (*in module dicee.static_funcs*), 226
`create_fsdp_sharded_model_class()` (*in module dicee.models.fsdp_models*), 112
`create_hits_dict()` (*in module dicee.evaluation.utils*), 55
`create_random_data()` (*dicee.callbacks.PseudoLabellingCallback method*), 25
`create_recipriocal_triples()` (*in module dicee.read_preprocess_save_load_kg.util*), 218
`create_recipriocal_triples()` (*in module dicee.static_funcs*), 223
`create_vector_database()` (*dicee.KGE method*), 247
`create_vector_database()` (*dicee.knowledge_graph_embeddings.KGE method*), 72
`crop_block_size()` (*dicee.models.transformers.GPT method*), 139
`ctx` (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 236
`ctx` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
`current_epoch` (*dicee.weight_averaging.EMA attribute*), 243
`current_epoch` (*dicee.weight_averaging.SWA attribute*), 241
`current_epoch` (*dicee.weight_averaging.SWAG attribute*), 242
`current_epoch` (*dicee.weight_averaging.TWA attribute*), 243
`cycle_length` (*dicee.callbacks.LRScheduler attribute*), 27

D

`data_module` (*dicee.callbacks.PseudoLabellingCallback attribute*), 25
`data_property_embeddings` (*dicee.models.literal.LiteralEmbeddings attribute*), 118
`dataset_dir` (*dicee.config.Namespace attribute*), 28
`dataset_dir` (*dicee.knowledge_graph.KG attribute*), 69, 70
`dataset_sanity_checking()` (*in module dicee.read_preprocess_save_load_kg.util*), 218
`DeCaL` (*class in dicee.models*), 193
`DeCaL` (*class in dicee.models.clifford*), 100
`decay` (*dicee.weight_averaging.EMA attribute*), 242

- `decide()` (*dicee.weight_averaging.ASWA method*), 240
- `default_eval_model` (*dicee.callbacks.PeriodicEvalCallback attribute*), 27
- `DEFAULT_HITS_RANGE` (*in module dicee.evaluation.utils*), 54
- `defer_large_embeddings` (*dicee.models.base_model.BaseKGE attribute*), 91
- `defer_large_embeddings` (*dicee.models.BaseKGE attribute*), 149, 155, 167, 176, 183, 198, 203
- `degree` (*dicee.models.function_space.LFMult attribute*), 115
- `degree` (*dicee.models.LFMult attribute*), 209
- `denormalize()` (*dicee.dataset_classes.LiteralDataset static method*), 38
- `describe()` (*dicee.knowledge_graph.KG method*), 71
- `description_of_input` (*dicee.knowledge_graph.KG attribute*), 71
- `deviations` (*dicee.weight_averaging.SWAG attribute*), 242
- `device` (*dicee.models.fsdp_models.RowWiseShardedEmbedding attribute*), 111
- `device` (*dicee.models.literal.LiteralEmbeddings property*), 118
- `device` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
- `DICE_Trainer` (*class in dicee*), 253
- `DICE_Trainer` (*class in dicee.trainer*), 238
- `DICE_Trainer` (*class in dicee.trainer.dice_trainer*), 231
- `dicee`
 - module, 12
- `dicee.__main__`
 - module, 12
- `dicee.abstracts`
 - module, 12
- `dicee.analyse_experiments`
 - module, 20
- `dicee.callbacks`
 - module, 22
- `dicee.config`
 - module, 28
- `dicee.dataset_classes`
 - module, 32
- `dicee.eval_static_funcs`
 - module, 41
- `dicee.evaluation`
 - module, 44
- `dicee.evaluation.ensemble`
 - module, 45
- `dicee.evaluation.evaluator`
 - module, 46
- `dicee.evaluation.link_prediction`
 - module, 49
- `dicee.evaluation.literal_prediction`
 - module, 52
- `dicee.evaluation.utils`
 - module, 53
- `dicee.evaluator`
 - module, 62
- `dicee.executer`
 - module, 66
- `dicee.knowledge_graph`
 - module, 69
- `dicee.knowledge_graph_embeddings`
 - module, 71
- `dicee.models`
 - module, 77
- `dicee.models.adopt`
 - module, 77
- `dicee.models.base_model`
 - module, 86
- `dicee.models.clifford`
 - module, 95
- `dicee.models.complex`
 - module, 103
- `dicee.models.dualE`
 - module, 108
- `dicee.models.ensemble`
 - module, 109
- `dicee.models.fsdp_models`

- module, 110
- dicee.models.function_space
 - module, 112
- dicee.models.literal
 - module, 117
- dicee.models.octonion
 - module, 118
- dicee.models.pykeen_models
 - module, 121
- dicee.models.quaternion
 - module, 122
- dicee.models.real
 - module, 125
- dicee.models.static_funcs
 - module, 133
- dicee.models.transformers
 - module, 134
- dicee.query_generator
 - module, 211
- dicee.read_preprocess_save_load_kg
 - module, 213
- dicee.read_preprocess_save_load_kg.preprocess
 - module, 213
- dicee.read_preprocess_save_load_kg.read_from_disk
 - module, 214
- dicee.read_preprocess_save_load_kg.save_load_disk
 - module, 214
- dicee.read_preprocess_save_load_kg.util
 - module, 215
- dicee.sanity_checkers
 - module, 220
- dicee.scripts
 - module, 220
- dicee.scripts.index_serve
 - module, 220
- dicee.scripts.run
 - module, 222
- dicee.static_funcs
 - module, 222
- dicee.static_funcs_training
 - module, 227
- dicee.static_preprocess_funcs
 - module, 228
- dicee.trainer
 - module, 229
- dicee.trainer.auto_batch_finder
 - module, 229
- dicee.trainer.dice_trainer
 - module, 230
- dicee.trainer.model_parallelism
 - module, 232
- dicee.trainer.torch_trainer
 - module, 233
- dicee.trainer.torch_trainer_ddp
 - module, 235
- dicee.trainer.torch_trainer_fsdp
 - module, 236
- dicee.weight_averaging
 - module, 239
- discrete_points (*dicee.models.FMult2 attribute*), 208
- discrete_points (*dicee.models.function_space.FMult2 attribute*), 114
- dist_func (*dicee.models.Pyke attribute*), 164
- dist_func (*dicee.models.real.Pyke attribute*), 131
- DistMult (*class in dicee.models*), 158
- DistMult (*class in dicee.models.real*), 126
- distributed (*dicee.Execute attribute*), 245
- distributed (*dicee.executer.Execute attribute*), 66, 67
- domain_constraints_per_rel (*dicee.evaluation.Evaluator attribute*), 57

domain_constraints_per_rel (*dicee.evaluation.evaluator.Evaluator attribute*), 47
 domain_constraints_per_rel (*dicee.Evaluator attribute*), 255
 domain_constraints_per_rel (*dicee.evaluator.Evaluator attribute*), 64
 download_file() (in module *dicee.static_funcs*), 226
 download_files_from_url() (in module *dicee.static_funcs*), 226
 download_pretrained_model() (in module *dicee.static_funcs*), 226
 dropout (*dicee.models.clifford.TransformerMLP attribute*), 100
 dropout (*dicee.models.clifford.TransformerSelfAttention attribute*), 100
 dropout (*dicee.models.CoKEConfig attribute*), 164, 165
 dropout (*dicee.models.literal.LiteralEmbeddings attribute*), 117, 118
 dropout (*dicee.models.real.CoKEConfig attribute*), 132
 dropout (*dicee.models.transformers.GPTConfig attribute*), 138
 dropout (*dicee.models.transformers.MLP attribute*), 137
 dropout (*dicee.models.transformers.SelfAttention attribute*), 136
 DualE (class in *dicee.models*), 210
 DualE (class in *dicee.models.dualE*), 108
 dummy_eval() (*dicee.evaluation.Evaluator method*), 59
 dummy_eval() (*dicee.evaluation.evaluator.Evaluator method*), 49
 dummy_eval() (*dicee.Evaluator method*), 257
 dummy_eval() (*dicee.evaluator.Evaluator method*), 65
 dummy_id (*dicee.knowledge_graph.KG attribute*), 71
 during_training (*dicee.evaluation.Evaluator attribute*), 57
 during_training (*dicee.evaluation.evaluator.Evaluator attribute*), 46, 47
 during_training (*dicee.Evaluator attribute*), 255
 during_training (*dicee.evaluator.Evaluator attribute*), 63, 64

E

ee_vocab (*dicee.evaluation.Evaluator attribute*), 56, 57
 ee_vocab (*dicee.evaluation.evaluator.Evaluator attribute*), 46, 47
 ee_vocab (*dicee.Evaluator attribute*), 255
 ee_vocab (*dicee.evaluator.Evaluator attribute*), 63
 efficient_zero_grad() (in module *dicee.evaluation.utils*), 55
 efficient_zero_grad() (in module *dicee.static_funcs_training*), 228
 EMA (class in *dicee.weight_averaging*), 242
 ema (*dicee.config.Namespace attribute*), 31
 ema_c_epochs (*dicee.weight_averaging.EMA attribute*), 242
 ema_model (*dicee.weight_averaging.EMA attribute*), 242
 ema_start_epoch (*dicee.weight_averaging.EMA attribute*), 242
 ema_update() (*dicee.weight_averaging.EMA static method*), 243
 embedding_dim (*dicee.analyse_experiments.Experiment attribute*), 21
 embedding_dim (*dicee.config.Namespace attribute*), 29
 embedding_dim (*dicee.models.base_model.BaseKGE attribute*), 90
 embedding_dim (*dicee.models.BaseKGE attribute*), 148, 154, 166, 175, 182, 197, 202
 embedding_dim (*dicee.models.fsdps_models.RowWiseShardedEmbedding attribute*), 111
 embedding_dim (*dicee.models.literal.LiteralEmbeddings attribute*), 117
 embedding_dims (*dicee.models.literal.LiteralEmbeddings attribute*), 117
 enable_log (in module *dicee.static_preprocess_funcs*), 229
 enc (*dicee.knowledge_graph.KG attribute*), 70
 end() (*dicee.Execute method*), 246
 end() (*dicee.executer.Execute method*), 68
 end_row (*dicee.models.fsdps_models.RowWiseShardedEmbedding attribute*), 111
 EnsembleKGE (class in *dicee.models.ensemble*), 109
 ent2id (*dicee.query_generator.QueryGenerator attribute*), 212
 ent2id (*dicee.QueryGenerator attribute*), 252
 ent_in (*dicee.query_generator.QueryGenerator attribute*), 212
 ent_in (*dicee.QueryGenerator attribute*), 252
 ent_out (*dicee.query_generator.QueryGenerator attribute*), 212
 ent_out (*dicee.QueryGenerator attribute*), 252
 entities_str (*dicee.knowledge_graph.KG property*), 71
 entity_embeddings (*dicee.models.AConvQ attribute*), 181
 entity_embeddings (*dicee.models.ConvQ attribute*), 181
 entity_embeddings (*dicee.models.DualE attribute*), 210
 entity_embeddings (*dicee.models.dualE.DualE attribute*), 108
 entity_embeddings (*dicee.models.FMult attribute*), 206
 entity_embeddings (*dicee.models.FMult2 attribute*), 208
 entity_embeddings (*dicee.models.function_space.FMult attribute*), 113
 entity_embeddings (*dicee.models.function_space.FMult2 attribute*), 114

entity_embeddings (*dicee.models.function_space.GFMult attribute*), 113
entity_embeddings (*dicee.models.function_space.LFMult attribute*), 115
entity_embeddings (*dicee.models.function_space.LFMult1 attribute*), 115
entity_embeddings (*dicee.models.GFMult attribute*), 207
entity_embeddings (*dicee.models.LFMult attribute*), 209
entity_embeddings (*dicee.models.LFMult1 attribute*), 208
entity_embeddings (*dicee.models.literal.LiteralEmbeddings attribute*), 117, 118
entity_embeddings (*dicee.models.pykeen_models.PykeenKGE attribute*), 121
entity_embeddings (*dicee.models.PykeenKGE attribute*), 201
entity_embeddings (*dicee.models.quaternion.AConvQ attribute*), 125
entity_embeddings (*dicee.models.quaternion.ConvQ attribute*), 124
entity_optim_device (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
entity_to_idx (*dicee.dataset_classes.LiteralDataset attribute*), 38
entity_to_idx (*dicee.knowledge_graph.KG attribute*), 69, 70
entity_to_idx (*dicee.scripts.index_serve.NeuralSearcher attribute*), 221
epoch_count (*dicee.abstracts.AbstractPPECallback attribute*), 19
epoch_count (*dicee.weight_averaging.ASWA attribute*), 239
epoch_counter (*dicee.callbacks.Eval attribute*), 25
epoch_counter (*dicee.callbacks.KGESaveCallback attribute*), 24
epoch_counter (*dicee.callbacks.PeriodicEvalCallback attribute*), 27
epoch_ratio (*dicee.callbacks.Eval attribute*), 25
er_vocab (*dicee.evaluation.Evaluator attribute*), 56, 57
er_vocab (*dicee.evaluation.evaluator.Evaluator attribute*), 46, 47
er_vocab (*dicee.Evaluator attribute*), 255
er_vocab (*dicee.evaluator.Evaluator attribute*), 63
estimate_mfu() (*dicee.models.transformers.GPT method*), 139
estimate_q() (*in module dicee.callbacks*), 25
Eval (*class in dicee.callbacks*), 25
eval() (*dicee.evaluation.Evaluator method*), 57
eval() (*dicee.evaluation.evaluator.Evaluator method*), 47
eval() (*dicee.Evaluator method*), 256
eval() (*dicee.evaluator.Evaluator method*), 64
eval() (*dicee.models.ensemble.EnsembleKGE method*), 110
eval_at_epochs (*dicee.config.Namespace attribute*), 31
eval_epochs (*dicee.callbacks.PeriodicEvalCallback attribute*), 27
eval_every_n_epochs (*dicee.config.Namespace attribute*), 31
eval_lp_performance() (*dicee.KGE method*), 247
eval_lp_performance() (*dicee.knowledge_graph_embeddings.KGE method*), 72
eval_model (*dicee.config.Namespace attribute*), 30
eval_model (*dicee.knowledge_graph.KG attribute*), 70
eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.evaluation.Evaluator method*), 58
eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.evaluation.evaluator.Evaluator method*), 47
eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.Evaluator method*), 256
eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.evaluator.Evaluator method*), 64
eval_rank_of_head_and_tail_entity() (*dicee.evaluation.Evaluator method*), 58
eval_rank_of_head_and_tail_entity() (*dicee.evaluation.evaluator.Evaluator method*), 47
eval_rank_of_head_and_tail_entity() (*dicee.Evaluator method*), 256
eval_rank_of_head_and_tail_entity() (*dicee.evaluator.Evaluator method*), 64
eval_with_bpe_vs_all() (*dicee.evaluation.Evaluator method*), 58
eval_with_bpe_vs_all() (*dicee.evaluation.evaluator.Evaluator method*), 48
eval_with_bpe_vs_all() (*dicee.Evaluator method*), 256
eval_with_bpe_vs_all() (*dicee.evaluator.Evaluator method*), 64
eval_with_byte() (*dicee.evaluation.Evaluator method*), 58
eval_with_byte() (*dicee.evaluation.evaluator.Evaluator method*), 47
eval_with_byte() (*dicee.Evaluator method*), 256
eval_with_byte() (*dicee.evaluator.Evaluator method*), 64
eval_with_data() (*dicee.evaluation.Evaluator method*), 59
eval_with_data() (*dicee.evaluation.evaluator.Evaluator method*), 49
eval_with_data() (*dicee.Evaluator method*), 257
eval_with_data() (*dicee.evaluator.Evaluator method*), 65
eval_with_vs_all() (*dicee.evaluation.Evaluator method*), 58
eval_with_vs_all() (*dicee.evaluation.evaluator.Evaluator method*), 48
eval_with_vs_all() (*dicee.Evaluator method*), 256
eval_with_vs_all() (*dicee.evaluator.Evaluator method*), 64
evaluate() (*in module dicee.static_funcs*), 226
evaluate_bpe_lp() (*in module dicee.evaluation*), 59
evaluate_bpe_lp() (*in module dicee.evaluation.link_prediction*), 51
evaluate_bpe_lp() (*in module dicee.static_funcs_training*), 227

evaluate_ensemble_link_prediction_performance() (in module *dicce.eval_static_funcs*), 41
 evaluate_ensemble_link_prediction_performance() (in module *dicce.evaluation*), 56
 evaluate_ensemble_link_prediction_performance() (in module *dicce.evaluation.ensemble*), 45
 evaluate_link_prediction_performance() (in module *dicce.eval_static_funcs*), 42
 evaluate_link_prediction_performance() (in module *dicce.evaluation*), 60
 evaluate_link_prediction_performance() (in module *dicce.evaluation.link_prediction*), 50
 evaluate_link_prediction_performance_with_bpe() (in module *dicce.eval_static_funcs*), 42
 evaluate_link_prediction_performance_with_bpe() (in module *dicce.evaluation*), 60
 evaluate_link_prediction_performance_with_bpe() (in module *dicce.evaluation.link_prediction*), 50
 evaluate_link_prediction_performance_with_bpe_reciprocals() (in module *dicce.eval_static_funcs*), 42
 evaluate_link_prediction_performance_with_bpe_reciprocals() (in module *dicce.evaluation*), 60
 evaluate_link_prediction_performance_with_bpe_reciprocals() (in module *dicce.evaluation.link_prediction*), 50
 evaluate_link_prediction_performance_with_reciprocals() (in module *dicce.eval_static_funcs*), 43
 evaluate_link_prediction_performance_with_reciprocals() (in module *dicce.evaluation*), 60
 evaluate_link_prediction_performance_with_reciprocals() (in module *dicce.evaluation.link_prediction*), 50
 evaluate_literal_prediction() (in module *dicce.eval_static_funcs*), 43
 evaluate_literal_prediction() (in module *dicce.evaluation*), 62
 evaluate_literal_prediction() (in module *dicce.evaluation.literal_prediction*), 52
 evaluate_lp() (*dicce.evaluation.Evaluator* method), 59
 evaluate_lp() (*dicce.evaluation.evaluator.Evaluator* method), 48
 evaluate_lp() (*dicce.Evaluator* method), 257
 evaluate_lp() (*dicce.evaluator.Evaluator* method), 65
 evaluate_lp() (in module *dicce.evaluation*), 61
 evaluate_lp() (in module *dicce.evaluation.link_prediction*), 51
 evaluate_lp() (in module *dicce.static_funcs_training*), 227
 evaluate_lp_bpe_k_vs_all() (*dicce.evaluation.Evaluator* method), 58
 evaluate_lp_bpe_k_vs_all() (*dicce.evaluation.evaluator.Evaluator* method), 48
 evaluate_lp_bpe_k_vs_all() (*dicce.Evaluator* method), 257
 evaluate_lp_bpe_k_vs_all() (*dicce.evaluator.Evaluator* method), 65
 evaluate_lp_bpe_k_vs_all() (in module *dicce.eval_static_funcs*), 43
 evaluate_lp_bpe_k_vs_all() (in module *dicce.evaluation*), 61
 evaluate_lp_bpe_k_vs_all() (in module *dicce.evaluation.link_prediction*), 52
 evaluate_lp_k_vs_all() (*dicce.evaluation.Evaluator* method), 58
 evaluate_lp_k_vs_all() (*dicce.evaluation.evaluator.Evaluator* method), 48
 evaluate_lp_k_vs_all() (*dicce.Evaluator* method), 256
 evaluate_lp_k_vs_all() (*dicce.evaluator.Evaluator* method), 64
 evaluate_lp_with_byte() (*dicce.evaluation.Evaluator* method), 58
 evaluate_lp_with_byte() (*dicce.evaluation.evaluator.Evaluator* method), 48
 evaluate_lp_with_byte() (*dicce.Evaluator* method), 256
 evaluate_lp_with_byte() (*dicce.evaluator.Evaluator* method), 65
 Evaluator (class in *dicce*), 254
 Evaluator (class in *dicce.evaluation*), 56
 Evaluator (class in *dicce.evaluation.evaluator*), 46
 Evaluator (class in *dicce.evaluator*), 63
 evaluator (*dicce.DICE_Trainer* attribute), 254
 evaluator (*dicce.Execute* attribute), 245
 evaluator (*dicce.executer.Execute* attribute), 67
 evaluator (*dicce.trainer.DICE_Trainer* attribute), 238
 evaluator (*dicce.trainer.dice_trainer.DICE_Trainer* attribute), 231
 every_x_epoch (*dicce.callbacks.KGESaveCallback* attribute), 24
 example_input_array (*dicce.models.ensemble.EnsembleKGE* property), 110
 Execute (class in *dicce*), 245
 Execute (class in *dicce.executer*), 66
 exists() (*dicce.knowledge_graph.KG* method), 71
 Experiment (class in *dicce.analyse_experiments*), 21
 experiment_dir (*dicce.callbacks.LRScheduler* attribute), 27
 experiment_dir (*dicce.callbacks.PeriodicEvalCallback* attribute), 27
 explicit (*dicce.models.QMult* attribute), 180
 explicit (*dicce.models.quaternion.QMult* attribute), 123
 exponential_function() (in module *dicce.static_funcs*), 226
 extract_input_outputs() (*dicce.trainer.torch_trainer_ddp.NodeTrainer* method), 236
 extract_input_outputs() (*dicce.trainer.torch_trainer_fsdp.TorchFSDPTrainer* method), 237
 extract_input_outputs() (in module *dicce.trainer.model_parallelism*), 233
 extract_input_outputs_set_device() (*dicce.trainer.torch_trainer.TorchTrainer* method), 234

F

f (*dicce.callbacks.KronE* attribute), 26

`fc` (*dicee.models.literal.LiteralEmbeddings attribute*), 118
`fc1` (*dicee.models.AConEx attribute*), 171
`fc1` (*dicee.models.AConvO attribute*), 188
`fc1` (*dicee.models.AConvQ attribute*), 181
`fc1` (*dicee.models.complex.AConEx attribute*), 105
`fc1` (*dicee.models.complex.ConEx attribute*), 104
`fc1` (*dicee.models.ConEx attribute*), 170
`fc1` (*dicee.models.ConvO attribute*), 188
`fc1` (*dicee.models.ConvQ attribute*), 181
`fc1` (*dicee.models.octonion.AConvO attribute*), 120
`fc1` (*dicee.models.octonion.ConvO attribute*), 120
`fc1` (*dicee.models.quaternion.AConvQ attribute*), 125
`fc1` (*dicee.models.quaternion.ConvQ attribute*), 124
`fc_num_input` (*dicee.models.AConEx attribute*), 171
`fc_num_input` (*dicee.models.AConvO attribute*), 188
`fc_num_input` (*dicee.models.AConvQ attribute*), 181
`fc_num_input` (*dicee.models.complex.AConEx attribute*), 105
`fc_num_input` (*dicee.models.complex.ConEx attribute*), 104
`fc_num_input` (*dicee.models.ConEx attribute*), 170
`fc_num_input` (*dicee.models.ConvO attribute*), 188
`fc_num_input` (*dicee.models.ConvQ attribute*), 181
`fc_num_input` (*dicee.models.octonion.AConvO attribute*), 120
`fc_num_input` (*dicee.models.octonion.ConvO attribute*), 120
`fc_num_input` (*dicee.models.quaternion.AConvQ attribute*), 125
`fc_num_input` (*dicee.models.quaternion.ConvQ attribute*), 124
`fc_out` (*dicee.models.literal.LiteralEmbeddings attribute*), 118
`feature_map_dropout` (*dicee.models.AConEx attribute*), 171
`feature_map_dropout` (*dicee.models.AConvO attribute*), 188
`feature_map_dropout` (*dicee.models.AConvQ attribute*), 181
`feature_map_dropout` (*dicee.models.complex.AConEx attribute*), 105
`feature_map_dropout` (*dicee.models.complex.ConEx attribute*), 104
`feature_map_dropout` (*dicee.models.ConEx attribute*), 170
`feature_map_dropout` (*dicee.models.ConvO attribute*), 188
`feature_map_dropout` (*dicee.models.ConvQ attribute*), 181
`feature_map_dropout` (*dicee.models.octonion.AConvO attribute*), 120
`feature_map_dropout` (*dicee.models.octonion.ConvO attribute*), 120
`feature_map_dropout` (*dicee.models.quaternion.AConvQ attribute*), 125
`feature_map_dropout` (*dicee.models.quaternion.ConvQ attribute*), 124
`feature_map_dropout_rate` (*dicee.config.Namespace attribute*), 31
`feature_map_dropout_rate` (*dicee.models.base_model.BaseKGE attribute*), 91
`feature_map_dropout_rate` (*dicee.models.BaseKGE attribute*), 149, 154, 166, 175, 182, 197, 203
`fetch_worker()` (in module *dicee.read_preprocess_save_load_kg.util*), 217
`fill_query()` (*dicee.query_generator.QueryGenerator method*), 212
`fill_query()` (*dicee.QueryGenerator method*), 253
`find_good_batch_size()` (in module *dicee.trainer.auto_batch_finder*), 229
`find_missing_triples()` (*dicee.KGE method*), 251
`find_missing_triples()` (*dicee.knowledge_graph_embeddings.KGE method*), 76
`fit()` (*dicee.trainer.model_parallelism.TensorParallel method*), 233
`fit()` (*dicee.trainer.torch_trainer_ddp.TorchDDPTrainer method*), 235
`fit()` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer method*), 237
`fit()` (*dicee.trainer.torch_trainer.TorchTrainer method*), 234
`FixedNegSampleDataset` (class in *dicee.dataset_classes*), 39
`flash` (*dicee.models.transformers.SelfAttention attribute*), 136
`FMult` (class in *dicee.models*), 206
`FMult` (class in *dicee.models.function_space*), 113
`FMult2` (class in *dicee.models*), 208
`FMult2` (class in *dicee.models.function_space*), 114
`form_of_labelling` (*dicee.DICE_Trainer attribute*), 254
`form_of_labelling` (*dicee.trainer.DICE_Trainer attribute*), 238
`form_of_labelling` (*dicee.trainer.dice_trainer.DICE_Trainer attribute*), 231
`forward()` (*dicee.models.base_model.BaseKGE method*), 92
`forward()` (*dicee.models.base_model.IdentityClass static method*), 95
`forward()` (*dicee.models.BaseKGE method*), 150, 156, 168, 177, 184, 199, 204
`forward()` (*dicee.models.Block method*), 154
`forward()` (*dicee.models.clifford.TransformerBlock method*), 99
`forward()` (*dicee.models.clifford.TransformerMLP method*), 100
`forward()` (*dicee.models.clifford.TransformerSelfAttention method*), 100
`forward()` (*dicee.models.fsdp_models.RowWiseShardedEmbedding method*), 111

`forward()` (*dicee.models.IdentityClass* static method), 153, 179, 187
`forward()` (*dicee.models.literal.LiteralEmbeddings* method), 118
`forward()` (*dicee.models.transformers.Block* method), 138
`forward()` (*dicee.models.transformers.BytE* method), 135
`forward()` (*dicee.models.transformers.GPT* method), 139
`forward()` (*dicee.models.transformers.LayerNorm* method), 136
`forward()` (*dicee.models.transformers.MLP* method), 137
`forward()` (*dicee.models.transformers.SelfAttention* method), 137
`forward_backward_update()` (*dicee.trainer.torch_trainer.TorchTrainer* method), 234
`forward_backward_update_loss()` (in module *dicee.trainer.model_parallelism*), 233
`forward_byte_pair_encoded_k_vs_all()` (*dicee.models.base_model.BaseKGE* method), 91
`forward_byte_pair_encoded_k_vs_all()` (*dicee.models.BaseKGE* method), 150, 155, 167, 176, 183, 198, 203
`forward_byte_pair_encoded_triple()` (*dicee.models.base_model.BaseKGE* method), 92
`forward_byte_pair_encoded_triple()` (*dicee.models.BaseKGE* method), 150, 155, 167, 176, 184, 198, 204
`forward_k_vs_all()` (*dicee.models.AConEx* method), 171
`forward_k_vs_all()` (*dicee.models.AConvO* method), 189
`forward_k_vs_all()` (*dicee.models.AConvQ* method), 182
`forward_k_vs_all()` (*dicee.models.base_model.BaseKGE* method), 93
`forward_k_vs_all()` (*dicee.models.BaseKGE* method), 151, 156, 169, 177, 185, 199, 205
`forward_k_vs_all()` (*dicee.models.clifford.DeCaL* method), 101
`forward_k_vs_all()` (*dicee.models.clifford.Keci* method), 98
`forward_k_vs_all()` (*dicee.models.clifford.KeciTransformer* method), 99
`forward_k_vs_all()` (*dicee.models.CoKE* method), 165
`forward_k_vs_all()` (*dicee.models.ComplEx* method), 173
`forward_k_vs_all()` (*dicee.models.complex.AConEx* method), 105
`forward_k_vs_all()` (*dicee.models.complex.ComplEx* method), 106
`forward_k_vs_all()` (*dicee.models.complex.ConEx* method), 104
`forward_k_vs_all()` (*dicee.models.complex.RotatE* method), 107
`forward_k_vs_all()` (*dicee.models.ConEx* method), 170
`forward_k_vs_all()` (*dicee.models.ConvO* method), 188
`forward_k_vs_all()` (*dicee.models.ConvQ* method), 181
`forward_k_vs_all()` (*dicee.models.DeCaL* method), 194
`forward_k_vs_all()` (*dicee.models.DistMult* method), 159
`forward_k_vs_all()` (*dicee.models.DualE* method), 211
`forward_k_vs_all()` (*dicee.models.dualE.DualE* method), 109
`forward_k_vs_all()` (*dicee.models.fsdp_models.FSDPShardedEntityModel* method), 112
`forward_k_vs_all()` (*dicee.models.Keci* method), 191
`forward_k_vs_all()` (*dicee.models.KeciTransformer* method), 196
`forward_k_vs_all()` (*dicee.models.MuRE* method), 162
`forward_k_vs_all()` (*dicee.models.octonion.AConvO* method), 121
`forward_k_vs_all()` (*dicee.models.octonion.ConvO* method), 120
`forward_k_vs_all()` (*dicee.models.octonion.OMult* method), 119
`forward_k_vs_all()` (*dicee.models.OMult* method), 187
`forward_k_vs_all()` (*dicee.models.pykeen_models.PykeenKGE* method), 121
`forward_k_vs_all()` (*dicee.models.PykeenKGE* method), 201
`forward_k_vs_all()` (*dicee.models.QMult* method), 180
`forward_k_vs_all()` (*dicee.models.quaternion.AConvQ* method), 125
`forward_k_vs_all()` (*dicee.models.quaternion.ConvQ* method), 124
`forward_k_vs_all()` (*dicee.models.quaternion.QMult* method), 124
`forward_k_vs_all()` (*dicee.models.real.CoKE* method), 133
`forward_k_vs_all()` (*dicee.models.real.DistMult* method), 126
`forward_k_vs_all()` (*dicee.models.real.MuRE* method), 130
`forward_k_vs_all()` (*dicee.models.real.Shallom* method), 131
`forward_k_vs_all()` (*dicee.models.real.TransE* method), 128
`forward_k_vs_all()` (*dicee.models.real.TransH* method), 129
`forward_k_vs_all()` (*dicee.models.RotatE* method), 174
`forward_k_vs_all()` (*dicee.models.Shallom* method), 163
`forward_k_vs_all()` (*dicee.models.TransE* method), 160
`forward_k_vs_all()` (*dicee.models.TransH* method), 161
`forward_k_vs_sample()` (*dicee.models.AConEx* method), 172
`forward_k_vs_sample()` (*dicee.models.base_model.BaseKGE* method), 93
`forward_k_vs_sample()` (*dicee.models.BaseKGE* method), 151, 157, 169, 177, 185, 199, 205
`forward_k_vs_sample()` (*dicee.models.clifford.Keci* method), 98
`forward_k_vs_sample()` (*dicee.models.CoKE* method), 165
`forward_k_vs_sample()` (*dicee.models.ComplEx* method), 173
`forward_k_vs_sample()` (*dicee.models.complex.AConEx* method), 105
`forward_k_vs_sample()` (*dicee.models.complex.ComplEx* method), 106
`forward_k_vs_sample()` (*dicee.models.complex.ConEx* method), 104

`forward_k_vs_sample()` (*dicee.models.complex.RotatE method*), 108
`forward_k_vs_sample()` (*dicee.models.ConEx method*), 171
`forward_k_vs_sample()` (*dicee.models.DistMult method*), 159
`forward_k_vs_sample()` (*dicee.models.Keci method*), 192
`forward_k_vs_sample()` (*dicee.models.MuRE method*), 162
`forward_k_vs_sample()` (*dicee.models.pykeen_models.PykeenKGE method*), 122
`forward_k_vs_sample()` (*dicee.models.PykeenKGE method*), 202
`forward_k_vs_sample()` (*dicee.models.QMult method*), 180
`forward_k_vs_sample()` (*dicee.models.quaternion.QMult method*), 124
`forward_k_vs_sample()` (*dicee.models.real.CoKE method*), 133
`forward_k_vs_sample()` (*dicee.models.real.DistMult method*), 126
`forward_k_vs_sample()` (*dicee.models.real.MuRE method*), 130
`forward_k_vs_sample()` (*dicee.models.real.TransH method*), 129
`forward_k_vs_sample()` (*dicee.models.RotatE method*), 174
`forward_k_vs_sample()` (*dicee.models.TransH method*), 161
`forward_k_vs_with_explicit()` (*dicee.models.clifford.Keci method*), 97
`forward_k_vs_with_explicit()` (*dicee.models.Keci method*), 191
`forward_triples()` (*dicee.models.AConEx method*), 172
`forward_triples()` (*dicee.models.AConvO method*), 188
`forward_triples()` (*dicee.models.AConvQ method*), 182
`forward_triples()` (*dicee.models.base_model.BaseKGE method*), 93
`forward_triples()` (*dicee.models.BaseKGE method*), 151, 156, 168, 177, 185, 199, 205
`forward_triples()` (*dicee.models.clifford.DeCaL method*), 101
`forward_triples()` (*dicee.models.clifford.Keci method*), 99
`forward_triples()` (*dicee.models.complex.AConEx method*), 105
`forward_triples()` (*dicee.models.complex.ConEx method*), 104
`forward_triples()` (*dicee.models.ConEx method*), 171
`forward_triples()` (*dicee.models.ConvO method*), 188
`forward_triples()` (*dicee.models.ConvQ method*), 181
`forward_triples()` (*dicee.models.DeCaL method*), 193
`forward_triples()` (*dicee.models.DualE method*), 211
`forward_triples()` (*dicee.models.dualE.DualE method*), 109
`forward_triples()` (*dicee.models.FMult method*), 207
`forward_triples()` (*dicee.models.FMult2 method*), 208
`forward_triples()` (*dicee.models.function_space.FMult method*), 113
`forward_triples()` (*dicee.models.function_space.FMult2 method*), 114
`forward_triples()` (*dicee.models.function_space.GFMult method*), 114
`forward_triples()` (*dicee.models.function_space.LFMult method*), 115
`forward_triples()` (*dicee.models.function_space.LFMult1 method*), 115
`forward_triples()` (*dicee.models.GFMult method*), 207
`forward_triples()` (*dicee.models.Keci method*), 192
`forward_triples()` (*dicee.models.LFMult method*), 209
`forward_triples()` (*dicee.models.LFMult1 method*), 209
`forward_triples()` (*dicee.models.MuRE method*), 162
`forward_triples()` (*dicee.models.octonion.AConvO method*), 120
`forward_triples()` (*dicee.models.octonion.ConvO method*), 120
`forward_triples()` (*dicee.models.Pyke method*), 164
`forward_triples()` (*dicee.models.pykeen_models.PykeenKGE method*), 122
`forward_triples()` (*dicee.models.PykeenKGE method*), 202
`forward_triples()` (*dicee.models.quaternion.AConvQ method*), 125
`forward_triples()` (*dicee.models.quaternion.ConvQ method*), 124
`forward_triples()` (*dicee.models.real.MuRE method*), 130
`forward_triples()` (*dicee.models.real.Pyke method*), 131
`forward_triples()` (*dicee.models.real.Shallom method*), 131
`forward_triples()` (*dicee.models.real.TransH method*), 128
`forward_triples()` (*dicee.models.Shallom method*), 163
`forward_triples()` (*dicee.models.TransH method*), 161
`freeze_entity_embeddings` (*dicee.models.literal.LiteralEmbeddings attribute*), 117, 118
`frequency` (*dicee.callbacks.Perturb attribute*), 27
`from_pretrained()` (*dicee.models.transformers.GPT class method*), 139
`from_pretrained_model_write_embeddings_into_csv()` (*in module dicee.static_funcs*), 226
`FSDPlvsSampleDataset` (*class in dicee.dataset_classes*), 35
`FSDPShardedEntityModel` (*class in dicee.models.fsdp_models*), 111
`full_storage_path` (*dicee.analyse_experiments.Experiment attribute*), 21
`func_triple_to_bpe_representation` (*dicee.evaluation.Evaluator attribute*), 57
`func_triple_to_bpe_representation` (*dicee.evaluation.evaluator.Evaluator attribute*), 47
`func_triple_to_bpe_representation` (*dicee.Evaluator attribute*), 255
`func_triple_to_bpe_representation` (*dicee.evaluator.Evaluator attribute*), 63

func_triple_to_bpe_representation() (*dicее.knowledge_graph.KG method*), 71
function() (*dicее.models.FMult2 method*), 208
function() (*dicее.models.function_space.FMult2 method*), 114

G

gamma (*dicее.models.FMult attribute*), 207
gamma (*dicее.models.function_space.FMult attribute*), 113
gate_residual (*dicее.models.literal.LiteralEmbeddings attribute*), 117, 118
gated_residual_proj (*dicее.models.literal.LiteralEmbeddings attribute*), 118
gather_entity_embeddings_on_rank_zero() (*dicее.models.fsdp_models.FSDPShardedEntityModel method*), 112
gelu (*dicее.models.clifford.TransformerMLP attribute*), 100
gelu (*dicее.models.transformers.MLP attribute*), 137
gen_test (*dicее.query_generator.QueryGenerator attribute*), 212
gen_test (*dicее.QueryGenerator attribute*), 252
gen_valid (*dicее.query_generator.QueryGenerator attribute*), 212
gen_valid (*dicее.QueryGenerator attribute*), 252
generate() (*dicее.KGE method*), 247
generate() (*dicее.knowledge_graph_embeddings.KGE method*), 72
generate() (*dicее.models.transformers.BytE method*), 135
generate_queries() (*dicее.query_generator.QueryGenerator method*), 212
generate_queries() (*dicее.QueryGenerator method*), 253
get_aswa_state_dict() (*dicее.weight_averaging.ASWA method*), 240
get_bpe_head_and_relation_representation() (*dicее.models.base_model.BaseKGE method*), 94
get_bpe_head_and_relation_representation() (*dicее.models.BaseKGE method*), 152, 158, 170, 178, 186, 200, 206
get_bpe_token_representation() (*dicее.abstracts.BaseInteractiveKGE method*), 14
get_callbacks() (*in module dicее.trainer.dice_trainer*), 231
get_default_arguments() (*in module dicее.analyse_experiments*), 21
get_default_arguments() (*in module dicее.scripts.index_serve*), 221
get_default_arguments() (*in module dicее.scripts.run*), 222
get_ee_vocab() (*in module dicее.read_preprocess_save_load_kg.util*), 218
get_ee_vocab() (*in module dicее.static_funcs*), 224
get_ee_vocab() (*in module dicее.static_preprocess_funcs*), 229
get_embeddings() (*dicее.models.base_model.BaseKGE method*), 94
get_embeddings() (*dicее.models.BaseKGE method*), 152, 158, 170, 179, 186, 201, 206
get_embeddings() (*dicее.models.ensemble.EnsembleKGE method*), 110
get_embeddings() (*dicее.models.fsdp_models.FSDPShardedEntityModel method*), 112
get_embeddings() (*dicее.models.real.Shallom method*), 131
get_embeddings() (*dicее.models.Shallom method*), 163
get_entity_embeddings() (*dicее.abstracts.BaseInteractiveKGE method*), 16
get_entity_index() (*dicее.abstracts.BaseInteractiveKGE method*), 15
get_er_vocab() (*in module dicее.read_preprocess_save_load_kg.util*), 218
get_er_vocab() (*in module dicее.static_funcs*), 223
get_er_vocab() (*in module dicее.static_preprocess_funcs*), 229
get_eval_report() (*dicее.abstracts.BaseInteractiveKGE method*), 14
get_head_relation_representation() (*dicее.models.base_model.BaseKGE method*), 93
get_head_relation_representation() (*dicее.models.BaseKGE method*), 152, 157, 169, 178, 185, 200, 205
get_kronecker_triple_representation() (*dicее.callbacks.KronE method*), 26
get_mean_and_var() (*dicее.weight_averaging.SWAG method*), 242
get_num_params() (*dicее.models.transformers.GPT method*), 139
get_padded_bpe_triple_representation() (*dicее.abstracts.BaseInteractiveKGE method*), 14
get_queries() (*dicее.query_generator.QueryGenerator method*), 212
get_queries() (*dicее.QueryGenerator method*), 253
get_re_vocab() (*in module dicее.read_preprocess_save_load_kg.util*), 218
get_re_vocab() (*in module dicее.static_funcs*), 223
get_re_vocab() (*in module dicее.static_preprocess_funcs*), 229
get_relation_embeddings() (*dicее.abstracts.BaseInteractiveKGE method*), 16
get_relation_index() (*dicее.abstracts.BaseInteractiveKGE method*), 15
get_sentence_representation() (*dicее.models.base_model.BaseKGE method*), 94
get_sentence_representation() (*dicее.models.BaseKGE method*), 152, 157, 169, 178, 186, 200, 206
get_transductive_entity_embeddings() (*dicее.KGE method*), 247
get_transductive_entity_embeddings() (*dicее.knowledge_graph_embeddings.KGE method*), 72
get_triple_representation() (*dicее.models.base_model.BaseKGE method*), 93
get_triple_representation() (*dicее.models.BaseKGE method*), 151, 157, 169, 178, 185, 200, 205
GFMult (*class in dicее.models*), 207
GFMult (*class in dicее.models.function_space*), 113
global_rank (*dicее.abstracts.AbstractTrainer attribute*), 13
global_rank (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 235

global_rank (*diccee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
 GPT (*class in diccee.models.transformers*), 138
 GPTConfig (*class in diccee.models.transformers*), 138
 gpus (*diccee.config.Namespace attribute*), 29
 gradient_accumulation_steps (*diccee.config.Namespace attribute*), 30
 gradient_clip_val (*diccee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
 ground_queries () (*diccee.query_generator.QueryGenerator method*), 212
 ground_queries () (*diccee.QueryGenerator method*), 253
 gswa_n (*diccee.weight_averaging.SWAG attribute*), 242

H

half_dim (*diccee.models.complex.RotatE attribute*), 107
 half_dim (*diccee.models.RotatE attribute*), 173
 hidden_dim (*diccee.models.literal.LiteralEmbeddings attribute*), 118
 hidden_dropout (*diccee.models.base_model.BaseKGE attribute*), 91
 hidden_dropout (*diccee.models.BaseKGE attribute*), 149, 155, 167, 175, 183, 197, 203
 hidden_dropout_rate (*diccee.config.Namespace attribute*), 31
 hidden_dropout_rate (*diccee.models.base_model.BaseKGE attribute*), 90
 hidden_dropout_rate (*diccee.models.BaseKGE attribute*), 149, 154, 166, 175, 182, 197, 202
 hidden_normalizer (*diccee.models.base_model.BaseKGE attribute*), 91
 hidden_normalizer (*diccee.models.BaseKGE attribute*), 149, 155, 167, 175, 183, 197, 203

I

IdentityClass (*class in diccee.models*), 153, 179, 186
 IdentityClass (*class in diccee.models.base_model*), 94
 idx_entity_to_bpe_shaped (*diccee.knowledge_graph.KG attribute*), 70
 index () (*in module diccee.scripts.index_serve*), 221
 index_triple () (*diccee.abstracts.BaseInteractiveKGE method*), 16
 init_data_loader () (*diccee.DICE_Trainer method*), 254
 init_data_loader () (*diccee.trainer.DICE_Trainer method*), 238
 init_data_loader () (*diccee.trainer.dice_trainer.DICE_Trainer method*), 232
 init_dataset () (*diccee.DICE_Trainer method*), 254
 init_dataset () (*diccee.trainer.DICE_Trainer method*), 238
 init_dataset () (*diccee.trainer.dice_trainer.DICE_Trainer method*), 232
 init_entity_embeddings () (*diccee.models.base_model.BaseKGE method*), 91
 init_entity_embeddings () (*diccee.models.BaseKGE method*), 149, 155, 167, 176, 183, 198, 203
 init_param (*diccee.config.Namespace attribute*), 30
 init_params_with_sanity_checking () (*diccee.models.base_model.BaseKGE method*), 92
 init_params_with_sanity_checking () (*diccee.models.BaseKGE method*), 150, 156, 168, 176, 184, 198, 204
 init_relation_embeddings () (*diccee.models.base_model.BaseKGE method*), 91
 init_relation_embeddings () (*diccee.models.BaseKGE method*), 150, 155, 167, 176, 183, 198, 203
 initial_eval_setting (*diccee.weight_averaging.ASWA attribute*), 239
 initialize_or_load_model () (*diccee.DICE_Trainer method*), 254
 initialize_or_load_model () (*diccee.trainer.DICE_Trainer method*), 238
 initialize_or_load_model () (*diccee.trainer.dice_trainer.DICE_Trainer method*), 232
 initialize_trainer () (*diccee.DICE_Trainer method*), 254
 initialize_trainer () (*diccee.trainer.DICE_Trainer method*), 238
 initialize_trainer () (*diccee.trainer.dice_trainer.DICE_Trainer method*), 232
 initialize_trainer () (*in module diccee.trainer.dice_trainer*), 230
 input_dp_ent_real (*diccee.models.base_model.BaseKGE attribute*), 91
 input_dp_ent_real (*diccee.models.BaseKGE attribute*), 149, 155, 167, 175, 183, 197, 203
 input_dp_rel_real (*diccee.models.base_model.BaseKGE attribute*), 91
 input_dp_rel_real (*diccee.models.BaseKGE attribute*), 149, 155, 167, 175, 183, 197, 203
 input_dropout_rate (*diccee.config.Namespace attribute*), 31
 input_dropout_rate (*diccee.models.base_model.BaseKGE attribute*), 90
 input_dropout_rate (*diccee.models.BaseKGE attribute*), 149, 154, 166, 175, 182, 197, 202
 InteractiveQueryDecomposition (*class in diccee.abstracts*), 17
 initialize_model () (*in module diccee.static_funcs*), 225
 is_continual_training (*diccee.DICE_Trainer attribute*), 253
 is_continual_training (*diccee.evaluation.Evaluator attribute*), 57
 is_continual_training (*diccee.evaluation.evaluator.Evaluator attribute*), 47
 is_continual_training (*diccee.Evaluator attribute*), 255
 is_continual_training (*diccee.evaluator.Evaluator attribute*), 63
 is_continual_training (*diccee.Execute attribute*), 245
 is_continual_training (*diccee.executer.Execute attribute*), 67
 is_continual_training (*diccee.trainer.DICE_Trainer attribute*), 238
 is_continual_training (*diccee.trainer.dice_trainer.DICE_Trainer attribute*), 231

`is_global_rank_zero()` (*dicee.Execute method*), 246
`is_global_rank_zero()` (*dicee.executer.Execute method*), 67
`is_global_zero` (*dicee.abstracts.AbstractTrainer attribute*), 13
`is_global_zero` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
`is_local_rank_zero()` (*dicee.Execute method*), 246
`is_local_rank_zero()` (*dicee.executer.Execute method*), 67
`is_seen()` (*dicee.abstracts.BaseInteractiveKGE method*), 15
`is_sparql_endpoint_alive()` (*in module dicee.sanity_checkers*), 220

K

`k` (*dicee.models.FMult attribute*), 207
`k` (*dicee.models.FMult2 attribute*), 208
`k` (*dicee.models.function_space.FMult attribute*), 113
`k` (*dicee.models.function_space.FMult2 attribute*), 114
`k` (*dicee.models.function_space.GFMult attribute*), 114
`k` (*dicee.models.GFMult attribute*), 207
`k_fold_cross_validation()` (*dicee.DICE_Trainer method*), 254
`k_fold_cross_validation()` (*dicee.trainer.DICE_Trainer method*), 239
`k_fold_cross_validation()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 232
`k_vs_all_score()` (*dicee.models.clifford.Keci method*), 97
`k_vs_all_score()` (*dicee.models.ComplEx static method*), 172
`k_vs_all_score()` (*dicee.models.complex.ComplEx static method*), 106
`k_vs_all_score()` (*dicee.models.DistMult method*), 158
`k_vs_all_score()` (*dicee.models.Keci method*), 191
`k_vs_all_score()` (*dicee.models.octonion.OMult method*), 119
`k_vs_all_score()` (*dicee.models.OMult method*), 187
`k_vs_all_score()` (*dicee.models.QMult method*), 180
`k_vs_all_score()` (*dicee.models.quaternion.QMult method*), 123
`k_vs_all_score()` (*dicee.models.real.DistMult method*), 126
`Keci` (*class in dicee.models*), 189
`Keci` (*class in dicee.models.clifford*), 95
`KeciTransformer` (*class in dicee.models*), 196
`KeciTransformer` (*class in dicee.models.clifford*), 99
`kernel_size` (*dicee.config.Namespace attribute*), 30
`kernel_size` (*dicee.models.base_model.BaseKGE attribute*), 91
`kernel_size` (*dicee.models.BaseKGE attribute*), 149, 154, 166, 175, 182, 197, 203
`KG` (*class in dicee.knowledge_graph*), 69
`kg` (*dicee.callbacks.PseudoLabellingCallback attribute*), 25
`kg` (*dicee.read_preprocess_save_load_kg.LoadSaveToDisk attribute*), 219
`kg` (*dicee.read_preprocess_save_load_kg.PreprocessKG attribute*), 219
`kg` (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG attribute*), 213
`kg` (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk attribute*), 214
`kg` (*dicee.read_preprocess_save_load_kg.ReadFromDisk attribute*), 219
`kg` (*dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk attribute*), 215
`KGE` (*class in dicee*), 247
`KGE` (*class in dicee.knowledge_graph_embeddings*), 72
`KGESaveCallback` (*class in dicee.callbacks*), 24
`knowledge_graph` (*dicee.Execute attribute*), 245
`knowledge_graph` (*dicee.executer.Execute attribute*), 67
`KronE` (*class in dicee.callbacks*), 26
`KvsAll` (*class in dicee.dataset_classes*), 36
`kvsall_score()` (*dicee.models.DualE method*), 210
`kvsall_score()` (*dicee.models.dualE.DualE method*), 108
`KvsSampleDataset` (*class in dicee.dataset_classes*), 36

L

`label_smoothing_rate` (*dicee.config.Namespace attribute*), 30
`label_smoothing_rate` (*dicee.dataset_classes.AllvsAll attribute*), 35
`label_smoothing_rate` (*dicee.dataset_classes.FixedNegSampleDataset attribute*), 39
`label_smoothing_rate` (*dicee.dataset_classes.FSDP1vsSampleDataset attribute*), 36
`label_smoothing_rate` (*dicee.dataset_classes.KvsAll attribute*), 36
`label_smoothing_rate` (*dicee.dataset_classes.KvsSampleDataset attribute*), 37
`label_smoothing_rate` (*dicee.dataset_classes.OnevsSample attribute*), 40
`label_smoothing_rate` (*dicee.dataset_classes.TriplePredictionDataset attribute*), 40
`layer_norm` (*dicee.models.literal.LiteralEmbeddings attribute*), 118
`LayerNorm` (*class in dicee.models.transformers*), 135
`learning_rate` (*dicee.models.base_model.BaseKGE attribute*), 90

`learning_rate` (*dicee.models.BaseKGE* attribute), 149, 154, 166, 175, 182, 197, 202
`length` (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 39
`length` (*dicee.dataset_classes.TriplePredictionDataset* attribute), 40
`level` (*dicee.callbacks.Perturb* attribute), 26
`LFMult` (class in *dicee.models*), 209
`LFMult` (class in *dicee.models.function_space*), 115
`LFMult1` (class in *dicee.models*), 208
`LFMult1` (class in *dicee.models.function_space*), 115
`linear()` (*dicee.models.function_space.LFMult* method), 116
`linear()` (*dicee.models.LFMult* method), 209
`list2tuple()` (*dicee.query_generator.QueryGenerator* method), 212
`list2tuple()` (*dicee.QueryGenerator* method), 252
`LiteralDataset` (class in *dicee.dataset_classes*), 38
`LiteralEmbeddings` (class in *dicee.models.literal*), 117
`lm_head` (*dicee.models.clifford.KeciTransformer* attribute), 99
`lm_head` (*dicee.models.KeciTransformer* attribute), 196
`lm_head` (*dicee.models.transformers.BytE* attribute), 135
`lm_head` (*dicee.models.transformers.GPT* attribute), 139
`ln_1` (*dicee.models.Block* attribute), 153
`ln_1` (*dicee.models.clifford.TransformerBlock* attribute), 99
`ln_1` (*dicee.models.transformers.Block* attribute), 138
`ln_2` (*dicee.models.Block* attribute), 153
`ln_2` (*dicee.models.clifford.TransformerBlock* attribute), 99
`ln_2` (*dicee.models.transformers.Block* attribute), 138
`ln_f` (*dicee.models.CoKE* attribute), 165
`ln_f` (*dicee.models.real.CoKE* attribute), 133
`load()` (*dicee.read_preprocess_save_load_kg.LoadSaveToDisk* method), 219
`load()` (*dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk* method), 215
`load_and_validate_literal_data()` (*dicee.dataset_classes.LiteralDataset* static method), 38
`load_from_mmap()` (*dicee.Execute* method), 246
`load_from_mmap()` (*dicee.executer.Execute* method), 67
`load_json()` (in module *dicee.static_funcs*), 225
`load_local_shard_checkpoint()` (*dicee.models.fsdp_models.FSDPShardedEntityModel* method), 112
`load_model()` (in module *dicee.static_funcs*), 225
`load_model_ensemble()` (in module *dicee.static_funcs*), 225
`load_numpy()` (in module *dicee.static_funcs*), 226
`load_numpy_ndarray()` (in module *dicee.read_preprocess_save_load_kg.util*), 218
`load_pickle()` (in module *dicee.read_preprocess_save_load_kg.util*), 218
`load_pickle()` (in module *dicee.static_funcs*), 224
`load_queries()` (*dicee.query_generator.QueryGenerator* method), 212
`load_queries()` (*dicee.QueryGenerator* method), 253
`load_queries_and_answers()` (*dicee.query_generator.QueryGenerator* static method), 213
`load_queries_and_answers()` (*dicee.QueryGenerator* static method), 253
`load_state_dict()` (*dicee.models.ensemble.EnsembleKGE* method), 110
`load_term_mapping()` (in module *dicee.static_funcs*), 224
`load_term_mapping()` (in module *dicee.trainer.dice_trainer*), 230
`load_with_pandas()` (in module *dicee.read_preprocess_save_load_kg.util*), 218
`loader_backend` (*dicee.dataset_classes.LiteralDataset* attribute), 38
`LoadSaveToDisk` (class in *dicee.read_preprocess_save_load_kg*), 219
`LoadSaveToDisk` (class in *dicee.read_preprocess_save_load_kg.save_load_disk*), 215
`local_rank` (*dicee.abstracts.AbstractTrainer* attribute), 13
`local_rank` (*dicee.Execute* attribute), 245
`local_rank` (*dicee.executer.Execute* attribute), 67
`local_rank` (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 235
`local_rank` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237
`local_rows` (*dicee.models.fsdp_models.RowWiseShardedEmbedding* attribute), 111
`log_level` (*dicee.config.Namespace* attribute), 30
`logger` (in module *dicee.abstracts*), 13
`logger` (in module *dicee.analyse_experiments*), 21
`logger` (in module *dicee.callbacks*), 23
`logger` (in module *dicee.evaluation.evaluator*), 46
`logger` (in module *dicee.evaluation.link_prediction*), 50
`logger` (in module *dicee.evaluation.literal_prediction*), 52
`logger` (in module *dicee.executer*), 66
`logger` (in module *dicee.knowledge_graph_embeddings*), 72
`logger` (in module *dicee.models*), 144, 201, 206
`logger` (in module *dicee.models.base_model*), 86
`logger` (in module *dicee.models.clifford*), 95

- logger (in module *dicee.models.function_space*), 113
- logger (in module *dicee.models.pykeen_models*), 121
- logger (in module *dicee.models.transformers*), 134
- logger (in module *dicee.query_generator*), 211
- logger (in module *dicee.read_preprocess_save_load_kg.preprocess*), 213
- logger (in module *dicee.read_preprocess_save_load_kg.read_from_disk*), 214
- logger (in module *dicee.read_preprocess_save_load_kg.save_load_disk*), 215
- logger (in module *dicee.read_preprocess_save_load_kg.util*), 216
- logger (in module *dicee.sanity_checkers*), 220
- logger (in module *dicee.scripts.index_serve*), 221
- logger (in module *dicee.static_funcs*), 223
- logger (in module *dicee.static_preprocess_funcs*), 229
- logger (in module *dicee.trainer.auto_batch_finder*), 229
- logger (in module *dicee.trainer.dice_trainer*), 230
- logger (in module *dicee.trainer.model_parallelism*), 233
- logger (in module *dicee.trainer.torch_trainer*), 234
- logger (in module *dicee.trainer.torch_trainer_ddp*), 235
- logger (in module *dicee.trainer.torch_trainer_fsdp*), 236
- loss (*dicee.models.base_model.BaseKGE* attribute), 91
- loss (*dicee.models.BaseKGE* attribute), 149, 154, 166, 175, 183, 197, 203
- loss_func (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 235
- loss_func (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237
- loss_function (*dicee.trainer.torch_trainer.TorchTrainer* attribute), 234
- loss_function() (*dicee.models.base_model.BaseKGE* Lightning method), 87
- loss_function() (*dicee.models.BaseKGE* Lightning method), 145
- loss_function() (*dicee.models.transformers.BytE* method), 135
- loss_history (*dicee.models.base_model.BaseKGE* attribute), 91
- loss_history (*dicee.models.BaseKGE* attribute), 149, 155, 167, 175, 183, 197, 203
- loss_history (*dicee.models.pykeen_models.PykeenKGE* attribute), 121
- loss_history (*dicee.models.PykeenKGE* attribute), 201
- loss_history (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 236
- lr (*dicee.analyse_experiments.Experiment* attribute), 21
- lr (*dicee.config.Namespace* attribute), 29
- lr_init (*dicee.weight_averaging.SWA* attribute), 241
- lr_init (*dicee.weight_averaging.SWAG* attribute), 241
- lr_init (*dicee.weight_averaging.TWA* attribute), 243
- lr_lambda (*dicee.callbacks.LRScheduler* attribute), 28
- LRScheduler (class in *dicee.callbacks*), 27

M

- m (*dicee.models.function_space.LFMult* attribute), 115
- m (*dicee.models.LFMult* attribute), 209
- main() (in module *dicee.scripts.index_serve*), 221
- main() (in module *dicee.scripts.run*), 222
- make_iterable_verbose() (in module *dicee.evaluation.utils*), 54
- make_iterable_verbose() (in module *dicee.static_funcs_training*), 228
- make_iterable_verbose() (in module *dicee.trainer.torch_trainer_ddp*), 235
- mapping_from_first_two_cols_to_third() (in module *dicee.static_preprocess_funcs*), 229
- margin (*dicee.models.complex.RotatE* attribute), 107
- margin (*dicee.models.Pyke* attribute), 164
- margin (*dicee.models.real.Pyke* attribute), 131
- margin (*dicee.models.real.TransE* attribute), 127
- margin (*dicee.models.RotatE* attribute), 173
- margin (*dicee.models.TransE* attribute), 160
- mask_emb (*dicee.models.CoKE* attribute), 165
- mask_emb (*dicee.models.real.CoKE* attribute), 133
- max_ans_num (*dicee.query_generator.QueryGenerator* attribute), 212
- max_ans_num (*dicee.QueryGenerator* attribute), 252
- max_epochs (*dicee.callbacks.KGESaveCallback* attribute), 24
- max_epochs (*dicee.callbacks.PeriodicEvalCallback* attribute), 27
- max_epochs (*dicee.weight_averaging.EMA* attribute), 242
- max_epochs (*dicee.weight_averaging.SWA* attribute), 241
- max_epochs (*dicee.weight_averaging.SWAG* attribute), 241
- max_epochs (*dicee.weight_averaging.TWA* attribute), 243
- max_length_subword_tokens (*dicee.knowledge_graph.KG* attribute), 71
- max_length_subword_tokens (*dicee.models.base_model.BaseKGE* attribute), 91
- max_length_subword_tokens (*dicee.models.BaseKGE* attribute), 149, 155, 167, 176, 183, 198, 203

- `max_num_models` (*dicee.weight_averaging.SWAG attribute*), 242
- `max_num_of_classes` (*dicee.dataset_classes.FSDP1vsSampleDataset attribute*), 36
- `max_num_of_classes` (*dicee.dataset_classes.KvsSampleDataset attribute*), 37
- `mean` (*dicee.weight_averaging.SWAG attribute*), 242
- `mem_of_model()` (*dicee.models.base_model.BaseKGELightning method*), 86
- `mem_of_model()` (*dicee.models.BaseKGELightning method*), 145
- `mem_of_model()` (*dicee.models.ensemble.EnsembleKGE method*), 110
- `method` (*dicee.callbacks.Perturb attribute*), 26
- `MLP` (*class in dicee.models.transformers*), 137
- `mlp` (*dicee.models.Block attribute*), 153
- `mlp` (*dicee.models.clifford.TransformerBlock attribute*), 99
- `mlp` (*dicee.models.transformers.Block attribute*), 138
- `mode` (*dicee.query_generator.QueryGenerator attribute*), 212
- `mode` (*dicee.QueryGenerator attribute*), 252
- `model` (*dicee.config.Namespace attribute*), 29
- `model` (*dicee.models.pykeen_models.PykeenKGE attribute*), 121
- `model` (*dicee.models.PykeenKGE attribute*), 201
- `model` (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 236
- `model` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
- `model` (*dicee.trainer.torch_trainer.TorchTrainer attribute*), 234
- `model_kwargs` (*dicee.models.pykeen_models.PykeenKGE attribute*), 121
- `model_kwargs` (*dicee.models.PykeenKGE attribute*), 201
- `model_name` (*dicee.analyse_experiments.Experiment attribute*), 21
- `MODEL_REGISTRY` (*in module dicee.static_funcs*), 223
- `module`
 - `dicee`, 12
 - `dicee.__main__`, 12
 - `dicee.abstracts`, 12
 - `dicee.analyse_experiments`, 20
 - `dicee.callbacks`, 22
 - `dicee.config`, 28
 - `dicee.dataset_classes`, 32
 - `dicee.eval_static_funcs`, 41
 - `dicee.evaluation`, 44
 - `dicee.evaluation.ensemble`, 45
 - `dicee.evaluation.evaluator`, 46
 - `dicee.evaluation.link_prediction`, 49
 - `dicee.evaluation.literal_prediction`, 52
 - `dicee.evaluation.utils`, 53
 - `dicee.evaluator`, 62
 - `dicee.executer`, 66
 - `dicee.knowledge_graph`, 69
 - `dicee.knowledge_graph_embeddings`, 71
 - `dicee.models`, 77
 - `dicee.models.adopt`, 77
 - `dicee.models.base_model`, 86
 - `dicee.models.clifford`, 95
 - `dicee.models.complex`, 103
 - `dicee.models.dualE`, 108
 - `dicee.models.ensemble`, 109
 - `dicee.models.fsdp_models`, 110
 - `dicee.models.function_space`, 112
 - `dicee.models.literal`, 117
 - `dicee.models.octonion`, 118
 - `dicee.models.pykeen_models`, 121
 - `dicee.models.quaternion`, 122
 - `dicee.models.real`, 125
 - `dicee.models.static_funcs`, 133
 - `dicee.models.transformers`, 134
 - `dicee.query_generator`, 211
 - `dicee.read_preprocess_save_load_kg`, 213
 - `dicee.read_preprocess_save_load_kg.preprocess`, 213
 - `dicee.read_preprocess_save_load_kg.read_from_disk`, 214
 - `dicee.read_preprocess_save_load_kg.save_load_disk`, 214
 - `dicee.read_preprocess_save_load_kg.util`, 215
 - `dicee.sanity_checkers`, 220
 - `dicee.scripts`, 220
 - `dicee.scripts.index_serve`, 220

- `dicee.scripts.run`, 222
- `dicee.static_funcs`, 222
- `dicee.static_funcs_training`, 227
- `dicee.static_preprocess_funcs`, 228
- `dicee.trainer`, 229
- `dicee.trainer.auto_batch_finder`, 229
- `dicee.trainer.dice_trainer`, 230
- `dicee.trainer.model_parallelism`, 232
- `dicee.trainer.torch_trainer`, 233
- `dicee.trainer.torch_trainer_ddp`, 235
- `dicee.trainer.torch_trainer_fsdp`, 236
- `dicee.weight_averaging`, 239
- `modules()` (*dicee.models.ensemble.EnsembleKGE method*), 110
- `moving_average()` (*dicee.weight_averaging.SWA static method*), 241
- `MultiClassClassificationDataset` (*class in dicee.dataset_classes*), 33
- `MultiLabelDataset` (*class in dicee.dataset_classes*), 34
- `MuRE` (*class in dicee.models*), 162
- `MuRE` (*class in dicee.models.real*), 129

N

- `n` (*dicee.models.FMult2 attribute*), 208
- `n` (*dicee.models.function_space.FMult2 attribute*), 114
- `n_embd` (*dicee.models.clifford.TransformerSelfAttention attribute*), 100
- `n_embd` (*dicee.models.CoKEConfig attribute*), 164, 165
- `n_embd` (*dicee.models.real.CoKEConfig attribute*), 132
- `n_embd` (*dicee.models.transformers.GPTConfig attribute*), 138
- `n_embd` (*dicee.models.transformers.SelfAttention attribute*), 136
- `n_epochs_eval_model` (*dicee.callbacks.PeriodicEvalCallback attribute*), 27
- `n_epochs_eval_model` (*dicee.config.Namespace attribute*), 31
- `n_head` (*dicee.models.clifford.TransformerSelfAttention attribute*), 100
- `n_head` (*dicee.models.CoKEConfig attribute*), 164, 165
- `n_head` (*dicee.models.real.CoKEConfig attribute*), 132
- `n_head` (*dicee.models.transformers.GPTConfig attribute*), 138
- `n_head` (*dicee.models.transformers.SelfAttention attribute*), 136
- `n_layer` (*dicee.models.CoKEConfig attribute*), 164, 165
- `n_layer` (*dicee.models.real.CoKEConfig attribute*), 132
- `n_layer` (*dicee.models.transformers.GPTConfig attribute*), 138
- `n_layers` (*dicee.models.FMult2 attribute*), 208
- `n_layers` (*dicee.models.function_space.FMult2 attribute*), 114
- `name` (*dicee.abstracts.BaseInteractiveKGE property*), 14
- `name` (*dicee.models.AConEx attribute*), 171
- `name` (*dicee.models.AConvO attribute*), 188
- `name` (*dicee.models.AConvQ attribute*), 181
- `name` (*dicee.models.CKeci attribute*), 193
- `name` (*dicee.models.clifford.CKeci attribute*), 100
- `name` (*dicee.models.clifford.DeCaL attribute*), 100
- `name` (*dicee.models.clifford.Keci attribute*), 95
- `name` (*dicee.models.clifford.KeciTransformer attribute*), 99
- `name` (*dicee.models.CoKE attribute*), 165
- `name` (*dicee.models.ComplEx attribute*), 172
- `name` (*dicee.models.complex.AConEx attribute*), 105
- `name` (*dicee.models.complex.ComplEx attribute*), 106
- `name` (*dicee.models.complex.ConEx attribute*), 104
- `name` (*dicee.models.complex.RotatE attribute*), 107
- `name` (*dicee.models.ConEx attribute*), 170
- `name` (*dicee.models.ConvO attribute*), 187
- `name` (*dicee.models.ConvQ attribute*), 180
- `name` (*dicee.models.DeCaL attribute*), 193
- `name` (*dicee.models.DistMult attribute*), 158
- `name` (*dicee.models.DualE attribute*), 210
- `name` (*dicee.models.dualE.DualE attribute*), 108
- `name` (*dicee.models.ensemble.EnsembleKGE attribute*), 109
- `name` (*dicee.models.FMult attribute*), 206
- `name` (*dicee.models.FMult2 attribute*), 208
- `name` (*dicee.models.function_space.FMult attribute*), 113
- `name` (*dicee.models.function_space.FMult2 attribute*), 114
- `name` (*dicee.models.function_space.GFMult attribute*), 113

name (*dicdee.models.function_space.LFMult* attribute), 115
 name (*dicdee.models.function_space.LFMult1* attribute), 115
 name (*dicdee.models.GFMult* attribute), 207
 name (*dicdee.models.Keci* attribute), 189
 name (*dicdee.models.KeciTransformer* attribute), 196
 name (*dicdee.models.LFMult* attribute), 209
 name (*dicdee.models.LFMult1* attribute), 208
 name (*dicdee.models.MuRE* attribute), 162
 name (*dicdee.models.octonion.AConvO* attribute), 120
 name (*dicdee.models.octonion.ConvO* attribute), 119
 name (*dicdee.models.octonion.OMult* attribute), 119
 name (*dicdee.models.OMult* attribute), 187
 name (*dicdee.models.Pyke* attribute), 164
 name (*dicdee.models.pykeen_models.PykeenKGE* attribute), 121
 name (*dicdee.models.PykeenKGE* attribute), 201
 name (*dicdee.models.QMult* attribute), 179
 name (*dicdee.models.quaternion.AConvQ* attribute), 125
 name (*dicdee.models.quaternion.ConvQ* attribute), 124
 name (*dicdee.models.quaternion.QMult* attribute), 123
 name (*dicdee.models.real.CoKE* attribute), 133
 name (*dicdee.models.real.DistMult* attribute), 126
 name (*dicdee.models.real.MuRE* attribute), 130
 name (*dicdee.models.real.Pyke* attribute), 131
 name (*dicdee.models.real.Shallom* attribute), 131
 name (*dicdee.models.real.TransE* attribute), 127
 name (*dicdee.models.real.TransH* attribute), 128
 name (*dicdee.models.RotatE* attribute), 173
 name (*dicdee.models.Shallom* attribute), 163
 name (*dicdee.models.TransE* attribute), 160
 name (*dicdee.models.transformers.BytE* attribute), 134
 name (*dicdee.models.TransH* attribute), 161
 named_children() (*dicdee.models.ensemble.EnsembleKGE* method), 110
 Namespace (class in *dicdee.config*), 28
 neg_ratio (*dicdee.config.Namespace* attribute), 29
 neg_ratio (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 33
 neg_ratio (*dicdee.dataset_classes.FSDP1vsSampleDataset* attribute), 36
 neg_ratio (*dicdee.dataset_classes.KvsSampleDataset* attribute), 37
 neg_sample_ratio (*dicdee.dataset_classes.FixedNegSampleDataset* attribute), 39
 neg_sample_ratio (*dicdee.dataset_classes.OnevsSample* attribute), 40
 neg_sample_ratio (*dicdee.dataset_classes.TriplePredictionDataset* attribute), 40
 negnorm() (*dicdee.abstracts.InteractiveQueryDecomposition* method), 18
 neural_searcher (in module *dicdee.scripts.index_serve*), 221
 NeuralSearcher (class in *dicdee.scripts.index_serve*), 221
 NodeTrainer (class in *dicdee.trainer.torch_trainer_ddp*), 235
 norm_fc1 (*dicdee.models.AConEx* attribute), 171
 norm_fc1 (*dicdee.models.AConvO* attribute), 188
 norm_fc1 (*dicdee.models.complex.AConEx* attribute), 105
 norm_fc1 (*dicdee.models.complex.ConEx* attribute), 104
 norm_fc1 (*dicdee.models.ConEx* attribute), 170
 norm_fc1 (*dicdee.models.ConvO* attribute), 188
 norm_fc1 (*dicdee.models.octonion.AConvO* attribute), 120
 norm_fc1 (*dicdee.models.octonion.ConvO* attribute), 120
 normalization (*dicdee.analyse_experiments.Experiment* attribute), 22
 normalization (*dicdee.config.Namespace* attribute), 30
 normalization_params (*dicdee.dataset_classes.LiteralDataset* attribute), 38
 normalization_type (*dicdee.dataset_classes.LiteralDataset* attribute), 38
 normalize_head_entity_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 91
 normalize_head_entity_embeddings (*dicdee.models.BaseKGE* attribute), 149, 154, 166, 175, 183, 197, 203
 normalize_relation_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 91
 normalize_relation_embeddings (*dicdee.models.BaseKGE* attribute), 149, 154, 166, 175, 183, 197, 203
 normalize_tail_entity_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 91
 normalize_tail_entity_embeddings (*dicdee.models.BaseKGE* attribute), 149, 154, 167, 175, 183, 197, 203
 normalizer_class (*dicdee.models.base_model.BaseKGE* attribute), 91
 normalizer_class (*dicdee.models.BaseKGE* attribute), 149, 154, 166, 175, 183, 197, 203
 num_bpe_entities (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 33
 num_bpe_entities (*dicdee.knowledge_graph.KG* attribute), 70
 num_core (*dicdee.config.Namespace* attribute), 30
 num_datapoints (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 33

num_datapoints (*dicее.dataset_classes.MultiLabelDataset attribute*), 34
 num_embeddings (*dicее.models.fsdп_models.RowWiseShardedEmbedding attribute*), 111
 num_ent (*dicее.models.DualE attribute*), 210
 num_ent (*dicее.models.dualE.DualE attribute*), 108
 num_entities (*dicее.dataset_classes.FixedNegSampleDataset attribute*), 39
 num_entities (*dicее.dataset_classes.FSDP1vsSampleDataset attribute*), 36
 num_entities (*dicее.dataset_classes.KvsSampleDataset attribute*), 37
 num_entities (*dicее.dataset_classes.LiteralDataset attribute*), 38
 num_entities (*dicее.dataset_classes.OnevsSample attribute*), 40
 num_entities (*dicее.dataset_classes.TriplePredictionDataset attribute*), 40
 num_entities (*dicее.evaluation.Evaluator attribute*), 56, 57
 num_entities (*dicее.evaluation.evaluator.Evaluator attribute*), 46, 47
 num_entities (*dicее.Evaluator attribute*), 255
 num_entities (*dicее.evaluator.Evaluator attribute*), 63
 num_entities (*dicее.knowledge_graph.KG attribute*), 69, 70
 num_entities (*dicее.models.base_model.BaseKGE attribute*), 90
 num_entities (*dicее.models.BaseKGE attribute*), 148, 154, 166, 175, 182, 197, 202
 num_epochs (*dicее.abstracts.AbstractPPECallback attribute*), 19
 num_epochs (*dicее.analyse_experiments.Experiment attribute*), 21
 num_epochs (*dicее.config.Namespace attribute*), 29
 num_epochs (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 236
 num_epochs (*dicее.weight_averaging.ASWA attribute*), 239
 num_folds_for_cv (*dicее.config.Namespace attribute*), 30
 num_negatives (*dicее.dataset_classes.FSDP1vsSampleDataset attribute*), 36
 num_of_data_points (*dicее.dataset_classes.MultiClassClassificationDataset attribute*), 33
 num_of_data_properties (*dicее.models.literal.LiteralEmbeddings attribute*), 117
 num_of_epochs (*dicее.callbacks.PseudoLabellingCallback attribute*), 25
 num_of_output_channels (*dicее.config.Namespace attribute*), 30
 num_of_output_channels (*dicее.models.base_model.BaseKGE attribute*), 91
 num_of_output_channels (*dicее.models.BaseKGE attribute*), 149, 154, 166, 175, 183, 197, 203
 num_params (*dicее.analyse_experiments.Experiment attribute*), 21
 num_relations (*dicее.dataset_classes.FixedNegSampleDataset attribute*), 39
 num_relations (*dicее.dataset_classes.FSDP1vsSampleDataset attribute*), 36
 num_relations (*dicее.dataset_classes.OnevsSample attribute*), 40
 num_relations (*dicее.dataset_classes.TriplePredictionDataset attribute*), 40
 num_relations (*dicее.evaluation.Evaluator attribute*), 57
 num_relations (*dicее.evaluation.evaluator.Evaluator attribute*), 46, 47
 num_relations (*dicее.Evaluator attribute*), 255
 num_relations (*dicее.evaluator.Evaluator attribute*), 63
 num_relations (*dicее.knowledge_graph.KG attribute*), 69, 70
 num_relations (*dicее.models.base_model.BaseKGE attribute*), 90
 num_relations (*dicее.models.BaseKGE attribute*), 148, 154, 166, 175, 182, 197, 202
 num_sample (*dicее.models.FMult attribute*), 207
 num_sample (*dicее.models.function_space.FMult attribute*), 113
 num_sample (*dicее.models.function_space.GFMult attribute*), 114
 num_sample (*dicее.models.GFMult attribute*), 207
 num_samples (*dicее.weight_averaging.TWA attribute*), 243
 num_tokens (*dicее.knowledge_graph.KG attribute*), 70
 num_tokens (*dicее.models.base_model.BaseKGE attribute*), 90
 num_tokens (*dicее.models.BaseKGE attribute*), 149, 154, 166, 175, 182, 197, 202
 num_workers (*dicее.trainer.torch_trainer_fsdп.TorchFSDPTrainer attribute*), 237
 numpy_data_type_changer() (in module *dicее.static_funcs*), 225

O

octonion_mul() (in module *dicее.models*), 187
 octonion_mul() (in module *dicее.models.octonion*), 119
 octonion_mul_norm() (in module *dicее.models*), 187
 octonion_mul_norm() (in module *dicее.models.octonion*), 119
 octonion_normalizer() (*dicее.models.AConvO static method*), 188
 octonion_normalizer() (*dicее.models.ConvO static method*), 188
 octonion_normalizer() (*dicее.models.octonion.AConvO static method*), 120
 octonion_normalizer() (*dicее.models.octonion.ConvO static method*), 120
 octonion_normalizer() (*dicее.models.octonion.OMult static method*), 119
 octonion_normalizer() (*dicее.models.OMult static method*), 187
 OMult (class in *dicее.models*), 187
 OMult (class in *dicее.models.octonion*), 119
 on_epoch_end() (*dicее.callbacks.KGESaveCallback method*), 24

- `on_epoch_end()` (*dicee.callbacks.PseudoLabellingCallback* method), 25
- `on_fit_end()` (*dicee.abstracts.AbstractCallback* method), 19
- `on_fit_end()` (*dicee.abstracts.AbstractPPECallback* method), 19
- `on_fit_end()` (*dicee.abstracts.AbstractTrainer* method), 13
- `on_fit_end()` (*dicee.callbacks.AccumulateEpochLossCallback* method), 23
- `on_fit_end()` (*dicee.callbacks.Eval* method), 25
- `on_fit_end()` (*dicee.callbacks.KGESaveCallback* method), 24
- `on_fit_end()` (*dicee.callbacks.LRScheduler* method), 28
- `on_fit_end()` (*dicee.callbacks.PeriodicEvalCallback* method), 27
- `on_fit_end()` (*dicee.callbacks.PrintCallback* method), 23
- `on_fit_end()` (*dicee.weight_averaging.ASWA* method), 240
- `on_fit_end()` (*dicee.weight_averaging.EMA* method), 243
- `on_fit_end()` (*dicee.weight_averaging.SWA* method), 241
- `on_fit_end()` (*dicee.weight_averaging.SWAG* method), 242
- `on_fit_end()` (*dicee.weight_averaging.TWA* method), 244
- `on_fit_start()` (*dicee.abstracts.AbstractCallback* method), 18
- `on_fit_start()` (*dicee.abstracts.AbstractPPECallback* method), 19
- `on_fit_start()` (*dicee.abstracts.AbstractTrainer* method), 13
- `on_fit_start()` (*dicee.callbacks.Eval* method), 25
- `on_fit_start()` (*dicee.callbacks.KGESaveCallback* method), 24
- `on_fit_start()` (*dicee.callbacks.KronE* method), 26
- `on_fit_start()` (*dicee.callbacks.PrintCallback* method), 23
- `on_init_end()` (*dicee.abstracts.AbstractCallback* method), 18
- `on_init_start()` (*dicee.abstracts.AbstractCallback* method), 18
- `on_train_batch_end()` (*dicee.abstracts.AbstractCallback* method), 19
- `on_train_batch_end()` (*dicee.abstracts.AbstractTrainer* method), 13
- `on_train_batch_end()` (*dicee.callbacks.Eval* method), 26
- `on_train_batch_end()` (*dicee.callbacks.KGESaveCallback* method), 24
- `on_train_batch_end()` (*dicee.callbacks.LRScheduler* method), 28
- `on_train_batch_end()` (*dicee.callbacks.PrintCallback* method), 23
- `on_train_batch_start()` (*dicee.callbacks.Perturb* method), 27
- `on_train_epoch_end()` (*dicee.abstracts.AbstractCallback* method), 19
- `on_train_epoch_end()` (*dicee.abstracts.AbstractTrainer* method), 13
- `on_train_epoch_end()` (*dicee.callbacks.Eval* method), 25
- `on_train_epoch_end()` (*dicee.callbacks.KGESaveCallback* method), 24
- `on_train_epoch_end()` (*dicee.callbacks.PeriodicEvalCallback* method), 27
- `on_train_epoch_end()` (*dicee.callbacks.PrintCallback* method), 23
- `on_train_epoch_end()` (*dicee.models.base_model.BaseKGELightning* method), 87
- `on_train_epoch_end()` (*dicee.models.BaseKGELightning* method), 145
- `on_train_epoch_end()` (*dicee.weight_averaging.ASWA* method), 240
- `on_train_epoch_end()` (*dicee.weight_averaging.EMA* method), 243
- `on_train_epoch_end()` (*dicee.weight_averaging.SWA* method), 241
- `on_train_epoch_end()` (*dicee.weight_averaging.SWAG* method), 242
- `on_train_epoch_end()` (*dicee.weight_averaging.TWA* method), 244
- `on_train_epoch_start()` (*dicee.abstracts.AbstractTrainer* method), 13
- `on_train_epoch_start()` (*dicee.weight_averaging.EMA* method), 243
- `on_train_epoch_start()` (*dicee.weight_averaging.SWA* method), 241
- `on_train_epoch_start()` (*dicee.weight_averaging.SWAG* method), 242
- `on_train_epoch_start()` (*dicee.weight_averaging.TWA* method), 244
- `on_train_start()` (*dicee.callbacks.LRScheduler* method), 28
- `OnevsAllDataset` (class in *dicee.dataset_classes*), 37
- `OnevsSample` (class in *dicee.dataset_classes*), 39
- `optim` (*dicee.config.Namespace* attribute), 29
- `OptimizedModule` (in module *dicee.trainer.torch_trainer_fsdp*), 236
- `optimizer` (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 235
- `optimizer` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237
- `optimizer` (*dicee.trainer.torch_trainer.TorchTrainer* attribute), 234
- `optimizer_name` (*dicee.models.base_model.BaseKGE* attribute), 90
- `optimizer_name` (*dicee.models.BaseKGE* attribute), 149, 154, 166, 175, 182, 197, 203
- `ordered_bpe_entities` (*dicee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 33
- `ordered_bpe_entities` (*dicee.knowledge_graph.KG* attribute), 71
- `ordered_shaped_bpe_tokens` (*dicee.knowledge_graph.KG* attribute), 70

P

- `p` (*dicee.config.Namespace* attribute), 30
- `p` (*dicee.models.clifford.DeCaL* attribute), 100
- `p` (*dicee.models.clifford.Keci* attribute), 95

p (*dicee.models.DeCaL* attribute), 193
p (*dicee.models.Keci* attribute), 189
P (*dicee.weight_averaging.TWA* attribute), 244
padding (*dicee.knowledge_graph.KG* attribute), 71
pandas_dataframe_indexer() (in module *dicee.read_preprocess_save_load_kg.util*), 217
param_init (*dicee.models.base_model.BaseKGE* attribute), 91
param_init (*dicee.models.BaseKGE* attribute), 149, 155, 167, 175, 183, 197, 203
parameters() (*dicee.abstracts.BaseInteractiveKGE* method), 17
parameters() (*dicee.models.ensemble.EnsembleKGE* method), 110
path (*dicee.abstracts.AbstractPPECallback* attribute), 19
path (*dicee.callbacks.AccumulateEpochLossCallback* attribute), 23
path (*dicee.callbacks.Eval* attribute), 25
path (*dicee.callbacks.KGESaveCallback* attribute), 24
path (*dicee.weight_averaging.ASWA* attribute), 239
path_dataset_folder (*dicee.analyse_experiments.Experiment* attribute), 21
path_for_deserialization (*dicee.knowledge_graph.KG* attribute), 70
path_for_serialization (*dicee.knowledge_graph.KG* attribute), 70
path_single_kg (*dicee.config.Namespace* attribute), 28
path_single_kg (*dicee.knowledge_graph.KG* attribute), 70
path_to_store_single_run (*dicee.config.Namespace* attribute), 28
PeriodicEvalCallback (class in *dicee.callbacks*), 27
Perturb (class in *dicee.callbacks*), 26
pl_trainer_kwargs (*dicee.config.Namespace* attribute), 30
polars_dataframe_indexer() (in module *dicee.read_preprocess_save_load_kg.util*), 216
poly_NN() (*dicee.models.function_space.LFMMult* method), 116
poly_NN() (*dicee.models.LFMMult* method), 209
polynomial() (*dicee.models.function_space.LFMMult* method), 116
polynomial() (*dicee.models.LFMMult* method), 210
pop() (*dicee.models.function_space.LFMMult* method), 116
pop() (*dicee.models.LFMMult* method), 210
pos_emb (*dicee.models.CoKE* attribute), 165
pos_emb (*dicee.models.real.CoKE* attribute), 133
pq (*dicee.analyse_experiments.Experiment* attribute), 21
predict() (*dicee.KGE* method), 248
predict() (*dicee.knowledge_graph_embeddings.KGE* method), 73
predict_data_loader() (*dicee.models.base_model.BaseKGELightning* method), 89
predict_data_loader() (*dicee.models.BaseKGELightning* method), 147
predict_literals() (*dicee.KGE* method), 252
predict_literals() (*dicee.knowledge_graph_embeddings.KGE* method), 77
predict_missing_head_entity() (*dicee.KGE* method), 247
predict_missing_head_entity() (*dicee.knowledge_graph_embeddings.KGE* method), 72
predict_missing_relations() (*dicee.KGE* method), 247
predict_missing_relations() (*dicee.knowledge_graph_embeddings.KGE* method), 72
predict_missing_tail_entity() (*dicee.KGE* method), 248
predict_missing_tail_entity() (*dicee.knowledge_graph_embeddings.KGE* method), 73
predict_topk() (*dicee.KGE* method), 249
predict_topk() (*dicee.knowledge_graph_embeddings.KGE* method), 74
prefetch_factor (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237
preprocess_with_byte_pair_encoding() (*dicee.read_preprocess_save_load_kg.PreprocessKG* method), 219
preprocess_with_byte_pair_encoding() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 213
preprocess_with_byte_pair_encoding_with_padding() (*dicee.read_preprocess_save_load_kg.PreprocessKG* method), 219
preprocess_with_byte_pair_encoding_with_padding() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 213
preprocess_with_pandas() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 213
preprocess_with_polars() (*dicee.read_preprocess_save_load_kg.PreprocessKG* method), 219
preprocess_with_polars() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 213
preprocesses_input_args() (in module *dicee.static_preprocess_funcs*), 229
PreprocessKG (class in *dicee.read_preprocess_save_load_kg*), 219
PreprocessKG (class in *dicee.read_preprocess_save_load_kg.preprocess*), 213
PrintCallback (class in *dicee.callbacks*), 23
process (*dicee.trainer.torch_trainer.TorchTrainer* attribute), 234
PseudoLabellingCallback (class in *dicee.callbacks*), 24
ptdtype (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237
Pyke (class in *dicee.models*), 164
Pyke (class in *dicee.models.real*), 131
pykeen_model_kwargs (*dicee.config.Namespace* attribute), 30
PykeenKGE (class in *dicee.models*), 201
PykeenKGE (class in *dicee.models.pykeen_models*), 121

Q

`q` (*dicee.config.Namespace* attribute), 30
`q` (*dicee.models.clifford.DeCaL* attribute), 100
`q` (*dicee.models.clifford.Keci* attribute), 95
`q` (*dicee.models.DeCaL* attribute), 193
`q` (*dicee.models.Keci* attribute), 189
`qdrant_client` (*dicee.scripts.index_serve.NeuralSearcher* attribute), 221
`QMult` (class in *dicee.models*), 179
`QMult` (class in *dicee.models.quaternion*), 123
`quaternion_mul()` (in module *dicee.models*), 179
`quaternion_mul()` (in module *dicee.models.static_funcs*), 134
`quaternion_mul_with_unit_norm()` (in module *dicee.models*), 179
`quaternion_mul_with_unit_norm()` (in module *dicee.models.quaternion*), 123
`quaternion_multiplication_followed_by_inner_product()` (*dicee.models.QMult* method), 180
`quaternion_multiplication_followed_by_inner_product()` (*dicee.models.quaternion.QMult* method), 123
`quaternion_normalizer()` (*dicee.models.QMult* static method), 180
`quaternion_normalizer()` (*dicee.models.quaternion.QMult* static method), 123
`queries` (*dicee.scripts.index_serve.StringListRequest* attribute), 221
`query_name_to_struct` (*dicee.query_generator.QueryGenerator* attribute), 212
`query_name_to_struct` (*dicee.QueryGenerator* attribute), 252
`QueryGenerator` (class in *dicee*), 252
`QueryGenerator` (class in *dicee.query_generator*), 211

R

`r` (*dicee.models.clifford.DeCaL* attribute), 101
`r` (*dicee.models.clifford.Keci* attribute), 95
`r` (*dicee.models.DeCaL* attribute), 193
`r` (*dicee.models.Keci* attribute), 189
`random_seed` (*dicee.config.Namespace* attribute), 30
`range_constraints_per_rel` (*dicee.evaluation.Evaluator* attribute), 57
`range_constraints_per_rel` (*dicee.evaluation.evaluator.Evaluator* attribute), 47
`range_constraints_per_rel` (*dicee.Evaluator* attribute), 255
`range_constraints_per_rel` (*dicee.evaluator.Evaluator* attribute), 64
`rank` (*dicee.Execute* attribute), 245
`rank` (*dicee.executer.Execute* attribute), 66, 67
`rank` (*dicee.models.fsdp_models.RowWiseShardedEmbedding* attribute), 111
`ratio` (*dicee.callbacks.Perturb* attribute), 26
`raw_model` (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237
`raw_test_set` (*dicee.knowledge_graph.KG* attribute), 70
`raw_train_set` (*dicee.knowledge_graph.KG* attribute), 70
`raw_valid_set` (*dicee.knowledge_graph.KG* attribute), 70
`re` (*dicee.models.clifford.DeCaL* attribute), 101
`re` (*dicee.models.DeCaL* attribute), 193
`re_vocab` (*dicee.evaluation.Evaluator* attribute), 56, 57
`re_vocab` (*dicee.evaluation.evaluator.Evaluator* attribute), 46, 47
`re_vocab` (*dicee.Evaluator* attribute), 255
`re_vocab` (*dicee.evaluator.Evaluator* attribute), 63
`read_from_disk()` (in module *dicee.read_preprocess_save_load_kg.util*), 217
`read_from_triple_store_with_pandas()` (in module *dicee.read_preprocess_save_load_kg.util*), 218
`read_from_triple_store_with_polars()` (in module *dicee.read_preprocess_save_load_kg.util*), 218
`read_only_few` (*dicee.config.Namespace* attribute), 30
`read_only_few` (*dicee.knowledge_graph.KG* attribute), 70
`read_or_load_kg()` (in module *dicee.static_funcs*), 225
`read_with_pandas()` (in module *dicee.read_preprocess_save_load_kg.util*), 217
`read_with_polars()` (in module *dicee.read_preprocess_save_load_kg.util*), 217
`ReadFromDisk` (class in *dicee.read_preprocess_save_load_kg*), 219
`ReadFromDisk` (class in *dicee.read_preprocess_save_load_kg.read_from_disk*), 214
`reducer` (*dicee.scripts.index_serve.StringListRequest* attribute), 221
`reg_lambda` (*dicee.weight_averaging.TWA* attribute), 243
`rel2id` (*dicee.query_generator.QueryGenerator* attribute), 212
`rel2id` (*dicee.QueryGenerator* attribute), 252
`relation_embeddings` (*dicee.models.AConvQ* attribute), 181
`relation_embeddings` (*dicee.models.complex.RotatE* attribute), 107
`relation_embeddings` (*dicee.models.ConvQ* attribute), 181
`relation_embeddings` (*dicee.models.DualE* attribute), 210
`relation_embeddings` (*dicee.models.dualE.DualE* attribute), 108
`relation_embeddings` (*dicee.models.FMult* attribute), 207

- `relation_embeddings` (*dicee.models.FMult2* attribute), 208
- `relation_embeddings` (*dicee.models.function_space.FMult* attribute), 113
- `relation_embeddings` (*dicee.models.function_space.FMult2* attribute), 114
- `relation_embeddings` (*dicee.models.function_space.GFMult* attribute), 113
- `relation_embeddings` (*dicee.models.function_space.LFMult* attribute), 115
- `relation_embeddings` (*dicee.models.function_space.LFMult1* attribute), 115
- `relation_embeddings` (*dicee.models.GFMult* attribute), 207
- `relation_embeddings` (*dicee.models.LFMult* attribute), 209
- `relation_embeddings` (*dicee.models.LFMult1* attribute), 208
- `relation_embeddings` (*dicee.models.pykeen_models.PykeenKGE* attribute), 121
- `relation_embeddings` (*dicee.models.PykeenKGE* attribute), 201
- `relation_embeddings` (*dicee.models.quaternion.AConvQ* attribute), 125
- `relation_embeddings` (*dicee.models.quaternion.ConvQ* attribute), 124
- `relation_embeddings` (*dicee.models.RotatE* attribute), 173
- `relation_normal_translations` (*dicee.models.real.TransH* attribute), 128
- `relation_normal_translations` (*dicee.models.TransH* attribute), 161
- `relation_to_idx` (*dicee.knowledge_graph.KG* attribute), 70
- `relation_translations` (*dicee.models.MuRE* attribute), 162
- `relation_translations` (*dicee.models.real.MuRE* attribute), 130
- `relations_str` (*dicee.knowledge_graph.KG* property), 71
- `reload_dataset()` (in module *dicee.dataset_classes*), 35
- `report` (*dicee.DICE_Trainer* attribute), 253
- `report` (*dicee.evaluation.Evaluator* attribute), 57
- `report` (*dicee.evaluation.evaluator.Evaluator* attribute), 46, 47
- `report` (*dicee.Evaluator* attribute), 255
- `report` (*dicee.evaluator.Evaluator* attribute), 63, 64
- `report` (*dicee.Execute* attribute), 245
- `report` (*dicee.executer.Execute* attribute), 67
- `report` (*dicee.trainer.DICE_Trainer* attribute), 238
- `report` (*dicee.trainer.dice_trainer.DICE_Trainer* attribute), 231
- `reports` (*dicee.callbacks.Eval* attribute), 25
- `reports` (*dicee.callbacks.PeriodicEvalCallback* attribute), 27
- `requires_grad_for_interactions` (*dicee.models.CKeci* attribute), 193
- `requires_grad_for_interactions` (*dicee.models.clifford.CKeci* attribute), 100
- `requires_grad_for_interactions` (*dicee.models.clifford.Keci* attribute), 96
- `requires_grad_for_interactions` (*dicee.models.Keci* attribute), 189
- `resid_dropout` (*dicee.models.clifford.TransformerSelfAttention* attribute), 100
- `resid_dropout` (*dicee.models.transformers.SelfAttention* attribute), 136
- `residual_convolution()` (*dicee.models.AConEx* method), 171
- `residual_convolution()` (*dicee.models.AConvO* method), 188
- `residual_convolution()` (*dicee.models.AConvQ* method), 181
- `residual_convolution()` (*dicee.models.complex.AConEx* method), 105
- `residual_convolution()` (*dicee.models.complex.ConEx* method), 104
- `residual_convolution()` (*dicee.models.ConEx* method), 170
- `residual_convolution()` (*dicee.models.ConvO* method), 188
- `residual_convolution()` (*dicee.models.ConvQ* method), 181
- `residual_convolution()` (*dicee.models.octonion.AConvO* method), 120
- `residual_convolution()` (*dicee.models.octonion.ConvO* method), 120
- `residual_convolution()` (*dicee.models.quaternion.AConvQ* method), 125
- `residual_convolution()` (*dicee.models.quaternion.ConvQ* method), 124
- `retrieve_embedding()` (*dicee.scripts.index_serve.NeuralSearcher* method), 221
- `retrieve_embeddings()` (in module *dicee.scripts.index_serve*), 221
- `return_multi_hop_query_results()` (*dicee.KGE* method), 250
- `return_multi_hop_query_results()` (*dicee.knowledge_graph_embeddings.KGE* method), 75
- `root()` (in module *dicee.scripts.index_serve*), 221
- `roots` (*dicee.models.FMult* attribute), 207
- `roots` (*dicee.models.function_space.FMult* attribute), 113
- `roots` (*dicee.models.function_space.GFMult* attribute), 114
- `roots` (*dicee.models.GFMult* attribute), 207
- `RotatE` (class in *dicee.models*), 173
- `RotatE` (class in *dicee.models.complex*), 107
- `RowWiseShardedEmbedding` (class in *dicee.models.fsdp_models*), 111
- `runtime` (*dicee.analyse_experiments.Experiment* attribute), 22

S

- `sample()` (*dicee.weight_averaging.SWAG* method), 242
- `sample_counter` (*dicee.abstracts.AbstractPPECallback* attribute), 19

`sample_entity()` (*dicее.abstracts.BaseInteractiveKGE method*), 15
`sample_relation()` (*dicее.abstracts.BaseInteractiveKGE method*), 15
`sample_triples_ratio` (*dicее.config.Namespace attribute*), 30
`sample_triples_ratio` (*dicее.knowledge_graph.KG attribute*), 70
`sample_weights()` (*dicее.weight_averaging.TWA method*), 244
`sampling_ratio` (*dicее.dataset_classes.LiteralDataset attribute*), 38
`sanity_check_callback_args()` (*in module dicее.sanity_checkers*), 220
`sanity_checking_with_arguments()` (*in module dicее.sanity_checkers*), 220
`save()` (*dicее.abstracts.BaseInteractiveKGE method*), 15
`save()` (*dicее.read_preprocess_save_load_kg.LoadSaveToDisk method*), 219
`save()` (*dicее.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk method*), 215
`save_checkpoint()` (*dicее.abstracts.AbstractTrainer static method*), 13
`save_checkpoint_model()` (*in module dicее.static_funcs*), 225
`save_embeddings()` (*in module dicее.static_funcs*), 226
`save_embeddings_as_csv` (*dicее.config.Namespace attribute*), 28
`save_every_n_epochs` (*dicее.config.Namespace attribute*), 31
`save_experiment()` (*dicее.analyse_experiments.Experiment method*), 22
`save_local_shard_checkpoint()` (*dicее.models.fsdп_models.FSDPShardedEntityModel method*), 112
`save_model_at_every_epoch` (*dicее.config.Namespace attribute*), 30
`save_model_every_n_epoch` (*dicее.callbacks.PeriodicEvalCallback attribute*), 27
`save_numpy_ndarray()` (*in module dicее.read_preprocess_save_load_kg.util*), 218
`save_numpy_ndarray()` (*in module dicее.static_funcs*), 225
`save_pickle()` (*in module dicее.read_preprocess_save_load_kg.util*), 218
`save_pickle()` (*in module dicее.static_funcs*), 224
`save_queries()` (*dicее.query_generator.QueryGenerator method*), 212
`save_queries()` (*dicее.QueryGenerator method*), 253
`save_queries_and_answers()` (*dicее.query_generator.QueryGenerator static method*), 213
`save_queries_and_answers()` (*dicее.QueryGenerator static method*), 253
`save_trained_model()` (*dicее.Execute method*), 246
`save_trained_model()` (*dicее.executer.Execute method*), 68
`scalar_batch_NN()` (*dicее.models.function_space.LFMult method*), 116
`scalar_batch_NN()` (*dicее.models.LFMult method*), 209
`scaler` (*dicее.callbacks.Perturb attribute*), 27
`scaler` (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 236
`scaler` (*dicее.trainer.torch_trainer_fsdп.TorchFSDPTrainer attribute*), 237
`scheduler` (*dicее.callbacks.LRScheduler attribute*), 28
`score()` (*dicее.models.clifford.Keci method*), 98
`score()` (*dicее.models.CoKE method*), 165
`score()` (*dicее.models.ComplEx static method*), 172
`score()` (*dicее.models.complex.ComplEx static method*), 106
`score()` (*dicее.models.complex.RotatE method*), 107
`score()` (*dicее.models.DistMult method*), 159
`score()` (*dicее.models.Keci method*), 192
`score()` (*dicее.models.octonion.OMult method*), 119
`score()` (*dicее.models.OMult method*), 187
`score()` (*dicее.models.QMult method*), 180
`score()` (*dicее.models.quaternion.QMult method*), 123
`score()` (*dicее.models.real.CoKE method*), 133
`score()` (*dicее.models.real.DistMult method*), 127
`score()` (*dicее.models.real.TransE method*), 127
`score()` (*dicее.models.RotatE method*), 174
`score()` (*dicее.models.TransE method*), 160
`score_func` (*dicее.models.FMult2 attribute*), 208
`score_func` (*dicее.models.function_space.FMult2 attribute*), 114
`scoring_technique` (*dicее.analyse_experiments.Experiment attribute*), 22
`scoring_technique` (*dicее.config.Namespace attribute*), 29
`search()` (*dicее.scripts.index_serve.NeuralSearcher method*), 221
`search_embeddings()` (*in module dicее.scripts.index_serve*), 221
`search_embeddings_batch()` (*in module dicее.scripts.index_serve*), 221
`seed` (*dicее.dataset_classes.FixedNegSampleDataset attribute*), 39
`seed` (*dicее.dataset_classes.TriplePredictionDataset attribute*), 40
`seed` (*dicее.query_generator.QueryGenerator attribute*), 212
`seed` (*dicее.QueryGenerator attribute*), 252
`select_model()` (*in module dicее.static_funcs*), 224
`selected_optimizer` (*dicее.models.base_model.BaseKGE attribute*), 91
`selected_optimizer` (*dicее.models.BaseKGE attribute*), 149, 154, 166, 175, 183, 197, 203
`SelfAttention` (*class in dicее.models.transformers*), 136
`separator` (*dicее.config.Namespace attribute*), 29

separator (*dicee.knowledge_graph.KG attribute*), 70
 seq_len (*dicee.models.clifford.KeciTransformer attribute*), 99
 seq_len (*dicee.models.KeciTransformer attribute*), 196
 sequential_vocabulary_construction() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 219
 sequential_vocabulary_construction() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 213
 serve() (*in module dicee.scripts.index_serve*), 221
 set_global_seed() (*dicee.query_generator.QueryGenerator method*), 212
 set_global_seed() (*dicee.QueryGenerator method*), 252
 set_model_eval_mode() (*dicee.abstracts.BaseInteractiveKGE method*), 14
 set_model_train_mode() (*dicee.abstracts.BaseInteractiveKGE method*), 14
 setup_distributed_training() (*in module dicee.static_funcs*), 224
 setup_executor() (*dicee.Execute method*), 246
 setup_executor() (*dicee.executor.Execute method*), 67
 setup_fsdp_training() (*dicee.models.fsdp_models.FSDPShardedEntityModel method*), 111
 Shallom (*class in dicee.models*), 163
 Shallom (*class in dicee.models.real*), 130
 shallom (*dicee.models.real.Shallom attribute*), 131
 shallom (*dicee.models.Shallom attribute*), 163
 shard_size (*dicee.models.fsdp_models.RowWiseShardedEmbedding attribute*), 111
 sharding_strategy (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
 single_hop_query_answering() (*dicee.KGE method*), 250
 single_hop_query_answering() (*dicee.knowledge_graph_embeddings.KGE method*), 75
 snapshot_dir (*dicee.callbacks.LRScheduler attribute*), 27
 snapshot_loss (*dicee.callbacks.LRScheduler attribute*), 28
 sparql_endpoint (*dicee.config.Namespace attribute*), 29
 sparql_endpoint (*dicee.knowledge_graph.KG attribute*), 70
 sq_mean (*dicee.weight_averaging.SWAG attribute*), 242
 start() (*dicee.DICE_Trainer method*), 254
 start() (*dicee.Execute method*), 246
 start() (*dicee.executor.Execute method*), 68
 start() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 219
 start() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 213
 start() (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk method*), 214
 start() (*dicee.read_preprocess_save_load_kg.ReadFromDisk method*), 219
 start() (*dicee.trainer.DICE_Trainer method*), 238
 start() (*dicee.trainer.dice_trainer.DICE_Trainer method*), 232
 start_row (*dicee.models.fsdp_models.RowWiseShardedEmbedding attribute*), 111
 start_time (*dicee.callbacks.PrintCallback attribute*), 23
 start_time (*dicee.Execute attribute*), 245
 start_time (*dicee.executor.Execute attribute*), 67
 state_dict() (*dicee.models.ensemble.EnsembleKGE method*), 110
 step() (*dicee.models.ADOPT method*), 143
 step() (*dicee.models.adopt.ADOPT method*), 81
 step() (*dicee.models.ensemble.EnsembleKGE method*), 110
 step_count (*dicee.callbacks.LRScheduler attribute*), 28
 storage_path (*dicee.config.Namespace attribute*), 28
 storage_path (*dicee.DICE_Trainer attribute*), 253
 storage_path (*dicee.trainer.DICE_Trainer attribute*), 238
 storage_path (*dicee.trainer.dice_trainer.DICE_Trainer attribute*), 231
 store() (*in module dicee.static_funcs*), 225
 store_ensemble() (*dicee.abstracts.AbstractPPECallback method*), 20
 strategy (*dicee.abstracts.AbstractTrainer attribute*), 13
 StringListRequest (*class in dicee.scripts.index_serve*), 221
 SWA (*class in dicee.weight_averaging*), 240
 swa (*dicee.config.Namespace attribute*), 31
 swa_c_epochs (*dicee.config.Namespace attribute*), 32
 swa_c_epochs (*dicee.weight_averaging.SWA attribute*), 240
 swa_c_epochs (*dicee.weight_averaging.SWAG attribute*), 241
 swa_lr (*dicee.weight_averaging.SWA attribute*), 240
 swa_lr (*dicee.weight_averaging.SWAG attribute*), 241
 swa_model (*dicee.weight_averaging.SWA attribute*), 241
 swa_n (*dicee.weight_averaging.SWA attribute*), 241
 swa_start_epoch (*dicee.config.Namespace attribute*), 31
 swa_start_epoch (*dicee.weight_averaging.SWA attribute*), 240
 swa_start_epoch (*dicee.weight_averaging.SWAG attribute*), 241
 SWAG (*class in dicee.weight_averaging*), 241
 swag (*dicee.config.Namespace attribute*), 31

T

`T()` (*dicее.models.DualE method*), 211
`T()` (*dicее.models.dualE.DualE method*), 109
`t_conorm()` (*dicее.abstracts.InteractiveQueryDecomposition method*), 17
`t_norm()` (*dicее.abstracts.InteractiveQueryDecomposition method*), 17
`target_dim` (*dicее.dataset_classes.AllvsAll attribute*), 35
`target_dim` (*dicее.dataset_classes.MultiLabelDataset attribute*), 34
`target_dim` (*dicее.dataset_classes.OnevsAllDataset attribute*), 37
`target_dim` (*dicее.knowledge_graph.KG attribute*), 71
`temperature` (*dicее.models.transformers.BytE attribute*), 135
`tensor_t_norm()` (*dicее.abstracts.InteractiveQueryDecomposition method*), 17
`TensorParallel` (*class in dicее.trainer.model_parallelism*), 233
`test_dataloader()` (*dicее.models.base_model.BaseKGELightning method*), 87
`test_dataloader()` (*dicее.models.BaseKGELightning method*), 146
`test_epoch_end()` (*dicее.models.base_model.BaseKGELightning method*), 87
`test_epoch_end()` (*dicее.models.BaseKGELightning method*), 146
`test_h1` (*dicее.analyse_experiments.Experiment attribute*), 22
`test_h3` (*dicее.analyse_experiments.Experiment attribute*), 22
`test_h10` (*dicее.analyse_experiments.Experiment attribute*), 22
`test_mrr` (*dicее.analyse_experiments.Experiment attribute*), 22
`test_path` (*dicее.query_generator.QueryGenerator attribute*), 212
`test_path` (*dicее.QueryGenerator attribute*), 252
`test_set` (*dicее.knowledge_graph.KG attribute*), 69, 70
`timeit()` (*in module dicее.read_preprocess_save_load_kg.util*), 217
`timeit()` (*in module dicее.static_funcs*), 224
`timeit()` (*in module dicее.static_preprocess_funcs*), 229
`to()` (*dicее.KGE method*), 247
`to()` (*dicее.knowledge_graph_embeddings.KGE method*), 72
`to()` (*dicее.models.ensemble.EnsembleKGE method*), 110
`to_df()` (*dicее.analyse_experiments.Experiment method*), 22
`topk` (*dicее.models.transformers.BytE attribute*), 135
`topk` (*dicее.scripts.index_serve.NeuralSearcher attribute*), 221
`torch_ordered_shaped_bpe_entities` (*dicее.dataset_classes.MultiLabelDataset attribute*), 34
`TorchDDPTrainer` (*class in dicее.trainer.torch_trainer_ddp*), 235
`TorchFSDPTrainer` (*class in dicее.trainer.torch_trainer_fsdp*), 236
`TorchTrainer` (*class in dicее.trainer.torch_trainer*), 234
`total_epochs` (*dicее.callbacks.LRScheduler attribute*), 27
`total_steps` (*dicее.callbacks.LRScheduler attribute*), 27
`train()` (*dicее.abstracts.BaseInteractiveTrainKGE method*), 20
`train()` (*dicее.trainer.torch_trainer_ddp.NodeTrainer method*), 236
`train_data` (*dicее.dataset_classes.AllvsAll attribute*), 35
`train_data` (*dicее.dataset_classes.FSDPIvsSampleDataset attribute*), 35
`train_data` (*dicее.dataset_classes.KvsAll attribute*), 36
`train_data` (*dicее.dataset_classes.KvsSampleDataset attribute*), 37
`train_data` (*dicее.dataset_classes.MultiClassClassificationDataset attribute*), 33
`train_data` (*dicее.dataset_classes.OnevsAllDataset attribute*), 37
`train_data` (*dicее.dataset_classes.OnevsSample attribute*), 40
`train_dataloader()` (*dicее.models.base_model.BaseKGELightning method*), 89
`train_dataloader()` (*dicее.models.BaseKGELightning method*), 147
`train_dataloaders` (*dicее.trainer.torch_trainer.TorchTrainer attribute*), 234
`train_dataset_loader` (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 235
`train_dataset_loader` (*dicее.trainer.torch_trainer_fsdp.TorchFSDPTrainer attribute*), 237
`train_file_path` (*dicее.dataset_classes.LiteralDataset attribute*), 38
`train_h1` (*dicее.analyse_experiments.Experiment attribute*), 21
`train_h3` (*dicее.analyse_experiments.Experiment attribute*), 21
`train_h10` (*dicее.analyse_experiments.Experiment attribute*), 21
`train_indices_target` (*dicее.dataset_classes.MultiLabelDataset attribute*), 34
`train_k_vs_all()` (*dicее.abstracts.BaseInteractiveTrainKGE method*), 20
`train_literals()` (*dicее.abstracts.BaseInteractiveTrainKGE method*), 20
`train_mode` (*dicее.models.ensemble.EnsembleKGE attribute*), 109
`train_mrr` (*dicее.analyse_experiments.Experiment attribute*), 21
`train_path` (*dicее.query_generator.QueryGenerator attribute*), 211
`train_path` (*dicее.QueryGenerator attribute*), 252
`train_set` (*dicее.dataset_classes.BPE_NegativeSamplingDataset attribute*), 33
`train_set` (*dicее.dataset_classes.MultiLabelDataset attribute*), 34
`train_set` (*dicее.knowledge_graph.KG attribute*), 69, 70
`train_set_target` (*dicее.knowledge_graph.KG attribute*), 71
`train_target` (*dicее.dataset_classes.AllvsAll attribute*), 35

train_target (*dicee.dataset_classes.KvsAll* attribute), 36
 train_target (*dicee.dataset_classes.KvsSampleDataset* attribute), 37
 train_target_indices (*dicee.knowledge_graph.KG* attribute), 71
 train_triples (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 39
 train_triples() (*dicee.abstracts.BaseInteractiveTrainKGE* method), 20
 trained_model (*dicee.Execute* attribute), 245
 trained_model (*dicee.executer.Execute* attribute), 67
 trainer (*dicee.config.Namespace* attribute), 29
 trainer (*dicee.DICE_Trainer* attribute), 253
 trainer (*dicee.Execute* attribute), 245
 trainer (*dicee.executer.Execute* attribute), 67
 trainer (*dicee.trainer.DICE_Trainer* attribute), 238
 trainer (*dicee.trainer.dice_trainer.DICE_Trainer* attribute), 231
 trainer (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 235
 training_step (*dicee.trainer.torch_trainer.TorchTrainer* attribute), 234
 training_step() (*dicee.models.base_model.BaseKGELightning* method), 86
 training_step() (*dicee.models.BaseKGELightning* method), 145
 training_step() (*dicee.models.transformers.BytE* method), 135
 training_step_outputs (*dicee.models.base_model.BaseKGELightning* attribute), 86
 training_step_outputs (*dicee.models.BaseKGELightning* attribute), 145
 training_technique (*dicee.knowledge_graph.KG* attribute), 70
 TransE (*class in dicee.models*), 159
 TransE (*class in dicee.models.real*), 127
 transformer (*dicee.models.clifford.KeciTransformer* attribute), 99
 transformer (*dicee.models.KeciTransformer* attribute), 196
 transformer (*dicee.models.transformers.BytE* attribute), 135
 transformer (*dicee.models.transformers.GPT* attribute), 139
 TransformerBlock (*class in dicee.models.clifford*), 99
 TransformerMLP (*class in dicee.models.clifford*), 100
 TransformerSelfAttention (*class in dicee.models.clifford*), 99
 TransH (*class in dicee.models*), 160
 TransH (*class in dicee.models.real*), 128
 trapezoid() (*dicee.models.FMult2* method), 208
 trapezoid() (*dicee.models.function_space.FMult2* method), 114
 tri_score() (*dicee.models.function_space.LFMult* method), 116
 tri_score() (*dicee.models.function_space.LFMult1* method), 115
 tri_score() (*dicee.models.LFMult* method), 210
 tri_score() (*dicee.models.LFMult1* method), 209
 triple_score() (*dicee.KGE* method), 249
 triple_score() (*dicee.knowledge_graph_embeddings.KGE* method), 74
 TriplePredictionDataset (*class in dicee.dataset_classes*), 40
 tuple2list() (*dicee.query_generator.QueryGenerator* method), 212
 tuple2list() (*dicee.QueryGenerator* method), 252
 TWA (*class in dicee.weight_averaging*), 243
 twa (*dicee.config.Namespace* attribute), 31
 twa_c_epochs (*dicee.weight_averaging.TWA* attribute), 243
 twa_model (*dicee.weight_averaging.TWA* attribute), 243
 twa_start_epoch (*dicee.weight_averaging.TWA* attribute), 243

U

unlabelled_size (*dicee.callbacks.PseudoLabellingCallback* attribute), 25
 unmap() (*dicee.query_generator.QueryGenerator* method), 212
 unmap() (*dicee.QueryGenerator* method), 253
 unmap_query() (*dicee.query_generator.QueryGenerator* method), 212
 unmap_query() (*dicee.QueryGenerator* method), 253
 update_hits() (*in module dicee.evaluation.utils*), 54
 use_clifford_mul (*dicee.models.clifford.KeciTransformer* attribute), 99
 use_clifford_mul (*dicee.models.KeciTransformer* attribute), 196
 use_compile (*dicee.trainer.torch_trainer_fsdp.TorchFSDPTrainer* attribute), 237

V

val_aswa (*dicee.weight_averaging.ASWA* attribute), 239
 val_dataloader() (*dicee.models.base_model.BaseKGELightning* method), 88
 val_dataloader() (*dicee.models.BaseKGELightning* method), 146
 val_h1 (*dicee.analyse_experiments.Experiment* attribute), 21
 val_h3 (*dicee.analyse_experiments.Experiment* attribute), 21
 val_h10 (*dicee.analyse_experiments.Experiment* attribute), 22

val_mrr (*dicee.analyse_experiments.Experiment attribute*), 21
 val_path (*dicee.query_generator.QueryGenerator attribute*), 212
 val_path (*dicee.QueryGenerator attribute*), 252
 VALID_SCORING_TECHNIQUES (*in module dicee.evaluation.evaluator*), 46
 valid_set (*dicee.knowledge_graph.KG attribute*), 69, 70
 validate_knowledge_graph() (*in module dicee.sanity_checkers*), 220
 var_clamp (*dicee.weight_averaging.SWAG attribute*), 242
 vocab_preparation() (*dicee.evaluation.Evaluator method*), 57
 vocab_preparation() (*dicee.evaluation.evaluator.Evaluator method*), 47
 vocab_preparation() (*dicee.Evaluator method*), 256
 vocab_preparation() (*dicee.evaluator.Evaluator method*), 64
 vocab_size (*dicee.models.CoKEConfig attribute*), 164, 165
 vocab_size (*dicee.models.real.CoKEConfig attribute*), 132
 vocab_size (*dicee.models.transformers.GPTConfig attribute*), 138
 vocab_to_parquet() (*in module dicee.static_funcs*), 226
 vtp_score() (*dicee.models.function_space.LFMult method*), 116
 vtp_score() (*dicee.models.function_space.LFMult method*), 115
 vtp_score() (*dicee.models.LFMult method*), 210
 vtp_score() (*dicee.models.LFMult method*), 209

W

warmup_steps (*dicee.callbacks.LRScheduler attribute*), 28
 weight (*dicee.models.fsdg_models.RowWiseShardedEmbedding attribute*), 111
 weight (*dicee.models.transformers.LayerNorm attribute*), 135
 weight_decay (*dicee.config.Namespace attribute*), 29
 weight_decay (*dicee.models.base_model.BaseKGE attribute*), 91
 weight_decay (*dicee.models.BaseKGE attribute*), 149, 154, 166, 175, 183, 197, 203
 weight_samples (*dicee.weight_averaging.TWA attribute*), 243
 weights (*dicee.models.FMult attribute*), 207
 weights (*dicee.models.function_space.FMult attribute*), 113
 weights (*dicee.models.function_space.GFMult attribute*), 114
 weights (*dicee.models.GFMult attribute*), 207
 world_size (*dicee.Execute attribute*), 245
 world_size (*dicee.executer.Execute attribute*), 66, 67
 world_size (*dicee.models.fsdg_models.RowWiseShardedEmbedding attribute*), 111
 write_csv_from_model_parallel() (*in module dicee.static_funcs*), 226
 write_links() (*dicee.query_generator.QueryGenerator method*), 212
 write_links() (*dicee.QueryGenerator method*), 253
 write_report() (*dicee.Execute method*), 246
 write_report() (*dicee.executer.Execute method*), 68

X

x_values (*dicee.models.function_space.LFMult attribute*), 115
 x_values (*dicee.models.LFMult attribute*), 209