
DICE Embeddings

Release 0.1.3.2

Caglar Demir

Feb 11, 2026

Contents:

1	Dicee Manual	2
2	Installation	3
2.1	Installation from Source	3
3	Download Knowledge Graphs	3
4	Knowledge Graph Embedding Models	3
5	How to Train	3
6	Creating an Embedding Vector Database	5
6.1	Learning Embeddings	5
6.2	Loading Embeddings into Qdrant Vector Database	6
6.3	Launching Webservice	6
7	Answering Complex Queries	6
8	Predicting Missing Links	8
9	Downloading Pretrained Models	8
10	How to Deploy	8
11	Docker	8
12	Coverage Report	8
13	How to cite	10
14	dicee	12
14.1	Submodules	12
14.2	Attributes	209
14.3	Classes	209
14.4	Package Contents	210
	Python Module Index	222

DICE Embeddings¹: Hardware-agnostic Framework for Large-scale Knowledge Graph Embeddings:

1 Dicee Manual

Version: dicee 0.3.2

GitHub repository: <https://github.com/dice-group/dice-embeddings>

Publisher and maintainer: Caglar Demir²

Contact: caglar.demir@upb.de

License: OSI Approved :: MIT License

Dicee is a hardware-agnostic framework for large-scale knowledge graph embeddings.

Knowledge graph embedding research has mainly focused on learning continuous representations of knowledge graphs towards the link prediction problem. Recently developed frameworks can be effectively applied in a wide range of research-related applications. Yet, using these frameworks in real-world applications becomes more challenging as the size of the knowledge graph grows

We developed the DICE Embeddings framework (dicee) to compute embeddings for large-scale knowledge graphs in a hardware-agnostic manner. To achieve this goal, we rely on

1. **Pandas**³ & Co. to use parallelism at preprocessing a large knowledge graph,
2. **PyTorch**⁴ & Co. to learn knowledge graph embeddings via multi-CPU, GPUs, TPUs or computing cluster, and
3. **Huggingface**⁵ to ease the deployment of pre-trained models.

Why Pandas⁶ & Co. ? A large knowledge graph can be read and preprocessed (e.g. removing literals) by pandas, modin, or polars in parallel. Through polars, a knowledge graph having more than 1 billion triples can be read in parallel fashion. Importantly, using these frameworks allow us to perform all necessary computations on a single CPU as well as a cluster of computers.

Why PyTorch⁷ & Co. ? PyTorch is one of the most popular machine learning frameworks available at the time of writing. PytorchLightning facilitates scaling the training procedure of PyTorch without boilerplate. In our framework, we combine **PyTorch**⁸ & **PytorchLightning**⁹. Users can choose the trainer class (e.g., DDP by Pytorch) to train large knowledge graph embedding models with billions of parameters. PytorchLightning allows us to use state-of-the-art model parallelism techniques (e.g. Fully Sharded Training, FairScale, or DeepSpeed) without extra effort. With our framework, practitioners can directly use PytorchLightning for model parallelism to train gigantic embedding models.

Why Huggingface¹⁰? Seamlessly deploy and share pre-trained embedding models through the Huggingface ecosystem.

¹ <https://github.com/dice-group/dice-embeddings>

² <https://github.com/Demirrr>

³ <https://pandas.pydata.org/>

⁴ <https://pytorch.org/>

⁵ <https://huggingface.co/>

⁶ <https://pandas.pydata.org/>

⁷ <https://pytorch.org/>

⁸ <https://pytorch.org/>

⁹ <https://www.pytorchlightning.ai/>

¹⁰ <https://huggingface.co/>

2 Installation

2.1 Installation from Source

```
git clone https://github.com/dice-group/dice-embeddings.git
conda create -n dice python=3.10.13 --no-default-packages && conda activate dice &&
↪ cd dice-embeddings &&
pip3 install -e .
```

or

```
pip install dicee
```

3 Download Knowledge Graphs

```
wget https://files.dice-research.org/datasets/dice-embeddings/KGs.zip --no-check-
↪ certificate && unzip KGs.zip
```

To test the Installation

```
python -m pytest -p no:warnings -x # Runs >114 tests leading to > 15 mins
python -m pytest -p no:warnings --lf # run only the last failed test
python -m pytest -p no:warnings --ff # to run the failures first and then the rest of
↪ the tests.
```

4 Knowledge Graph Embedding Models

1. TransE, DistMult, ComplEx, ConEx, QMult, OMult, ConvO, ConvQ, Keci
2. All 44 models available in <https://github.com/pykeen/pykeen#models>

For more, please refer to examples.

5 How to Train

To Train a KGE model (KECI) and evaluate it on the train, validation, and test sets of the UMLS benchmark dataset.

```
from dicee.executer import Execute
from dicee.config import Namespace
args = Namespace()
args.model = 'Keci'
args.scoring_technique = "KvsAll" # 1vsAll, or AllvsAll, or NegSample
args.dataset_dir = "KGs/UMLS"
args.path_to_store_single_run = "Keci_UMLS"
args.num_epochs = 100
args.embedding_dim = 32
args.batch_size = 1024
reports = Execute(args).start()
print(reports["Train"]["MRR"]) # => 0.9912
print(reports["Test"]["MRR"]) # => 0.8155
# See the Keci_UMLS folder embeddings and all other files
```

where the data is in the following form

```
$ head -3 KGs/UMLS/train.txt
acquired_abnormality    location_of      experimental_model_of_disease
anatomical_abnormality  manifestation_of physiologic_function
alga    isa        entity
```

A KGE model can also be trained from the command line

```
dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
```

dicee automatically detects available GPUs and trains a model with distributed data parallels technique. Under the hood, dicee uses lightning as a default trainer.

```
# Train a model by only using the GPU-0
CUDA_VISIBLE_DEVICES=0 dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model
↪ "train_val_test"
# Train a model by only using GPU-1
CUDA_VISIBLE_DEVICES=1 dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model
↪ "train_val_test"
NCCL_P2P_DISABLE=1 CUDA_VISIBLE_DEVICES=0,1 python dicee/scripts/run.py --trainer PL -
↪ --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
```

Under the hood, dicee executes run.py script and uses lightning as a default trainer

```
# Two equivalent executions
# (1)
dicee --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
↪ 9753123402351737}
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
↪ 'MRR': 0.8072362996241839}
# Evaluate Keci on Test set: Evaluate Keci on Test set
# {'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
↪ 'MRR': 0.8064032293278861}

# (2)
CUDA_VISIBLE_DEVICES=0,1 python dicee/scripts/run.py --trainer PL --dataset_dir "KGs/
↪ UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
↪ 9753123402351737}
# Evaluate Keci on Train set: Evaluate Keci on Train set
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
↪ 'MRR': 0.8072362996241839}
# Evaluate Keci on Test set: Evaluate Keci on Test set
# {'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
↪ 'MRR': 0.8064032293278861}
```

Similarly, models can be easily trained with torchrun

```
torchrun --standalone --nnodes=1 --nproc_per_node=gpu dicee/scripts/run.py --trainer_
→torchDDP --dataset_dir "KGs/UMLS" --model Keci --eval_model "train_val_test"
# Evaluate Keci on Train set: Evaluate Keci on Train set: Evaluate Keci on Train set
# {'H@1': 0.9518788343558282, 'H@3': 0.9988496932515337, 'H@10': 1.0, 'MRR': 0.
→9753123402351737}
# Evaluate Keci on Validation set: Evaluate Keci on Validation set
# {'H@1': 0.6932515337423313, 'H@3': 0.9041411042944786, 'H@10': 0.9754601226993865,
→'MRR': 0.8072499937521418}
# Evaluate Keci on Test set: Evaluate Keci on Test set
{'H@1': 0.6951588502269289, 'H@3': 0.9039334341906202, 'H@10': 0.9750378214826021,
→'MRR': 0.8064032293278861}
```

You can also train a model in multi-node multi-gpu setting.

```
torchrun --nnodes 2 --nproc_per_node=gpu --node_rank 0 --rdzv_id 455 --rdzv_backend_
→c10d --rdzv_endpoint=nebula dicee/scripts/run.py --trainer torchDDP --dataset_dir_
→KGs/UMLS
torchrun --nnodes 2 --nproc_per_node=gpu --node_rank 1 --rdzv_id 455 --rdzv_backend_
→c10d --rdzv_endpoint=nebula dicee/scripts/run.py --trainer torchDDP --dataset_dir_
→KGs/UMLS
```

Train a KGE model by providing the path of a single file and store all parameters under newly created directory called KeciFamilyRun.

```
dicee --path_single_kg "KGs/Family/family-benchmark_rich_background.owl" --model Keci_
→--path_to_store_single_run KeciFamilyRun --backend rdflib
```

where the data is in the following form

```
$ head -3 KGs/Family/train.txt
_:1 <http://www.w3.org/1999/02/22-rdf-syntax-ns#type> <http://www.w3.org/2002/07/owl
→#Ontology> .
<http://www.benchmark.org/family#hasChild> <http://www.w3.org/1999/02/22-rdf-syntax-ns
→#type> <http://www.w3.org/2002/07/owl#ObjectProperty> .
<http://www.benchmark.org/family#hasParent> <http://www.w3.org/1999/02/22-rdf-syntax-
→ns#type> <http://www.w3.org/2002/07/owl#ObjectProperty> .
```

Apart from n-triples or standard link prediction dataset formats, we support ["owl", "nt", "turtle", "rdf/xml", "n3"]*. Moreover, a KGE model can be also trained by providing an endpoint of a triple store.

```
dicee --sparql_endpoint "http://localhost:3030/mutagenesis/" --model Keci
```

For more, please refer to examples.

6 Creating an Embedding Vector Database

6.1 Learning Embeddings

```
# Train an embedding model
dicee --dataset_dir KGs/Countries-S1 --path_to_store_single_run CountryEmbeddings --
→model Keci --p 0 --q 1 --embedding_dim 32 --adaptive_swa
```

6.2 Loading Embeddings into Qdrant Vector Database

```
# Ensure that Qdrant available
# docker pull qdrant/qdrant && docker run -p 6333:6333 -p 6334:6334 -v $(pwd)/
↪qdrant_storage:/qdrant/storage:z qdrant/qdrant
diceeindex --path_model "CountryEmbeddings" --collection_name "dummy" --location
↪"localhost"
```

6.3 Launching Webservice

```
diceeserve --path_model "CountryEmbeddings" --collection_name "dummy" --collection_
↪location "localhost"
```

Retrieve and Search

Get embedding of germany

```
curl -X 'GET' 'http://0.0.0.0:8000/api/get?q=germany' -H 'accept: application/json'
```

Get most similar things to europe

```
curl -X 'GET' 'http://0.0.0.0:8000/api/search?q=europe' -H 'accept: application/json'
{"result":[{"hit":"europe","score":1.0},
{"hit":"northern_europe","score":0.67126536},
{"hit":"western_europe","score":0.6010134},
{"hit":"puerto_rico","score":0.5051694},
{"hit":"southern_europe","score":0.4829831}]}
```

7 Answering Complex Queries

```
# pip install dicee
# wget https://files.dice-research.org/datasets/dice-embeddings/KGs.zip --no-check-
↪certificate & unzip KGs.zip
from dicee.executer import Execute
from dicee.config import Namespace
from dicee.knowledge_graph_embeddings import KGE
# (1) Train a KGE model
args = Namespace()
args.model = 'Keci'
args.p=0
args.q=1
args.optim = 'Adam'
args.scoring_technique = "AllvsAll"
args.path_single_kg = "KGs/Family/family-benchmark_rich_background.owl"
args.backend = "rdflib"
args.num_epochs = 200
args.batch_size = 1024
args.lr = 0.1
args.embedding_dim = 512
result = Execute(args).start()
# (2) Load the pre-trained model
```

(continues on next page)

(continued from previous page)

```
pre_trained_kge = KGE(path=result['path_experiment_folder'])
# (3) Single-hop query answering
# Query: ?E : \exist E.hasSibling(E, F9M167)
# Question: Who are the siblings of F9M167?
# Answer: [F9M157, F9F141], as (F9M167, hasSibling, F9M157) and (F9M167, hasSibling, ↵
↵F9F141)
predictions = pre_trained_kge.answer_multi_hop_query(query_type="1p",
                                                       query=('http://www.benchmark.org/
↵family#F9M167',
                                                         ('http://www.benchmark.
↵org/family#hasSibling',)),
                                                       tnorm="min", k=3)
top_entities = [topk_entity for topk_entity, query_score in predictions]
assert "http://www.benchmark.org/family#F9F141" in top_entities
assert "http://www.benchmark.org/family#F9M157" in top_entities
# (2) Two-hop query answering
# Query: ?D : \exist E.Married(D, E) \land hasSibling(E, F9M167)
# Question: To whom a sibling of F9M167 is married to?
# Answer: [F9F158, F9M142] as (F9M157 #married F9F158) and (F9F141 #married F9M142)
predictions = pre_trained_kge.answer_multi_hop_query(query_type="2p",
                                                       query=("http://www.benchmark.org/
↵family#F9M167",
                                                         ("http://www.benchmark.
↵org/family#hasSibling",
                                                         "http://www.benchmark.
↵org/family#married")),
                                                       tnorm="min", k=3)
top_entities = [topk_entity for topk_entity, query_score in predictions]
assert "http://www.benchmark.org/family#F9M142" in top_entities
assert "http://www.benchmark.org/family#F9F158" in top_entities
# (3) Three-hop query answering
# Query: ?T : \exist D.type(D,T) \land Married(D,E) \land hasSibling(E, F9M167)
# Question: What are the type of people who are married to a sibling of F9M167?
# (3) Answer: [Person, Male, Father] since F9M157 is [Brother Father Grandfather ↵
↵Male] and F9M142 is [Male Grandfather Father]
predictions = pre_trained_kge.answer_multi_hop_query(query_type="3p", query=("http://
↵www.benchmark.org/family#F9M167",
                                                         ("http://
↵www.benchmark.org/family#hasSibling",
                                                         "http://
↵www.benchmark.org/family#married",
                                                         "http://
↵www.w3.org/1999/02/22-rdf-syntax-ns#type")),
                                                       tnorm="min", k=5)
top_entities = [topk_entity for topk_entity, query_score in predictions]
print(top_entities)
assert "http://www.benchmark.org/family#Person" in top_entities
assert "http://www.benchmark.org/family#Father" in top_entities
assert "http://www.benchmark.org/family#Male" in top_entities
```

For more, please refer to examples/multi_hop_query_answering.

8 Predicting Missing Links

```
from dicee import KGE
# (1) Train a knowledge graph embedding model..
# (2) Load a pretrained model
pre_trained_kge = KGE(path='../')
# (3) Predict missing links through head entity rankings
pre_trained_kge.predict_topk(h=[".."],r=[".."],topk=10)
# (4) Predict missing links through relation rankings
pre_trained_kge.predict_topk(h=[".."],t=[".."],topk=10)
# (5) Predict missing links through tail entity rankings
pre_trained_kge.predict_topk(r=[".."],t=[".."],topk=10)
```

9 Downloading Pretrained Models

```
from dicee import KGE
# (1) Load a pretrained ConEx on DBpedia
model = KGE(url="https://files.dice-research.org/projects/DiceEmbeddings/KINSHIP-Keci-
↳dim128-epoch256-KvsAll")
```

- For more please look at dice-research.org/projects/DiceEmbeddings/¹¹

10 How to Deploy

```
from dicee import KGE
KGE(path='...').deploy(share=True, top_k=10)
```

11 Docker

To build the Docker image:

```
docker build -t dice-embeddings .
```

To test the Docker image:

```
docker run --rm -v ~/.local/share/dicee/KGs:/dicee/KGs dice-embeddings ./main.py --
↳model AConEx --embedding_dim 16
```

12 Coverage Report

The coverage report is generated using `coverage.py`¹²:

Name	Stmts	Miss	Cover	Missing
-----	-----	-----	-----	-----
dicee/___init___ .py	7	0	100%	
dicee/abstracts.py	338	115	66%	112-113, 115

(continues on next page)

¹¹ <https://files.dice-research.org/projects/DiceEmbeddings/>

¹² <https://coverage.readthedocs.io/en/7.6.0/>

(continued from previous page)

↪131, 154-155, 160, 173, 197, 240-254, 290, 303-306, 309-313, 353-364, 379-387, 402, ↪				
↪413-417, 427-428, 434-436, 442-445, 448-453, 576-596, 602-606, 610-612, 631, 658-696				
dicee/callbacks.py	248	103	58%	50-55, ↪
↪67-73, 76, 88-93, 98-103, 106-109, 116-133, 138-142, 146-147, 247, 281-285, 291-292, ↪				
↪ 310-316, 319, 324-325, 337-343, 349-358, 363-365, 410, 421-434, 438-473, 485-491				
dicee/config.py	97	2	98%	146-147
dicee/dataset_classes.py	430	146	66%	16, 44, ↪
↪57, 89-98, 104, 111-118, 121, 124, 127-151, 207-213, 216, 219-221, 324, 335-338, ↪				
↪354, 420-421, 439, 562-581, 583, 587-599, 606-615, 618, 622-636, 780-787, 790-794, ↪				
↪845, 866-878, 902-915, 937, 941-954, 964-967, 973, 985, 987, 989, 1012-1022				
dicee/eval_static_funcs.py	256	100	61%	104, 109, ↪
↪ 114, 261-356, 363-414, 442, 465-468				
dicee/evaluator.py	267	48	82%	48, 53, ↪
↪58, 77, 82-83, 86, 102, 119, 130, 134, 139, 173-184, 191-202, 310, 340-358, 452, ↪				
↪462, 480-485				
dicee/executer.py	134	16	88%	53-57, ↪
↪166-176, 235-236, 283				
dicee/knowledge_graph.py	82	10	88%	84, 94- ↪
↪95, 124, 128, 132-134, 137-138, 140				
dicee/knowledge_graph_embeddings.py	654	415	37%	25, 28- ↪
↪29, 37-50, 55-88, 91-125, 129-137, 171, 173-229, 261, 265, 276-277, 301-303, 311, ↪				
↪339-362, 493, 497-519, 523-547, 580, 656, 665, 710-716, 748, 806-1171, 1202-1263, ↪				
↪1267-1295, 1326, 1332				
dicee/models/___init___py	9	0	100%	
dicee/models/adopt.py	187	172	8%	50-86, ↪
↪99-110, 129-185, 195-242, 266-322, 346-448, 484-517				
dicee/models/base_model.py	240	35	85%	30-35, ↪
↪64, 66, 92, 99-116, 171, 204, 244, 250, 259, 262, 266, 273, 277, 279, 294, 307-308, ↪				
↪362, 365, 438, 450				
dicee/models/clifford.py	470	278	41%	10, 12, ↪
↪16, 24-25, 52-56, 79-87, 101-103, 108-109, 140-160, 184, 191, 195-256, 273-277, 289, ↪				
↪ 292, 297, 302, 346-361, 377-444, 464-470, 483, 486, 491, 496, 525-531, 544, 547, ↪				
↪552, 557, 567-576, 592-593, 613-685, 696-699, 724-749, 773-806, 842-846, 859, 869, ↪				
↪872, 877, 882, 887, 891, 895, 904-905, 935, 942, 947, 975-979, 1007-1016, 1026-1034, ↪				
↪ 1052-1054, 1072-1074, 1090-1092				
dicee/models/complex.py	162	25	85%	86-109, ↪
↪273-287				
dicee/models/dualE.py	59	10	83%	93-102, ↪
↪142-156				
dicee/models/ensemble.py	89	67	25%	7-29, 31, ↪
↪ 34, 37, 40, 43, 46, 49, 52-54, 56-58, 64-68, 71-90, 93-94, 97-112, 131				
dicee/models/function_space.py	262	221	16%	10-23, ↪
↪27-36, 39-48, 52-69, 76-87, 90-99, 102-111, 115-127, 135-157, 160-166, 169-186, 189- ↪				
↪195, 198-206, 209, 214-235, 244-247, 251-255, 259-268, 272-293, 302-308, 312-329, ↪				
↪333-336, 345-353, 356, 367-373, 393-407, 425-439, 444-454, 462-466, 475-479				
dicee/models/literal.py	33	1	97%	82
dicee/models/octonion.py	227	83	63%	21-44, ↪
↪320-329, 334-345, 348-370, 374-416, 426-474				
dicee/models/pykeen_models.py	55	5	91%	77-80, ↪
↪135				
dicee/models/quaternion.py	192	69	64%	7-21, 30- ↪
↪55, 68-72, 107, 185, 328-342, 345-364, 368-389, 399-426				

(continues on next page)

(continued from previous page)

dicee/models/real.py	61	12	80%	37-42, ↵
↪70-73, 91, 107-110				
dicee/models/static_funcs.py	10	0	100%	
dicee/models/transformers.py	234	189	19%	20-39, ↵
↪42, 56-71, 80-98, 101-112, 119-121, 124, 130-147, 151-176, 182-186, 189-193, 199-				
↪203, 206-208, 225-252, 261-264, 267-272, 275-300, 306-311, 315-368, 372-394, 400-410				
dicee/query_generator.py	374	346	7%	17-51, ↵
↪55, 61-64, 68-69, 77-91, 99-146, 154-187, 191-205, 211-268, 273-302, 306-442, 452-				
↪471, 479-502, 509-513, 518, 523-529				
dicee/read_preprocess_save_load_kg/___init___py	3	0	100%	
dicee/read_preprocess_save_load_kg/preprocess.py	243	40	84%	33, 39, ↵
↪76, 100-125, 131, 136-149, 175, 205, 380-381				
dicee/read_preprocess_save_load_kg/read_from_disk.py	36	11	69%	34, 38-
↪40, 47, 55, 58-72				
dicee/read_preprocess_save_load_kg/save_load_disk.py	53	21	60%	29-30, ↵
↪38, 47-68				
dicee/read_preprocess_save_load_kg/util.py	236	125	47%	159, 173-
↪175, 179-180, 198-204, 207-209, 214-216, 230, 244-247, 252-260, 265-271, 276-281, ↵				
↪286-291, 303-324, 330-386, 390-394, 398-399, 403, 407-408, 436, 441, 448-449				
dicee/sanity_checkers.py	47	19	60%	8-12, 21-
↪31, 46, 51, 58, 69-79				
dicee/static_funcs.py	483	194	60%	42, 52, ↵
↪58-63, 85, 92-96, 109-119, 129-131, 136, 143, 167, 172, 184, 190, 198, 202, 229-233,				
↪295, 303-309, 320-330, 341-361, 389, 413-414, 419-420, 437-438, 440-441, 443-444, ↵				
↪452, 470-474, 491-494, 498-503, 507-511, 515-516, 522-524, 539-553, 558-561, 566-				
↪569, 578-629, 634-646, 663-680, 683-691, 695-713, 724				
dicee/static_funcs_training.py	155	66	57%	7-10, ↵
↪222-319, 327-328				
dicee/static_preprocess_funcs.py	98	43	56%	17-25, ↵
↪50, 57, 59, 70, 83-107, 112-115, 120-123, 128-131				
dicee/trainer/___init___py	1	0	100%	
dicee/trainer/dice_trainer.py	151	18	88%	22, 30-
↪31, 33-35, 97, 104, 109-114, 152, 237, 280-283				
dicee/trainer/model_parallelism.py	99	87	12%	10-25, ↵
↪30-116, 121-132, 136, 141-197				
dicee/trainer/torch_trainer.py	77	6	92%	31, 102, ↵
↪168, 179-181				
dicee/trainer/torch_trainer_ddp.py	89	71	20%	11-14, ↵
↪43, 47-67, 78-94, 113-122, 126-136, 151-158, 168-191				

TOTAL	6948	3169	54%	

13 How to cite

Currently, we are working on our manuscript describing our framework. If you really like our work and want to cite it now, feel free to chose one :)

```
# Keci
@inproceedings{demir2023clifford,
  title={Clifford Embeddings--A Generalized Approach for Embedding in Normed Algebras}
  ↪,
```

(continues on next page)

```

    author={Demir, Caglar and Ngonga Ngomo, Axel-Cyrille},
    booktitle={Joint European Conference on Machine Learning and Knowledge Discovery in
↪Databases},
    pages={567--582},
    year={2023},
    organization={Springer}
}
# LitCQD
@inproceedings{demir2023litcq,
    title={LitCQD: Multi-Hop Reasoning in Incomplete Knowledge Graphs with Numeric
↪Literals},
    author={Demir, Caglar and Wiebesiek, Michel and Lu, Renzhong and Ngonga Ngomo, Axel-
↪Cyrille and Heindorf, Stefan},
    booktitle={Joint European Conference on Machine Learning and Knowledge Discovery in
↪Databases},
    pages={617--633},
    year={2023},
    organization={Springer}
}
# DICE Embedding Framework
@article{demir2022hardware,
    title={Hardware-agnostic computation for large-scale knowledge graph embeddings},
    author={Demir, Caglar and Ngomo, Axel-Cyrille Ngonga},
    journal={Software Impacts},
    year={2022},
    publisher={Elsevier}
}
# KronE
@inproceedings{demir2022kronecker,
    title={Kronecker decomposition for knowledge graph embeddings},
    author={Demir, Caglar and Lienen, Julian and Ngonga Ngomo, Axel-Cyrille},
    booktitle={Proceedings of the 33rd ACM Conference on Hypertext and Social Media},
    pages={1--10},
    year={2022}
}
# QMult, OMult, ConvQ, ConvO
@InProceedings{pmlr-v157-demir21a,
    title = {Convolutional Hypercomplex Embeddings for Link Prediction},
    author = {Demir, Caglar and Moussallem, Diego and Heindorf, Stefan and Ngonga
↪Ngomo, Axel-Cyrille},
    booktitle = {Proceedings of The 13th Asian Conference on Machine Learning},
    pages = {656--671},
    year = {2021},
    editor = {Balasubramanian, Vineeth N. and Tsang, Ivor},
    volume = {157},
    series = {Proceedings of Machine Learning Research},
    month = {17--19 Nov},
    publisher = {PMLR},
    pdf = {https://proceedings.mlr.press/v157/demir21a/demir21a.pdf},
    url = {https://proceedings.mlr.press/v157/demir21a.html},
}
# ConEx

```

```
@inproceedings{demir2021convolutional,
  title={Convolutional Complex Knowledge Graph Embeddings},
  author={Caglar Demir and Axel-Cyrille Ngonga Ngomo},
  booktitle={Eighteenth Extended Semantic Web Conference - Research Track},
  year={2021},
  url={https://openreview.net/forum?id=6T45-4TFqaX}}
# Shallom
@inproceedings{demir2021shallow,
  title={A shallow neural model for relation prediction},
  author={Demir, Caglar and Moussallem, Diego and Ngomo, Axel-Cyrille Ngonga},
  booktitle={2021 IEEE 15th International Conference on Semantic Computing (ICSC)},
  pages={179--182},
  year={2021},
  organization={IEEE}}
```

For any questions or wishes, please contact: caglar.demir@upb.de

14 dicee

DICE Embeddings - Knowledge Graph Embedding Library.

A library for training and using knowledge graph embedding models with support for various scoring techniques and training strategies.

Submodules:

evaluation: Model evaluation functions and Evaluator class models: KGE model implementations trainer: Training orchestration scripts: Utility scripts

14.1 Submodules

`dicee.__main__`

`dicee.abstracts`

Classes

<i>AbstractTrainer</i>	Abstract class for Trainer class for knowledge graph embedding models
<i>BaseInteractiveKGE</i>	Abstract/base class for using knowledge graph embedding models interactively.
<i>InteractiveQueryDecomposition</i>	
<i>AbstractCallback</i>	Abstract class for Callback class for knowledge graph embedding models
<i>AbstractPPECallback</i>	Abstract class for Callback class for knowledge graph embedding models
<i>BaseInteractiveTrainKGE</i>	Abstract/base class for training knowledge graph embedding models interactively.

Module Contents

class `dicee.abstracts.AbstractTrainer` (*args, callbacks*)

Abstract class for Trainer class for knowledge graph embedding models

Parameter

args

[str] ?

callbacks: list

?

attributes

callbacks

is_global_zero = True

global_rank = 0

local_rank = 0

strategy = None

on_fit_start (**args, **kwargs*)

A function to call callbacks before the training starts.

Parameter

args

kwargs

rtype

None

on_fit_end (**args, **kwargs*)

A function to call callbacks at the end of the training.

Parameter

args

kwargs

rtype

None

on_train_epoch_start (**args, **kwargs*)

A function to call callbacks at the start of an epoch.

Parameter

args

kwargs

rtype

None

on_train_epoch_end (*args, **kwargs)

A function to call callbacks at the end of an epoch.

Parameter

args

kwargs

rtype

None

on_train_batch_end (*args, **kwargs)

A function to call callbacks at the end of each mini-batch during training.

Parameter

args

kwargs

rtype

None

static save_checkpoint (full_path: str, model) → None

A static function to save a model into disk

Parameter

full_path : str

model:

rtype

None

class dicee.abstracts.**BaseInteractiveKGE** (path: str = None, url: str = None,
construct_ensemble: bool = False, model_name: str = None,
apply_semantic_constraint: bool = False)

Abstract/base class for using knowledge graph embedding models interactively.

Parameter

path_of_pretrained_model_dir

[str] ?

construct_ensemble: boolean

?

model_name: str apply_semantic_constraint : boolean

construct_ensemble = False

apply_semantic_constraint = False

configs

get_eval_report() → dict

get_bpe_token_representation (*str_entity_or_relation: List[str] | str*) → List[List[int]] | List[int]

Parameters

str_entity_or_relation (*corresponds to a str or a list of strings to be tokenized via BPE and shaped.*)

Return type

A list integer(s) or a list of lists containing integer(s)

get_padded_bpe_triple_representation (*triples: List[List[str]]*) → Tuple[List, List, List]

Parameters

triples

set_model_train_mode() → None

Setting the model into training mode

Parameter

set_model_eval_mode() → None

Setting the model into eval mode

Parameter

property name

sample_entity (*n: int*) → List[str]

sample_relation (*n: int*) → List[str]

is_seen (*entity: str = None, relation: str = None*) → bool

save() → None

get_entity_index (*x: str*)

get_relation_index (*x: str*)

index_triple (*head_entity: List[str], relation: List[str], tail_entity: List[str]*)
→ Tuple[torch.LongTensor, torch.LongTensor, torch.LongTensor]

Index Triple

Parameter

head_entity: List[str]

String representation of selected entities.

relation: List[str]

String representation of selected relations.

tail_entity: List[str]

String representation of selected entities.

Returns: Tuple

pytorch tensor of triple score

add_new_entity_embeddings (entity_name: str = None, embeddings: torch.FloatTensor = None)

get_entity_embeddings (items: List[str])

Return embedding of an entity given its string representation

Parameter

items:

entities

get_relation_embeddings (items: List[str])

Return embedding of a relation given its string representation

Parameter

items:

relations

construct_input_and_output (head_entity: List[str], relation: List[str], tail_entity: List[str], labels)

Construct a data point :param head_entity: :param relation: :param tail_entity: :param labels: :return:

parameters ()

class dicee.abstracts.**InteractiveQueryDecomposition**

t_norm (tens_1: torch.Tensor, tens_2: torch.Tensor, tnorm: str = 'min') → torch.Tensor

tensor_t_norm (subquery_scores: torch.FloatTensor, tnorm: str = 'min') → torch.FloatTensor

Compute T-norm over $[0,1]^{n \times d}$ where n denotes the number of hops and d denotes number of entities

t_conorm (tens_1: torch.Tensor, tens_2: torch.Tensor, tconorm: str = 'min') → torch.Tensor

negnorm (tens_1: torch.Tensor, lambda_: float, neg_norm: str = 'standard') → torch.Tensor

class dicee.abstracts.**AbstractCallback**

Bases: abc.ABC, lightning.pytorch.callbacks.Callback

Abstract class for Callback class for knowledge graph embedding models

Parameter

on_init_start (*args, **kwargs)

Parameter

trainer:

model:

rtype

None

on_init_end (**args, **kwargs*)

Call at the beginning of the training.

Parameter

trainer:

model:

rtype

None

on_fit_start (*trainer, model*)

Call at the beginning of the training.

Parameter

trainer:

model:

rtype

None

on_train_epoch_end (*trainer, model*)

Call at the end of each epoch during training.

Parameter

trainer:

model:

rtype

None

on_train_batch_end (**args, **kwargs*)

Call at the end of each mini-batch during the training.

Parameter

trainer:

model:

rtype

None

on_fit_end (**args, **kwargs*)

Call at the end of the training.

Parameter

trainer:

model:

rtype

None

```
class dicee.abstracts.AbstractPPECallback (num_epochs, path, epoch_to_start,  
last_percent_to_consider)
```

Bases: [AbstractCallback](#)

Abstract class for Callback class for knowledge graph embedding models

Parameter

num_epochs

path

sample_counter = 0

epoch_count = 0

alphas = None

on_fit_start (trainer, model)

Call at the beginning of the training.

Parameter

trainer:

model:

rtype

None

on_fit_end (trainer, model)

Call at the end of the training.

Parameter

trainer:

model:

rtype

None

store_ensemble (param_ensemble) → None

```
class dicee.abstracts.BaseInteractiveTrainKGE
```

Abstract/base class for training knowledge graph embedding models interactively. This class provides methods for re-training KGE models and also Literal Embedding model.

train_triples (h: List[str], r: List[str], t: List[str], labels: List[float], iteration=2, optimizer=None)

```
train_k_vs_all (h, r, iteration=1, lr=0.001)
```

Train k vs all :param head_entity: :param relation: :param iteration: :param lr: :return:

```
train (kg, lr=0.1, epoch=10, batch_size=32, neg_sample_ratio=10, num_workers=1) → None
```

Retrained a pretrain model on an input KG via negative sampling.

```
train_literals (train_file_path: str = None, num_epochs: int = 100, lit_lr: float = 0.001,  
lit_normalization_type: str = 'z-norm', batch_size: int = 1024, sampling_ratio: float = None,  
random_seed=1, loader_backend: str = 'pandas', freeze_entity_embeddings: bool = True,  
gate_residual: bool = True, device: str = None, suffle_data: bool = True)
```

Trains the Literal Embeddings model using literal data.

Parameters

- **train_file_path** (*str*) – Path to the training data file.
- **num_epochs** (*int*) – Number of training epochs.
- **lit_lr** (*float*) – Learning rate for the literal model.
- **norm_type** (*str*) – Normalization type to use ('z-norm', 'min-max', or None).
- **batch_size** (*int*) – Batch size for training.
- **sampling_ratio** (*float*) – Ratio of training triples to use.
- **loader_backend** (*str*) – Backend for loading the dataset ('pandas' or 'rdflib').
- **freeze_entity_embeddings** (*bool*) – If True, freeze the entity embeddings during training.
- **gate_residual** (*bool*) – If True, use gate residual connections in the model.
- **device** (*str*) – Device to use for training ('cuda' or 'cpu'). If None, will use available GPU or CPU.
- **suffle_data** (*bool*) – If True, shuffle the dataset before training.

dicee.analyse_experiments

This script should be moved to dicee/scripts Example: python dicee/analyse_experiments.py --dir Experiments --features "model" "trainMRR" "testMRR"

Classes

Experiment

Functions

```
get_default_arguments()  
analyse(args)
```

Module Contents

```
dicee.analyse_experiments.get_default_arguments()
```

```
class dicee.analyse_experiments.Experiment
```

```

model_name = []
callbacks = []
embedding_dim = []
num_params = []
num_epochs = []
batch_size = []
lr = []
byte_pair_encoding = []
aswa = []
path_dataset_folder = []
full_storage_path = []
pq = []
train_mrr = []
train_h1 = []
train_h3 = []
train_h10 = []
val_mrr = []
val_h1 = []
val_h3 = []
val_h10 = []
test_mrr = []
test_h1 = []
test_h3 = []
test_h10 = []
runtime = []
normalization = []
scoring_technique = []
save_experiment(x)
to_df()

dicee.analyse_experiments.analyse(args)

```

dicee.callbacks

Callbacks for training lifecycle events.

Provides callback classes for various training events including epoch end, model saving, weight averaging, and evaluation.

Classes

<i>AccumulateEpochLossCallback</i>	Callback to store epoch losses to a CSV file.
<i>PrintCallback</i>	Abstract class for Callback class for knowledge graph embedding models
<i>KGESaveCallback</i>	Abstract class for Callback class for knowledge graph embedding models
<i>PseudoLabellingCallback</i>	Abstract class for Callback class for knowledge graph embedding models
<i>Eval</i>	Abstract class for Callback class for knowledge graph embedding models
<i>KronE</i>	Abstract class for Callback class for knowledge graph embedding models
<i>Perturb</i>	A callback for a three-Level Perturbation
<i>PeriodicEvalCallback</i>	Callback to periodically evaluate the model and optionally save checkpoints during training.
<i>LRScheduler</i>	Callback for managing learning rate scheduling and model snapshots.

Functions

<i>estimate_q</i> (eps)	estimate rate of convergence q from sequence esp
<i>compute_convergence</i> (seq, i)	

Module Contents

```
class dicee.callbacks.AccumulateEpochLossCallback (path: str)
```

Bases: *dicee.abstracts.AbstractCallback*

Callback to store epoch losses to a CSV file.

Parameters

path – Directory path where the loss file will be saved.

path

on_fit_end (trainer, model) → None

Store epoch loss history to CSV file.

Parameters

- **trainer** – The trainer instance.
- **model** – The model being trained.

```
class dicee.callbacks.PrintCallback
```

Bases: *dicee.abstracts.AbstractCallback*

Abstract class for Callback class for knowledge graph embedding models

Parameter

start_time

on_fit_start (*trainer, pl_module*)

Call at the beginning of the training.

Parameter

trainer:

model:

rtype

None

on_fit_end (*trainer, pl_module*)

Call at the end of the training.

Parameter

trainer:

model:

rtype

None

on_train_batch_end (**args, **kwargs*)

Call at the end of each mini-batch during the training.

Parameter

trainer:

model:

rtype

None

on_train_epoch_end (**args, **kwargs*)

Call at the end of each epoch during training.

Parameter

trainer:

model:

rtype

None

class dicee.callbacks.KGESaveCallback (*every_x_epoch: int, max_epochs: int, path: str*)

Bases: *dicee.abstracts.AbstractCallback*

Abstract class for Callback class for knowledge graph embedding models

Parameter

every_x_epoch

max_epochs

epoch_counter = 0

path

on_train_batch_end (**args, **kwargs*)

Call at the end of each mini-batch during the training.

Parameter

trainer:

model:

rtype

None

on_fit_start (*trainer, pl_module*)

Call at the beginning of the training.

Parameter

trainer:

model:

rtype

None

on_train_epoch_end (**args, **kwargs*)

Call at the end of each epoch during training.

Parameter

trainer:

model:

rtype

None

on_fit_end (**args, **kwargs*)

Call at the end of the training.

Parameter

trainer:

model:

rtype

None

on_epoch_end (*model, trainer, **kwargs*)

```
class dicee.callbacks.PseudoLabellingCallback(data_module, kg, batch_size)
```

Bases: *dicee.abstracts.AbstractCallback*

Abstract class for Callback class for knowledge graph embedding models

Parameter

data_module

kg

num_of_epochs = 0

unlabelled_size

batch_size

create_random_data()

on_epoch_end(*trainer, model*)

```
dicee.callbacks.estimate_q(eps)
```

estimate rate of convergence q from sequence esp

```
dicee.callbacks.compute_convergence(seq, i)
```

```
class dicee.callbacks.Eval(path, epoch_ratio: int = None)
```

Bases: *dicee.abstracts.AbstractCallback*

Abstract class for Callback class for knowledge graph embedding models

Parameter

path

reports = []

epoch_ratio = None

epoch_counter = 0

on_fit_start(*trainer, model*)

Call at the beginning of the training.

Parameter

trainer:

model:

rtype

None

on_fit_end(*trainer, model*)

Call at the end of the training.

Parameter

trainer:

model:

rtype

None

on_train_epoch_end (*trainer, model*)

Call at the end of each epoch during training.

Parameter

trainer:

model:

rtype

None

on_train_batch_end (**args, **kwargs*)

Call at the end of each mini-batch during the training.

Parameter

trainer:

model:

rtype

None

class dicee.callbacks.**KronE**

Bases: *dicee.abstracts.AbstractCallback*

Abstract class for Callback class for knowledge graph embedding models

Parameter

f = None

static batch_kronecker_product (*a, b*)

Kronecker product of matrices a and b with leading batch dimensions. Batch dimensions are broadcast. The number of them must be the same. :type a: torch.Tensor :type b: torch.Tensor :rtype: torch.Tensor

get_kronecker_triple_representation (*indexed_triple: torch.LongTensor*)

Get kronecker embeddings

on_fit_start (*trainer, model*)

Call at the beginning of the training.

Parameter

trainer:

model:

rtype

None

```
class dicee.callbacks.Perturb (level: str = 'input', ratio: float = 0.0, method: str = None,  

                             scaler: float = None, frequency=None)
```

Bases: `dicee.abstracts.AbstractCallback`

A callback for a three-Level Perturbation

Input Perturbation: During training an input *x* is perturbed by randomly replacing its element. In the context of knowledge graph embedding models, *x* can denote a triple, a tuple of an entity and a relation, or a tuple of two entities. A perturbation means that a component of *x* is randomly replaced by an entity or a relation.

Parameter Perturbation:

Output Perturbation:

```
level = 'input'
```

```
ratio = 0.0
```

```
method = None
```

```
scaler = None
```

```
frequency = None
```

```
on_train_batch_start (trainer, model, batch, batch_idx)
```

Called when the train batch begins.

```
class dicee.callbacks.PeriodicEvalCallback (experiment_path: str, max_epochs: int,  

                                           eval_every_n_epoch: int = 0, eval_at_epochs: list = None,  

                                           save_model_every_n_epoch: bool = True, n_epochs_eval_model: str = 'val_test')
```

Bases: `dicee.abstracts.AbstractCallback`

Callback to periodically evaluate the model and optionally save checkpoints during training.

Evaluates at regular intervals (every *N* epochs) or at explicitly specified epochs. Stores evaluation reports and model checkpoints.

```
experiment_dir
```

```
max_epochs
```

```
epoch_counter = 0
```

```
save_model_every_n_epoch = True
```

```
reports
```

```
n_epochs_eval_model = 'val_test'
```

```
default_eval_model = None
```

```
eval_epochs
```

```
on_fit_end (trainer, model)
```

Called at the end of training. Saves final evaluation report.

```
on_train_epoch_end (trainer, model)
```

Called at the end of each training epoch. Performs evaluation and checkpointing if scheduled.

```

class dicee.callbacks.LRScheduler(adaptive_lr_config: dict, total_epochs: int, experiment_dir: str,
                                eta_max: float = 0.1, snapshot_dir: str = 'snapshots')
    Bases: dicee.abstracts.AbstractCallback

    Callback for managing learning rate scheduling and model snapshots.

    Supports cosine annealing (“cca”), MMCCLR (“mmcclr”), and their deferred (warmup) variants: - “deferred_cca”
    - “deferred_mmcclr”

    At the end of each learning rate cycle, the model can optionally be saved as a snapshot.

    total_epochs

    experiment_dir

    snapshot_dir

    batches_per_epoch = None

    total_steps = None

    cycle_length = None

    warmup_steps = None

    lr_lambda = None

    scheduler = None

    step_count = 0

    snapshot_loss

    on_train_start(trainer, model)
        Initialize training parameters and LR scheduler at start of training.

    on_train_batch_end(trainer, model, outputs, batch, batch_idx)
        Step the LR scheduler and save model snapshot if needed after each batch.

    on_fit_end(trainer, model)
        Call at the end of the training.

    Parameter

    trainer:

    model:

    rtype
        None

```

`dicee.config`

Configuration module for DICE embeddings.

Provides the Namespace class with default configuration values for training knowledge graph embedding models.

Classes

<i>Namespace</i>	Extended Namespace with default KGE training configuration.
------------------	---

Module Contents

```
class dicee.config.Namespace (**kwargs)
```

Bases: argparse.Namespace

Extended Namespace with default KGE training configuration.

Provides sensible defaults for all training parameters while allowing easy customization through command-line arguments or direct assignment.

dataset_dir: str = None

The path of a folder containing train.txt, and/or valid.txt and/or test.txt

save_embeddings_as_csv: bool = False

Embeddings of entities and relations are stored into CSV files to facilitate easy usage.

storage_path: str = 'Experiments'

A directory named with time of execution under `storage_path` that contains related data about embeddings.

path_to_store_single_run: str = None

A single directory created that contains related data about embeddings.

path_single_kg = None

Path of a file corresponding to the input knowledge graph

sparql_endpoint = None

An endpoint of a triple store.

model: str = 'Keci'

KGE model

optim: str = 'Adam'

Optimizer

embedding_dim: int = 64

Size of continuous vector representation of an entity/relation

num_epochs: int = 150

Number of pass over the training data

batch_size: int = 1024

Mini-batch size if it is None, an automatic batch finder technique applied

lr: float = 0.1

Learning rate

add_noise_rate: float = None

The ratio of added random triples into training dataset

gpus = None

Number GPUs to be used during training

callbacks
 10}}

Type
 Callbacks, e.g., {"PPE"}

Type
 {"last_percent_to_consider"}

backend: str = 'pandas'
 Backend to read, process, and index input knowledge graph. pandas, polars and rdflib available

separator: str = '\\s+'
 separator for extracting head, relation and tail from a triple

trainer: str = 'torchCPUTrainer'
 Trainer for knowledge graph embedding model

scoring_technique: str = 'KvsAll'
 Scoring technique for knowledge graph embedding models

neg_ratio: int = 0
 Negative ratio for a true triple in NegSample training_technique

weight_decay: float = 0.0
 Weight decay for all trainable params

normalization: str = 'None'
 LayerNorm, BatchNorm1d, or None

init_param: str = None
 xavier_normal or None

gradient_accumulation_steps: int = 0
 Not tested e

num_folds_for_cv: int = 0
 Number of folds for CV

eval_model: str = 'train_val_test'
 ["None", "train", "train_val", "train_val_test", "test"]

Type
 Evaluate trained model choices

save_model_at_every_epoch: int = None
 Not tested

label_smoothing_rate: float = 0.0

num_core: int = 0
 Number of CPUs to be used in the mini-batch loading process

random_seed: int = 0
 Random Seed

sample_triples_ratio: float = None
 Read some triples that are uniformly at random sampled. Ratio being between 0 and 1

read_only_few: int = None
 Read only first few triples

pykeen_model_kwargs
 Additional keyword arguments for pykeen models

kernel_size: int = 3
 Size of a square kernel in a convolution operation

num_of_output_channels: int = 32
 Number of slices in the generated feature map by convolution.

p: int = 0
 P parameter of Clifford Embeddings

q: int = 1
 Q parameter of Clifford Embeddings

input_dropout_rate: float = 0.0
 Dropout rate on embeddings of input triples

hidden_dropout_rate: float = 0.0
 Dropout rate on hidden representations of input triples

feature_map_dropout_rate: float = 0.0
 Dropout rate on a feature map generated by a convolution operation

byte_pair_encoding: bool = False
 Byte pair encoding

Type
 WIP

adaptive_swa: bool = False
 Adaptive stochastic weight averaging

swa: bool = False
 Stochastic weight averaging

swag: bool = False
 Stochastic weight averaging - Gaussian

ema: bool = False
 Exponential Moving Average

twa: bool = False
 Trainable weight averaging

block_size: int = None
 block size of LLM

continual_learning = None
 Path of a pretrained model size of LLM

auto_batch_finding = False
 A flag for using auto batch finding

eval_every_n_epochs: int = 0
 Evaluate model every n epochs. If 0, no evaluation is applied.

save_every_n_epochs: bool = False
 Save model every n epochs. If True, save model at every epoch.

eval_at_epochs: list = None
 List of epoch numbers at which to evaluate the model (e.g., 1 5 10).

n_epochs_eval_model: str = 'val_test'
 Evaluating link prediction performance on data splits while performing periodic evaluation.

adaptive_lr
 “cca”}

Type
 Adaptive learning rate parameters, e.g., {“scheduler_name”

swa_start_epoch: int = None
 Epoch at which to start applying stochastic weight averaging.

swa_c_epochs: int = 1
 Number of epochs to average over for SWA, SWAG, EMA, TWA.

__iter__()

dicee.dataset_classes

Dataset classes for knowledge graph embedding training.

This package groups the various `torch.utils.data.Dataset` implementations used throughout DICE Embeddings into thematic sub-modules:

- `_bpe` – Byte-pair-encoding related datasets.
- `_negative_sampling` – Negative sampling based datasets.
- `_label_based` – Multi-label / multi-class scoring datasets.
- `_literal` – Literal (numeric) embedding dataset.
- `_factory` – `construct_dataset` / `reload_dataset` helpers.

All public names are re-exported here so that existing `from dicee.dataset_classes import ...` statements continue to work unchanged.

Classes

<i>BPE_NegativeSamplingDataset</i>	Dataset for negative sampling with byte-pair encoded triples.
<i>MultiClassClassificationDataset</i>	Dataset for autoregressive multi-class classification on sub-word units.
<i>MultiLabelDataset</i>	Multi-label dataset for BPE-based KvsAll / AllvsAll training.
<i>AllvsAll</i>	Dataset for AllvsAll training (multi-label, exhaustive).
<i>KvsAll</i>	Dataset for KvsAll training (multi-label).
<i>KvsSampleDataset</i>	Dataset for KvsSample training (dynamic multi-label).
<i>OnevsAllDataset</i>	Dataset for the 1-vs-All training strategy (multi-class).
<i>FixedNegSampleDataset</i>	Pre-computed (fixed) negative sampling dataset.

continues on next page

Table 7 – continued from previous page

<i>OnevsSample</i>	Dataset for 1-vs-Sample training (dynamic multi-class with negatives).
<i>TriplePredictionDataset</i>	Dataset for triple prediction with on-the-fly negative sampling.
<i>LiteralDataset</i>	Dataset for loading and processing literal data for Literal Embedding models.

Functions

<i>construct_dataset</i> ($\rightarrow$ torch.utils.data.Dataset)	Build the appropriate dataset for the given training configuration.
<i>reload_dataset</i> (path, form_of_labelling, ...)	Reload training data from disk and construct a Pytorch dataset.

Package Contents

class dicee.dataset_classes.**BPE_NegativeSamplingDataset** (*train_set: torch.LongTensor*,
ordered_shaped_bpe_entities: torch.LongTensor, *neg_ratio: int*)

Bases: torch.utils.data.Dataset

Dataset for negative sampling with byte-pair encoded triples.

Each sample is a BPE-encoded triple. The custom `collate_fn` constructs negatives by corrupting head or tail entities with random BPE entities.

Parameters

- **train_set** (*torch.LongTensor*) – Integer-encoded triples of shape $(N, 3)$.
- **ordered_shaped_bpe_entities** (*torch.LongTensor*) – All BPE entity representations, ordered by entity index.
- **neg_ratio** (*int*) – Number of negative samples per positive triple.

train_set

ordered_bpe_entities

num_bpe_entities

neg_ratio

num_datapoints

__len__()

__getitem__ (*idx*)

collate_fn (*batch_shaped_bpe_triples: List[Tuple[torch.Tensor, torch.Tensor]]*)

class dicee.dataset_classes.**MultiClassClassificationDataset** (
subword_units: numpy.ndarray, *block_size: int = 8*)

Bases: torch.utils.data.Dataset

Dataset for autoregressive multi-class classification on sub-word units.

Splits a flat sequence of sub-word token ids into overlapping windows of size `block_size` for next-token prediction.

Parameters

- **subword_units** (*numpy.ndarray*) – 1-D array of sub-word token ids.
- **block_size** (*int, optional*) – Context window length (default 8).

`train_data`

`block_size = 8`

`num_of_data_points`

`collate_fn = None`

`__len__()`

`__getitem__(idx)`

```
class dicee.dataset_classes.MultiLabelDataset (train_set: torch.LongTensor,  
        train_indices_target: torch.LongTensor, target_dim: int,  
        torch_ordered_shaped_bpe_entities: torch.LongTensor)
```

Bases: `torch.utils.data.Dataset`

Multi-label dataset for BPE-based KvsAll / AllvsAll training.

Each sample is a BPE-encoded (head, relation) pair together with a binary multi-label target vector over all entities.

Parameters

- **train_set** (*torch.LongTensor*) – BPE-encoded input pairs of shape (N, 2, token_length).
- **train_indices_target** (*torch.LongTensor*) – Per-sample lists of positive target entity indices.
- **target_dim** (*int*) – Dimensionality of the target vector (number of entities).
- **torch_ordered_shaped_bpe_entities** (*torch.LongTensor*) – Ordered BPE entity representations.

`train_set`

`train_indices_target`

`target_dim`

`num_datapoints`

`torch_ordered_shaped_bpe_entities`

`collate_fn = None`

`__len__()`

`__getitem__(idx)`

```
class dicee.dataset_classes.AllvsAll (train_set_idx: numpy.ndarray, entity_idxes, relation_idxes,  
    label_smoothing_rate=0.0)
```

Bases: torch.utils.data.Dataset

Dataset for AllvsAll training (multi-label, exhaustive).

Extends the KvsAll idea: every *possible* (entity, relation) combination is included — not just those observed in the KG. Pairs without any known tail entities receive an all-zeros label vector.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **entity_idxes** (*dict*) – Entity-name → index mapping.
- **relation_idxes** (*dict*) – Relation-name → index mapping.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data = None

train_target = None

label_smoothing_rate

collate_fn = None

target_dim

__len__ ()

__getitem__ (*idx*)

```
class dicee.dataset_classes.KvsAll (train_set_idx: numpy.ndarray, entity_idxes, relation_idxes, form,  
    store=None, label_smoothing_rate: float = 0.0)
```

Bases: torch.utils.data.Dataset

Dataset for KvsAll training (multi-label).

D := {(x, y)_i }_{i=1}^N where

- x = (h, r) is a unique (entity, relation) pair observed in the KG,
- y ∈ [0, 1]^{|IEI|} is a multi-label vector with y_j = 1 iff (h, r, e_j) ∈ KG.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **entity_idxes** (*dict*) – Entity-name → index mapping.
- **relation_idxes** (*dict*) – Relation-name → index mapping.
- **form** (*str*) – 'EntityPrediction' or 'RelationPrediction'.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data = None

train_target = None

`label_smoothing_rate`

`collate_fn = None`

`__len__()`

`__getitem__(idx)`

class `dicee.dataset_classes.KvsSampleDataset` (*train_set_idx: numpy.ndarray, entity_idxes, relation_idxes, form, store=None, neg_ratio=None, label_smoothing_rate: float = 0.0*)

Bases: `torch.utils.data.Dataset`

Dataset for KvsSample training (dynamic multi-label).

Like `KvsAll` but sub-samples the target vector at each access to keep mini-batch sizes manageable when the entity set is large.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **entity_idxes** (*dict*) – Entity-name → index mapping.
- **relation_idxes** (*dict*) – Relation-name → index mapping.
- **form** (*str*) – 'EntityPrediction'.
- **neg_ratio** (*int*) – Number of negative samples per positive target.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

`train_data = None`

`train_target = None`

`neg_ratio = None`

`num_entities`

`label_smoothing_rate`

`collate_fn = None`

`max_num_of_classes`

`__len__()`

`__getitem__(idx)`

class `dicee.dataset_classes.OnevsAllDataset` (*train_set_idx: numpy.ndarray, entity_idxes*)

Bases: `torch.utils.data.Dataset`

Dataset for the 1-vs-All training strategy (multi-class).

Each sample is a (head, relation) pair with a one-hot target vector whose single active position corresponds to the true tail entity.

Parameters

- **train_set_idx** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **entity_idxes** (*dict*) – Entity-name → index mapping (used to determine the target dimension).

```

train_data

target_dim

collate_fn = None

__len__()

__getitem__(idx)

```

```

class dicee.dataset_classes.FixedNegSampleDataset (train_set: numpy.ndarray, num_entities: int,
                                                    num_relations: int, neg_sample_ratio: int = 1, label_smoothing_rate: float = 0.0, seed: int = None)

```

Bases: torch.utils.data.Dataset

Pre-computed (fixed) negative sampling dataset.

At construction time every positive triple is paired with one random negative (head- or tail-corrupted) using vectorized operations for efficiency. The pairs are stored so that `__getitem__` is a simple lookup.

This is useful when you want deterministic negatives across epochs (e.g., for reproducibility or debugging).

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int, optional*) – Number of negative samples per positive triple (default 1).
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

```
neg_sample_ratio = 1
```

```
num_entities
```

```
num_relations
```

```
label_smoothing_rate = 0.0
```

```
collate_fn = None
```

```
seed = None
```

```
train_triples
```

```
length
```

```
__len__() → int
```

```
__getitem__(idx: int) → Tuple[torch.Tensor, torch.Tensor]
```

```

class dicee.dataset_classes.OnevsSample (train_set: numpy.ndarray, num_entities: int,
                                           num_relations: int, neg_sample_ratio: int = None, label_smoothing_rate: float = 0.0)

```

Bases: torch.utils.data.Dataset

Dataset for 1-vs-Sample training (dynamic multi-class with negatives).

For every positive triple (h, r, t) the dataset draws `neg_sample_ratio` random entities as negatives and returns a label vector that marks the true tail and the negatives.

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int*) – Number of negative samples per positive.
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

train_data

num_entities

num_relations

neg_sample_ratio = None

label_smoothing_rate

collate_fn = None

__len__()

__getitem__(*idx*)

```
class dicee.dataset_classes.TriplePredictionDataset (train_set: numpy.ndarray,  
    num_entities: int, num_relations: int, neg_sample_ratio: int = 1, label_smoothing_rate: float = 0.0,  
    seed: int = None)
```

Bases: `torch.utils.data.Dataset`

Dataset for triple prediction with on-the-fly negative sampling.

Each item is a single positive triple; the custom `collate_fn` generates a batch of mixed positive and negative triples.

Parameters

- **train_set** (*numpy.ndarray*) – (N, 3) integer-indexed triples.
- **num_entities** (*int*) – Total number of entities.
- **num_relations** (*int*) – Total number of relations.
- **neg_sample_ratio** (*int, optional*) – Number of negative samples per positive triple (default 1).
- **label_smoothing_rate** (*float, optional*) – Label smoothing coefficient (default 0.0).

label_smoothing_rate

neg_sample_ratio

seed = None

train_set

length

num_entities

```

num_relations

__len__()

__getitem__(idx)

collate_fn(batch: List[torch.Tensor])

class dicee.dataset_classes.LiteralDataset (file_path: str, ent_idx: dict = None,
      normalization_type: str = 'z-norm', sampling_ratio: float = None, loader_backend: str = 'pandas')
Bases: torch.utils.data.Dataset

Dataset for loading and processing literal data for Literal Embedding models.

Handles loading, normalization, and preparation of (entity, attribute, value) triples. Supports z-score
and min-max normalization as well as optional sub-sampling for ablation studies.

```

Parameters

- **file_path** (*str*) – Path to the training data file (CSV / TSV / RDF).
- **ent_idx** (*dict*) – Entity-name → index mapping.
- **normalization_type** (*str, optional*) – 'z-norm', 'min-max', or None (default 'z-norm').
- **sampling_ratio** (*float or None, optional*) – Fraction of the training set to keep (default None = use all).
- **loader_backend** (*str, optional*) – 'pandas' or 'rdflib' (default 'pandas').

```

train_file_path

loader_backend = 'pandas'

normalization_type = 'z-norm'

normalization_params

sampling_ratio = None

entity_to_idx = None

num_entities

__getitem__(index)

__len__()

static load_and_validate_literal_data (file_path: str = None, loader_backend: str = 'pandas')
      → pandas.DataFrame

Load and validate a literal data file.

```

Parameters

- **file_path** (*str*) – Path to the data file.
- **loader_backend** (*str*) – 'pandas' or 'rdflib'.

Returns

Three-column DataFrame with columns head, attribute, value.

Return type

pandas.DataFrame

static denormalize (*preds_norm, attributes, normalization_params*) → *numpy.ndarray*

Reverse the normalization applied during training.

Parameters

- **preds_norm** (*numpy.ndarray*) – Normalized predictions.
- **attributes** (*list*) – Attribute names corresponding to each prediction.
- **normalization_params** (*dict*) – Parameters stored during training.

Returns

Denormalized predictions.

Return type

numpy.ndarray

`dicee.dataset_classes.construct_dataset` (* (*Keyword-only parameters separator (PEP 3102)*)),
train_set: numpy.ndarray | list, valid_set=None, test_set=None, ordered_bpe_entities=None,
train_target_indices=None, target_dim: int = None, entity_to_idx: dict, relation_to_idx: dict,
form_of_labelling: str, scoring_technique: str, neg_ratio: int, label_smoothing_rate: float,
byte_pair_encoding=None, block_size: int = None, seed: int = None) → *torch.utils.data.Dataset*

Build the appropriate dataset for the given training configuration.

Parameters

- **train_set** (*numpy.ndarray or list*) – Raw integer-indexed triples.
- **entity_to_idx** (*dict*) – Name → index mappings.
- **relation_to_idx** (*dict*) – Name → index mappings.
- **form_of_labelling** (*str*) – 'EntityPrediction' or 'RelationPrediction'.
- **scoring_technique** (*str*) – One of 'NegSample', 'FixedNegSample', '1vsAll', '1vsSample', 'KvsAll', 'AllvsAll', 'KvsSample'.
- **neg_ratio** (*int*) – Negative sample ratio.
- **label_smoothing_rate** (*float*) – Label smoothing coefficient.

Return type

torch.utils.data.Dataset

`dicee.dataset_classes.reload_dataset` (*path: str, form_of_labelling, scoring_technique, neg_ratio,*
label_smoothing_rate)

Reload training data from disk and construct a Pytorch dataset.

dicee.eval_static_funcs

Static evaluation functions for KGE models.

This module provides backward compatibility by re-exporting from the new `dicee.evaluation` module.

Deprecated since version Use: `dicee.evaluation` submodules instead. This module will be removed in a future version.

Functions

`evaluate_link_prediction_performance`(→
Dict[str, float])

Evaluate link prediction performance with head and tail prediction.

continues on next page

Table 9 – continued from previous page

<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with reciprocal relations.
<code>evaluate_link_prediction_performance_with_j</code>	Evaluate link prediction with BPE encoding (head and tail).
<code>evaluate_link_prediction_performance_with_j</code>	Evaluate link prediction with BPE encoding and reciprocals.
<code>evaluate_lp_bpe_k_vs_all(→ Dict[str, float])</code>	Evaluate BPE link prediction with KvsAll scoring.
<code>evaluate_literal_prediction(→ Optional[pandas.DataFrame])</code>	Evaluate trained literal prediction model on a test file.
<code>evaluate_ensemble_link_prediction_performance</code>	Evaluate link prediction performance of an ensemble of KGE models.

Module Contents

`dicee.eval_static_funcs.evaluate_link_prediction_performance(model, triples, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.eval_static_funcs.evaluate_link_prediction_performance_with_reciprocals(model, triples, er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.eval_static_funcs.evaluate_link_prediction_performance_with_bpe(model, within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.

- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.eval_static_funcs.evaluate_link_prediction_performance_with_bpe_reciprocals(
    model, within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List])
    → Dict[str, float]
```

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.eval_static_funcs.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]],
    er_vocab: Dict | None = None, batch_size: int | None = None,
    func_triple_to_bpe_representation: Callable | None = None,
    str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

```
dicee.eval_static_funcs.evaluate_literal_prediction(kge_model, eval_file_path: str = None,
    store_lit_preds: bool = True, eval_literals: bool = True, loader_backend: str = 'pandas',
    return_attr_error_metrics: bool = False) → pandas.DataFrame | None
```

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.

- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is True, otherwise None.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

`dicee.eval_static_funcs.evaluate_ensemble_link_prediction_performance` (*models: List, triples, er_vocab: Dict[Tuple, List], weights: List[float] | None = None, batch_size: int = 512, weighted_averaging: bool = True, normalize_scores: bool = True*) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is True but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

dicee.evaluation

Evaluation module for knowledge graph embedding models.

This module provides comprehensive evaluation capabilities for KGE models, including link prediction, literal prediction, and ensemble evaluation.

Modules:

link_prediction: Functions for evaluating link prediction performance
literal_prediction: Functions for evaluating literal/attribute prediction
ensemble: Functions for ensemble model evaluation
evaluator: Main Evaluator class for integrated evaluation
utils: Shared utility functions for evaluation

Example

```
>>> from dicee.evaluation import Evaluator
>>> from dicee.evaluation.link_prediction import evaluate_link_prediction_performance
>>> from dicee.evaluation.ensemble import evaluate_ensemble_link_prediction_
↪performance
```

Submodules

dicee.evaluation.ensemble

Ensemble evaluation functions.

This module provides functions for evaluating ensemble models, including weighted averaging and score normalization.

Functions

<code>evaluate_ensemble_link_prediction_performance</code>	Evaluate link prediction performance of an ensemble of KGE models.
--	--

Module Contents

`dicee.evaluation.ensemble.evaluate_ensemble_link_prediction_performance` (*models*: List, *triples*, *er_vocab*: Dict[Tuple, List], *weights*: List[float] | None = None, *batch_size*: int = 512, *weighted_averaging*: bool = True, *normalize_scores*: bool = True) → Dict[str, float]

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is True but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

dicee.evaluation.evaluator

Main Evaluator class for KGE model evaluation.

This module provides the Evaluator class which orchestrates evaluation of knowledge graph embedding models across different datasets and scoring techniques.

Attributes

VALID_SCORING_TECHNIQUES

Classes

Evaluator

Evaluator class for KGE models in various downstream tasks.

Module Contents

`dicee.evaluation.evaluator.VALID_SCORING_TECHNIQUES`

class `dicee.evaluation.evaluator.Evaluator` (*args*, *is_continual_training*: *bool* = *False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset, trained_model, form_of_labelling: str, during_training: bool = False*) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set, valid_set=None, test_set=None, trained_model*) → None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set=None, valid_set=None, test_set=None, ordered_bpe_entities, trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, *raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate Byte model with generation.

eval_with_bpe_vs_all (*, *raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, *train_set, valid_set=None, test_set=None, trained_model, form_of_labelling*) → None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all (*model, triple_idx, info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – 'EntityPrediction' or 'RelationPrediction'.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp_with_byte (*model*, *triples*: List[List[str]], *info*: str | None = None) → Dict[str, float]

Evaluate Byte model with text generation.

Parameters

- **model** – Byte model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (*model*, *triples*: List[List[str]], *info*: str | None = None, *form_of_labelling*: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (*model*, *triple_idx*, *info*: str) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: str) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: numpy.ndarray, *form_of_labelling*: str) → Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.

- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

dicee.evaluation.link_prediction

Link prediction evaluation functions.

This module provides various functions for evaluating link prediction performance of knowledge graph embedding models.

Functions

<code>evaluate_link_prediction_performance(→ Dict[str, float])</code>	Evaluate link prediction performance with head and tail prediction.
<code>evaluate_link_prediction_performance_with_reciprocals(→ Dict[str, float])</code>	Evaluate link prediction with reciprocal relations.
<code>evaluate_link_prediction_performance_with_reciprocals_and_bpe(→ Dict[str, float])</code>	Evaluate link prediction with BPE encoding and reciprocals.
<code>evaluate_link_prediction_performance_with_reciprocals_and_bpe_head_tail(→ Dict[str, float])</code>	Evaluate link prediction with BPE encoding (head and tail).
<code>evaluate_lp(→ Dict[str, float])</code>	Evaluate link prediction with batched processing.
<code>evaluate_bpe_lp(→ Dict[str, float])</code>	Evaluate link prediction with BPE-encoded entities.
<code>evaluate_lp_bpe_k_vs_all(→ Dict[str, float])</code>	Evaluate BPE link prediction with KvsAll scoring.

Module Contents

`dicee.evaluation.link_prediction.evaluate_link_prediction_performance(model, triples, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.link_prediction.evaluate_link_prediction_performance_with_reciprocals(model, triples, er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.link_prediction.`

`evaluate_link_prediction_performance_with_bpe_reciprocals(model, within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.link_prediction.evaluate_link_prediction_performance_with_bpe(`

`model, within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List]) → Dict[str, float]`

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.link_prediction.evaluate_lp(model, triple_idx, num_entities: int,`

`er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List], info: str = 'Eval Starts', batch_size: int = 128, chunk_size: int = 1000) → Dict[str, float]`

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.

- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_bpe_lp(model, triple_idx: List[Tuple],
all_bpe_shaped_entities, er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List],
info: str = 'Eval Starts') → Dict[str, float]
```

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.link_prediction.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]],
er_vocab: Dict | None = None, batch_size: int | None = None,
func_triple_to_bpe_representation: Callable | None = None,
str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

dicee.evaluation.literal_prediction

Literal prediction evaluation functions.

This module provides functions for evaluating literal/attribute prediction performance of knowledge graph embedding models.

Functions

<code>evaluate_literal_prediction</code> (→ Optional[pandas.DataFrame])	Evaluate trained literal prediction model on a test file.
---	---

Module Contents

```
dicee.evaluation.literal_prediction.evaluate_literal_prediction(kge_model,  
    eval_file_path: str = None, store_lit_preds: bool = True, eval_literals: bool = True,  
    loader_backend: str = 'pandas', return_attr_error_metrics: bool = False)  
    → pandas.DataFrame | None
```

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.
- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is True, otherwise None.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE  
>>> from dicee.evaluation import evaluate_literal_prediction  
>>> model = KGE(path="pretrained_model")  
>>> metrics = evaluate_literal_prediction(  
...     model,  
...     eval_file_path="test_literals.csv",  
...     return_attr_error_metrics=True  
... )  
>>> print(metrics)
```

`dicee.evaluation.utils`

Utility functions for evaluation module.

This module contains shared helper functions used across different evaluation components.

Attributes

<code>DEFAULT_HITS_RANGE</code>
<code>ALL_HITS_RANGE</code>

Functions

<code>make_iterable_verbose(→ Iterable)</code>	Wrap an iterable with tqdm progress bar if verbose is True.
<code>compute_metrics_from_ranks(→ Dict[str, float])</code>	Compute standard link prediction metrics from ranks.
<code>compute_metrics_from_ranks_simple(→ Dict[str, float])</code>	Compute link prediction metrics without scaling factor.
<code>update_hits(→ None)</code>	Update hits dictionary based on rank.
<code>create_hits_dict(→ Dict[int, List[float]])</code>	Create an initialized hits dictionary.
<code>efficient_zero_grad(→ None)</code>	Efficiently zero gradients using <code>parameter.grad = None</code> .

Module Contents

`dicee.evaluation.utils.DEFAULT_HITS_RANGE: List[int] = [1, 3, 10]`

`dicee.evaluation.utils.ALL_HITS_RANGE: List[int]`

`dicee.evaluation.utils.make_iterable_verbose(iterable_object: Iterable, verbose: bool, desc: str = 'Default', position: int | None = None, leave: bool = True) → Iterable`

Wrap an iterable with tqdm progress bar if verbose is True.

Parameters

- **iterable_object** – The iterable to potentially wrap.
- **verbose** – Whether to show progress bar.
- **desc** – Description for the progress bar.
- **position** – Position of the progress bar.
- **leave** – Whether to leave the progress bar after completion.

Returns

The original iterable or a tqdm-wrapped version.

`dicee.evaluation.utils.compute_metrics_from_ranks(ranks: List[int], num_triples: int, hits_dict: Dict[int, List[float]], scale_factor: int = 1) → Dict[str, float]`

Compute standard link prediction metrics from ranks.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.
- **scale_factor** – Factor to scale the denominator (e.g., 2 for head+tail).

Returns

Dictionary containing H@1, H@3, H@10, and MRR metrics.

`dicee.evaluation.utils.compute_metrics_from_ranks_simple` (*ranks*: List[int], *num_triples*: int, *hits_dict*: Dict[int, List[float]]) → Dict[str, float]

Compute link prediction metrics without scaling factor.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.

Returns

Dictionary containing H@1, H@3, H@10, and MRR metrics.

`dicee.evaluation.utils.update_hits` (*hits*: Dict[int, List[float]], *rank*: int, *hits_range*: List[int] | None = None) → None

Update hits dictionary based on rank.

Parameters

- **hits** – Dictionary to update in-place.
- **rank** – The rank to check against hit levels.
- **hits_range** – List of hit levels to check (default: ALL_HITS_RANGE).

`dicee.evaluation.utils.create_hits_dict` (*hits_range*: List[int] | None = None) → Dict[int, List[float]]

Create an initialized hits dictionary.

Parameters

hits_range – List of hit levels to initialize (default: ALL_HITS_RANGE).

Returns

Dictionary with empty lists for each hit level.

`dicee.evaluation.utils.efficient_zero_grad` (*model*) → None

Efficiently zero gradients using `parameter.grad = None`.

This is more efficient than `optimizer.zero_grad()` as it avoids memory operations.

See: https://pytorch.org/tutorials/recipes/recipes/tuning_guide.html

Parameters

model – PyTorch model to zero gradients for.

Attributes

`DEFAULT_HITS_RANGE`

`ALL_HITS_RANGE`

Classes

`Evaluator`

Evaluator class for KGE models in various downstream tasks.

Functions

<code>evaluate_link_prediction_performance(→ Dict[str, float])</code>	Evaluate link prediction performance with head and tail prediction.
<code>evaluate_link_prediction_performance_with_.</code>	Evaluate link prediction with reciprocal relations.
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with BPE encoding (head and tail).
<code>evaluate_link_prediction_performance_with_</code>	Evaluate link prediction with BPE encoding and reciprocals.
<code>evaluate_lp(→ Dict[str, float])</code>	Evaluate link prediction with batched processing.
<code>evaluate_lp_bpe_k_vs_all(→ Dict[str, float])</code>	Evaluate BPE link prediction with KvsAll scoring.
<code>evaluate_bpe_lp(→ Dict[str, float])</code>	Evaluate link prediction with BPE-encoded entities.
<code>evaluate_literal_prediction(→ Optional[pandas.DataFrame])</code>	Evaluate trained literal prediction model on a test file.
<code>evaluate_ensemble_link_prediction_performance</code>	Evaluate link prediction performance of an ensemble of KGE models.
<code>compute_metrics_from_ranks(→ Dict[str, float])</code>	Compute standard link prediction metrics from ranks.
<code>compute_metrics_from_ranks_simple(→ Dict[str, float])</code>	Compute link prediction metrics without scaling factor.
<code>make_iterable_verbose(→ Iterable)</code>	Wrap an iterable with tqdm progress bar if verbose is True.
<code>update_hits(→ None)</code>	Update hits dictionary based on rank.
<code>create_hits_dict(→ Dict[int, List[float]])</code>	Create an initialized hits dictionary.

Package Contents

class `dicee.evaluation.Evaluator` (*args, is_continual_training: bool = False*)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset*, *trained_model*, *form_of_labelling*: str, *during_training*: bool = False) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set*, *valid_set*=None, *test_set*=None, *trained_model*) → None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set*=None, *valid_set*=None, *test_set*=None, *ordered_bpe_entities*, *trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte(**, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate ByteE model with generation.

eval_with_bpe_vs_all(**, raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all(**, train_set, valid_set=None, test_set=None, trained_model, form_of_labelling*) → None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all(*model, triple_idx, info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp_with_byte(*model, triples: List[List[str]], info: str | None = None*) → Dict[str, float]

Evaluate ByteE model with text generation.

Parameters

- **model** – ByteE model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all(*model, triples: List[List[str]], info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp(*model, triple_idx, info: str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: *str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: *numpy.ndarray*, *form_of_labelling*: *str*)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

`dicee.evaluation.evaluate_link_prediction_performance` (*model*, *triples*,
er_vocab: Dict[Tuple, List], *re_vocab*: Dict[Tuple, List]) → Dict[str, float]

Evaluate link prediction performance with head and tail prediction.

Performs filtered evaluation where known correct answers are filtered out before computing ranks.

Parameters

- **model** – KGE model wrapper with entity/relation mappings.
- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.evaluation.evaluate_link_prediction_performance_with_reciprocals` (*model*, *triples*,
er_vocab: Dict[Tuple, List]) → Dict[str, float]

Evaluate link prediction with reciprocal relations.

Optimized for models trained with reciprocal triples where only tail prediction is needed.

Parameters

- **model** – KGE model wrapper.

- **triples** – Test triples as list of (head, relation, tail) strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance_with_bpe(model,
    within_entities: List[str], triples: List[Tuple[str]], er_vocab: Dict[Tuple, List],
    re_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding (head and tail).

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of (head, relation, tail) tuples.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.
- **re_vocab** – Mapping (relation, entity) -> list of valid head entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_link_prediction_performance_with_bpe_reciprocals(model,
    within_entities: List[str], triples: List[List[str]], er_vocab: Dict[Tuple, List]) → Dict[str, float]
```

Evaluate link prediction with BPE encoding and reciprocals.

Parameters

- **model** – KGE model wrapper with BPE support.
- **within_entities** – List of entities to evaluate within.
- **triples** – Test triples as list of [head, relation, tail] strings.
- **er_vocab** – Mapping (entity, relation) -> list of valid tail entities.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_lp(model, triple_idx, num_entities: int, er_vocab: Dict[Tuple, List],
    re_vocab: Dict[Tuple, List], info: str = 'Eval Starts', batch_size: int = 128, chunk_size: int = 1000)
    → Dict[str, float]
```

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.
- **batch_size** – Batch size for triple processing.

- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_lp_bpe_k_vs_all(model, triples: List[List[str]],
    er_vocab: Dict | None = None, batch_size: int | None = None,
    func_triple_to_bpe_representation: Callable | None = None,
    str_to_bpe_entity_to_idx: Dict | None = None) → Dict[str, float]
```

Evaluate BPE link prediction with KvsAll scoring.

Parameters

- **model** – The KGE model wrapper.
- **triples** – List of string triples.
- **er_vocab** – Entity-relation vocabulary for filtering.
- **batch_size** – Batch size for processing.
- **func_triple_to_bpe_representation** – Function to convert triples to BPE.
- **str_to_bpe_entity_to_idx** – Mapping from string entities to BPE indices.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

Raises

ValueError – If batch_size is not provided.

```
dicee.evaluation.evaluate_bpe_lp(model, triple_idx: List[Tuple], all_bpe_shaped_entities,
    er_vocab: Dict[Tuple, List], re_vocab: Dict[Tuple, List], info: str = 'Eval Starts') → Dict[str, float]
```

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.
- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

```
dicee.evaluation.evaluate_literal_prediction(kge_model, eval_file_path: str = None,
    store_lit_preds: bool = True, eval_literals: bool = True, loader_backend: str = 'pandas',
    return_attr_error_metrics: bool = False) → pandas.DataFrame | None
```

Evaluate trained literal prediction model on a test file.

Evaluates the literal prediction capabilities of a KGE model by computing MAE and RMSE metrics for each attribute.

Parameters

- **kge_model** – Trained KGE model with literal prediction capability.
- **eval_file_path** – Path to the evaluation file containing test literals.

- **store_lit_preds** – If True, stores predictions to CSV file.
- **eval_literals** – If True, evaluates and prints error metrics.
- **loader_backend** – Backend for loading dataset ('pandas' or 'rdflib').
- **return_attr_error_metrics** – If True, returns the metrics DataFrame.

Returns

DataFrame with per-attribute MAE and RMSE if `return_attr_error_metrics` is True, otherwise None.

Raises

- **RuntimeError** – If the KGE model doesn't have a trained literal model.
- **AssertionError** – If model is invalid or test set has no valid data.

Example

```
>>> from dicee import KGE
>>> from dicee.evaluation import evaluate_literal_prediction
>>> model = KGE(path="pretrained_model")
>>> metrics = evaluate_literal_prediction(
...     model,
...     eval_file_path="test_literals.csv",
...     return_attr_error_metrics=True
... )
>>> print(metrics)
```

`dicee.evaluation.evaluate_ensemble_link_prediction_performance(models: List, triples, er_vocab: Dict[Tuple, List], weights: List[float] | None = None, batch_size: int = 512, weighted_averaging: bool = True, normalize_scores: bool = True) → Dict[str, float]`

Evaluate link prediction performance of an ensemble of KGE models.

Combines predictions from multiple models using weighted or simple averaging, with optional score normalization.

Parameters

- **models** – List of KGE models (e.g., snapshots from training).
- **triples** – Test triples as numpy array or list, shape (N, 3), with integer indices (head, relation, tail).
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices for filtered evaluation.
- **weights** – Weights for model averaging. Required if `weighted_averaging` is True. Must sum to 1 for proper averaging.
- **batch_size** – Batch size for processing triples.
- **weighted_averaging** – If True, use weighted averaging of predictions. If False, use simple mean.
- **normalize_scores** – If True, normalize scores to [0, 1] range per sample before averaging.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

Raises

AssertionError – If `weighted_averaging` is True but weights are not provided or have wrong length.

Example

```
>>> from dicee.evaluation import evaluate_ensemble_link_prediction_performance
>>> models = [model1, model2, model3]
>>> weights = [0.5, 0.3, 0.2]
>>> results = evaluate_ensemble_link_prediction_performance(
...     models, test_triples, er_vocab,
...     weights=weights, weighted_averaging=True
... )
>>> print(f"MRR: {results['MRR']:.4f}")
```

`dicee.evaluation.compute_metrics_from_ranks` (*ranks: List[int], num_triples: int, hits_dict: Dict[int, List[float]], scale_factor: int = 1*) → Dict[str, float]

Compute standard link prediction metrics from ranks.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.
- **scale_factor** – Factor to scale the denominator (e.g., 2 for head+tail).

Returns

Dictionary containing H@1, H@3, H@10, and MRR metrics.

`dicee.evaluation.compute_metrics_from_ranks_simple` (*ranks: List[int], num_triples: int, hits_dict: Dict[int, List[float]]*) → Dict[str, float]

Compute link prediction metrics without scaling factor.

Parameters

- **ranks** – List of ranks for each prediction.
- **num_triples** – Total number of triples evaluated.
- **hits_dict** – Dictionary mapping hit levels to lists of hits.

Returns

Dictionary containing H@1, H@3, H@10, and MRR metrics.

`dicee.evaluation.make_iterable_verbose` (*iterable_object: Iterable, verbose: bool, desc: str = 'Default', position: int | None = None, leave: bool = True*) → Iterable

Wrap an iterable with tqdm progress bar if verbose is True.

Parameters

- **iterable_object** – The iterable to potentially wrap.
- **verbose** – Whether to show progress bar.
- **desc** – Description for the progress bar.
- **position** – Position of the progress bar.
- **leave** – Whether to leave the progress bar after completion.

Returns

The original iterable or a tqdm-wrapped version.

`dicee.evaluation.update_hits` (*hits*: Dict[int, List[float]], *rank*: int, *hits_range*: List[int] | None = None) → None

Update hits dictionary based on rank.

Parameters

- **hits** – Dictionary to update in-place.
- **rank** – The rank to check against hit levels.
- **hits_range** – List of hit levels to check (default: ALL_HITS_RANGE).

`dicee.evaluation.create_hits_dict` (*hits_range*: List[int] | None = None) → Dict[int, List[float]]

Create an initialized hits dictionary.

Parameters

hits_range – List of hit levels to initialize (default: ALL_HITS_RANGE).

Returns

Dictionary with empty lists for each hit level.

`dicee.evaluation.DEFAULT_HITS_RANGE`: List[int] = [1, 3, 10]

`dicee.evaluation.ALL_HITS_RANGE`: List[int]

dicee.evaluator

Evaluator module for knowledge graph embedding models.

This module provides backward compatibility by re-exporting from the new `dicee.evaluation` module.

Deprecated since version Use: `dicee.evaluation.Evaluator` instead. This module will be removed in a future version.

Classes

Evaluator

Evaluator class for KGE models in various downstream tasks.

Module Contents

class `dicee.evaluator.Evaluator` (*args*, *is_continual_training*: bool = False)

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations
Total number of relations in the knowledge graph.

args
Configuration arguments.

report
Dictionary storing evaluation results.

during_training
Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

```
re_vocab: Dict | None = None
er_vocab: Dict | None = None
ee_vocab: Dict | None = None
func_triple_to_bpe_representation = None
is_continual_training = False
num_entities: int | None = None
num_relations: int | None = None
domain_constraints_per_rel = None
range_constraints_per_rel = None
args
report: Dict
during_training = False
vocab_preparation(dataset) → None
    Prepare vocabularies from the dataset for evaluation.
    Resolves any future objects and saves vocabularies to disk.
```

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (dataset, trained_model, form_of_labelling: str, during_training: bool = False) → Dict | None
Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').

- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set*, *valid_set=None*, *test_set=None*, *trained_model*)
→ None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set=None*, *valid_set=None*,
test_set=None, *ordered_bpe_entities*, *trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, *raw_train_set*, *raw_valid_set=None*, *raw_test_set=None*, *trained_model*,
form_of_labelling) → None

Evaluate Byte model with generation.

eval_with_bpe_vs_all (*, *raw_train_set*, *raw_valid_set=None*, *raw_test_set=None*, *trained_model*,
form_of_labelling) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, *train_set*, *valid_set=None*, *test_set=None*, *trained_model*, *form_of_labelling*)
→ None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all (*model*, *triple_idx*, *info: str | None = None*,
form_of_labelling: str | None = None) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, H@10, and MRR metrics.

evaluate_lp_with_byte (*model*, *triples: List[List[str]]*, *info: str | None = None*) → Dict[str, float]

Evaluate Byte model with text generation.

Parameters

- **model** – Byte model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (*model*, *triples: List[List[str]]*, *info: str | None = None*,
form_of_labelling: str | None = None) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.

- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (*model*, *triple_idx*, *info*: *str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model*, *form_of_labelling*: *str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset*, *trained_model*, *triple_idx*: *numpy.ndarray*, *form_of_labelling*: *str*)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

dicee.executer

Executor module for training, retraining and evaluating KGE models.

This module provides the Execute and ContinuousExecute classes for managing the full lifecycle of knowledge graph embedding model training.

Classes

<i>Execute</i>	Executor class for training, retraining and evaluating KGE models.
<i>ContinuousExecute</i>	A subclass of Execute Class for retraining

Module Contents

class `dicee.executer.Execute` (*args*, *continuous_training: bool = False*)

Executor class for training, retraining and evaluating KGE models.

Handles the complete workflow: 1. Loading & Preprocessing & Serializing input data 2. Training & Validation & Testing 3. Storing all necessary information

args

Processed input arguments.

distributed

Whether distributed training is enabled.

rank

Process rank in distributed training.

world_size

Total number of processes.

local_rank

Local GPU rank.

trainer

Training handler instance.

trained_model

The trained model after training completes.

knowledge_graph

The loaded knowledge graph.

report

Dictionary storing training metrics and results.

evaluator

Model evaluation handler.

distributed

args

is_continual_training = False

trainer: `dicee.trainer.DICE_Trainer` | None = None

trained_model = None

knowledge_graph: `dicee.knowledge_graph.KG` | None = None

report: Dict

evaluator: `dicee.evaluator.Evaluator` | None = None

start_time: float | None = None

is_rank_zero() → bool

cleanup()

setup_executor() → None

Set up storage directories for the experiment.

Creates or reuses experiment directories based on configuration. Saves the configuration to a JSON file.

create_and_store_kg() → None

Create knowledge graph and store as memory-mapped file.

Only executed on rank 0 in distributed training. Skips if memmap already exists.

load_from_memmap() → None

Load knowledge graph from memory-mapped file.

save_trained_model() → None

Save a knowledge graph embedding model

(1) Send model to eval mode and cpu.

(2) Store the memory footprint of the model.

(3) Save the model into disk.

(4) Update the stats of KG again ?

Parameter

rtype

None

end(*form_of_labelling: str*) → dict

End training

(1) Store trained model.

(2) Report runtimes.

(3) Eval model if required.

Parameter

rtype

A dict containing information about the training and/or evaluation

write_report() → None

Report training related information in a report.json file

start() → dict

Start training

(1) Loading the Data # (2) Create an evaluator object. # (3) Create a trainer object. # (4) Start the training

Parameter

rtype

A dict containing information about the training and/or evaluation

class dicee.executer.**ContinuousExecute** (*args*)

Bases: *Execute*

A subclass of Execute Class for retraining

- (1) Loading & Preprocessing & Serializing input data.
- (2) Training & Validation & Testing
- (3) Storing all necessary info

During the continual learning we can only modify * **num_epochs** * parameter. Trained model stored in the same folder as the seed model for the training. Trained model is noted with the current time.

continual_start () → dict

Start Continual Training

- (1) Initialize training.
- (2) Start continual training.
- (3) Save trained model.

Parameter

rtype

A dict containing information about the training and/or evaluation

dicee.knowledge_graph

Knowledge Graph module for data loading and preprocessing.

Provides the KG class for handling knowledge graph data including loading, preprocessing, and indexing operations.

Classes

<i>KG</i>	Knowledge Graph container and processor.
-----------	--

Module Contents

```
class dicee.knowledge_graph.KG (dataset_dir: str | None = None, byte_pair_encoding: bool = False,  
    padding: bool = False, add_noise_rate: float | None = None, sparql_endpoint: str | None = None,  
    path_single_kg: str | None = None, path_for_deserialization: str | None = None,  
    add_reciprocal: bool | None = None, eval_model: str | None = None,  
    read_only_few: int | None = None, sample_triples_ratio: float | None = None,  
    path_for_serialization: str | None = None, entity_to_idx: Dict | None = None,  
    relation_to_idx: Dict | None = None, backend: str | None = None,  
    training_technique: str | None = None, separator: str | None = None)
```

Knowledge Graph container and processor.

Handles loading, preprocessing, and indexing of knowledge graph data from various sources including files, SPARQL endpoints, and serialized formats.

dataset_dir
Path to directory containing train/valid/test files.

num_entities
Total number of unique entities.

num_relations
Total number of unique relations.

train_set
Indexed training triples as numpy array.

valid_set
Indexed validation triples (optional).

test_set
Indexed test triples (optional).

entity_to_idx
Mapping from entity strings to indices.

relation_to_idx
Mapping from relation strings to indices.

dataset_dir = None

sparql_endpoint = None

path_single_kg = None

byte_pair_encoding = False

ordered_shaped_bpe_tokens = None

add_noise_rate = None

num_entities: int | None = None

num_relations: int | None = None

path_for_deserialization = None

add_reciprocal = None

eval_model = None

read_only_few = None

sample_triples_ratio = None

path_for_serialization = None

entity_to_idx = None

relation_to_idx = None

backend = 'pandas'

training_technique = None

```

separator = None

raw_train_set = None

raw_valid_set = None

raw_test_set = None

train_set = None

valid_set = None

test_set = None

idx_entity_to_bpe_shaped: Dict

enc

num_tokens

num_bpe_entities: int | None = None

padding = False

dummy_id

max_length_subword_tokens: int | None = None

train_set_target = None

target_dim: int | None = None

train_target_indices = None

ordered_bpe_entities = None

description_of_input = None

describe() → None
    Generate a description string of the dataset statistics.

property_entities_str: List[str]
    Get list of all entity strings.

property_relations_str: List[str]
    Get list of all relation strings.

exists(h: str, r: str, t: str) → bool
    Check if a triple exists in the training set.

```

Parameters

- **h** – Head entity string.
- **r** – Relation string.
- **t** – Tail entity string.

Returns

True if the triple exists, False otherwise.

```

__iter__() → Iterator[Tuple[str, str, str]]
    Iterate over training triples as string tuples.

__len__() → int
    Return number of triples in the raw training set.

func_triple_to_bpe_representation(triple: List[str])

```

dicee.knowledge_graph_embeddings

Classes

<i>KGE</i>	Knowledge Graph Embedding Class for interactive usage of pre-trained models
------------	---

Module Contents

```

class dicee.knowledge_graph_embeddings.KGE (path=None, url=None, construct_ensemble=False,
      model_name=None)

    Bases: dicee.abstracts.BaseInteractiveKGE, dicee.abstracts.InteractiveQueryDecomposition, dicee.abstracts.BaseInteractiveTrainKGE

    Knowledge Graph Embedding Class for interactive usage of pre-trained models

    __str__()

    to(device: str) → None

    get_transductive_entity_embeddings(indices: torch.LongTensor | List[str], as_pytorch=False,
      as_numpy=False, as_list=True) → torch.FloatTensor | numpy.ndarray | List[float]

    create_vector_database(collection_name: str, distance: str, location: str = 'localhost',
      port: int = 6333)

    generate(h="", r="")

    eval_lp_performance(dataset=List[Tuple[str, str, str]], filtered=True)

    predict_missing_head_entity(relation: List[str] | str, tail_entity: List[str] | str, within=None,
      batch_size=2, topk=1, return_indices=False) → Tuple

    Given a relation and a tail entity, return top k ranked head entity.

     $\operatorname{argmax}_{\{e \in E\}} f(e, r, t)$ , where  $r \in R$ ,  $t \in E$ .

```

Parameter

relation: Union[List[str], str]
String representation of selected relations.

tail_entity: Union[List[str], str]
String representation of selected entities.

k: int
Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_relations (*head_entity: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False*) → Tuple

Given a head entity and a tail entity, return top k ranked relations.

$\text{argmax}_{\{r \in R\}} f(h,r,t)$, where $h, t \in E$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_tail_entity (*head_entity: List[str] | str, relation: List[str] | str, within: List[str] = None, batch_size=2, topk=1, return_indices=False*) → torch.FloatTensor

Given a head entity and a relation, return top k ranked entities

$\text{argmax}_{\{e \in E\}} f(h,r,e)$, where $h \in E$ and $r \in R$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

Returns: Tuple

scores

predict (*, *h: List[str] | str = None, r: List[str] | str = None, t: List[str] | str = None, within=None, logits=True*) → torch.FloatTensor

Parameters

- **logits**
- **h**
- **r**
- **t**
- **within**

predict_topk (*, h: str | List[str] = None, r: str | List[str] = None, t: str | List[str] = None, topk: int = 10, within: List[str] = None, batch_size: int = 1024)

Predict missing item in a given triple.

Returns

- If you query a single (h, r, ?) or (?, r, t) or (h, ?, t), returns List[(item, score)]
- If you query a batch of B, returns List of B such lists.

triple_score (h: List[str] | str = None, r: List[str] | str = None, t: List[str] | str = None, logits=False)
→ torch.FloatTensor

Predict triple score

Parameter

head_entity: List[str]

String representation of selected entities.

relation: List[str]

String representation of selected relations.

tail_entity: List[str]

String representation of selected entities.

logits: bool

If logits is True, unnormalized score returned

Returns: Tuple

pytorch tensor of triple score

return_multi_hop_query_results (aggregated_query_for_all_entities, k: int, only_scores)

single_hop_query_answering (query: tuple, only_scores: bool = True, k: int = None)

answer_multi_hop_query (query_type: str = None, query: Tuple[str | Tuple[str, str], Ellipsis] = None, queries: List[Tuple[str | Tuple[str, str], Ellipsis]] = None, tnorm: str = 'prod', neg_norm: str = 'standard', lambda_: float = 0.0, k: int = 10, only_scores=False)
→ List[Tuple[str, torch.Tensor]]

@TODO: Refactoring is needed # @TODO: Score computation for each query type should be done in a static function

Find an answer set for EPFO queries including negation and disjunction

Parameter

query_type: str The type of the query, e.g., “2p”.

query: Union[str, Tuple[str, Tuple[str, str]]] The query itself, either a string or a nested tuple.

queries: List of Tuple[Union[str, Tuple[str, str]], ...]

tnorm: str The t-norm operator.

neg_norm: str The negation norm.

lambda_: float lambda parameter for sugeno and yager negation norms

k: int The top-k substitutions for intermediate variables.

returns

- *List[Tuple[str, torch.Tensor]]*
- *Entities and corresponding scores sorted in the descening order of scores*

find_missing_triples (*confidence: float, entities: List[str] = None, relations: List[str] = None, topk: int = 10, at_most: int = sys.maxsize*) → Set

Find missing triples

Iterative over a set of entities E and a set of relation R :

orall e in E and orall r in R f(e,r,x)

Return (e,r,x)

otin G and f(e,r,x) > confidence

confidence: float

A threshold for an output of a sigmoid function given a triple.

topk: int

Highest ranked k item to select triples with f(e,r,x) > confidence .

at_most: int

Stop after finding at_most missing triples

{(e,r,x) | f(e,r,x) > confidence land (e,r,x)}

otin G

predict_literals (*entity: List[str] | str = None, attribute: List[str] | str = None, denormalize_preds: bool = True*) → numpy.ndarray

Predicts literal values for given entities and attributes.

Parameters

- **entity** (*Union[List[str], str]*) – Entity or list of entities to predict literals for.
- **attribute** (*Union[List[str], str]*) – Attribute or list of attributes to predict literals for.
- **denormalize_preds** (*bool*) – If True, denormalizes the predictions.

Returns

Predictions for the given entities and attributes.

Return type

numpy ndarray

dicee.models

Submodules

dicee.models.adopt

ADOPT Optimizer Implementation.

This module implements the ADOPT (Adaptive Optimization with Precise Tracking) algorithm, an advanced optimization method for training neural networks.

ADOPT Overview:

ADOPT is an adaptive learning rate optimization algorithm that combines the benefits of momentum-based methods with per-parameter learning rate adaptation. Unlike Adam, which applies momentum to raw gradients, ADOPT normalizes gradients first and then applies momentum, leading to more stable training dynamics.

Key Features: - Gradient normalization before momentum application - Adaptive per-parameter learning rates - Optional gradient clipping that grows with training steps - Support for decoupled weight decay (AdamW-style) - Multiple execution modes: single-tensor, multi-tensor (foreach), and fused (planned)

Algorithm Comparison:

Adam: $m = \beta_1 * m + (1 - \beta_1) * g$, $\theta = \theta - \alpha * m / \sqrt{v}$ ADOPT: $m = \beta_1 * m + (1 - \beta_1) * g / \sqrt{v}$, $\theta = \theta - \alpha * m$

The key difference is that ADOPT normalizes gradients before momentum, which provides better stability and can lead to improved convergence.

Classes:

- ADOPT: Main optimizer class (extends torch.optim.Optimizer)

Functions:

- adopt: Functional API for ADOPT algorithm computation
- _single_tensor_adopt: Single-tensor implementation (TorchScript compatible)
- _multi_tensor_adopt: Multi-tensor implementation using foreach operations

Performance:

- Single-tensor: Default, compatible with torch.jit.script
- Multi-tensor (foreach): 2-3x faster on GPU through vectorization
- Fused (planned): Would provide maximum performance via specialized kernels

Example:

```
>>> import torch
>>> from dicee.models.adopt import ADOPT
>>>
>>> model = torch.nn.Linear(10, 1)
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01, decouple=True)
>>>
>>> # Training loop
>>> for epoch in range(num_epochs):
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     optimizer.step()
```

References:

Original implementation: <https://github.com/iShohei220/adopt>

Notes:

This implementation is based on the original ADOPT implementation and adapted to work with the PyTorch optimizer interface and the dicee framework.

Classes

<code>ADOPT</code>	ADOPT Optimizer.
--------------------	------------------

Functions

<code>adopt(params, grads, exp_avgs, exp_avg_sqs, state_steps)</code>	Functional API that performs ADOPT algorithm computation.
---	---

Module Contents

```
class dicee.models.adopt.ADOPT (params: torch.optim.optimizer.ParamsT,  
    lr: float | torch.Tensor = 0.001, betas: Tuple[float, float] = (0.9, 0.9999), eps: float = 1e-06,  
    clip_lambda: Callable[[int], float] | None = lambda step: ..., weight_decay: float = 0.0,  
    decouple: bool = False, *, foreach: bool | None = None, maximize: bool = False,  
    capturable: bool = False, differentiable: bool = False, fused: bool | None = None)
```

Bases: torch.optim.optimizer.Optimizer

ADOPT Optimizer.

ADOPT is an adaptive learning rate optimization algorithm that combines momentum-based updates with adaptive per-parameter learning rates. It uses exponential moving averages of gradients and squared gradients, with gradient clipping for stability.

The algorithm performs the following key operations: 1. Normalizes gradients by the square root of the second moment estimate 2. Applies optional gradient clipping based on the training step 3. Updates parameters using momentum-smoothed normalized gradients 4. Supports decoupled weight decay (AdamW-style) or L2 regularization

Mathematical formulation:

$$m_t = \beta_1 * m_{t-1} + (1 - \beta_1) * \text{clip}(g_t / \sqrt{v_t}) \quad v_t = \beta_2 * v_{t-1} + (1 - \beta_2) * g_t^2 \quad \theta_t = \theta_{t-1} - \alpha * m_t$$

where:

- θ_t : parameter at step t
- g_t : gradient at step t
- m_t : first moment estimate (momentum)
- v_t : second moment estimate (variance)
- α : learning rate
- β_1, β_2 : exponential decay rates
- $\text{clip}()$: optional gradient clipping function

Reference:

Original implementation: <https://github.com/iShohei220/adopt>

Parameters

- **params** (*ParamsT*) – Iterable of parameters to optimize or dicts defining parameter groups.
- **lr** (*float or Tensor, optional*) – Learning rate. Can be a float or 1-element Tensor. Default: 1e-3
- **betas** (*Tuple[float, float], optional*) – Coefficients (β_1, β_2) for computing running averages of gradient and its square. β_1 controls momentum, β_2 controls variance. Default: (0.9, 0.9999)
- **eps** (*float, optional*) – Term added to denominator to improve numerical stability. Default: 1e-6
- **clip_lambda** (*Callable[[int], float], optional*) – Function that takes the step number and returns the gradient clipping threshold. Common choices: - lambda step: step**0.25 (default, gradually increases clipping threshold) - lambda step: 1.0 (constant clipping) - None (no clipping) Default: lambda step: step**0.25
- **weight_decay** (*float, optional*) – Weight decay coefficient (L2 penalty). Default: 0.0
- **decouple** (*bool, optional*) – If True, uses decoupled weight decay (AdamW-style), applying weight decay directly to parameters. If False, adds weight decay to gradients (L2 regularization). Default: False
- **foreach** (*bool, optional*) – If True, uses the faster foreach implementation for multi-tensor operations. Default: None (auto-select)
- **maximize** (*bool, optional*) – If True, maximizes parameters instead of minimizing. Useful for reinforcement learning. Default: False
- **capturable** (*bool, optional*) – If True, the optimizer is safe to capture in a CUDA graph. Requires learning rate as Tensor. Default: False
- **differentiable** (*bool, optional*) – If True, the optimization step can be differentiated. Useful for meta-learning. Default: False
- **fused** (*bool, optional*) – If True, uses fused kernel implementation (currently not supported). Default: None

Raises

- **ValueError** – If learning rate, epsilon, betas, or weight_decay are invalid.
- **RuntimeError** – If fused is enabled (not currently supported).
- **RuntimeError** – If lr is a Tensor with foreach=True and capturable=False.

Example

```
>>> # Basic usage
>>> optimizer = ADOPT(model.parameters(), lr=0.001)
>>> optimizer.zero_grad()
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With decoupled weight decay
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01,
↳decouple=True)
```

```
>>> # Custom gradient clipping
>>> optimizer = ADOPT(model.parameters(), clip_lambda=lambda step: max(1.0,
↳step**0.5))
```

Note

- For most use cases, the default hyperparameters work well
- Consider using decouple=True for better generalization (similar to AdamW)
- The clip_lambda function helps stabilize training in early steps

clip_lambda

__setstate__ (state)

Restore optimizer state from a checkpoint.

This method handles backward compatibility when loading optimizer state from older versions. It ensures all required fields are present with default values and properly converts step counters to tensors if needed.

Key responsibilities: 1. Set default values for newly added hyperparameters 2. Convert old-style scalar step counters to tensor format 3. Place step tensors on appropriate devices based on capturable/fused modes

Parameters

state (dict) – Optimizer state dictionary (typically from torch.load()).

Note

- This enables loading checkpoints saved with older ADOPT versions
- Step counters are converted to appropriate device/dtype for compatibility
- Capturable and fused modes require step tensors on parameter devices

step (closure=None)

Perform a single optimization step.

This method executes one iteration of the ADOPT optimization algorithm across all parameter groups. It orchestrates the following workflow:

1. Optionally evaluates a closure to recompute the loss (useful for algorithms like LBFGS or when loss needs multiple evaluations)
2. For each parameter group: - Collects parameters with gradients and their associated state - Extracts hyperparameters (betas, learning rate, etc.) - Calls the functional adopt() API to perform the actual update
3. Returns the loss value if a closure was provided

The functional API (adopt()) handles three execution modes: - Single-tensor: Updates one parameter at a time (default, JIT-compatible) - Multi-tensor (foreach): Batches operations for better performance - Fused: Uses fused CUDA kernels (not yet implemented)

Gradient scaling support: This method is compatible with automatic mixed precision (AMP) training. It can access `grad_scale` and `found_inf` attributes for gradient unscaling and inf/nan detection when used with `GradScaler`.

Parameters

closure (*Callable, optional*) – A callable that reevaluates the model and returns the loss. The closure should: - Enable gradients (`torch.enable_grad()`) - Compute forward pass - Compute loss - Compute backward pass - Return the loss value Example: `lambda: (loss := model(x), loss.backward(), loss)[-1]` Default: `None`

Returns

The loss value returned by the closure, or `None` if no closure was provided.

Return type

Optional[`Tensor`]

Example

```
>>> # Standard usage
>>> loss = criterion(model(input), target)
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With closure (e.g., for line search)
>>> def closure():
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     return loss
>>> loss = optimizer.step(closure)
```

Note

- Call `zero_grad()` before computing gradients for the next step
- CUDA graph capture is checked for safety when `capturable=True`
- The method is thread-safe for different parameter groups

`dicee.models.adopt.adopt` (*params: List[torch.Tensor], grads: List[torch.Tensor], exp_avgs: List[torch.Tensor], exp_avg_sqs: List[torch.Tensor], state_steps: List[torch.Tensor], foreach: bool | None = None, capturable: bool = False, differentiable: bool = False, fused: bool | None = None, grad_scale: torch.Tensor | None = None, found_inf: torch.Tensor | None = None, has_complex: bool = False, *, beta1: float, beta2: float, lr: float | torch.Tensor, clip_lambda: Callable[[int], float] | None, weight_decay: float, decouple: bool, eps: float, maximize: bool*)

Functional API that performs ADOPT algorithm computation.

This is the main functional interface for the ADOPT optimization algorithm. It dispatches to one of three implementations based on the execution mode:

1. **Single-tensor mode** (default): Updates parameters one at a time - Compatible with `torch.jit.script` - More flexible but slower - Used when `foreach=False` or automatically for small models

2. **Multi-tensor (foreach) mode:** Batches operations across tensors - 2-3x faster on GPU through vectorization
- Groups tensors by device/dtype automatically - Used when foreach=True
3. **Fused mode:** Uses specialized fused kernels (not yet implemented) - Would provide maximum performance
- Currently raises RuntimeError if enabled

Algorithm overview (ADOPT):

ADOPT adapts learning rates per-parameter while using momentum on normalized gradients. The key innovation is normalizing gradients before momentum, which provides more stable training than standard Adam.

Mathematical formulation:

```
# Normalize gradient by its historical variance  $\text{normed\_g\_t} = \text{g\_t} / \sqrt{(\text{v\_t} + \epsilon)}$ 

# Optional gradient clipping for stability  $\text{normed\_g\_t} = \text{clip}(\text{normed\_g\_t}, \text{threshold}(t))$ 

# Momentum on normalized gradients (key difference from Adam)  $\text{m\_t} = \beta_1 * \text{m}_{\{t-1\}} + (1 - \beta_1) * \text{normed\_g\_t}$ 

# Parameter update  $\theta\_t = \theta_{\{t-1\}} - \alpha * \text{m\_t}$ 

# Update variance estimate  $\text{v\_t} = \beta_2 * \text{v}_{\{t-1\}} + (1 - \beta_2) * \text{g\_t}^2$ 
```

where:

- θ : parameters
- g : gradients
- m : first moment (momentum of normalized gradients)
- v : second moment (variance of raw gradients)
- α : learning rate
- β_1, β_2 : exponential decay rates
- ϵ : numerical stability constant
- $\text{clip}()$: gradient clipping function based on step

Automatic mode selection:

When foreach and fused are both None (default), the function automatically selects the best implementation based on:
- Parameter types and devices - Whether differentiable mode is enabled - Learning rate type (float vs Tensor)
- Capturable mode requirements

param params

Parameters to optimize.

type params

List[Tensor]

param grads

Gradients for each parameter.

type grads

List[Tensor]

param exp_avgs

First moment estimates (momentum).

type exp_avgs

List[Tensor]

param exp_avg_sqs
Second moment estimates (variance).

type exp_avg_sqs
List[Tensor]

param state_steps
Step counters (must be singleton tensors).

type state_steps
List[Tensor]

param foreach
Whether to use multi-tensor implementation. None: auto-select based on configuration (default).

type foreach
Optional[bool]

param capturable
If True, ensure CUDA graph capture safety.

type capturable
bool

param differentiable
If True, allow gradients through optimization step.

type differentiable
bool

param fused
If True, use fused kernels (not implemented).

type fused
Optional[bool]

param grad_scale
Gradient scaler for AMP training.

type grad_scale
Optional[Tensor]

param found_inf
Flag for inf/nan detection in AMP.

type found_inf
Optional[Tensor]

param has_complex
Whether any parameters are complex-valued.

type has_complex
bool

param beta1
Exponential decay rate for first moment (momentum). Typical range: 0.9-0.95.

type beta1
float

param beta2
Exponential decay rate for second moment (variance). Typical range: 0.999-0.9999 (higher than Adam).

type beta2
float

param lr
Learning rate. Can be a scalar Tensor for dynamic learning rate with capturable=True.

type lr
Union[float, Tensor]

param clip_lambda
Function that maps step number to gradient clipping threshold. None disables clipping.

type clip_lambda
Optional[Callable[[int], float]]

param weight_decay
Weight decay coefficient (L2 penalty).

type weight_decay
float

param decouple
If True, use decoupled weight decay (AdamW-style). Recommended for better generalization.

type decouple
bool

param eps
Small constant for numerical stability in normalization.

type eps
float

param maximize
If True, maximize objective instead of minimize.

type maximize
bool

raises RuntimeError
If torch.jit.script is used with foreach or fused.

raises RuntimeError
If state_steps contains non-tensor elements.

raises RuntimeError
If fused=True (not yet implemented).

raises RuntimeError
If lr is Tensor with foreach=True and capturable=False.

Example

```
>>> # Typically called by ADOPT optimizer, not directly
>>> adopt (
...     params=[p1, p2],
...     grads=[g1, g2],
...     exp_avgs=[m1, m2],
...     exp_avg_sqs=[v1, v2],
...     state_steps=[step1, step2],
...     beta1=0.9,
```

(continues on next page)

(continued from previous page)

```
...     beta2=0.9999,
...     lr=0.001,
...     clip_lambda=lambda s: s**0.25,
...     weight_decay=0.01,
...     decouple=True,
...     eps=1e-6,
...     maximize=False,
... )
```

Note

- For distributed training, this API is compatible with torch/distributed/optim
- The foreach mode is generally preferred for GPU training
- Complex parameters are handled transparently by viewing as real
- First optimization step only initializes variance, doesn't update parameters

See also

- ADOPT class: High-level optimizer interface
- `_single_tensor_adopt`: Single-tensor implementation details
- `_multi_tensor_adopt`: Multi-tensor implementation details

dicee.models.base_model

Classes

<code>BaseKGELightning</code>	Base class for all neural network modules.
<code>BaseKGE</code>	Base class for all neural network modules.
<code>IdentityClass</code>	Base class for all neural network modules.

Module Contents

class `dicee.models.base_model.BaseKGELightning` (*args, **kwargs)

Bases: `lightning.LightningModule`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F
```

(continues on next page)

(continued from previous page)

```
class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`training_step_outputs = []`

`mem_of_model()` → Dict

Size of model in MB and number of params

`training_step (batch, batch_idx=None)`

Here you compute and return the training loss and some additional metrics for e.g. the progress bar or logger.

Parameters

- **batch** – The output of your data iterable, normally a `DataLoader`.
- **batch_idx** – The index of this batch.
- **dataloader_idx** – The index of the dataloader that produced this batch. (only if multiple dataloaders used)

Returns

- `Tensor` - The loss tensor
- `dict` - A dictionary which can include any keys, but must include the key `'loss'` in the case of automatic optimization.
- `None` - In automatic optimization, this will skip to the next batch (but is not supported for multi-GPU, TPU, or DeepSpeed). For manual optimization, this has no special meaning, as returning the loss is not required.

In this step you'd normally do the forward pass and calculate the loss for a batch. You can also do fancier things like multiple forward passes or something model specific.

Example:

```
def training_step(self, batch, batch_idx):
    x, y, z = batch
    out = self.encoder(x)
    loss = self.loss(out, x)
    return loss
```

To use multiple optimizers, you can switch to ‘manual optimization’ and control their stepping:

```
def __init__(self):
    super().__init__()
    self.automatic_optimization = False

# Multiple optimizers (e.g.: GANs)
def training_step(self, batch, batch_idx):
    opt1, opt2 = self.optimizers()

    # do training_step with encoder
    ...
    opt1.step()
    # do training_step with decoder
    ...
    opt2.step()
```

Note

When `accumulate_grad_batches > 1`, the loss returned here will be automatically normalized by `accumulate_grad_batches` internally.

loss_function (*yhat_batch: torch.FloatTensor, y_batch: torch.FloatTensor*)

Parameters

- **yhat_batch**
- **y_batch**

on_train_epoch_end (*args, **kwargs)

Called in the training loop at the very end of the epoch.

To access all batch outputs at the end of the epoch, you can cache step outputs as an attribute of the `LightningModule` and access them in this hook:

```
class MyLightningModule(L.LightningModule):
    def __init__(self):
        super().__init__()
        self.training_step_outputs = []

    def training_step(self):
        loss = ...
        self.training_step_outputs.append(loss)
        return loss

    def on_train_epoch_end(self):
```

(continues on next page)

(continued from previous page)

```
# do something with all training_step outputs, for example:
epoch_mean = torch.stack(self.training_step_outputs).mean()
self.log("training_epoch_mean", epoch_mean)
# free up the memory
self.training_step_outputs.clear()
```

test_epoch_end(*outputs: List[Any]*)

test_dataloader() → None

An iterable or collection of iterables specifying test samples.

For more information about multiple dataloaders, see this section.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

Warning

do not assign state in `prepare_data`

- `test()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Note

If you don't need a test dataset and a `test_step()`, you don't need to implement this method.

val_dataloader() → None

An iterable or collection of iterables specifying validation samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `fit()`
- `validate()`

- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Note

If you don't need a validation dataset and a `validation_step()`, you don't need to implement this method.

`predict_dataloader()` → None

An iterable or collection of iterables specifying prediction samples.

For more information about multiple dataloaders, see this section.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `predict()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Returns

A `torch.utils.data.DataLoader` or a sequence of them specifying prediction samples.

`train_dataloader()` → None

An iterable or collection of iterables specifying training samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **`:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs``** to a positive integer.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

Warning

do not assign state in `prepare_data`

- `fit()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

`configure_optimizers` (*parameters=None*)

Choose what optimizers and learning-rate schedulers to use in your optimization. Normally you'd need one. But in the case of GANs or similar you might have multiple. Optimization with multiple optimizers only works in the manual optimization mode.

Returns

Any of these 6 options.

- **Single optimizer.**
- **List or Tuple** of optimizers.
- **Two lists** - The first list has multiple optimizers, and the second has multiple LR schedulers (or multiple `lr_scheduler_config`).
- **Dictionary**, with an "optimizer" key, and (optionally) a "lr_scheduler" key whose value is a single LR scheduler or `lr_scheduler_config`.
- **None** - Fit will run without any optimizer.

The `lr_scheduler_config` is a dictionary which contains the scheduler and its associated configuration. The default configuration is shown below.

```
lr_scheduler_config = {
    # REQUIRED: The scheduler instance
    "scheduler": lr_scheduler,
    # The unit of the scheduler's step size, could also be 'step'.
    # 'epoch' updates the scheduler on epoch end whereas 'step'
    # updates it after a optimizer update.
    "interval": "epoch",
    # How many epochs/steps should pass between calls to
    # `scheduler.step()`. 1 corresponds to updating the learning
    # rate after every epoch/step.
    "frequency": 1,
    # Metric to monitor for schedulers like `ReduceLROnPlateau`
    "monitor": "val_loss",
    # If set to `True`, will enforce that the value specified 'monitor'
    # is available when the scheduler is updated, thus stopping
    # training if not found. If set to `False`, it will only produce a warning
```

(continues on next page)

(continued from previous page)

```
"strict": True,
# If using the `LearningRateMonitor` callback to monitor the
# learning rate progress, this keyword can be used to specify
# a custom logged name
"name": None,
}
```

When there are schedulers in which the `.step()` method is conditioned on a value, such as the `torch.optim.lr_scheduler.ReduceLROnPlateau` scheduler, Lightning requires that the `lr_scheduler_config` contains the keyword "monitor" set to the metric name that the scheduler should be conditioned on.

Metrics can be made available to monitor by simply logging it using `self.log('metric_to_track', metric_val)` in your `LightningModule`.

Note

Some things to know:

- Lightning calls `.backward()` and `.step()` automatically in case of automatic optimization.
- If a learning rate scheduler is specified in `configure_optimizers()` with key "interval" (default "epoch") in the scheduler configuration, Lightning will call the scheduler's `.step()` method automatically in case of automatic optimization.
- If you use 16-bit precision (`precision=16`), Lightning will automatically handle the optimizer.
- If you use `torch.optim.LBFGS`, Lightning handles the closure function automatically for you.
- If you use multiple optimizers, you will have to switch to 'manual optimization' mode and step them yourself.
- If you need to control how often the optimizer steps, override the `optimizer_step()` hook.

```
class dicee.models.base_model.BaseKGE (args: dict)
```

Bases: `BaseKGELightning`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
```

(continues on next page)

```
x = F.relu(self.conv1(x))
return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

args

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

loss

`selected_optimizer = None`

`normalizer_class = None`

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

```

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

    Parameters
         $\mathbf{x} (B \times 2 \times T)$ 

forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])
    byte pair encoded neural link predictors

    Parameters
    -----

init_params_with_sanity_checking ()

forward (x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
        y_idx: torch.LongTensor = None)

    Parameters

        •  $\mathbf{x}$ 

        •  $\mathbf{y\_idx}$ 

        •  $\mathbf{ordered\_bpe\_entities}$ 

forward_triples (x: torch.LongTensor)  $\rightarrow$  torch.Tensor

    Parameters

         $\mathbf{x}$ 

forward_k_vs_all (*args, **kwargs)

forward_k_vs_sample (*args, **kwargs)

get_triple_representation (idx_hrt)

get_head_relation_representation (indexed_triple)

get_sentence_representation (x: torch.LongTensor)

    Parameters

        • ( $\mathbf{b} (x \text{ shape})$ )

        • 3

        •  $\mathbf{t}$ )

```

```
get_bpe_head_and_relation_representation(x: torch.LongTensor)
    → Tuple[torch.FloatTensor, torch.FloatTensor]
```

Parameters

$\mathbf{x} (B \times 2 \times T)$

```
get_embeddings() → Tuple[numpy.ndarray, numpy.ndarray]
```

```
class dicee.models.base_model.IdentityClass (args=None)
```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`args = None`

`__call__(x)`

`static forward(x)`

`dicee.models.clifford`

Classes

Keci

Base class for all neural network modules.

continues on next page

Table 27 – continued from previous page

<i>KeciTransformer</i>	Keci with Transformer architecture.
<i>TransformerBlock</i>	A single transformer block with self-attention and MLP.
<i>TransformerSelfAttention</i>	Multi-head self-attention for the Keci Transformer.
<i>TransformerMLP</i>	MLP for the Keci Transformer.
<i>CKeci</i>	Without learning dimension scaling
<i>DeCaL</i>	Base class for all neural network modules.

Module Contents

class dicee.models.clifford.**Keci**(args)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'Keci'

p

q

r

requires_grad_for_interactions = True

compute_sigma_pp (*hp, rp*)

Compute $\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{k=i+1}^p (h_{i,r_k} - h_{k,r_i}) e_i e_k$

σ_{pp} captures the interactions between along p bases For instance, let $p = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_qq (*hq, rq*)

Compute $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_{j,r_k} - h_{k,r_j}) e_j e_k$ σ_{qq} captures the interactions between along q bases For instance, let $q = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_pq (*, *hp, hq, rp, rq*)

$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_{i,r_j} - h_{j,r_i}) e_i e_j$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

apply_coefficients (*hp, hq, rp, rq*)

Multiplying a base vector with its scalar coefficient

clifford_multiplication (*h0, hp, hq, r0, rp, rq*)

Compute our CL multiplication

$h = h_0 + \sum_{i=1}^p h_i e_i + \sum_{j=p+1}^{p+q} h_j e_j$ $r = r_0 + \sum_{i=1}^p r_i e_i + \sum_{j=p+1}^{p+q} r_j e_j$

$e_i^2 = +1$ for $i \leq p$ $e_j^2 = -1$ for $p < j \leq p+q$ $e_i e_j = -e_j e_i$ for i

e_j

$h r = \sigma_0 + \sigma_p + \sigma_q + \sigma_{pp} + \sigma_{qq} + \sigma_{pq}$ where

(1) $\sigma_0 = h_0 r_0 + \sum_{i=1}^p (h_0 r_i - h_i r_0) e_i - \sum_{j=p+1}^{p+q} (h_j r_j) e_j$

- (2) $\sigma_p = \sum_{i=1}^p (h_{0r_i} + h_{ir_0}) e_i$
- (3) $\sigma_q = \sum_{j=p+1}^{p+q} (h_{0r_j} + h_{jr_0}) e_j$
- (4) $\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{k=i+1}^p (h_{ir_k} - h_{kr_i}) e_i e_k$
- (5) $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_{jr_k} - h_{kr_j}) e_j e_k$
- (6) $\sigma_{pq} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_{ir_j} - h_{jr_i}) e_i e_j$

construct_cl_multivector (*x: torch.FloatTensor, r: int, p: int, q: int*)
 $\rightarrow$ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]
Construct a batch of multivectors $Cl_{\{p,q\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,d) shape

returns

- **a0** (*torch.FloatTensor with (n,r) shape*)
- **ap** (*torch.FloatTensor with (n,r,p) shape*)
- **aq** (*torch.FloatTensor with (n,r,q) shape*)

forward_k_vs_with_explicit (*x: torch.Tensor*)
k_vs_all_score (*bpe_head_ent_emb, bpe_rel_ent_emb, E*)
forward_k_vs_all (*x: torch.Tensor*) $\rightarrow$ torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations $\mathbb{R}^d$.
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— x: torch.LongTensor with (n,2) shape :rtype: torch.FloatTensor with (n, **IEI**) shape

construct_batch_selected_cl_multivector (*x: torch.FloatTensor, r: int, p: int, q: int*)
 $\rightarrow$ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]
Construct a batch of batchs multivectors $Cl_{\{p,q\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,k, d) shape

returns

- **a0** (*torch.FloatTensor with (n,k, m) shape*)
- **ap** (*torch.FloatTensor with (n,k, m, p) shape*)
- **aq** (*torch.FloatTensor with (n,k, m, q) shape*)

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*) $\rightarrow$ torch.FloatTensor

Parameter

x: torch.LongTensor with (n,2) shape

target_entity_idx: torch.LongTensor with (n, k) shape k denotes the selected number of examples.

rtype

torch.FloatTensor with (n, k) shape

score (*h, r, t*)

forward_triples (*x: torch.Tensor*) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,3) shape

rtype

torch.FloatTensor with (n) shape

class dicee.models.clifford.**KeciTransformer** (*args*)

Bases: *Keci*

Keci with Transformer architecture.

Concatenates h0, hp, hq, r0, rp, rq into a single embedding vector and processes through transformer.

name = 'KeciTransformer'

input_dim

seq_len = 1

transformer

lm_head

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

Kvsall training

Parameter

x: torch.LongTensor with (n,2) shape

rtype

torch.FloatTensor with (n, **IEI**) shape

class dicee.models.clifford.**TransformerBlock** (*n_embd, n_head, dropout=0.0, bias=False*)

Bases: torch.nn.Module

A single transformer block with self-attention and MLP.

ln_1

attn

ln_2

mlp

forward (*x*)


```
class dicee.models.clifford.TransformerSelfAttention(n_embd, n_head, dropout=0.0,  

bias=False)
```

Bases: `torch.nn.Module`

Multi-head self-attention for the Keci Transformer.

c_attn

c_proj

attn_dropout

resid_dropout

n_head

n_embd

dropout = 0.0

forward(*x*)

```
class dicee.models.clifford.TransformerMLP(n_embd, dropout=0.0, bias=False)
```

Bases: `torch.nn.Module`

MLP for the Keci Transformer.

c_fc

gelu

c_proj

dropout

forward(*x*)

```
class dicee.models.clifford.CKeci(args)
```

Bases: *Keci*

Without learning dimension scaling

name = 'CKeci'

requires_grad_for_interactions = False

```
class dicee.models.clifford.DeCaL(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
```

(continues on next page)

(continued from previous page)

```
def __init__(self) -> None:
    super().__init__()
    self.conv1 = nn.Conv2d(1, 20, 5)
    self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'DeCaL'

entity_embeddings

relation_embeddings

p

q

r

re

forward_triples (*x: torch.Tensor*) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,) shape

rtype

torch.FloatTensor with (n) shape

cl_pqr (*a: torch.tensor*) → torch.tensor

Input: tensor(batch_size, emb_dim) → output: tensor with 1+p+q+r components with size (batch_size, emb_dim/(1+p+q+r)) each.

1) takes a tensor of size (batch_size, emb_dim), split it into 1 + p + q + r components, hence 1+p+q+r must be a divisor of the emb_dim. 2) Return a list of the 1+p+q+r components vectors, each are tensors of size (batch_size, emb_dim/(1+p+q+r))

compute_sigmas_single (*list_h_emb, list_r_emb, list_t_emb*)

here we compute all the sums with no others vectors interaction taken with the scalar product with t, that is,

$$s0 = h_0 r_0 t_0 s1 = \sum_{i=1}^p h_i r_i t_0 s2 = \sum_{j=p+1}^{p+q} h_j r_j t_0 s3 = \sum_{i=1}^q (h_0 r_i t_i + h_i r_0 t_i) s4 = \sum_{i=p+1}^{p+q} (h_0 r_i t_i + h_i r_0 t_i) s5 = \sum_{i=p+q+1}^{p+q+r} (h_0 r_i t_i + h_i r_0 t_i)$$

and return:

$$sigma_0 t = \sigma_0 \cdot t_0 = s0 + s1 - s2 s3, s4 \text{ and } s5$$

compute_sigmas_multivect (*list_h_emb, list_r_emb*)

Here we compute and return all the sums with vectors interaction for the same and different bases.

For same bases vectors interaction we have

$$\sigma_p p = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (h_i r_{i'} - h_{i'} r_i) (model \text{ the interactions between } e_i \text{ and } e_{i'} \text{ for } 1 \leq i, i' \leq p) \sigma_q q = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (h_j r_{j'} - h_{j'} r_j)$$

For different base vector interactions, we have

$$\sigma_p q = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) (interactions \text{ between } e_i \text{ and } e_j \text{ for } 1 \leq i \leq p \text{ and } p+1 \leq j \leq p+q) \sigma_p r = \sum_{i=1}^p \sum_{j=p+q+1}^{p+q+r} (h_i r_j - h_j r_i)$$

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q,r\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— x: torch.LongTensor with (n,) shape :rtype: torch.FloatTensor with (n, **IEI**) shape

apply_coefficients (*h0, hp, hq, hk, r0, rp, rq, rk*)

Multiplying a base vector with its scalar coefficient

construct_cl_multivector (*x: torch.FloatTensor, re: int, p: int, q: int, r: int*)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Construct a batch of multivectors $Cl_{\{p,q,r\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,d) shape

returns

- **a0** (*torch.FloatTensor*)
- **ap** (*torch.FloatTensor*)
- **aq** (*torch.FloatTensor*)
- **ar** (*torch.FloatTensor*)

compute_sigma_pp (*hp, rp*)

Compute .. math:

$$\sigma_{p,p} = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (x_i y_{i'} - x_{i'} y_i)$$

σ_{pp} captures the interactions between along p bases For instance, let p e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3, and e_2 e_3 This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_qq (*hq, rq*)

Compute

$$\sigma_{q,q}^* = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (x_j y_{j'} - x_{j'} y_j) Eq.16$$

σ_{qq} captures the interactions between along q bases For instance, let q e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3, and e_2 e_3 This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_rr (*hk, rk*)

$$\sigma_{r,r}^* = \sum_{k=p+q+1}^{p+q+r-1} \sum_{k'=k+1}^p (x_k y_{k'} - x_{k'} y_k)$$

compute_sigma_pq (*, *hp, hq, rp, rq*)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

```

print(sigma_pq.shape)
compute_sigma_pr(*, hp, hk, rp, rk)
Compute

```

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

```

results = []
sigma_pq = torch.zeros(b, r, p, q)
for i in range(p):
    for j in range(q):
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
print(sigma_pq.shape)
compute_sigma_qr(*, hq, hk, rq, rk)

```

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

```

results = []
sigma_pq = torch.zeros(b, r, p, q)
for i in range(p):
    for j in range(q):
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
print(sigma_pq.shape)

```

dicее.models.complex

Classes

<i>ConEx</i>	Convolutional ComplEx Knowledge Graph Embeddings
<i>AConEx</i>	Additive Convolutional ComplEx Knowledge Graph Embeddings
<i>ComplEx</i>	Base class for all neural network modules.

Module Contents

```

class dicее.models.complex.ConEx(args)
    Bases: dicее.models.base_model.BaseKGE
    Convolutional ComplEx Knowledge Graph Embeddings
    name = 'ConEx'
    conv2d
    fc_num_input
    fc1
    norm_fc1
    bn_conv2d
    feature_map_dropout

```

```
residual_convolution (C_1: Tuple[torch.Tensor, torch.Tensor],
                       C_2: Tuple[torch.Tensor, torch.Tensor]) → torch.FloatTensor
```

Compute residual score of two complex-valued embeddings. :param *C_1*: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param *C_2*: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

```
forward_k_vs_all (x: torch.Tensor) → torch.FloatTensor
```

```
forward_triples (x: torch.Tensor) → torch.FloatTensor
```

Parameters

x

```
forward_k_vs_sample (x: torch.Tensor, target_entity_idx: torch.Tensor)
```

```
class dicee.models.complex.AConEx (args)
```

Bases: [dicee.models.base_model.BaseKGE](#)

Additive Convolutional ComplEx Knowledge Graph Embeddings

```
name = 'AConEx'
```

```
conv2d
```

```
fc_num_input
```

```
fc1
```

```
norm_fc1
```

```
bn_conv2d
```

```
feature_map_dropout
```

```
residual_convolution (C_1: Tuple[torch.Tensor, torch.Tensor],
                       C_2: Tuple[torch.Tensor, torch.Tensor]) → torch.FloatTensor
```

Compute residual score of two complex-valued embeddings. :param *C_1*: a tuple of two pytorch tensors that corresponds complex-valued embeddings :param *C_2*: a tuple of two pytorch tensors that corresponds complex-valued embeddings :return:

```
forward_k_vs_all (x: torch.Tensor) → torch.FloatTensor
```

```
forward_triples (x: torch.Tensor) → torch.FloatTensor
```

Parameters

x

```
forward_k_vs_sample (x: torch.Tensor, target_entity_idx: torch.Tensor)
```

```
class dicee.models.complex.Complex (args)
```

Bases: [dicee.models.base_model.BaseKGE](#)

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'Complex'

static score (*head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor, tail_ent_emb: torch.FloatTensor*)

static k_vs_all_score (*emb_h: torch.FloatTensor, emb_r: torch.FloatTensor, emb_E: torch.FloatTensor*)

Parameters

- **emb_h**
- **emb_r**
- **emb_E**

forward_k_vs_all (*x: torch.LongTensor*) $\rightarrow$ *torch.FloatTensor*

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*)

dicee.models.dualE

Classes

DualE

Dual Quaternion Knowledge Graph Embeddings
(<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

Module Contents

class dicee.models.dualE.DualE(*args*)

Bases: *dicee.models.base_model.BaseKGE*

Dual Quaternion Knowledge Graph Embeddings (<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

name = 'DualE'

entity_embeddings

relation_embeddings

num_ent = None

kvsall_score(*e_1_h, e_2_h, e_3_h, e_4_h, e_5_h, e_6_h, e_7_h, e_8_h, e_1_t, e_2_t, e_3_t, e_4_t, e_5_t, e_6_t, e_7_t, e_8_t, r_1, r_2, r_3, r_4, r_5, r_6, r_7, r_8*) → torch.tensor

KvsAll scoring function

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_triples(*idx_triple: torch.tensor*) → torch.tensor

Negative Sampling forward pass:

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

forward_k_vs_all(*x*)

KvsAll forward pass

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

T(*x: torch.tensor*) → torch.tensor

Transpose function

Input: Tensor with shape (nxm) Output: Tensor with shape (mxn)

dicee.models.ensemble

Classes

<i>EnsembleKGE</i>

Module Contents

class `dicee.models.ensemble.EnsembleKGE` (*models: list = None, seed_model=None, pretrained_models: List = None*)

`name`

`train_mode = True`

`args`

`named_children()`

`property example_input_array`

`parameters()`

`modules()`

`__iter__()`

`__len__()`

`eval()`

`to(device)`

`state_dict()`
Return the state dict of the ensemble.

`load_state_dict(state_dict, strict=True)`
Load the state dict into the ensemble.

`mem_of_model()`

`__call__(x_batch)`

`step()`

`get_embeddings()`

`__str__()`

dicee.models.function_space

Classes

<i>FMult</i>	Learning Knowledge Neural Graphs
<i>GFMult</i>	Learning Knowledge Neural Graphs

continues on next page

Table 31 – continued from previous page

<i>FMult2</i>	Learning Knowledge Neural Graphs
<i>LFMult1</i>	Embedding with trigonometric functions. We represent all entities and relations in the complex number space as:
<i>LFMult</i>	Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:

Module Contents

```
class dicee.models.function_space.FMult(args)
    Bases: dicee.models.base_model.BaseKGE
    Learning Knowledge Neural Graphs
    name = 'FMult'
    entity_embeddings
    relation_embeddings
    k
    num_sample = 50
    gamma
    roots
    weights
    compute_func(weights: torch.FloatTensor, x) → torch.FloatTensor
    chain_func(weights, x: torch.FloatTensor)
    forward_triples(idx_triple: torch.Tensor) → torch.Tensor
```

Parameters

x

```
class dicee.models.function_space.GFMult(args)
    Bases: dicee.models.base_model.BaseKGE
    Learning Knowledge Neural Graphs
    name = 'GFMult'
    entity_embeddings
    relation_embeddings
    k
    num_sample = 250
    roots
    weights
    compute_func(weights: torch.FloatTensor, x) → torch.FloatTensor
```

`chain_func(weights, x: torch.FloatTensor)`

`forward_triples(idx_triple: torch.Tensor) → torch.Tensor`

Parameters

x

`class dicee.models.function_space.FMult2(args)`

Bases: `dicee.models.base_model.BaseKGE`

Learning Knowledge Neural Graphs

`name = 'FMult2'`

`n_layers = 3`

`k`

`n = 50`

`score_func = 'compositional'`

`discrete_points`

`entity_embeddings`

`relation_embeddings`

`build_func(Vec)`

`build_chain_funcs(list_Vec)`

`compute_func(W, b, x) → torch.FloatTensor`

`function(list_W, list_b)`

`trapezoid(list_W, list_b)`

`forward_triples(idx_triple: torch.Tensor) → torch.Tensor`

Parameters

x

`class dicee.models.function_space.LFMult1(args)`

Bases: `dicee.models.base_model.BaseKGE`

Embedding with trigonometric functions. We represent all entities and relations in the complex number space as: $f(x) = \sum_{k=0}^{d-1} w_k e^{kix}$. and use the three differents scoring function as in the paper to evaluate the score

`name = 'LFMult1'`

`entity_embeddings`

`relation_embeddings`

`forward_triples(idx_triple)`

Parameters

x

```

tri_score(h, r, t)

vtp_score(h, r, t)

class dicee.models.function_space.LFMult(args)
    Bases: dicee.models.base_model.BaseKGE

    Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:  $f(x) = \sum_{i=0}^{d-1} a_i x^i$  and use the three differents scoring function as in the paper to evaluate the score. We also consider combining with Neural Networks.

    name = 'LFMult'

    entity_embeddings

    relation_embeddings

    degree

    m

    x_values

    forward_triples(idx_triple)

    Parameters
    x

    construct_multi_coeff(x)

    poly_NN(x, coefh, coefr, coeft)
        Constructing a 2 layers NN to represent the embeddings.  $h = \sigma(w_h^T x + b_h)$ ,  $r = \sigma(w_r^T x + b_r)$ ,  $t = \sigma(w_t^T x + b_t)$ 

    linear(x, w, b)

    scalar_batch_NN(a, b, c)
        element wise multiplication between a,b and c: Inputs : a, b, c ==> torch.tensor of size batch_size x m x d
        Output : a tensor of size batch_size x d

    tri_score(coeff_h, coeff_r, coeff_t)
        this part implement the trilinear scoring techniques:

        
$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \text{dfrac}\{a_i*b_j*c_k\}\{1+(i+j+k)\%d\}$$


        1. generate the range for i,j and k from [0 d-1]
        2. perform  $\text{dfrac}\{a_i*b_j*c_k\}\{1+(i+j+k)\%d\}$  in parallel for every batch
        3. take the sum over each batch

    vtp_score(h, r, t)
        this part implement the vector triple product scoring techniques:

        
$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \text{dfrac}\{a_i*c_j*b_k - b_i*c_j*a_k\}\{(1+(i+j)\%d)(1+k)\}$$


        1. generate the range for i,j and k from [0 d-1]
        2. Compute the first and second terms of the sum
        3. Multiply with then denominator and take the sum

```

4. take the sum over each batch

comp_func (*h, r, t*)

this part implement the function composition scoring techniques: i.e. score = <hor, t>

polynomial (*coeff, x, degree*)

This function takes a matrix tensor of coefficients (*coeff*), a tensor vector of points *x* and range of integer $[0, 1, \dots, d]$ and return a vector tensor ($\text{coeff}[0][0] + \text{coeff}[0][1]x + \dots + \text{coeff}[0][d]x^d$,

$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d$)

pop (*coeff, x, degree*)

This function allow us to evaluate the composition of two polynomes without for loops :) it takes a matrix tensor of coefficients (*coeff*), a matrix tensor of points *x* and range of integer $[0, 1, \dots, d]$

and return a tensor ($\text{coeff}[0][0] + \text{coeff}[0][1]x + \dots + \text{coeff}[0][d]x^d$,

$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d$)

dicee.models.literal

Classes

<i>LiteralEmbeddings</i>	A model for learning and predicting numerical literals using pre-trained KGE.
--------------------------	---

Module Contents

class `dicee.models.literal.LiteralEmbeddings` (*num_of_data_properties: int, embedding_dims: int, entity_embeddings: torch.tensor, dropout: float = 0.3, gate_residual=True, freeze_entity_embeddings=True*)

Bases: `torch.nn.Module`

A model for learning and predicting numerical literals using pre-trained KGE.

num_of_data_properties

Number of data properties (attributes).

Type

int

embedding_dims

Dimension of the embeddings.

Type

int

entity_embeddings

Pre-trained entity embeddings.

Type

torch.tensor

dropout

Dropout rate for regularization.

Type

float

gate_residual

Whether to use gated residual connections.

Type

bool

freeze_entity_embeddings

Whether to freeze the entity embeddings during training.

Type

bool

embedding_dim

num_of_data_properties

hidden_dim

gate_residual = True

freeze_entity_embeddings = True

entity_embeddings

data_property_embeddings

fc

fc_out

dropout

gated_residual_proj

layer_norm

forward(*entity_idx*, *attr_idx*)

Parameters

- **entity_idx** (*Tensor*) – Entity indices (batch).
- **attr_idx** (*Tensor*) – Attribute (Data property) indices (batch).

Returns

scalar predictions.

Return type

Tensor

property device

dicee.models.octonion

Classes

<i>OMult</i>	Base class for all neural network modules.
<i>ConvO</i>	Base class for all neural network modules.

continues on next page

Table 33 – continued from previous page

<i>AConvO</i>	Additive Convolutional Octonion Knowledge Graph Embeddings
---------------	--

Functions

<code>octonion_mul(*, O_1, O_2)</code>
<code>octonion_mul_norm(*, O_1, O_2)</code>

Module Contents

`dicee.models.octonion.octonion_mul(*, O_1, O_2)`

`dicee.models.octonion.octonion_mul_norm(*, O_1, O_2)`

class `dicee.models.octonion.OMult(args)`

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'OMult'

```
static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                           emb_rel_e5, emb_rel_e6, emb_rel_e7)
```

```
score(head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor,
      tail_ent_emb: torch.FloatTensor)
```

```
k_vs_all_score(bpe_head_ent_emb, bpe_rel_ent_emb, E)
```

```
forward_k_vs_all(x)
```

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e.,
[score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and
relations => shape (size of batch,| Entities|)

```
class dicee.models.octonion.ConvO(args: dict)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'ConvO'

conv2d

fc_num_input

fc1


```

bn_conv2d

norm_fc1

feature_map_dropout

static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                           emb_rel_e5, emb_rel_e6, emb_rel_e7)

residual_convolution(O_1, O_2)

forward_triples(x: torch.Tensor) → torch.Tensor

```

Parameters

x

forward_k_vs_all (x: torch.Tensor)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch, **Entities**)

```
class dicee.models.octonion.AConvO(args: dict)
```

Bases: `dicee.models.base_model.BaseKGE`

Additive Convolutional Octonion Knowledge Graph Embeddings

name = 'AConvO'

conv2d

fc_num_input

fc1

bn_conv2d

norm_fc1

feature_map_dropout

```
static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                           emb_rel_e5, emb_rel_e6, emb_rel_e7)
```

```
residual_convolution(O_1, O_2)
```

```
forward_triples(x: torch.Tensor) → torch.Tensor
```

Parameters

x

forward_k_vs_all (x: torch.Tensor)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch, **Entities**)

dicee.models.pykeen_models

Classes

Module Contents

class `dicee.models.pykeen_models.PykeenKGE` (*args: dict*)

Bases: `dicee.models.base_model.BaseKGE`

A class for using knowledge graph embedding models implemented in Pykeen

Notes: Pykeen_DistMult: C Pykeen_ComplEx: Pykeen_QuatE: Pykeen_MuRE: Pykeen_CP: Pykeen_HolE: Pykeen_HolE: Pykeen_HolE: Pykeen_TransD: Pykeen_TransE: Pykeen_TransF: Pykeen_TransH: Pykeen_TransR:

model_kwargs

name

model

loss_history = []

args

entity_embeddings = None

relation_embeddings = None

forward_k_vs_all (*x: torch.LongTensor*)

=> Explicit version by this we can apply bn and dropout

(1) Retrieve embeddings of heads and relations + apply Dropout & Normalization if given. h, r = self.get_head_relation_representation(x) # (2) Reshape (1). if self.last_dim > 0:

h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim, self.last_dim)

(3) Reshape all entities. if self.last_dim > 0:

t = self.entity_embeddings.weight.reshape(self.num_entities, self.embedding_dim, self.last_dim)

else:

t = self.entity_embeddings.weight

(4) Call the score_t from interactions to generate triple scores. return self.interaction.score_t(h=h, r=r, all_entities=t, slice_size=1)

forward_triples (*x: torch.LongTensor*) → torch.FloatTensor

=> Explicit version by this we can apply bn and dropout

(1) Retrieve embeddings of heads, relations and tails and apply Dropout & Normalization if given. h, r, t = self.get_triple_representation(x) # (2) Reshape (1). if self.last_dim > 0:

h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim, self.last_dim) t = t.reshape(len(x), self.embedding_dim, self.last_dim)

(3) Compute the triple score return self.interaction.score(h=h, r=r, t=t, slice_size=None, slice_dim=0)

abstractmethod forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx*)

dicee.models.quaternion

Classes

<i>QMult</i>	Base class for all neural network modules.
<i>ConvQ</i>	Convolutional Quaternion Knowledge Graph Embeddings
<i>AConvQ</i>	Additive Convolutional Quaternion Knowledge Graph Embeddings

Functions

<i>quaternion_mul_with_unit_norm</i> (*, Q_1, Q_2)
--

Module Contents

dicee.models.quaternion.**quaternion_mul_with_unit_norm**(*, Q_1, Q_2)

class dicee.models.quaternion.**QMult** (args)

Bases: *dicee.models.base_model.BaseKGE*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the submodules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

`name = 'QMult'`

`explicit = True`

`quaternion_multiplication_followed_by_inner_product (h, r, t)`

Parameters

- **h** – shape: (*batch_dims, dim) The head representations.
- **r** – shape: (*batch_dims, dim) The head representations.
- **t** – shape: (*batch_dims, dim) The tail representations.

Returns

Triple scores.

static quaternion_normalizer (*x*: torch.FloatTensor) → torch.FloatTensor

Normalize the length of relation vectors, if the forward constraint has not been applied yet.

Absolute value of a quaternion

$$|a + bi + cj + dk| = \sqrt{a^2 + b^2 + c^2 + d^2}$$

L2 norm of quaternion vector:

$$\|x\|^2 = \sum_{i=1}^d |x_i|^2 = \sum_{i=1}^d (x_i.re^2 + x_i.im_1^2 + x_i.im_2^2 + x_i.im_3^2)$$

Parameters

x – The vector.

Returns

The normalized vector.

score (*head_ent_emb*: torch.FloatTensor, *rel_ent_emb*: torch.FloatTensor, *tail_ent_emb*: torch.FloatTensor)

k_vs_all_score (*bpe_head_ent_emb*, *bpe_rel_ent_emb*, *E*)

Parameters

- **bpe_head_ent_emb**
- **bpe_rel_ent_emb**
- **E**

forward_k_vs_all (*x*)

Parameters

x

forward_k_vs_sample (*x*, *target_entity_idx*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e., [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,| Entities|)

class dicee.models.quaternion.ConvQ (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Convolutional Quaternion Knowledge Graph Embeddings

```

name = 'ConvQ'

entity_embeddings

relation_embeddings

conv2d

fc_num_input

fc1

bn_conv1

bn_conv2

feature_map_dropout

residual_convolution(Q_1, Q_2)

forward_triples(indexed_triple: torch.Tensor) → torch.Tensor

```

Parameters

x

```
forward_k_vs_all(x: torch.Tensor)
```

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities)

```
class dicee.models.quaternion.AConvQ(args)
```

Bases: *[dicee.models.base_model.BaseKGE](#)*

Additive Convolutional Quaternion Knowledge Graph Embeddings

```
name = 'AConvQ'
```

```
entity_embeddings
```

```
relation_embeddings
```

```
conv2d
```

```
fc_num_input
```

```
fc1
```

```
bn_conv1
```

```
bn_conv2
```

```
feature_map_dropout
```

```
residual_convolution(Q_1, Q_2)
```

```
forward_triples(indexed_triple: torch.Tensor) → torch.Tensor
```

Parameters

x

forward_k_vs_all (*x*: *torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities)

dicee.models.real

Classes

<i>DistMult</i>	Embedding Entities and Relations for Learning and Inference in Knowledge Bases
<i>TransE</i>	Translating Embeddings for Modeling
<i>Shallom</i>	A shallow neural model for relation prediction (https://arxiv.org/abs/2101.09090)
<i>Pyke</i>	A Physical Embedding Model for Knowledge Graphs
<i>CoKEConfig</i>	Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.
<i>CoKE</i>	Contextualized Knowledge Graph Embedding (CoKE) model.

Module Contents

class dicee.models.real.**DistMult** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Embedding Entities and Relations for Learning and Inference in Knowledge Bases <https://arxiv.org/abs/1412.6575>

name = 'DistMult'

k_vs_all_score (*emb_h*: *torch.FloatTensor*, *emb_r*: *torch.FloatTensor*, *emb_E*: *torch.FloatTensor*)

Parameters

- **emb_h**
- **emb_r**
- **emb_E**

forward_k_vs_all (*x*: *torch.LongTensor*)

forward_k_vs_sample (*x*: *torch.LongTensor*, *target_entity_idx*: *torch.LongTensor*)

score (*h*, *r*, *t*)

class dicee.models.real.**TransE** (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Translating Embeddings for Modeling Multi-relational Data <https://proceedings.neurips.cc/paper/2013/file/1cecc7a77928ca8133fa24680a88d2f9-Paper.pdf>

name = 'TransE'

margin = 4

score (*head_ent_emb*, *rel_ent_emb*, *tail_ent_emb*)

```

forward_k_vs_all (x: torch.Tensor) → torch.FloatTensor

class dicee.models.real.Shallom(args)
    Bases: dicee.models.base_model.BaseKGE
    A shallow neural model for relation prediction (https://arxiv.org/abs/2101.09090)
    name = 'Shallom'

    shallom

    get_embeddings () → Tuple[numpy.ndarray, None]

    forward_k_vs_all (x) → torch.FloatTensor

    forward_triples (x) → torch.FloatTensor

        Parameters
        x

        Returns

class dicee.models.real.Pyke(args)
    Bases: dicee.models.base_model.BaseKGE
    A Physical Embedding Model for Knowledge Graphs
    name = 'Pyke'

    dist_func

    margin = 1.0

    forward_triples (x: torch.LongTensor)

        Parameters
        x

class dicee.models.real.CoKEConfig
    Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.

    block_size
        Sequence length for transformer (3 for triples: head, relation, tail)

    vocab_size
        Total vocabulary size (num_entities + num_relations)

    n_layer
        Number of transformer layers

    n_head
        Number of attention heads per layer

    n_embd
        Embedding dimension (set to match model embedding_dim)

    dropout
        Dropout rate applied throughout the model

    bias
        Whether to use bias in linear layers

```

```

causal
    Whether to use causal masking (False for bidirectional attention)

block_size: int = 3

vocab_size: int = None

n_layer: int = 6

n_head: int = 8

n_embd: int = None

dropout: float = 0.3

bias: bool = True

causal: bool = False

```

```

class dicee.models.real.CoKE (args, config: CoKEConfig = CoKEConfig())

```

Bases: `dicee.models.base_model.BaseKGE`

Contextualized Knowledge Graph Embedding (CoKE) model. Based on: <https://arxiv.org/pdf/1911.02168>.

CoKE uses a transformer encoder to learn contextualized representations of entities and relations. For link prediction, it predicts masked elements in (head, relation, tail) triples using bidirectional attention, similar to BERT's masked language modeling approach.

The model creates a sequence [head_emb, relation_emb, mask_emb], adds positional embeddings, and processes it through transformer layers to predict the tail entity.

```

name = 'CoKE'

```

```

config

```

```

pos_emb

```

```

mask_emb

```

```

blocks

```

```

ln_f

```

```

coke_dropout

```

```

forward_k_vs_all (x: torch.Tensor)

```

```

score (emb_h, emb_r, emb_t)

```

```

forward_k_vs_sample (x: torch.LongTensor, target_entity_idx: torch.LongTensor)

```

`dicee.models.static_funcs`

Functions

```

quaternion_mul(→ Tuple[torch.Tensor, torch.Tensor,    Perform quaternion multiplication
...)
```

Module Contents

```
dicee.models.static_funcs.quaternion_mul(* , Q_1, Q_2)  
→ Tuple[torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor]
```

Perform quaternion multiplication :param Q_1: :param Q_2: :return:

`dicee.models.transformers`

Full definition of a GPT Language Model, all of it in this single file. References: 1) the official GPT-2 TensorFlow implementation released by OpenAI: <https://github.com/openai/gpt-2/blob/master/src/model.py> 2) huggingface/transformers PyTorch implementation: https://github.com/huggingface/transformers/blob/main/src/transformers/models/gpt2/modeling_gpt2.py

Classes

<i>ByteE</i>	Base class for all neural network modules.
<i>LayerNorm</i>	LayerNorm but with an optional bias. PyTorch doesn't support simply bias=False
<i>SelfAttention</i>	Base class for all neural network modules.
<i>MLP</i>	Base class for all neural network modules.
<i>Block</i>	Base class for all neural network modules.
<i>GPTConfig</i>	
<i>GPT</i>	Base class for all neural network modules.

Module Contents

```
class dicee.models.transformers.ByteE(*args, **kwargs)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the submodules as regular attributes:

```
import torch.nn as nn  
import torch.nn.functional as F  
  
class Model(nn.Module):  
    def __init__(self) -> None:  
        super() __init__()  
        self.conv1 = nn.Conv2d(1, 20, 5)  
        self.conv2 = nn.Conv2d(20, 20, 5)  
  
    def forward(self, x):  
        x = F.relu(self.conv1(x))  
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'ByteE'

config

temperature = 0.5

topk = 2

transformer

lm_head

loss_function (*yhat_batch, y_batch*)

Parameters

- **yhat_batch**
- **y_batch**

forward (*x: torch.LongTensor*)

Parameters

x (*B by T tensor*)

generate (*idx, max_new_tokens, temperature=1.0, top_k=None*)

Take a conditioning sequence of indices *idx* (LongTensor of shape (b,t)) and complete the sequence *max_new_tokens* times, feeding the predictions back into the model each time. Most likely you'll want to make sure to be in `model.eval()` mode of operation for this.

training_step (*batch, batch_idx=None*)

Here you compute and return the training loss and some additional metrics for e.g. the progress bar or logger.

Parameters

- **batch** – The output of your data iterable, normally a `DataLoader`.
- **batch_idx** – The index of this batch.
- **dataloader_idx** – The index of the dataloader that produced this batch. (only if multiple dataloaders used)

Returns

- **Tensor** - The loss tensor
- **dict** - A dictionary which can include any keys, but must include the key 'loss' in the case of automatic optimization.
- **None** - In automatic optimization, this will skip to the next batch (but is not supported for multi-GPU, TPU, or DeepSpeed). For manual optimization, this has no special meaning, as returning the loss is not required.

In this step you'd normally do the forward pass and calculate the loss for a batch. You can also do fancier things like multiple forward passes or something model specific.

Example:

```
def training_step(self, batch, batch_idx):
    x, y, z = batch
    out = self.encoder(x)
    loss = self.loss(out, x)
    return loss
```

To use multiple optimizers, you can switch to 'manual optimization' and control their stepping:

```
def __init__(self):
    super().__init__()
    self.automatic_optimization = False

# Multiple optimizers (e.g.: GANs)
def training_step(self, batch, batch_idx):
    opt1, opt2 = self.optimizers()

    # do training_step with encoder
    ...
    opt1.step()
    # do training_step with decoder
    ...
    opt2.step()
```

Note

When `accumulate_grad_batches > 1`, the loss returned here will be automatically normalized by `accumulate_grad_batches` internally.

```
class dicee.models.transformers.LayerNorm(ndim, bias)
```

Bases: `torch.nn.Module`

LayerNorm but with an optional bias. PyTorch doesn't support simply `bias=False`

weight

bias

forward(input)

```
class dicee.models.transformers.SelfAttention(config)
```

Bases: `torch.nn.Module`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`c_attn`

`c_proj`

`attn_dropout`

`resid_dropout`

`n_head`

`n_embd`

`dropout`

`causal`

`flash = True`

`forward(x)`

```
class dicee.models.transformers.MLP(config)
```

Bases: `torch.nn.Module`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the submodules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

c_fc

gelu

c_proj

dropout

forward (*x*)

```
class dicee.models.transformers.Block(config)
```

Bases: `torch.nn.Module`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

```

(continues on next page)

```
def forward(self, x):
    x = F.relu(self.conv1(x))
    return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

ln_1

attn

ln_2

mlp

forward (*x*)

```
class dicee.models.transformers.GPTConfig
```

```
    block_size: int = 1024
```

```
    vocab_size: int = 50304
```

```
    n_layer: int = 12
```

```
    n_head: int = 12
```

```
    n_embd: int = 768
```

```
    dropout: float = 0.0
```

```
    bias: bool = False
```

```
    causal: bool = True
```

```
class dicee.models.transformers.GPT(config)
```

```
    Bases: torch.nn.Module
```

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

config

transformer

lm_head

get_num_params (*non_embedding=True*)

Return the number of parameters in the model. For non-embedding count (default), the position embeddings get subtracted. The token embeddings would too, except due to the parameter sharing these params are actually used as weights in the final layer, so we include them.

forward (*idx, targets=None*)

crop_block_size (*block_size*)

classmethod from_pretrained (*model_type, override_args=None*)

configure_optimizers (*weight_decay, learning_rate, betas, device_type*)

estimate_mfu (*fwdbwd_per_iter, dt*)

estimate model flops utilization (MFU) in units of A100 bfloat16 peak FLOPS

Classes

<i>ADOPT</i>	ADOPT Optimizer.
<i>BaseKGELightning</i>	Base class for all neural network modules.
<i>BaseKGE</i>	Base class for all neural network modules.

continues on next page

Table 41 – continued from previous page

<i>IdentityClass</i>	Base class for all neural network modules.
<i>BaseKGE</i>	Base class for all neural network modules.
<i>Block</i>	Base class for all neural network modules.
<i>DistMult</i>	Embedding Entities and Relations for Learning and Inference in Knowledge Bases
<i>TransE</i>	Translating Embeddings for Modeling
<i>Shallom</i>	A shallow neural model for relation prediction (https://arxiv.org/abs/2101.09090)
<i>Pyke</i>	A Physical Embedding Model for Knowledge Graphs
<i>CoKEConfig</i>	Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.
<i>CoKE</i>	Contextualized Knowledge Graph Embedding (CoKE) model.
<i>BaseKGE</i>	Base class for all neural network modules.
<i>ConEx</i>	Convolutional ComplEx Knowledge Graph Embeddings
<i>AConEx</i>	Additive Convolutional ComplEx Knowledge Graph Embeddings
<i>ComplEx</i>	Base class for all neural network modules.
<i>BaseKGE</i>	Base class for all neural network modules.
<i>IdentityClass</i>	Base class for all neural network modules.
<i>QMult</i>	Base class for all neural network modules.
<i>ConvQ</i>	Convolutional Quaternion Knowledge Graph Embeddings
<i>AConvQ</i>	Additive Convolutional Quaternion Knowledge Graph Embeddings
<i>BaseKGE</i>	Base class for all neural network modules.
<i>IdentityClass</i>	Base class for all neural network modules.
<i>OMult</i>	Base class for all neural network modules.
<i>ConvO</i>	Base class for all neural network modules.
<i>AConvO</i>	Additive Convolutional Octonion Knowledge Graph Embeddings
<i>Keci</i>	Base class for all neural network modules.
<i>CKeci</i>	Without learning dimension scaling
<i>DeCaL</i>	Base class for all neural network modules.
<i>KeciTransformer</i>	Keci with Transformer architecture.
<i>BaseKGE</i>	Base class for all neural network modules.
<i>PykeenKGE</i>	A class for using knowledge graph embedding models implemented in Pykeen
<i>BaseKGE</i>	Base class for all neural network modules.
<i>FMult</i>	Learning Knowledge Neural Graphs
<i>GFMult</i>	Learning Knowledge Neural Graphs
<i>FMult2</i>	Learning Knowledge Neural Graphs
<i>LFMult1</i>	Embedding with trigonometric functions. We represent all entities and relations in the complex number space as:
<i>LFMult</i>	Embedding with polynomial functions. We represent all entities and relations in the polynomial space as:
<i>DualE</i>	Dual Quaternion Knowledge Graph Embeddings (https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657)

Functions

```
quaternion_mul(→ Tuple[torch.Tensor, torch.Tensor, Perform quaternion multiplication
...))
quaternion_mul_with_unit_norm(*, Q_1, Q_2)
octonion_mul(*, O_1, O_2)
octonion_mul_norm(*, O_1, O_2)
```

Package Contents

```
class dicee.models.ADOPT (params: torch.optim.optimizer.ParamsT, lr: float | torch.Tensor = 0.001,
    betas: Tuple[float, float] = (0.9, 0.9999), eps: float = 1e-06,
    clip_lambda: Callable[[int], float] | None = lambda step: ..., weight_decay: float = 0.0,
    decouple: bool = False, *, foreach: bool | None = None, maximize: bool = False,
    capturable: bool = False, differentiable: bool = False, fused: bool | None = None)
```

Bases: `torch.optim.optimizer.Optimizer`

ADOPT Optimizer.

ADOPT is an adaptive learning rate optimization algorithm that combines momentum-based updates with adaptive per-parameter learning rates. It uses exponential moving averages of gradients and squared gradients, with gradient clipping for stability.

The algorithm performs the following key operations: 1. Normalizes gradients by the square root of the second moment estimate 2. Applies optional gradient clipping based on the training step 3. Updates parameters using momentum-smoothed normalized gradients 4. Supports decoupled weight decay (AdamW-style) or L2 regularization

Mathematical formulation:

$$m_t = \beta_1 * m_{t-1} + (1 - \beta_1) * \text{clip}(g_t / \sqrt{v_t}) \quad v_t = \beta_2 * v_{t-1} + (1 - \beta_2) * g_t^2 \quad \theta_t = \theta_{t-1} - \alpha * m_t$$

where:

- θ_t : parameter at step t
- g_t : gradient at step t
- m_t : first moment estimate (momentum)
- v_t : second moment estimate (variance)
- α : learning rate
- β_1, β_2 : exponential decay rates
- `clip()`: optional gradient clipping function

Reference:

Original implementation: <https://github.com/iShohei220/adopt>

Parameters

- **params** (*ParamsT*) – Iterable of parameters to optimize or dicts defining parameter groups.
- **lr** (*float or Tensor, optional*) – Learning rate. Can be a float or 1-element Tensor. Default: 1e-3
- **betas** (*Tuple[float, float], optional*) – Coefficients (β_1, β_2) for computing running averages of gradient and its square. β_1 controls momentum, β_2 controls variance. Default: (0.9, 0.9999)

- **eps** (*float, optional*) – Term added to denominator to improve numerical stability. Default: 1e-6
- **clip_lambda** (*Callable[[int], float], optional*) – Function that takes the step number and returns the gradient clipping threshold. Common choices: - lambda step: step**0.25 (default, gradually increases clipping threshold) - lambda step: 1.0 (constant clipping) - None (no clipping) Default: lambda step: step**0.25
- **weight_decay** (*float, optional*) – Weight decay coefficient (L2 penalty). Default: 0.0
- **decouple** (*bool, optional*) – If True, uses decoupled weight decay (AdamW-style), applying weight decay directly to parameters. If False, adds weight decay to gradients (L2 regularization). Default: False
- **foreach** (*bool, optional*) – If True, uses the faster foreach implementation for multi-tensor operations. Default: None (auto-select)
- **maximize** (*bool, optional*) – If True, maximizes parameters instead of minimizing. Useful for reinforcement learning. Default: False
- **capturable** (*bool, optional*) – If True, the optimizer is safe to capture in a CUDA graph. Requires learning rate as Tensor. Default: False
- **differentiable** (*bool, optional*) – If True, the optimization step can be differentiated. Useful for meta-learning. Default: False
- **fused** (*bool, optional*) – If True, uses fused kernel implementation (currently not supported). Default: None

Raises

- **ValueError** – If learning rate, epsilon, betas, or weight_decay are invalid.
- **RuntimeError** – If fused is enabled (not currently supported).
- **RuntimeError** – If lr is a Tensor with foreach=True and capturable=False.

Example

```
>>> # Basic usage
>>> optimizer = ADOPT(model.parameters(), lr=0.001)
>>> optimizer.zero_grad()
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With decoupled weight decay
>>> optimizer = ADOPT(model.parameters(), lr=0.001, weight_decay=0.01,
↳ decouple=True)
```

```
>>> # Custom gradient clipping
>>> optimizer = ADOPT(model.parameters(), clip_lambda=lambda step: max(1.0,
↳ step**0.5))
```

Note

- For most use cases, the default hyperparameters work well
- Consider using decouple=True for better generalization (similar to AdamW)

- The `clip_lambda` function helps stabilize training in early steps

`clip_lambda`

`__setstate__` (*state*)

Restore optimizer state from a checkpoint.

This method handles backward compatibility when loading optimizer state from older versions. It ensures all required fields are present with default values and properly converts step counters to tensors if needed.

Key responsibilities: 1. Set default values for newly added hyperparameters 2. Convert old-style scalar step counters to tensor format 3. Place step tensors on appropriate devices based on capturable/fused modes

Parameters

state (*dict*) – Optimizer state dictionary (typically from `torch.load()`).

Note

- This enables loading checkpoints saved with older ADOPT versions
- Step counters are converted to appropriate device/dtype for compatibility
- Capturable and fused modes require step tensors on parameter devices

`step` (*closure=None*)

Perform a single optimization step.

This method executes one iteration of the ADOPT optimization algorithm across all parameter groups. It orchestrates the following workflow:

1. Optionally evaluates a closure to recompute the loss (useful for algorithms like LBFGS or when loss needs multiple evaluations)
2. For each parameter group: - Collects parameters with gradients and their associated state - Extracts hyperparameters (betas, learning rate, etc.) - Calls the functional `adopt()` API to perform the actual update
3. Returns the loss value if a closure was provided

The functional API (`adopt()`) handles three execution modes: - Single-tensor: Updates one parameter at a time (default, JIT-compatible) - Multi-tensor (foreach): Batches operations for better performance - Fused: Uses fused CUDA kernels (not yet implemented)

Gradient scaling support: This method is compatible with automatic mixed precision (AMP) training. It can access `grad_scale` and `found_inf` attributes for gradient unscaling and inf/nan detection when used with `GradScaler`.

Parameters

closure (*Callable, optional*) – A callable that reevaluates the model and returns the loss. The closure should: - Enable gradients (`torch.enable_grad()`) - Compute forward pass - Compute loss - Compute backward pass - Return the loss value Example: `lambda: (loss := model(x), loss.backward(), loss)[-1]` Default: `None`

Returns

The loss value returned by the closure, or `None` if no closure was provided.

Return type
Optional[Tensor]

Example

```
>>> # Standard usage
>>> loss = criterion(model(input), target)
>>> loss.backward()
>>> optimizer.step()
```

```
>>> # With closure (e.g., for line search)
>>> def closure():
...     optimizer.zero_grad()
...     output = model(input)
...     loss = criterion(output, target)
...     loss.backward()
...     return loss
>>> loss = optimizer.step(closure)
```

Note

- Call `zero_grad()` before computing gradients for the next step
- CUDA graph capture is checked for safety when `capturable=True`
- The method is thread-safe for different parameter groups

```
class dicee.models.BaseKGE Lightning(*args, **kwargs)
```

Bases: `lightning.LightningModule`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

training_step_outputs = []

mem_of_model() → Dict

Size of model in MB and number of params

training_step (*batch*, *batch_idx=None*)

Here you compute and return the training loss and some additional metrics for e.g. the progress bar or logger.

Parameters

- **batch** – The output of your data iterable, normally a `DataLoader`.
- **batch_idx** – The index of this batch.
- **dataloader_idx** – The index of the dataloader that produced this batch. (only if multiple dataloaders used)

Returns

- **Tensor** – The loss tensor
- **dict** – A dictionary which can include any keys, but must include the key 'loss' in the case of automatic optimization.
- **None** – In automatic optimization, this will skip to the next batch (but is not supported for multi-GPU, TPU, or DeepSpeed). For manual optimization, this has no special meaning, as returning the loss is not required.

In this step you'd normally do the forward pass and calculate the loss for a batch. You can also do fancier things like multiple forward passes or something model specific.

Example:

```
def training_step(self, batch, batch_idx):
    x, y, z = batch
    out = self.encoder(x)
    loss = self.loss(out, x)
    return loss
```

To use multiple optimizers, you can switch to 'manual optimization' and control their stepping:

```
def __init__(self):
    super().__init__()
    self.automatic_optimization = False

# Multiple optimizers (e.g.: GANs)
def training_step(self, batch, batch_idx):
    opt1, opt2 = self.optimizers()
```

(continues on next page)

(continued from previous page)

```
# do training_step with encoder
...
opt1.step()
# do training_step with decoder
...
opt2.step()
```

Note

When `accumulate_grad_batches > 1`, the loss returned here will be automatically normalized by `accumulate_grad_batches` internally.

loss_function (*yhat_batch*: *torch.FloatTensor*, *y_batch*: *torch.FloatTensor*)

Parameters

- **yhat_batch**
- **y_batch**

on_train_epoch_end (*args, **kwargs)

Called in the training loop at the very end of the epoch.

To access all batch outputs at the end of the epoch, you can cache step outputs as an attribute of the `LightningModule` and access them in this hook:

```
class MyLightningModule(L.LightningModule):
    def __init__(self):
        super().__init__()
        self.training_step_outputs = []

    def training_step(self):
        loss = ...
        self.training_step_outputs.append(loss)
        return loss

    def on_train_epoch_end(self):
        # do something with all training_step outputs, for example:
        epoch_mean = torch.stack(self.training_step_outputs).mean()
        self.log("training_epoch_mean", epoch_mean)
        # free up the memory
        self.training_step_outputs.clear()
```

test_epoch_end (outputs: *List[Any]*)

test_dataloader () → None

An iterable or collection of iterables specifying test samples.

For more information about multiple dataloaders, see this section.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

 Warning

do not assign state in `prepare_data`

- `test()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

 Note

If you don't need a test dataset and a `test_step()`, you don't need to implement this method.

`val_dataloader()` → None

An iterable or collection of iterables specifying validation samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `fit()`
- `validate()`
- `prepare_data()`
- `setup()`

 Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware There is no need to set it yourself.

 Note

If you don't need a validation dataset and a `validation_step()`, you don't need to implement this method.

`predict_dataloader()` → None

An iterable or collection of iterables specifying prediction samples.

For more information about multiple dataloaders, see this section.

It's recommended that all data downloads and preparation happen in `prepare_data()`.

- `predict()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

Returns

A `torch.utils.data.DataLoader` or a sequence of them specifying prediction samples.

`train_dataloader()` → None

An iterable or collection of iterables specifying training samples.

For more information about multiple dataloaders, see this section.

The dataloader you return will not be reloaded unless you set **:param-ref:~lightning.pytorch.trainer.trainer.Trainer.reload_dataloaders_every_n_epochs** to a positive integer.

For data processing use the following pattern:

- download in `prepare_data()`
- process and split in `setup()`

However, the above are only necessary for distributed processing.

Warning

do not assign state in `prepare_data`

- `fit()`
- `prepare_data()`
- `setup()`

Note

Lightning tries to add the correct sampler for distributed and arbitrary hardware. There is no need to set it yourself.

configure_optimizers (*parameters=None*)

Choose what optimizers and learning-rate schedulers to use in your optimization. Normally you'd need one. But in the case of GANs or similar you might have multiple. Optimization with multiple optimizers only works in the manual optimization mode.

Returns

Any of these 6 options.

- **Single optimizer.**
- **List or Tuple** of optimizers.
- **Two lists** - The first list has multiple optimizers, and the second has multiple LR schedulers (or multiple `lr_scheduler_config`).
- **Dictionary**, with an "optimizer" key, and (optionally) a "lr_scheduler" key whose value is a single LR scheduler or `lr_scheduler_config`.
- **None** - Fit will run without any optimizer.

The `lr_scheduler_config` is a dictionary which contains the scheduler and its associated configuration. The default configuration is shown below.

```
lr_scheduler_config = {  
    # REQUIRED: The scheduler instance  
    "scheduler": lr_scheduler,  
    # The unit of the scheduler's step size, could also be 'step'.  
    # 'epoch' updates the scheduler on epoch end whereas 'step'  
    # updates it after a optimizer update.  
    "interval": "epoch",  
    # How many epochs/steps should pass between calls to  
    # `scheduler.step()`. 1 corresponds to updating the learning  
    # rate after every epoch/step.  
    "frequency": 1,  
    # Metric to monitor for schedulers like `ReduceLROnPlateau`  
    "monitor": "val_loss",  
    # If set to `True`, will enforce that the value specified 'monitor'  
    # is available when the scheduler is updated, thus stopping  
    # training if not found. If set to `False`, it will only produce a warning  
    "strict": True,  
    # If using the `LearningRateMonitor` callback to monitor the  
    # learning rate progress, this keyword can be used to specify  
    # a custom logged name  
    "name": None,  
}
```

When there are schedulers in which the `.step()` method is conditioned on a value, such as the `torch.optim.lr_scheduler.ReduceLROnPlateau` scheduler, Lightning requires that the `lr_scheduler_config` contains the keyword "monitor" set to the metric name that the scheduler should be conditioned on.

Metrics can be made available to monitor by simply logging it using `self.log('metric_to_track', metric_val)` in your `LightningModule`.

Note

Some things to know:

- Lightning calls `.backward()` and `.step()` automatically in case of automatic optimization.
- If a learning rate scheduler is specified in `configure_optimizers()` with key "interval" (default "epoch") in the scheduler configuration, Lightning will call the scheduler's `.step()` method automatically in case of automatic optimization.
- If you use 16-bit precision (`precision=16`), Lightning will automatically handle the optimizer.
- If you use `torch.optim.LBFGS`, Lightning handles the closure function automatically for you.
- If you use multiple optimizers, you will have to switch to 'manual optimization' mode and step them yourself.
- If you need to control how often the optimizer steps, override the `optimizer_step()` hook.

```
class dicee.models.BaseKGE(args: dict)
```

Bases: *BaseKGELightning*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

args

```

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

```

Parameters

$\mathbf{x} (B \times 2 \times T)$

forward_byte_pair_encoded_triple (*x*: *Tuple[torch.LongTensor, torch.LongTensor]*)
 byte pair encoded neural link predictors

Parameters

init_params_with_sanity_checking ()

forward (*x*: *torch.LongTensor* | *Tuple[torch.LongTensor, torch.LongTensor]*,
y_idx: *torch.LongTensor* = *None*)

Parameters

- **x**
- **y_idx**
- **ordered_bpe_entities**

forward_triples (*x*: *torch.LongTensor*) → *torch.Tensor*

Parameters
x

forward_k_vs_all (**args*, ***kwargs*)

forward_k_vs_sample (**args*, ***kwargs*)

get_triple_representation (*idx_hrt*)

get_head_relation_representation (*indexed_triple*)

get_sentence_representation (*x*: *torch.LongTensor*)

Parameters

- (**b** (*x* *shape*))
- **3**
- **t**)

get_bpe_head_and_relation_representation (*x*: *torch.LongTensor*)
 → *Tuple[torch.FloatTensor, torch.FloatTensor]*

Parameters

x (*B* × 2 × *T*)

get_embeddings () → *Tuple[numpy.ndarray, numpy.ndarray]*

class *dicee.models.IdentityClass* (*args=None*)

Bases: *torch.nn.Module*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

args = None

__call__ (*x*)

static forward (*x*)

class `dicdee.models.BaseKGE` (*args: dict*)

Bases: `BaseKGELightning`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training` (*bool*) – Boolean represents whether this module is in training or evaluation mode.

`args`

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

`loss`

`selected_optimizer = None`

`normalizer_class = None`

`normalize_head_entity_embeddings`

`normalize_relation_embeddings`

`normalize_tail_entity_embeddings`

`hidden_normalizer`

`param_init`

`input_dp_ent_real`

`input_dp_rel_real`

```

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

    Parameters
         $\mathbf{x}$  ( $B \times 2 \times T$ )

forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])
    byte pair encoded neural link predictors

    Parameters
    -----

init_params_with_sanity_checking()

forward (x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
        y_idx: torch.LongTensor = None)

    Parameters

    •  $\mathbf{x}$ 

    •  $\mathbf{y\_idx}$ 

    • ordered_bpe_entities

forward_triples (x: torch.LongTensor) → torch.Tensor

    Parameters
         $\mathbf{x}$ 

forward_k_vs_all (*args, **kwargs)

forward_k_vs_sample (*args, **kwargs)

get_triple_representation (idx_hrt)

get_head_relation_representation (indexed_triple)

get_sentence_representation (x: torch.LongTensor)

    Parameters

    • ( $\mathbf{b}$  ( $x$  shape))

    • 3

    •  $\mathbf{t}$ )

get_bpe_head_and_relation_representation (x: torch.LongTensor)
    → Tuple[torch.FloatTensor, torch.FloatTensor]

    Parameters
         $\mathbf{x}$  ( $B \times 2 \times T$ )

```

`get_embeddings()` → Tuple[numpy.ndarray, numpy.ndarray]

`class dicee.models.Block(config)`

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`ln_1`

`attn`

`ln_2`

`mlp`

`forward(x)`

`class dicee.models.DistMult(args)`

Bases: `dicee.models.base_model.BaseKGE`

Embedding Entities and Relations for Learning and Inference in Knowledge Bases <https://arxiv.org/abs/1412.6575>

`name = 'DistMult'`

k_vs_all_score (*emb_h: torch.FloatTensor, emb_r: torch.FloatTensor, emb_E: torch.FloatTensor*)

Parameters

- **emb_h**
- **emb_r**
- **emb_E**

forward_k_vs_all (*x: torch.LongTensor*)

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*)

score (*h, r, t*)

class `dicee.models.TransE` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

Translating Embeddings for Modeling Multi-relational Data <https://proceedings.neurips.cc/paper/2013/file/1cecc7a77928ca8133fa24680a88d2f9-Paper.pdf>

name = 'TransE'

margin = 4

score (*head_ent_emb, rel_ent_emb, tail_ent_emb*)

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

class `dicee.models.Shallom` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

A shallow neural model for relation prediction (<https://arxiv.org/abs/2101.09090>)

name = 'Shallom'

shallom

get_embeddings () → Tuple[numpy.ndarray, None]

forward_k_vs_all (*x*) → torch.FloatTensor

forward_triples (*x*) → torch.FloatTensor

Parameters

x

Returns

class `dicee.models.Pyke` (*args*)

Bases: `dicee.models.base_model.BaseKGE`

A Physical Embedding Model for Knowledge Graphs

name = 'Pyke'

dist_func

margin = 1.0

forward_triples (*x*: *torch.LongTensor*)

Parameters

x

class `dicee.models.CoKEConfig`

Configuration for the CoKE (Contextualized Knowledge Graph Embedding) model.

block_size

Sequence length for transformer (3 for triples: head, relation, tail)

vocab_size

Total vocabulary size (num_entities + num_relations)

n_layer

Number of transformer layers

n_head

Number of attention heads per layer

n_embd

Embedding dimension (set to match model embedding_dim)

dropout

Dropout rate applied throughout the model

bias

Whether to use bias in linear layers

causal

Whether to use causal masking (False for bidirectional attention)

block_size: `int = 3`

vocab_size: `int = None`

n_layer: `int = 6`

n_head: `int = 8`

n_embd: `int = None`

dropout: `float = 0.3`

bias: `bool = True`

causal: `bool = False`

class `dicee.models.CoKE` (*args*, *config*: *CoKEConfig* = *CoKEConfig*())

Bases: `dicee.models.base_model.BaseKGE`

Contextualized Knowledge Graph Embedding (CoKE) model. Based on: <https://arxiv.org/pdf/1911.02168>.

CoKE uses a transformer encoder to learn contextualized representations of entities and relations. For link prediction, it predicts masked elements in (head, relation, tail) triples using bidirectional attention, similar to BERT's masked language modeling approach.

The model creates a sequence [head_emb, relation_emb, mask_emb], adds positional embeddings, and processes it through transformer layers to predict the tail entity.

```

name = 'CoKE'

config

pos_emb

mask_emb

blocks

ln_f

coke_dropout

forward_k_vs_all(x: torch.Tensor)

score(emb_h, emb_r, emb_t)

forward_k_vs_sample(x: torch.LongTensor, target_entity_idx: torch.LongTensor)

```

```
class dicee.models.BaseKGE(args: dict)
```

Bases: *BaseKGELightning*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

args

```

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

```

Parameters

$\mathbf{x} (B \times 2 \times T)$

```

forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])
    byte pair encoded neural link predictors

    Parameters
    -----

init_params_with_sanity_checking()

forward (x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
         y_idx: torch.LongTensor = None)

    Parameters
    • x
    • y_idx
    • ordered_bpe_entities

forward_triples (x: torch.LongTensor) → torch.Tensor

    Parameters
    • x

forward_k_vs_all (*args, **kwargs)

forward_k_vs_sample (*args, **kwargs)

get_triple_representation (idx_hrt)

get_head_relation_representation (indexed_triple)

get_sentence_representation (x: torch.LongTensor)

    Parameters
    • (b (x shape)
    • 3
    • t)

get_bpe_head_and_relation_representation (x: torch.LongTensor)
    → Tuple[torch.FloatTensor, torch.FloatTensor]

    Parameters
    • x (B × 2 × T)

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

class dicee.models.ConEx (args)
    Bases: dicee.models.base_model.BaseKGE
    Convolutional ComplEx Knowledge Graph Embeddings
    name = 'ConEx'

    conv2d

    fc_num_input

    fc1

```

```

norm_fc1

bn_conv2d

feature_map_dropout

residual_convolution (C_1: Tuple[torch.Tensor, torch.Tensor],
                      C_2: Tuple[torch.Tensor, torch.Tensor]) → torch.FloatTensor
    Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors
    that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds
    complex-valued embeddings :return:

forward_k_vs_all (x: torch.Tensor) → torch.FloatTensor

forward_triples (x: torch.Tensor) → torch.FloatTensor

    Parameters
    x

forward_k_vs_sample (x: torch.Tensor, target_entity_idx: torch.Tensor)

class dicee.models.AConEx (args)
    Bases: dicee.models.base_model.BaseKGE
    Additive Convolutional ComplEx Knowledge Graph Embeddings
    name = 'AConEx'

    conv2d

    fc_num_input

    fc1

    norm_fc1

    bn_conv2d

    feature_map_dropout

    residual_convolution (C_1: Tuple[torch.Tensor, torch.Tensor],
                          C_2: Tuple[torch.Tensor, torch.Tensor]) → torch.FloatTensor
        Compute residual score of two complex-valued embeddings. :param C_1: a tuple of two pytorch tensors
        that corresponds complex-valued embeddings :param C_2: a tuple of two pytorch tensors that corresponds
        complex-valued embeddings :return:

    forward_k_vs_all (x: torch.Tensor) → torch.FloatTensor

    forward_triples (x: torch.Tensor) → torch.FloatTensor

    Parameters
    x

    forward_k_vs_sample (x: torch.Tensor, target_entity_idx: torch.Tensor)

class dicee.models.ComplEx (args)
    Bases: dicee.models.base_model.BaseKGE
    Base class for all neural network modules.
    Your models should also subclass this class.

```

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'Complex'

static score (*head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor, tail_ent_emb: torch.FloatTensor*)

static k_vs_all_score (*emb_h: torch.FloatTensor, emb_r: torch.FloatTensor, emb_E: torch.FloatTensor*)

Parameters

- **emb_h**
- **emb_r**
- **emb_E**

forward_k_vs_all (*x: torch.LongTensor*) → torch.FloatTensor

forward_k_vs_sample (*x: torch.LongTensor, target_entity_idx: torch.LongTensor*)

dicee.models.quaternion_mul (*, *Q_1, Q_2*)
→ Tuple[torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor]

Perform quaternion multiplication :param Q_1: :param Q_2: :return:

class dicee.models.**BaseKGE** (*args: dict*)

Bases: *BaseKGELightning*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None


```

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

    Parameters
         $\mathbf{x}$  ( $B \times 2 \times T$ )

forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])
    byte pair encoded neural link predictors

    Parameters
    -----

init_params_with_sanity_checking()

forward (x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
        y_idx: torch.LongTensor = None)

    Parameters

        •  $\mathbf{x}$ 

        •  $\mathbf{y\_idx}$ 

        • ordered_bpe_entities

forward_triples (x: torch.LongTensor)  $\rightarrow$  torch.Tensor

    Parameters
         $\mathbf{x}$ 

forward_k_vs_all (*args, **kwargs)

forward_k_vs_sample (*args, **kwargs)

```

```

get_triple_representation(idx_hrt)

get_head_relation_representation(indexed_triple)

get_sentence_representation(x: torch.LongTensor)

```

Parameters

- **b** (x shape)
- 3
- **t**)

```

get_bpe_head_and_relation_representation(x: torch.LongTensor)
→ Tuple[torch.FloatTensor, torch.FloatTensor]

```

Parameters

x ($B \times 2 \times T$)

```

get_embeddings() → Tuple[numpy.ndarray, numpy.ndarray]

```

```

class dicee.models.IdentityClass(args=None)

```

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

```
args = None
```

```
__call__(x)
```

```
static forward(x)
```

```
dicee.models.quaternion_mul_with_unit_norm(*, Q_1, Q_2)
```

```
class dicee.models.QMult(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

```
name = 'QMult'
```

```
explicit = True
```

```
quaternion_multiplication_followed_by_inner_product (h, r, t)
```

Parameters

- `h` – shape: `(*batch_dims, dim)` The head representations.
- `r` – shape: `(*batch_dims, dim)` The head representations.
- `t` – shape: `(*batch_dims, dim)` The tail representations.

Returns

Triple scores.

static quaternion_normalizer (x : *torch.FloatTensor*) $\rightarrow$ *torch.FloatTensor*

Normalize the length of relation vectors, if the forward constraint has not been applied yet.

Absolute value of a quaternion

$$|a + bi + cj + dk| = \sqrt{a^2 + b^2 + c^2 + d^2}$$

L2 norm of quaternion vector:

$$\|x\|^2 = \sum_{i=1}^d |x_i|^2 = \sum_{i=1}^d (x_i.re^2 + x_i.im_1^2 + x_i.im_2^2 + x_i.im_3^2)$$

Parameters

x – The vector.

Returns

The normalized vector.

score (*head_ent_emb*: *torch.FloatTensor*, *rel_ent_emb*: *torch.FloatTensor*,
tail_ent_emb: *torch.FloatTensor*)

k_vs_all_score (*bpe_head_ent_emb*, *bpe_rel_ent_emb*, *E*)

Parameters

- **bpe_head_ent_emb**
- **bpe_rel_ent_emb**
- **E**

forward_k_vs_all (*x*)

Parameters

x

forward_k_vs_sample (*x*, *target_entity_idx*)

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e.,
[score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and
relations => shape (size of batch,| Entities|)

class *dicee.models.ConvQ* (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Convolutional Quaternion Knowledge Graph Embeddings

name = 'ConvQ'

entity_embeddings

relation_embeddings

conv2d

fc_num_input

fc1

```

bn_conv1

bn_conv2

feature_map_dropout

residual_convolution(Q_1, Q_2)

forward_triples(indexed_triple: torch.Tensor) → torch.Tensor

```

Parameters

x

```
forward_k_vs_all(x: torch.Tensor)
```

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

```
class dicee.models.AConvQ(args)
```

Bases: *dicee.models.base_model.BaseKGE*

Additive Convolutional Quaternion Knowledge Graph Embeddings

name = 'AConvQ'

entity_embeddings

relation_embeddings

conv2d

fc_num_input

fc1

bn_conv1

bn_conv2

feature_map_dropout

residual_convolution(Q_1, Q_2)

forward_triples(indexed_triple: torch.Tensor) → torch.Tensor

Parameters

x

```
forward_k_vs_all(x: torch.Tensor)
```

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

```
class dicee.models.BaseKGE(args: dict)
```

Bases: *BaseKGELightning*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`args`

`embedding_dim = None`

`num_entities = None`

`num_relations = None`

`num_tokens = None`

`learning_rate = None`

`apply_unit_norm = None`

`input_dropout_rate = None`

`hidden_dropout_rate = None`

`optimizer_name = None`

`feature_map_dropout_rate = None`

`kernel_size = None`

`num_of_output_channels = None`

`weight_decay = None`

`loss`

```

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

    Parameters
         $\mathbf{x}$  ( $B \times 2 \times T$ )

forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])
    byte pair encoded neural link predictors

    Parameters
    -----

init_params_with_sanity_checking ()

forward (x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
        y_idx: torch.LongTensor = None)

    Parameters

        •  $\mathbf{x}$ 

        •  $\mathbf{y\_idx}$ 

        • ordered_bpe_entities

forward_triples (x: torch.LongTensor)  $\rightarrow$  torch.Tensor

    Parameters
         $\mathbf{x}$ 

forward_k_vs_all (*args, **kwargs)

forward_k_vs_sample (*args, **kwargs)

get_triple_representation (idx_hrt)

```

`get_head_relation_representation(indexed_triple)`

`get_sentence_representation(x: torch.LongTensor)`

Parameters

- **b** (x shape)
- 3
- **t**)

`get_bpe_head_and_relation_representation(x: torch.LongTensor)`
→ Tuple[torch.FloatTensor, torch.FloatTensor]

Parameters

x ($B \times 2 \times T$)

`get_embeddings()` → Tuple[numpy.ndarray, numpy.ndarray]

`class dicee.models.IdentityClass(args=None)`

Bases: torch.nn.Module

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

args = None


```
__call__(x)
```

```
static forward(x)
```

```
dicee.models.octonion_mul(*, O_1, O_2)
```

```
dicee.models.octonion_mul_norm(*, O_1, O_2)
```

```
class dicee.models.OMult(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

```
name = 'OMult'
```

```
static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4,
                           emb_rel_e5, emb_rel_e6, emb_rel_e7)
```

```
score(head_ent_emb: torch.FloatTensor, rel_ent_emb: torch.FloatTensor,
      tail_ent_emb: torch.FloatTensor)
```

```
k_vs_all_score(bpe_head_ent_emb, bpe_rel_ent_emb, E)
```

```
forward_k_vs_all(x)
```

Completed. Given a head entity and a relation (h,r), we compute scores for all possible triples,i.e., $[\text{score}(h,r,x) | x \text{ in Entities}] \Rightarrow [0.0,0.1,\dots,0.8]$, shape= $\Rightarrow (1, \text{Entities!})$ Given a batch of head entities and relations $\Rightarrow$ shape (size of batch,| Entities!)

```
class dicee.models.ConvO(args: dict)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

`training (bool)` – Boolean represents whether this module is in training or evaluation mode.

`name = 'ConvO'`

`conv2d`

`fc_num_input`

`fc1`

`bn_conv2d`

`norm_fc1`

`feature_map_dropout`

`static octonion_normalizer(emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4, emb_rel_e5, emb_rel_e6, emb_rel_e7)`

`residual_convolution(O_1, O_2)`

forward_triples (*x: torch.Tensor*) → torch.Tensor

Parameters

x

forward_k_vs_all (*x: torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

```
class dicee.models.AConvO(args: dict)
```

Bases: *dicee.models.base_model.BaseKGE*

Additive Convolutional Octonion Knowledge Graph Embeddings

name = 'AConvO'

conv2d

fc_num_input

fc1

bn_conv2d

norm_fc1

feature_map_dropout

static octonion_normalizer (*emb_rel_e0, emb_rel_e1, emb_rel_e2, emb_rel_e3, emb_rel_e4, emb_rel_e5, emb_rel_e6, emb_rel_e7*)

residual_convolution (*O_1, O_2*)

forward_triples (*x: torch.Tensor*) → torch.Tensor

Parameters

x

forward_k_vs_all (*x: torch.Tensor*)

Given a head entity and a relation (h,r), we compute scores for all entities. [score(h,r,x)|x in Entities] => [0.0,0.1,...,0.8], shape=> (1, **Entities**) Given a batch of head entities and relations => shape (size of batch,|Entities|)

```
class dicee.models.Keci(args)
```

Bases: *dicee.models.base_model.BaseKGE*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
```

(continues on next page)

(continued from previous page)

```
super().__init__()
self.conv1 = nn.Conv2d(1, 20, 5)
self.conv2 = nn.Conv2d(20, 20, 5)

def forward(self, x):
    x = F.relu(self.conv1(x))
    return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'Keci'

p

q

r

requires_grad_for_interactions = True

compute_sigma_pp (*hp, rp*)

Compute $\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{k=i+1}^p (h_i r_k - h_k r_i) e_i e_k$

σ_{pp} captures the interactions between along p bases For instance, let $p = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

 for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., $e_1 e_1, e_1 e_2, e_1 e_3$,

$e_2 e_1, e_2 e_2, e_2 e_3, e_3 e_1, e_3 e_2, e_3 e_3$

Then select the triangular matrix without diagonals: $e_1 e_2, e_1 e_3, e_2 e_3$.

compute_sigma_qq (*hq, rq*)

Compute $\sigma_{qq} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_j r_k - h_k r_j) e_j e_k \sigma_{qq}$ captures the interactions between along q bases For instance, let $q = e_1, e_2, e_3$, we compute interactions between $e_1 e_2, e_1 e_3$, and $e_2 e_3$ This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

 for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

```
sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))
```

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p , e.g., e_1e_1 , e_1e_2 , e_1e_3 ,

e_2e_1 , e_2e_2 , e_2e_3 , e_3e_1 , e_3e_2 , e_3e_3

Then select the triangular matrix without diagonals: e_1e_2 , e_1e_3 , e_2e_3 .

```
compute_sigma_pq(*, hp, hq, rp, rq)
```

```
sum_{i=1}^p sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j
```

```
results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):
```

```
    for j in range(q):
```

```
        sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]
```

```
    print(sigma_pq.shape)
```

```
apply_coefficients(hp, hq, rp, rq)
```

Multiplying a base vector with its scalar coefficient

```
clifford_multiplication(h0, hp, hq, r0, rp, rq)
```

Compute our CL multiplication

```
h = h_0 + sum_{i=1}^p h_i e_i + sum_{j=p+1}^{p+q} h_j e_j
r = r_0 + sum_{i=1}^p r_i e_i + sum_{j=p+1}^{p+q} r_j e_j
```

$e_i^2 = +1$ for $i \leq p$ $e_j^2 = -1$ for $p < j \leq p+q$ $e_i e_j = -e_j e_i$ for i

e_j

$h r = \text{sigma}_0 + \text{sigma}_p + \text{sigma}_q + \text{sigma}_{\{pp\}} + \text{sigma}_{\{q\}} + \text{sigma}_{\{pq\}}$ where

(1) $\text{sigma}_0 = h_0 r_0 + \sum_{i=1}^p (h_0 r_i - h_i r_0) e_i - \sum_{j=p+1}^{p+q} (h_j r_j) e_j$

(2) $\text{sigma}_p = \sum_{i=1}^p (h_0 r_i + h_i r_0) e_i$

(3) $\text{sigma}_q = \sum_{j=p+1}^{p+q} (h_0 r_j + h_j r_0) e_j$

(4) $\text{sigma}_{\{pp\}} = \sum_{i=1}^p \sum_{k=i+1}^p (h_i r_k - h_k r_i) e_i e_k$

(5) $\text{sigma}_{\{qq\}} = \sum_{j=1}^{p+q-1} \sum_{k=j+1}^{p+q} (h_j r_k - h_k r_j) e_j e_k$

(6) $\text{sigma}_{\{pq\}} = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$

```
construct_cl_multivector(x: torch.FloatTensor, r: int, p: int, q: int)
```

```
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]
```

Construct a batch of multivectors $\text{Cl}_{\{p,q\}}(\mathbb{R}^d)$

Parameter

x : torch.FloatTensor with (n,d) shape

returns

- **a0** (torch.FloatTensor with (n,r) shape)
- **ap** (torch.FloatTensor with (n,r,p) shape)
- **aq** (torch.FloatTensor with (n,r,q) shape)

```
forward_k_vs_with_explicit(x: torch.Tensor)
```

```
k_vs_all_score(bpe_head_ent_emb, bpe_rel_ent_emb, E)
```

forward_k_vs_all (*x*: torch.Tensor) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations $\mathbb{R}^d$.
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this function are identical Parameter ——— *x*: torch.LongTensor with (n,2) shape :rtype: torch.FloatTensor with (n, **IE**) shape

construct_batch_selected_cl_multivector (*x*: torch.FloatTensor, *r*: int, *p*: int, *q*: int)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Construct a batch of batchs multivectors $Cl_{\{p,q\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,k, d) shape

returns

- **a0** (torch.FloatTensor with (n,k, m) shape)
- **ap** (torch.FloatTensor with (n,k, m, p) shape)
- **aq** (torch.FloatTensor with (n,k, m, q) shape)

forward_k_vs_sample (*x*: torch.LongTensor, *target_entity_idx*: torch.LongTensor) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,2) shape

target_entity_idx: torch.LongTensor with (n, k) shape k denotes the selected number of examples.

rtype

torch.FloatTensor with (n, k) shape

score (*h*, *r*, *t*)

forward_triples (*x*: torch.Tensor) → torch.FloatTensor

Parameter

x: torch.LongTensor with (n,3) shape

rtype

torch.FloatTensor with (n) shape

class dicee.models.CKeci (*args*)

Bases: *Keci*

Without learning dimension scaling

name = 'CKeci'

requires_grad_for_interactions = False

```
class dicee.models.DeCaL(args)
```

Bases: `dicee.models.base_model.BaseKGE`

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the submodules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

name = 'DeCaL'

entity_embeddings

relation_embeddings

p

q

r

re

forward_triples (*x: torch.Tensor*) $\rightarrow$ torch.FloatTensor

Parameter

x: torch.LongTensor with (n,) shape

rtype

torch.FloatTensor with (n) shape

cl_pqr (*a: torch.tensor*) → torch.tensor

Input: tensor(batch_size, emb_dim) → output: tensor with 1+p+q+r components with size (batch_size, emb_dim/(1+p+q+r)) each.

1) takes a tensor of size (batch_size, emb_dim), split it into 1 + p + q + r components, hence 1+p+q+r must be a divisor of the emb_dim. 2) Return a list of the 1+p+q+r components vectors, each are tensors of size (batch_size, emb_dim/(1+p+q+r))

compute_sigmas_single (*list_h_emb, list_r_emb, list_t_emb*)

here we compute all the sums with no others vectors interaction taken with the scalar product with t, that is,

$$s0 = h_0 r_0 t_0 s1 = \sum_{i=1}^p h_i r_i t_0 s2 = \sum_{j=p+1}^{p+q} h_j r_j t_0 s3 = \sum_{i=1}^q (h_0 r_i t_i + h_i r_0 t_i) s4 = \sum_{i=p+1}^{p+q} (h_0 r_i t_i + h_i r_0 t_i) s5 = \sum_{i=p+q+1}^{p+q+r} (h_0 r_i t_i + h_i r_0 t_i)$$

and return:

$$sigma_0 t = \sigma_0 \cdot t_0 = s0 + s1 - s2 s3, s4 \text{ and } s5$$

compute_sigmas_multivect (*list_h_emb, list_r_emb*)

Here we compute and return all the sums with vectors interaction for the same and different bases.

For same bases vectors interaction we have

$$\sigma_p p = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (h_i r_{i'} - h_{i'} r_i) (model \text{ the interactions between } e_i \text{ and } e_{i'} \text{ for } 1 \leq i, i' \leq p) \sigma_q q = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (h_j r_{j'} - h_{j'} r_j)$$

For different base vector interactions, we have

$$\sigma_p q = \sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) (interactions \text{ between } e_i \text{ and } e_j \text{ for } 1 \leq i \leq p \text{ and } p+1 \leq j \leq p+q) \sigma_p r = \sum_{i=1}^p \sum_{j=p+q+1}^{p+q+r} (h_i r_j - h_j r_i)$$

forward_k_vs_all (*x: torch.Tensor*) → torch.FloatTensor

Kvsall training

- (1) Retrieve real-valued embedding vectors for heads and relations
- (2) Construct head entity and relation embeddings according to $Cl_{\{p,q,r\}}(\mathbb{R}^d)$.
- (3) Perform Cl multiplication
- (4) Inner product of (3) and all entity embeddings

forward_k_vs_with_explicit and this functions are identical Parameter ——— x: torch.LongTensor with (n,) shape :rtype: torch.FloatTensor with (n, **IE**) shape

apply_coefficients (*h0, hp, hq, hk, r0, rp, rq, rk*)

Multiplying a base vector with its scalar coefficient

construct_cl_multivector (*x: torch.FloatTensor, re: int, p: int, q: int, r: int*)
→ tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]

Construct a batch of multivectors $Cl_{\{p,q,r\}}(\mathbb{R}^d)$

Parameter

x: torch.FloatTensor with (n,d) shape

returns

- **a0** (*torch.FloatTensor*)
- **ap** (*torch.FloatTensor*)
- **aq** (*torch.FloatTensor*)
- **ar** (*torch.FloatTensor*)

compute_sigma_pp (*hp, rp*)

Compute .. math:

$$\sigma_{pp} = \sum_{i=1}^{p-1} \sum_{i'=i+1}^p (x_{iy_{i'}} - x_{i'}y_i)$$

σ_{pp} captures the interactions between along p bases For instance, let p e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3, and e_2 e_3 This can be implemented with a nested two for loops

results = [] for i in range(p - 1):

for k in range(i + 1, p):

 results.append(hp[:, :, i] * rp[:, :, k] - hp[:, :, k] * rp[:, :, i])

sigma_pp = torch.stack(results, dim=2) assert sigma_pp.shape == (b, r, int((p * (p - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_qq (*hq, rq*)

Compute

$$\sigma_{qq}^* = \sum_{j=p+1}^{p+q-1} \sum_{j'=j+1}^{p+q} (x_j y_{j'} - x_{j'} y_j) E q.16$$

σ_{qq} captures the interactions between along q bases For instance, let q e_1, e_2, e_3, we compute interactions between e_1 e_2, e_1 e_3, and e_2 e_3 This can be implemented with a nested two for loops

results = [] for j in range(q - 1):

for k in range(j + 1, q):

 results.append(hq[:, :, j] * rq[:, :, k] - hq[:, :, k] * rq[:, :, j])

sigma_qq = torch.stack(results, dim=2) assert sigma_qq.shape == (b, r, int((q * (q - 1)) / 2))

Yet, this computation would be quite inefficient. Instead, we compute interactions along all p, e.g., e1e1, e1e2, e1e3,

e2e1, e2e2, e2e3, e3e1, e3e2, e3e3

Then select the triangular matrix without diagonals: e1e2, e1e3, e2e3.

compute_sigma_rr (*hk, rk*)

$$\sigma_{rr}^* = \sum_{k=p+q+1}^{p+q+r-1} \sum_{k'=k+1}^p (x_k y_{k'} - x_{k'} y_k)$$

compute_sigma_pq(* , hp, hq, rp, rq)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

compute_sigma_pr(* , hp, hk, rp, rk)

Compute

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

compute_sigma_qr(* , hq, hk, rq, rk)

$$\sum_{i=1}^p \sum_{j=p+1}^{p+q} (h_i r_j - h_j r_i) e_i e_j$$

results = [] sigma_pq = torch.zeros(b, r, p, q) for i in range(p):

for j in range(q):

 sigma_pq[:, :, i, j] = hp[:, :, i] * rq[:, :, j] - hq[:, :, j] * rp[:, :, i]

print(sigma_pq.shape)

class dicee.models.**KeciTransformer**(args)

Bases: [Keci](#)

Keci with Transformer architecture.

Concatenates h0, hp, hq, r0, rp, rq into a single embedding vector and processes through transformer.

name = 'KeciTransformer'

input_dim

seq_len = 1

transformer

lm_head

forward_k_vs_all(x: torch.Tensor) → torch.FloatTensor

Kvsall training

Parameter

x: torch.LongTensor with (n,2) shape

rtype

torch.FloatTensor with (n, **IE**) shape

class dicee.models.**BaseKGE** (args: dict)

Bases: *BaseKGELightning*

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (*bool*) – Boolean represents whether this module is in training or evaluation mode.

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

```

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

    Parameters
         $\mathbf{x} (B \times 2 \times T)$ 

forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])
    byte pair encoded neural link predictors

    Parameters
    -----

init_params_with_sanity_checking()

forward (x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
        y_idx: torch.LongTensor = None)

    Parameters

    •  $\mathbf{x}$ 

    •  $\mathbf{y\_idx}$ 

```

```

        • ordered_bpe_entities
forward_triples (x: torch.LongTensor) → torch.Tensor

    Parameters
    x

forward_k_vs_all (*args, **kwargs)

forward_k_vs_sample (*args, **kwargs)

get_triple_representation (idx_hrt)

get_head_relation_representation (indexed_triple)

get_sentence_representation (x: torch.LongTensor)

    Parameters
        • (b (x shape)
        • 3
        • t)

get_bpe_head_and_relation_representation (x: torch.LongTensor)
    → Tuple[torch.FloatTensor, torch.FloatTensor]

    Parameters
        x (B × 2 × T)

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

class dicee.models.PykeenKGE (args: dict)
    Bases: dicee.models.base_model.BaseKGE

    A class for using knowledge graph embedding models implemented in Pykeen

    Notes: Pykeen_DistMult: C Pykeen_ComplEx: Pykeen_QuatE: Pykeen_MuRE: Pykeen_CP: Pykeen_HolE: Py-
    keen_HolE: Pykeen_HolE: Pykeen_TransD: Pykeen_TransE: Pykeen_TransF: Pykeen_TransH: Pykeen_TransR:

    model_kwargs

    name

    model

    loss_history = []

    args

    entity_embeddings = None

    relation_embeddings = None

forward_k_vs_all (x: torch.LongTensor)
    # => Explicit version by this we can apply bn and dropout
    # (1) Retrieve embeddings of heads and relations + apply Dropout & Normalization if given. h, r =
    self.get_head_relation_representation(x) # (2) Reshape (1). if self.last_dim > 0:
        h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim,
        self.last_dim)

```

```

# (3) Reshape all entities. if self.last_dim > 0:
    t = self.entity_embeddings.weight.reshape(self.num_entities, self.embedding_dim, self.last_dim)

else:
    t = self.entity_embeddings.weight

# (4) Call the score_t from interactions to generate triple scores. return self.interaction.score_t(h=h, r=r,
all_entities=t, slice_size=1)

forward_triples (x: torch.LongTensor) → torch.FloatTensor
# => Explicit version by this we can apply bn and dropout

# (1) Retrieve embeddings of heads, relations and tails and apply Dropout & Normalization if given. h, r, t =
self.get_triple_representation(x) # (2) Reshape (1). if self.last_dim > 0:
    h = h.reshape(len(x), self.embedding_dim, self.last_dim) r = r.reshape(len(x), self.embedding_dim,
    self.last_dim) t = t.reshape(len(x), self.embedding_dim, self.last_dim)

# (3) Compute the triple score return self.interaction.score(h=h, r=r, t=t, slice_size=None, slice_dim=0)

abstractmethod forward_k_vs_sample (x: torch.LongTensor, target_entity_idx)

```

class dicee.models.BaseKGE (args: dict)

Bases: BaseKGELightning

Base class for all neural network modules.

Your models should also subclass this class.

Modules can also contain other Modules, allowing them to be nested in a tree structure. You can assign the sub-modules as regular attributes:

```

import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))

```

Submodules assigned in this way will be registered, and will also have their parameters converted when you call `to()`, etc.

Note

As per the example above, an `__init__()` call to the parent class must be made before assignment on the child.

Variables

training (bool) – Boolean represents whether this module is in training or evaluation mode.

```

args

embedding_dim = None

num_entities = None

num_relations = None

num_tokens = None

learning_rate = None

apply_unit_norm = None

input_dropout_rate = None

hidden_dropout_rate = None

optimizer_name = None

feature_map_dropout_rate = None

kernel_size = None

num_of_output_channels = None

weight_decay = None

loss

selected_optimizer = None

normalizer_class = None

normalize_head_entity_embeddings

normalize_relation_embeddings

normalize_tail_entity_embeddings

hidden_normalizer

param_init

input_dp_ent_real

input_dp_rel_real

hidden_dropout

loss_history = []

byte_pair_encoding

max_length_subword_tokens

block_size

forward_byte_pair_encoded_k_vs_all (x: torch.LongTensor)

```

Parameters

$\mathbf{x} (B \times 2 \times T)$

```

forward_byte_pair_encoded_triple (x: Tuple[torch.LongTensor, torch.LongTensor])
    byte pair encoded neural link predictors

    Parameters
    -----

init_params_with_sanity_checking()

forward (x: torch.LongTensor | Tuple[torch.LongTensor, torch.LongTensor],
         y_idx: torch.LongTensor = None)

    Parameters
    -----
    • x
    • y_idx
    • ordered_bpe_entities

forward_triples (x: torch.LongTensor) → torch.Tensor

    Parameters
    -----
    • x

forward_k_vs_all (*args, **kwargs)

forward_k_vs_sample (*args, **kwargs)

get_triple_representation (idx_hrt)

get_head_relation_representation (indexed_triple)

get_sentence_representation (x: torch.LongTensor)

    Parameters
    -----
    • (b (x shape)
    • 3
    • t)

get_bpe_head_and_relation_representation (x: torch.LongTensor)
    → Tuple[torch.FloatTensor, torch.FloatTensor]

    Parameters
    -----
    • x (B × 2 × T)

get_embeddings () → Tuple[numpy.ndarray, numpy.ndarray]

class dicee.models.FMult (args)
    Bases: dicee.models.base_model.BaseKGE
    Learning Knowledge Neural Graphs
    name = 'FMult'
    entity_embeddings
    relation_embeddings
    k

```



```

num_sample = 50

gamma

roots

weights

compute_func (weights: torch.FloatTensor, x) → torch.FloatTensor

chain_func (weights, x: torch.FloatTensor)

forward_triples (idx_triple: torch.Tensor) → torch.Tensor

```

Parameters

x

```

class dicee.models.GFMult (args)
    Bases: dicee.models.base_model.BaseKGE
    Learning Knowledge Neural Graphs
    name = 'GFMult'
    entity_embeddings
    relation_embeddings
    k
    num_sample = 250
    roots
    weights
    compute_func (weights: torch.FloatTensor, x) → torch.FloatTensor
    chain_func (weights, x: torch.FloatTensor)
    forward_triples (idx_triple: torch.Tensor) → torch.Tensor

```

Parameters

x

```

class dicee.models.FMult2 (args)
    Bases: dicee.models.base_model.BaseKGE
    Learning Knowledge Neural Graphs
    name = 'FMult2'
    n_layers = 3
    k
    n = 50
    score_func = 'compositional'
    discrete_points

```

```

entity_embeddings

relation_embeddings

build_func (Vec)

build_chain_funcs (list_Vec)

compute_func (W, b, x) → torch.FloatTensor

function (list_W, list_b)

trapezoid (list_W, list_b)

forward_triples (idx_triple: torch.Tensor) → torch.Tensor

```

Parameters

x

```
class dicee.models.LFMult1 (args)
```

Bases: *[dicee.models.base_model.BaseKGE](#)*

Embedding with trigonometric functions. We represent all entities and relations in the complex number space as: $f(x) = \sum_{k=0}^{d-1} w_k e^{kix}$. and use the three different scoring function as in the paper to evaluate the score

```
name = 'LFMult1'
```

```
entity_embeddings
```

```
relation_embeddings
```

```
forward_triples (idx_triple)
```

Parameters

x

```
tri_score (h, r, t)
```

```
vtp_score (h, r, t)
```

```
class dicee.models.LFMult (args)
```

Bases: *[dicee.models.base_model.BaseKGE](#)*

Embedding with polynomial functions. We represent all entities and relations in the polynomial space as: $f(x) = \sum_{i=0}^{d-1} a_i x^i$ and use the three different scoring function as in the paper to evaluate the score. We also consider combining with Neural Networks.

```
name = 'LFMult'
```

```
entity_embeddings
```

```
relation_embeddings
```

```
degree
```

```
m
```

```
x_values
```

forward_triples (*idx_triple*)

Parameters

x

construct_multi_coeff (*x*)

poly_NN (*x, coefh, coefr, coeft*)

Constructing a 2 layers NN to represent the embeddings. $h = \text{sigma}(w_h^T x + b_h)$, $r = \text{sigma}(w_r^T x + b_r)$,
 $t = \text{sigma}(w_t^T x + b_t)$

linear (*x, w, b*)

scalar_batch_NN (*a, b, c*)

element wise multiplication between a,b and c: Inputs : a, b, c ==> torch.tensor of size batch_size x m x d
Output : a tensor of size batch_size x d

tri_score (*coeff_h, coeff_r, coeff_t*)

this part implement the trilinear scoring techniques:

$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \text{dfrac}\{a_i*b_j*c_k\} \{1+(i+j+k)*d\}$$

1. generate the range for i,j and k from [0 d-1]
2. perform $\text{dfrac}\{a_i*b_j*c_k\} \{1+(i+j+k)*d\}$ in parallel for every batch
3. take the sum over each batch

vtp_score (*h, r, t*)

this part implement the vector triple product scoring techniques:

$$\text{score}(h,r,t) = \int_0^1 h(x)r(x)t(x) dx = \sum_{i,j,k=0}^{d-1} \text{dfrac}\{a_i*c_j*b_k - b_i*c_j*a_k\} \{(1+(i+j)*d)(1+k)\}$$

1. generate the range for i,j and k from [0 d-1]
2. Compute the first and second terms of the sum
3. Multiply with then denominator and take the sum
4. take the sum over each batch

comp_func (*h, r, t*)

this part implement the function composition scoring techniques: i.e. score = <hor, t>

polynomial (*coeff, x, degree*)

This function takes a matrix tensor of coefficients (coeff), a tensor vector of points x and range of integer [0,1,...d] and return a vector tensor (coeff[0][0] + coeff[0][1]x +...+ coeff[0][d]x^d,

$$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d)$$

pop (*coeff, x, degree*)

This function allow us to evaluate the composition of two polynomes without for loops :) it takes a matrix tensor of coefficients (coeff), a matrix tensor of points x and range of integer [0,1,...d]

and return a tensor (coeff[0][0] + coeff[0][1]x +...+ coeff[0][d]x^d,

$$\text{coeff}[1][0] + \text{coeff}[1][1]x + \dots + \text{coeff}[1][d]x^d)$$

class dicee.models.DualE (*args*)

Bases: *dicee.models.base_model.BaseKGE*

Dual Quaternion Knowledge Graph Embeddings (<https://ojs.aaai.org/index.php/AAAI/article/download/16850/16657>)

```

name = 'DualE'

entity_embeddings

relation_embeddings

num_ent = None

kvsall_score(e_1_h, e_2_h, e_3_h, e_4_h, e_5_h, e_6_h, e_7_h, e_8_h, e_1_t, e_2_t, e_3_t, e_4_t,
            e_5_t, e_6_t, e_7_t, e_8_t, r_1, r_2, r_3, r_4, r_5, r_6, r_7, r_8) → torch.tensor
KvsAll scoring function

```

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

```

forward_triples(idx_triple: torch.tensor) → torch.tensor
Negative Sampling forward pass:

```

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

```

forward_k_vs_all(x)
KvsAll forward pass

```

Input

x: torch.LongTensor with (n,) shape

Output

torch.FloatTensor with (n) shape

```

T(x: torch.tensor) → torch.tensor
Transpose function

```

Input: Tensor with shape (nxm) Output: Tensor with shape (mxn)

dicee.query_generator

Classes

QueryGenerator

Module Contents

```
class dicee.query_generator.QueryGenerator(train_path: str, val_path: str, test_path: str,
    ent2id: Dict = None, rel2id: Dict = None, seed: int = 1, gen_valid: bool = False,
    gen_test: bool = True)

    train_path

    val_path

    test_path

    gen_valid = False

    gen_test = True

    seed = 1

    max_ans_num = 1000000.0

    mode

    ent2id = None

    rel2id: Dict = None

    ent_in: Dict

    ent_out: Dict

    query_name_to_struct

    list2tuple(list_data)

    tuple2list(x: List | Tuple) → List | Tuple
        Convert a nested tuple to a nested list.

    set_global_seed(seed: int)
        Set seed

    construct_graph(paths: List[str]) → Tuple[Dict, Dict]
        Construct graph from triples Returns dicts with incoming and outgoing edges

    fill_query(query_structure: List[str | List], ent_in: Dict, ent_out: Dict, answer: int) → bool
        Private method for fill_query logic.

    achieve_answer(query: List[str | List], ent_in: Dict, ent_out: Dict) → set
        Private method for achieve_answer logic. @TODO: Document the code

    write_links(ent_out, small_ent_out)

    ground_queries(query_structure: List[str | List], ent_in: Dict, ent_out: Dict, small_ent_in: Dict,
        small_ent_out: Dict, gen_num: int, query_name: str)
        Generating queries and achieving answers

    unmap(query_type, queries, tp_answers, fp_answers, fn_answers)

    unmap_query(query_structure, query, id2ent, id2rel)
```

generate_queries (*query_struct: List, gen_num: int, query_type: str*)

Passing incoming and outgoing edges to ground queries depending on mode [train valid or text] and getting queries and answers in return @ TODO: create a class for each single query struct

save_queries (*query_type: str, gen_num: int, save_path: str*)

abstractmethod load_queries (*path*)

get_queries (*query_type: str, gen_num: int*)

static save_queries_and_answers (*path: str, data: List[Tuple[str, Tuple[collections.defaultdict]]]*)
→ None

Save Queries into Disk

static load_queries_and_answers (*path: str*) → List[Tuple[str, Tuple[collections.defaultdict]]]

Load Queries from Disk to Memory

dicee.read_preprocess_save_load_kg

Submodules

dicee.read_preprocess_save_load_kg.preprocess

Classes

PreprocessKG

Preprocess the data in memory

Module Contents

class dicee.read_preprocess_save_load_kg.preprocess.**PreprocessKG** (*kg*)

Preprocess the data in memory

kg

start () → None

Preprocess train, valid and test datasets stored in knowledge graph instance

Parameter

rtype

None

preprocess_with_byte_pair_encoding ()

preprocess_with_byte_pair_encoding_with_padding () → None

Preprocess with byte pair encoding and add padding

preprocess_with_pandas () → None

Preprocess with pandas: add reciprocal triples, construct vocabulary, and index datasets

preprocess_with_polars () → None

Preprocess with polars: add reciprocal triples and create indexed datasets

`sequential_vocabulary_construction()` → None

- (1) Read input data into memory
- (2) Remove triples with a condition
- (3) **Serialize vocabularies in a pandas dataframe where**
=> the index is integer and => a single column is string (e.g. URI)

`dicee.read_preprocess_save_load_kg.read_from_disk`

Classes

ReadFromDisk

Read the data from disk into memory

Module Contents

`class dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk(kg)`

Read the data from disk into memory

kg

`start()` → None

Read a knowledge graph from disk into memory

Data will be available at the `train_set`, `test_set`, `valid_set` attributes.

Parameter

None

rtype

None

`add_noisy_triples_into_training()`

`dicee.read_preprocess_save_load_kg.save_load_disk`

Classes

LoadSaveToDisk

Module Contents

`class dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk(kg)`

kg

`save()`

`load()`

dicee.read_preprocess_save_load_kg.util

Functions

<code>polars_dataframe_indexer</code> ($\rightarrow$ polars.DataFrame)	Replaces 'subject', 'relation', and 'object' columns in the input Polars DataFrame with their corresponding index values
<code>pandas_dataframe_indexer</code> ($\rightarrow$ pandas.DataFrame)	Replaces 'subject', 'relation', and 'object' columns in the input Pandas DataFrame with their corresponding index values
<code>apply_reciprocal_or_noise</code> (add_reciprocal, eval_model)	Add reciprocal triples if conditions are met
<code>timeit</code> (func)	
<code>read_with_polars</code> ($\rightarrow$ polars.DataFrame)	Load and Preprocess via Polars
<code>read_with_pandas</code> (data_path[, read_only_few, ...])	Load and Preprocess via Pandas
<code>read_from_disk</code> ($\rightarrow$ Tuple[polars.DataFrame, pandas.DataFrame])	
<code>count_triples</code> ($\rightarrow$ int)	Returns the total number of triples in the triple store.
<code>fetch_worker</code> (endpoint, offsets, chunk_size, ...)	Worker process: fetch assigned chunks and save to disk with per-worker tqdm.
<code>read_from_triple_store_with_polars</code> (endpoint[, ...])	Main function to read all triples in parallel, save as Parquet, and load into Polars dataframe.
<code>read_from_triple_store_with_pandas</code> ([endpoint])	Read triples from triple store into pandas dataframe
<code>get_er_vocab</code> (data[, file_path])	
<code>get_re_vocab</code> (data[, file_path])	
<code>get_ee_vocab</code> (data[, file_path])	
<code>create_constraints</code> (triples[, file_path])	
<code>load_with_pandas</code> ($\rightarrow$ None)	Deserialize data
<code>save_numpy_ndarray</code> (*, data, file_path)	
<code>load_numpy_ndarray</code> (*, file_path)	
<code>save_pickle</code> (*, data[, file_path])	
<code>load_pickle</code> (*[, file_path])	
<code>create_recipriocal_triples</code> (x)	Add inverse triples into dask dataframe
<code>dataset_sanity_checking</code> ($\rightarrow$ None)	

Module Contents

`dicee.read_preprocess_save_load_kg.util.polars_dataframe_indexer` (
 df_polars: polars.DataFrame, idx_entity: polars.DataFrame, idx_relation: polars.DataFrame
 $\rightarrow$ polars.DataFrame

Replaces 'subject', 'relation', and 'object' columns in the input Polars DataFrame with their corresponding index values from the entity and relation index DataFrames.

This function processes the DataFrame in three main steps: 1. Replace the 'relation' values with the corresponding index from *idx_relation*. 2. Replace the 'subject' values with the corresponding index from *idx_entity*. 3. Replace the 'object' values with the corresponding index from *idx_entity*.

Parameters:

df_polars

[polars.DataFrame] The input Polars DataFrame containing columns: 'subject', 'relation', and 'object'.

idx_entity

[polars.DataFrame] A Polars DataFrame that contains the mapping between entity names and their corresponding indices. Must have columns: 'entity' and 'index'.

idx_relation

[polars.DataFrame] A Polars DataFrame that contains the mapping between relation names and their corresponding indices. Must have columns: 'relation' and 'index'.

Returns:

polars.DataFrame

A DataFrame with the 'subject', 'relation', and 'object' columns replaced by their corresponding indices.

Example Usage:

```
>>> df_polars = pl.DataFrame({
    "subject": ["Alice", "Bob", "Charlie"],
    "relation": ["knows", "works_with", "lives_in"],
    "object": ["Dave", "Eve", "Frank"]
})
>>> idx_entity = pl.DataFrame({
    "entity": ["Alice", "Bob", "Charlie", "Dave", "Eve", "Frank"],
    "index": [0, 1, 2, 3, 4, 5]
})
>>> idx_relation = pl.DataFrame({
    "relation": ["knows", "works_with", "lives_in"],
    "index": [0, 1, 2]
})
>>> polars_dataframe_indexer(df_polars, idx_entity, idx_relation)
```

Steps:

1. Join the input DataFrame *df_polars* on the 'relation' column with *idx_relation* to replace the relations with their indices.
2. Join on 'subject' to replace it with the corresponding entity index using a left join on *idx_entity*.
3. Join on 'object' to replace it with the corresponding entity index using a left join on *idx_entity*.
4. Select only the 'subject', 'relation', and 'object' columns to return the final result.

```
dicee.read_preprocess_save_load_kg.util.pandas_dataframe_indexer(
    df_pandas: pandas.DataFrame, idx_entity: pandas.DataFrame, idx_relation: pandas.DataFrame
) → pandas.DataFrame
```

Replaces 'subject', 'relation', and 'object' columns in the input Pandas DataFrame with their corresponding index values from the entity and relation index DataFrames.

Parameters:

df_pandas

[pd.DataFrame] The input Pandas DataFrame containing columns: 'subject', 'relation', and 'object'.

idx_entity

[pd.DataFrame] A Pandas DataFrame that contains the mapping between entity names and their corresponding indices. Must have columns: 'entity' and 'index'.

idx_relation

[pd.DataFrame] A Pandas DataFrame that contains the mapping between relation names and their corresponding indices. Must have columns: 'relation' and 'index'.

Returns:

pd.DataFrame

A DataFrame with the 'subject', 'relation', and 'object' columns replaced by their corresponding indices.

```
dicee.read_preprocess_save_load_kg.util.apply_reciprocal_or_noise (add_reciprocal: bool,  
    eval_model: str, df: object = None, info: str = None)
```

Add reciprocal triples if conditions are met

```
dicee.read_preprocess_save_load_kg.util.timeit (func)
```

```
dicee.read_preprocess_save_load_kg.util.read_with_polars (data_path,  
    read_only_few: int = None, sample_triples_ratio: float = None, separator: str = None)  
    → polars.DataFrame
```

Load and Preprocess via Polars

```
dicee.read_preprocess_save_load_kg.util.read_with_pandas (data_path,  
    read_only_few: int = None, sample_triples_ratio: float = None, separator: str = None)
```

Load and Preprocess via Pandas

```
dicee.read_preprocess_save_load_kg.util.read_from_disk (data_path: str,  
    read_only_few: int = None, sample_triples_ratio: float = None, backend: str = None,  
    separator: str = None) → Tuple[polars.DataFrame, pandas.DataFrame]
```

```
dicee.read_preprocess_save_load_kg.util.count_triples (endpoint: str) → int
```

Returns the total number of triples in the triple store.

```
dicee.read_preprocess_save_load_kg.util.fetch_worker (endpoint: str, offsets: list[int],  
    chunk_size: int, output_dir: str, worker_id: int)
```

Worker process: fetch assigned chunks and save to disk with per-worker tqdm.

```
dicee.read_preprocess_save_load_kg.util.read_from_triple_store_with_polars (  
    endpoint: str, chunk_size: int = 500000, output_dir: str = 'triples_parquet')
```

Main function to read all triples in parallel, save as Parquet, and load into Polars dataframe.

```
dicee.read_preprocess_save_load_kg.util.read_from_triple_store_with_pandas (  
    endpoint: str = None)
```

Read triples from triple store into pandas dataframe

```
dicee.read_preprocess_save_load_kg.util.get_er_vocab (data, file_path: str = None)
```

```
dicee.read_preprocess_save_load_kg.util.get_re_vocab (data, file_path: str = None)
```

```
dicee.read_preprocess_save_load_kg.util.get_ee_vocab (data, file_path: str = None)
```

```
dicee.read_preprocess_save_load_kg.util.create_constraints (triples, file_path: str = None)
```

(1) Extract domains and ranges of relations

(2) Store a mapping from relations to entities that are outside of the domain and range. Create constrained entities based on the range of relations :param triples: :return: Tuple[dict, dict]

```
dicee.read_preprocess_save_load_kg.util.load_with_pandas (self) → None
```

Deserialize data

```
dicee.read_preprocess_save_load_kg.util.save_numpy_ndarray(* , data: numpy.ndarray,  
    file_path: str)
```

```
dicee.read_preprocess_save_load_kg.util.load_numpy_ndarray(* , file_path: str)
```

```
dicee.read_preprocess_save_load_kg.util.save_pickle(* , data: object, file_path=str)
```

```
dicee.read_preprocess_save_load_kg.util.load_pickle(* , file_path=str)
```

```
dicee.read_preprocess_save_load_kg.util.create_recipriocal_triples(x)
```

Add inverse triples into dask dataframe :param x: :return:

```
dicee.read_preprocess_save_load_kg.util.dataset_sanity_checking(  
    train_set: numpy.ndarray, num_entities: int, num_relations: int) → None
```

Parameters

- **train_set**
- **num_entities**
- **num_relations**

Returns

Classes

<i>PreprocessKG</i>	Preprocess the data in memory
<i>LoadSaveToDisk</i>	
<i>ReadFromDisk</i>	Read the data from disk into memory

Package Contents

```
class dicee.read_preprocess_save_load_kg.PreprocessKG(kg)
```

Preprocess the data in memory

kg

start() → None

Preprocess train, valid and test datasets stored in knowledge graph instance

Parameter

rtype

None

preprocess_with_byte_pair_encoding()

preprocess_with_byte_pair_encoding_with_padding() → None

Preprocess with byte pair encoding and add padding

preprocess_with_pandas() → None

Preprocess with pandas: add reciprocal triples, construct vocabulary, and index datasets

preprocess_with_polars() → None

Preprocess with polars: add reciprocal triples and create indexed datasets

sequential_vocabulary_construction() → None

- (1) Read input data into memory
- (2) Remove triples with a condition
- (3) **Serialize vocabularies in a pandas dataframe where**
=> the index is integer and => a single column is string (e.g. URI)

class dicee.read_preprocess_save_load_kg.**LoadSaveToDisk**(kg)

kg

save()

load()

class dicee.read_preprocess_save_load_kg.**ReadFromDisk**(kg)

Read the data from disk into memory

kg

start() → None

Read a knowledge graph from disk into memory

Data will be available at the train_set, test_set, valid_set attributes.

Parameter

None

rtype

None

add_noisy_triples_into_training()

dicee.sanity_checkers

Functions

<code>is_sparql_endpoint_alive([sparql_endpoint])</code>	
<code>validate_knowledge_graph(args)</code>	Validating the source of knowledge graph
<code>sanity_checking_with_arguments(args)</code>	
<code>sanity_check_callback_args(args)</code>	Perform sanity checks on callback-related arguments.

Module Contents

`dicee.sanity_checkers.is_sparql_endpoint_alive(sparql_endpoint: str = None)`

`dicee.sanity_checkers.validate_knowledge_graph(args)`

Validating the source of knowledge graph

`dicee.sanity_checkers.sanity_checking_with_arguments(args)`

`dicee.sanity_checkers.sanity_check_callback_args(args)`

Perform sanity checks on callback-related arguments.

dicee.scripts

Submodules

dicee.scripts.index_serve

```
$ docker pull qdrant/qdrant && docker run -p 6333:6333 -p 6334:6334 -v $(pwd)/qdrant_storage:/qdrant/storage:z  
qdrant/qdrant $ dicee_vector_db -index -serve -path CountryEmbeddings -collection "countries_vdb"
```

Attributes

<code>app</code>
<code>neural_searcher</code>

Classes

<code>NeuralSearcher</code>
<code>StringListRequest</code>

Functions

<code>get_default_arguments()</code>
<code>index(args)</code>
<code>root()</code>
<code>search_embeddings(q)</code>
<code>retrieve_embeddings(q)</code>
<code>search_embeddings_batch(request)</code>
<code>serve(args)</code>
<code>main()</code>

Module Contents

```
dicee.scripts.index_serve.get_default_arguments()  
  
dicee.scripts.index_serve.index(args)  
  
dicee.scripts.index_serve.app  
  
dicee.scripts.index_serve.neural_searcher = None  
  
class dicee.scripts.index_serve.NeuralSearcher(args)  
    collection_name  
  
    entity_to_idx = None  
  
    qdrant_client  
  
    topk = 5  
  
    retrieve_embedding(entity: str = None, entities: List[str] = None) → List
```

```

    search(entity: str)

async dicee.scripts.index_serve.root()

async dicee.scripts.index_serve.search_embeddings(q: str)

async dicee.scripts.index_serve.retrieve_embeddings(q: str)

class dicee.scripts.index_serve.StringListRequest
    Bases: pydantic.BaseModel

    queries: List[str]

    reducer: str | None = None

async dicee.scripts.index_serve.search_embeddings_batch(request: StringListRequest)

dicee.scripts.index_serve.serve(args)

dicee.scripts.index_serve.main()

```

dicee.scripts.run

Functions

<code>get_default_arguments([description])</code> <code>main()</code>	Extends pytorch_lightning Trainer's arguments with ours
--	---

Module Contents

```

dicee.scripts.run.get_default_arguments(description=None)
    Extends pytorch_lightning Trainer's arguments with ours

dicee.scripts.run.main()

```

dicee.static_funcs

Static utility functions for DICE embeddings.

This module provides utility functions for model initialization, data loading, serialization, and various helper operations.

Attributes

<code>MODEL_REGISTRY</code>	
-----------------------------	--

Functions

<code>create_recipriocal_triples(→ pandas.DataFrame)</code>	Add inverse triples to a DataFrame.
<code>get_er_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build entity-relation to tail vocabulary.
<code>get_re_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build relation-entity (tail) to head vocabulary.

continues on next page

Table 55 – continued from previous page

<code>get_ee_vocab(→ Dict[Tuple[int, int], List[int]])</code>	Build entity-entity to relation vocabulary.
<code>timeit(→ Callable)</code>	Decorator to measure and print execution time and memory usage.
<code>save_pickle(→ None)</code>	Save data to a pickle file.
<code>load_pickle(→ object)</code>	Load data from a pickle file.
<code>load_term_mapping(→ Union[dict, polars.DataFrame])</code>	Load term-to-index mapping from pickle or CSV file.
<code>select_model(args[, is_continual_training, storage_path])</code>	
<code>load_model(→ Tuple[object, Tuple[dict, dict]])</code>	Load weights and initialize pytorch module from namespace arguments
<code>load_model_ensemble(...)</code>	Construct Ensemble Of weights and initialize pytorch module from namespace arguments
<code>save_numpy_ndarray(*, data, file_path)</code>	
<code>numpy_data_type_changer(→ numpy.ndarray)</code>	Detect most efficient data type for a given triples
<code>save_checkpoint_model(→ None)</code>	Store Pytorch model into disk
<code>store(→ None)</code>	
<code>add_noisy_triples(→ pandas.DataFrame)</code>	Add randomly constructed triples
<code>read_or_load_kg(args, cls)</code>	
<code>intialize_model(...)</code>	Initialize a knowledge graph embedding model.
<code>load_json(→ Dict)</code>	Load JSON file into a dictionary.
<code>save_embeddings(→ None)</code>	Save embeddings to a CSV file.
<code>vocab_to_parquet(vocab_to_idx, name, ...)</code>	
<code>create_experiment_folder(→ str)</code>	Create a timestamped experiment folder.
<code>continual_training_setup_executor(→ None)</code>	
<code>exponential_function(→ torch.FloatTensor)</code>	
<code>load_numpy(→ numpy.ndarray)</code>	
<code>evaluate(entity_to_idx, scores, easy_answers, hard_answers)</code>	# @TODO: CD: Renamed this function
<code>download_file(url[, destination_folder])</code>	
<code>download_files_from_url(→ None)</code>	
<code>download_pretrained_model(→ str)</code>	
<code>write_csv_from_model_parallel(path)</code>	Create
<code>from_pretrained_model_write_embeddings_into_csv(path, None)</code>	

Module Contents

`dicee.static_funcs.MODEL_REGISTRY: Dict[str, Tuple[Type, str]]`

`dicee.static_funcs.create_recipriocal_triples(df: pandas.DataFrame) → pandas.DataFrame`

Add inverse triples to a DataFrame.

For each triple (s, p, o), creates an inverse triple (o, p_inverse, s).

Parameters

df – DataFrame with ‘subject’, ‘relation’, and ‘object’ columns.

Returns

DataFrame with original and inverse triples concatenated.

`dicee.static_funcs.get_er_vocab(data: numpy.ndarray, file_path: str | None = None)`

→ Dict[Tuple[int, int], List[int]]

Build entity-relation to tail vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (head, relation) pairs to list of tail entities.

```
dicee.static_funcs.get_re_vocab(data: numpy.ndarray, file_path: str | None = None)  
→ Dict[Tuple[int, int], List[int]]
```

Build relation-entity (tail) to head vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (relation, tail) pairs to list of head entities.

```
dicee.static_funcs.get_ee_vocab(data: numpy.ndarray, file_path: str | None = None)  
→ Dict[Tuple[int, int], List[int]]
```

Build entity-entity to relation vocabulary.

Parameters

- **data** – Array of triples with shape (n, 3) where columns are (head, relation, tail).
- **file_path** – Optional path to save the vocabulary as pickle.

Returns

Dictionary mapping (head, tail) pairs to list of relations.

```
dicee.static_funcs.timeit(func: Callable) → Callable
```

Decorator to measure and print execution time and memory usage.

Parameters

func – Function to be timed.

Returns

Wrapped function that prints timing information.

```
dicee.static_funcs.save_pickle(*, data: object | None = None, file_path: str) → None
```

Save data to a pickle file.

Note: Consider using more portable formats (JSON, Parquet) for new code.

Parameters

- **data** – Object to serialize. If None, nothing is saved.
- **file_path** – Path where the pickle file will be saved.

```
dicee.static_funcs.load_pickle(file_path: str) → object
```

Load data from a pickle file.

Note: Consider using more portable formats (JSON, Parquet) for new code.

Parameters

file_path – Path to the pickle file.

Returns

Deserialized object from the pickle file.

`dicee.static_funcs.load_term_mapping (file_path: str) → dict | polars.DataFrame`

Load term-to-index mapping from pickle or CSV file.

Attempts to load from pickle first, falls back to CSV if not found.

Parameters

file_path – Base path without extension.

Returns

Dictionary or Polars DataFrame containing the mapping.

`dicee.static_funcs.select_model (args: dict, is_continual_training: bool = None, storage_path: str = None)`

`dicee.static_funcs.load_model (path_of_experiment_folder: str, model_name='model.pt', verbose=0) → Tuple[object, Tuple[dict, dict]]`

Load weights and initialize pytorch module from namespace arguments

`dicee.static_funcs.load_model_ensemble (path_of_experiment_folder: str) → Tuple[dicee.models.base_model.BaseKGE, Tuple[pandas.DataFrame, pandas.DataFrame]]`

Construct Ensemble Of weights and initialize pytorch module from namespace arguments

- (1) Detect models under given path
- (2) Accumulate parameters of detected models
- (3) Normalize parameters
- (4) Insert (3) into model.

`dicee.static_funcs.save_numpy_ndarray (*, data: numpy.ndarray, file_path: str)`

`dicee.static_funcs.numpy_data_type_changer (train_set: numpy.ndarray, num: int) → numpy.ndarray`

Detect most efficient data type for a given triples :param train_set: :param num: :return:

`dicee.static_funcs.save_checkpoint_model (model, path: str) → None`

Store Pytorch model into disk

`dicee.static_funcs.store (trained_model, model_name: str = 'model', full_storage_path: str = None, save_embeddings_as_csv=False) → None`

`dicee.static_funcs.add_noisy_triples (train_set: pandas.DataFrame, add_noise_rate: float) → pandas.DataFrame`

Add randomly constructed triples :param train_set: :param add_noise_rate: :return:

`dicee.static_funcs.read_or_load_kg (args, cls)`

`dicee.static_funcs.intialize_model (args: Dict, verbose: int = 0) → Tuple[dicee.models.base_model.BaseKGE, str]`

Initialize a knowledge graph embedding model.

Parameters

- **args** – Dictionary containing model configuration including ‘model’ key.
- **verbose** – Verbosity level. If > 0, prints initialization message.

Returns

Tuple of (initialized model, form of labelling string).

Raises

ValueError – If the model name is not recognized.

`dicee.static_funcs.load_json(path: str) → Dict`

Load JSON file into a dictionary.

Parameters

path – Path to the JSON file.

Returns

Dictionary containing the JSON data.

Raises

- **FileNotFoundError** – If the file does not exist.
- **json.JSONDecodeError** – If the file contains invalid JSON.

`dicee.static_funcs.save_embeddings(embeddings: numpy.ndarray, indexes: List, path: str) → None`

Save embeddings to a CSV file.

Parameters

- **embeddings** – NumPy array of embeddings with shape (n_items, embedding_dim).
- **indexes** – List of index labels (entity/relation names).
- **path** – Output file path.

`dicee.static_funcs.vocab_to_parquet(vocab_to_idx, name, path_for_serialization, print_into)`

`dicee.static_funcs.create_experiment_folder(folder_name: str = 'Experiments') → str`

Create a timestamped experiment folder.

Parameters

folder_name – Base directory name for experiments.

Returns

Full path to the created experiment folder.

`dicee.static_funcs.continual_training_setup_executor(executor) → None`

`dicee.static_funcs.exponential_function(x: numpy.ndarray, lam: float, ascending_order=True)`
→ torch.FloatTensor

`dicee.static_funcs.load_numpy(path) → numpy.ndarray`

`dicee.static_funcs.evaluate(entity_to_idx, scores, easy_answers, hard_answers)`

@TODO: CD: Renamed this function Evaluate multi hop query answering on different query types

`dicee.static_funcs.download_file(url, destination_folder='.')`

`dicee.static_funcs.download_files_from_url(base_url: str, destination_folder='.') → None`

Parameters

- **base_url** (e.g. <https://files.dice-research.org/projects/DiceEmbeddings/KINSHIP-Keci-dim128-epoch256-KvsAll>)
- **destination_folder** (e.g. `"KINSHIP-Keci-dim128-epoch256-KvsAll"`)

`dicee.static_funcs.download_pretrained_model(url: str) → str`

`dicee.static_funcs.write_csv_from_model_parallel(path: str)`

Create

`dicee.static_funcs.from_pretrained_model_write_embeddings_into_csv(path: str) → None`

dicee.static_funcs_training

Training-related static functions.

This module provides backward compatibility by re-exporting evaluation functions from the new `dicee.evaluation` module, along with training utilities.

Deprecated since version Evaluation: functions have moved to `dicee.evaluation`. Use that module for new code. This module will continue to export training utilities.

Functions

<code>evaluate_lp</code> ($\rightarrow$ Dict[str, float])	Evaluate link prediction with batched processing.
<code>evaluate_bpe_lp</code> ($\rightarrow$ Dict[str, float])	Evaluate link prediction with BPE-encoded entities.
<code>make_iterable_verbose</code> ($\rightarrow$ Iterable)	Wrap an iterable with tqdm progress bar if verbose is True.
<code>efficient_zero_grad</code> ($\rightarrow$ None)	Efficiently zero gradients using <code>parameter.grad = None</code> .

Module Contents

`dicee.static_funcs_training.evaluate_lp`(*model*, *triple_idx*, *num_entities*: int, *er_vocab*: Dict[Tuple, List], *re_vocab*: Dict[Tuple, List], *info*: str = 'Eval Starts', *batch_size*: int = 128, *chunk_size*: int = 1000) $\rightarrow$ Dict[str, float]

Evaluate link prediction with batched processing.

Memory-efficient evaluation using chunked entity scoring.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – Integer-indexed triples as numpy array.
- **num_entities** – Total number of entities.
- **er_vocab** – Mapping (head_idx, rel_idx) -> list of tail indices.
- **re_vocab** – Mapping (rel_idx, tail_idx) -> list of head indices.
- **info** – Description to print.
- **batch_size** – Batch size for triple processing.
- **chunk_size** – Chunk size for entity scoring.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.static_funcs_training.evaluate_bpe_lp`(*model*, *triple_idx*: List[Tuple], *all_bpe_shaped_entities*, *er_vocab*: Dict[Tuple, List], *re_vocab*: Dict[Tuple, List], *info*: str = 'Eval Starts') $\rightarrow$ Dict[str, float]

Evaluate link prediction with BPE-encoded entities.

Parameters

- **model** – The KGE model to evaluate.
- **triple_idx** – List of BPE-encoded triple tuples.
- **all_bpe_shaped_entities** – All entities with BPE representations.
- **er_vocab** – Mapping for tail filtering.

- **re_vocab** – Mapping for head filtering.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

`dicee.static_funcs_training.make_iterable_verbose(iterable_object: Iterable, verbose: bool, desc: str = 'Default', position: int | None = None, leave: bool = True) → Iterable`

Wrap an iterable with tqdm progress bar if verbose is True.

Parameters

- **iterable_object** – The iterable to potentially wrap.
- **verbose** – Whether to show progress bar.
- **desc** – Description for the progress bar.
- **position** – Position of the progress bar.
- **leave** – Whether to leave the progress bar after completion.

Returns

The original iterable or a tqdm-wrapped version.

`dicee.static_funcs_training.efficient_zero_grad(model) → None`

Efficiently zero gradients using `parameter.grad = None`.

This is more efficient than `optimizer.zero_grad()` as it avoids memory operations.

See: https://pytorch.org/tutorials/recipes/recipes/tuning_guide.html

Parameters

model – PyTorch model to zero gradients for.

dicee.static_preprocess_funcs

Attributes

<code>enable_log</code>

Functions

<code>timeit(func)</code>	
<code>preprocesses_input_args(args)</code>	Sanity Checking in input arguments
<code>create_constraints(→ Tuple[dict, dict, dict, dict])</code>	
<code>get_er_vocab(data)</code>	
<code>get_re_vocab(data)</code>	
<code>get_ee_vocab(data)</code>	
<code>mapping_from_first_two_cols_to_third(train_se</code>	

Module Contents

`dicee.static_preprocess_funcs.enable_log = False`

`dicee.static_preprocess_funcs.timeit(func)`

```
dicee.static_preprocess_funcs.preprocesses_input_args (args)
```

Sanity Checking in input arguments

```
dicee.static_preprocess_funcs.create_constraints (triples: numpy.ndarray)
```

→ Tuple[dict, dict, dict, dict]

(1) Extract domains and ranges of relations

(2) Store a mapping from relations to entities that are outside of the domain and range. Create constraints entities based on the range of relations :param triples: :return:

```
dicee.static_preprocess_funcs.get_er_vocab (data)
```

```
dicee.static_preprocess_funcs.get_re_vocab (data)
```

```
dicee.static_preprocess_funcs.get_ee_vocab (data)
```

```
dicee.static_preprocess_funcs.mapping_from_first_two_cols_to_third (train_set_idx)
```

dicee.trainer

Submodules

dicee.trainer.dice_trainer

DICE Trainer module for knowledge graph embedding training.

Provides the DICE_Trainer class which supports multiple training backends including PyTorch Lightning, DDP, and custom CPU/GPU trainers.

Classes

<code>DICE_Trainer</code>	DICE_Trainer implement
---------------------------	------------------------

Functions

<code>load_term_mapping</code> (→ polars.DataFrame)	Load term-to-index mapping from CSV file.
<code>initialize_trainer</code> (...)	Initialize the appropriate trainer based on configuration.
<code>get_callbacks</code> (→ List)	Create list of callbacks based on configuration.

Module Contents

```
dicee.trainer.dice_trainer.load_term_mapping (file_path: str) → polars.DataFrame
```

Load term-to-index mapping from CSV file.

Parameters

file_path – Base path without extension.

Returns

Polars DataFrame containing the mapping.

```
dicee.trainer.dice_trainer.initialize_trainer (args, callbacks: List)
```

→ `dicee.trainer.torch_trainer.TorchTrainer` | `dicee.trainer.model_parallelism.TensorParallel` | `dicee.trainer.torch_trainer_ddp`

Initialize the appropriate trainer based on configuration.

Parameters

- **args** – Configuration arguments containing trainer type.
- **callbacks** – List of training callbacks.

Returns

Initialized trainer instance.

Raises

AssertionError – If trainer is None after initialization.

`dicee.trainer.dice_trainer.get_callbacks(args) → List`

Create list of callbacks based on configuration.

Parameters

args – Configuration arguments.

Returns

List of callback instances.

```
class dicee.trainer.dice_trainer.DICE_Trainer (args, is_continual_training: bool, storage_path,
        evaluator=None)
```

DICE_Trainer implement

- 1- Pytorch Lightning trainer (<https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html>)
- 2- Multi-GPU Trainer(<https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html>)
- 3- CPU Trainer

args

is_continual_training:bool

storage_path:str

evaluator:

report:dict

report

args

trainer = None

is_continual_training

storage_path

evaluator = None

form_of_labelling = None

continual_start (*knowledge_graph*)

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

initialize_trainer (*callbacks: List*)

→ `lightning.Trainer` | `diccee.trainer.model_parallelism.TensorParallel` | `diccee.trainer.torch_trainer.TorchTrainer` | `diccee.t`

Initialize Trainer from input arguments

initialize_or_load_model ()

init_dataloader (*dataset: torch.utils.data.Dataset*) → `torch.utils.data.DataLoader`

init_dataset () → `torch.utils.data.Dataset`

start (*knowledge_graph: diccee.knowledge_graph.KG* | *numpy.memmap*)

→ `Tuple[diccee.models.base_model.BaseKGE, str]`

Start the training

- (1) Initialize Trainer
- (2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

k_fold_cross_validation (*dataset*) → `Tuple[diccee.models.base_model.BaseKGE, str]`

Perform K-fold Cross-Validation

1. Obtain K train and test splits.
2. **For each split,**
 - 2.1 initialize trainer and model
 - 2.2. Train model with configuration provided in args.
 - 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.
3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

diccee.trainer.model_parallelism

Classes

TensorParallel

Abstract class for Trainer class for knowledge graph embedding models

Functions

<code>extract_input_outputs(z[, device])</code>
<code>find_good_batch_size(train_loader,</code> <code>tp_ensemble_model)</code>
<code>forward_backward_update_loss(→ float)</code>

Module Contents

```
dicee.trainer.model_parallelism.extract_input_outputs (z: list, device=None)
dicee.trainer.model_parallelism.find_good_batch_size (train_loader, tp_ensemble_model)
dicee.trainer.model_parallelism.forward_backward_update_loss (z: Tuple, ensemble_model)
    → float
```

```
class dicee.trainer.model_parallelism.TensorParallel (args, callbacks)
```

Bases: `dicee.abstracts.AbstractTrainer`

Abstract class for Trainer class for knowledge graph embedding models

Parameter

args

[str] ?

callbacks: list

?

fit (*args, **kwargs)

Train model

`dicee.trainer.torch_trainer`

Classes

<code>TorchTrainer</code>	TorchTrainer for using single GPU or multi CPUs on a single node
---------------------------	--

Module Contents

```
class dicee.trainer.torch_trainer.TorchTrainer (args, callbacks)
```

Bases: `dicee.abstracts.AbstractTrainer`

TorchTrainer for using single GPU or multi CPUs on a single node

Arguments

callbacks: list of Abstract callback instances

loss_function = None

optimizer = None

model = None

train_dataloaders = None

training_step = None

process

fit (*args, train_dataloaders, **kwargs) → None

Training starts

Arguments

kwargs:Tuple

empty dictionary

Return type

batch loss (float)

forward_backward_update (x_batch: torch.Tensor, y_batch: torch.Tensor) → torch.Tensor

Compute forward, loss, backward, and parameter update

Arguments

Return type

batch loss (float)

extract_input_outputs_set_device (batch: list) → Tuple

Construct inputs and outputs from a batch of inputs with outputs From a batch of inputs and put

Arguments

Return type

(tuple) mini-batch on select device

dicee.trainer.torch_trainer_ddp

Classes

TorchDDPTrainer

A Trainer based on torch.nn.parallel.DistributedDataParallel

NodeTrainer

Functions

make_iterable_verbose(→ Iterable)

Module Contents

dicee.trainer.torch_trainer_ddp.**make_iterable_verbose** (iterable_object, verbose,
desc='Default', position=None, leave=True) → Iterable

```

class dicee.trainer.torch_trainer_ddp.TorchDDPTrainer (args, callbacks)
    Bases: dicee.abstracts.AbstractTrainer

    A Trainer based on torch.nn.parallel.DistributedDataParallel

    Arguments

    entity_idx
        mapping.

    relation_idx
        mapping.

    form
        ?

    store
        ?

    label_smoothing_rate
        Using hard targets (0,1) drives weights to infinity. An outlier produces enormous gradients.

    Return type
        torch.utils.data.Dataset

    fit (*args, **kwargs)
        Train model

class dicee.trainer.torch_trainer_ddp.NodeTrainer (trainer, model: torch.nn.Module,
    train_dataset_loader: torch.utils.data.DataLoader, callbacks, num_epochs: int)

    trainer

    local_rank

    global_rank

    optimizer

    train_dataset_loader

    loss_func

    callbacks

    model

    num_epochs

    loss_history = []

    ctx

    scaler

    extract_input_outputs (z: list)

    train ()
        Training loop for DDP

```

Classes

DICE_Trainer

DICE_Trainer implement

Package Contents

```
class dicee.trainer.DICE_Trainer (args, is_continual_training: bool, storage_path, evaluator=None)
```

DICE_Trainer implement

- 1- Pytorch Lightning trainer (<https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html>)
- 2- Multi-GPU Trainer(<https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html>)
- 3- CPU Trainer

args

is_continual_training:bool

storage_path:str

evaluator:

report:dict

report

args

trainer = None

is_continual_training

storage_path

evaluator = None

form_of_labelling = None

continual_start (*knowledge_graph*)

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

initialize_trainer (*callbacks: List*)

→ *lightning.Trainer* | *dicee.trainer.model_parallelism.TensorParallel* | *dicee.trainer.torch_trainer.TorchTrainer* | *dicee.t*

Initialize Trainer from input arguments

initialize_or_load_model ()

init_dataloader (*dataset: torch.utils.data.Dataset*) → torch.utils.data.DataLoader

init_dataset () → torch.utils.data.Dataset

start (*knowledge_graph: dicee.knowledge_graph.KG | numpy.memmap*)
→ Tuple[*dicee.models.base_model.BaseKGE*, str]

Start the training

(1) Initialize Trainer

(2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

k_fold_cross_validation (*dataset*) → Tuple[*dicee.models.base_model.BaseKGE*, str]

Perform K-fold Cross-Validation

1. Obtain K train and test splits.

2. **For each split,**

2.1 initialize trainer and model 2.2. Train model with configuration provided in args. 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.

3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

dicee.weight_averaging

Classes

<i>ASWA</i>	Adaptive stochastic weight averaging
<i>SWA</i>	Stochastic Weight Averaging callback.
<i>SWAG</i>	Stochastic Weight Averaging - Gaussian (SWAG).
<i>EMA</i>	Exponential Moving Average (EMA) callback.
<i>TWA</i>	Train with Weight Averaging (TWA) using subspace projection + averaging.

Module Contents

class `dicee.weight_averaging.ASWA` (*num_epochs, path*)

Bases: `dicee.abstracts.AbstractCallback`

Adaptive stochastic weight averaging ASWE keeps track of the validation performance and update s the ensemble model accordingly.

path

num_epochs

initial_eval_setting = None

```
epoch_count = 0

alphas = []

val_aswa = -1

on_fit_end(trainer, model)
    Call at the end of the training.
```

Parameter

trainer:

model:

rtype

None

```
static compute_mrr(trainer, model) → float
```

```
get_aswa_state_dict(model)
```

```
decide(running_model_state_dict, ensemble_state_dict, val_running_model,
        mrr_updated_ensemble_model)
```

Perform Hard Update, software or rejection

Parameters

- **running_model_state_dict**
- **ensemble_state_dict**
- **val_running_model**
- **mrr_updated_ensemble_model**

```
on_train_epoch_end(trainer, model)
```

Call at the end of each epoch during training.

Parameter

trainer:

model:

rtype

None

```
class dicee.weight_averaging.SWA(swa_start_epoch, swa_c_epochs: int = 1, lr_init: float = 0.1,
    swa_lr: float = 0.05, max_epochs: int = None)
```

Bases: *[dicee.abstracts.AbstractCallback](#)*

Stochastic Weight Averaging callback.

Initialize SWA callback.

swa_start_epoch: int

The epoch at which to start SWA.

swa_c_epochs: int

The number of epochs to use for SWA.

```

lr_init: float
    The initial learning rate.

swa_lr: float
    The learning rate to use during SWA.

max_epochs: int
    The maximum number of epochs. args.num_epochs

swa_start_epoch

swa_c_epochs = 1

swa_lr = 0.05

lr_init = 0.1

max_epochs = None

swa_model = None

swa_n = 0

current_epoch = -1

static moving_average (swa_model, running_model, alpha)
    Update SWA model with moving average of current model. Math: # SWA update:  $\theta_{\text{swa}} \leftarrow (1 - \alpha) * \theta_{\text{swa}} + \alpha * \theta$  #  $\alpha = 1 / (n + 1)$ , where n = number of models already averaged # alpha is tracked via self.swa_n in code

on_train_epoch_start (trainer, model)
    Update learning rate according to SWA schedule.

on_train_epoch_end (trainer, model)
    Apply SWA averaging if conditions are met.

on_fit_end (trainer, model)
    Replace main model with SWA model at the end of training.

class dicee.weight_averaging.SWAG (swa_start_epoch: int = 1, lr_init: float = 0.1,
    swa_lr: float = 0.05, max_epochs: int = None, max_num_models: int = 20,
    var_clamp: float = 1e-30)

Bases: dicee.abstracts.AbstractCallback

Stochastic Weight Averaging - Gaussian (SWAG). Parameters

swa_start_epoch
    [int] Epoch at which to start collecting weights.

swa_c_epochs
    [int] Interval of epochs between updates.

lr_init
    [float] Initial LR.

swa_lr
    [float] LR in SWA / GSWA phase.

max_epochs
    [int] Total number of epochs.

```

```

    max_num_models
        [int] Number of models to keep for low-rank covariance approx.

    var_clamp
        [float] Clamp low variance for stability.

swa_start_epoch

swa_c_epochs = 1

swa_lr = 0.05

lr_init = 0.1

max_epochs = None

max_num_models = 20

var_clamp = 1e-30

mean = None

sq_mean = None

deviations = []

gswa_n = 0

current_epoch = -1

get_mean_and_var()
    Return mean + variance (diagonal part).

sample(base_model, scale=0.5)
    Sample new model from SWAG posterior distribution.

    Math: # From SWAG, posterior is approximated as:  $\theta \sim N(\text{mean}, \Sigma)$  # where  $\Sigma \approx \text{diag}(\text{var}) + (1/(K-1)) * D D^T$  # - mean = running average of weights # - var = elementwise variance (sq_mean - mean^2) # - D = [dev_1, dev_2, ..., dev_K], deviations from mean (low-rank approx) # - K = number of collected models

    # Sampling step: # 1.  $\theta_{\text{diag}} = \text{mean} + \text{scale} * \text{std} \odot \varepsilon$ , where  $\varepsilon \sim N(0, I)$  # 2.  $\theta_{\text{lowrank}} = \theta_{\text{diag}} + (D z) / \text{sqrt}(K-1)$ , where  $z \sim N(0, I_K)$  # Final sample =  $\theta_{\text{lowrank}}$ 

on_train_epoch_start(trainer, model)
    Update LR schedule (same as SWA).

on_train_epoch_end(trainer, model)
    Collect Gaussian stats at the end of epochs after swa_start.

on_fit_end(trainer, model)
    Set model weights to the collected SWAG mean at the end of training.

class dicee.weight_averaging.EMA(ema_start_epoch: int, decay: float = 0.999,
    max_epochs: int = None, ema_c_epochs: int = 1)

    Bases: dicee.abstracts.AbstractCallback

    Exponential Moving Average (EMA) callback.

    Parameters
        • ema_start_epoch (int) – Epoch to start EMA.

```

- **decay** (*float*) – EMA decay rate (typical: 0.99 - 0.9999) Math: $\theta_{\text{ema}} \leftarrow \text{decay} * \theta_{\text{ema}} + (1 - \text{decay}) * \theta$
- **max_epochs** (*int*) – Maximum number of epochs.

ema_start_epoch

decay = 0.999

max_epochs = None

ema_c_epochs = 1

ema_model = None

current_epoch = -1

static ema_update (*ema_model, running_model, decay: float*)

Update EMA model with exponential moving average of current model. Math: # EMA update: $\theta_{\text{ema}} \leftarrow (1 - \alpha) * \theta_{\text{ema}} + \alpha * \theta$ # $\alpha = 1 - \text{decay}$, where decay is the EMA smoothing factor (typical 0.99 - 0.999) # α controls how much of the current model θ contributes to the EMA # decay is fixed in code -> can be extended to scheduled

on_train_epoch_start (*trainer, model*)

Track current epoch.

on_train_epoch_end (*trainer, model*)

Update EMA if past start epoch.

on_fit_end (*trainer, model*)

Replace main model with EMA model at the end of training.

class dicee.weight_averaging.**TWA** (*twa_start_epoch: int, lr_init: float, num_samples: int = 5, reg_lambda: float = 0.0, max_epochs: int = None, twa_c_epochs: int = 1*)

Bases: *dicee.abstracts.AbstractCallback*

Train with Weight Averaging (TWA) using subspace projection + averaging.

Parameters

twa_start_epoch

[int] Epoch to start TWA.

lr_init

[float] Learning rate used for β updates.

num_samples

[int] Number of sampled weight snapshots to build projection subspace.

reg_lambda

[float] Regularization coefficient for β updates.

max_epochs

[int] Total number of training epochs.

twa_c_epochs

[int] Interval of epochs between TWA updates.

twa_start_epoch

num_samples = 5


```

reg_lambda = 0.0

max_epochs = None

lr_init

twa_c_epochs = 1

current_epoch = -1

weight_samples = []

twa_model = None

base_weights = None

P = None

beta = None

sample_weights(model)
    Collect sampled weights from the current model and maintain rolling buffer.

build_projection(weight_samples, k=None)
    Build projection subspace from collected weight samples. :param weight_samples: list of flat weight tensors
    [(D,), ...] :param k: number of basis vectors to keep. Defaults to min(N, D).

    Returns
        (D,) base weight vector (average) P: (D, k) projection matrix with top-k basis directions

    Return type
        mean_w

on_train_epoch_start(trainer, model)
    Track epoch.

on_train_epoch_end(trainer, model)
    Main TWA logic: build subspace and update in  $\beta$  space.

    # Math: # TWA weight update: #  $w_{twa} = \text{mean\_w} + P * \beta$  #  $\text{mean\_w} = (1/n) * \sum_i w_i$  (SWA baseline)
    #  $\beta \leftarrow (1 - \eta * \lambda) * \beta - \eta * P^T * g$  #  $g$  = gradient of training loss w.r.t. full model weights
    #  $\eta$  = learning rate,  $\lambda$  = ridge regularization #  $P$  = orthonormal basis spanning sampled checkpoints
    { $w_i$ }

on_fit_end(trainer, model)
    Replace with TWA model at the end.

```

14.2 Attributes

`__version__`

14.3 Classes

`Execute`

Executor class for training, retraining and evaluating KGE models.

continues on next page

Table 69 – continued from previous page

<i>KGE</i>	Knowledge Graph Embedding Class for interactive usage of pre-trained models
<i>QueryGenerator</i>	
<i>DICE_Trainer</i>	DICE_Trainer implement
<i>Evaluator</i>	Evaluator class for KGE models in various downstream tasks.

14.4 Package Contents

class `dicee.Execute` (*args*, *continuous_training*: *bool = False*)

Executor class for training, retraining and evaluating KGE models.

Handles the complete workflow: 1. Loading & Preprocessing & Serializing input data 2. Training & Validation & Testing 3. Storing all necessary information

args

Processed input arguments.

distributed

Whether distributed training is enabled.

rank

Process rank in distributed training.

world_size

Total number of processes.

local_rank

Local GPU rank.

trainer

Training handler instance.

trained_model

The trained model after training completes.

knowledge_graph

The loaded knowledge graph.

report

Dictionary storing training metrics and results.

evaluator

Model evaluation handler.

distributed

args

is_continual_training = `False`

trainer: `dicee.trainer.DICE_Trainer` | `None` = `None`

trained_model = `None`

knowledge_graph: `dicee.knowledge_graph.KG` | `None` = `None`

report: Dict

evaluator: `dicee.evaluator.Evaluator` | None = None

start_time: float | None = None

is_rank_zero() → bool

cleanup()

setup_executor() → None
 Set up storage directories for the experiment.
 Creates or reuses experiment directories based on configuration. Saves the configuration to a JSON file.

create_and_store_kg() → None
 Create knowledge graph and store as memory-mapped file.
 Only executed on rank 0 in distributed training. Skips if memmap already exists.

load_from_memmap() → None
 Load knowledge graph from memory-mapped file.

save_trained_model() → None
 Save a knowledge graph embedding model
 (1) Send model to eval mode and cpu.
 (2) Store the memory footprint of the model.
 (3) Save the model into disk.
 (4) Update the stats of KG again ?

Parameter

rtype
 None

end(*form_of_labelling: str*) → dict
 End training
 (1) Store trained model.
 (2) Report runtimes.
 (3) Eval model if required.

Parameter

rtype
 A dict containing information about the training and/or evaluation

write_report() → None
 Report training related information in a report.json file

start() → dict
 Start training
 # (1) Loading the Data # (2) Create an evaluator object. # (3) Create a trainer object. # (4) Start the training

Parameter

rtype

A dict containing information about the training and/or evaluation

```
class dicee.KGE (path=None, url=None, construct_ensemble=False, model_name=None)
```

Bases: `dicee.abstracts.BaseInteractiveKGE`, `dicee.abstracts.InteractiveQueryDecomposition`, `dicee.abstracts.BaseInteractiveTrainKGE`

Knowledge Graph Embedding Class for interactive usage of pre-trained models

`__str__()`

`to(device: str) → None`

`get_transductive_entity_embeddings(indices: torch.LongTensor | List[str], as_pytorch=False, as_numpy=False, as_list=True) → torch.FloatTensor | numpy.ndarray | List[float]`

`create_vector_database(collection_name: str, distance: str, location: str = 'localhost', port: int = 6333)`

`generate(h="", r="")`

`eval_lp_performance(dataset=List[Tuple[str, str, str]], filtered=True)`

`predict_missing_head_entity(relation: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False) → Tuple`

Given a relation and a tail entity, return top k ranked head entity.

$\operatorname{argmax}_{\{e \in E\}} f(e, r, t)$, where $r \in R$, $t \in E$.

Parameter

relation: Union[List[str], str]

String representation of selected relations.

tail_entity: Union[List[str], str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

`predict_missing_relations(head_entity: List[str] | str, tail_entity: List[str] | str, within=None, batch_size=2, topk=1, return_indices=False) → Tuple`

Given a head entity and a tail entity, return top k ranked relations.

$\operatorname{argmax}_{\{r \in R\}} f(h, r, t)$, where $h, t \in E$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

k: int

Highest ranked k entities.

Returns: Tuple

Highest K scores and entities

predict_missing_tail_entity (*head_entity: List[str] | str, relation: List[str] | str,*
within: List[str] = None, batch_size=2, topk=1, return_indices=False) → torch.FloatTensor

Given a head entity and a relation, return top k ranked entities

$\text{argmax}_{\{e \in E\}} f(h, r, e)$, where $h \in E$ and $r \in R$.

Parameter

head_entity: List[str]

String representation of selected entities.

tail_entity: List[str]

String representation of selected entities.

Returns: Tuple

scores

predict (*, *h: List[str] | str = None, r: List[str] | str = None, t: List[str] | str = None, within=None,*
logits=True) → torch.FloatTensor

Parameters

- **logits**
- **h**
- **r**
- **t**
- **within**

predict_topk (*, *h: str | List[str] = None, r: str | List[str] = None, t: str | List[str] = None, topk: int = 10,*
within: List[str] = None, batch_size: int = 1024)

Predict missing item in a given triple.

Returns

- If you query a single (h, r, ?) or (?, r, t) or (h, ?, t), returns List[(item, score)]
- If you query a batch of B, returns List of B such lists.

triple_score (*h: List[str] | str = None, r: List[str] | str = None, t: List[str] | str = None, logits=False*)
→ torch.FloatTensor

Predict triple score

Parameter

head_entity: List[str]

String representation of selected entities.

relation: List[str]

String representation of selected relations.

tail_entity: List[str]

String representation of selected entities.

logits: bool

If logits is True, unnormalized score returned

Returns: Tuple

pytorch tensor of triple score

return_multi_hop_query_results (*aggregated_query_for_all_entities*, *k*: int, *only_scores*)

single_hop_query_answering (*query*: tuple, *only_scores*: bool = True, *k*: int = None)

answer_multi_hop_query (*query_type*: str = None, *query*: Tuple[str | Tuple[str, str], Ellipsis] = None, *queries*: List[Tuple[str | Tuple[str, str], Ellipsis]] = None, *tnorm*: str = 'prod', *neg_norm*: str = 'standard', *lambda_*: float = 0.0, *k*: int = 10, *only_scores*=False) → List[Tuple[str, torch.Tensor]]

@TODO: Refactoring is needed # @TODO: Score computation for each query type should be done in a static function

Find an answer set for EPFO queries including negation and disjunction

Parameter

query_type: str The type of the query, e.g., “2p”.

query: Union[str, Tuple[str, Tuple[str, str]]] The query itself, either a string or a nested tuple.

queries: List of Tuple[Union[str, Tuple[str, str]], ...]

tnorm: str The t-norm operator.

neg_norm: str The negation norm.

lambda_: float lambda parameter for sugeno and yager negation norms

k: int The top-k substitutions for intermediate variables.

returns

- List[Tuple[str, torch.Tensor]]
- Entities and corresponding scores sorted in the descening order of scores

find_missing_triples (*confidence*: float, *entities*: List[str] = None, *relations*: List[str] = None, *topk*: int = 10, *at_most*: int = sys.maxsize) → Set

Find missing triples

Iterative over a set of entities E and a set of relation R :

orall e in E and orall r in R f(e,r,x)

```

    Return (e,r,x)
otin G and f(e,r,x) > confidence
    confidence: float
    A threshold for an output of a sigmoid function given a triple.
    topk: int
    Highest ranked k item to select triples with f(e,r,x) > confidence .
    at_most: int
    Stop after finding at_most missing triples
    {(e,r,x) | f(e,r,x) > confidence land (e,r,x)
otin G
predict_literals (entity: List[str] | str = None, attribute: List[str] | str = None,
    denormalize_preds: bool = True) → numpy.ndarray
    Predicts literal values for given entities and attributes.

Parameters
    • entity (Union[List[str], str]) – Entity or list of entities to predict literals for.
    • attribute (Union[List[str], str]) – Attribute or list of attributes to predict literals
        for.
    • denormalize_preds (bool) – If True, denormalizes the predictions.

Returns
    Predictions for the given entities and attributes.

Return type
    numpy ndarray

class dicee.QueryGenerator (train_path: str, val_path: str, test_path: str, ent2id: Dict = None,
    rel2id: Dict = None, seed: int = 1, gen_valid: bool = False, gen_test: bool = True)

    train_path
    val_path
    test_path
    gen_valid = False
    gen_test = True
    seed = 1
    max_ans_num = 1000000.0
    mode
    ent2id = None
    rel2id: Dict = None
    ent_in: Dict

```

```

ent_out: Dict

query_name_to_struct

list2tuple(list_data)

tuple2list(x: List | Tuple) → List | Tuple
    Convert a nested tuple to a nested list.

set_global_seed(seed: int)
    Set seed

construct_graph(paths: List[str]) → Tuple[Dict, Dict]
    Construct graph from triples Returns dicts with incoming and outgoing edges

fill_query(query_structure: List[str | List], ent_in: Dict, ent_out: Dict, answer: int) → bool
    Private method for fill_query logic.

achieve_answer(query: List[str | List], ent_in: Dict, ent_out: Dict) → set
    Private method for achieve_answer logic. @TODO: Document the code

write_links(ent_out, small_ent_out)

ground_queries(query_structure: List[str | List], ent_in: Dict, ent_out: Dict, small_ent_in: Dict,
               small_ent_out: Dict, gen_num: int, query_name: str)
    Generating queries and achieving answers

unmap(query_type, queries, tp_answers, fp_answers, fn_answers)

unmap_query(query_structure, query, id2ent, id2rel)

generate_queries(query_struct: List, gen_num: int, query_type: str)
    Passing incoming and outgoing edges to ground queries depending on mode [train valid or text] and getting
    queries and answers in return @ TODO: create a class for each single query struct

save_queries(query_type: str, gen_num: int, save_path: str)

abstractmethod load_queries(path)

get_queries(query_type: str, gen_num: int)

static save_queries_and_answers(path: str, data: List[Tuple[str, Tuple[collections.defaultdict]]])
    → None
    Save Queries into Disk

static load_queries_and_answers(path: str) → List[Tuple[str, Tuple[collections.defaultdict]]]
    Load Queries from Disk to Memory

class dicee.DICE_Trainer(args, is_continual_training: bool, storage_path, evaluator=None)

DICE_Trainer implement
1- Pytorch Lightning trainer (https://pytorch-lightning.readthedocs.io/en/stable/common/trainer.html)
2- Multi-GPU Trainer(https://pytorch.org/docs/stable/generated/torch.nn.parallel.DistributedDataParallel.html)
3- CPU Trainer

args
is_continual_training:bool
storage_path:str

```



```

    evaluator:
    report:dict
report
args
trainer = None
is_continual_training
storage_path
evaluator = None
form_of_labelling = None
continual_start (knowledge_graph)

```

- (1) Initialize training.
- (2) Load model
- (3) Load trainer (3) Fit model

Parameter

returns

- *model*
- **form_of_labelling** (*str*)

```

initialize_trainer (callbacks: List)
    → lightning.Trainer | dicke.trainer.model_parallelism.TensorParallel | dicke.trainer.torch_trainer.TorchTrainer | dicke.

```

Initialize Trainer from input arguments

```

initialize_or_load_model ()

```

```

init_dataloader (dataset: torch.utils.data.Dataset) → torch.utils.data.DataLoader

```

```

init_dataset () → torch.utils.data.Dataset

```

```

start (knowledge_graph: dicke.knowledge_graph.KG | numpy.memmap)
    → Tuple[dicke.models.base_model.BaseKGE, str]

```

Start the training

- (1) Initialize Trainer
- (2) Initialize or load a pretrained KGE model

in DDP setup, we need to load the memory map of already read/index KG.

```

k_fold_cross_validation (dataset) → Tuple[dicke.models.base_model.BaseKGE, str]

```

Perform K-fold Cross-Validation

1. Obtain K train and test splits.
2. **For each split,**
 - 2.1 initialize trainer and model
 - 2.2. Train model with configuration provided in args.
 - 2.3. Compute the mean reciprocal rank (MRR) score of the model on the test respective split.

3. Report the mean and average MRR .

Parameters

- **self**
- **dataset**

Returns

model

```
class dicee.Evaluator(args, is_continual_training: bool = False)
```

Evaluator class for KGE models in various downstream tasks.

Orchestrates link prediction evaluation with different scoring techniques including standard evaluation and byte-pair encoding based evaluation.

er_vocab

Entity-relation to tail vocabulary for filtered ranking.

re_vocab

Relation-entity (tail) to head vocabulary.

ee_vocab

Entity-entity to relation vocabulary.

num_entities

Total number of entities in the knowledge graph.

num_relations

Total number of relations in the knowledge graph.

args

Configuration arguments.

report

Dictionary storing evaluation results.

during_training

Whether evaluation is happening during training.

Example

```
>>> from dicee.evaluation import Evaluator
>>> evaluator = Evaluator(args)
>>> results = evaluator.eval(dataset, model, 'EntityPrediction')
>>> print(f"Test MRR: {results['Test']['MRR']:.4f}")
```

re_vocab: Dict | None = None

er_vocab: Dict | None = None

ee_vocab: Dict | None = None

func_triple_to_bpe_representation = None

is_continual_training = False

```

num_entities: int | None = None

num_relations: int | None = None

domain_constraints_per_rel = None

range_constraints_per_rel = None

args

report: Dict

during_training = False

```

vocab_preparation (*dataset*) → None

Prepare vocabularies from the dataset for evaluation.

Resolves any future objects and saves vocabularies to disk.

Parameters

dataset – Knowledge graph dataset with vocabulary attributes.

eval (*dataset, trained_model, form_of_labelling: str, during_training: bool = False*) → Dict | None

Evaluate the trained model on the dataset.

Parameters

- **dataset** – Knowledge graph dataset (KG instance).
- **trained_model** – The trained KGE model.
- **form_of_labelling** – Type of labelling ('EntityPrediction' or 'RelationPrediction').
- **during_training** – Whether evaluation is during training.

Returns

Dictionary of evaluation metrics, or None if evaluation is skipped.

eval_rank_of_head_and_tail_entity (*, *train_set, valid_set=None, test_set=None, trained_model*) → None

Evaluate with negative sampling scoring.

eval_rank_of_head_and_tail_byte_pair_encoded_entity (*, *train_set=None, valid_set=None, test_set=None, ordered_bpe_entities, trained_model*) → None

Evaluate with BPE-encoded entities and negative sampling.

eval_with_byte (*, *raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate Byte model with generation.

eval_with_bpe_vs_all (*, *raw_train_set, raw_valid_set=None, raw_test_set=None, trained_model, form_of_labelling*) → None

Evaluate with BPE and KvsAll scoring.

eval_with_vs_all (*, *train_set, valid_set=None, test_set=None, trained_model, form_of_labelling*) → None

Evaluate with KvsAll or 1vsAll scoring.

evaluate_lp_k_vs_all (*model, triple_idx, info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Filtered link prediction evaluation with KvsAll scoring.

Parameters

- **model** – The trained model to evaluate.
- **triple_idx** – Integer-indexed test triples.
- **info** – Description to print.
- **form_of_labelling** – ‘EntityPrediction’ or ‘RelationPrediction’.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp_with_byte (*model, triples: List[List[str]], info: str | None = None*) → Dict[str, float]

Evaluate Byte model with text generation.

Parameters

- **model** – Byte model.
- **triples** – String triples.
- **info** – Description to print.

Returns

Dictionary with placeholder metrics (-1 values).

evaluate_lp_bpe_k_vs_all (*model, triples: List[List[str]], info: str | None = None, form_of_labelling: str | None = None*) → Dict[str, float]

Evaluate BPE model with KvsAll scoring.

Parameters

- **model** – BPE-enabled model.
- **triples** – String triples.
- **info** – Description to print.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

evaluate_lp (*model, triple_idx, info: str*) → Dict[str, float]

Evaluate link prediction with negative sampling.

Parameters

- **model** – The model to evaluate.
- **triple_idx** – Integer-indexed triples.
- **info** – Description to print.

Returns

Dictionary with H@1, H@3, **H@10**, and MRR metrics.

dummy_eval (*trained_model, form_of_labelling: str*) → None

Run evaluation from saved data (for continual training).

Parameters

- **trained_model** – The trained model.
- **form_of_labelling** – Type of labelling.

eval_with_data (*dataset, trained_model, triple_idx: numpy.ndarray, form_of_labelling: str*)
→ Dict[str, float]

Evaluate a trained model on a given dataset.

Parameters

- **dataset** – Knowledge graph dataset.
- **trained_model** – The trained model.
- **triple_idx** – Integer-indexed triples to evaluate.
- **form_of_labelling** – Type of labelling.

Returns

Dictionary with evaluation metrics.

Raises

ValueError – If scoring technique is invalid.

`dicee.__version__ = '0.3.2'`

Python Module Index

d

- `dicee`, 12
- `dicee.__main__`, 12
- `dicee.abstracts`, 12
- `dicee.analyse_experiments`, 19
- `dicee.callbacks`, 21
- `dicee.config`, 27
- `dicee.dataset_classes`, 31
- `dicee.eval_static_funcs`, 39
- `dicee.evaluation`, 43
 - `dicee.evaluation.ensemble`, 43
 - `dicee.evaluation.evaluator`, 44
 - `dicee.evaluation.link_prediction`, 48
 - `dicee.evaluation.literal_prediction`, 50
 - `dicee.evaluation.utils`, 51
- `dicee.evaluator`, 62
- `dicee.executer`, 65
- `dicee.knowledge_graph`, 68
- `dicee.knowledge_graph_embeddings`, 71
- `dicee.models`, 74
 - `dicee.models.adopt`, 74
 - `dicee.models.base_model`, 83
 - `dicee.models.clifford`, 92
 - `dicee.models.complex`, 101
 - `dicee.models.dualE`, 103
 - `dicee.models.ensemble`, 105
 - `dicee.models.function_space`, 105
 - `dicee.models.literal`, 109
 - `dicee.models.octonion`, 110
 - `dicee.models.pykeen_models`, 113
 - `dicee.models.quaternion`, 115
 - `dicee.models.real`, 118
 - `dicee.models.static_funcs`, 120
 - `dicee.models.transformers`, 121
- `dicee.query_generator`, 180
- `dicee.read_preprocess_save_load_kg`, 182
- `dicee.read_preprocess_save_load_kg.preprocess`, 182
- `dicee.read_preprocess_save_load_kg.read_from_disk`, 183
- `dicee.read_preprocess_save_load_kg.save_load_disk`, 183
- `dicee.read_preprocess_save_load_kg.util`, 184
- `dicee.sanity_checkers`, 188
- `dicee.scripts`, 189
 - `dicee.scripts.index_serve`, 189
 - `dicee.scripts.run`, 190
- `dicee.static_funcs`, 190
- `dicee.static_funcs_training`, 195
- `dicee.static_preprocess_funcs`, 196
- `dicee.trainer`, 197
 - `dicee.trainer.dice_trainer`, 197
 - `dicee.trainer.model_parallelism`, 199
 - `dicee.trainer.torch_trainer`, 200
 - `dicee.trainer.torch_trainer_ddp`, 201
- `dicee.weight_averaging`, 204

Index

Non-alphabetical

`__call__()` (*dicee.models.base_model.IdentityClass method*), 92
`__call__()` (*dicee.models.ensemble.EnsembleKGE method*), 105
`__call__()` (*dicee.models.IdentityClass method*), 141, 155, 160
`__getitem__()` (*dicee.dataset_classes.AllvsAll method*), 34
`__getitem__()` (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 32
`__getitem__()` (*dicee.dataset_classes.FixedNegSampleDataset method*), 36
`__getitem__()` (*dicee.dataset_classes.KvsAll method*), 35
`__getitem__()` (*dicee.dataset_classes.KvsSampleDataset method*), 35
`__getitem__()` (*dicee.dataset_classes.LiteralDataset method*), 38
`__getitem__()` (*dicee.dataset_classes.MultiClassClassificationDataset method*), 33
`__getitem__()` (*dicee.dataset_classes.MultiLabelDataset method*), 33
`__getitem__()` (*dicee.dataset_classes.OnevsAllDataset method*), 36
`__getitem__()` (*dicee.dataset_classes.OnevsSample method*), 37
`__getitem__()` (*dicee.dataset_classes.TriplePredictionDataset method*), 38
`__iter__()` (*dicee.config.Namespace method*), 31
`__iter__()` (*dicee.knowledge_graph.KG method*), 70
`__iter__()` (*dicee.models.ensemble.EnsembleKGE method*), 105
`__len__()` (*dicee.dataset_classes.AllvsAll method*), 34
`__len__()` (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 32
`__len__()` (*dicee.dataset_classes.FixedNegSampleDataset method*), 36
`__len__()` (*dicee.dataset_classes.KvsAll method*), 35
`__len__()` (*dicee.dataset_classes.KvsSampleDataset method*), 35
`__len__()` (*dicee.dataset_classes.LiteralDataset method*), 38
`__len__()` (*dicee.dataset_classes.MultiClassClassificationDataset method*), 33
`__len__()` (*dicee.dataset_classes.MultiLabelDataset method*), 33
`__len__()` (*dicee.dataset_classes.OnevsAllDataset method*), 36
`__len__()` (*dicee.dataset_classes.OnevsSample method*), 37
`__len__()` (*dicee.dataset_classes.TriplePredictionDataset method*), 38
`__len__()` (*dicee.knowledge_graph.KG method*), 71
`__len__()` (*dicee.models.ensemble.EnsembleKGE method*), 105
`__setstate__()` (*dicee.models.ADOPT method*), 131
`__setstate__()` (*dicee.models.adopt.ADOPT method*), 78
`__str__()` (*dicee.KGE method*), 212
`__str__()` (*dicee.knowledge_graph_embeddings.KGE method*), 71
`__str__()` (*dicee.models.ensemble.EnsembleKGE method*), 105
`__version__` (*in module dicee*), 221

A

`AbstractCallback` (*class in dicee.abstracts*), 16
`AbstractPPECallback` (*class in dicee.abstracts*), 18
`AbstractTrainer` (*class in dicee.abstracts*), 13
`AccumulateEpochLossCallback` (*class in dicee.callbacks*), 21
`achieve_answer()` (*dicee.query_generator.QueryGenerator method*), 181
`achieve_answer()` (*dicee.QueryGenerator method*), 216
`AConEx` (*class in dicee.models*), 150
`AConEx` (*class in dicee.models.complex*), 102
`AConvO` (*class in dicee.models*), 163
`AConvO` (*class in dicee.models.octonion*), 113
`AConvQ` (*class in dicee.models*), 157
`AConvQ` (*class in dicee.models.quaternion*), 117
`adaptive_lr` (*dicee.config.Namespace attribute*), 31
`adaptive_swa` (*dicee.config.Namespace attribute*), 30
`add_new_entity_embeddings()` (*dicee.abstracts.BaseInteractiveKGE method*), 16
`add_noise_rate` (*dicee.config.Namespace attribute*), 28
`add_noise_rate` (*dicee.knowledge_graph.KG attribute*), 69
`add_noisy_triples()` (*in module dicee.static_funcs*), 193
`add_noisy_triples_into_training()` (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk method*), 183
`add_noisy_triples_into_training()` (*dicee.read_preprocess_save_load_kg.ReadFromDisk method*), 188
`add_reciprocal` (*dicee.knowledge_graph.KG attribute*), 69
`ADOPT` (*class in dicee.models*), 129
`ADOPT` (*class in dicee.models.adopt*), 76
`adopt()` (*in module dicee.models.adopt*), 79
`ALL_HITS_RANGE` (*in module dicee.evaluation*), 62
`ALL_HITS_RANGE` (*in module dicee.evaluation.utils*), 52

AllvsAll (class in *dicce.dataset_classes*), 33
 alphas (*dicce.abstracts.AbstractPPECallback* attribute), 18
 alphas (*dicce.weight_averaging.ASWA* attribute), 205
 analyse () (in module *dicce.analyse_experiments*), 20
 answer_multi_hop_query () (*dicce.KGE* method), 214
 answer_multi_hop_query () (*dicce.knowledge_graph_embeddings.KGE* method), 73
 app (in module *dicce.scripts.index_serve*), 189
 apply_coefficients () (*dicce.models.clifford.DeCaL* method), 99
 apply_coefficients () (*dicce.models.clifford.Keci* method), 94
 apply_coefficients () (*dicce.models.DeCaL* method), 168
 apply_coefficients () (*dicce.models.Keci* method), 165
 apply_reciprocal_or_noise () (in module *dicce.read_preprocess_save_load_kg.util*), 186
 apply_semantic_constraint (*dicce.abstracts.BaseInteractiveKGE* attribute), 15
 apply_unit_norm (*dicce.models.base_model.BaseKGE* attribute), 90
 apply_unit_norm (*dicce.models.BaseKGE* attribute), 139, 142, 148, 152, 158, 171, 175
 args (*dicce.DICE_Trainer* attribute), 217
 args (*dicce.evaluation.Evaluator* attribute), 54, 55
 args (*dicce.evaluation.evaluator.Evaluator* attribute), 45
 args (*dicce.Evaluator* attribute), 218, 219
 args (*dicce.evaluator.Evaluator* attribute), 63
 args (*dicce.Execute* attribute), 210
 args (*dicce.executer.Execute* attribute), 66
 args (*dicce.models.base_model.BaseKGE* attribute), 90
 args (*dicce.models.base_model.IdentityClass* attribute), 92
 args (*dicce.models.BaseKGE* attribute), 138, 142, 147, 152, 158, 171, 174
 args (*dicce.models.ensemble.EnsembleKGE* attribute), 105
 args (*dicce.models.IdentityClass* attribute), 141, 154, 160
 args (*dicce.models.pykeen_models.PykeenKGE* attribute), 114
 args (*dicce.models.PykeenKGE* attribute), 173
 args (*dicce.trainer.DICE_Trainer* attribute), 203
 args (*dicce.trainer.dice_trainer.DICE_Trainer* attribute), 198
 ASWA (class in *dicce.weight_averaging*), 204
 aswa (*dicce.analyse_experiments.Experiment* attribute), 20
 attn (*dicce.models.Block* attribute), 144
 attn (*dicce.models.clifford.TransformerBlock* attribute), 96
 attn (*dicce.models.transformers.Block* attribute), 126
 attn_dropout (*dicce.models.clifford.TransformerSelfAttention* attribute), 97
 attn_dropout (*dicce.models.transformers.SelfAttention* attribute), 124
 attributes (*dicce.abstracts.AbstractTrainer* attribute), 13
 auto_batch_finding (*dicce.config.Namespace* attribute), 30

B

backend (*dicce.config.Namespace* attribute), 29
 backend (*dicce.knowledge_graph.KG* attribute), 69
 base_weights (*dicce.weight_averaging.TWA* attribute), 209
 BaseInteractiveKGE (class in *dicce.abstracts*), 14
 BaseInteractiveTrainKGE (class in *dicce.abstracts*), 18
 BaseKGE (class in *dicce.models*), 138, 141, 147, 151, 157, 171, 174
 BaseKGE (class in *dicce.models.base_model*), 89
 BaseKGELightning (class in *dicce.models*), 132
 BaseKGELightning (class in *dicce.models.base_model*), 83
 batch_kronecker_product () (*dicce.callbacks.KronE* static method), 25
 batch_size (*dicce.analyse_experiments.Experiment* attribute), 20
 batch_size (*dicce.callbacks.PseudoLabellingCallback* attribute), 24
 batch_size (*dicce.config.Namespace* attribute), 28
 batches_per_epoch (*dicce.callbacks.LRScheduler* attribute), 27
 beta (*dicce.weight_averaging.TWA* attribute), 209
 bias (*dicce.models.CoKEConfig* attribute), 146
 bias (*dicce.models.real.CoKEConfig* attribute), 119, 120
 bias (*dicce.models.transformers.GPTConfig* attribute), 126
 bias (*dicce.models.transformers.LayerNorm* attribute), 123
 Block (class in *dicce.models*), 144
 Block (class in *dicce.models.transformers*), 125
 block_size (*dicce.config.Namespace* attribute), 30
 block_size (*dicce.dataset_classes.MultiClassClassificationDataset* attribute), 33
 block_size (*dicce.models.base_model.BaseKGE* attribute), 91
 block_size (*dicce.models.BaseKGE* attribute), 139, 143, 148, 153, 159, 172, 175

block_size (*dicee.models.CoKEConfig* attribute), 146
 block_size (*dicee.models.real.CoKEConfig* attribute), 119, 120
 block_size (*dicee.models.transformers.GPTConfig* attribute), 126
 blocks (*dicee.models.CoKE* attribute), 147
 blocks (*dicee.models.real.CoKE* attribute), 120
 bn_conv1 (*dicee.models.AConvQ* attribute), 157
 bn_conv1 (*dicee.models.ConvQ* attribute), 156
 bn_conv1 (*dicee.models.quaternion.AConvQ* attribute), 117
 bn_conv1 (*dicee.models.quaternion.ConvQ* attribute), 117
 bn_conv2 (*dicee.models.AConvQ* attribute), 157
 bn_conv2 (*dicee.models.ConvQ* attribute), 157
 bn_conv2 (*dicee.models.quaternion.AConvQ* attribute), 117
 bn_conv2 (*dicee.models.quaternion.ConvQ* attribute), 117
 bn_conv2d (*dicee.models.AConEx* attribute), 150
 bn_conv2d (*dicee.models.AConvO* attribute), 163
 bn_conv2d (*dicee.models.complex.AConEx* attribute), 102
 bn_conv2d (*dicee.models.complex.ConEx* attribute), 101
 bn_conv2d (*dicee.models.ConEx* attribute), 150
 bn_conv2d (*dicee.models.ConvO* attribute), 162
 bn_conv2d (*dicee.models.octonion.AConvO* attribute), 113
 bn_conv2d (*dicee.models.octonion.ConvO* attribute), 112
 BPE_NegativeSamplingDataset (*class in dicee.dataset_classes*), 32
 build_chain_funcs () (*dicee.models.FMult2* method), 178
 build_chain_funcs () (*dicee.models.function_space.FMult2* method), 107
 build_func () (*dicee.models.FMult2* method), 178
 build_func () (*dicee.models.function_space.FMult2* method), 107
 build_projection () (*dicee.weight_averaging.TWA* method), 209
 Byte (*class in dicee.models.transformers*), 121
 byte_pair_encoding (*dicee.analyse_experiments.Experiment* attribute), 20
 byte_pair_encoding (*dicee.config.Namespace* attribute), 30
 byte_pair_encoding (*dicee.knowledge_graph.KG* attribute), 69
 byte_pair_encoding (*dicee.models.base_model.BaseKGE* attribute), 91
 byte_pair_encoding (*dicee.models.BaseKGE* attribute), 139, 143, 148, 153, 159, 172, 175

C

c_attn (*dicee.models.clifford.TransformerSelfAttention* attribute), 97
 c_attn (*dicee.models.transformers.SelfAttention* attribute), 124
 c_fc (*dicee.models.clifford.TransformerMLP* attribute), 97
 c_fc (*dicee.models.transformers.MLP* attribute), 125
 c_proj (*dicee.models.clifford.TransformerMLP* attribute), 97
 c_proj (*dicee.models.clifford.TransformerSelfAttention* attribute), 97
 c_proj (*dicee.models.transformers.MLP* attribute), 125
 c_proj (*dicee.models.transformers.SelfAttention* attribute), 124
 callbacks (*dicee.abstracts.AbstractTrainer* attribute), 13
 callbacks (*dicee.analyse_experiments.Experiment* attribute), 20
 callbacks (*dicee.config.Namespace* attribute), 28
 callbacks (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 202
 causal (*dicee.models.CoKEConfig* attribute), 146
 causal (*dicee.models.real.CoKEConfig* attribute), 119, 120
 causal (*dicee.models.transformers.GPTConfig* attribute), 126
 causal (*dicee.models.transformers.SelfAttention* attribute), 124
 chain_func () (*dicee.models.FMult* method), 177
 chain_func () (*dicee.models.function_space.FMult* method), 106
 chain_func () (*dicee.models.function_space.GFMult* method), 106
 chain_func () (*dicee.models.GFMult* method), 177
 CKeci (*class in dicee.models*), 166
 CKeci (*class in dicee.models.clifford*), 97
 cl_pqr () (*dicee.models.clifford.DeCaL* method), 98
 cl_pqr () (*dicee.models.DeCaL* method), 168
 cleanup () (*dicee.Execute* method), 211
 cleanup () (*dicee.executer.Execute* method), 67
 clifford_multiplication () (*dicee.models.clifford.Keci* method), 94
 clifford_multiplication () (*dicee.models.Keci* method), 165
 clip_lambda (*dicee.models.ADOPT* attribute), 131
 clip_lambda (*dicee.models.adopt.ADOPT* attribute), 78
 CoKE (*class in dicee.models*), 146
 CoKE (*class in dicee.models.real*), 120

coke_dropout (*dicee.models.CoKE attribute*), 147
 coke_dropout (*dicee.models.real.CoKE attribute*), 120
 CoKEConfig (*class in dicee.models*), 146
 CoKEConfig (*class in dicee.models.real*), 119
 collate_fn (*dicee.dataset_classes.AllvsAll attribute*), 34
 collate_fn (*dicee.dataset_classes.FixedNegSampleDataset attribute*), 36
 collate_fn (*dicee.dataset_classes.KvsAll attribute*), 35
 collate_fn (*dicee.dataset_classes.KvsSampleDataset attribute*), 35
 collate_fn (*dicee.dataset_classes.MultiClassClassificationDataset attribute*), 33
 collate_fn (*dicee.dataset_classes.MultiLabelDataset attribute*), 33
 collate_fn (*dicee.dataset_classes.OnevsAllDataset attribute*), 36
 collate_fn (*dicee.dataset_classes.OnevsSample attribute*), 37
 collate_fn () (*dicee.dataset_classes.BPE_NegativeSamplingDataset method*), 32
 collate_fn () (*dicee.dataset_classes.TriplePredictionDataset method*), 38
 collection_name (*dicee.scripts.index_serve.NeuralSearcher attribute*), 189
 comp_func () (*dicee.models.function_space.LFMMult method*), 109
 comp_func () (*dicee.models.LFMMult method*), 179
 ComplEx (*class in dicee.models*), 150
 ComplEx (*class in dicee.models.complex*), 102
 compute_convergence () (*in module dicee.callbacks*), 24
 compute_func () (*dicee.models.FMMult method*), 177
 compute_func () (*dicee.models.FMMult2 method*), 178
 compute_func () (*dicee.models.function_space.FMMult method*), 106
 compute_func () (*dicee.models.function_space.FMMult2 method*), 107
 compute_func () (*dicee.models.function_space.GFMMult method*), 106
 compute_func () (*dicee.models.GFMMult method*), 177
 compute_metrics_from_ranks () (*in module dicee.evaluation*), 61
 compute_metrics_from_ranks () (*in module dicee.evaluation.utils*), 52
 compute_metrics_from_ranks_simple () (*in module dicee.evaluation*), 61
 compute_metrics_from_ranks_simple () (*in module dicee.evaluation.utils*), 52
 compute_mrr () (*dicee.weight_averaging.ASWA static method*), 205
 compute_sigma_pp () (*dicee.models.clifford.DeCaL method*), 99
 compute_sigma_pp () (*dicee.models.clifford.Keci method*), 93
 compute_sigma_pp () (*dicee.models.DeCaL method*), 169
 compute_sigma_pp () (*dicee.models.Keci method*), 164
 compute_sigma_pq () (*dicee.models.clifford.DeCaL method*), 100
 compute_sigma_pq () (*dicee.models.clifford.Keci method*), 94
 compute_sigma_pq () (*dicee.models.DeCaL method*), 170
 compute_sigma_pq () (*dicee.models.Keci method*), 165
 compute_sigma_pr () (*dicee.models.clifford.DeCaL method*), 101
 compute_sigma_pr () (*dicee.models.DeCaL method*), 170
 compute_sigma_qq () (*dicee.models.clifford.DeCaL method*), 100
 compute_sigma_qq () (*dicee.models.clifford.Keci method*), 94
 compute_sigma_qq () (*dicee.models.DeCaL method*), 169
 compute_sigma_qq () (*dicee.models.Keci method*), 164
 compute_sigma_qr () (*dicee.models.clifford.DeCaL method*), 101
 compute_sigma_qr () (*dicee.models.DeCaL method*), 170
 compute_sigma_rr () (*dicee.models.clifford.DeCaL method*), 100
 compute_sigma_rr () (*dicee.models.DeCaL method*), 169
 compute_sigmas_multivect () (*dicee.models.clifford.DeCaL method*), 99
 compute_sigmas_multivect () (*dicee.models.DeCaL method*), 168
 compute_sigmas_single () (*dicee.models.clifford.DeCaL method*), 98
 compute_sigmas_single () (*dicee.models.DeCaL method*), 168
 ConEx (*class in dicee.models*), 149
 ConEx (*class in dicee.models.complex*), 101
 config (*dicee.models.CoKE attribute*), 147
 config (*dicee.models.real.CoKE attribute*), 120
 config (*dicee.models.transformers.BytE attribute*), 122
 config (*dicee.models.transformers.GPT attribute*), 127
 configs (*dicee.abstracts.BaseInteractiveKGE attribute*), 15
 configure_optimizers () (*dicee.models.base_model.BaseKGELightning method*), 88
 configure_optimizers () (*dicee.models.BaseKGELightning method*), 136
 configure_optimizers () (*dicee.models.transformers.GPT method*), 127
 construct_batch_selected_cl_multivector () (*dicee.models.clifford.Keci method*), 95
 construct_batch_selected_cl_multivector () (*dicee.models.Keci method*), 166
 construct_cl_multivector () (*dicee.models.clifford.DeCaL method*), 99
 construct_cl_multivector () (*dicee.models.clifford.Keci method*), 95
 construct_cl_multivector () (*dicee.models.DeCaL method*), 168

`construct_cl_multivector()` (*dicee.models.Keci method*), 165
`construct_dataset()` (*in module dicee.dataset_classes*), 39
`construct_ensemble` (*dicee.abstracts.BaseInteractiveKGE attribute*), 14
`construct_graph()` (*dicee.query_generator.QueryGenerator method*), 181
`construct_graph()` (*dicee.QueryGenerator method*), 216
`construct_input_and_output()` (*dicee.abstracts.BaseInteractiveKGE method*), 16
`construct_multi_coeff()` (*dicee.models.function_space.LFMMult method*), 108
`construct_multi_coeff()` (*dicee.models.LFMMult method*), 179
`continual_learning` (*dicee.config.Namespace attribute*), 30
`continual_start()` (*dicee.DICE_Trainer method*), 217
`continual_start()` (*dicee.executer.ContinuousExecute method*), 68
`continual_start()` (*dicee.trainer.DICE_Trainer method*), 203
`continual_start()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 198
`continual_training_setup_executor()` (*in module dicee.static_funcs*), 194
`ContinuousExecute` (*class in dicee.executer*), 68
`conv2d` (*dicee.models.AConEx attribute*), 150
`conv2d` (*dicee.models.AConvO attribute*), 163
`conv2d` (*dicee.models.AConvQ attribute*), 157
`conv2d` (*dicee.models.complex.AConEx attribute*), 102
`conv2d` (*dicee.models.complex.ConEx attribute*), 101
`conv2d` (*dicee.models.ConEx attribute*), 149
`conv2d` (*dicee.models.ConvO attribute*), 162
`conv2d` (*dicee.models.ConvQ attribute*), 156
`conv2d` (*dicee.models.octonion.AConvO attribute*), 113
`conv2d` (*dicee.models.octonion.ConvO attribute*), 112
`conv2d` (*dicee.models.quaternion.AConvQ attribute*), 117
`conv2d` (*dicee.models.quaternion.ConvQ attribute*), 117
`ConvO` (*class in dicee.models*), 162
`ConvO` (*class in dicee.models.octonion*), 112
`ConvQ` (*class in dicee.models*), 156
`ConvQ` (*class in dicee.models.quaternion*), 116
`count_triples()` (*in module dicee.read_preprocess_save_load_kg.util*), 186
`create_and_store_kg()` (*dicee.Execute method*), 211
`create_and_store_kg()` (*dicee.executer.Execute method*), 67
`create_constraints()` (*in module dicee.read_preprocess_save_load_kg.util*), 186
`create_constraints()` (*in module dicee.static_preprocess_funcs*), 197
`create_experiment_folder()` (*in module dicee.static_funcs*), 194
`create_hits_dict()` (*in module dicee.evaluation*), 62
`create_hits_dict()` (*in module dicee.evaluation.utils*), 53
`create_random_data()` (*dicee.callbacks.PseudoLabellingCallback method*), 24
`create_recipriocal_triples()` (*in module dicee.read_preprocess_save_load_kg.util*), 187
`create_recipriocal_triples()` (*in module dicee.static_funcs*), 191
`create_vector_database()` (*dicee.KGE method*), 212
`create_vector_database()` (*dicee.knowledge_graph_embeddings.KGE method*), 71
`crop_block_size()` (*dicee.models.transformers.GPT method*), 127
`ctx` (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
`current_epoch` (*dicee.weight_averaging.EMA attribute*), 208
`current_epoch` (*dicee.weight_averaging.SWA attribute*), 206
`current_epoch` (*dicee.weight_averaging.SWAG attribute*), 207
`current_epoch` (*dicee.weight_averaging.TWA attribute*), 209
`cycle_length` (*dicee.callbacks.LRScheduler attribute*), 27

D

`data_module` (*dicee.callbacks.PseudoLabellingCallback attribute*), 24
`data_property_embeddings` (*dicee.models.literal.LiteralEmbeddings attribute*), 110
`dataset_dir` (*dicee.config.Namespace attribute*), 28
`dataset_dir` (*dicee.knowledge_graph.KG attribute*), 68, 69
`dataset_sanity_checking()` (*in module dicee.read_preprocess_save_load_kg.util*), 187
`DeCaL` (*class in dicee.models*), 166
`DeCaL` (*class in dicee.models.clifford*), 97
`decay` (*dicee.weight_averaging.EMA attribute*), 208
`decide()` (*dicee.weight_averaging.ASWA method*), 205
`default_eval_model` (*dicee.callbacks.PeriodicEvalCallback attribute*), 26
`DEFAULT_HITS_RANGE` (*in module dicee.evaluation*), 62
`DEFAULT_HITS_RANGE` (*in module dicee.evaluation.utils*), 52
`degree` (*dicee.models.function_space.LFMMult attribute*), 108
`degree` (*dicee.models.LFMMult attribute*), 178

- `denormalize()` (*dicee.dataset_classes.LiteralDataset static method*), 38
- `describe()` (*dicee.knowledge_graph.KG method*), 70
- `description_of_input` (*dicee.knowledge_graph.KG attribute*), 70
- `deviations` (*dicee.weight_averaging.SWAG attribute*), 207
- `device` (*dicee.models.literal.LiteralEmbeddings property*), 110
- `DICE_Trainer` (*class in dicee*), 216
- `DICE_Trainer` (*class in dicee.trainer*), 203
- `DICE_Trainer` (*class in dicee.trainer.dice_trainer*), 198
- `dicee`
 - module, 12
- `dicee.__main__`
 - module, 12
- `dicee.abstracts`
 - module, 12
- `dicee.analyse_experiments`
 - module, 19
- `dicee.callbacks`
 - module, 21
- `dicee.config`
 - module, 27
- `dicee.dataset_classes`
 - module, 31
- `dicee.eval_static_funcs`
 - module, 39
- `dicee.evaluation`
 - module, 43
- `dicee.evaluation.ensemble`
 - module, 43
- `dicee.evaluation.evaluator`
 - module, 44
- `dicee.evaluation.link_prediction`
 - module, 48
- `dicee.evaluation.literal_prediction`
 - module, 50
- `dicee.evaluation.utils`
 - module, 51
- `dicee.evaluator`
 - module, 62
- `dicee.executer`
 - module, 65
- `dicee.knowledge_graph`
 - module, 68
- `dicee.knowledge_graph_embeddings`
 - module, 71
- `dicee.models`
 - module, 74
- `dicee.models.adopt`
 - module, 74
- `dicee.models.base_model`
 - module, 83
- `dicee.models.clifford`
 - module, 92
- `dicee.models.complex`
 - module, 101
- `dicee.models.dualE`
 - module, 103
- `dicee.models.ensemble`
 - module, 105
- `dicee.models.function_space`
 - module, 105
- `dicee.models.literal`
 - module, 109
- `dicee.models.octonion`
 - module, 110
- `dicee.models.pykeen_models`
 - module, 113
- `dicee.models.quaternion`
 - module, 115

- `dicee.models.real`
 - module, 118
- `dicee.models.static_funcs`
 - module, 120
- `dicee.models.transformers`
 - module, 121
- `dicee.query_generator`
 - module, 180
- `dicee.read_preprocess_save_load_kg`
 - module, 182
- `dicee.read_preprocess_save_load_kg.preprocess`
 - module, 182
- `dicee.read_preprocess_save_load_kg.read_from_disk`
 - module, 183
- `dicee.read_preprocess_save_load_kg.save_load_disk`
 - module, 183
- `dicee.read_preprocess_save_load_kg.util`
 - module, 184
- `dicee.sanity_checkers`
 - module, 188
- `dicee.scripts`
 - module, 189
- `dicee.scripts.index_serve`
 - module, 189
- `dicee.scripts.run`
 - module, 190
- `dicee.static_funcs`
 - module, 190
- `dicee.static_funcs_training`
 - module, 195
- `dicee.static_preprocess_funcs`
 - module, 196
- `dicee.trainer`
 - module, 197
- `dicee.trainer.dice_trainer`
 - module, 197
- `dicee.trainer.model_parallelism`
 - module, 199
- `dicee.trainer.torch_trainer`
 - module, 200
- `dicee.trainer.torch_trainer_ddp`
 - module, 201
- `dicee.weight_averaging`
 - module, 204
- `discrete_points (dicee.models.FMult2 attribute)`, 177
- `discrete_points (dicee.models.function_space.FMult2 attribute)`, 107
- `dist_func (dicee.models.Pyke attribute)`, 145
- `dist_func (dicee.models.real.Pyke attribute)`, 119
- `DistMult (class in dicee.models)`, 144
- `DistMult (class in dicee.models.real)`, 118
- `distributed (dicee.Execute attribute)`, 210
- `distributed (dicee.executer.Execute attribute)`, 66
- `domain_constraints_per_rel (dicee.evaluation.Evaluator attribute)`, 55
- `domain_constraints_per_rel (dicee.evaluation.evaluator.Evaluator attribute)`, 45
- `domain_constraints_per_rel (dicee.Evaluator attribute)`, 219
- `domain_constraints_per_rel (dicee.evaluator.Evaluator attribute)`, 63
- `download_file()` (in module `dicee.static_funcs`), 194
- `download_files_from_url()` (in module `dicee.static_funcs`), 194
- `download_pretrained_model()` (in module `dicee.static_funcs`), 194
- `dropout (dicee.models.clifford.TransformerMLP attribute)`, 97
- `dropout (dicee.models.clifford.TransformerSelfAttention attribute)`, 97
- `dropout (dicee.models.CoKEConfig attribute)`, 146
- `dropout (dicee.models.literal.LiteralEmbeddings attribute)`, 109, 110
- `dropout (dicee.models.real.CoKEConfig attribute)`, 119, 120
- `dropout (dicee.models.transformers.GPTConfig attribute)`, 126
- `dropout (dicee.models.transformers.MLP attribute)`, 125
- `dropout (dicee.models.transformers.SelfAttention attribute)`, 124
- `DualE (class in dicee.models)`, 179

DualE (class in *dicее.models.dualE*), 104
 dummy_eval () (*dicее.evaluation.Evaluator* method), 57
 dummy_eval () (*dicее.evaluation.evaluator.Evaluator* method), 47
 dummy_eval () (*dicее.Evaluator* method), 220
 dummy_eval () (*dicее.evaluator.Evaluator* method), 65
 dummy_id (*dicее.knowledge_graph.KG* attribute), 70
 during_training (*dicее.evaluation.Evaluator* attribute), 54, 55
 during_training (*dicее.evaluation.evaluator.Evaluator* attribute), 45, 46
 during_training (*dicее.Evaluator* attribute), 218, 219
 during_training (*dicее.evaluator.Evaluator* attribute), 63

E

ee_vocab (*dicее.evaluation.Evaluator* attribute), 54, 55
 ee_vocab (*dicее.evaluation.evaluator.Evaluator* attribute), 45
 ee_vocab (*dicее.Evaluator* attribute), 218
 ee_vocab (*dicее.evaluator.Evaluator* attribute), 62, 63
 efficient_zero_grad () (in module *dicее.evaluation.utils*), 53
 efficient_zero_grad () (in module *dicее.static_funcs_training*), 196
 EMA (class in *dicее.weight_averaging*), 207
 ema (*dicее.config.Namespace* attribute), 30
 ema_c_epochs (*dicее.weight_averaging.EMA* attribute), 208
 ema_model (*dicее.weight_averaging.EMA* attribute), 208
 ema_start_epoch (*dicее.weight_averaging.EMA* attribute), 208
 ema_update () (*dicее.weight_averaging.EMA* static method), 208
 embedding_dim (*dicее.analyse_experiments.Experiment* attribute), 20
 embedding_dim (*dicее.config.Namespace* attribute), 28
 embedding_dim (*dicее.models.base_model.BaseKGE* attribute), 90
 embedding_dim (*dicее.models.BaseKGE* attribute), 138, 142, 147, 152, 158, 171, 175
 embedding_dim (*dicее.models.literal.LiteralEmbeddings* attribute), 110
 embedding_dims (*dicее.models.literal.LiteralEmbeddings* attribute), 109
 enable_log (in module *dicее.static_preprocess_funcs*), 196
 enc (*dicее.knowledge_graph.KG* attribute), 70
 end () (*dicее.Execute* method), 211
 end () (*dicее.executer.Execute* method), 67
 EnsembleKGE (class in *dicее.models.ensemble*), 105
 ent2id (*dicее.query_generator.QueryGenerator* attribute), 181
 ent2id (*dicее.QueryGenerator* attribute), 215
 ent_in (*dicее.query_generator.QueryGenerator* attribute), 181
 ent_in (*dicее.QueryGenerator* attribute), 215
 ent_out (*dicее.query_generator.QueryGenerator* attribute), 181
 ent_out (*dicее.QueryGenerator* attribute), 215
 entities_str (*dicее.knowledge_graph.KG* property), 70
 entity_embeddings (*dicее.models.AConvQ* attribute), 157
 entity_embeddings (*dicее.models.clifford.DeCaL* attribute), 98
 entity_embeddings (*dicее.models.ConvQ* attribute), 156
 entity_embeddings (*dicее.models.DeCaL* attribute), 167
 entity_embeddings (*dicее.models.DualE* attribute), 180
 entity_embeddings (*dicее.models.dualE.DualE* attribute), 104
 entity_embeddings (*dicее.models.FMult* attribute), 176
 entity_embeddings (*dicее.models.FMult2* attribute), 177
 entity_embeddings (*dicее.models.function_space.FMult* attribute), 106
 entity_embeddings (*dicее.models.function_space.FMult2* attribute), 107
 entity_embeddings (*dicее.models.function_space.GFMult* attribute), 106
 entity_embeddings (*dicее.models.function_space.LFMult* attribute), 108
 entity_embeddings (*dicее.models.function_space.LFMult1* attribute), 107
 entity_embeddings (*dicее.models.GFMult* attribute), 177
 entity_embeddings (*dicее.models.LFMult* attribute), 178
 entity_embeddings (*dicее.models.LFMult1* attribute), 178
 entity_embeddings (*dicее.models.literal.LiteralEmbeddings* attribute), 109, 110
 entity_embeddings (*dicее.models.pykeen_models.PykeenKGE* attribute), 114
 entity_embeddings (*dicее.models.PykeenKGE* attribute), 173
 entity_embeddings (*dicее.models.quaternion.AConvQ* attribute), 117
 entity_embeddings (*dicее.models.quaternion.ConvQ* attribute), 117
 entity_to_idx (*dicее.dataset_classes.LiteralDataset* attribute), 38
 entity_to_idx (*dicее.knowledge_graph.KG* attribute), 69
 entity_to_idx (*dicее.scripts.index_serve.NeuralSearcher* attribute), 189
 epoch_count (*dicее.abstracts.AbstractPPECallback* attribute), 18

epoch_count (*dicee.weight_averaging.ASWA attribute*), 204
 epoch_counter (*dicee.callbacks.Eval attribute*), 24
 epoch_counter (*dicee.callbacks.KGESaveCallback attribute*), 23
 epoch_counter (*dicee.callbacks.PeriodicEvalCallback attribute*), 26
 epoch_ratio (*dicee.callbacks.Eval attribute*), 24
 er_vocab (*dicee.evaluation.Evaluator attribute*), 54, 55
 er_vocab (*dicee.evaluation.evaluator.Evaluator attribute*), 45
 er_vocab (*dicee.Evaluator attribute*), 218
 er_vocab (*dicee.evaluator.Evaluator attribute*), 62, 63
 estimate_mfu() (*dicee.models.transformers.GPT method*), 127
 estimate_q() (*in module dicee.callbacks*), 24
 Eval (*class in dicee.callbacks*), 24
 eval() (*dicee.evaluation.Evaluator method*), 55
 eval() (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval() (*dicee.Evaluator method*), 219
 eval() (*dicee.evaluator.Evaluator method*), 63
 eval() (*dicee.models.ensemble.EnsembleKGE method*), 105
 eval_at_epochs (*dicee.config.Namespace attribute*), 31
 eval_epochs (*dicee.callbacks.PeriodicEvalCallback attribute*), 26
 eval_every_n_epochs (*dicee.config.Namespace attribute*), 30
 eval_lp_performance() (*dicee.KGE method*), 212
 eval_lp_performance() (*dicee.knowledge_graph_embeddings.KGE method*), 71
 eval_model (*dicee.config.Namespace attribute*), 29
 eval_model (*dicee.knowledge_graph.KG attribute*), 69
 eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.evaluation.Evaluator method*), 55
 eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.Evaluator method*), 219
 eval_rank_of_head_and_tail_byte_pair_encoded_entity() (*dicee.evaluator.Evaluator method*), 64
 eval_rank_of_head_and_tail_entity() (*dicee.evaluation.Evaluator method*), 55
 eval_rank_of_head_and_tail_entity() (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_rank_of_head_and_tail_entity() (*dicee.Evaluator method*), 219
 eval_rank_of_head_and_tail_entity() (*dicee.evaluator.Evaluator method*), 64
 eval_with_bpe_vs_all() (*dicee.evaluation.Evaluator method*), 56
 eval_with_bpe_vs_all() (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_with_bpe_vs_all() (*dicee.Evaluator method*), 219
 eval_with_bpe_vs_all() (*dicee.evaluator.Evaluator method*), 64
 eval_with_byte() (*dicee.evaluation.Evaluator method*), 55
 eval_with_byte() (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_with_byte() (*dicee.Evaluator method*), 219
 eval_with_byte() (*dicee.evaluator.Evaluator method*), 64
 eval_with_data() (*dicee.evaluation.Evaluator method*), 57
 eval_with_data() (*dicee.evaluation.evaluator.Evaluator method*), 47
 eval_with_data() (*dicee.Evaluator method*), 220
 eval_with_data() (*dicee.evaluator.Evaluator method*), 65
 eval_with_vs_all() (*dicee.evaluation.Evaluator method*), 56
 eval_with_vs_all() (*dicee.evaluation.evaluator.Evaluator method*), 46
 eval_with_vs_all() (*dicee.Evaluator method*), 219
 eval_with_vs_all() (*dicee.evaluator.Evaluator method*), 64
 evaluate() (*in module dicee.static_funcs*), 194
 evaluate_bpe_lp() (*in module dicee.evaluation*), 59
 evaluate_bpe_lp() (*in module dicee.evaluation.link_prediction*), 50
 evaluate_bpe_lp() (*in module dicee.static_funcs_training*), 195
 evaluate_ensemble_link_prediction_performance() (*in module dicee.eval_static_funcs*), 42
 evaluate_ensemble_link_prediction_performance() (*in module dicee.evaluation*), 60
 evaluate_ensemble_link_prediction_performance() (*in module dicee.evaluation.ensemble*), 43
 evaluate_link_prediction_performance() (*in module dicee.eval_static_funcs*), 40
 evaluate_link_prediction_performance() (*in module dicee.evaluation*), 57
 evaluate_link_prediction_performance() (*in module dicee.evaluation.link_prediction*), 48
 evaluate_link_prediction_performance_with_bpe() (*in module dicee.eval_static_funcs*), 40
 evaluate_link_prediction_performance_with_bpe() (*in module dicee.evaluation*), 58
 evaluate_link_prediction_performance_with_bpe() (*in module dicee.evaluation.link_prediction*), 49
 evaluate_link_prediction_performance_with_bpe_reciprocals() (*in module dicee.eval_static_funcs*), 41
 evaluate_link_prediction_performance_with_bpe_reciprocals() (*in module dicee.evaluation*), 58
 evaluate_link_prediction_performance_with_bpe_reciprocals() (*in module dicee.evaluation.link_prediction*), 49
 evaluate_link_prediction_performance_with_reciprocals() (*in module dicee.eval_static_funcs*), 40
 evaluate_link_prediction_performance_with_reciprocals() (*in module dicee.evaluation*), 57
 evaluate_link_prediction_performance_with_reciprocals() (*in module dicee.evaluation.link_prediction*), 48
 evaluate_literal_prediction() (*in module dicee.eval_static_funcs*), 41

`evaluate_literal_prediction()` (in module `dicee.evaluation`), 59
`evaluate_literal_prediction()` (in module `dicee.evaluation.literal_prediction`), 51
`evaluate_lp()` (`dicee.evaluation.Evaluator` method), 56
`evaluate_lp()` (`dicee.evaluation.evaluator.Evaluator` method), 47
`evaluate_lp()` (`dicee.Evaluator` method), 220
`evaluate_lp()` (`dicee.evaluator.Evaluator` method), 65
`evaluate_lp()` (in module `dicee.evaluation`), 58
`evaluate_lp()` (in module `dicee.evaluation.link_prediction`), 49
`evaluate_lp()` (in module `dicee.static_funcs_training`), 195
`evaluate_lp_bpe_k_vs_all()` (`dicee.evaluation.Evaluator` method), 56
`evaluate_lp_bpe_k_vs_all()` (`dicee.evaluation.evaluator.Evaluator` method), 47
`evaluate_lp_bpe_k_vs_all()` (`dicee.Evaluator` method), 220
`evaluate_lp_bpe_k_vs_all()` (`dicee.evaluator.Evaluator` method), 64
`evaluate_lp_bpe_k_vs_all()` (in module `dicee.eval_static_funcs`), 41
`evaluate_lp_bpe_k_vs_all()` (in module `dicee.evaluation`), 59
`evaluate_lp_bpe_k_vs_all()` (in module `dicee.evaluation.link_prediction`), 50
`evaluate_lp_k_vs_all()` (`dicee.evaluation.Evaluator` method), 56
`evaluate_lp_k_vs_all()` (`dicee.evaluation.evaluator.Evaluator` method), 46
`evaluate_lp_k_vs_all()` (`dicee.Evaluator` method), 219
`evaluate_lp_k_vs_all()` (`dicee.evaluator.Evaluator` method), 64
`evaluate_lp_with_byte()` (`dicee.evaluation.Evaluator` method), 56
`evaluate_lp_with_byte()` (`dicee.evaluation.evaluator.Evaluator` method), 47
`evaluate_lp_with_byte()` (`dicee.Evaluator` method), 220
`evaluate_lp_with_byte()` (`dicee.evaluator.Evaluator` method), 64
`Evaluator` (class in `dicee`), 218
`Evaluator` (class in `dicee.evaluation`), 54
`Evaluator` (class in `dicee.evaluation.evaluator`), 45
`Evaluator` (class in `dicee.evaluator`), 62
`evaluator` (`dicee.DICE_Trainer` attribute), 217
`evaluator` (`dicee.Execute` attribute), 210, 211
`evaluator` (`dicee.executer.Execute` attribute), 66
`evaluator` (`dicee.trainer.DICE_Trainer` attribute), 203
`evaluator` (`dicee.trainer.dice_trainer.DICE_Trainer` attribute), 198
`every_x_epoch` (`dicee.callbacks.KGESaveCallback` attribute), 23
`example_input_array` (`dicee.models.ensemble.EnsembleKGE` property), 105
`Execute` (class in `dicee`), 210
`Execute` (class in `dicee.executer`), 66
`exists()` (`dicee.knowledge_graph.KG` method), 70
`Experiment` (class in `dicee.analyse_experiments`), 19
`experiment_dir` (`dicee.callbacks.LRScheduler` attribute), 27
`experiment_dir` (`dicee.callbacks.PeriodicEvalCallback` attribute), 26
`explicit` (`dicee.models.QMult` attribute), 155
`explicit` (`dicee.models.quaternion.QMult` attribute), 116
`exponential_function()` (in module `dicee.static_funcs`), 194
`extract_input_outputs()` (`dicee.trainer.torch_trainer_ddp.NodeTrainer` method), 202
`extract_input_outputs()` (in module `dicee.trainer.model_parallelism`), 200
`extract_input_outputs_set_device()` (`dicee.trainer.torch_trainer.TorchTrainer` method), 201

F

`f` (`dicee.callbacks.KronE` attribute), 25
`fc` (`dicee.models.literal.LiteralEmbeddings` attribute), 110
`fc1` (`dicee.models.AConEx` attribute), 150
`fc1` (`dicee.models.AConvO` attribute), 163
`fc1` (`dicee.models.AConvQ` attribute), 157
`fc1` (`dicee.models.complex.AConEx` attribute), 102
`fc1` (`dicee.models.complex.ConEx` attribute), 101
`fc1` (`dicee.models.ConEx` attribute), 149
`fc1` (`dicee.models.ConvO` attribute), 162
`fc1` (`dicee.models.ConvQ` attribute), 156
`fc1` (`dicee.models.octonion.AConvO` attribute), 113
`fc1` (`dicee.models.octonion.ConvO` attribute), 112
`fc1` (`dicee.models.quaternion.AConvQ` attribute), 117
`fc1` (`dicee.models.quaternion.ConvQ` attribute), 117
`fc_num_input` (`dicee.models.AConEx` attribute), 150
`fc_num_input` (`dicee.models.AConvO` attribute), 163
`fc_num_input` (`dicee.models.AConvQ` attribute), 157
`fc_num_input` (`dicee.models.complex.AConEx` attribute), 102

`fc_num_input` (*dicее.models.complex.ConEx attribute*), 101
`fc_num_input` (*dicее.models.ConEx attribute*), 149
`fc_num_input` (*dicее.models.ConvO attribute*), 162
`fc_num_input` (*dicее.models.ConvQ attribute*), 156
`fc_num_input` (*dicее.models.octonion.AConvO attribute*), 113
`fc_num_input` (*dicее.models.octonion.ConvO attribute*), 112
`fc_num_input` (*dicее.models.quaternion.AConvQ attribute*), 117
`fc_num_input` (*dicее.models.quaternion.ConvQ attribute*), 117
`fc_out` (*dicее.models.literal.LiteralEmbeddings attribute*), 110
`feature_map_dropout` (*dicее.models.AConEx attribute*), 150
`feature_map_dropout` (*dicее.models.AConvO attribute*), 163
`feature_map_dropout` (*dicее.models.AConvQ attribute*), 157
`feature_map_dropout` (*dicее.models.complex.AConEx attribute*), 102
`feature_map_dropout` (*dicее.models.complex.ConEx attribute*), 101
`feature_map_dropout` (*dicее.models.ConEx attribute*), 150
`feature_map_dropout` (*dicее.models.ConvO attribute*), 162
`feature_map_dropout` (*dicее.models.ConvQ attribute*), 157
`feature_map_dropout` (*dicее.models.octonion.AConvO attribute*), 113
`feature_map_dropout` (*dicее.models.octonion.ConvO attribute*), 113
`feature_map_dropout` (*dicее.models.quaternion.AConvQ attribute*), 117
`feature_map_dropout` (*dicее.models.quaternion.ConvQ attribute*), 117
`feature_map_dropout_rate` (*dicее.config.Namespace attribute*), 30
`feature_map_dropout_rate` (*dicее.models.base_model.BaseKGE attribute*), 90
`feature_map_dropout_rate` (*dicее.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 172, 175
`fetch_worker()` (in module *dicее.read_preprocess_save_load_kg.util*), 186
`fill_query()` (*dicее.query_generator.QueryGenerator method*), 181
`fill_query()` (*dicее.QueryGenerator method*), 216
`find_good_batch_size()` (in module *dicее.trainer.model_parallelism*), 200
`find_missing_triples()` (*dicее.KGE method*), 214
`find_missing_triples()` (*dicее.knowledge_graph_embeddings.KGE method*), 74
`fit()` (*dicее.trainer.model_parallelism.TensorParallel method*), 200
`fit()` (*dicее.trainer.torch_trainer.TorchDDPTrainer method*), 202
`fit()` (*dicее.trainer.torch_trainer.TorchTrainer method*), 201
`FixedNegSampleDataset` (class in *dicее.dataset_classes*), 36
`flash` (*dicее.models.transformers.SelfAttention attribute*), 124
`FMult` (class in *dicее.models*), 176
`FMult` (class in *dicее.models.function_space*), 106
`FMult2` (class in *dicее.models*), 177
`FMult2` (class in *dicее.models.function_space*), 107
`form_of_labelling` (*dicее.DICE_Trainer attribute*), 217
`form_of_labelling` (*dicее.trainer.DICE_Trainer attribute*), 203
`form_of_labelling` (*dicее.trainer.dice_trainer.DICE_Trainer attribute*), 198
`forward()` (*dicее.models.base_model.BaseKGE method*), 91
`forward()` (*dicее.models.base_model.IdentityClass static method*), 92
`forward()` (*dicее.models.BaseKGE method*), 140, 143, 149, 153, 159, 172, 176
`forward()` (*dicее.models.Block method*), 144
`forward()` (*dicее.models.clifford.TransformerBlock method*), 96
`forward()` (*dicее.models.clifford.TransformerMLP method*), 97
`forward()` (*dicее.models.clifford.TransformerSelfAttention method*), 97
`forward()` (*dicее.models.IdentityClass static method*), 141, 155, 161
`forward()` (*dicее.models.literal.LiteralEmbeddings method*), 110
`forward()` (*dicее.models.transformers.Block method*), 126
`forward()` (*dicее.models.transformers.BytE method*), 122
`forward()` (*dicее.models.transformers.GPT method*), 127
`forward()` (*dicее.models.transformers.LayerNorm method*), 123
`forward()` (*dicее.models.transformers.MLP method*), 125
`forward()` (*dicее.models.transformers.SelfAttention method*), 124
`forward_backward_update()` (*dicее.trainer.torch_trainer.TorchTrainer method*), 201
`forward_backward_update_loss()` (in module *dicее.trainer.model_parallelism*), 200
`forward_byte_pair_encoded_k_vs_all()` (*dicее.models.base_model.BaseKGE method*), 91
`forward_byte_pair_encoded_k_vs_all()` (*dicее.models.BaseKGE method*), 139, 143, 148, 153, 159, 172, 175
`forward_byte_pair_encoded_triple()` (*dicее.models.base_model.BaseKGE method*), 91
`forward_byte_pair_encoded_triple()` (*dicее.models.BaseKGE method*), 139, 143, 148, 153, 159, 172, 175
`forward_k_vs_all()` (*dicее.models.AConEx method*), 150
`forward_k_vs_all()` (*dicее.models.AConvO method*), 163
`forward_k_vs_all()` (*dicее.models.AConvQ method*), 157
`forward_k_vs_all()` (*dicее.models.base_model.BaseKGE method*), 91
`forward_k_vs_all()` (*dicее.models.BaseKGE method*), 140, 143, 149, 153, 159, 173, 176

`forward_k_vs_all()` (*dicee.models.clifford.DeCaL method*), 99
`forward_k_vs_all()` (*dicee.models.clifford.Keci method*), 95
`forward_k_vs_all()` (*dicee.models.clifford.KeciTransformer method*), 96
`forward_k_vs_all()` (*dicee.models.CoKE method*), 147
`forward_k_vs_all()` (*dicee.models.ComplEx method*), 151
`forward_k_vs_all()` (*dicee.models.complex.AConEx method*), 102
`forward_k_vs_all()` (*dicee.models.complex.ComplEx method*), 103
`forward_k_vs_all()` (*dicee.models.complex.ConEx method*), 102
`forward_k_vs_all()` (*dicee.models.ConEx method*), 150
`forward_k_vs_all()` (*dicee.models.ConvO method*), 163
`forward_k_vs_all()` (*dicee.models.ConvQ method*), 157
`forward_k_vs_all()` (*dicee.models.DeCaL method*), 168
`forward_k_vs_all()` (*dicee.models.DistMult method*), 145
`forward_k_vs_all()` (*dicee.models.DualE method*), 180
`forward_k_vs_all()` (*dicee.models.dualE.DualE method*), 104
`forward_k_vs_all()` (*dicee.models.Keci method*), 165
`forward_k_vs_all()` (*dicee.models.KeciTransformer method*), 170
`forward_k_vs_all()` (*dicee.models.octonion.AConvO method*), 113
`forward_k_vs_all()` (*dicee.models.octonion.ConvO method*), 113
`forward_k_vs_all()` (*dicee.models.octonion.OMult method*), 112
`forward_k_vs_all()` (*dicee.models.OMult method*), 161
`forward_k_vs_all()` (*dicee.models.pykeen_models.PykeenKGE method*), 114
`forward_k_vs_all()` (*dicee.models.PykeenKGE method*), 173
`forward_k_vs_all()` (*dicee.models.QMult method*), 156
`forward_k_vs_all()` (*dicee.models.quaternion.AConvQ method*), 117
`forward_k_vs_all()` (*dicee.models.quaternion.ConvQ method*), 117
`forward_k_vs_all()` (*dicee.models.quaternion.QMult method*), 116
`forward_k_vs_all()` (*dicee.models.real.CoKE method*), 120
`forward_k_vs_all()` (*dicee.models.real.DistMult method*), 118
`forward_k_vs_all()` (*dicee.models.real.Shallom method*), 119
`forward_k_vs_all()` (*dicee.models.real.TransE method*), 118
`forward_k_vs_all()` (*dicee.models.Shallom method*), 145
`forward_k_vs_all()` (*dicee.models.TransE method*), 145
`forward_k_vs_sample()` (*dicee.models.AConEx method*), 150
`forward_k_vs_sample()` (*dicee.models.base_model.BaseKGE method*), 91
`forward_k_vs_sample()` (*dicee.models.BaseKGE method*), 140, 143, 149, 153, 159, 173, 176
`forward_k_vs_sample()` (*dicee.models.clifford.Keci method*), 95
`forward_k_vs_sample()` (*dicee.models.CoKE method*), 147
`forward_k_vs_sample()` (*dicee.models.ComplEx method*), 151
`forward_k_vs_sample()` (*dicee.models.complex.AConEx method*), 102
`forward_k_vs_sample()` (*dicee.models.complex.ComplEx method*), 103
`forward_k_vs_sample()` (*dicee.models.complex.ConEx method*), 102
`forward_k_vs_sample()` (*dicee.models.ConEx method*), 150
`forward_k_vs_sample()` (*dicee.models.DistMult method*), 145
`forward_k_vs_sample()` (*dicee.models.Keci method*), 166
`forward_k_vs_sample()` (*dicee.models.pykeen_models.PykeenKGE method*), 114
`forward_k_vs_sample()` (*dicee.models.PykeenKGE method*), 174
`forward_k_vs_sample()` (*dicee.models.QMult method*), 156
`forward_k_vs_sample()` (*dicee.models.quaternion.QMult method*), 116
`forward_k_vs_sample()` (*dicee.models.real.CoKE method*), 120
`forward_k_vs_sample()` (*dicee.models.real.DistMult method*), 118
`forward_k_vs_with_explicit()` (*dicee.models.clifford.Keci method*), 95
`forward_k_vs_with_explicit()` (*dicee.models.Keci method*), 165
`forward_triples()` (*dicee.models.AConEx method*), 150
`forward_triples()` (*dicee.models.AConvO method*), 163
`forward_triples()` (*dicee.models.AConvQ method*), 157
`forward_triples()` (*dicee.models.base_model.BaseKGE method*), 91
`forward_triples()` (*dicee.models.BaseKGE method*), 140, 143, 149, 153, 159, 173, 176
`forward_triples()` (*dicee.models.clifford.DeCaL method*), 98
`forward_triples()` (*dicee.models.clifford.Keci method*), 96
`forward_triples()` (*dicee.models.complex.AConEx method*), 102
`forward_triples()` (*dicee.models.complex.ConEx method*), 102
`forward_triples()` (*dicee.models.ConEx method*), 150
`forward_triples()` (*dicee.models.ConvO method*), 162
`forward_triples()` (*dicee.models.ConvQ method*), 157
`forward_triples()` (*dicee.models.DeCaL method*), 167
`forward_triples()` (*dicee.models.DualE method*), 180
`forward_triples()` (*dicee.models.dualE.DualE method*), 104

`forward_triples()` (*dicee.models.FMult method*), 177
`forward_triples()` (*dicee.models.FMult2 method*), 178
`forward_triples()` (*dicee.models.function_space.FMult method*), 106
`forward_triples()` (*dicee.models.function_space.FMult2 method*), 107
`forward_triples()` (*dicee.models.function_space.GFMult method*), 107
`forward_triples()` (*dicee.models.function_space.LFMult method*), 108
`forward_triples()` (*dicee.models.function_space.LFMult1 method*), 107
`forward_triples()` (*dicee.models.GFMult method*), 177
`forward_triples()` (*dicee.models.Keci method*), 166
`forward_triples()` (*dicee.models.LFMult method*), 178
`forward_triples()` (*dicee.models.LFMult1 method*), 178
`forward_triples()` (*dicee.models.octonion.AConvO method*), 113
`forward_triples()` (*dicee.models.octonion.ConvO method*), 113
`forward_triples()` (*dicee.models.Pyke method*), 145
`forward_triples()` (*dicee.models.pykeen_models.PykeenKGE method*), 114
`forward_triples()` (*dicee.models.PykeenKGE method*), 174
`forward_triples()` (*dicee.models.quaternion.AConvQ method*), 117
`forward_triples()` (*dicee.models.quaternion.ConvQ method*), 117
`forward_triples()` (*dicee.models.real.Pyke method*), 119
`forward_triples()` (*dicee.models.real.Shallom method*), 119
`forward_triples()` (*dicee.models.Shallom method*), 145
`freeze_entity_embeddings` (*dicee.models.literal.LiteralEmbeddings attribute*), 110
`frequency` (*dicee.callbacks.Perturb attribute*), 26
`from_pretrained()` (*dicee.models.transformers.GPT class method*), 127
`from_pretrained_model_write_embeddings_into_csv()` (*in module dicee.static_funcs*), 194
`full_storage_path` (*dicee.analyse_experiments.Experiment attribute*), 20
`func_triple_to_bpe_representation` (*dicee.evaluation.Evaluator attribute*), 55
`func_triple_to_bpe_representation` (*dicee.evaluation.evaluator.Evaluator attribute*), 45
`func_triple_to_bpe_representation` (*dicee.Evaluator attribute*), 218
`func_triple_to_bpe_representation` (*dicee.evaluator.Evaluator attribute*), 63
`func_triple_to_bpe_representation()` (*dicee.knowledge_graph.KG method*), 71
`function()` (*dicee.models.FMult2 method*), 178
`function()` (*dicee.models.function_space.FMult2 method*), 107

G

`gamma` (*dicee.models.FMult attribute*), 177
`gamma` (*dicee.models.function_space.FMult attribute*), 106
`gate_residual` (*dicee.models.literal.LiteralEmbeddings attribute*), 109, 110
`gated_residual_proj` (*dicee.models.literal.LiteralEmbeddings attribute*), 110
`gelu` (*dicee.models.clifford.TransformerMLP attribute*), 97
`gelu` (*dicee.models.transformers.MLP attribute*), 125
`gen_test` (*dicee.query_generator.QueryGenerator attribute*), 181
`gen_test` (*dicee.QueryGenerator attribute*), 215
`gen_valid` (*dicee.query_generator.QueryGenerator attribute*), 181
`gen_valid` (*dicee.QueryGenerator attribute*), 215
`generate()` (*dicee.KGE method*), 212
`generate()` (*dicee.knowledge_graph_embeddings.KGE method*), 71
`generate()` (*dicee.models.transformers.BytE method*), 122
`generate_queries()` (*dicee.query_generator.QueryGenerator method*), 181
`generate_queries()` (*dicee.QueryGenerator method*), 216
`get_aswa_state_dict()` (*dicee.weight_averaging.ASWA method*), 205
`get_bpe_head_and_relation_representation()` (*dicee.models.base_model.BaseKGE method*), 91
`get_bpe_head_and_relation_representation()` (*dicee.models.BaseKGE method*), 140, 143, 149, 154, 160, 173, 176
`get_bpe_token_representation()` (*dicee.abstracts.BaseInteractiveKGE method*), 15
`get_callbacks()` (*in module dicee.trainer.dice_trainer*), 198
`get_default_arguments()` (*in module dicee.analyse_experiments*), 19
`get_default_arguments()` (*in module dicee.scripts.index_serve*), 189
`get_default_arguments()` (*in module dicee.scripts.run*), 190
`get_ee_vocab()` (*in module dicee.read_preprocess_save_load_kg.util*), 186
`get_ee_vocab()` (*in module dicee.static_funcs*), 192
`get_ee_vocab()` (*in module dicee.static_preprocess_funcs*), 197
`get_embeddings()` (*dicee.models.base_model.BaseKGE method*), 92
`get_embeddings()` (*dicee.models.BaseKGE method*), 140, 143, 149, 154, 160, 173, 176
`get_embeddings()` (*dicee.models.ensemble.EnsembleKGE method*), 105
`get_embeddings()` (*dicee.models.real.Shallom method*), 119
`get_embeddings()` (*dicee.models.Shallom method*), 145
`get_entity_embeddings()` (*dicee.abstracts.BaseInteractiveKGE method*), 16

[get_entity_index\(\)](#) (*dicee.abstracts.BaseInteractiveKGE method*), 15
[get_er_vocab\(\)](#) (in module *dicee.read_preprocess_save_load_kg.util*), 186
[get_er_vocab\(\)](#) (in module *dicee.static_funcs*), 191
[get_er_vocab\(\)](#) (in module *dicee.static_preprocess_funcs*), 197
[get_eval_report\(\)](#) (*dicee.abstracts.BaseInteractiveKGE method*), 15
[get_head_relation_representation\(\)](#) (*dicee.models.base_model.BaseKGE method*), 91
[get_head_relation_representation\(\)](#) (*dicee.models.BaseKGE method*), 140, 143, 149, 154, 159, 173, 176
[get_kronecker_triple_representation\(\)](#) (*dicee.callbacks.KronE method*), 25
[get_mean_and_var\(\)](#) (*dicee.weight_averaging.SWAG method*), 207
[get_num_params\(\)](#) (*dicee.models.transformers.GPT method*), 127
[get_padded_bpe_triple_representation\(\)](#) (*dicee.abstracts.BaseInteractiveKGE method*), 15
[get_queries\(\)](#) (*dicee.query_generator.QueryGenerator method*), 182
[get_queries\(\)](#) (*dicee.QueryGenerator method*), 216
[get_re_vocab\(\)](#) (in module *dicee.read_preprocess_save_load_kg.util*), 186
[get_re_vocab\(\)](#) (in module *dicee.static_funcs*), 192
[get_re_vocab\(\)](#) (in module *dicee.static_preprocess_funcs*), 197
[get_relation_embeddings\(\)](#) (*dicee.abstracts.BaseInteractiveKGE method*), 16
[get_relation_index\(\)](#) (*dicee.abstracts.BaseInteractiveKGE method*), 15
[get_sentence_representation\(\)](#) (*dicee.models.base_model.BaseKGE method*), 91
[get_sentence_representation\(\)](#) (*dicee.models.BaseKGE method*), 140, 143, 149, 154, 160, 173, 176
[get_transductive_entity_embeddings\(\)](#) (*dicee.KGE method*), 212
[get_transductive_entity_embeddings\(\)](#) (*dicee.knowledge_graph_embeddings.KGE method*), 71
[get_triple_representation\(\)](#) (*dicee.models.base_model.BaseKGE method*), 91
[get_triple_representation\(\)](#) (*dicee.models.BaseKGE method*), 140, 143, 149, 153, 159, 173, 176
[GFMult](#) (class in *dicee.models*), 177
[GFMult](#) (class in *dicee.models.function_space*), 106
[global_rank](#) (*dicee.abstracts.AbstractTrainer attribute*), 13
[global_rank](#) (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
[GPT](#) (class in *dicee.models.transformers*), 126
[GPTConfig](#) (class in *dicee.models.transformers*), 126
[gpus](#) (*dicee.config.Namespace attribute*), 28
[gradient_accumulation_steps](#) (*dicee.config.Namespace attribute*), 29
[ground_queries\(\)](#) (*dicee.query_generator.QueryGenerator method*), 181
[ground_queries\(\)](#) (*dicee.QueryGenerator method*), 216
[gswa_n](#) (*dicee.weight_averaging.SWAG attribute*), 207

H

[hidden_dim](#) (*dicee.models.literal.LiteralEmbeddings attribute*), 110
[hidden_dropout](#) (*dicee.models.base_model.BaseKGE attribute*), 91
[hidden_dropout](#) (*dicee.models.BaseKGE attribute*), 139, 142, 148, 153, 159, 172, 175
[hidden_dropout_rate](#) (*dicee.config.Namespace attribute*), 30
[hidden_dropout_rate](#) (*dicee.models.base_model.BaseKGE attribute*), 90
[hidden_dropout_rate](#) (*dicee.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 171, 175
[hidden_normalizer](#) (*dicee.models.base_model.BaseKGE attribute*), 90
[hidden_normalizer](#) (*dicee.models.BaseKGE attribute*), 139, 142, 148, 153, 159, 172, 175

I

[IdentityClass](#) (class in *dicee.models*), 140, 154, 160
[IdentityClass](#) (class in *dicee.models.base_model*), 92
[idx_entity_to_bpe_shaped](#) (*dicee.knowledge_graph.KG attribute*), 70
[index\(\)](#) (in module *dicee.scripts.index_serve*), 189
[index_triple\(\)](#) (*dicee.abstracts.BaseInteractiveKGE method*), 15
[init_data_loader\(\)](#) (*dicee.DICE_Trainer method*), 217
[init_data_loader\(\)](#) (*dicee.trainer.DICE_Trainer method*), 203
[init_data_loader\(\)](#) (*dicee.trainer.dice_trainer.DICE_Trainer method*), 199
[init_dataset\(\)](#) (*dicee.DICE_Trainer method*), 217
[init_dataset\(\)](#) (*dicee.trainer.DICE_Trainer method*), 204
[init_dataset\(\)](#) (*dicee.trainer.dice_trainer.DICE_Trainer method*), 199
[init_param](#) (*dicee.config.Namespace attribute*), 29
[init_params_with_sanity_checking\(\)](#) (*dicee.models.base_model.BaseKGE method*), 91
[init_params_with_sanity_checking\(\)](#) (*dicee.models.BaseKGE method*), 140, 143, 149, 153, 159, 172, 176
[initial_eval_setting](#) (*dicee.weight_averaging.ASWA attribute*), 204
[initialize_or_load_model\(\)](#) (*dicee.DICE_Trainer method*), 217
[initialize_or_load_model\(\)](#) (*dicee.trainer.DICE_Trainer method*), 203
[initialize_or_load_model\(\)](#) (*dicee.trainer.dice_trainer.DICE_Trainer method*), 199
[initialize_trainer\(\)](#) (*dicee.DICE_Trainer method*), 217
[initialize_trainer\(\)](#) (*dicee.trainer.DICE_Trainer method*), 203

`initialize_trainer()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 199
`initialize_trainer()` (*in module dicee.trainer.dice_trainer*), 197
`input_dim` (*dicee.models.clifford.KeciTransformer attribute*), 96
`input_dim` (*dicee.models.KeciTransformer attribute*), 170
`input_dp_ent_real` (*dicee.models.base_model.BaseKGE attribute*), 91
`input_dp_ent_real` (*dicee.models.BaseKGE attribute*), 139, 142, 148, 153, 159, 172, 175
`input_dp_rel_real` (*dicee.models.base_model.BaseKGE attribute*), 91
`input_dp_rel_real` (*dicee.models.BaseKGE attribute*), 139, 142, 148, 153, 159, 172, 175
`input_dropout_rate` (*dicee.config.Namespace attribute*), 30
`input_dropout_rate` (*dicee.models.base_model.BaseKGE attribute*), 90
`input_dropout_rate` (*dicee.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 171, 175
`InteractiveQueryDecomposition` (*class in dicee.abstracts*), 16
`intialize_model()` (*in module dicee.static_funcs*), 193
`is_continual_training` (*dicee.DICE_Trainer attribute*), 217
`is_continual_training` (*dicee.evaluation.Evaluator attribute*), 55
`is_continual_training` (*dicee.evaluation.evaluator.Evaluator attribute*), 45
`is_continual_training` (*dicee.Evaluator attribute*), 218
`is_continual_training` (*dicee.evaluator.Evaluator attribute*), 63
`is_continual_training` (*dicee.Execute attribute*), 210
`is_continual_training` (*dicee.executer.Execute attribute*), 66
`is_continual_training` (*dicee.trainer.DICE_Trainer attribute*), 203
`is_continual_training` (*dicee.trainer.dice_trainer.DICE_Trainer attribute*), 198
`is_global_zero` (*dicee.abstracts.AbstractTrainer attribute*), 13
`is_rank_zero()` (*dicee.Execute method*), 211
`is_rank_zero()` (*dicee.executer.Execute method*), 67
`is_seen()` (*dicee.abstracts.BaseInteractiveKGE method*), 15
`is_sparql_endpoint_alive()` (*in module dicee.sanity_checkers*), 188

K

`k` (*dicee.models.FMult attribute*), 176
`k` (*dicee.models.FMult2 attribute*), 177
`k` (*dicee.models.function_space.FMult attribute*), 106
`k` (*dicee.models.function_space.FMult2 attribute*), 107
`k` (*dicee.models.function_space.GFMult attribute*), 106
`k` (*dicee.models.GFMult attribute*), 177
`k_fold_cross_validation()` (*dicee.DICE_Trainer method*), 217
`k_fold_cross_validation()` (*dicee.trainer.DICE_Trainer method*), 204
`k_fold_cross_validation()` (*dicee.trainer.dice_trainer.DICE_Trainer method*), 199
`k_vs_all_score()` (*dicee.models.clifford.Keci method*), 95
`k_vs_all_score()` (*dicee.models.ComplEx static method*), 151
`k_vs_all_score()` (*dicee.models.complex.ComplEx static method*), 103
`k_vs_all_score()` (*dicee.models.DistMult method*), 144
`k_vs_all_score()` (*dicee.models.Keci method*), 165
`k_vs_all_score()` (*dicee.models.octonion.OMult method*), 112
`k_vs_all_score()` (*dicee.models.OMult method*), 161
`k_vs_all_score()` (*dicee.models.QMult method*), 156
`k_vs_all_score()` (*dicee.models.quaternion.QMult method*), 116
`k_vs_all_score()` (*dicee.models.real.DistMult method*), 118
`Keci` (*class in dicee.models*), 163
`Keci` (*class in dicee.models.clifford*), 93
`KeciTransformer` (*class in dicee.models*), 170
`KeciTransformer` (*class in dicee.models.clifford*), 96
`kernel_size` (*dicee.config.Namespace attribute*), 30
`kernel_size` (*dicee.models.base_model.BaseKGE attribute*), 90
`kernel_size` (*dicee.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 172, 175
`KG` (*class in dicee.knowledge_graph*), 68
`kg` (*dicee.callbacks.PseudoLabellingCallback attribute*), 24
`kg` (*dicee.read_preprocess_save_load_kg.LoadSaveToDisk attribute*), 188
`kg` (*dicee.read_preprocess_save_load_kg.PreprocessKG attribute*), 187
`kg` (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG attribute*), 182
`kg` (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk attribute*), 183
`kg` (*dicee.read_preprocess_save_load_kg.ReadFromDisk attribute*), 188
`kg` (*dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk attribute*), 183
`KGE` (*class in dicee*), 212
`KGE` (*class in dicee.knowledge_graph_embeddings*), 71
`KGESaveCallback` (*class in dicee.callbacks*), 22
`knowledge_graph` (*dicee.Execute attribute*), 210

knowledge_graph (*dicee.executer.Execute* attribute), 66
 KronE (*class in dicee.callbacks*), 25
 KvsAll (*class in dicee.dataset_classes*), 34
 kvsall_score () (*dicee.models.DualE* method), 180
 kvsall_score () (*dicee.models.dualE.DualE* method), 104
 KvsSampleDataset (*class in dicee.dataset_classes*), 35

L

label_smoothing_rate (*dicee.config.Namespace* attribute), 29
 label_smoothing_rate (*dicee.dataset_classes.AllvsAll* attribute), 34
 label_smoothing_rate (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 36
 label_smoothing_rate (*dicee.dataset_classes.KvsAll* attribute), 34
 label_smoothing_rate (*dicee.dataset_classes.KvsSampleDataset* attribute), 35
 label_smoothing_rate (*dicee.dataset_classes.OnevsSample* attribute), 37
 label_smoothing_rate (*dicee.dataset_classes.TriplePredictionDataset* attribute), 37
 layer_norm (*dicee.models.literal.LiteralEmbeddings* attribute), 110
 LayerNorm (*class in dicee.models.transformers*), 123
 learning_rate (*dicee.models.base_model.BaseKGE* attribute), 90
 learning_rate (*dicee.models.BaseKGE* attribute), 139, 142, 148, 152, 158, 171, 175
 length (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 36
 length (*dicee.dataset_classes.TriplePredictionDataset* attribute), 37
 level (*dicee.callbacks.Perturb* attribute), 26
 LFMult (*class in dicee.models*), 178
 LFMult (*class in dicee.models.function_space*), 108
 LFMult1 (*class in dicee.models*), 178
 LFMult1 (*class in dicee.models.function_space*), 107
 linear () (*dicee.models.function_space.LFMult* method), 108
 linear () (*dicee.models.LFMult* method), 179
 list2tuple () (*dicee.query_generator.QueryGenerator* method), 181
 list2tuple () (*dicee.QueryGenerator* method), 216
 LiteralDataset (*class in dicee.dataset_classes*), 38
 LiteralEmbeddings (*class in dicee.models.literal*), 109
 lm_head (*dicee.models.clifford.KeciTransformer* attribute), 96
 lm_head (*dicee.models.KeciTransformer* attribute), 170
 lm_head (*dicee.models.transformers.BytE* attribute), 122
 lm_head (*dicee.models.transformers.GPT* attribute), 127
 ln_1 (*dicee.models.Block* attribute), 144
 ln_1 (*dicee.models.clifford.TransformerBlock* attribute), 96
 ln_1 (*dicee.models.transformers.Block* attribute), 126
 ln_2 (*dicee.models.Block* attribute), 144
 ln_2 (*dicee.models.clifford.TransformerBlock* attribute), 96
 ln_2 (*dicee.models.transformers.Block* attribute), 126
 ln_f (*dicee.models.CoKE* attribute), 147
 ln_f (*dicee.models.real.CoKE* attribute), 120
 load () (*dicee.read_preprocess_save_load_kg.LoadSaveToDisk* method), 188
 load () (*dicee.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk* method), 183
 load_and_validate_literal_data () (*dicee.dataset_classes.LiteralDataset* static method), 38
 load_from_memmap () (*dicee.Execute* method), 211
 load_from_memmap () (*dicee.executer.Execute* method), 67
 load_json () (*in module dicee.static_funcs*), 194
 load_model () (*in module dicee.static_funcs*), 193
 load_model_ensemble () (*in module dicee.static_funcs*), 193
 load_numpy () (*in module dicee.static_funcs*), 194
 load_numpy_ndarray () (*in module dicee.read_preprocess_save_load_kg.util*), 187
 load_pickle () (*in module dicee.read_preprocess_save_load_kg.util*), 187
 load_pickle () (*in module dicee.static_funcs*), 192
 load_queries () (*dicee.query_generator.QueryGenerator* method), 182
 load_queries () (*dicee.QueryGenerator* method), 216
 load_queries_and_answers () (*dicee.query_generator.QueryGenerator* static method), 182
 load_queries_and_answers () (*dicee.QueryGenerator* static method), 216
 load_state_dict () (*dicee.models.ensemble.EnsembleKGE* method), 105
 load_term_mapping () (*in module dicee.static_funcs*), 192
 load_term_mapping () (*in module dicee.trainer.dice_trainer*), 197
 load_with_pandas () (*in module dicee.read_preprocess_save_load_kg.util*), 186
 loader_backend (*dicee.dataset_classes.LiteralDataset* attribute), 38
 LoadSaveToDisk (*class in dicee.read_preprocess_save_load_kg*), 188
 LoadSaveToDisk (*class in dicee.read_preprocess_save_load_kg.save_load_disk*), 183

local_rank (*dicее.abstracts.AbstractTrainer attribute*), 13
 local_rank (*dicее.Execute attribute*), 210
 local_rank (*dicее.executer.Execute attribute*), 66
 local_rank (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
 loss (*dicее.models.base_model.BaseKGE attribute*), 90
 loss (*dicее.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 172, 175
 loss_func (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
 loss_function (*dicее.trainer.torch_trainer.TorchTrainer attribute*), 200
 loss_function() (*dicее.models.base_model.BaseKGELightning method*), 85
 loss_function() (*dicее.models.BaseKGELightning method*), 134
 loss_function() (*dicее.models.transformers.BytE method*), 122
 loss_history (*dicее.models.base_model.BaseKGE attribute*), 91
 loss_history (*dicее.models.BaseKGE attribute*), 139, 143, 148, 153, 159, 172, 175
 loss_history (*dicее.models.pykeen_models.PykeenKGE attribute*), 114
 loss_history (*dicее.models.PykeenKGE attribute*), 173
 loss_history (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
 lr (*dicее.analyse_experiments.Experiment attribute*), 20
 lr (*dicее.config.Namespace attribute*), 28
 lr_init (*dicее.weight_averaging.SWA attribute*), 206
 lr_init (*dicее.weight_averaging.SWAG attribute*), 207
 lr_init (*dicее.weight_averaging.TWA attribute*), 209
 lr_lambda (*dicее.callbacks.LRScheduler attribute*), 27
 LRScheduler (*class in dicее.callbacks*), 26

M

m (*dicее.models.function_space.LFMult attribute*), 108
 m (*dicее.models.LFMult attribute*), 178
 main() (*in module dicее.scripts.index_serve*), 190
 main() (*in module dicее.scripts.run*), 190
 make_iterable_verbose() (*in module dicее.evaluation*), 61
 make_iterable_verbose() (*in module dicее.evaluation.utils*), 52
 make_iterable_verbose() (*in module dicее.static_funcs_training*), 196
 make_iterable_verbose() (*in module dicее.trainer.torch_trainer_ddp*), 201
 mapping_from_first_two_cols_to_third() (*in module dicее.static_preprocess_funcs*), 197
 margin (*dicее.models.Pyke attribute*), 145
 margin (*dicее.models.real.Pyke attribute*), 119
 margin (*dicее.models.real.TransE attribute*), 118
 margin (*dicее.models.TransE attribute*), 145
 mask_emb (*dicее.models.CoKE attribute*), 147
 mask_emb (*dicее.models.real.CoKE attribute*), 120
 max_ans_num (*dicее.query_generator.QueryGenerator attribute*), 181
 max_ans_num (*dicее.QueryGenerator attribute*), 215
 max_epochs (*dicее.callbacks.KGESaveCallback attribute*), 23
 max_epochs (*dicее.callbacks.PeriodicEvalCallback attribute*), 26
 max_epochs (*dicее.weight_averaging.EMA attribute*), 208
 max_epochs (*dicее.weight_averaging.SWA attribute*), 206
 max_epochs (*dicее.weight_averaging.SWAG attribute*), 207
 max_epochs (*dicее.weight_averaging.TWA attribute*), 209
 max_length_subword_tokens (*dicее.knowledge_graph.KG attribute*), 70
 max_length_subword_tokens (*dicее.models.base_model.BaseKGE attribute*), 91
 max_length_subword_tokens (*dicее.models.BaseKGE attribute*), 139, 143, 148, 153, 159, 172, 175
 max_num_models (*dicее.weight_averaging.SWAG attribute*), 207
 max_num_of_classes (*dicее.dataset_classes.KvsSampleDataset attribute*), 35
 mean (*dicее.weight_averaging.SWAG attribute*), 207
 mem_of_model() (*dicее.models.base_model.BaseKGELightning method*), 84
 mem_of_model() (*dicее.models.BaseKGELightning method*), 133
 mem_of_model() (*dicее.models.ensemble.EnsembleKGE method*), 105
 method (*dicее.callbacks.Perturb attribute*), 26
 MLP (*class in dicее.models.transformers*), 124
 mlp (*dicее.models.Block attribute*), 144
 mlp (*dicее.models.clifford.TransformerBlock attribute*), 96
 mlp (*dicее.models.transformers.Block attribute*), 126
 mode (*dicее.query_generator.QueryGenerator attribute*), 181
 mode (*dicее.QueryGenerator attribute*), 215
 model (*dicее.config.Namespace attribute*), 28
 model (*dicее.models.pykeen_models.PykeenKGE attribute*), 114
 model (*dicее.models.PykeenKGE attribute*), 173

- model (*dicee.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
- model (*dicee.trainer.torch_trainer.TorchTrainer attribute*), 200
- model_kwargs (*dicee.models.pykeen_models.PykeenKGE attribute*), 114
- model_kwargs (*dicee.models.PykeenKGE attribute*), 173
- model_name (*dicee.analyse_experiments.Experiment attribute*), 19
- MODEL_REGISTRY (*in module dicee.static_funcs*), 191
- module
 - dicee, 12
 - dicee.__main__, 12
 - dicee.abstracts, 12
 - dicee.analyse_experiments, 19
 - dicee.callbacks, 21
 - dicee.config, 27
 - dicee.dataset_classes, 31
 - dicee.eval_static_funcs, 39
 - dicee.evaluation, 43
 - dicee.evaluation.ensemble, 43
 - dicee.evaluation.evaluator, 44
 - dicee.evaluation.link_prediction, 48
 - dicee.evaluation.literal_prediction, 50
 - dicee.evaluation.utils, 51
 - dicee.evaluator, 62
 - dicee.executer, 65
 - dicee.knowledge_graph, 68
 - dicee.knowledge_graph_embeddings, 71
 - dicee.models, 74
 - dicee.models.adopt, 74
 - dicee.models.base_model, 83
 - dicee.models.clifford, 92
 - dicee.models.complex, 101
 - dicee.models.dualE, 103
 - dicee.models.ensemble, 105
 - dicee.models.function_space, 105
 - dicee.models.literal, 109
 - dicee.models.octonion, 110
 - dicee.models.pykeen_models, 113
 - dicee.models.quaternion, 115
 - dicee.models.real, 118
 - dicee.models.static_funcs, 120
 - dicee.models.transformers, 121
 - dicee.query_generator, 180
 - dicee.read_preprocess_save_load_kg, 182
 - dicee.read_preprocess_save_load_kg.preprocess, 182
 - dicee.read_preprocess_save_load_kg.read_from_disk, 183
 - dicee.read_preprocess_save_load_kg.save_load_disk, 183
 - dicee.read_preprocess_save_load_kg.util, 184
 - dicee.sanity_checkers, 188
 - dicee.scripts, 189
 - dicee.scripts.index_serve, 189
 - dicee.scripts.run, 190
 - dicee.static_funcs, 190
 - dicee.static_funcs_training, 195
 - dicee.static_preprocess_funcs, 196
 - dicee.trainer, 197
 - dicee.trainer.dice_trainer, 197
 - dicee.trainer.model_parallelism, 199
 - dicee.trainer.torch_trainer, 200
 - dicee.trainer.torch_trainer_ddp, 201
 - dicee.weight_averaging, 204
- modules () (*dicee.models.ensemble.EnsembleKGE method*), 105
- moving_average () (*dicee.weight_averaging.SWA static method*), 206
- MultiClassClassificationDataset (*class in dicee.dataset_classes*), 32
- MultiLabelDataset (*class in dicee.dataset_classes*), 33

N

- n (*dicee.models.FMult2 attribute*), 177
- n (*dicee.models.function_space.FMult2 attribute*), 107

n_embd (dicee.models.clifford.TransformerSelfAttention attribute), 97
n_embd (dicee.models.CoKEConfig attribute), 146
n_embd (dicee.models.real.CoKEConfig attribute), 119, 120
n_embd (dicee.models.transformers.GPTConfig attribute), 126
n_embd (dicee.models.transformers.SelfAttention attribute), 124
n_epochs_eval_model (dicee.callbacks.PeriodicEvalCallback attribute), 26
n_epochs_eval_model (dicee.config.Namespace attribute), 31
n_head (dicee.models.clifford.TransformerSelfAttention attribute), 97
n_head (dicee.models.CoKEConfig attribute), 146
n_head (dicee.models.real.CoKEConfig attribute), 119, 120
n_head (dicee.models.transformers.GPTConfig attribute), 126
n_head (dicee.models.transformers.SelfAttention attribute), 124
n_layer (dicee.models.CoKEConfig attribute), 146
n_layer (dicee.models.real.CoKEConfig attribute), 119, 120
n_layer (dicee.models.transformers.GPTConfig attribute), 126
n_layers (dicee.models.FMult2 attribute), 177
n_layers (dicee.models.function_space.FMult2 attribute), 107
name (dicee.abstracts.BaseInteractiveKGE property), 15
name (dicee.models.AConEx attribute), 150
name (dicee.models.AConvO attribute), 163
name (dicee.models.AConvQ attribute), 157
name (dicee.models.CKeci attribute), 166
name (dicee.models.clifford.CKeci attribute), 97
name (dicee.models.clifford.DeCaL attribute), 98
name (dicee.models.clifford.Keci attribute), 93
name (dicee.models.clifford.KeciTransformer attribute), 96
name (dicee.models.CoKE attribute), 146
name (dicee.models.ComplEx attribute), 151
name (dicee.models.complex.AConEx attribute), 102
name (dicee.models.complex.ComplEx attribute), 103
name (dicee.models.complex.ConEx attribute), 101
name (dicee.models.ConEx attribute), 149
name (dicee.models.ConvO attribute), 162
name (dicee.models.ConvQ attribute), 156
name (dicee.models.DeCaL attribute), 167
name (dicee.models.DistMult attribute), 144
name (dicee.models.DualE attribute), 179
name (dicee.models.dualE.DualE attribute), 104
name (dicee.models.ensemble.EnsembleKGE attribute), 105
name (dicee.models.FMult attribute), 176
name (dicee.models.FMult2 attribute), 177
name (dicee.models.function_space.FMult attribute), 106
name (dicee.models.function_space.FMult2 attribute), 107
name (dicee.models.function_space.GFMult attribute), 106
name (dicee.models.function_space.LFMult attribute), 108
name (dicee.models.function_space.LFMult1 attribute), 107
name (dicee.models.GFMult attribute), 177
name (dicee.models.Keci attribute), 164
name (dicee.models.KeciTransformer attribute), 170
name (dicee.models.LFMult attribute), 178
name (dicee.models.LFMult1 attribute), 178
name (dicee.models.octonion.AConvO attribute), 113
name (dicee.models.octonion.ConvO attribute), 112
name (dicee.models.octonion.OMult attribute), 111
name (dicee.models.OMult attribute), 161
name (dicee.models.Pyke attribute), 145
name (dicee.models.pykeen_models.PykeenKGE attribute), 114
name (dicee.models.PykeenKGE attribute), 173
name (dicee.models.QMult attribute), 155
name (dicee.models.quaternion.AConvQ attribute), 117
name (dicee.models.quaternion.ConvQ attribute), 116
name (dicee.models.quaternion.QMult attribute), 116
name (dicee.models.real.CoKE attribute), 120
name (dicee.models.real.DistMult attribute), 118
name (dicee.models.real.Pyke attribute), 119
name (dicee.models.real.Shallom attribute), 119
name (dicee.models.real.TransE attribute), 118
name (dicee.models.Shallom attribute), 145

name (*dicdee.models.TransE* attribute), 145
 name (*dicdee.models.transformers.ByT* attribute), 122
 named_children() (*dicdee.models.ensemble.EnsembleKGE* method), 105
 Namespace (class in *dicdee.config*), 28
 neg_ratio (*dicdee.config.Namespace* attribute), 29
 neg_ratio (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 32
 neg_ratio (*dicdee.dataset_classes.KvsSampleDataset* attribute), 35
 neg_sample_ratio (*dicdee.dataset_classes.FixedNegSampleDataset* attribute), 36
 neg_sample_ratio (*dicdee.dataset_classes.OnevsSample* attribute), 37
 neg_sample_ratio (*dicdee.dataset_classes.TriplePredictionDataset* attribute), 37
 negnorm() (*dicdee.abstracts.InteractiveQueryDecomposition* method), 16
 neural_searcher (in module *dicdee.scripts.index_serve*), 189
 NeuralSearcher (class in *dicdee.scripts.index_serve*), 189
 NodeTrainer (class in *dicdee.trainer.torch_trainer_ddp*), 202
 norm_fc1 (*dicdee.models.AConEx* attribute), 150
 norm_fc1 (*dicdee.models.AConvO* attribute), 163
 norm_fc1 (*dicdee.models.complex.AConEx* attribute), 102
 norm_fc1 (*dicdee.models.complex.ConEx* attribute), 101
 norm_fc1 (*dicdee.models.ConEx* attribute), 149
 norm_fc1 (*dicdee.models.ConvO* attribute), 162
 norm_fc1 (*dicdee.models.octonion.AConvO* attribute), 113
 norm_fc1 (*dicdee.models.octonion.ConvO* attribute), 113
 normalization (*dicdee.analyse_experiments.Experiment* attribute), 20
 normalization (*dicdee.config.Namespace* attribute), 29
 normalization_params (*dicdee.dataset_classes.LiteralDataset* attribute), 38
 normalization_type (*dicdee.dataset_classes.LiteralDataset* attribute), 38
 normalize_head_entity_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 90
 normalize_head_entity_embeddings (*dicdee.models.BaseKGE* attribute), 139, 142, 148, 153, 159, 172, 175
 normalize_relation_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 90
 normalize_relation_embeddings (*dicdee.models.BaseKGE* attribute), 139, 142, 148, 153, 159, 172, 175
 normalize_tail_entity_embeddings (*dicdee.models.base_model.BaseKGE* attribute), 90
 normalize_tail_entity_embeddings (*dicdee.models.BaseKGE* attribute), 139, 142, 148, 153, 159, 172, 175
 normalizer_class (*dicdee.models.base_model.BaseKGE* attribute), 90
 normalizer_class (*dicdee.models.BaseKGE* attribute), 139, 142, 148, 153, 159, 172, 175
 num_bpe_entities (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 32
 num_bpe_entities (*dicdee.knowledge_graph.KG* attribute), 70
 num_core (*dicdee.config.Namespace* attribute), 29
 num_datapoints (*dicdee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 32
 num_datapoints (*dicdee.dataset_classes.MultiLabelDataset* attribute), 33
 num_ent (*dicdee.models.DualE* attribute), 180
 num_ent (*dicdee.models.dualE.DualE* attribute), 104
 num_entities (*dicdee.dataset_classes.FixedNegSampleDataset* attribute), 36
 num_entities (*dicdee.dataset_classes.KvsSampleDataset* attribute), 35
 num_entities (*dicdee.dataset_classes.LiteralDataset* attribute), 38
 num_entities (*dicdee.dataset_classes.OnevsSample* attribute), 37
 num_entities (*dicdee.dataset_classes.TriplePredictionDataset* attribute), 37
 num_entities (*dicdee.evaluation.Evaluator* attribute), 54, 55
 num_entities (*dicdee.evaluation.evaluator.Evaluator* attribute), 45
 num_entities (*dicdee.Evaluator* attribute), 218
 num_entities (*dicdee.evaluator.Evaluator* attribute), 62, 63
 num_entities (*dicdee.knowledge_graph.KG* attribute), 69
 num_entities (*dicdee.models.base_model.BaseKGE* attribute), 90
 num_entities (*dicdee.models.BaseKGE* attribute), 139, 142, 148, 152, 158, 171, 175
 num_epochs (*dicdee.abstracts.AbstractPPECallback* attribute), 18
 num_epochs (*dicdee.analyse_experiments.Experiment* attribute), 20
 num_epochs (*dicdee.config.Namespace* attribute), 28
 num_epochs (*dicdee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 202
 num_epochs (*dicdee.weight_averaging.ASWA* attribute), 204
 num_folds_for_cv (*dicdee.config.Namespace* attribute), 29
 num_of_data_points (*dicdee.dataset_classes.MultiClassClassificationDataset* attribute), 33
 num_of_data_properties (*dicdee.models.literal.LiteralEmbeddings* attribute), 109, 110
 num_of_epochs (*dicdee.callbacks.PseudoLabellingCallback* attribute), 24
 num_of_output_channels (*dicdee.config.Namespace* attribute), 30
 num_of_output_channels (*dicdee.models.base_model.BaseKGE* attribute), 90
 num_of_output_channels (*dicdee.models.BaseKGE* attribute), 139, 142, 148, 152, 158, 172, 175
 num_params (*dicdee.analyse_experiments.Experiment* attribute), 20
 num_relations (*dicdee.dataset_classes.FixedNegSampleDataset* attribute), 36
 num_relations (*dicdee.dataset_classes.OnevsSample* attribute), 37

num_relations (*dicee.dataset_classes.TriplePredictionDataset attribute*), 37
 num_relations (*dicee.evaluation.Evaluator attribute*), 54, 55
 num_relations (*dicee.evaluation.evaluator.Evaluator attribute*), 45
 num_relations (*dicee.Evaluator attribute*), 218, 219
 num_relations (*dicee.evaluator.Evaluator attribute*), 62, 63
 num_relations (*dicee.knowledge_graph.KG attribute*), 69
 num_relations (*dicee.models.base_model.BaseKGE attribute*), 90
 num_relations (*dicee.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 171, 175
 num_sample (*dicee.models.FMult attribute*), 176
 num_sample (*dicee.models.function_space.FMult attribute*), 106
 num_sample (*dicee.models.function_space.GFMult attribute*), 106
 num_sample (*dicee.models.GFMult attribute*), 177
 num_samples (*dicee.weight_averaging.TWA attribute*), 208
 num_tokens (*dicee.knowledge_graph.KG attribute*), 70
 num_tokens (*dicee.models.base_model.BaseKGE attribute*), 90
 num_tokens (*dicee.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 171, 175
 numpy_data_type_changer() (in module *dicee.static_funcs*), 193

O

octonion_mul() (in module *dicee.models*), 161
 octonion_mul() (in module *dicee.models.octonion*), 111
 octonion_mul_norm() (in module *dicee.models*), 161
 octonion_mul_norm() (in module *dicee.models.octonion*), 111
 octonion_normalizer() (*dicee.models.AConvO static method*), 163
 octonion_normalizer() (*dicee.models.ConvO static method*), 162
 octonion_normalizer() (*dicee.models.octonion.AConvO static method*), 113
 octonion_normalizer() (*dicee.models.octonion.ConvO static method*), 113
 octonion_normalizer() (*dicee.models.octonion.OMult static method*), 111
 octonion_normalizer() (*dicee.models.OMult static method*), 161
 OMult (class in *dicee.models*), 161
 OMult (class in *dicee.models.octonion*), 111
 on_epoch_end() (*dicee.callbacks.KGESaveCallback method*), 23
 on_epoch_end() (*dicee.callbacks.PseudoLabellingCallback method*), 24
 on_fit_end() (*dicee.abstracts.AbstractCallback method*), 17
 on_fit_end() (*dicee.abstracts.AbstractPPECallback method*), 18
 on_fit_end() (*dicee.abstracts.AbstractTrainer method*), 13
 on_fit_end() (*dicee.callbacks.AccumulateEpochLossCallback method*), 21
 on_fit_end() (*dicee.callbacks.Eval method*), 24
 on_fit_end() (*dicee.callbacks.KGESaveCallback method*), 23
 on_fit_end() (*dicee.callbacks.LRScheduler method*), 27
 on_fit_end() (*dicee.callbacks.PeriodicEvalCallback method*), 26
 on_fit_end() (*dicee.callbacks.PrintCallback method*), 22
 on_fit_end() (*dicee.weight_averaging.ASWA method*), 205
 on_fit_end() (*dicee.weight_averaging.EMA method*), 208
 on_fit_end() (*dicee.weight_averaging.SWA method*), 206
 on_fit_end() (*dicee.weight_averaging.SWAG method*), 207
 on_fit_end() (*dicee.weight_averaging.TWA method*), 209
 on_fit_start() (*dicee.abstracts.AbstractCallback method*), 17
 on_fit_start() (*dicee.abstracts.AbstractPPECallback method*), 18
 on_fit_start() (*dicee.abstracts.AbstractTrainer method*), 13
 on_fit_start() (*dicee.callbacks.Eval method*), 24
 on_fit_start() (*dicee.callbacks.KGESaveCallback method*), 23
 on_fit_start() (*dicee.callbacks.KronE method*), 25
 on_fit_start() (*dicee.callbacks.PrintCallback method*), 22
 on_init_end() (*dicee.abstracts.AbstractCallback method*), 17
 on_init_start() (*dicee.abstracts.AbstractCallback method*), 16
 on_train_batch_end() (*dicee.abstracts.AbstractCallback method*), 17
 on_train_batch_end() (*dicee.abstracts.AbstractTrainer method*), 14
 on_train_batch_end() (*dicee.callbacks.Eval method*), 25
 on_train_batch_end() (*dicee.callbacks.KGESaveCallback method*), 23
 on_train_batch_end() (*dicee.callbacks.LRScheduler method*), 27
 on_train_batch_end() (*dicee.callbacks.PrintCallback method*), 22
 on_train_batch_start() (*dicee.callbacks.Perturb method*), 26
 on_train_epoch_end() (*dicee.abstracts.AbstractCallback method*), 17
 on_train_epoch_end() (*dicee.abstracts.AbstractTrainer method*), 14
 on_train_epoch_end() (*dicee.callbacks.Eval method*), 25
 on_train_epoch_end() (*dicee.callbacks.KGESaveCallback method*), 23

on_train_epoch_end() (dicee.callbacks.PeriodicEvalCallback method), 26
 on_train_epoch_end() (dicee.callbacks.PrintCallback method), 22
 on_train_epoch_end() (dicee.models.base_model.BaseKGELightning method), 85
 on_train_epoch_end() (dicee.models.BaseKGELightning method), 134
 on_train_epoch_end() (dicee.weight_averaging.ASWA method), 205
 on_train_epoch_end() (dicee.weight_averaging.EMA method), 208
 on_train_epoch_end() (dicee.weight_averaging.SWA method), 206
 on_train_epoch_end() (dicee.weight_averaging.SWAG method), 207
 on_train_epoch_end() (dicee.weight_averaging.TWA method), 209
 on_train_epoch_start() (dicee.abstracts.AbstractTrainer method), 13
 on_train_epoch_start() (dicee.weight_averaging.EMA method), 208
 on_train_epoch_start() (dicee.weight_averaging.SWA method), 206
 on_train_epoch_start() (dicee.weight_averaging.SWAG method), 207
 on_train_epoch_start() (dicee.weight_averaging.TWA method), 209
 on_train_start() (dicee.callbacks.LRScheduler method), 27
 OnevsAllDataset (class in dicee.dataset_classes), 35
 OnevsSample (class in dicee.dataset_classes), 36
 optim (dicee.config.Namespace attribute), 28
 optimizer (dicee.trainer.torch_trainer_ddp.NodeTrainer attribute), 202
 optimizer (dicee.trainer.torch_trainer.TorchTrainer attribute), 200
 optimizer_name (dicee.models.base_model.BaseKGE attribute), 90
 optimizer_name (dicee.models.BaseKGE attribute), 139, 142, 148, 152, 158, 172, 175
 ordered_bpe_entities (dicee.dataset_classes.BPE_NegativeSamplingDataset attribute), 32
 ordered_bpe_entities (dicee.knowledge_graph.KG attribute), 70
 ordered_shaped_bpe_tokens (dicee.knowledge_graph.KG attribute), 69

P

p (dicee.config.Namespace attribute), 30
 p (dicee.models.clifford.DeCaL attribute), 98
 p (dicee.models.clifford.Keci attribute), 93
 p (dicee.models.DeCaL attribute), 167
 p (dicee.models.Keci attribute), 164
 P (dicee.weight_averaging.TWA attribute), 209
 padding (dicee.knowledge_graph.KG attribute), 70
 pandas_dataframe_indexer() (in module dicee.read_preprocess_save_load_kg.util), 185
 param_init (dicee.models.base_model.BaseKGE attribute), 90
 param_init (dicee.models.BaseKGE attribute), 139, 142, 148, 153, 159, 172, 175
 parameters() (dicee.abstracts.BaseInteractiveKGE method), 16
 parameters() (dicee.models.ensemble.EnsembleKGE method), 105
 path (dicee.abstracts.AbstractPPECallback attribute), 18
 path (dicee.callbacks.AccumulateEpochLossCallback attribute), 21
 path (dicee.callbacks.Eval attribute), 24
 path (dicee.callbacks.KGESaveCallback attribute), 23
 path (dicee.weight_averaging.ASWA attribute), 204
 path_dataset_folder (dicee.analyse_experiments.Experiment attribute), 20
 path_for_deserialization (dicee.knowledge_graph.KG attribute), 69
 path_for_serialization (dicee.knowledge_graph.KG attribute), 69
 path_single_kg (dicee.config.Namespace attribute), 28
 path_single_kg (dicee.knowledge_graph.KG attribute), 69
 path_to_store_single_run (dicee.config.Namespace attribute), 28
 PeriodicEvalCallback (class in dicee.callbacks), 26
 Perturb (class in dicee.callbacks), 25
 polars_dataframe_indexer() (in module dicee.read_preprocess_save_load_kg.util), 184
 poly_NN() (dicee.models.function_space.LFMMult method), 108
 poly_NN() (dicee.models.LFMMult method), 179
 polynomial() (dicee.models.function_space.LFMMult method), 109
 polynomial() (dicee.models.LFMMult method), 179
 pop() (dicee.models.function_space.LFMMult method), 109
 pop() (dicee.models.LFMMult method), 179
 pos_emb (dicee.models.CoKE attribute), 147
 pos_emb (dicee.models.real.CoKE attribute), 120
 pq (dicee.analyse_experiments.Experiment attribute), 20
 predict() (dicee.KGE method), 213
 predict() (dicee.knowledge_graph_embeddings.KGE method), 72
 predict_data_loader() (dicee.models.base_model.BaseKGELightning method), 87
 predict_data_loader() (dicee.models.BaseKGELightning method), 135
 predict_literals() (dicee.KGE method), 215

predict_literals() (*dicee.knowledge_graph_embeddings.KGE method*), 74
 predict_missing_head_entity() (*dicee.KGE method*), 212
 predict_missing_head_entity() (*dicee.knowledge_graph_embeddings.KGE method*), 71
 predict_missing_relations() (*dicee.KGE method*), 212
 predict_missing_relations() (*dicee.knowledge_graph_embeddings.KGE method*), 72
 predict_missing_tail_entity() (*dicee.KGE method*), 213
 predict_missing_tail_entity() (*dicee.knowledge_graph_embeddings.KGE method*), 72
 predict_topk() (*dicee.KGE method*), 213
 predict_topk() (*dicee.knowledge_graph_embeddings.KGE method*), 72
 preprocess_with_byte_pair_encoding() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 187
 preprocess_with_byte_pair_encoding() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 182
 preprocess_with_byte_pair_encoding_with_padding() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 187
 preprocess_with_byte_pair_encoding_with_padding() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 182
 preprocess_with_pandas() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 187
 preprocess_with_pandas() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 182
 preprocess_with_polars() (*dicee.read_preprocess_save_load_kg.PreprocessKG method*), 187
 preprocess_with_polars() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG method*), 182
 preprocesses_input_args() (*in module dicee.static_preprocess_funcs*), 196
 PreprocessKG (*class in dicee.read_preprocess_save_load_kg*), 187
 PreprocessKG (*class in dicee.read_preprocess_save_load_kg.preprocess*), 182
 PrintCallback (*class in dicee.callbacks*), 21
 process (*dicee.trainer.torch_trainer.TorchTrainer attribute*), 201
 PseudoLabellingCallback (*class in dicee.callbacks*), 23
 Pyke (*class in dicee.models*), 145
 Pyke (*class in dicee.models.real*), 119
 pykeen_model_kwargs (*dicee.config.Namespace attribute*), 30
 PykeenKGE (*class in dicee.models*), 173
 PykeenKGE (*class in dicee.models.pykeen_models*), 114

Q

q (*dicee.config.Namespace attribute*), 30
 q (*dicee.models.clifford.DeCaL attribute*), 98
 q (*dicee.models.clifford.Keci attribute*), 93
 q (*dicee.models.DeCaL attribute*), 167
 q (*dicee.models.Keci attribute*), 164
 qdrant_client (*dicee.scripts.index_serve.NeuralSearcher attribute*), 189
 QMult (*class in dicee.models*), 155
 QMult (*class in dicee.models.quaternion*), 115
 quaternion_mul() (*in module dicee.models*), 151
 quaternion_mul() (*in module dicee.models.static_funcs*), 121
 quaternion_mul_with_unit_norm() (*in module dicee.models*), 155
 quaternion_mul_with_unit_norm() (*in module dicee.models.quaternion*), 115
 quaternion_multiplication_followed_by_inner_product() (*dicee.models.QMult method*), 155
 quaternion_multiplication_followed_by_inner_product() (*dicee.models.quaternion.QMult method*), 116
 quaternion_normalizer() (*dicee.models.QMult static method*), 156
 quaternion_normalizer() (*dicee.models.quaternion.QMult static method*), 116
 queries (*dicee.scripts.index_serve.StringListRequest attribute*), 190
 query_name_to_struct (*dicee.query_generator.QueryGenerator attribute*), 181
 query_name_to_struct (*dicee.QueryGenerator attribute*), 216
 QueryGenerator (*class in dicee*), 215
 QueryGenerator (*class in dicee.query_generator*), 181

R

r (*dicee.models.clifford.DeCaL attribute*), 98
 r (*dicee.models.clifford.Keci attribute*), 93
 r (*dicee.models.DeCaL attribute*), 167
 r (*dicee.models.Keci attribute*), 164
 random_seed (*dicee.config.Namespace attribute*), 29
 range_constraints_per_rel (*dicee.evaluation.Evaluator attribute*), 55
 range_constraints_per_rel (*dicee.evaluation.evaluator.Evaluator attribute*), 45
 range_constraints_per_rel (*dicee.Evaluator attribute*), 219
 range_constraints_per_rel (*dicee.evaluator.Evaluator attribute*), 63
 rank (*dicee.Execute attribute*), 210
 rank (*dicee.executer.Execute attribute*), 66
 ratio (*dicee.callbacks.Perturb attribute*), 26
 raw_test_set (*dicee.knowledge_graph.KG attribute*), 70
 raw_train_set (*dicee.knowledge_graph.KG attribute*), 70

`raw_valid_set` (*dicee.knowledge_graph.KG attribute*), 70
`re` (*dicee.models.clifford.DeCaL attribute*), 98
`re` (*dicee.models.DeCaL attribute*), 167
`re_vocab` (*dicee.evaluation.Evaluator attribute*), 54, 55
`re_vocab` (*dicee.evaluation.evaluator.Evaluator attribute*), 45
`re_vocab` (*dicee.Evaluator attribute*), 218
`re_vocab` (*dicee.evaluator.Evaluator attribute*), 62, 63
`read_from_disk()` (in module *dicee.read_preprocess_save_load_kg.util*), 186
`read_from_triple_store_with_pandas()` (in module *dicee.read_preprocess_save_load_kg.util*), 186
`read_from_triple_store_with_polars()` (in module *dicee.read_preprocess_save_load_kg.util*), 186
`read_only_few` (*dicee.config.Namespace attribute*), 29
`read_only_few` (*dicee.knowledge_graph.KG attribute*), 69
`read_or_load_kg()` (in module *dicee.static_funcs*), 193
`read_with_pandas()` (in module *dicee.read_preprocess_save_load_kg.util*), 186
`read_with_polars()` (in module *dicee.read_preprocess_save_load_kg.util*), 186
`ReadFromDisk` (class in *dicee.read_preprocess_save_load_kg*), 188
`ReadFromDisk` (class in *dicee.read_preprocess_save_load_kg.read_from_disk*), 183
`reducer` (*dicee.scripts.index_serve.StringListRequest attribute*), 190
`reg_lambda` (*dicee.weight_averaging.TWA attribute*), 208
`rel2id` (*dicee.query_generator.QueryGenerator attribute*), 181
`rel2id` (*dicee.QueryGenerator attribute*), 215
`relation_embeddings` (*dicee.models.AConvQ attribute*), 157
`relation_embeddings` (*dicee.models.clifford.DeCaL attribute*), 98
`relation_embeddings` (*dicee.models.ConvQ attribute*), 156
`relation_embeddings` (*dicee.models.DeCaL attribute*), 167
`relation_embeddings` (*dicee.models.DualE attribute*), 180
`relation_embeddings` (*dicee.models.dualE.DualE attribute*), 104
`relation_embeddings` (*dicee.models.FMult attribute*), 176
`relation_embeddings` (*dicee.models.FMult2 attribute*), 178
`relation_embeddings` (*dicee.models.function_space.FMult attribute*), 106
`relation_embeddings` (*dicee.models.function_space.FMult2 attribute*), 107
`relation_embeddings` (*dicee.models.function_space.GFMult attribute*), 106
`relation_embeddings` (*dicee.models.function_space.LFMult attribute*), 108
`relation_embeddings` (*dicee.models.function_space.LFMult1 attribute*), 107
`relation_embeddings` (*dicee.models.GFMult attribute*), 177
`relation_embeddings` (*dicee.models.LFMult attribute*), 178
`relation_embeddings` (*dicee.models.LFMult1 attribute*), 178
`relation_embeddings` (*dicee.models.pykeen_models.PykeenKGE attribute*), 114
`relation_embeddings` (*dicee.models.PykeenKGE attribute*), 173
`relation_embeddings` (*dicee.models.quaternion.AConvQ attribute*), 117
`relation_embeddings` (*dicee.models.quaternion.ConvQ attribute*), 117
`relation_to_idx` (*dicee.knowledge_graph.KG attribute*), 69
`relations_str` (*dicee.knowledge_graph.KG property*), 70
`reload_dataset()` (in module *dicee.dataset_classes*), 39
`report` (*dicee.DICE_Trainer attribute*), 217
`report` (*dicee.evaluation.Evaluator attribute*), 54, 55
`report` (*dicee.evaluation.evaluator.Evaluator attribute*), 45, 46
`report` (*dicee.Evaluator attribute*), 218, 219
`report` (*dicee.evaluator.Evaluator attribute*), 63
`report` (*dicee.Execute attribute*), 210
`report` (*dicee.executer.Execute attribute*), 66
`report` (*dicee.trainer.DICE_Trainer attribute*), 203
`report` (*dicee.trainer.dice_trainer.DICE_Trainer attribute*), 198
`reports` (*dicee.callbacks.Eval attribute*), 24
`reports` (*dicee.callbacks.PeriodicEvalCallback attribute*), 26
`requires_grad_for_interactions` (*dicee.models.CKeci attribute*), 166
`requires_grad_for_interactions` (*dicee.models.clifford.CKeci attribute*), 97
`requires_grad_for_interactions` (*dicee.models.clifford.Keci attribute*), 93
`requires_grad_for_interactions` (*dicee.models.Keci attribute*), 164
`resid_dropout` (*dicee.models.clifford.TransformerSelfAttention attribute*), 97
`resid_dropout` (*dicee.models.transformers.SelfAttention attribute*), 124
`residual_convolution()` (*dicee.models.AConEx method*), 150
`residual_convolution()` (*dicee.models.AConvO method*), 163
`residual_convolution()` (*dicee.models.AConvQ method*), 157
`residual_convolution()` (*dicee.models.complex.AConEx method*), 102
`residual_convolution()` (*dicee.models.complex.ConEx method*), 101
`residual_convolution()` (*dicee.models.ConEx method*), 150
`residual_convolution()` (*dicee.models.ConvO method*), 162

`residual_convolution()` (*dicее.models.ConvQ method*), 157
`residual_convolution()` (*dicее.models.octonion.AConvO method*), 113
`residual_convolution()` (*dicее.models.octonion.ConvO method*), 113
`residual_convolution()` (*dicее.models.quaternion.AConvQ method*), 117
`residual_convolution()` (*dicее.models.quaternion.ConvQ method*), 117
`retrieve_embedding()` (*dicее.scripts.index_serve.NeuralSearcher method*), 189
`retrieve_embeddings()` (*in module dicее.scripts.index_serve*), 190
`return_multi_hop_query_results()` (*dicее.KGE method*), 214
`return_multi_hop_query_results()` (*dicее.knowledge_graph_embeddings.KGE method*), 73
`root()` (*in module dicее.scripts.index_serve*), 190
`roots` (*dicее.models.FMult attribute*), 177
`roots` (*dicее.models.function_space.FMult attribute*), 106
`roots` (*dicее.models.function_space.GFMult attribute*), 106
`roots` (*dicее.models.GFMult attribute*), 177
`runtime` (*dicее.analyse_experiments.Experiment attribute*), 20

S

`sample()` (*dicее.weight_averaging.SWAG method*), 207
`sample_counter` (*dicее.abstracts.AbstractPPECallback attribute*), 18
`sample_entity()` (*dicее.abstracts.BaseInteractiveKGE method*), 15
`sample_relation()` (*dicее.abstracts.BaseInteractiveKGE method*), 15
`sample_triples_ratio` (*dicее.config.Namespace attribute*), 29
`sample_triples_ratio` (*dicее.knowledge_graph.KG attribute*), 69
`sample_weights()` (*dicее.weight_averaging.TWA method*), 209
`sampling_ratio` (*dicее.dataset_classes.LiteralDataset attribute*), 38
`sanity_check_callback_args()` (*in module dicее.sanity_checkers*), 188
`sanity_checking_with_arguments()` (*in module dicее.sanity_checkers*), 188
`save()` (*dicее.abstracts.BaseInteractiveKGE method*), 15
`save()` (*dicее.read_preprocess_save_load_kg.LoadSaveToDisk method*), 188
`save()` (*dicее.read_preprocess_save_load_kg.save_load_disk.LoadSaveToDisk method*), 183
`save_checkpoint()` (*dicее.abstracts.AbstractTrainer static method*), 14
`save_checkpoint_model()` (*in module dicее.static_funcs*), 193
`save_embeddings()` (*in module dicее.static_funcs*), 194
`save_embeddings_as_csv` (*dicее.config.Namespace attribute*), 28
`save_every_n_epochs` (*dicее.config.Namespace attribute*), 30
`save_experiment()` (*dicее.analyse_experiments.Experiment method*), 20
`save_model_at_every_epoch` (*dicее.config.Namespace attribute*), 29
`save_model_every_n_epoch` (*dicее.callbacks.PeriodicEvalCallback attribute*), 26
`save_numpy_ndarray()` (*in module dicее.read_preprocess_save_load_kg.util*), 186
`save_numpy_ndarray()` (*in module dicее.static_funcs*), 193
`save_pickle()` (*in module dicее.read_preprocess_save_load_kg.util*), 187
`save_pickle()` (*in module dicее.static_funcs*), 192
`save_queries()` (*dicее.query_generator.QueryGenerator method*), 182
`save_queries()` (*dicее.QueryGenerator method*), 216
`save_queries_and_answers()` (*dicее.query_generator.QueryGenerator static method*), 182
`save_queries_and_answers()` (*dicее.QueryGenerator static method*), 216
`save_trained_model()` (*dicее.Execute method*), 211
`save_trained_model()` (*dicее.executer.Execute method*), 67
`scalar_batch_NN()` (*dicее.models.function_space.LFMult method*), 108
`scalar_batch_NN()` (*dicее.models.LFMult method*), 179
`scaler` (*dicее.callbacks.Perturb attribute*), 26
`scaler` (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
`scheduler` (*dicее.callbacks.LRScheduler attribute*), 27
`score()` (*dicее.models.clifford.Keci method*), 96
`score()` (*dicее.models.CoKE method*), 147
`score()` (*dicее.models.ComplEx static method*), 151
`score()` (*dicее.models.complex.ComplEx static method*), 103
`score()` (*dicее.models.DistMult method*), 145
`score()` (*dicее.models.Keci method*), 166
`score()` (*dicее.models.octonion.OMult method*), 112
`score()` (*dicее.models.OMult method*), 161
`score()` (*dicее.models.QMult method*), 156
`score()` (*dicее.models.quaternion.QMult method*), 116
`score()` (*dicее.models.real.CoKE method*), 120
`score()` (*dicее.models.real.DistMult method*), 118
`score()` (*dicее.models.real.TransE method*), 118
`score()` (*dicее.models.TransE method*), 145

score_func (*dicee.models.FMulti2* attribute), 177
 score_func (*dicee.models.function_space.FMulti2* attribute), 107
 scoring_technique (*dicee.analyse_experiments.Experiment* attribute), 20
 scoring_technique (*dicee.config.Namespace* attribute), 29
 search() (*dicee.scripts.index_serve.NeuralSearcher* method), 189
 search_embeddings() (in module *dicee.scripts.index_serve*), 190
 search_embeddings_batch() (in module *dicee.scripts.index_serve*), 190
 seed (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 36
 seed (*dicee.dataset_classes.TriplePredictionDataset* attribute), 37
 seed (*dicee.query_generator.QueryGenerator* attribute), 181
 seed (*dicee.QueryGenerator* attribute), 215
 select_model() (in module *dicee.static_funcs*), 193
 selected_optimizer (*dicee.models.base_model.BaseKGE* attribute), 90
 selected_optimizer (*dicee.models.BaseKGE* attribute), 139, 142, 148, 153, 158, 172, 175
 SelfAttention (class in *dicee.models.transformers*), 123
 separator (*dicee.config.Namespace* attribute), 29
 separator (*dicee.knowledge_graph.KG* attribute), 69
 seq_len (*dicee.models.clifford.KeciTransformer* attribute), 96
 seq_len (*dicee.models.KeciTransformer* attribute), 170
 sequential_vocabulary_construction() (*dicee.read_preprocess_save_load_kg.PreprocessKG* method), 187
 sequential_vocabulary_construction() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 182
 serve() (in module *dicee.scripts.index_serve*), 190
 set_global_seed() (*dicee.query_generator.QueryGenerator* method), 181
 set_global_seed() (*dicee.QueryGenerator* method), 216
 set_model_eval_mode() (*dicee.abstracts.BaseInteractiveKGE* method), 15
 set_model_train_mode() (*dicee.abstracts.BaseInteractiveKGE* method), 15
 setup_executor() (*dicee.Execute* method), 211
 setup_executor() (*dicee.executer.Execute* method), 67
 Shallom (class in *dicee.models*), 145
 Shallom (class in *dicee.models.real*), 119
 shallom (*dicee.models.real.Shallom* attribute), 119
 shallom (*dicee.models.Shallom* attribute), 145
 single_hop_query_answering() (*dicee.KGE* method), 214
 single_hop_query_answering() (*dicee.knowledge_graph_embeddings.KGE* method), 73
 snapshot_dir (*dicee.callbacks.LRScheduler* attribute), 27
 snapshot_loss (*dicee.callbacks.LRScheduler* attribute), 27
 sparql_endpoint (*dicee.config.Namespace* attribute), 28
 sparql_endpoint (*dicee.knowledge_graph.KG* attribute), 69
 sq_mean (*dicee.weight_averaging.SWAG* attribute), 207
 start() (*dicee.DICE_Trainer* method), 217
 start() (*dicee.Execute* method), 211
 start() (*dicee.executer.Execute* method), 67
 start() (*dicee.read_preprocess_save_load_kg.PreprocessKG* method), 187
 start() (*dicee.read_preprocess_save_load_kg.preprocess.PreprocessKG* method), 182
 start() (*dicee.read_preprocess_save_load_kg.read_from_disk.ReadFromDisk* method), 183
 start() (*dicee.read_preprocess_save_load_kg.ReadFromDisk* method), 188
 start() (*dicee.trainer.DICE_Trainer* method), 204
 start() (*dicee.trainer.dice_trainer.DICE_Trainer* method), 199
 start_time (*dicee.callbacks.PrintCallback* attribute), 22
 start_time (*dicee.Execute* attribute), 211
 start_time (*dicee.executer.Execute* attribute), 66
 state_dict() (*dicee.models.ensemble.EnsembleKGE* method), 105
 step() (*dicee.models.ADOPT* method), 131
 step() (*dicee.models.adapt.ADOPT* method), 78
 step() (*dicee.models.ensemble.EnsembleKGE* method), 105
 step_count (*dicee.callbacks.LRScheduler* attribute), 27
 storage_path (*dicee.config.Namespace* attribute), 28
 storage_path (*dicee.DICE_Trainer* attribute), 217
 storage_path (*dicee.trainer.DICE_Trainer* attribute), 203
 storage_path (*dicee.trainer.dice_trainer.DICE_Trainer* attribute), 198
 store() (in module *dicee.static_funcs*), 193
 store_ensemble() (*dicee.abstracts.AbstractPPECallback* method), 18
 strategy (*dicee.abstracts.AbstractTrainer* attribute), 13
 StringListRequest (class in *dicee.scripts.index_serve*), 190
 SWA (class in *dicee.weight_averaging*), 205
 swa (*dicee.config.Namespace* attribute), 30
 swa_c_epochs (*dicee.config.Namespace* attribute), 31
 swa_c_epochs (*dicee.weight_averaging.SWA* attribute), 206

swa_c_epochs (*dicее.weight_averaging.SWAG attribute*), 207
 swa_lr (*dicее.weight_averaging.SWA attribute*), 206
 swa_lr (*dicее.weight_averaging.SWAG attribute*), 207
 swa_model (*dicее.weight_averaging.SWA attribute*), 206
 swa_n (*dicее.weight_averaging.SWA attribute*), 206
 swa_start_epoch (*dicее.config.Namespace attribute*), 31
 swa_start_epoch (*dicее.weight_averaging.SWA attribute*), 206
 swa_start_epoch (*dicее.weight_averaging.SWAG attribute*), 207
 SWAG (*class in dicее.weight_averaging*), 206
 swag (*dicее.config.Namespace attribute*), 30

T

T() (*dicее.models.DualE method*), 180
 T() (*dicее.models.dualE.DualE method*), 104
 t_conorm() (*dicее.abstracts.InteractiveQueryDecomposition method*), 16
 t_norm() (*dicее.abstracts.InteractiveQueryDecomposition method*), 16
 target_dim (*dicее.dataset_classes.AllvsAll attribute*), 34
 target_dim (*dicее.dataset_classes.MultiLabelDataset attribute*), 33
 target_dim (*dicее.dataset_classes.OnevsAllDataset attribute*), 36
 target_dim (*dicее.knowledge_graph.KG attribute*), 70
 temperature (*dicее.models.transformers.BytE attribute*), 122
 tensor_t_norm() (*dicее.abstracts.InteractiveQueryDecomposition method*), 16
 TensorParallel (*class in dicее.trainer.model_parallelism*), 200
 test_dataloader() (*dicее.models.base_model.BaseKGELightning method*), 86
 test_dataloader() (*dicее.models.BaseKGELightning method*), 134
 test_epoch_end() (*dicее.models.base_model.BaseKGELightning method*), 86
 test_epoch_end() (*dicее.models.BaseKGELightning method*), 134
 test_h1 (*dicее.analyse_experiments.Experiment attribute*), 20
 test_h3 (*dicее.analyse_experiments.Experiment attribute*), 20
 test_h10 (*dicее.analyse_experiments.Experiment attribute*), 20
 test_mrr (*dicее.analyse_experiments.Experiment attribute*), 20
 test_path (*dicее.query_generator.QueryGenerator attribute*), 181
 test_path (*dicее.QueryGenerator attribute*), 215
 test_set (*dicее.knowledge_graph.KG attribute*), 69, 70
 timeit() (*in module dicее.read_preprocess_save_load_kg.util*), 186
 timeit() (*in module dicее.static_funcs*), 192
 timeit() (*in module dicее.static_preprocess_funcs*), 196
 to() (*dicее.KGE method*), 212
 to() (*dicее.knowledge_graph_embeddings.KGE method*), 71
 to() (*dicее.models.ensemble.EnsembleKGE method*), 105
 to_df() (*dicее.analyse_experiments.Experiment method*), 20
 topk (*dicее.models.transformers.BytE attribute*), 122
 topk (*dicее.scripts.index_serve.NeuralSearcher attribute*), 189
 torch_ordered_shaped_bpe_entities (*dicее.dataset_classes.MultiLabelDataset attribute*), 33
 TorchDDPTrainer (*class in dicее.trainer.torch_trainer_ddp*), 201
 TorchTrainer (*class in dicее.trainer.torch_trainer*), 200
 total_epochs (*dicее.callbacks.LRScheduler attribute*), 27
 total_steps (*dicее.callbacks.LRScheduler attribute*), 27
 train() (*dicее.abstracts.BaseInteractiveTrainKGE method*), 19
 train() (*dicее.trainer.torch_trainer_ddp.NodeTrainer method*), 202
 train_data (*dicее.dataset_classes.AllvsAll attribute*), 34
 train_data (*dicее.dataset_classes.KvsAll attribute*), 34
 train_data (*dicее.dataset_classes.KvsSampleDataset attribute*), 35
 train_data (*dicее.dataset_classes.MultiClassClassificationDataset attribute*), 33
 train_data (*dicее.dataset_classes.OnevsAllDataset attribute*), 35
 train_data (*dicее.dataset_classes.OnevsSample attribute*), 37
 train_dataloader() (*dicее.models.base_model.BaseKGELightning method*), 87
 train_dataloader() (*dicее.models.BaseKGELightning method*), 136
 train_data_loaders (*dicее.trainer.torch_trainer.TorchTrainer attribute*), 201
 train_dataset_loader (*dicее.trainer.torch_trainer_ddp.NodeTrainer attribute*), 202
 train_file_path (*dicее.dataset_classes.LiteralDataset attribute*), 38
 train_h1 (*dicее.analyse_experiments.Experiment attribute*), 20
 train_h3 (*dicее.analyse_experiments.Experiment attribute*), 20
 train_h10 (*dicее.analyse_experiments.Experiment attribute*), 20
 train_indices_target (*dicее.dataset_classes.MultiLabelDataset attribute*), 33
 train_k_vs_all() (*dicее.abstracts.BaseInteractiveTrainKGE method*), 18
 train_literals() (*dicее.abstracts.BaseInteractiveTrainKGE method*), 19

train_mode (*dicee.models.ensemble.EnsembleKGE* attribute), 105
 train_mrr (*dicee.analyse_experiments.Experiment* attribute), 20
 train_path (*dicee.query_generator.QueryGenerator* attribute), 181
 train_path (*dicee.QueryGenerator* attribute), 215
 train_set (*dicee.dataset_classes.BPE_NegativeSamplingDataset* attribute), 32
 train_set (*dicee.dataset_classes.MultiLabelDataset* attribute), 33
 train_set (*dicee.dataset_classes.TriplePredictionDataset* attribute), 37
 train_set (*dicee.knowledge_graph.KG* attribute), 69, 70
 train_set_target (*dicee.knowledge_graph.KG* attribute), 70
 train_target (*dicee.dataset_classes.AllvsAll* attribute), 34
 train_target (*dicee.dataset_classes.KvsAll* attribute), 34
 train_target (*dicee.dataset_classes.KvsSampleDataset* attribute), 35
 train_target_indices (*dicee.knowledge_graph.KG* attribute), 70
 train_triples (*dicee.dataset_classes.FixedNegSampleDataset* attribute), 36
 train_triples() (*dicee.abstracts.BaseInteractiveTrainKGE* method), 18
 trained_model (*dicee.Execute* attribute), 210
 trained_model (*dicee.executer.Execute* attribute), 66
 trainer (*dicee.config.Namespace* attribute), 29
 trainer (*dicee.DICE_Trainer* attribute), 217
 trainer (*dicee.Execute* attribute), 210
 trainer (*dicee.executer.Execute* attribute), 66
 trainer (*dicee.trainer.DICE_Trainer* attribute), 203
 trainer (*dicee.trainer.dice_trainer.DICE_Trainer* attribute), 198
 trainer (*dicee.trainer.torch_trainer_ddp.NodeTrainer* attribute), 202
 training_step (*dicee.trainer.torch_trainer.TorchTrainer* attribute), 201
 training_step() (*dicee.models.base_model.BaseKGELightning* method), 84
 training_step() (*dicee.models.BaseKGELightning* method), 133
 training_step() (*dicee.models.transformers.BytE* method), 122
 training_step_outputs (*dicee.models.base_model.BaseKGELightning* attribute), 84
 training_step_outputs (*dicee.models.BaseKGELightning* attribute), 133
 training_technique (*dicee.knowledge_graph.KG* attribute), 69
 TransE (*class in dicee.models*), 145
 TransE (*class in dicee.models.real*), 118
 transformer (*dicee.models.clifford.KeciTransformer* attribute), 96
 transformer (*dicee.models.KeciTransformer* attribute), 170
 transformer (*dicee.models.transformers.BytE* attribute), 122
 transformer (*dicee.models.transformers.GPT* attribute), 127
 TransformerBlock (*class in dicee.models.clifford*), 96
 TransformerMLP (*class in dicee.models.clifford*), 97
 TransformerSelfAttention (*class in dicee.models.clifford*), 96
 trapezoid() (*dicee.models.FMult2* method), 178
 trapezoid() (*dicee.models.function_space.FMult2* method), 107
 tri_score() (*dicee.models.function_space.LFMult* method), 108
 tri_score() (*dicee.models.function_space.LFMult1* method), 107
 tri_score() (*dicee.models.LFMult* method), 179
 tri_score() (*dicee.models.LFMult1* method), 178
 triple_score() (*dicee.KGE* method), 213
 triple_score() (*dicee.knowledge_graph_embeddings.KGE* method), 73
 TriplePredictionDataset (*class in dicee.dataset_classes*), 37
 tuple2list() (*dicee.query_generator.QueryGenerator* method), 181
 tuple2list() (*dicee.QueryGenerator* method), 216
 TWA (*class in dicee.weight_averaging*), 208
 twa (*dicee.config.Namespace* attribute), 30
 twa_c_epochs (*dicee.weight_averaging.TWA* attribute), 209
 twa_model (*dicee.weight_averaging.TWA* attribute), 209
 twa_start_epoch (*dicee.weight_averaging.TWA* attribute), 208

U

unlabelled_size (*dicee.callbacks.PseudoLabellingCallback* attribute), 24
 unmap() (*dicee.query_generator.QueryGenerator* method), 181
 unmap() (*dicee.QueryGenerator* method), 216
 unmap_query() (*dicee.query_generator.QueryGenerator* method), 181
 unmap_query() (*dicee.QueryGenerator* method), 216
 update_hits() (*in module dicee.evaluation*), 61
 update_hits() (*in module dicee.evaluation.utils*), 53

V

`val_aswa` (*dicee.weight_averaging.ASWA attribute*), 205
`val_data_loader()` (*dicee.models.base_model.BaseKGELighting method*), 86
`val_data_loader()` (*dicee.models.BaseKGELighting method*), 135
`val_h1` (*dicee.analyse_experiments.Experiment attribute*), 20
`val_h3` (*dicee.analyse_experiments.Experiment attribute*), 20
`val_h10` (*dicee.analyse_experiments.Experiment attribute*), 20
`val_mrr` (*dicee.analyse_experiments.Experiment attribute*), 20
`val_path` (*dicee.query_generator.QueryGenerator attribute*), 181
`val_path` (*dicee.QueryGenerator attribute*), 215
`VALID_SCORING_TECHNIQUES` (*in module dicee.evaluation.evaluator*), 45
`valid_set` (*dicee.knowledge_graph.KG attribute*), 69, 70
`validate_knowledge_graph()` (*in module dicee.sanity_checkers*), 188
`var_clamp` (*dicee.weight_averaging.SWAG attribute*), 207
`vocab_preparation()` (*dicee.evaluation.Evaluator method*), 55
`vocab_preparation()` (*dicee.evaluation.evaluator.Evaluator method*), 46
`vocab_preparation()` (*dicee.Evaluator method*), 219
`vocab_preparation()` (*dicee.evaluator.Evaluator method*), 63
`vocab_size` (*dicee.models.CoKEConfig attribute*), 146
`vocab_size` (*dicee.models.real.CoKEConfig attribute*), 119, 120
`vocab_size` (*dicee.models.transformers.GPTConfig attribute*), 126
`vocab_to_parquet()` (*in module dicee.static_funcs*), 194
`vtp_score()` (*dicee.models.function_space.LFMult method*), 108
`vtp_score()` (*dicee.models.function_space.LFMultI method*), 108
`vtp_score()` (*dicee.models.LFMult method*), 179
`vtp_score()` (*dicee.models.LFMultI method*), 178

W

`warmup_steps` (*dicee.callbacks.LRScheduler attribute*), 27
`weight` (*dicee.models.transformers.LayerNorm attribute*), 123
`weight_decay` (*dicee.config.Namespace attribute*), 29
`weight_decay` (*dicee.models.base_model.BaseKGE attribute*), 90
`weight_decay` (*dicee.models.BaseKGE attribute*), 139, 142, 148, 152, 158, 172, 175
`weight_samples` (*dicee.weight_averaging.TWA attribute*), 209
`weights` (*dicee.models.FMult attribute*), 177
`weights` (*dicee.models.function_space.FMult attribute*), 106
`weights` (*dicee.models.function_space.GFMult attribute*), 106
`weights` (*dicee.models.GFMult attribute*), 177
`world_size` (*dicee.Execute attribute*), 210
`world_size` (*dicee.executer.Execute attribute*), 66
`write_csv_from_model_parallel()` (*in module dicee.static_funcs*), 194
`write_links()` (*dicee.query_generator.QueryGenerator method*), 181
`write_links()` (*dicee.QueryGenerator method*), 216
`write_report()` (*dicee.Execute method*), 211
`write_report()` (*dicee.executer.Execute method*), 67

X

`x_values` (*dicee.models.function_space.LFMult attribute*), 108
`x_values` (*dicee.models.LFMult attribute*), 178